

The Systems Engineer's Path

A Five-Volume Book — from senior engineer to systems architect: distributed systems, database internals, and platform design from first principles

A training book · with worked examples, build-alongs, exercises, and full solutions

Preface

This is the second volume. The first, *The Data Engineer's Path*, took a computer-science graduate and trained them into a **senior data engineer** — someone who can hold a company's data platform and reason about the systems they operate. This volume takes that engineer the next step: to a **data architect and staff-level systems engineer** who understands those systems *from the inside*, can *build* them, not just use them, and can *design platforms* and defend the trade-offs at scale. Where Volume 1 taught "the shuffle is expensive," A Five-Volume Book shows you the exchange-operator machinery underneath it, has you build a query engine that shuffles, and teaches you to architect the platform it runs on.

It assumes Volume 1 (or an equivalent senior data engineer) and cross-references it generously, so it is usable on its own. It is denser and goes deeper — but it is written for someone senior, so it respects your time: no re-teaching, heavy signposting, and every abstraction motivated by the problem that forced it into existence.

Four threads run through the whole volume.

- **Think from first principles.** Derive a system's behavior from its underlying model — the failure model, the storage model, the consistency model — not from a tool's documentation. You will see a *first-principles* callout that traces each mechanism to the constraint that produced it.
- **Reason about cost, complexity, and scale.** Now with formal tools — amplification and the RUM conjecture, queueing theory and Little's law, the Universal Scalability Law — and back-of-the-envelope estimation at architect scale. A *cost & scale* callout runs throughout.
- **Design under trade-offs.** Every architecture is a deliberate set of trade-offs against constraints and non-functional requirements; naming and defending them *is* the job. A *design trade-off* callout makes each one explicit, and recurring "design review" exercises have you critique and defend real architectures.
- **Communicate the design.** Architecture decision records, diagrams, and the ability to make a skeptical room trust a design — the third thread of Volume 1, elevated. (The leadership and technical-strategy half of the architect role is the subject of Volume 5; this volume stays technical.)

Correctness is non-negotiable, as in Volume 1. Every citation names a real paper, specification, book with free-legal access, or documented system, with the correct author and venue; numbers are dated because hardware and pricing age; fast-moving claims carry their date; and where the field is genuinely unsettled the text says so. The sources are the canon of systems: the distributed-systems and database papers, *Designing Data-Intensive Applications*, Petrov's *Database Internals*, the CMU and MIT graduate courses, and the official specs.

How to use this book. Read it front to back and do the work. Each chapter motivates its topic, develops it with worked examples and traces, and closes with a *key ideas* box, *exercises* with full solutions, and pointers into the primary literature. Several parts end with a **build-along** — a mini Raft module, an LSM store, a query engine, a Kubernetes operator — because at this level, reading is not enough: you understand a system when you have built a small one. The build-alongs are in Python for approachability, with pointers to how production engines differ in Rust and C++.

Contents

Vol 1 • Systems Programming & Performance • Part I — The Machine & the Cost Model

Chapter 1 — The Real Machine vs. the RAM Model.....	9
Chapter 2 — The CPU Up Close	19
Chapter 3 — The Memory Hierarchy & Caches	28
Chapter 4 — Virtual Memory & Address Translation	38
Chapter 5 — Mechanical Sympathy & Data-Oriented Design	47
Chapter 6 — The Cost Model, Applied (Capstone Lab)	55

Vol 1 • Systems Programming & Performance • Part II — C & Manual Memory

Chapter 7 — Pointers & Memory Layout.....	64
Chapter 8 — The Stack & the Heap.....	75
Chapter 9 — Manual Allocation: malloc/free and What Goes Wrong	84
Chapter 10 — Undefined Behavior & the Optimizing Compiler	94
Chapter 11 — Arrays, Strings & Buffer Safety	104
Chapter 12 — Custom Allocators: Arenas, Pools & Free Lists	113

Vol 1 • Systems Programming & Performance • Part III — The Operating System

Chapter 13 — The System-Call Boundary.....	124
Chapter 14 — Processes: fork, exec, and wait.....	133
Chapter 15 — Threads: Shared Memory and Its Hazards	141
Chapter 16 — CPU Scheduling.....	150
Chapter 17 — Virtual Memory: Demand Paging and Page Replacement	159
Chapter 18 — Signals: Asynchronous Control Flow	168
Chapter 19 — Inter-Process Communication: Pipes, Shared Memory, and Sockets	177

Vol 1 • Systems Programming & Performance • Part IV — Concurrency & Parallelism

Chapter 20 — Mutual Exclusion: Locks and the Critical Section.....	188
Chapter 21 — Coordinating Threads: Condition Variables and Semaphores	197
Chapter 22 — Atomics and Compare-and-Swap: Building Blocks Below the Lock.....	207
Chapter 23 — Lock-Free Data Structures	217
Chapter 24 — The Memory Model: memory_order and Happens-Before.....	227
Chapter 25 — Thread Pools and Task Parallelism	237
Chapter 26 — Scalability: Amdahl's Law, Contention, and False Sharing.....	247

Vol 1 • Systems Programming & Performance • Part V — Rust for Systems

Chapter 27 — Ownership and Moves	257
Chapter 28 — Borrowing and References	267
Chapter 29 — Lifetimes	277

Chapter 30 — Traits and Generics.....	287
Chapter 31 — Error Handling.....	297
Chapter 32 — Fearless Concurrency	306

Vol 1 · Systems Programming & Performance · Part VI — Performance Engineering

Chapter 33 — Measure First: Profiling and Benchmarking.....	317
Chapter 34 — Cache-Aware Design: Programming for Locality.....	327
Chapter 35 — Data-Oriented Design: Laying Out Data for the Cache	336
Chapter 36 — SIMD and Vectorization.....	346

Vol 1 · Systems Programming & Performance · Part VII — I/O & Networking, Low-Level

Vol 2 · Database Internals (Build a Database) · Part VIII — Storage & the Disk

Vol 2 · Database Internals (Build a Database) · Part IX — Storage Engines

Vol 2 · Database Internals (Build a Database) · Part X — Durability & Recovery

Vol 2 · Database Internals (Build a Database) · Part XI — Transactions & Concurrency Control

Vol 2 · Database Internals (Build a Database) · Part XII — Query Processing & Execution

Vol 2 · Database Internals (Build a Database) · Part XIII — Query Optimization

Vol 2 · Database Internals (Build a Database) · Part XIV — Indexing & Access Methods

Vol 2 · Database Internals (Build a Database) · Part XV — [BUILD] A Working Database

Vol 2 · Database Internals (Build a Database) · Part XVI — Toward Distributed Databases

Vol 3 · Distributed Systems, Deeply · Part XVII — Models & Foundations

Vol 3 · Distributed Systems, Deeply · Part XVIII — Replication & Consistency

Vol 3 · Distributed Systems, Deeply · Part XIX — Consensus

Vol 3 · Distributed Systems, Deeply · Part XX — Coordination & Transactions

Vol 3 · Distributed Systems, Deeply · Part XXI — Replicated Data & Convergence

Vol 3 · Distributed Systems, Deeply · Part XXII — Partitioning & Scale

Vol 3 · Distributed Systems, Deeply · Part XXIII — Distributed Algorithms & Patterns

Vol 3 · Distributed Systems, Deeply · Part XXIV — Formal Methods

Vol 3 · Distributed Systems, Deeply · Part XXV — Building & Operating Distributed Systems

Vol 4 · Security & Privacy Engineering · Part XXVI — Security Foundations & Threat Modeling

Vol 4 · Security & Privacy Engineering · Part XXVII — Applied Cryptography

Vol 4 · Security & Privacy Engineering · Part XXVIII — Authentication & Authorization

Vol 4 · Security & Privacy Engineering · Part XXIX — Software & Systems Security

Vol 4 · Security & Privacy Engineering · Part XXX — Data Security

Vol 4 · Security & Privacy Engineering · Part XXXI — Privacy Engineering

Vol 4 · Security & Privacy Engineering · Part XXXII — Secure Data Pipelines & Platforms

Vol 4 · Security & Privacy Engineering · Part XXXIII — Detection, Response & Operations

Vol 5 · The Staff/Principal Engineer & Tech Lead · Part XXXIV — The Staff+ Role & Scope

Vol 5 · The Staff/Principal Engineer & Tech Lead · Part XXXV — Technical Leadership & Strategy

Vol 5 · The Staff/Principal Engineer & Tech Lead · Part XXXVI — Architecture & Design at Scale

Vol 5 · The Staff/Principal Engineer & Tech Lead · Part XXXVII — Influence Without Authority

Vol 5 · The Staff/Principal Engineer & Tech Lead · Part XXXVIII — Communication & Writing

Vol 5 · The Staff/Principal Engineer & Tech Lead · Part XXXIX — Mentoring & Growing Engineers

Vol 5 · The Staff/Principal Engineer & Tech Lead · Part XXXX — Execution & Judgment

Vol 5 · The Staff/Principal Engineer & Tech Lead · Part XXXXI — Career & the Organization

Back Matter

Solutions to Exercises.....	391
-----------------------------	-----

VOL 1 · SYSTEMS PROGRAMMING & PERFORMANCE · PART I

The Machine & the Cost Model

What a program actually runs on: CPU, memory hierarchy, caches, and mechanical sympathy — the real cost model that the RAM abstraction hides.

Chapter 1 — The Real Machine vs. the RAM Model

Before you start: nothing — this is the beginning of Volume 1. · **Pairs with:** Chapter 2 (the CPU up close), Chapter 3 (the memory hierarchy and caches). Together these three install the cost model the rest of the volume optimizes against.

BEFORE YOU READ ON — PREDICT FIRST

No prior chapter to recall, so instead make three *guesses* and write them down; the point is to have a number in your head to be surprised by. (1) A modern CPU core can do roughly how many simple operations — additions, comparisons — per second: about a million, a billion, or a trillion? (2) Reading a value that happens to sit in the CPU's fastest cache versus reading one that must come from main memory (DRAM): same speed, about 10× apart, or about 100× apart? (3) When you write $x = a + b$ in C, is "fetch a from memory" the same cost as "add"? Answers resolve over the course of this chapter.

Every algorithms course you have ever taken quietly lied to you, and the lie was a good one. It said: assume the machine can read any piece of memory, do one arithmetic operation, and write the result, each in one fixed unit of time, the same unit regardless of *which* piece of memory or *which* operation. On that assumption you learned to count operations, and counting operations told you whether an algorithm scaled. That assumption is the **RAM model**, and it is the foundation of asymptotic analysis. It is also false about the machine on your desk in ways that span six orders of magnitude — and the whole of performance engineering, the whole reason this volume exists, lives in the gap between the model and the metal.

This chapter is about that gap. We will state precisely what the RAM model assumes, see exactly where a real processor and a real memory system diverge from it, and meet the single most important divergence — the **memory wall**, the enormous and growing distance between how fast a CPU can compute and how long it takes to fetch a byte from DRAM. The goal is not to make you distrust big-O; big-O is right about what it claims. The goal is to make you fluent in the second cost model that big-O deliberately hides — the one denominated in *where the data is*, not *how many operations you do* — because at the level this book operates, that second model is where the performance is won and lost.

1.1 What the RAM model assumes

The **Random Access Machine** (RAM) model is the idealized computer that theoretical analysis runs on. It has a processor that executes instructions one at a time, in program order, and a memory of cells that the processor can read or write by address. Its defining assumption is **uniform cost**: every basic operation — an arithmetic op, a comparison, a load from memory, a store to memory — takes the same constant amount of time, usually idealized as "one unit," and crucially that unit does not depend on which address you touch. Reading cell 5 costs the same as reading cell 5,000,000,000. There is no notion of "nearby" or "far away" memory; access is, as the name says, uniformly *random-access*.

This is a spectacularly useful abstraction, and you should not give it up. Because every operation costs one unit, the running time of an algorithm is just the *number of operations* it performs, which you can count by reading the code. That is what lets you say a linear scan is $O(n)$, a comparison sort is $O(n \log n)$, and a nested-loop join is $O(n^2)$, and to know — before running anything — that the last one will eventually outrun any machine you can buy. The RAM model throws away every detail of real hardware precisely so that *growth rate* survives, and growth rate is the thing that no amount of faster hardware can fix. Nothing in this chapter contradicts that. An $O(n^2)$ algorithm is a disaster on any real machine too.

What the model buys with that simplification is a blind spot, and the blind spot is the subject of this volume. By declaring every operation equal, the RAM model asserts three things that are simply untrue of the computer in front of you:

- **All memory accesses cost the same.** They do not. On real hardware the cost of a load depends enormously on *where the data currently sits* — in a register, in a small fast cache, in a larger slower cache, in DRAM, on an SSD, or across a network — and those costs differ by factors of ten to factors of ten million.
- **One instruction takes one unit of time.** It does not. A real CPU overlaps instructions, executes them out of order, guesses the outcome of branches, and can retire several instructions in a single clock tick or stall for hundreds of ticks on a single one — the subject of [Chapter 2](#).
- **Operations are independent and equal.** They are not. A multiply may cost more than an add; a division much more; and an instruction that depends on the result of the previous one cannot start until that result is ready, so *dependencies* between operations, not just their count, govern speed.

Each of these is a place where the abstraction leaks. This chapter takes the first one — the uniform-cost-memory assumption — because it is the largest leak by far, and because understanding it reorganizes how you think about all fast code. [Chapter 2](#) takes the second (the CPU), and [Chapter 3](#) takes the mechanism, caches, that turns the first assumption from "false" into "false but predictable, and therefore exploitable."

QUICK CHECK

State the RAM model's uniform-cost assumption in one sentence, and name the real-hardware fact that most directly contradicts it. (Answer below the next section — try it from memory first.)

1.2 The memory wall

Here is the historical fact that makes the uniform-cost assumption not just wrong but *catastrophically* wrong on a modern machine. For several decades, the speed at which a CPU can execute instructions improved much faster, year over year, than the speed at which main memory (DRAM) can deliver a requested byte. Processor throughput raced ahead; DRAM *latency* — the time from "I want the byte at this address" to "here is the byte" — improved only slowly. The two curves diverged, and the widening gap between them was named the **memory wall** by Wulf and McKee in 1995, whose short paper

"Hitting the Memory Wall" pointed out the then-uncomfortable consequence: as the gap grows, the average time to touch memory comes to dominate program performance, no matter how fast the processor's arithmetic is (Wulf & McKee, *ACM SIGARCH Computer Architecture News*, 1995).

The consequence for us is stark and it is worth stating as baldly as possible. On a modern processor, the time to fetch a single value from DRAM is on the order of a *hundred nanoseconds* — and in that same hundred nanoseconds, a single core can execute on the order of hundreds of instructions. So a load that misses every cache and goes to main memory is not "one unit" of work in any meaningful sense; measured against the arithmetic the core could otherwise be doing, it is hundreds of units. A program that spends its time chasing pointers through DRAM can leave the processor's arithmetic units idle more than 90% of the time, waiting. The CPU is not slow. It is *waiting for memory*. That single sentence explains an enormous fraction of the real-world performance problems this volume teaches you to see.

The hardware's answer to the memory wall is not faster DRAM — DRAM latency is hard physics and improves slowly. The answer is the **memory hierarchy**: small, fast, expensive memories (caches) placed close to the processor, holding copies of the data the program is most likely to touch next, backed by progressively larger, slower, cheaper memories behind them. The hierarchy does not repeal the memory wall; it *hides* it, but only for programs whose access patterns let the caches do their job. Learning to write those programs — to earn the cache's help rather than fight it — is a large part of what "systems performance" means, and it is the whole subject of [Chapter 3](#). For now the load-bearing idea is only this: *the cost of an operation is dominated by where its data lives, and the gap between "close" and "far" is the central fact of modern performance.*

(Answer to the quick check: the RAM model assumes every basic operation, including any memory access, costs one fixed unit of time independent of the address. The fact that most directly contradicts it is the memory hierarchy — accessing data in a nearby cache is one to two orders of magnitude faster than reaching main memory, and further tiers are slower still.)

1.3 Latency, as orders of magnitude only

To reason about the real cost model you need anchors: a rough sense of how far apart the tiers of the hierarchy are. But here we must be careful, because this is exactly the kind of place where a confident, precise-looking number is worse than useless. **The absolute latencies below are approximate, they are machine-dependent, and they drift as hardware changes.** What is stable — what you should actually memorize — is the *shape*: the ordering of the tiers and the rough number of orders of magnitude between them. Treat every specific figure as "as of the mid-2020s, on a typical server-class x86-64 core, give or take a large factor."

With that caveat stated as loudly as it can be, here is the ordering. It is the single most important picture in this volume.

Where the data is	Rough access latency (order of magnitude)	Relative to L1
CPU register	< 1 ns (available essentially at once, within the pipeline)	—
L1 cache	~1 ns (a few clock cycles)	1×
L2 cache	a few ns (~10 cycles)	~10×
L3 cache	~10 ns (tens of cycles)	~10-50×
Main memory (DRAM)	~100 ns (hundreds of cycles)	~100×
Local SSD (NVMe)	tens of μ s	~10 ⁴ -10 ⁵ ×
Same-datacenter network round trip	hundreds of μ s to ~1 ms	~10 ⁵ -10 ⁶ ×

The register-through-DRAM rows follow the memory-hierarchy access times given for a specific Intel Core i7 in Bryant and O'Hallaron's *Computer Systems: A Programmer's Perspective* (CS:APP, Chapter 6), which are themselves presented there as approximate and processor-specific; the SSD and network rows are conventional community orders of magnitude of the kind popularized by the widely circulated "latency numbers every programmer should know" list (attributed to Jeff Dean, and later made time-adjustable by Colin Scott). **Use them as ratios, never as guarantees.** If your program's performance depends on whether L2 is 8 cycles or 14 on your particular chip, you are measuring, not reading a book — and measuring is exactly what [Chapter 3](#) and the performance-engineering chapters teach.

Read the table for its shape and three facts fall out. First, the jumps are *multiplicative*, not additive: each tier is roughly 10× to 100× slower than the one above it, so the distance from register to network round trip is not "a lot" but a factor of around a million. Second, there is a cliff between on-chip (register, L1/L2/L3 — nanoseconds) and off-chip (DRAM and below — DRAM at ~100 ns, then a huge further step to SSD at microseconds), and the memory wall is the height of that cliff. Third — and this is the one that reorganizes how you write code — the difference between a cache hit and a cache miss to DRAM is the same ~100× that the difference between a good and a bad algorithm often is. You can lose two orders of magnitude of performance without changing your algorithm's big-O at all, purely by where your data lands.

THE SAME IDEA THREE WAYS: RAM MODEL VS. REAL MACHINE

1. In words. The RAM model charges one unit for every memory access. The real machine charges wildly different amounts depending on which tier of the hierarchy holds the data — from under a nanosecond in a register to ~ 100 ns in DRAM and far more below that.

2. As a picture. A pyramid: registers at the tiny fast apex, then L1, L2, L3, then DRAM, then SSD, then network at the wide slow base. Size grows and speed falls as you descend.




Figure 1.1: The memory hierarchy. Latency rises multiplicatively as you descend; the RAM model pretends the whole pyramid is a single flat tier.

3. As a trace. Consider summing an array of 8 million longs (64 MB), one value at a time. The RAM model says: 8×10^6 loads + 8×10^6 adds = 1.6×10^7 "units," done. The real machine says: the array does not fit in cache, so the cost is dominated by streaming 64 MB up from DRAM; the adds are essentially free, hidden under the wait for memory. Same operation count, completely different account of where the time goes — and only the second account predicts what you will actually measure.

1.4 Why "one operation = one unit of time" is a useful lie

It would be easy to draw the wrong conclusion from all this — to decide the RAM model is broken and should be discarded. That is exactly backwards, and getting the relationship right is the whole point of this chapter. The RAM model is a *useful lie*, in the precise sense that it is false in its literal claims but true and indispensable for the job it is meant to do. Its job is to predict how cost *scales*, and for that job the constant-factor differences between memory tiers are noise: whether a memory access costs 1 ns or 100 ns, an $O(n^2)$ algorithm still loses to an $O(n \log n)$ one once n is large enough. The uniform-cost lie is what makes asymptotic analysis *possible*, and asymptotic analysis is non-negotiable. You must know it and you must trust it, for what it is for.

The mistake is not using the RAM model. The mistake is using it for a question it was never built to answer: *how fast will this actually run?* The RAM model has nothing to say about constant factors, and on modern hardware the constant factors span a

hundredfold, all of it hidden inside the word "one" in "one unit of time." Two algorithms with identical $O(n)$ complexity can differ by $50\times$ in wall-clock time because one walks memory in an order the caches love and the other scatters its accesses across DRAM — and the RAM model, by construction, cannot see the difference. So the discipline is to hold two cost models at once and know which question each one answers:

Question	Right model	Denominated in
Will this scale? Which algorithm as $n \rightarrow \infty$?	RAM / big-O	operation count
How fast will it actually run at this size on this machine?	the real cost model	where the data lives; bytes moved across tiers

This is the through-line of the entire volume. Every technique that is coming — cache-friendly data layout, avoiding pointer chasing, keeping the CPU's pipeline fed, minimizing system calls, batching I/O — is a way of respecting the second model *after* you have already satisfied the first. You pick a good algorithm with big-O; then you make it fast with mechanical sympathy. Neither replaces the other. An engineer who knows only big-O writes asymptotically optimal code that runs ten times slower than it should; an engineer who knows only micro-optimization polishes an $O(n^2)$ disaster. You need both, in that order.

TRAP: TRUSTING THE RAM MODEL FOR PERFORMANCE

The classic error this chapter exists to prevent: reasoning that because two implementations have the same big-O, they will perform about the same — and therefore choosing between them on style or convenience rather than measuring. "They're both $O(n)$, so it doesn't matter" is true about *scaling* and can be false about *speed* by one to two orders of magnitude, because big-O deliberately discards the very constant factor — where the data lives — that dominates real time. A linked list and an array can both be $O(n)$ to traverse; the array is often many times faster because it streams contiguous memory the cache can prefetch, while the list chases pointers into random DRAM addresses (we will measure a version of this in [Chapter 3](#)). The fix is not to distrust big-O — it is right about scaling — but to *stop asking big-O a question it cannot answer*. When the question is "how fast," the answer comes from the memory hierarchy and, ultimately, from measurement, never from the operation count alone.

SAY IT LIKE A SYSTEMS ENGINEER

Try it first. A colleague says, "Both loops are $O(n)$, so they'll perform the same." Cover the paragraph below and articulate, in two or three sentences, why that reasoning is incomplete — using the words *cache*, *constant factor*, and *memory hierarchy*.

Model answer. "They have the same *asymptotic* cost, which tells us they scale the same way — but big-O throws away the constant factor, and on real hardware that constant is dominated by where the data lives in the memory hierarchy. If one loop streams contiguous memory and the other chases pointers into random DRAM addresses, they can differ by $10\times$ to $100\times$ in wall-clock time despite identical big-O, because a cache hit and a DRAM miss are about two orders of magnitude apart. So same complexity, potentially very different speed — we should measure before assuming they tie." This answer signals seniority precisely because it separates the two cost models and names which one governs the question being asked.

1.5 Mechanical sympathy: the thesis of this volume

There is a phrase for the discipline this chapter is pointing at, borrowed from motor racing and popularized in software by Martin Thompson: **mechanical sympathy**. A driver with mechanical sympathy understands enough about how the car works to get the most out of it without breaking it; a programmer with mechanical sympathy understands enough about how the machine works to write code that runs *with* the grain of the hardware rather than against it. It does not mean writing assembly, and it does not mean premature micro-optimization. It means knowing that the machine has a memory hierarchy, a pipelined out-of-order CPU, a cost for crossing into the operating system, and a real network underneath every abstraction — and letting that knowledge shape the code you write by default.

That is the promise of this volume, and this chapter is its foundation. We have established the one idea everything else builds on: *the RAM model's uniform cost is a useful lie, and real performance comes from the cost model it hides — the one where the price of an operation is set by where its data lives*. From here the volume descends into the machine to make that concrete. [Chapter 2](#) opens up the CPU itself — pipelining, out-of-order execution, branch prediction — to explain the second broken assumption, that one instruction takes one unit of time. [Chapter 3](#) opens up the memory hierarchy — cache lines, locality, the three kinds of miss — to turn "memory access cost varies" from a warning into a set of concrete, exploitable rules. By the end of Part 1 you will look at a loop and see not an operation count but a pattern of movement through the hierarchy, and that shift in vision is the beginning of writing fast code.

CAN YOU... WITHOUT LOOKING?

- State the RAM model's **uniform-cost assumption** and name the three real-hardware facts it gets wrong (memory is not uniform-cost; instructions do not take one unit each; operations are not independent and equal).
- Explain the **memory wall**: CPU throughput outgrew DRAM latency for decades, so average memory-access time comes to dominate performance (Wulf & McKee, 1995).
- Recite the memory hierarchy **in order** — register < L1 < L2 < L3 < DRAM < SSD < network — and say *why the numbers are orders of magnitude only* (machine- dependent, drift with hardware).
- Say why "one operation = one unit of time" is a **useful lie**: right about scaling, silent about the $\sim 100\times$ constant factor that decides real speed.
- Hold both cost models at once and state **which question each answers** (big-O: scaling; the real cost model: actual speed).
- Define **mechanical sympathy** and say why it comes *after* choosing a good algorithm, not instead of it.

RETRIEVAL — CLOSED BOOK

Cover everything above. Answer from memory; then check yourself against the section noted. Rate your confidence (low/med/high) before you look — an overconfident miss is your most valuable signal.

1. **R1.** What single assumption defines the RAM model, and what does adopting it buy you? (§1.1)
2. **R2.** What is the memory wall, who named it, and what is its consequence for program performance? (§1.2)
3. **R3.** Put these in order of increasing latency and give the rough order-of- magnitude gaps: DRAM, L1 cache, register, network round trip, L3 cache, SSD. Why must you present them as orders of magnitude? (§1.3)
4. **R4.** Two implementations have identical $O(n)$ complexity. Can one be $50\times$ slower? Explain in terms of what big-O discards. (§1.4)
5. **R5.** Which cost model answers "will it scale?" and which answers "how fast will it run?" — and in what units is each denominated? (§1.4)

FURTHER READING

- **Bryant & O'Hallaron, *Computer Systems: A Programmer's Perspective (CS:APP)***, Chapter 6 ("The Memory Hierarchy"). The canonical treatment; its Core i7 access- time figures are the source of this chapter's register-through-DRAM orders of magnitude, and it is careful to present them as processor-specific. The free CMU 15-213 course materials cover the same ground.
- **Wulf & McKee, "Hitting the Memory Wall: Implications of the Obvious"** (*ACM SIGARCH Computer Architecture News*, 1995). The two-page paper that named the memory wall; short, readable, and prescient.
- **Drepper, "What Every Programmer Should Know About Memory"** (2007). A long, free, seminal write-up on caches and DRAM. Dated in its specific numbers — read it for the mechanisms, not the figures — but unmatched in depth on why the hierarchy exists.
- **"Latency Numbers Every Programmer Should Know."** The community list attributed to Jeff Dean, and Colin Scott's interactive, year-adjustable version. Use strictly as orders of magnitude; the absolute numbers are dated hardware and are meant to be read as ratios.

Exercises

- 1.1** WARM-UP State the RAM model's uniform-cost assumption in one sentence. Then list the three things it asserts about real hardware that this chapter says are false.
- 1.2** WARM-UP Put these seven locations in order from fastest to slowest access, and give the rough order-of- magnitude latency for each: L2 cache, DRAM, register, network round trip (same datacenter), L1 cache, local NVMe SSD, L3 cache.
- 1.3** WARM-UP In two sentences, explain the memory wall and why it makes the RAM model's uniform-cost assumption worse over time rather than better.

1.4 **CORE** A DRAM access costs roughly 100 ns and a single core executes on the order of a few billion simple operations per second. (a) Roughly how many simple operations could the core have executed in the time of one DRAM access? (b) In one or two sentences, explain what this ratio implies for a program that spends most of its time following pointers into randomly located memory. State clearly which of your numbers are order-of-magnitude approximations.

1.5 **CORE** Your colleague replaced an $O(n \log n)$ sort with an $O(n)$ bucketing pass and is puzzled that the "faster" version runs slower on their real dataset. Give two distinct explanations grounded in this chapter, and say what single action would settle which one (if either) is right.

1.6 **FADED** *Worked, with the last step for you.* You must argue in a design review that a proposed change — switching a hot lookup table from a hash map spread across the heap to a sorted array searched by binary search — could be faster *despite* binary search's $O(\log n)$ being asymptotically worse than the hash map's $O(1)$. Steps 1-3 are done; complete step 4.

Step 1 (given): Both are cheap in big-O for realistic n , so scaling is not the deciding factor here — this is a constant-factor question.

Step 2 (given): The constant factor is dominated by the memory hierarchy: how many *slow* memory accesses each structure makes per lookup.

Step 3 (given): A hash map lookup typically jumps to one effectively random heap location (likely a cache miss); a binary search over a small contiguous sorted array touches a handful of locations, but they are close together and the array may sit largely in cache.

Step 4 (you): Complete the argument — state the conclusion about which is likely to make fewer costly trips to slow memory, and name the one thing you would do before trusting the argument. Explain *why that final step and not simply "pick the lower big-O."*

1.7 **MIXED** For each situation, decide whether the RAM/big-O model or the real cost model is the right tool, and say why in one line: (a) choosing between a nested-loop join and a hash join for tables that will grow to billions of rows; (b) explaining why two $O(n)$ array-processing loops on a fixed 100 MB array differ by $8\times$ in wall-clock time; (c) deciding whether an algorithm will still be viable when the input is $1000\times$ bigger; (d) explaining why a linked-list traversal is slower than an array traversal of the same length.

1.8 **MIXED** A teammate presents two claims: (i) "This algorithm is $O(n)$, so it's optimal." (ii) "This one's $O(n)$ too but ten times faster in practice." Are these claims in contradiction? Decide which cost model each claim belongs to, explain whether both can be true at once, and state what information you would still need to choose between the two implementations for production.

1.9 **CORE** Define *mechanical sympathy* in your own words, and explain the ordering rule this chapter gives: why you apply it *after* choosing a good algorithm, not instead of choosing one. Give a one-sentence example of each failure mode (knowing only big-O; knowing only micro-optimization).

1.10 **STRETCH** The latency table in §1.3 is labelled "orders of magnitude only, machine-dependent." Give three distinct reasons the specific numbers cannot be trusted as exact for *your* machine, and explain why the *ratios* between tiers are nonetheless far more stable than the absolute values. What would you do to get numbers you could actually trust for a performance decision?

1.11 **LAB** On your machine, discover the real numbers instead of trusting the book. Run `getconf LEVEL1_DCACHE_LINESIZE` and `lscpu` (or `sysctl hw` on macOS) and write down your L1/L2/L3 cache sizes and the cache-line size. Then, in any language, write a loop that sums a large array that comfortably exceeds your L3 size and one that fits inside L1, timing each per element. Report the two per-element times and their ratio. State clearly which numbers are yours (the cache sizes, the measured ratio) and which came from the book (nothing — that's the point of the exercise).

Chapter 2 — The CPU Up Close

Before you start: Chapter 1 (the real machine vs. the RAM model — the two cost models). · **Pairs with:** Chapter 3 (the memory hierarchy — the other reason "one instruction = one unit" fails). This chapter attacks the RAM model's *second* false assumption.

BEFORE YOU READ ON

From Chapter 1, recall and answer from memory: (1) The RAM model makes three false assertions about real hardware; §1.1 named them. One was "all memory accesses cost the same" (the memory-hierarchy leak). What was the *second* one — the one this chapter is about? (2) Why is "one operation = one unit of time" called a *useful lie* rather than simply a mistake? (3) *Predict:* a modern CPU executes instructions one after another, finishing each before starting the next — true or false? Answers resolve in §2.1-2.2.

In Chapter 1 we said the RAM model tells three lies, dealt with the first (memory is not uniform-cost), and promised to return for the second: that *one instruction takes one unit of time*. This is the chapter. A modern CPU does not execute your instructions one at a time, in order, each in a tidy tick of the clock. It executes many at once, frequently out of their written order, and it *guesses* — commits to work before it knows whether that work is needed. Understanding this is what separates code that keeps the processor busy from code that leaves it stalling, and the difference, for the same big-O, can again be a large constant factor.

We will stay conceptual and correct. This chapter describes *mechanisms* — the pipeline, superscalar and out-of-order execution, branch prediction, instruction-level parallelism — and gives their costs only as orders of magnitude with explicit caveats, because the exact pipeline depth and misprediction penalty are microarchitecture-specific numbers that differ from chip to chip and that this book will not invent for you. The mechanisms, though, are general and stable, and knowing them changes how you read a hot loop. The through-line: **straight-line, predictable, independent work is fast because it keeps the machine full; branchy, dependent, unpredictable work is slow because it forces the machine to wait or to throw work away.**

2.1 Fetch-decode-execute, and why it isn't one step


At the coarsest level a CPU repeats a cycle: **fetch** the next instruction from memory, **decode** it into control signals, **execute** it (compute, or read/write memory), and make its results available. This is the von Neumann fetch-decode-execute loop, and if a processor really did it one instruction at a time — finish all stages for instruction 1, then start instruction 2 — then "one instruction = one unit of time" would be roughly true and this chapter would be short. Early processors worked essentially that way.

The trouble is that the stages use *different parts* of the chip. While the execute unit is busy with instruction 1, the fetch unit is idle; while instruction 1 is being decoded, the fetch unit is idle again. Doing one instruction start-to-finish before beginning the next

wastes most of the hardware most of the time. The entire architecture of a fast CPU is a series of answers to the question *how do we stop leaving hardware idle?* — and each answer breaks the RAM model's tidy accounting a little more. We take them in order: pipelining (overlap the stages), superscalar (duplicate the units), out-of-order (reorder to avoid waiting), and speculation (guess so you never have to wait for a branch). Bryant and O'Hallaron develop this progression in detail in *Computer Systems: A Programmer's Perspective* (CS:APP), Chapters 4 and 5; we take the conceptual tour.

2.2 Pipelining: overlap the stages

The first idea is an assembly line. Split instruction processing into stages (conceptually: fetch, decode, execute, memory, write-back — real designs have more, and the exact count is a microarchitecture-specific number we will not invent), and let each stage work on a *different* instruction at the same time. While instruction 1 is in "execute," instruction 2 is in "decode," and instruction 3 is in "fetch." Once the pipeline is full, an instruction *completes* every clock cycle even though each individual instruction still takes several cycles to pass all the way through. Throughput goes up by roughly the number of stages, though the latency of any single instruction does not.



Once full, one instruction completes per cycle (F=fetch, D=decode, E=execute, M=memory, W=write-back).

Figure 2.1: A pipeline overlaps stages so different instructions occupy different units in the same cycle. Stage names and count are illustrative; real pipelines have more stages.

Pipelining is why "one instruction = one unit of time" is doubly wrong: an instruction takes several cycles of latency, yet the machine can *complete* one (or more) per cycle. But the assembly line only runs smoothly if the next instruction is always ready and its inputs are available. Two things break that smooth flow — a dependency where an instruction needs a result that is not ready yet (a **data hazard**, which forces a stall), and a branch where the CPU does not yet know *which* instruction comes next (a **control hazard**). The rest of the chapter is about how the hardware fights these two, and how your code can help or hurt that fight.

QUICK CHECK

In a filled pipeline, an instruction's *latency* is several cycles but the pipeline's *throughput* is about one instruction per cycle. In one sentence, why are those two numbers different? (Answer after the next section.)

2.3 Superscalar and out-of-order: more units, smarter order

If one assembly line is good, several are better. A **superscalar** processor has multiple execution units and can fetch, decode, and issue several instructions *per cycle*, so its peak throughput is more than one instruction per cycle — provided it can find independent instructions to feed the parallel units. That proviso is everything. Two instructions can run in the same cycle only if neither needs the other's result; if instruction 2 uses instruction 1's output, they must run in sequence no matter how many units are idle.

To keep those units fed, modern CPUs execute **out of order**. Rather than stalling the whole machine when the next instruction in program order is waiting on a slow input (say, a value still coming from DRAM — recall [Chapter 1](#): that is hundreds of cycles), the processor looks ahead in the instruction stream, finds later instructions whose inputs *are* ready, and executes those now, filling the gap. Results are then reassembled into the original program order before they become visible, so the program's observed behavior is still that of straightforward sequential execution — the reordering is invisible to correctness but decisive for speed. This is why a cache miss does not necessarily stall the CPU: if there is independent work nearby, the machine does that work while the miss is serviced, hiding the latency.

The practical lesson lands directly on how you write hot code. The processor can only overlap work it can prove is independent, and it can only look a bounded distance ahead. Code with lots of **instruction-level parallelism** (ILP) — many operations that do not depend on each other's results — gives the out-of-order engine plenty to chew on and runs near the machine's peak. Code that is a long *dependency chain* — each step needing the previous step's result — starves the parallel units no matter how many there are, because there is never more than one thing that can run next. Same instruction count, very different speed; and, as in [Chapter 1](#), the difference is exactly the constant factor big-O throws away.

(Answer to the quick check: latency counts the cycles for one instruction to traverse all stages; throughput counts how often a new instruction completes. Because stages overlap, a new instruction can finish every cycle even though each one individually took several cycles to get through.)

THE SAME IDEA THREE WAYS: DEPENDENCY CHAINS VS. INDEPENDENT WORK

1. In words. Summing an array into a single accumulator forces each add to wait for the previous add's result — a dependency chain, low ILP. Summing into several independent partial accumulators and combining them at the end lets the superscalar, out-of-order engine run several adds per cycle — high ILP.

2. As a picture. One long chain $a \rightarrow a \rightarrow a \rightarrow a$ (each arrow "must finish before") versus four short parallel chains that only merge at the very end.




Figure 2.2: A dependency chain serializes execution; independent operations expose parallelism the CPU can exploit. Combining partial sums at the end recovers the same result.

3. As a trace. Both compute the same sum of n numbers — identical operation count, identical $O(n)$. The chained version can retire roughly one add per dependency latency; the four-accumulator version can keep several execution units busy at once, finishing in a fraction of the wall-clock time. (This is a real and measurable effect; CS:APP Chapter 5 develops it as "loop unrolling with multiple accumulators," and the performance-engineering chapters of this volume return to it.)

2.4 Branch prediction and the cost of a miss

Now the control hazard. When the CPU reaches a conditional branch — an `if`, a loop test, a `switch` — it does not yet know which way execution will go, because that depends on a condition that may not be computed yet. But the pipeline needs to fetch *something* next; it cannot sit idle for the several cycles it takes to resolve the branch. So the processor **predicts** which way the branch will go and speculatively executes down the predicted path, doing real work on the assumption that its guess is right.

Modern **branch predictors** are remarkably good — they track the history of each branch and find patterns, and on predictable branches (a loop that runs a million times and exits once; a check that is almost always false) they are right the overwhelming majority of the time, so the speculation is nearly free and the pipeline never stalls. This is why a loop with a highly predictable exit condition runs fast: the CPU correctly guesses "keep looping" every time until the one time it does not.

The cost arrives on a **misprediction**. When the branch resolves and the guess was wrong, all the speculatively executed work on the wrong path must be discarded, and the pipeline must be refilled from the correct target — the instructions already in flight are thrown away and the machine restarts fetching from the right place. The penalty is roughly proportional to how deep the pipeline is: everything that was in flight is lost. On modern deeply pipelined CPUs this penalty is on the order of *ten to a few tens of cycles*

per misprediction — an order of magnitude, not a figure to quote precisely, since the exact number is specific to each microarchitecture and this book will not invent one for a particular chip. (CS:APP, Chapter 5, works a concrete example for one processor; treat any single number as machine-specific.) What matters is the shape: a mispredicted branch costs many times what a correctly predicted one does.

The consequence for your code is direct and sometimes surprising. A branch that is *unpredictable* — one that goes each way about half the time, in no pattern the predictor can learn — pays the misprediction penalty roughly half the time it executes. In a tight loop, that can dominate the loop's cost. This is the mechanism behind a famous and initially baffling observation: processing a *sorted* array can be markedly faster than processing the same data *unsorted*, when the loop body contains a data-dependent branch, because sorting makes the branch predictable (long runs of the same outcome) whereas the shuffled data makes it a coin flip the predictor cannot learn. The arithmetic is identical; only the predictability changed. When you see a performance mystery like that, the branch predictor is a prime suspect.

TRAP: THE UNPREDICTABLE BRANCH IN THE HOT LOOP

The error is to assume that because two versions of a loop do the same arithmetic, they cost the same — forgetting that a *data-dependent, unpredictable branch* inside a hot loop can add a misprediction penalty on a large fraction of iterations, dwarfing the arithmetic. It is the CPU-level cousin of the [Chapter 1](#) trap ("same big-O, so same speed"): same operation count, very different wall-clock time, because one version keeps the pipeline full and the other keeps flushing it. The tells: a branch whose outcome depends on the data and follows no pattern, in code that runs an enormous number of times. The fixes are the standard toolkit — make the branch predictable (e.g. by sorting), or remove it entirely by turning control flow into data flow (branchless code, conditional-move or arithmetic tricks, table lookups, or SIMD masks — the performance-engineering chapters cover these). But confirm with measurement first: on *predictable* branches the predictor already makes the branch nearly free, and a branchless rewrite there just adds complexity for no gain.

2.5 What this means for the code you write

Step back and the theme of the chapter is one sentence: **the CPU runs fastest when you feed it a steady stream of predictable, independent work, and stalls when you feed it dependencies and surprises.** That single idea reorganizes a pile of otherwise-disconnected performance folklore into consequences of how the machine actually works:

- **Straight-line code with predictable branches** keeps the pipeline full and speculation cheap — fast.
- **Independent operations (high ILP)** let the superscalar, out-of-order engine run several things per cycle — fast.
- **Long dependency chains** serialize execution and starve the parallel units — slow, even at the same instruction count.
- **Unpredictable, data-dependent branches** cause mispredictions that flush the pipeline — slow, sometimes dramatically.

- **A cache miss (Chapter 1) need not stall the CPU** if there is independent work nearby for the out-of-order engine to do while it waits — which is why ILP and memory behavior are entangled, and why both matter together.

None of this changes an algorithm's big-O, and none of it should tempt you to abandon big-O — a pipeline-friendly $O(n^2)$ is still a disaster. This is, exactly as in [Chapter 1](#), the *second* cost model at work: after you have chosen a good algorithm, the constant factor that decides real speed is set by how well your code fits the machine — its memory hierarchy (Chapter 1 and [Chapter 3](#)) and its pipelined, speculative, superscalar core (this chapter). [Chapter 3](#) now returns to the memory side and turns "data location matters" into the concrete, exploitable rules of caches and locality — the other half of mechanical sympathy.

SAY IT LIKE A SYSTEMS ENGINEER

Try it first. In a review, someone asks why sorting an array before a filtering loop made the loop faster, even though sorting is extra work and the loop does the same comparisons either way. Cover the paragraph below and explain it in two or three sentences, using *branch predictor*, *misprediction*, and *pipeline*.

Model answer. "The loop has a data-dependent branch. On the unsorted data that branch goes each way unpredictably, so the branch predictor is wrong about half the time, and every misprediction flushes the pipeline — a penalty on the order of tens of cycles per miss, paid on a huge fraction of iterations. Sorting makes the branch outcomes come in long predictable runs, so the predictor is almost always right and the pipeline stays full; the arithmetic is identical, but the misprediction cost nearly vanished. Net of the one-time sort, the loop got cheaper because we made it *predictable*, not because we did less arithmetic." Naming the mechanism — rather than saying "sorting is just faster" — is what marks the answer as senior.

CAN YOU... WITHOUT LOOKING?

- Explain **pipelining** and why it makes an instruction's *latency* (several cycles) differ from the pipeline's *throughput* (about one per cycle).
- Say what **superscalar** and **out-of-order** execution each add, and why both depend on finding *independent* instructions.
- Contrast a **dependency chain** (low ILP, serialized) with **independent work** (high ILP, parallel) and give the multiple-accumulator sum as the example.
- Describe **branch prediction** and the cost of a **misprediction** (flush the pipeline; penalty on the order of tens of cycles — machine-specific, an order of magnitude only).
- Explain why processing a **sorted** array can beat the same data unsorted when the loop has a data-dependent branch.
- State the chapter's one sentence: the CPU is fast on predictable, independent work and stalls on dependencies and surprises — and why this is the RAM model's *second* lie.

RETRIEVAL — CLOSED BOOK

Cover the chapter. Answer from memory; rate your confidence before checking. An overconfident miss is the signal to reread.

1. **R1.** Why does executing one instruction fully before starting the next waste the hardware, and what is pipelining's fix? (§2.1–2.2)
2. **R2.** What must be true of two instructions for a superscalar CPU to run them in the same cycle, and what does out-of-order execution do when the next in-order instruction is stalled? (§2.3)
3. **R3.** Define instruction-level parallelism. Why is a single-accumulator sum lower ILP than a multi-accumulator sum computing the same total? (§2.3)
4. **R4.** What does a CPU do at a conditional branch before the condition is known, and what exactly is discarded on a misprediction? Give the penalty as an order of magnitude with its caveat. (§2.4)
5. **R5.** Explain the sorted-vs-unsorted array speed difference in terms of the branch predictor. Which of the two cost models from Chapter 1 does this belong to? (§2.4–2.5)

FURTHER READING

- **Bryant & O'Hallaron, *Computer Systems: A Programmer's Perspective (CS:APP)***, Chapter 4 ("Processor Architecture") and Chapter 5 ("Optimizing Program Performance"). Chapter 4 builds a pipelined processor from scratch; Chapter 5 covers superscalar/out-of-order machines, branch misprediction, dependency chains, and multiple accumulators, all with concrete (processor-specific) numbers. The primary source for this chapter.
- **Agner Fog, *optimization manuals*** (agner.org, free). Deep, empirical detail on microarchitecture, pipelines, and branch prediction across real x86 CPUs. Dated in places — verify against current parts — but the mechanism-level explanations are excellent.
- **MIT 6.172, *Performance Engineering of Software Systems*** (free OCW). The lectures on modern processors, ILP, and branch behavior connect these mechanisms to measurable speedups in real code.

Exercises

- 2.1** **WARM-UP** Name the four techniques, in the order this chapter introduces them, that a fast CPU uses to stop leaving hardware idle, and give a one-line description of what each one does.
- 2.2** **WARM-UP** Distinguish *latency* and *throughput* for a single instruction passing through a pipeline. Why can a pipeline complete about one instruction per cycle while each instruction still takes several cycles from start to finish?
- 2.3** **WARM-UP** Define a *data hazard* and a *control hazard* in one sentence each, and say which of the two branch prediction is designed to address.

2.4 **CORE** Two loops compute the same sum of an array of doubles. Loop A accumulates into a single variable; loop B accumulates into four independent partial sums and adds them at the end. Both do the same number of additions. Explain, using ILP and the out-of-order engine, why B can be substantially faster, and name the one condition on the additions that makes B's parallelism possible.

2.5 **CORE** A hot loop contains `if (x[i] > threshold) count++;`. On dataset P the branch is taken in long predictable runs; on dataset Q it is taken about half the time in no pattern. The arithmetic is identical. Explain why the loop can run several times faster on P, and estimate qualitatively where the extra time on Q goes.

2.6 **COST** A branch mispredicts with probability p per execution, and each misprediction costs about k cycles (an order-of-magnitude, machine-specific number — state that caveat). The loop body otherwise takes about b cycles. (a) Write the approximate average cost per iteration. (b) For what range of p does the misprediction term dominate the body? (c) Explain why making the branch predictable attacks p , and a branchless rewrite attacks k (by removing the branch), and why you would measure before choosing.

2.7 **FADED** *Worked, with the last step for you.* Explain why a cache miss (Chapter 1) does not always stall the CPU. Steps 1–3 given; complete step 4.

Step 1 (given): A load that misses to DRAM takes on the order of hundreds of cycles to return its value (Chapter 1).

Step 2 (given): A modern CPU executes out of order: when the next in-order instruction is waiting, it looks ahead for later instructions whose inputs are already available.

Step 3 (given): If there is independent work near the missing load, the machine can execute that work during the hundreds of cycles the miss is being serviced.

Step 4 (you): State the conclusion about whether the miss stalls the CPU, and name the property of the surrounding code that determines whether the latency gets hidden or not. Explain *why that property, and not the miss latency alone, decides the outcome*.

2.8 **MIXED** For each observation, decide whether the dominant mechanism is the *memory hierarchy* (Chapter 1), *branch prediction*, or an *instruction-level dependency chain* (this chapter) — and justify in one line: (a) a loop got $3\times$ faster after the input array was sorted, with no change to the arithmetic; (b) a loop got faster after splitting one accumulator into four; (c) a traversal got much slower when the same data was stored as a linked list instead of an array; (d) a tight numeric loop with no branches and no memory pressure is already near peak throughput.

2.9 **MIXED** A colleague says: "My function is $O(n)$ and yours is $O(n)$, but yours runs $5\times$ faster — that's impossible, big-O says they're the same." Decide which cost model resolves the paradox, list three distinct mechanisms from Chapters 1–2 that could each independently produce a $5\times$ gap at equal big-O, and state the single action that would identify which one is responsible here.

2.10 **STRETCH** Speculation means the CPU does work it might have to throw away. (a) Explain why this is still a net performance win on typical code. (b) Give a case where heavy speculation on an unpredictable branch is a net *loss*. (c) Speculative execution has also had *security* consequences in real processors. Without inventing specifics, state at a high level why executing instructions before you know they should run could, in principle, leak information — and note that the precise mechanisms and mitigations are a real, separate topic you would need primary sources to describe accurately.

2.11 **LAB** Measure a dependency chain. In C (or any compiled language), write two loops that sum a large array of `double`: one with a single accumulator, one with four independent accumulators combined at the end. Compile with optimization *off* for the reduction (or use a language/setting that will not auto-vectorize away the difference), time both, and report the ratio. State which numbers are yours and note that the exact ratio depends on your CPU and compiler. If you cannot prevent the compiler from transforming the loop, say so and report what you observed instead — do not report a number you did not actually measure.

Chapter 3 — The Memory Hierarchy & Caches

Before you start: Chapter 1 (the memory wall and the latency hierarchy), Chapter 2 (the CPU; why a miss need not stall). · **Pairs with:** the rest of Part 1 and the performance-engineering chapters, which turn these rules into measured speedups. This chapter makes "data location matters" concrete and exploitable.

BEFORE YOU READ ON

From the previous two chapters, from memory: (1) Chapter 1 said accessing DRAM costs on the order of ~100 ns while a nearby cache hit costs ~1 ns — about how many orders of magnitude apart is that? (2) Chapter 2 said a cache miss does *not* always stall the CPU. Under what condition does the latency get hidden? (3) *Predict:* you sum every element of a large 2-D array. Does the *order* in which you visit the elements — row by row vs. column by column — change how long it takes, given that you touch exactly the same elements either way? By the end of this chapter you will have a measured answer.

Chapter 1 told you the memory wall exists and that hardware answers it with a hierarchy of caches. Chapter 2 told you the CPU can sometimes hide a miss behind other work. This chapter is where "memory access cost varies" stops being a warning and becomes a set of concrete, mechanical rules you can design against — because caches are not magic, they follow simple policies, and once you know the policies you can write code that hits the cache instead of missing it. The payoff is the largest, cheapest performance win available to most programs: not a better algorithm, just the same algorithm arranged so the data is where the CPU expects it.

We will finish with a measurement, not a claim. At the end of the chapter is a short C program that traverses the same array two ways — the cache-friendly way and the cache-hostile way — and the exact speed ratio it produced when run for this book, on a specific machine, so you can see the effect is real and roughly how big it is. The number is machine-dependent, and we will say so; but it is measured, not invented.

3.1 The cache line: memory moves in blocks, not bytes

The first and most important fact about caches overturns a picture you have probably carried since you first learned about memory: that you read one byte, or one `int`, at a time. You do not. When the CPU needs a value that is not already in cache, it does not fetch just that value — it fetches an entire aligned **cache line** (also called a cache block) containing it, and installs the whole line in the cache. On the common x86-64 processors this line is **64 bytes**; Bryant and O'Hallaron state this standard size in *Computer Systems: A Programmer's Perspective* (CS:APP, Chapter 6), and it is what `getconf LEVEL1_DCACHE_LINESIZE` reports on such machines. (It is a hardware parameter, not a law of nature — verify it on your target — but 64 bytes is the near-universal value on current x86-64.)

This single mechanism explains most of cache behavior. When you touch one byte, the seven-plus neighbors sharing its 64-byte line come along for free — they are now in cache, and touching them next costs a cheap hit instead of an expensive miss. So the cost of a memory access is not really "per access"; it is "per line brought in." A program that uses every byte of each line it fetches gets the full value of every expensive DRAM trip. A program that touches one byte per line and then moves far away wastes 63 of every 64 bytes it pays to move — it pays full price for the line and uses a fraction of it. That ratio, *useful bytes per line fetched*, is the hidden efficiency knob behind an enormous amount of performance.

QUICK CHECK

On a machine with 64-byte cache lines, you read one `int` (4 bytes). How many bytes does the hardware actually move into the cache, and how many neighboring `ints` are now cheap to access? (Answer after the next section.)

3.2 Locality: the property caches reward

Caches work — they successfully hide the memory wall — because real programs exhibit **locality of reference**, a strong tendency to access memory that is close in time or space to what they just accessed. There are two kinds, and CS:APP (Chapter 6) frames them as the two properties good code should have:

- **Temporal locality:** if you access a location, you are likely to access it *again soon*. A loop variable, an accumulator, a hot lookup table — accessed repeatedly, so after the first (missing) access they sit in cache and every later access is a hit.
- **Spatial locality:** if you access a location, you are likely to access *nearby* locations soon. Walking an array in order, reading struct fields together — because memory arrives a whole line at a time, spatial locality means the neighbors you fetched are exactly the ones you are about to use.

The hardware bets on both. The cache keeps recently used lines (exploiting temporal locality), and by fetching a full line it pre-loads spatial neighbors (exploiting spatial locality). Some CPUs also *prefetch* — detect a sequential access pattern and fetch lines ahead of the program's demand, so the data is already resident by the time the loop reaches it. All of this is automatic and invisible until you fight it. Your job is to write code whose access pattern *has* locality, so that the hardware's bets pay off. Code with good locality runs near the speed of cache; code that scatters its accesses runs near the speed of DRAM; and from [Chapter 1](#) you know that is about a 100× difference, available on the same algorithm.

(Answer to the quick check: the hardware moves the full 64-byte line, so 64 bytes; that line holds 16 `ints`, so the 15 neighbors sharing the line are now cheap hits.)

3.3 Hits, misses, and the three C's

An access that finds its line already in cache is a **hit**; one that does not is a **miss**, and a miss pays to fetch the line from the next level down (another cache, or DRAM). Since misses are where the time goes, it helps to classify *why* a miss happened. The classic taxonomy, the **three C's** (presented in CS:APP, Chapter 6), sorts every miss into one of three causes:

- **Compulsory (cold) miss:** the first-ever reference to a line. It was never in cache, so the first touch must miss. Unavoidable in principle — you have to load data at least once — but you can make each compulsory miss *count* by using the whole line (spatial locality) once you have paid for it.
- **Capacity miss:** the line was in cache but got evicted because the program's working set — the data it is actively using — is simply larger than the cache. If you stream through 512 MB repeatedly and the cache holds a few MB, lines you will need again get pushed out before you return to them. The fix is to shrink the working set (e.g. process data in cache-sized *blocks*, a technique called blocking or tiling).
- **Conflict miss:** the line was evicted not because the cache was full overall, but because too many of the addresses you are using map to the *same* small set of cache slots and crowd each other out. This one depends on the cache's *associativity*, which we take next.

3.4 Associativity, conceptually

A cache cannot let any line live in any slot — searching the whole cache on every access would be too slow. Instead each memory address maps to a restricted set of slots. In a **direct-mapped** cache, each address has exactly one slot it may occupy; simple and fast, but two hot addresses that happen to map to the same slot will endlessly evict each other even if the rest of the cache is empty — that is a conflict miss. At the other extreme, a **fully associative** cache lets a line go in any slot (no conflict misses, but expensive to build). Real caches are **set-associative**: a compromise where each address maps to a small *set* of slots (say 8, an "8-way" cache) and the line may occupy any slot within its set. This kills most conflict misses while staying cheap enough to build.

You rarely need to compute set indices by hand, and this book will not turn associativity into arithmetic drills. What you need is the intuition: conflict misses are why a data structure can perform badly at one size or memory alignment and fine at another, seemingly at random. A classic symptom is a loop that slows sharply when an array's row length (its "stride") is a power of two that aligns many accesses onto the same cache sets. When performance changes with the exact size or alignment of your data rather than the amount of it, suspect conflict misses — and know the deeper treatment (and the set-index arithmetic) is in CS:APP, Chapter 6.

QUICK CHECK

Classify each miss: (a) the very first read of a freshly allocated array; (b) re-reading a 1 GB array in a loop when the cache holds only a few MB; (c) two hot variables that keep evicting each other though the cache is mostly empty. (Answer below the demonstration.)

3.5 Why arrays are fast and pointer-chasing is slow

Now the payoff that ties the chapter together, and the reason "same big-O, different speed" keeps recurring. Consider two ways to store a sequence of integers you will walk from start to finish. An **array** stores them contiguously: element $i+1$ sits right after element i in memory. A **linked list** stores each element in its own separately allocated node, connected by pointers; consecutive nodes can be anywhere in the heap, often far apart. Both take $O(n)$ to traverse — the same number of "steps" — and the RAM model says they cost the same. The real machine disagrees, sometimes by more than $10\times$.

Walking the array has excellent spatial locality: each 64-byte line brought in holds many consecutive elements, so most accesses are hits, and the hardware prefetcher, seeing the sequential pattern, fetches lines ahead of you so the data is waiting. Walking the linked list is **pointer-chasing**: each node's address is only known once you have read the *previous* node's pointer, so the accesses are (a) too scattered, effectively random addresses — poor spatial locality, a likely cache miss per node — and (b) a strict dependency chain, since you cannot fetch node $i+1$ until node i arrives. That dependency defeats the latency-hiding trick from [Chapter 2](#): there is no independent work to overlap with the miss, so the CPU stalls on the full DRAM latency at every step. Poor locality and a serial dependency, together, are why pointer-chasing through memory is one of the slowest things a program can do, and why cache-friendly, array-based layouts underpin most high-performance data structures. (This is the mechanism behind the array-vs-list remark back in [Chapter 1](#)'s warning box — now you know exactly why.)

3.6 A measured demonstration: row-major vs. column-major

Let us prove the effect rather than assert it, with the cleanest possible experiment: the *same* arithmetic on the *same* data, differing only in access order. Store an $N\times N$ matrix of `long` in a single contiguous block in row-major order (row 0, then row 1, ...), and sum every element two ways. The **row-major** traversal walks memory in order — inner loop over columns, so consecutive accesses are adjacent in memory (stride 1), excellent spatial locality. The **column-major** traversal walks down columns — inner loop over rows, so each step jumps N elements forward in memory (stride N), touching a different cache line almost every time. Same elements, same additions, same result; only the order differs.


```

#define N 8192                /* an N x N matrix of long = 512 MiB, far larger than cache */


/* row-major: inner index j varies fastest -> contiguous, stride 1 */
for (int i = 0; i < N; i++)
    for (int j = 0; j < N; j++)
        sum_row += a[(size_t)i * N + j];

/* column-major: inner index i varies fastest -> stride of N longs per step */
for (int j = 0; j < N; j++)
    for (int i = 0; i < N; i++)
        sum_col += a[(size_t)i * N + j];


```

Both loops produced the identical sum (a correctness check that they really do the same work). Built with `gcc -O2` and run for this book on a 12th-generation Intel Core i7 (i7-12700H) laptop, on a 512 MiB matrix that dwarfs the cache, the row-major traversal took roughly **0.04 s** and the column-major traversal roughly **0.95-0.99 s** — the column-major version was about **24-26× slower** across repeated runs, for doing the very same additions in a different order. **This exact ratio is machine-dependent** — it depends on cache sizes, prefetcher, memory bandwidth, and compiler, and you should expect a different number on your hardware — but the *direction and rough magnitude* (a large multiplier, not a few percent) is the robust, reproducible result, and you can reproduce it yourself with the lab exercise. The mechanism is exactly §3.1–3.2: the row-major loop uses all 8 longs of every 64-byte line it fetches and lets the prefetcher run ahead; the column-major loop touches one long per line and throws the other 7 away, so it issues many times more line fetches and stalls on the misses it cannot hide.

row-major (stride 1): uses the whole line



column-major (stride N): one use per line, 7 wasted



Measured (i7-12700H, 512 MiB, gcc -O2): column-major ~24-26× slower. Machine-dependent.

Figure 3.1: The two traversals do identical arithmetic. Row-major consumes every element of each fetched line; column-major uses one element per line and moves on, wasting most of the memory bandwidth it pays for.

(Answer to the §3.4 quick check: (a) compulsory/cold — first-ever reference; (b) capacity — the working set exceeds the cache; (c) conflict — addresses collide on the same cache set despite free space elsewhere.)

THE SAME IDEA THREE WAYS: LOCALITY PAYS

1. In words. Memory arrives one cache line at a time, so code that uses all of each line it fetches (good spatial locality) moves far less memory and hits far more often than code that touches one item per line and jumps away.

2. As a picture. Figure 3.1: a fully-consumed line vs. a line touched once and abandoned.

3. As a trace / measurement. The row-major vs. column-major experiment: identical sum, identical operation count, ~24-26× wall-clock difference on the test machine because one order has spatial locality and the other does not. Prose says "locality matters"; the number says how much.

TRAP: IGNORING CACHE LOCALITY (THE "SAME BIG-O, SO SAME SPEED" FALLACY, MEMORY EDITION)

The recurring error of this Part, now at the cache level: judging two implementations equivalent because they have the same complexity and the same operation count, while ignoring that one has good locality and the other does not. Column-major matrix traversal, linked lists walked as pointer chains, arrays-of-pointers-to-objects scattered across the heap, hash tables with poor probe locality — all can be an order of magnitude slower than a cache-friendly layout doing the identical logical work. The tell is a hot loop over a data structure larger than cache whose access pattern is *not* sequential. The fixes are layout, not algorithm: store data contiguously in the order you traverse it; keep frequently-used fields together; process large data in cache-sized blocks (blocking/tiling); prefer arrays of values to arrays of pointers. And, as always in this Part: confirm by *measuring*, on your data and your machine, rather than trusting any single ratio — including this chapter's ~25×, which is one machine's number, not a constant of nature.

SAY IT LIKE A SYSTEMS ENGINEER

Try it first. A reviewer asks why you rewrote a nested loop to swap the order of the two indices when "it computes exactly the same thing." Cover the paragraph below and answer in two or three sentences, using *cache line*, *spatial locality*, and *stride*.

Model answer. "It computes the same result, but the original visited the array with a large stride, touching one element per 64-byte cache line and then jumping to a different line — so it moved roughly a line's worth of memory per element and got almost no reuse. Swapping the loop order makes the inner traversal stride-1, so each fetched line is fully consumed before we move on: far fewer misses, the prefetcher can run ahead, and we use the memory bandwidth we're already paying for. Same big-O and same arithmetic; the change is purely spatial locality, and on our data it measured about an order of magnitude faster." Naming the mechanism — line, stride, locality — rather than "I made it faster" is the senior register.

3.7 The rules, and where Part 1 goes next

Three chapters in, the cost model the RAM abstraction hides is now concrete. [Chapter 1](#) established that where data lives dominates cost and that this is a second cost model alongside big-O. [Chapter 2](#) opened the CPU and showed why predictable, independent work runs fast. This chapter turned "data location matters" into rules you can design against: memory moves in 64-byte lines, so use the whole line (spatial locality) and reuse

hot data (temporal locality); classify your misses (compulsory, capacity, conflict) to know which fix applies (use each line, shrink the working set, avoid alignment collisions); and remember that contiguous array access is fast while pointer-chasing is slow because it has neither locality nor independent work to hide the misses. These are not micro-optimizations to sprinkle on at the end — they are how you get an order of magnitude on the same algorithm, and you can now *see* them in code.

The rest of Part 1 builds directly on this. The next chapters take the model down to how data is actually laid out in memory — bytes, alignment, the stack and the heap — so you can control locality deliberately rather than by accident, and the performance-engineering chapters return with profiling tools that *show* you the misses and mispredictions these three chapters taught you to predict. The habit to carry forward is the one this Part has drilled from three angles: when two pieces of code have the same big-O but different speed, the difference is the machine — its memory hierarchy and its pipeline — and you now know where to look.

CAN YOU... WITHOUT LOOKING?

- State the **cache line** size on common x86-64 (64 bytes, per CS:APP) and explain why memory moving in lines, not bytes, is the key fact behind cache behavior.
- Define **temporal** and **spatial** locality and say which one a full-line fetch and which one a prefetcher exploits.
- Name the **three C's** — compulsory, capacity, conflict — and give the fix each one points to.
- Explain **associativity** conceptually (direct-mapped vs. set-associative vs. fully associative) and why conflict misses make performance depend on size/alignment, not just data volume.
- Explain why **array traversal is fast and pointer-chasing is slow** in terms of locality *and* the dependency chain from [Chapter 2](#).
- Describe the **row-major vs. column-major** demonstration, its measured order-of-magnitude result, *and* why the exact number is machine-dependent.

RETRIEVAL — CLOSED BOOK

Cover the chapter. Answer from memory; rate confidence first.

1. **R1.** What is a cache line, what is its common size on x86-64, and why does its existence mean cost is "per line fetched," not "per byte"? (§3.1)
2. **R2.** Define temporal and spatial locality and give a code example of each. (§3.2)
3. **R3.** Name the three C's and, for each, the design change that reduces it. (§3.3)
4. **R4.** Why can a loop's speed change sharply with the array's size or alignment even when the amount of data is the same? Which of the three C's is this? (§3.4)
5. **R5.** Give two independent reasons a linked-list traversal is slower than an array traversal of the same length, one about locality and one about dependencies. (§3.5)

FURTHER READING

- **Bryant & O'Hallaron, *Computer Systems: A Programmer's Perspective (CS:APP)***, Chapter 6 ("The Memory Hierarchy"). The source for this chapter's cache-line size, locality framing, three-C's taxonomy, and set-associativity arithmetic, plus the "memory mountain" experiment that visualizes read throughput across sizes and strides. The free CMU 15-213 materials cover it.
- **Drepper, "What Every Programmer Should Know About Memory"** (2007, free). The deepest freely available treatment of cache organization, associativity, prefetching, and access patterns. Its absolute numbers are dated — read for mechanism — but its analysis of why layout matters is unmatched.
- **MIT 6.172, *Performance Engineering of Software Systems*** (free OCW). Lectures on cache-efficient algorithms, blocking/tiling, and measuring cache behavior with real tools — the applied continuation of this chapter.

Exercises

- 3.1** **WARM-UP** On a machine with 64-byte cache lines, how many doubles (8 bytes each) fit in one line? If you read one double from a cold line, how many neighbors become cheap hits, and which kind of locality does that reward?
- 3.2** **WARM-UP** Give a one-sentence code example of temporal locality and one of spatial locality, and say which cache mechanism (line fill, retention of recent lines, prefetching) exploits each.
- 3.3** **WARM-UP** Classify each miss as compulsory, capacity, or conflict: (a) the first read of a just-allocated buffer; (b) repeatedly scanning a 4 GB dataset with an 8 MB last-level cache; (c) a loop that slows down only when an array's row length is a particular power of two.
- 3.4** **CORE** You must sum an $N \times N$ matrix stored in row-major order. One colleague writes the inner loop over columns, another over rows. (a) Which is faster and why, in terms of cache lines and stride? (b) Both are $O(N^2)$ and do the same additions — reconcile that with a large measured speed difference. (c) Name the kind of miss the slow version suffers most.
- 3.5** **CORE** A profile shows a data structure of 5 million small objects is slow to traverse. It is currently an array of *pointers*, each pointing to a separately allocated object. Propose a layout change that would improve locality, and explain both why it reduces misses (spatial locality) and why it helps the CPU hide the remaining misses (recall the dependency argument from Chapter 2).

3.6 **COST** A traversal touches one 8-byte value per 64-byte cache line (stride 8 values). (a) What fraction of each fetched line does it use? (b) Compared with a stride-1 traversal of the same values, roughly how many times more memory must it move? (c) Explain how this back-of-the-envelope ratio relates to (though it will not exactly equal) the measured row-major/column-major slowdown, and name two hardware factors that make the measured number differ from this simple estimate.

3.7 **FADED** *Worked, with the last step for you.* Explain why the column-major traversal in §3.6 stalls the CPU more than the row-major one, tying it to [Chapter 2](#). Steps 1–3 given; complete step 4.

Step 1 (given): Column-major access has stride N , so almost every access lands on a new cache line — a likely miss to DRAM (~100 ns, Chapter 1).

Step 2 (given): A cache miss need not stall the CPU if there is independent work nearby to overlap with it (Chapter 2 §2.3).

Step 3 (given): A summation into one accumulator, walking predictably, does have some independent memory-level parallelism — but the row-major version keeps far more useful data resident per fetched line, so it issues far fewer misses to begin with.

Step 4 (you): Complete the explanation: state which effect dominates the $\sim 25\times$ gap here — the sheer *number* of misses (locality) or the inability to hide each one (dependencies) — and justify your choice. Explain *why the number of misses is the primary lever in this particular experiment*, and what you would change about the code to test your claim.

3.8 **MIXED** For each speedup, decide whether the primary mechanism is *spatial locality / cache lines*, *a capacity miss fixed by blocking*, *branch prediction* (Chapter 2), or an *ILP / dependency chain* (Chapter 2), and justify in a line: (a) tiling a large matrix multiply into cache-sized blocks; (b) sorting the input before a data-dependent filter loop; (c) switching an array-of-structs traversal to consume each cache line fully; (d) summing into four accumulators instead of one.

3.9 **MIXED** Two implementations of the same $O(n)$ aggregation differ by $15\times$ in wall-clock time. Without running anything yet, list the candidate causes drawn from all of Chapters 1–3 (name at least three distinct mechanisms), then describe the *single* measurement you would take first to narrow it down and why that one is most informative.

3.10 **STRETCH** The chapter reports a specific $\sim 24\text{--}26\times$ row-major/column-major ratio and repeatedly insists it is machine-dependent. (a) List three properties of a machine that would change this ratio, and say for each whether it would tend to raise or lower it. (b) Explain why the *existence* and rough *magnitude* (a large multiplier) of the effect is nonetheless robust across machines, even though the exact number is not. (c) Why is reporting "about $25\times$ on this specific CPU, and reproduce it yourself" more honest and more useful than printing a single universal multiplier?

3.11 **LAB** Reproduce the demonstration on your own machine. First run `getconf LEVEL1_DCACHE_LINESIZE` and `lscpu` (or the macOS `sysctl hw` equivalents) and record your cache-line size and L1/L2/L3 sizes. Then write the two-loop matrix-sum program (size it so the matrix comfortably exceeds your last-level cache), verify both loops produce the identical sum, time each, and report *your* ratio. Note how it compares to the book's $\sim 25\times$ and to your \$3.6 stride estimate. Report only numbers you actually measured; if the compiler optimizes a loop away, change the code (e.g. consume the sum) and say what you did.

Chapter 4 — Virtual Memory & Address Translation

Before you start: Chapter 3 (caches; memory moves in lines; a miss to DRAM costs ~100 ns), Chapter 1 (the latency hierarchy). · **Pairs with:** Chapter 5 (data layout) and the operating-system Part, which builds the page tables this chapter only describes. This chapter adds one more cache to the picture — a cache for *address translations* — and shows why a huge, scattered working set can be slow for a reason the last chapter did not cover.

BEFORE YOU READ ON

From memory, reaching back: (1) Chapter 3 said the CPU keeps recently used cache lines so a repeated access is a cheap hit — what is the general trick there, in one word? (2) Chapter 1 put a DRAM access at ~100 ns and an L1 hit at ~1 ns — about how many orders of magnitude apart? (3) *Predict:* the pointer `p` in your C program holds an address like `0x7ffd3a08`. When the CPU reads `*p`, do you think that number is the actual electrical address of a cell in your DRAM chips — the physical location — or something the hardware has to *translate* first? By the end of the chapter you will know which, and what the translation costs.

Every address your program has ever used is a lie — a convenient, deliberate, hardware-enforced lie. When your code holds a pointer and dereferences it, the number in that pointer is a **virtual address**: a name in a private, per-process address space that the operating system and CPU maintain, not the **physical address** of a byte in the DRAM chips. Before the load can touch memory, the hardware must *translate* virtual to physical. That translation is invisible, automatic, and — this is the point for a performance book — *not free*. This chapter explains the mechanism (why virtual memory exists, how pages and page tables implement it) and then, in the register of Part 1, isolates its cost: the translation itself is cached, in a structure called the TLB, and when that cache misses, the machine pays a price that scales with how scattered your data is. The through-line from Chapter 3 is exact — **the TLB is a cache for the page table**, and everything you learned about hits, misses, and working sets applies to it too.

4.1 Virtual vs. physical addresses: why the lie exists

On a modern system, each process runs as if it owns a large, contiguous block of memory starting near address zero, all to itself. It does not. Physical DRAM is shared among every process and the kernel, is smaller than the address space the program is allowed to name, and is handed out in scattered fragments. The **virtual address space** is the illusion that papers over this reality: the hardware's memory-management unit (MMU), configured by the OS, maps each process's virtual addresses onto whatever physical frames happen to be free, and keeps different processes' identical-looking virtual addresses pointing at different physical memory. Bryant and O'Hallaron develop this fully in *Computer Systems: A Programmer's Perspective* (CS:APP, Chapter 9, "Virtual Memory"); Arpaci-Dusseau's *Operating Systems: Three Easy Pieces* (OSTEP) devotes its whole virtualization part to it.

The illusion buys three things at once, which is why every general-purpose OS pays for it. First, **isolation**: process A literally cannot name a byte belonging to process B, because A's page tables contain no mapping to B's frames — a security and stability boundary enforced in hardware. Second, **a simple programming model**: your linker can lay out code and data at fixed virtual addresses without knowing where in physical RAM the program will actually land, and every process gets the same clean layout. Third, **flexibility the OS exploits**: a virtual page can be mapped to a physical frame, or to a page of a file on disk, or to nothing yet (allocated lazily on first touch), or shared read-only between processes (as with a common library) — all invisibly to your code. For this chapter, hold onto the first fact and the cost it implies: because the mapping is per-process and arbitrary, *the hardware must consult a translation on essentially every memory access*, and that lookup is what we now follow.

4.2 Pages and page tables

Translating every individual byte address would need a map with one entry per byte — absurd. Instead memory is divided into fixed-size **pages**: the address space is chopped into aligned blocks, and translation happens one page at a time. The near-universal base page size on x86-64 (and many other architectures) is **4 KiB**; CS:APP uses this standard size, and it is what `getconf PAGE_SIZE` reports on such systems. A virtual address then splits cleanly into two parts: the high bits are the **virtual page number** (which page), and the low bits — 12 of them for a 4 KiB page, since $2^{12} = 4096$ — are the **offset** within the page. Only the page number needs translating; the offset is the same in virtual and physical addresses, because a page maps to a physical **frame** of the same size and a byte's position within the page does not move.

The **page table** is the data structure, held in memory and owned by the OS per process, that maps virtual page numbers to physical frame numbers (plus permission and status bits: present or not, readable/writable/executable, accessed, dirty). Conceptually it is one big array indexed by virtual page number. But a flat array is impractical: a 48-bit virtual address space with 4 KiB pages has 2^{36} pages, and a flat table with an entry per page would be enormous *per process*, the overwhelming majority of it empty because most programs use a tiny, sparse fraction of their address space. The standard fix is a **multi-level (hierarchical) page table**: a tree of tables where the virtual page number is split into several fields, each indexing one level. x86-64 uses a four-level page table (a documented architectural fact, described in CS:APP). Only the branches that correspond to actually-used regions of memory need to exist, so a sparse address space costs a sparse tree instead of a giant flat array.

The cost that matters shows up right here. To translate one virtual page number through a four-level table, the hardware must read *each level in turn*: read the top-level table to find the address of the next-level table, read that to find the next, and so on — up to **four dependent memory accesses** to complete one translation on x86-64. This descent is called a **page-table walk**. And notice it is a dependency chain of the exact kind [Chapter 3](#) warned about: each read's address comes from the previous read, so the walk cannot be parallelized, and if the table levels are not in cache, each step is its own trip toward

DRAM. Doing four dependent memory accesses *for every* load and store would erase all the performance we have been chasing — so, as always, the machine caches the result.

QUICK CHECK

A 4 KiB page needs how many offset bits, and why exactly that many? If a virtual address is split into a page number and an offset, which part does the page table translate, and which part is copied unchanged into the physical address? (Answer after the next section.)

4.3 The TLB: a cache for translations

The **translation lookaside buffer (TLB)** is a small, fast, on-CPU cache that stores recent virtual-page-to-physical-frame translations. It is a cache in exactly the sense of [Chapter 3](#) — same idea, different contents. Chapter 3's data caches store recently used *cache lines* so a repeated data access is a cheap hit; the TLB stores recently used *translations* so a repeated access to the same page skips the page-table walk. On each memory access the MMU first checks the TLB with the virtual page number:

- A **TLB hit** — the translation is present — returns the physical frame in roughly a cycle, effectively free, and the access proceeds. This is the common case, because programs have locality: consecutive accesses usually fall on the same few pages (a 4 KiB page holds a lot of data), so one page's translation, cached once, serves thousands of accesses.
- A **TLB miss** — the translation is not cached — triggers the **page-table walk** of §4.2: the several dependent memory accesses to descend the table, after which the freshly found translation is installed in the TLB (evicting an older one) so nearby future accesses hit. On modern x86-64 this walk is done by hardware; on some architectures the OS handles it in software, which is costlier still.

So the cost of a TLB miss is the cost of the walk: on the order of a handful of dependent memory accesses. If those accesses hit in the data caches (the upper page-table levels are hot and small, and are cached), a miss is cheap — a few cycles. If they miss to DRAM, each is a ~100 ns trip ([Chapter 1](#)), and a walk can cost on the order of tens to hundreds of nanoseconds. The exact penalty and the exact TLB capacity are hardware parameters that vary by CPU, and this book will not quote a specific chip's numbers as if they were universal — CS:APP (Chapter 9) gives a worked example TLB and walk, and Drepper's memory paper covers TLB behavior in depth. What you must carry is the *shape*: a TLB hit is ~free, a TLB miss costs a page-table walk whose price is one to a few memory accesses and, in the bad case, is a DRAM-latency event stacked *on top of* whatever the data access itself will cost.


(Answer to the quick check: 12 offset bits, because $2^{12} = 4096 = 4 \text{ KiB}$, so 12 bits address every byte within a page. The page table translates the page-number bits; the offset is copied unchanged into the physical address, since a page and its frame are the same size and a byte keeps its position within the page.)

THE SAME IDEA THREE WAYS: THE TLB IS A CACHE FOR THE PAGE TABLE

1. In words. Translating an address needs a multi-step page-table walk. Rather than walk on every access, the CPU caches recent translations in the TLB; a hit skips the walk, a miss pays for it and caches the result — identical logic to a data cache, applied to translations instead of data.

2. As a picture. Figure 4.1: two parallel lookups. The data cache maps *address* → *line*; the TLB maps *virtual page* → *physical frame*. Same hit/miss structure, different key and value.

3. As a trace. Access page P: TLB miss → walk the 4-level table (up to 4 memory reads) → install P→frame in the TLB. Next 1000 accesses to page P: TLB hits, ~free. The walk's cost is amortized over every access that stays on the page — which is why sequential code barely notices the TLB and scattered code does.



A TLB hit skips the whole lower box. A miss pays for the walk, then caches the result.

Figure 4.1: Address translation. The TLB caches virtual-page→physical-frame mappings; a hit is near-free, a miss triggers a multi-level page-table walk (a dependency chain of memory reads) whose result is then cached.

4.4 Why huge, scattered working sets thrash the TLB

Here is the performance consequence, and it is a new failure mode not covered by Chapter 3's data caches. The TLB holds only a limited number of translations — a hardware-fixed count, small relative to the number of pages a large program can touch. The total span of memory the TLB can map at once is its **reach**: roughly (number of entries) × (page size). With 4 KiB pages, even a few thousand TLB entries reach only a few megabytes of address space at a time. So a program whose hot data spans *hundreds* of megabytes, accessed in a scattered pattern that jumps from page to page, will keep missing the TLB: each jump lands on a page whose translation was long since evicted, forcing a fresh page-table walk. This is **TLB thrashing**, and it is precisely a *capacity miss* (Chapter 3's taxonomy) at the translation level — the working set of *pages* exceeds what the TLB can hold, just as a data working set can exceed the L2.

Notice the two costs now stack. A scattered access over a huge array can miss the *data* cache (pay ~100 ns to fetch the line) *and* miss the *TLB* (pay a page-table walk to even learn where the line is), and in the worst case the walk itself misses cache and goes to DRAM. This is a large part of why truly random access over a multi-hundred-megabyte structure is so much slower than the raw ~100 ns DRAM figure alone would predict — a fact the [Chapter 6](#) capstone lab measures directly, where a random pointer-chase over 512 MiB costs well above the bare DRAM latency, with the surplus owing partly to exactly

this page-walk cost. By contrast, *sequential* access barely touches the TLB: a 4 KiB page holds ~ 512 `long`s, so one translation serves hundreds of consecutive accesses before you cross into the next page — the same "use the whole thing you paid to bring in" principle as the cache line, one level up.

QUICK CHECK

Two loops touch the same 512 MiB array the same number of times: one walks it sequentially, one jumps to random pages. Which one thrashes the TLB, and why does the other one hardly use it at all? (Answer below the huge-pages section.)

4.5 Huge pages, conceptually

If TLB reach is (entries) $\times$ (page size) and you cannot easily add TLB entries (that is fixed hardware), the other lever is the page size. **Huge pages** do exactly that: instead of mapping memory in 4 KiB units, the OS can map it in much larger units — on x86-64, commonly **2 MiB**, and **1 GiB** where supported (both are real, documented x86-64 page sizes; the `pdpe1gb` CPU flag advertises 1 GiB support). One huge-page translation covers 2 MiB instead of 4 KiB, so the *same* number of TLB entries now reaches roughly $512\times$ more memory. A working set that thrashed the TLB with 4 KiB pages may fit entirely within the TLB's reach once mapped with 2 MiB pages, turning a storm of page-table walks into TLB hits. A second, smaller benefit: a huge page's walk descends fewer levels (the translation stops at a higher level of the tree), so even the misses are a bit cheaper.

Huge pages are a real mitigation for real workloads — large in-memory databases, big hash tables, scientific arrays — where TLB misses are a measured bottleneck. They are not free lunch: they can waste memory to internal fragmentation (a 2 MiB page allocated for a few live kilobytes), can be harder for the OS to find contiguously once memory is fragmented, and help *only* when TLB misses were actually a problem. That last clause is the discipline of this whole Part: huge pages are something you reach for *after* a profiler shows page-walk cost is significant, not a switch you flip on faith. This book keeps the treatment conceptual — the mechanics of enabling them (explicit `mmap` flags vs. transparent huge pages) belong to the operating-system Part — but the model is what matters here: *bigger pages, fewer translations, more TLB reach*.

(Answer to the §4.4 quick check: the random-page jumper thrashes the TLB — each jump lands on a different page whose translation has been evicted, forcing a walk. The sequential walker hardly uses the TLB because each 4 KiB page's single translation serves the hundreds of consecutive accesses that fall on that page before it crosses to the next one.)

TRAP: IGNORING TLB COST (THE "IT'S ALL JUST DRAM LATENCY" FALLACY)

Having learned Chapter 3, it is tempting to model every scattered access as a single ~100 ns DRAM miss and stop there. That undercounts. A random access over a large working set can pay *two* independent penalties: the data-cache miss to fetch the line, *and* a TLB miss whose page-table walk is itself one to several memory accesses. When your measured cost per random access sits well above the raw DRAM latency you expected, suspect the TLB — especially if the working set spans far more memory than TLB reach (a few thousand pages × 4 KiB). The tell is a hot, random-access structure of hundreds of megabytes or more. The fixes are the ones this chapter and the last give: prefer sequential or blocked access so one translation serves many accesses (the sequential-vs-random gap is partly a TLB gap); shrink the working set; and, where a profiler confirms page-walk cost, consider huge pages. As always: confirm by measuring page-walk and TLB-miss counters, not by guessing.

SAY IT LIKE A SYSTEMS ENGINEER

Try it first. A teammate says: "Our random lookup into the big in-memory index costs about 160 ns per access, but you told me DRAM is only ~100 ns — where does the extra time come from?" Cover the paragraph below and answer in two or three sentences, using *TLB*, *page-table walk*, and *reach*.

Model answer. "The ~100 ns is just the data fetch. Each random lookup also needs an address translation, and because the index spans far more memory than the TLB can map at once (its reach is only a few megabytes with 4 KiB pages), most lookups miss the TLB and trigger a page-table walk — several dependent memory accesses — *before* the data access even starts. That walk cost stacks on top of the DRAM latency, which is where the surplus comes from. If a profiler confirms it, mapping the index with huge pages would raise TLB reach enough to turn most of those walks back into hits." Naming the TLB and page walk — rather than "memory is just slow" — is the senior register, and it points at a concrete, measurable fix.

4.6 What to carry forward

This chapter added the last piece of the machine's cost model that Part 1 needs. The addresses your program uses are virtual, translated to physical one page at a time through a multi-level page table; the translation is cached in the TLB, so it is near-free when your accesses have page-level locality and expensive — a page-table walk of several dependent memory accesses — when they do not. A huge, scattered working set thrashes the TLB the same way a large data working set overflows the L2, and huge pages widen TLB reach as a conceptual mitigation. The unifying idea, straight from [Chapter 3](#), is that *the TLB is one more cache*, subject to the same hit/miss/working-set logic as every other level — which is why the single habit of accessing memory sequentially or in blocks pays off twice: once in the data cache, once in the TLB.

Part 1 now has the whole machine in view: the CPU that runs your instructions ([Chapter 2](#)), the cache hierarchy that feeds it data ([Chapter 3](#)), and the translation layer that turns your addresses into real memory (this chapter). [Chapter 5](#) spends that knowledge: it is where you deliberately *design* your data's layout — array-of-structs versus struct-of-arrays, hot/cold splitting, alignment — so that both the cache and the TLB work *for* you. Then [Chapter 6](#) puts numbers on all of it in a capstone lab. The posture to keep: when

memory is slow, ask not only "did the data miss the cache?" but also "did the *address* miss the TLB?"

CAN YOU... WITHOUT LOOKING?

- Explain the difference between a **virtual** and a **physical** address, and name the three things virtual memory buys (isolation, a simple layout, OS flexibility).
- Split a virtual address into **page number** and **offset**, say why a 4 KiB page has a 12-bit offset, and explain why page tables are **multi-level**.
- State what a **page-table walk** is, why it is a dependency chain, and roughly how many memory accesses it can take on x86-64 (up to four).
- Explain why the **TLB is a cache for the page table**, what a hit and a miss each cost (in shape, not a specific chip's numbers), and how it maps onto Chapter 3's hit/miss model.
- Define **TLB reach** and explain why a huge, scattered working set **thrashes** the TLB while sequential access barely uses it.
- Explain how **huge pages** raise TLB reach, and name a cost of using them.

RETRIEVAL — CLOSED BOOK

Cover the chapter. Answer from memory; rate confidence first.

1. **R1.** What is the difference between a virtual and a physical address, and why does a general-purpose OS use virtual memory at all? (§4.1)
2. **R2.** Why are page tables multi-level rather than a single flat array, and what is a page-table walk? (§4.2)
3. **R3.** In what precise sense is the TLB "a cache for the page table"? What does a hit cost and what does a miss cost, in shape? (§4.3)
4. **R4.** Define TLB reach and explain TLB thrashing using Chapter 3's miss taxonomy. Which of the three C's is it? (§4.4)
5. **R5.** How do huge pages reduce TLB misses, and when are they *not* worth it? (§4.5)

FURTHER READING

- **Bryant & O'Hallaron, *Computer Systems: A Programmer's Perspective (CS:APP)***, Chapter 9 ("Virtual Memory"). The source for this chapter's address-translation mechanism, multi-level page tables, TLB, and a worked example of a walk. The free CMU 15-213 materials cover it.
- **Arpaci-Dusseau, *Operating Systems: Three Easy Pieces (OSTEP)***, the virtualization part (free full text). Paging, TLBs, and page-table structures from the OS side — the complement to CS:APP's hardware view, and where the operating-system Part of this book will pick up.
- **Drepper, "What Every Programmer Should Know About Memory"** (2007, free). Includes TLB behavior, the cost of misses, and huge pages. Absolute numbers are dated — read it for mechanism.

Exercises

- 4.1** **WARM-UP** A system uses 4 KiB pages. (a) How many bits of a virtual address are the page offset? (b) If you access `a[0]` then `a[1]` (adjacent `longs`), do they need one translation or two? (c) Roughly how many consecutive `longs` can you touch before you cross a page boundary and (possibly) need a new translation?
- 4.2** **WARM-UP** In one sentence each, name the three benefits virtual memory provides, and state which one is the security/isolation boundary between processes.
- 4.3** **WARM-UP** Complete the analogy and justify it in a sentence: a data cache stores recently used ____ so a repeated data access is a cheap hit; the TLB stores recently used ____ so a repeated access to the same ____ skips the ____.
- 4.4** **CORE** Explain why a page-table walk is a *dependency chain* and why that matters for its cost. Tie it explicitly to the pointer-chasing argument from [Chapter 3](#): what do a page-table walk and a linked-list traversal have in common?
- 4.5** **CORE** A hot in-memory hash table of 400 MiB is probed at random keys. Measurements show ~150 ns per probe, noticeably more than the ~100 ns you expected for a single DRAM miss. (a) Give two independent memory-system costs a random probe pays. (b) Explain, using *TLB reach*, why this structure is prone to TLB misses. (c) Name one change that would reduce the translation cost specifically, and one that would reduce the data-miss cost.
- 4.6** **COST** A TLB has E entries and the system uses 4 KiB pages. (a) Write the TLB reach as a formula. (b) For a working set of W bytes accessed randomly by page, under what condition (in terms of W , E , and page size) do you expect heavy TLB thrashing? (c) If you switch to 2 MiB pages without changing E , by what factor does reach change, and how does that move the threshold in (b)?
- 4.7** **FADED** *Worked, with the last step for you.* Explain why sequential access over a 512 MiB array barely stresses the TLB while random-by-page access over the same array thrashes it. Steps 1–3 given; complete step 4.
- Step 1 (given):** Translation happens per page (4 KiB); a page of `longs` holds ~512 elements.
- Step 2 (given):** The TLB reaches only (entries $\times$ 4 KiB) of address space at once — a few megabytes, far less than 512 MiB.
- Step 3 (given):** Sequential access stays on one page for ~512 consecutive accesses before crossing to the next, so one translation is amortized over many accesses.
- Step 4 (you):** Complete the argument for the random case: state how often the random walker needs a *new, uncached* translation, why the TLB's small reach means those translations are usually missing, and therefore why nearly every random access pays a page-table walk. Then state the one number you would measure to confirm the TLB — not just the data cache — is the culprit.

4.8 **MIXED** For each symptom, decide whether the primary cause is most likely a *data-cache capacity miss* (Ch 3), a *TLB miss / page-walk* (Ch 4), a *branch misprediction* (Ch 2), or a *dependency chain / low ILP* (Ch 2), and justify in a line: (a) a random-access loop over 800 MiB is slower than ~100 ns/access predicts, and enabling 2 MiB pages speeds it up; (b) a loop over 64 MiB re-read many times is much slower than one over 1 MiB re-read the same total number of times; (c) a sum over a linked list is far slower than over an array of the same length; (d) a tight loop with a data-dependent `if` speeds up sharply after the input is sorted.

4.9 **MIXED** A colleague proposes enabling huge pages everywhere "to make memory faster." Without running anything, give (a) the specific condition under which huge pages actually help, (b) two costs or downsides of huge pages, and (c) the one measurement you would take first to decide whether they are worth it for a given workload. Note which chapter's discipline this reflects.

4.10 **STRETCH** The chapter says the TLB "is a capacity miss at the translation level." (a) Map each of Chapter 3's three C's onto a translation-level analogue if one exists (compulsory, capacity, conflict), or say why it does not. (b) Explain why huge pages address the *capacity* analogue specifically. (c) Argue why increasing the base page size is not a free win at the system level (think of programs with small, sparse memory footprints).

4.11 **LAB** On your own machine, record the base page size (`getconf PAGE_SIZE`) and confirm your architecture's page-table depth and huge-page sizes from its documentation (do *not* guess — cite what you find). Then write a program that sums the same large array (comfortably larger than your last-level cache, e.g. 512 MiB) two ways: strictly sequentially, and in a random-by-page order (jump to a random page, read a few values, jump again). Time both and report *your* ns-per-access for each. If your platform exposes TLB-miss or page-walk counters (e.g. `perf stat -e dTLB-load-misses` on Linux), record them for both runs and check that the random run has far more. Report only numbers you actually measured, and note anything that surprised you.

Chapter 5 — Mechanical Sympathy & Data-Oriented Design

Before you start: Chapter 3 (cache lines, locality, the three C's), Chapter 4 (the TLB), Chapter 1 (mechanical sympathy as this volume's thesis). · **Pairs with:** Chapter 6 (the capstone lab) and the performance-engineering Part. This is where the last four chapters stop being knowledge *about* the machine and become a way of *designing your data*.

BEFORE YOU READ ON

From memory, reaching back: (1) Chapter 3 — memory moves in 64-byte lines, so the cost of an access is really "per line fetched." If a loop uses only 8 of the 64 bytes in each line it brings in, roughly what fraction of the memory bandwidth it pays for is wasted? (2) Chapter 1 named this volume's thesis after a driving metaphor — what is the two-word term, and who is it attributed to? (3) *Predict*: you have 16 million "particles," each a struct with 8 fields, and you need to sum just *one* field across all of them. Does it matter for speed whether you store them as one array of whole structs, or as eight separate arrays — one per field? By the end you will have a measured answer.

Chapters 1 through 4 were reconnaissance: the CPU, the cache hierarchy, and the translation layer, each with its own cost model that the RAM abstraction hides. This chapter spends that reconnaissance. The term **mechanical sympathy** — popularized in software by Martin Thompson, and introduced back in Chapter 1 — names the posture: understand enough about how the machine works to get the most out of it. **Data-oriented design** is that posture applied to one specific and enormously consequential decision — *how you lay your data out in memory* — and its central claim, which the rest of this chapter defends, is blunt: *the fastest code respects the memory system*. You do not earn an order of magnitude here by a cleverer algorithm; you earn it by arranging the same data so the cache lines you fetch are full of bytes you actually use and the pages you touch are few. This chapter turns Chapter 3's rules into design choices — array-of-structs versus struct-of-arrays, splitting hot fields from cold, padding and alignment — and, as promised, measures one of them on real hardware.

5.1 The design question: lay data out for the access pattern

Data-oriented design inverts the habit most of us learned first. Object-oriented instinct groups data by *conceptual entity*: a `Particle` "has" a position, a velocity, a mass, so we put them in one struct and make an array of those. That models the domain cleanly — and it is often exactly wrong for performance, because *the hardware does not care about your conceptual entities; it cares about which bytes travel together in a cache line*. The design question data-oriented design asks first is not "what is the natural object?" but "**what is the hot loop's access pattern, and which bytes does it actually touch together?**" You then lay memory out to match that pattern, so every fetched line and every translated page is spent on bytes the loop needs. Everything in this chapter is a special case of that one move.

5.2 Array-of-structs vs. struct-of-arrays

The canonical layout choice. Suppose each record has several fields and you have many records. There are two natural layouts:

- **Array-of-structs (AoS):** one array, each element a full struct — all of record 0's fields, then all of record 1's, and so on. Fields of the *same record* are adjacent in memory.
- **Struct-of-arrays (SoA):** one array *per field* — every record's x in one contiguous array, every y in another, and so on. The same field across *different records* is adjacent in memory.

Which wins depends entirely on the access pattern, and neither is universally better — this is a genuine trade-off, not a rule. The deciding question is *which bytes your hot loop touches together*:

- If the hot loop **scans one field (or a few) across many records** — sum every x , filter by one attribute, do SIMD over one channel — **SoA wins**. In SoA those values are contiguous, so each 64-byte line is full of the values you want and the traversal is stride-1 with perfect spatial locality (§3.1–3.2). In AoS the same scan strides over the whole struct, touching a few bytes per line and wasting the rest — the column-major mistake of §3.6, in a new costume.
- If the hot loop **uses most fields of one record at a time** — process particle i completely, then particle $i+1$ — **AoS wins**. All of record i 's fields share a line or two, so one fetch brings the whole record; SoA would instead touch one element in each of many separate arrays, scattering the access across many lines and pages.

A measured demonstration: sum one field of 16 million records

Let us measure the SoA-favoring case. Define a record of 8 doubles — position, velocity, mass, charge — which is exactly **64 bytes**, one cache line, in AoS. Make 16 million of them and sum a single field, x , two ways: over the array-of-structs (each step strides 64 bytes to the next record's x , using 8 of every 64 bytes fetched) and over a contiguous per-field array (stride-1, every byte of every line used).

```
struct Particle { double x, y, z, vx, vy, vz, mass, charge; }; /* 64 bytes = 1 line */


/* AoS: sum one field -> stride of 64 bytes, 8 of 64 bytes used per line */
for (size_t i = 0; i < N; i++) sum += aos[i].x;

/* SoA: the x field is its own contiguous array -> stride 1, whole line used */
for (size_t i = 0; i < N; i++) sum += soa_x[i];
```

Built with `gcc -O2` and run for this book on a 12th-generation Intel Core i7 (i7-12700H), over 16 million records (best of 5 timed runs), the AoS scan of one field took roughly **0.12 s** and the SoA scan roughly **0.032 s** — the SoA layout was about **3.7–4× faster** for computing the identical sum, differing only in how the data was laid out. **This exact ratio is machine-dependent** and code-dependent, and you should expect a different number on your hardware. Note it is a clear multiple, but *less* than the naive 8× you might predict from "AoS moves 8× the memory": the SoA loop is not purely bandwidth-

bound — summing into a single accumulator has a dependency chain (§2.3) that caps its speed, so it cannot fully cash in the bandwidth it saves, which compresses the ratio. Remove that ceiling (multiple accumulators, vectorization) and the gap widens toward the bandwidth-bound 8×. The robust, portable result is the *direction and rough magnitude*: laying the scanned field out contiguously is a large, several-fold win, for the same arithmetic — because SoA turns a one-value-per-line scan into a whole-line scan.

AoS: sum one field → 8 of 64 bytes used per line



SoA: x is its own array → every byte of every line is an x



Measured (i7-12700H, 16M records, gcc -O2, best of 5): AoS one-field scan ~0.12 s vs SoA ~0.032 s.
SoA ~3.7–4× faster for the same sum. Machine-dependent.

Figure 5.1: Summing one field. AoS strides over whole records, using a fraction of each fetched line; SoA stores that field contiguously, so every fetched line is full of values the loop uses.

QUICK CHECK

You process each record fully — read all 8 fields of particle i , do the physics, then move to particle $i+1$ — for every record. Which layout, AoS or SoA, is better *here*, and why does the answer flip from the sum-one-field case? (Answer after the next section.)

5.3 Hot/cold field splitting

AoS-vs-SoA is the all-or-nothing version; hot/cold splitting is the surgical one. Often a record has a few fields touched by the hot loop (the **hot** fields) and several touched rarely — logging metadata, debug info, a rarely-read description (the **cold** fields). If they all live in one struct, every cache line the hot loop fetches is padded out with cold bytes it will not use, lowering useful-bytes-per-line exactly as the AoS one-field scan did. **Hot/cold splitting** divides the record: put the hot fields in one compact struct (so more hot records fit per line and per page) and move the cold fields to a separate structure, referenced by index or pointer, touched only on the rare path. The hot loop now streams through dense hot records; the cold data is out of its way. This is the same principle as SoA — group by access frequency instead of by conceptual entity — applied at the granularity of "hot vs. cold" rather than "field by field." The tell that you want it: a struct where a hot loop reads two of twelve fields, and the other ten are dead weight in every line.

(Answer to the §5.2 quick check: **AoS wins** when you use most fields of one record at a time. All of particle i 's fields are adjacent, so one or two line fetches bring the whole record and every byte is used; SoA would scatter that record's fields across eight separate arrays — eight different lines and possibly pages per record. The answer flips

because the access pattern flipped: SoA is contiguous across records for one field; AoS is contiguous across fields for one record. Match the layout to whichever direction your hot loop travels.)

5.4 Padding and alignment

Two low-level facts about how the compiler places your struct in memory, both from the C model (CS:APP, Chapter 3 covers data alignment). First, **alignment**: each type has an alignment requirement — an 8-byte `double` is typically placed at an 8-byte-aligned address — because aligned accesses are efficient and, on some architectures, misaligned ones are slow or forbidden. To honor this inside a struct, the compiler inserts **padding**: unused bytes between fields so each field lands at a legal offset, plus tail padding so an array of the struct keeps every element aligned. The consequence for layout: field *order* changes struct size. A struct declared `{ char a; double b; char c; }` can be substantially larger than its 10 bytes of real data because padding is inserted around the fields; reordering to put the large, strictly-aligned fields first and pack the small ones together can shrink it. A smaller struct means more records per cache line and per page — a direct locality win, for free, from field ordering alone.

Second, alignment interacts with the cache line. A value that straddles a 64-byte line boundary can require touching *two* lines to read one value; keeping hot structures aligned to line boundaries (and sized so they tile the line cleanly, as our 64-byte `Particle` did) avoids that. There is a tension worth naming: padding a struct *up* to a nice size can help alignment and, as we will see next, avoid a concurrency hazard — but padding also makes the struct *bigger*, which hurts the locality of a scan. Like everything in this Part, it is a trade-off resolved by knowing the access pattern and, when it matters, measuring.

THE SAME IDEA THREE WAYS: GROUP DATA BY HOW IT'S ACCESSED

1. In words. The hardware moves and translates memory in fixed blocks (lines, pages). Lay data out so the bytes a hot loop uses together are stored together — by field (SoA), by temperature (hot/cold), and by alignment — so every block you pay for is full of bytes you need.

2. As a picture. Figure 5.1: an AoS line mostly wasted on unused fields vs. an SoA line entirely full of the scanned field.

3. As a measurement. Summing one field of 16 M records: SoA ~0.032 s vs. AoS ~0.12 s on the test machine — ~3.7–4× for the identical arithmetic, purely from layout. Prose says "group by access pattern"; the number says the grouping is worth several-fold.

5.5 A preview: false sharing

One layout hazard belongs to the concurrency Part but is worth planting now, because it is the mirror image of everything above. So far, packing related data into the same cache line has been *good* — it raises useful-bytes-per-line. But when *multiple threads* write to *different* variables that happen to share one cache line, the line ping-pongs between the cores' caches: each write forces the line out of the other core's cache to maintain coherence, even though the threads touch logically independent data. This is **false**

sharing — no real data is shared, but the cache line is, and the hardware's coherence protocol pays for it. The fix runs *opposite* to the locality advice of this chapter: pad or align the per-thread variables so each lands on its *own* cache line, deliberately *spreading* them out. That single-threaded packing helps and multi-threaded sharing hurts is not a contradiction — it is the same fact (a whole line moves as a unit) cutting two ways depending on whether one core or several are writing it. The full treatment, with a measurement, waits for the concurrency Part; for now, file it: *when you add threads, revisit which variables share a line*.

TRAP: AOS WHEN YOU SCAN ONE FIELD (PACKING THE WRONG AXIS)

The recurring layout error: choosing a struct layout to model the domain cleanly, then running a hot loop that touches only one or two fields of each record — so every cache line and every page you fetch is mostly bytes the loop ignores. It is the column-major mistake (§3.6) wearing an object-oriented costume, and it silently caps a scan at a fraction of its potential bandwidth. The tells: a hot loop over a large array of multi-field structs that reads only a couple of fields; a SIMD kernel that wants one channel contiguous but finds it strided; a profiler showing high cache-miss rate on a sequential-looking loop. The fixes are the chapter's: switch the scanned field(s) to **SoA**, or **split hot from cold** so the hot loop streams dense records; and reorder fields to shrink padding. And the inverse trap for later: do *not* reflexively pack everything tight once threads are involved — that invites false sharing (§5.5). As always, confirm the layout change with a measurement on your data and machine — the $\sim 4\times$ here is one machine's number, not a law.

SAY IT LIKE A SYSTEMS ENGINEER

Try it first. In review, someone asks why you "denormalized a clean Particle struct into a bunch of parallel arrays — it's uglier and the objects made more sense." Cover the paragraph and answer in two or three sentences, using *struct-of-arrays*, *cache line*, and *access pattern*.

Model answer. "The hot loop scans a single field across all records, so with the struct layout each 64-byte line we fetched held one useful value and seven fields the loop never reads — we were paying for eight times the bandwidth we used. Splitting that field into its own contiguous array (struct-of- arrays) makes the scan stride-1, so every fetched line is full of values we sum, and it measured several times faster for the identical result. If a different loop needed whole records at once I'd keep them as structs — the layout follows the access pattern, not the domain model." Naming the access pattern and useful-bytes-per-line — rather than "arrays are faster" — is the senior register, and it makes the trade-off explicit so the reviewer can check it against the *other* loops.

5.6 The rule, and where Part 1 lands

Five chapters of machine knowledge reduce, at the design level, to one instruction: **lay your data out for the way the hot loop touches it**. Memory moves in cache lines and is translated in pages (Chapters 3–4), so the layout that keeps each fetched line and each touched page full of bytes the loop actually uses wins — array-of-structs when you use whole records, struct-of- arrays when you scan a field, hot/cold splitting when a record is mostly cold, tight field ordering to shrink padding, and (later, under threads) deliberate spreading to avoid false sharing. None of this changes an algorithm's big-O; all of it can change its speed by a large factor, as the AoS/SoA measurement showed. That is why "the

fastest code respects the memory system" is not a slogan but a design discipline — the practical cash value of mechanical sympathy.

Part 1 is almost complete. You now have the machine (CPU, caches, TLB) and the design posture (data-oriented layout) to exploit it. [Chapter 6](#) closes the Part with the skill that ties it together: **back-of-the-envelope cost estimation** — using the orders of magnitude from Chapters 1–5 to *predict* whether a workload is compute-, memory-, or I/O-bound, then measuring to confirm. It is where "respect the memory system" becomes a number you can estimate before you write the code — and the bridge into Part 2, where you take manual control of that memory in C.

CAN YOU... WITHOUT LOOKING?

- State the central question of **data-oriented design** (what does the hot loop touch together?) and connect **mechanical sympathy** (Martin Thompson) to it.
- Define **AoS** and **SoA** and say which wins for (a) scanning one field across records and (b) using whole records one at a time — and *why* each.
- Explain **hot/cold splitting** and the tell that you want it.
- Explain **padding and alignment**: why field order changes struct size, and why a smaller struct improves locality.
- Describe **false sharing** in one sentence and why its fix (spread data out) is the *opposite* of single-threaded packing.
- Recall the AoS/SoA **measured result** (~3.7–4× on the test machine), *why* it was less than the naive 8×, and why the exact number is machine-dependent.

RETRIEVAL — CLOSED BOOK

Cover the chapter. Answer from memory; rate confidence first.

1. **R1.** What single question does data-oriented design ask before choosing a layout, and why is "what is the natural object?" the wrong one for performance? (§5.1)
2. **R2.** Define AoS and SoA and give the access pattern that favors each. (§5.2)
3. **R3.** What is hot/cold field splitting, and how does it raise useful-bytes-per-line? (§5.3)
4. **R4.** Why does reordering a struct's fields change its size, and why can that be a locality win? (§5.4)
5. **R5.** What is false sharing, and why is its fix the opposite of the packing advice in the rest of the chapter? (§5.5)

FURTHER READING

- **Drepper, "What Every Programmer Should Know About Memory"** (2007, free). The deepest free treatment of why layout and access pattern dominate memory performance; its sections on data structures and access patterns are the direct background for this chapter. Read for mechanism; its absolute numbers are dated.
- **MIT 6.172, *Performance Engineering of Software Systems*** (free OCW). Lectures on cache-efficient data structures, data layout, and measuring the effect — the applied continuation of this chapter, and where the performance-engineering Part picks up.
- **Bryant & O'Hallaron, *Computer Systems: A Programmer's Perspective (CS:APP)***, Chapter 3 (data alignment and struct layout) and Chapter 6 (locality). The reference for how the compiler pads and aligns structs, and why alignment matters.

Exercises

- 5.1** **WARM-UP** Define array-of-structs and struct-of-arrays in one sentence each, and state which stores "the same field across different records" contiguously.
- 5.2** **WARM-UP** A record is 8 doubles (64 bytes). You sum one field across many records in AoS. (a) How many bytes of each 64-byte line does the sum use? (b) Roughly how many times more memory does the AoS one-field scan move than an SoA scan of the same field? (c) Which kind of locality (§3.2) does SoA restore?
- 5.3** **WARM-UP** Give the tell — in one sentence — that a struct is a candidate for hot/cold field splitting.
- 5.4** **CORE** For each workload, choose AoS or SoA and justify in terms of the access pattern and cache lines: (a) a physics step that reads and writes all fields of one particle before moving to the next; (b) a query that sums one numeric column over ten million rows; (c) a SIMD kernel that multiplies every record's `x` by a scalar; (d) rendering that, per object, reads its position, color, and size together.
- 5.5** **CORE** A struct is declared `struct S { char flag; double value; int id; };`. (a) Explain why its size is larger than $1+8+4 = 13$ bytes. (b) Reorder the fields to minimize padding and explain your ordering. (c) State the locality benefit of the smaller struct when you have an array of ten million of them.
- 5.6** **COST** You scan one 8-byte field of records that are S bytes each, laid out AoS, over a large array. (a) In terms of S and the 64-byte line, how many useful bytes per fetched line does the scan use, and what fraction is that? (b) As S grows, what happens to the AoS/SoA bandwidth ratio for this scan? (c) Why might the *measured* speedup be smaller than this bandwidth ratio predicts (name one reason from Chapter 2 and relate it to why this chapter's demo measured $\sim 4\times$, not $8\times$)?

5.7 **FADED** *Worked, with the last step for you.* Argue which layout to choose for a mixed workload that has *two* hot loops: loop A sums one field over all records; loop B updates all fields of each record in turn. Steps 1–3 given; complete step 4.

Step 1 (given): Loop A (scan one field) favors SoA — contiguous field, whole lines used.

Step 2 (given): Loop B (whole record at a time) favors AoS — one record's fields adjacent, one fetch per record.

Step 3 (given): A single layout cannot be optimal for both; the choice depends on which loop dominates the runtime and by how much.

Step 4 (you): State how you would *decide*: what you would measure to learn which loop dominates, how hot/cold splitting or a hybrid layout might serve both partly, and why "pick the layout that helps the hotter loop, then re-measure" is the right posture rather than reasoning it out on paper. Name the discipline this reflects.

5.8 **MIXED** For each speedup, decide whether the primary mechanism is *SoA / whole-line use* (Ch 5), *hot/cold splitting* (Ch 5), a *TLB / huge-page effect* (Ch 4), a *capacity miss fixed by blocking* (Ch 3), or *ILP / multiple accumulators* (Ch 2), and justify in a line: (a) moving a rarely-read description field out of a hot record; (b) summing a column after switching from row-structs to column arrays; (c) tiling a matrix multiply; (d) enabling 2 MiB pages for a random-access index; (e) summing into four accumulators instead of one.

5.9 **MIXED** Two versions of the same aggregation over ten million records differ by $5\times$ in wall time; both are $O(n)$. Without running anything, list at least three distinct candidate causes drawn from Chapters 2–5 (include at least one layout cause and one from an earlier chapter), then name the single measurement you would take first and why it best narrows the field.

5.10 **STRETCH** This chapter both tells you to *pack* related data into shared lines (SoA, hot/cold, tight padding) and warns that packing per-thread data into a shared line causes *false sharing*. (a) Explain why these are not contradictory — reduce both to the one fact about cache lines that drives them. (b) Give a single scenario where you would deliberately pad a struct *up* to a full cache line, and say what you are trading away. (c) Why does the right choice depend on how many cores write the data?

5.11 **LAB** Reproduce the AoS/SoA demonstration on your own machine. Define a record of several fields (size it to a whole number of your cache line, e.g. 64 bytes), make enough records to comfortably exceed your last-level cache, and sum one field two ways: over the array-of-structs and over a contiguous per-field array. Verify both produce the identical sum, time each (best of several runs), and report *your* ratio. Then try to widen it toward the bandwidth ceiling by removing the accumulator dependency (several accumulators, or let the compiler vectorize) and report what changed. Report only numbers you actually measured; if a loop is optimized away, consume the sum and say what you did.

Chapter 6 — The Cost Model, Applied (Capstone Lab)

Before you start: all of Part 1 — Chapter 1 (the latency hierarchy), Chapter 2 (the CPU), Chapter 3 (caches), Chapter 4 (the TLB), Chapter 5 (data layout). · **Pairs with:** the performance-engineering Part, which turns these estimates into profiled, tool-verified analysis. This chapter is the capstone: you *predict* from the cost model, then *measure*, then *reconcile*.

BEFORE YOU READ ON

From memory, across the whole Part: (1) Chapter 1's orders of magnitude — an L1 hit ~1 ns, an L2 hit a few ns, a DRAM access ~100 ns. About how many times slower is DRAM than L2? (2) Chapter 3 — what does the hardware *prefetcher* do for a sequential access pattern? (3) *Predict, and write it down:* you sum a 1 MiB array (fits in L2) and a 512 MiB array (lives in DRAM), touching the same total number of elements. From the latency numbers alone, how many times slower would you guess the DRAM sum is — 2×? 10×? 100×? Hold your number; this chapter measures it, and the answer is a lesson.

Everything in Part 1 has been building one instrument: a **cost model** — a set of orders of magnitude and mechanisms that lets you estimate, before writing or profiling a line, roughly how fast a piece of code *can* go and what will limit it. This closing chapter uses that instrument for real. First we assemble the estimation toolkit and the single most useful classification in performance work — is this workload **compute-bound**, **memory-bound**, or **I/O-bound**? Then we run the capstone lab: predict the cost of two array sums from the model, measure them on real hardware, and — this is the whole point — *reconcile the surprise* when the measurement defies the naive prediction. The reconciliation is where the model earns its keep, because a cost model you cannot correct against reality is just a table of numbers to misquote.

6.1 Back-of-the-envelope estimation

Back-of-the-envelope estimation means getting the right *order of magnitude* with arithmetic you can do in your head, using the ratios from this Part — not predicting a benchmark to three digits. The toolkit is small, and every number in it is one you already have (treat all as approximate and machine-dependent; the register-through-DRAM figures anchor to CS:APP, the SSD/network ones to the community "latency numbers" list, both introduced in Chapter 1):

- **Latencies** (time to get *one* scattered item): L1 ~1 ns, L2 a few ns, L3 ~10 ns, DRAM ~100 ns, local SSD tens of μ s, network round trip hundreds of μ s and up — multiplicative jumps, ~10-100× per tier.
- **Bandwidths** (throughput for *streaming* many contiguous items): memory delivers on the order of *gigabytes per second* per core when accessed sequentially, because the prefetcher and full-line fetches keep the pipe full. Latency and bandwidth are

different numbers answering different questions — the distinction is the hinge of this whole chapter.

- **Compute:** a core executes on the order of a few instructions per cycle at a few GHz — so $\sim 10^9$ – 10^{10} simple operations per second per core (Chapters 1–2).
- **Sizes that gate tier:** which tier your data lives in is set by working-set size versus cache sizes (measure your own with `lscpu/getconf`). Fits in L2 → L2 speed; spills past L3 → DRAM speed; larger than RAM → SSD/I-O speed.

The method: estimate how many operations and how many bytes the work requires, divide by the relevant rate, and take the *larger* of the compute-time and memory-time estimates as your floor — because the resource in shorter supply is the bottleneck. That single comparison is the next section.

6.2 Compute-bound, memory-bound, or I/O-bound

The most useful question you can ask about a workload is *what resource is it waiting on?* Three broad answers, each pointing at a different toolkit:

- **Compute-bound:** the bottleneck is the CPU's execution units — lots of arithmetic per byte of data, data already in cache. Time scales with operations, not memory traffic. The levers are Chapter 2's: better ILP, fewer/ predictable branches, vectorization, a lower-complexity algorithm.
- **Memory-bound:** the bottleneck is moving data between the CPU and memory — little arithmetic per byte, or a working set that spills to DRAM. Time scales with bytes moved and miss counts, not with the arithmetic. The levers are Chapters 3–5's: locality, layout (AoS/SoA), blocking, and fewer, fuller cache lines and pages.
- **I/O-bound:** the bottleneck is the disk or network — the workload spends its time waiting on tens-of-microseconds-and-up operations. Time scales with I/O operations and their latency; the levers are batching, asynchrony, and reducing round trips (the low-level I/O Part).

The deciding ratio is *work per byte*: an operation that does much arithmetic on each byte it loads (a dense matrix multiply) tends compute-bound; one that does almost nothing per byte (summing an array, copying memory) tends memory-bound; one dominated by waiting on storage or a socket is I/O-bound. Crucially, *the same algorithm can change class with its input size*: a sum whose array fits in cache may be limited by the add pipeline (compute-ish), while the identical sum over an array that spills to DRAM becomes memory-bound. This is why "profile before you optimize" is not a platitude — apply the memory toolkit to a compute-bound loop and you get nothing. Estimate the class first; the estimate tells you which chapter to open.

QUICK CHECK

Classify each as most likely compute-, memory-, or I/O-bound: (a) summing a 500 MiB array once; (b) a dense $N \times N$ matrix multiply with N small enough to fit in cache; (c) reading 10,000 small records from an SSD one at a time. (Answer after the lab.)

6.3 The capstone lab, part 1: predict, measure, reconcile a surprise

Now the payoff. We will estimate, then measure, the cost of summing an array that fits in L2 versus one that spills to DRAM — and confront the gap between the naive prediction and the result.

The prediction (naive, from latency numbers). A 1 MiB array of `long` fits in this machine's per-core L2 (1.25 MiB); a 512 MiB array is far past the 24 MiB L3, so it lives in DRAM. If you reach for the *latency* numbers — L2 a few ns, DRAM ~100 ns — the tempting estimate is that the DRAM sum should be roughly **an order of magnitude or more slower**, maybe 10–100×. Write that prediction down; it is the one most readers make, and it is wrong — instructively.

The measurement. Summing into a `long`, touching the same total element count in both cases (the 1 MiB array swept many times, the 512 MiB array once), built with `gcc -O2` and run for this book on an i7-12700H (best of 5):

```
/* both sum the SAME number of elements; only where the data lives differs */
for (pass ...) for (i < n) sum += a[i]; /* small n: L2-resident; big n: DRAM */
```

Working set	Time (same total elements)	ns / element	Effective GB/s
L2-resident (1 MiB, many passes)	~0.048–0.050 s	~0.71–0.75	~10.7–11.2
DRAM-resident (512 MiB, one pass)	~0.082–0.086 s	~1.22–1.28	~6.3–6.6

The DRAM sum was only about **1.7×** slower than the L2 sum — *not* the 10–100× the latency numbers suggested. (These figures are machine-dependent; the point is the *ratio's* surprise, not the digits.)

The reconciliation — why the prediction was wrong. The naive estimate used the wrong number. The ~100 ns DRAM figure is a *latency* — the time to get *one* scattered item when you must wait for it. But a sequential sum does not wait for one item at a time: the access pattern is stride-1, so the **hardware prefetcher** (Chapter 3) sees it coming and streams lines ahead of the loop, and the CPU has many independent loads in flight. The workload is therefore **bandwidth-bound, not latency-bound** — its speed is set by how fast memory can stream contiguous data, not by the latency of a single miss. And streaming bandwidth from DRAM, while lower than from L2, is within a small factor of it; here the loop achieved ~11 GB/s out of L2 and ~6 GB/s out of DRAM (both also partly capped by the single-accumulator add dependency of §2.3, which is why neither is dramatic). Prefetching converted a would-be 100× latency gap into a ~1.7× bandwidth gap. *That* is the lesson of the capstone: **use latency for scattered access and bandwidth for streaming access** — reach for the wrong one and your estimate is off by orders of magnitude.

6.4 The capstone lab, part 2: when the memory wall really bites


So does the ~100 ns DRAM latency ever show up? Yes — precisely when the prefetcher *cannot* help, i.e. when access is scattered and dependent, so there is no stream to run

ahead of and no independent work to overlap. That is the pointer-chasing pattern of [Chapter 3](#). To expose it, chase a pointer around a random cyclic permutation (each step's address comes from the previous read) over the same two working sets — L2-resident and DRAM-resident — timing an equal number of dependent steps. Measured on the same machine (best of 3):

Working set (random pointer-chase)	ns / dependent step	DRAM/L2 ratio
L2-resident (1 MiB)	~6 ns	~24-25×
DRAM-resident (512 MiB)	~150-158 ns	

Now the latency gap appears in full: ~24-25×, and the DRAM chase costs ~150 ns per step — comfortably in the ~100 ns order of magnitude Chapter 1 gave for a DRAM access, and in fact somewhat *more*. That surplus over the bare ~100 ns is itself a Part-1 lesson: a random chase over 512 MiB also **thrashes the TLB** ([Chapter 4](#)) — each jump lands on a new page whose translation is likely evicted — so part of the ~150 ns is the page-table walk stacked on top of the data miss. Same hardware, same data sizes as §6.3; the *only* change is sequential→random access, and the cost went from a ~1.7× bandwidth gap to a ~25× latency gap. This is the whole Part in one contrast: *where your data lives matters, but how you walk it* decides whether the memory wall is hidden or paid in full.

Same L2 vs DRAM data, two access patterns



Prefetching hides latency for sequential access; a dependent random walk cannot be prefetched, so it pays full DRAM latency. Measured, i7-12700H, gcc -O2. Machine-dependent; the contrast, not the digits, is the point.

Figure 6.1: The same data sizes, two access patterns. Sequential streaming is bandwidth-bound (DRAM only ~1.7× slower than L2); a dependent random chase is latency-bound (~25×), exposing the full memory wall plus TLB cost.

THE SAME IDEA THREE WAYS: LATENCY VS. BANDWIDTH DECIDES THE ESTIMATE

1. In words. DRAM's ~100 ns is a *latency* — the wait for one scattered item. Sequential access is not a series of waits; the prefetcher streams data, so it runs at memory *bandwidth*, within a small factor of cache. Predict streaming with bandwidth, scattered access with latency.

2. As a picture. Figure 6.1: a small sequential gap (bandwidth) beside a large random gap (latency) over the identical data sizes.

3. As measurements. Sequential sum: DRAM ~1.7× L2. Random chase: DRAM ~25× L2, ~150 ns/step. The model predicted 100× from latency; the streaming case corrected it to ~1.7×; the scattered case confirmed the latency really is there when nothing hides it.

(Answer to the §6.2 quick check: (a) **memory-bound** — one add per 8 bytes, streams from DRAM; (b) **compute-bound** — much arithmetic per byte, data resident in cache; (c) **I/O-bound** — dominated by per-record SSD latency, and batching/async would help.)

TRAP: GUESSING INSTEAD OF PROFILING (AND ESTIMATING WITH THE WRONG NUMBER)

Two failures this Part has hammered, together at last. The first: *optimizing without measuring* — rewriting a loop's layout or algorithm on a hunch, when a five-minute profile would have shown the time is somewhere else entirely. The second, subtler: *estimating with the wrong model* — as in §6.3, where the ~100 ns latency number predicts a 100× gap for a workload that is actually bandwidth-bound and shows ~1.7×. A back-of-the-envelope estimate is a hypothesis about the bottleneck, not a verdict; it tells you where to *look* and which class you expect, and then you measure to confirm the class and find the real limiter. The tell that you have mis-estimated is exactly the §6.3 surprise: a measurement off from your prediction by more than a small factor — treat that gap as information (usually you used latency where bandwidth applies, or vice versa), reconcile it, and update the model. Guessing the cause and skipping the measurement is how good engineers waste days; the discipline of Part 1 is estimate, measure, reconcile, in that order.

SAY IT LIKE A SYSTEMS ENGINEER

Try it first. A colleague says: "DRAM is 100× slower than cache, so our big sequential scan over the 500 MiB table must be crippled — let's cache it in a fancy structure." Cover the paragraph and answer in three or four sentences, using *latency*, *bandwidth*, *prefetcher*, and *bound*.

Model answer. "That 100× is a *latency* figure — the wait for one scattered access. Our scan is sequential, so the prefetcher streams the data and we run at memory *bandwidth*, which is within a small factor of cache; the scan is bandwidth-bound, not latency-bound, so it's probably only a couple of times slower than if it fit in cache, not 100×. Before we build anything, let's measure the scan's GB/s and confirm it's near memory bandwidth — if it is, a fancy cache buys little and the win, if any, is fewer bytes touched, not lower latency. Where latency *would* bite is if we switched to scattered, pointer-chasing access — then we'd pay the full ~100 ns-plus per step." Naming which resource bounds the workload, and measuring before building, is the senior register.

6.5 Closing Part 1: from the machine to the code that respects it

Six chapters ago the RAM model told a comforting lie: every operation costs one unit, memory is uniform, and speed follows big-O. Part 1 replaced it with the real cost model. You know the machine now — the CPU that overlaps and predicts (Chapter 2), the cache hierarchy that moves memory in lines and rewards locality (Chapter 3), the TLB that translates addresses and punishes scattered pages (Chapter 4) — and you know how to spend that knowledge: lay data out for the access pattern (Chapter 5), and estimate before you optimize, distinguishing latency from bandwidth and compute-bound from memory-bound from I/O-bound (this chapter). The habit that unifies all of it is the one the whole Part drilled: when two pieces of code have the same big-O but different speed, the difference is the machine — and you now know where to look, what to estimate, and how to confirm it by measuring.

This closes Part 1, "The Machine & the Cost Model," and opens the rest of the volume. Part 2 takes you into **C and manual memory**: pointers, the stack and heap, alignment and layout under your direct control, and undefined behavior — programming with no safety net so you understand the net. Everything you estimated here in the abstract, you will there arrange in bytes by hand. The cost model is no longer something the machine hides from you; it is a tool you carry into every line you write.

CAN YOU... WITHOUT LOOKING?

- Do a **back-of-the-envelope** estimate: turn a workload into operations and bytes, divide by the right rate, and take the larger estimate as the floor.
- Distinguish **latency** from **bandwidth** and say which applies to scattered vs. streaming access — and why using the wrong one mispredicts by orders of magnitude.
- Classify a workload as **compute-, memory-, or I/O-bound** from its work-per-byte, and name which chapter's levers each class points to.
- Explain the **\$6.3 surprise**: why the sequential DRAM sum was only $\sim 1.7\times$ slower than the L2 sum, not $100\times$ (prefetcher $\rightarrow$ bandwidth-bound).
- Explain the **\$6.4 contrast**: why the random pointer-chase *did* show $\sim 25\times$ and ~ 150 ns/step, and what part of that surplus is TLB cost.
- State the discipline in order — **estimate, measure, reconcile** — and why a prediction-vs-measurement gap is information, not failure.

RETRIEVAL — CLOSED BOOK

Cover the chapter. Answer from memory; rate confidence first.

1. **R1.** Give the back-of-the-envelope method for deciding whether code is compute- or memory-bound, and the "work per byte" intuition behind it. (§6.1–6.2)
2. **R2.** What is the difference between latency and bandwidth, and which one governs a sequential scan? (§6.1, §6.3)
3. **R3.** Why was the sequential DRAM sum only $\sim 1.7\times$ slower than the L2 sum when the latency numbers suggested far more? (§6.3)
4. **R4.** Why did the random pointer-chase show $\sim 25\times$ and ~ 150 ns/step, and why is that *more* than the bare DRAM latency? (§6.4)
5. **R5.** State the estimate-measure-reconcile discipline and explain why a gap between prediction and measurement is useful rather than embarrassing. (§6.3, warning box)

FURTHER READING

- **MIT 6.172, *Performance Engineering of Software Systems*** (free OCW). The applied home of this chapter: profiling, distinguishing compute- from memory-bound work, and measuring bandwidth vs. latency with real tools — the direct continuation into the performance-engineering Part.
- **Agner Fog, *optimization manuals*** (free, *agner.org*). Microarchitectural detail behind the compute-bound levers — throughput vs. latency of instructions, and how many can be in flight. Dated in specifics; verify against your CPU.
- **Drepper, "What Every Programmer Should Know About Memory"** (2007, free) and **Bryant & O'Hallaron, *CS:APP***, Chapters 6 and 9. The mechanisms this chapter estimates against — cache and memory bandwidth, prefetching, and translation cost. Read Drepper for mechanism, not its dated absolute numbers; the "latency numbers" list (Jeff Dean / Colin Scott) supplies the SSD/network orders of magnitude, strictly as ratios.

Exercises

- 6.1** **WARM-UP** In one sentence each, define latency and bandwidth, and say which one you would use to estimate the time to (a) fetch one random element from a huge array, and (b) stream through a huge array in order.
- 6.2** **WARM-UP** Classify each workload as compute-, memory-, or I/O-bound and give the one-line reason: (a) `memcpy` of 1 GiB; (b) computing SHA-256 over data already in L1; (c) fetching 1,000 rows one-by-one from a remote database.
- 6.3** **WARM-UP** State the estimate-measure-reconcile discipline in order, and say in one sentence why a gap between your prediction and the measurement is useful information.
- 6.4** **CORE** You must sum a 512 MiB array of `long` once. (a) Estimate the time two ways: from DRAM *latency* (~ 100 ns per element) and from streaming *bandwidth* (assume ~ 6 GB/s for this loop). (b) Which estimate is right for a sequential sum, and why? (c) By roughly what factor do the two estimates differ, and what does that factor tell you about the cost of using the wrong model?
- 6.5** **CORE** Explain the §6.3 result to a teammate: the DRAM sum was only $\sim 1.7\times$ slower than the L2 sum. (a) Why did the prefetcher make this bandwidth-bound rather than latency-bound? (b) Why was even the L2 sum not much faster than a naive "few-ns-per-access" guess would allow (name the §2.3 factor)? (c) What single change to the code would let it approach memory bandwidth more closely?
- 6.6** **COST** A kernel does F floating-point operations and moves B bytes. Your core can do $\sim 10^{10}$ ops/s and stream $\sim 10^{10}$ bytes/s. (a) Write the compute-time and memory-time estimates. (b) Give the condition on F/B (work per byte) under which the kernel is compute-bound. (c) A sum does one add per 8 bytes — is it compute- or memory-bound by this test, and does the measurement in §6.3 agree?

6.7 **FADED** *Worked, with the last step for you.* Reconcile why the random pointer-chase over 512 MiB cost ~ 150 ns/step while the sequential sum over the same array cost ~ 1.25 ns/element. Steps 1–3 given; complete step 4.

Step 1 (given): Both touch DRAM-resident data of the same size; the only difference is access pattern (sequential vs. dependent-random).

Step 2 (given): Sequential access is prefetchable and has many independent loads in flight, so it runs at bandwidth (§6.3).

Step 3 (given): The random chase is a dependency chain — each step's address comes from the previous read — so it cannot be prefetched or overlapped and pays full latency per step (§3.5).

Step 4 (you): Complete the reconciliation: explain why ~ 150 ns/step is *more* than the ~ 100 ns bare DRAM latency (name the Chapter 4 mechanism), and state the one measurement that would confirm that extra cost is translation, not data. Then say what this pair of results teaches about when to estimate with latency vs. bandwidth.

6.8 **MIXED** For each observation, name the most likely bottleneck class (compute / memory / I/O) *and* the specific mechanism from Chapters 1–6, and justify in a line: (a) a sequential scan runs at ~ 6 GB/s and adding threads barely helps; (b) a loop with a data-dependent branch speeds up $3\times$ after sorting the input; (c) a random lookup into a 2 GiB index costs ~ 160 ns and huge pages help; (d) a dense matrix multiply that fits in cache is limited by how many multiply-adds per cycle the core issues.

6.9 **MIXED** A report claims a data pipeline stage is "slow because DRAM is $100\times$ slower than cache." Given only that the stage streams sequentially through a 4 GiB file-backed array, (a) say why the $100\times$ framing is probably the wrong model, (b) predict the rough class and the rough slowdown vs. an in-cache version, and (c) name the two measurements you would take to confirm your reframing before anyone optimizes.

6.10 **STRETCH** The chapter measured a sequential DRAM/L2 ratio of $\sim 1.7\times$ and a random one of $\sim 25\times$ on one machine. (a) Explain why the *random* ratio is far more portable across machines than the sequential one (think about what each is bounded by). (b) Give one machine change that would raise the sequential ratio and one that would lower the random ns-per-step. (c) Argue why reporting both ratios — with the access pattern that produced each — is more honest and more useful than a single "DRAM is $N\times$ slower" headline number.

6.11 **LAB Capstone.** On your own machine: (1) record your L2 and L3 sizes and cache-line size. (2) *Predict*, from the cost model and before running, the time to sum an L2-sized array vs. a DRAM-sized array (same total elements), stating whether you expect it bandwidth- or latency-bound and your predicted ratio — write it down first. (3) Measure both (verify identical sums; best of several runs) and report ns/element and effective GB/s. (4) Then measure the *random pointer-chase* version over the same two sizes and report ns/step. (5) *Reconcile*: compare each measurement to your prediction, explain any surprise using latency-vs-bandwidth and (for the random case) the TLB, and if you can, capture cache-miss and TLB-miss counters (e.g. `perf stat`) to support your explanation. Report only numbers you actually measured; if a loop is elided, consume the result and say what you did.

VOL 1 · SYSTEMS PROGRAMMING & PERFORMANCE · PART II

C & Manual Memory

Pointers, layout, the stack/heap, manual allocation, and undefined behavior — programming with no safety net so you understand the net.

Chapter 7 — Pointers & Memory Layout

Before you start: Chapter 4 (virtual addresses; a pointer holds a *virtual* address the hardware translates), Chapter 5 (alignment and struct padding, met there as a layout lever), Chapter 3 (cache lines and locality). · **Opens Part 2** and pairs with Chapter 8 (stack vs. heap) and Chapter 9 (manual allocation). Part 1 told you what memory *costs*; Part 2 hands you the controls. This chapter is the first: what a pointer actually is, how arithmetic on it works, and how your program's bytes are arranged in the address space.

BEFORE YOU READ ON

From memory, reaching back: (1) Chapter 4 — when your C pointer holds a number like 0x7ffd3a08 and you dereference it, is that the physical address of a DRAM cell, or a *virtual* address the hardware translates first? (2) Chapter 5 said a struct like { char a; double b; char c; } is bigger than its 10 bytes of real data — what is the compiler inserting, and why? (3) *Predict:* p is a pointer to an int and q is a pointer to a double. You write p + 1 and q + 1. Do both advance the address by the same number of bytes? If not, what decides each step? By the end of the chapter you will be able to compute the exact byte a pointer expression lands on.

A pointer is the most demystifiable idea in systems programming, because underneath the syntax it is just a number: **a pointer is a variable that holds a memory address.** That is the whole secret. The type decoration — `int *`, `double *`, `struct Particle *` — does not change what is stored (an address); it tells the compiler what *kind* of thing lives at that address, which is exactly what it needs to know to read the right number of bytes when you dereference, and to scale arithmetic correctly when you do math on the pointer. This chapter builds the model from that one fact. We will see what a pointer holds and what dereferencing does; how pointer arithmetic silently multiplies by the pointee's size; why arrays and pointers are close cousins but *not* the same thing; how your program's memory is carved into regions (text, data, bss, heap, stack); and how `sizeof`, alignment, and struct padding decide the exact byte layout — with every size on this page *verified with the compiler* rather than asserted, because in C those sizes are implementation-defined and you must never guess them.

7.1 A pointer is an address

When you declare `int x = 42;`, the compiler arranges for `x` to live at some address in memory — say byte 0x7ffd...a0. When you then write `int *p = &x;`, the **address-of** operator `&` yields that address, and `p` stores it. So `p` holds a number that names a location; `*p` — the **dereference** — means "go to the location named by `p` and read (or write) the object there." Because `p` was declared `int *`, `*p` reads `sizeof(int)` bytes and interprets them as an `int`. Reconnecting to Chapter 4: the number in `p` is a *virtual* address in this process's address space; dereferencing it triggers the translation-and-fetch machinery we already dissected. The pointer is the programmer-visible handle on the very addresses that chapter was translating.

How big is a pointer? Big enough to hold an address on this machine. On the x86-64 machine used throughout this book, compiled with gcc, a data pointer is **8 bytes** (64 bits) — verified, not assumed:

```
printf("%zu\n", sizeof(void*)); /* prints 8 on this machine (x86-64, LP64) */
```

The C standard does *not* fix the size of a pointer — it is implementation-defined, and different pointer types need not even be the same size in general (a subtlety that rarely bites on flat-address-space machines like x86-64, where all object pointers are the same width). What the standard *does* guarantee is narrower and worth stating precisely: `void*` can hold any object pointer and round-trip it back unchanged, and `sizeof(char) == 1` by definition — a `char` is one *byte*, the unit `sizeof` counts in (c)ppreference, "sizeof operator"; the C standard, §6.5.3.4). Everything else about sizes we will measure.

QUICK CHECK

`int *p = &x;`. In one sentence each: what does `p` hold, and what does `*p` do? And why does the *type* `int *` matter when you write `*p`, if the stored value is "just an address"? (Answer after the next section.)

7.2 Pointer arithmetic scales by the pointee size

Here is the feature that trips up everyone once. When you add an integer to a pointer, C does *not* add that integer to the raw address. It scales: `p + n` advances the address by `n * sizeof(*p)` bytes, so that `p + n` points at the *n*-th object of that type past `p`, not the *n*-th byte. This is deliberate — it makes `p + 1` mean "the next element," which is what you almost always want — but it means the byte step depends entirely on the pointer's type. Verified on this machine:

```
int    arr[4]; /* &arr[1] - &arr[0] == 4 bytes (sizeof(int) == 4) */
double darr[4]; /* &darr[1] - &darr[0] == 8 bytes (sizeof(double) == 8) */
```

Running exactly that on the book's machine, the `int` pointer stepped **4 bytes** and the `double` pointer stepped **8 bytes** — each equal to the pointee's `sizeof`, which is what "answers your prediction from the opener": no, `p + 1` and `q + 1` do *not* advance by the same number of bytes; each advances by its own element size. The mirror image is pointer *subtraction*: `q - p` for two pointers into the same array yields the number of *elements* between them (of type `ptrdiff_t`), again dividing out the element size. Indexing is just sugar for this: `arr[i]` is defined as `*(arr + i)`, so the subscript is scaled arithmetic wearing brackets.

Two precision points the C standard is strict about, both worth carrying into the UB chapter (Chapter 9). First, pointer arithmetic is only defined *within* a single array object, including the special "one past the last element" address, which you may compute and compare against but must *not* dereference (C standard §6.5.6, additive operators). Forming or dereferencing a pointer further out than one-past-the-end is undefined behavior — not "wraps around," not "reads the next variable," but genuinely undefined.

Second, the scaling is by the *static* type of the pointer, so a mismatched type computes the wrong address silently. The takeaway is mechanical and exact: *to find where $p + n$ points, multiply n by $\text{sizeof}(*p)$ and add that many bytes to p .*

THE SAME IDEA THREE WAYS: POINTER ARITHMETIC SCALES BY THE POUNTEE SIZE

- 1. In words.** Adding n to a pointer moves it by n elements, i.e. $n \times \text{sizeof}(*p)$ bytes — so the byte step is the element size, and `arr[i]` is exactly `*(arr + i)`.
- 2. As a picture.** Figure 7.1: two rulers over the same bytes. On the `int*` ruler each tick is 4 bytes apart; on the `double*` ruler each tick is 8 bytes apart. "+1" means "next tick," and the tick spacing is `sizeof(*p)`.
- 3. As a trace.** `double *q` at address 1000: $q+1 \rightarrow 1000 + 1 \times 8 = 1008$; $q+3 \rightarrow 1000 + 3 \times 8 = 1024$. An `int *p` at 1000: $p+3 \rightarrow 1000 + 3 \times 4 = 1012$. Same "+3", different bytes, because the element size differs.

A pointer holds an address; "+1" moves by one *element*, not one byte.

`int*` (element = 4 bytes)



`double*` (element = 8 bytes)



Verified on this machine: `int` step = 4 bytes, `double` step = 8 bytes. Sizes are implementation-defined.

Figure 7.1: Pointer arithmetic scales by the pointee's size. "+1" advances one element — 4 bytes for `int*`, 8 for `double*` on this x86-64 machine.

(Answer to the §7.1 quick check: `p` holds the address of `x`; `*p` goes to that address and reads/writes the object there. The type `int *` matters because it tells the compiler how many bytes to read at that address and how to interpret them — four bytes as an `int` here — and it sets the scale for pointer arithmetic.)

7.3 Arrays vs. pointers, and array-to-pointer decay

Arrays and pointers feel interchangeable in C, and beginners often conclude they are the same. They are not, and the difference is a frequent source of bugs. An **array** is a block of contiguous elements; the array's name refers to that whole block, and the block is the storage. A **pointer** is a separate variable that holds an address, which may point at an array's first element but is its own object. The two behave alike in expressions only because of a specific rule: **array-to-pointer decay**. In most expression contexts, an array name is automatically converted ("decays") to a pointer to its first element (C standard §6.3.2.1). That is why you can write `int *p = arr;` and `arr[i]` and pass `arr` to a function expecting `int *` — the array decayed to `&arr[0]`.

The decay does *not* happen in three contexts, and those exceptions expose the difference. The sharpest is `sizeof`. Applied to an array, `sizeof` yields the *whole array's* size in bytes; applied to a pointer, it yields the *pointer's* size. Verified on this machine:

```
int a10[10];
int *p = a10;
sizeof(a10); /* == 40 : 10 ints × 4 bytes each, the whole array */
sizeof(p);   /* == 8  : just the pointer, regardless of what it points to */
```

On the book's machine those printed **40** and **8**. This is the origin of a notorious trap: passing an array to a function makes it decay to a pointer, so *inside* the function `sizeof(param)` is the pointer size (8), not the array size — the length information is lost at the call boundary, which is why C array-handling functions take an explicit length parameter. (The other two non-decay contexts: the operand of unary `&`, where `&arr` has type "pointer to array"; and a string literal used to initialize a `char` array, where it initializes the array rather than setting a pointer.) The mental model: an array is the land; a pointer is a signpost that can point at it. Usually the signpost is what travels around your program, but the two are not the same object — and `sizeof` is where the difference becomes visible.

TRAP: SIZEOF ON A DECAYED ARRAY (MEASURING THE SIGNPOST, NOT THE LAND)

The classic bug: `void f(int a[]) { size_t n = sizeof(a) / sizeof(a[0]); ... }` intending to recover the element count. It cannot work: the parameter `a` is a *pointer* (the array decayed at the call), so `sizeof(a)` is the pointer size — 8 on this machine — and the "count" comes out as `8 / sizeof(int) = 2` regardless of the real length. The `sizeof(array) / sizeof(array[0])` idiom is valid *only* where the true array type is in scope (same function it was declared in, not a decayed parameter). The tell: computing a length from `sizeof` on a function parameter, or on anything you received as a pointer. The fix is the one the standard library uses everywhere — pass the length explicitly alongside the pointer. Confirm your assumption by printing `sizeof` at the point of use if you are ever unsure; on this machine the decayed case prints 8, the in-scope array prints its full byte size (40 for `int[10]`).

7.4 The address-space layout: text, data, bss, heap, stack

Where do all these addresses live? A running process's virtual address space (the illusion of Chapter 4) is conventionally divided into regions, each holding a different kind of data, described in CS:APP (Chapter 9) and OSTEP's address-space chapters. From low addresses upward, the classic picture is:

- **Text** (a.k.a. code): your compiled machine instructions, plus read-only constants like string literals (often a separate `.rodata`). Typically mapped read-only and executable.
- **Data** (`.data`): global and static variables that have a non-zero initial value — the initial bytes are stored in the executable and loaded here.
- **BSS** (`.bss`): global and static variables that are zero-initialized. They take *no* space in the executable file (only a size is recorded); the OS provides zeroed pages at load. (The name is a historical assembler mnemonic.)

- **Heap:** the region for *dynamic* allocation (`malloc`), which grows upward (toward higher addresses) as you request more. This is [Chapter 9](#)'s subject.
- **Stack:** automatic storage for function calls — local variables, saved return addresses, arguments — which grows *downward* (toward lower addresses) as calls nest. This is [Chapter 8](#)'s subject.

You can see the ordering directly. A small program that prints the address of a function (text), an initialized global (data), a zero global (bss), a `malloc`'d block (heap), and a local (stack), built `-no-pie` so the classic fixed layout is visible, printed on this machine:

```
text (&function)      = 0x401196
.data (&global_init)  = 0x404048
.bss (&global_bss)    = 0x40405c
heap (malloc'd block) = 0x351b2a0
stack (&local)        = 0x7ffa35e9a3c
```

The ordering is exactly the textbook one: text < data < bss < heap, all low, and the stack far up near the top of the space — with the heap growing up and the stack growing down toward each other, which is why the address space leaves a large gap between them. **These specific numbers are machine-, OS-, and build-dependent, not universal.** A default position-independent-executable (PIE) build on the same machine printed entirely different, higher base addresses, because **address-space layout randomization (ASLR)** relocates the segments each run as a security measure — the *relative* structure holds, the absolute numbers do not. What to carry: your program's memory is not one undifferentiated pool; it is regions with different purposes, permissions, and growth directions, and knowing which region a pointer points into tells you a lot about its lifetime — the through-line into the next two chapters.

QUICK CHECK

Which region holds each of these: a `malloc`'d buffer, a function's local `int`, a static `int counter = 5;`, and the machine code of `main`? Which two regions grow *toward* each other, and in which directions? (Answer after the next section.)

7.5 `sizeof`, alignment, and struct padding — computed and verified

Now the exact byte layout. Three rules from the C object model (CS:APP Chapter 3; C standard §6.7.2.1) govern how a struct is laid out, and together they let you compute a struct's size by hand — which we will then check against the compiler.

1. **Each type has a size and an alignment.** `sizeof(T)` is how many bytes an object of type `T` occupies; its alignment is an address requirement — an object of type `T` must be placed at an address that is a multiple of `_Alignof(T)`. On x86-64 the alignment of the scalar types equals their size. Verified on this machine (gcc, x86-64 LP64):

type	char	short	int	long	long long	float	double	void*
sizeof (bytes)	1	2	4	8	8	4	8	8
_Alignof	1	2	4	8	8	4	8	8

These sizes are implementation-/ABI-defined, not guaranteed by the language.

The C standard fixes only `sizeof(char) == 1` and minimum ranges (an `int` is *at least* 16 bits, a `long` at least 32); it does not promise `int` is 4 bytes. These are the values for this machine's ABI, measured — on a 16-bit or ILP64 platform they would differ, which is precisely why you print them rather than hard-code them.

1. **Members are laid out in declaration order, at increasing offsets, with the first member at offset 0.** Before each member the compiler inserts as much **padding** as needed to put that member at an offset that is a multiple of *its* alignment.
2. **The struct's own alignment is the largest alignment among its members, and its size is rounded up to a multiple of that alignment** (tail padding), so that in an array of the struct every element stays aligned.

Let us apply this to a concrete struct and predict its size before measuring:

```
struct Mix { char a; int b; char c; double d; };
```

Hand computation, using the verified alignments (char 1, int 4, double 8):

member	align	size	offset	note
a (char)	1	1	0	first member at 0; occupies byte 0
— padding —		3	1-3	to reach the next multiple of 4 for b
b (int)	4	4	4	4 is a multiple of 4; occupies 4-7
c (char)	1	1	8	occupies byte 8
— padding —		7	9-15	to reach the next multiple of 8 for d
d (double)	8	8	16	16 is a multiple of 8; occupies 16-23

Struct alignment = $\max(1, 4, 1, 8) = 8$; the running size is 24, already a multiple of 8, so no tail padding is needed. Predicted `sizeof(struct Mix) = 24`, of which only $1 + 4 + 1 + 8 = 14$ bytes are real data and **10 bytes are padding**. Now the check — running `offsetof` and `sizeof` on this machine printed *exactly*:

```
offsetof(a)=0  offsetof(b)=4  offsetof(c)=8  offsetof(d)=16
sizeof(struct Mix)=24  _Alignof(struct Mix)=8
```

The hand computation matches the compiler byte for byte. And it exposes the layout lever from [Chapter 5](#): field *order* changes size. Reorder the same fields largest-alignment first —

```
struct Packed { double d; int b; char a; char c; };
```

— and the layout becomes `d` at 0, `b` at 8, `a` at 12, `c` at 13, needing only 2 bytes of tail padding to round 14 up to the next multiple of 8. Verified: `sizeof(struct Packed) = 16` on this machine. Same four fields, same data, **24 vs. 16 bytes** — a third smaller — purely from ordering, which (as Chapter 5 argued) means more records per cache line and per page for free. The point of this section is method as much as result: *you can compute a struct's size from the alignment rules, and you should verify it, because the underlying sizes are implementation-defined.*

(Answer to the §7.4 quick check: `malloc'd buffer` → heap; `local int` → stack; `static int counter = 5;` → data (non-zero initializer); machine code of `main` → text. The heap (growing up) and the stack (growing down) grow toward each other.)

7.6 Endianness: which byte comes first

One more layout fact, invisible until you look at individual bytes. When a multi-byte integer sits in memory, in what order are its bytes stored? **Little-endian** stores the least-significant byte at the lowest address; **big-endian** stores the most-significant byte first. The C standard does *not* mandate either — byte order is part of the implementation-defined representation of integer types (C standard §6.2.6) — so portable code must not assume one. x86-64 is **little-endian**, and rather than assert it, we verify it: store the 32-bit value `0x01020304` and read its bytes one at a time.

```
uint32_t x = 0x01020304u;
unsigned char *b = (unsigned char*)&x;
/* print b[0] b[1] b[2] b[3] */
```

On this machine that printed the bytes **04 03 02 01** — the least-significant byte (`0x04`) first, at the lowest address — confirming **little-endian**, as expected for x86-64. On a big-endian machine the same program would print `01 02 03 04`. Endianness matters the moment bytes cross a boundary the CPU does not control: reading a binary file or network packet written by a different machine, or reinterpreting memory through a different-width pointer. Within a single program on one machine you rarely notice it — arithmetic and printing hide it — but at the byte level it is real and, like the sizes above, something you *check* rather than assume. (Network protocols standardize on big-endian, "network byte order," precisely so machines of either kind can interoperate; Beej's Guide covers the `htons/ntohl` conversions the sockets Part will use.)

SAY IT LIKE A SYSTEMS ENGINEER

Try it first. A teammate is debugging and says: "I have a struct { char a; int b; char c; double d; }, that's 14 bytes of data, but sizeof says 24 — is the compiler broken?" Cover the paragraph and answer in two or three sentences, using *alignment*, *padding*, and *field order*.

Model answer. "It's not broken — the compiler inserts padding so each field lands on an address that's a multiple of its alignment: the int needs a 4-aligned offset and the double an 8-aligned one, so there are 3 bytes of padding after a and 7 after c, and the whole struct rounds up to a multiple of 8 — 24 bytes, 10 of them padding. If you reorder the fields largest-alignment-first (double, int, char, char) it packs to 16, same data, because you stop paying to re-align. You can confirm any layout with offsetof and sizeof rather than guessing." Naming alignment and offsets — and showing you can recompute the size — is the senior register, and it turns "the compiler is weird" into a controllable layout decision.

7.7 What to carry forward

A pointer is an address in a variable; dereferencing it reads or writes the object there, using the pointer's type to know how many bytes and how to interpret them. Arithmetic on a pointer scales by the pointee's size, so `p + n` lands `n * sizeof(*p)` bytes along and `arr[i]` is just `*(arr + i)` — and stepping outside an array's bounds is undefined, not merely adventurous. Arrays and pointers are close but distinct: an array decays to a pointer to its first element in most contexts, but `sizeof` (and unary `&`) sees the real array, which is why length is lost at a function boundary. Your program's memory is regioned — text, data, bss, heap, stack — with different lifetimes and growth directions, verified here by printing addresses (and remembering ASLR moves the absolute numbers). And a struct's exact size follows from alignment and padding, which you can compute by hand and *must* verify, because the scalar sizes and byte order underneath are implementation-defined — measured here as `sizeof(struct Mix) = 24` and little-endian on this x86-64 machine.

That is the layer of control Part 1 was pointing at. [Chapter 8](#) takes the two regions with the most interesting lifetimes — the **stack** and the **heap** — and asks where a given object should live, why stack allocation is nearly free while the heap is managed, and how a pointer to a stack local becomes a dangling bug the instant its function returns. [Chapter 9](#) then hands you the heap's controls, `malloc` and `free`, and the precise ways manual memory management goes wrong. Everything in those chapters is pointers into these regions — which is why "a pointer is an address, and addresses live in regions with lifetimes" is the model to keep.

CAN YOU... WITHOUT LOOKING?

- State what a **pointer** holds, what **dereferencing** does, and why the pointer's **type** matters even though the stored value is "just an address."
- Compute where $p + n$ points for a given pointer type (multiply n by `sizeof(*p)`), and say why stepping past one-past-the-end of an array is undefined.
- Explain **array-to-pointer decay**, and why `sizeof` on a function's array *parameter* gives the pointer size, not the array size.
- Name the address-space regions — **text, data, bss, heap, stack** — say what each holds, and which two grow toward each other and in which directions.
- Compute a struct's size by hand from **alignment and padding**, and explain why field *order* changes it — and why you *verify* sizes rather than assume them.
- Define **little- vs. big-endian**, state which x86-64 is, and describe the one-line experiment that checks it.

RETRIEVAL — CLOSED BOOK

Cover the chapter. Answer from memory; rate confidence first.

1. **R1.** What does a pointer hold, and what two things does its *type* tell the compiler? (§7.1–7.2)
2. **R2.** How many bytes does $p + 3$ advance for an `int *p` vs. a `double *p` on this machine, and what general rule gives the answer? (§7.2)
3. **R3.** What is array-to-pointer decay, and what does `sizeof` report for an array vs. for a pointer (or a decayed parameter)? (§7.3)
4. **R4.** Name the five address-space regions low-to-high and say what each holds; which two grow toward each other? (§7.4)
5. **R5.** Give the three struct-layout rules and use them to explain why `{ char; int; char; double; }` is 24 bytes on this machine. (§7.5)

FURTHER READING

- **Bryant & O'Hallaron, *Computer Systems: A Programmer's Perspective (CS:APP)***, Chapter 3 (machine-level data, pointers, arrays, struct alignment) and Chapter 9 (the process address space and its regions). The reference for this chapter's layout and padding rules; the free CMU 15-213 materials cover them.
- **cppreference.com** and the **C standard** — the `sizeof` operator, pointer arithmetic (§6.5.6), array-to-pointer conversion (§6.3.2.1), and object representation (§6.2.6). The authority for what the language *guarantees* vs. leaves implementation-defined; consult these before assuming any size or byte order.
- **Arpaci-Dusseau, *Operating Systems: Three Easy Pieces (OSTEP)***, the address-space chapters (free full text). The OS view of how text/data/heap/stack are arranged and why the stack and heap grow toward each other.
- **Beej's Guide to C** (free). Practical, readable coverage of pointers, arrays, and byte order, including the network-byte-order conversions used in the low-level I/O Part.

Exercises

- 7.1** **WARM-UP** In one sentence each: (a) what does a pointer variable store? (b) what does `*p` do? (c) what does `&x` produce?
- 7.2** **WARM-UP** A `double *q` holds address 2000 (`sizeof(double)==8` here). What address is `q + 4`? What address is `q + 4` if `q` is instead an `int *` (`sizeof(int)==4`)? Show the arithmetic.
- 7.3** **WARM-UP** `int a[10]; int *p = a;` on this machine. State the value of `sizeof(a)` and of `sizeof(p)`, and explain the difference in one sentence.
- 7.4** **CORE** Explain array-to-pointer decay. Then explain precisely why `size_t len(int a[]) { return sizeof(a) / sizeof(a[0]); }` does *not* return the element count, and what the returned value actually is on this machine.
- 7.5** **CORE** For `struct S { char a; short b; char c; int d; };`, compute by hand (using `char` 1, `short` 2, `int` 4, with alignment = size for each) the offset of every member, the total `sizeof`, and how many bytes are padding. Then state how you would verify your answer with the compiler.
- 7.6** **CORE** Reorder the fields of `struct S` from 7.5 to minimize its size, give the new `sizeof`, and explain the ordering rule you applied and why it works.
- 7.7** **COST** You have an array of $N = 10$ million records, and a hot loop that scans them. (a) If the record is `struct Mix` (24 bytes here), how many bytes does the array occupy? (b) If reordering shrinks it to 16 bytes, how many bytes now, and roughly how many *more* records fit in a 64-byte cache line (§3.1) and a 4 KiB page (§4.2)? (c) In one sentence, tie the size reduction to the memory-system cost model of Part 1.
- 7.8** **FADED** *Worked, with the last step for you.* Compute `sizeof(struct T { double d; char a; int b; });` on this machine's ABI. Steps 1–3 given; complete step 4.
Step 1 (given): `d` (double, align 8) at offset 0, occupies 0–7.
Step 2 (given): `a` (char, align 1) at offset 8, occupies byte 8.
Step 3 (given): `b` (int, align 4) needs an offset that is a multiple of 4; the next such offset after byte 8 is 12, so 3 bytes of padding go at 9–11 and `b` occupies 12–15.
Step 4 (you): State the struct's alignment, whether any tail padding is needed to round the size up to a multiple of that alignment, the final `sizeof`, and the total padding bytes. Then name the one function you would call to confirm each member's offset.
- 7.9** **MIXED** For each symptom, decide which chapter's mechanism is the primary cause and justify in a line: (a) `sizeof(param)` inside a function returns 8 no matter how large the caller's array was; (b) a random-access loop over an 800 MiB table is far slower than ~100 ns/access predicts, and huge pages speed it up; (c) a struct is unexpectedly 24 bytes when its fields sum to 14; (d) reading a binary file written on a different architecture yields byte-swapped integers.

7.10 **MIXED** Without running code, decide for each what value (or category of value) is produced and why, drawing on Chapters 4–7: (a) the number stored in a pointer — is it a physical or virtual address? (b) `&arr[5] - &arr[2]` for an `int arr[10]`; (c) whether the absolute address printed by a PIE build is stable across runs, and what causes the answer; (d) whether `sizeof(int)` is guaranteed by the C standard to be 4.

7.11 **STRETCH** The chapter says pointer arithmetic is only defined within a single array object (plus one-past-the-end). (a) Explain why `&arr[10]` for `int arr[10]` is legal to *form* but not to *dereference*. (b) Give a plausible reason the standard makes going further undefined rather than defining it as "wraps" or "reads adjacent memory" (think about the optimizer and about [Chapter 4](#)'s segmented address space). (c) Connect this to the out-of-bounds bug you will meet in [Chapter 9](#).

7.12 **LAB** On your own machine, write and run a program that prints: (a) `sizeof` and `_Alignof` for `char`, `short`, `int`, `long`, `double`, `void*`; (b) the `offsetof` of every member and the `sizeof` of a struct of your choice with at least three differently-aligned fields, and verify it against your hand computation; (c) the byte order of a known 32-bit value to determine your machine's endianness. Report *your* measured numbers and label the machine/ABI; note any that differ from this chapter's x86-64 values and why they might.

Chapter 8 — The Stack & the Heap

Before you start: Chapter 7 (a pointer is an address; the address-space regions, including stack and heap), Chapter 4 (virtual memory), Chapter 5 (locality). · **Pairs with:** Chapter 9 (the heap's controls, malloc/free). Chapter 7 named the regions; this chapter takes the two with real lifetimes — the **stack** and the **heap** — and answers where an object should live, what each costs, and how a pointer that outlives its storage becomes one of C's classic bugs.

BEFORE YOU READ ON

From memory, reaching back: (1) Chapter 7 — of the five address-space regions, which two grow *toward* each other, and in which directions? (2) Chapter 4 — a local variable and a malloc'd block both have *virtual* addresses; does that mean they have the same *lifetime*? Guess what determines when each stops being valid. (3) *Predict:* a function creates a local `int n`, and returns `&n` (its address) to the caller. The caller then reads through that pointer. Do you expect that to reliably give back the value of `n` — and if not, what has gone wrong by the time the caller looks? By the end you will be able to state precisely why this is a bug.

Every object in a C program has a **storage duration** — a lifetime — and getting lifetimes wrong is the source of a whole family of bugs that the previous chapter's "a pointer is just an address" model does not, by itself, warn you about. A pointer stays a valid handle only as long as the object it points at is alive; the moment that object's lifetime ends, the pointer becomes **dangling**, and using it is undefined behavior. This chapter is about the two storage durations you manage most directly. **Automatic** storage — the stack — is created and destroyed for you as functions run, is essentially free to allocate, and has a lifetime tied strictly to a block's execution. **Dynamic** storage — the heap — you allocate and free by hand, is managed by a real allocator with real cost, and lives exactly as long as you keep it alive. We will see how the call stack works and why pushing a frame is nearly free; why the heap is different and must be managed; what stack overflow is; and then the payoff — the lifetime rules that make "return a pointer to a local" a bug, shown correct and broken, and compiled to see what actually happens.

8.1 Automatic storage: the call stack and stack frames

When a function is called, the machine needs somewhere to put that call's local variables, its arguments, the address to return to when it finishes, and some bookkeeping. That somewhere is a **stack frame** (or activation record), pushed onto the **call stack** — the stack region from Chapter 7. The stack is a last-in-first-out structure that mirrors the call structure of the program: calling a function *pushes* a new frame on top; returning *pops* it off. Because calls nest (a function calls a function calls a function) and returns unwind in the reverse order, LIFO is exactly the right discipline, and the region grows *downward* (toward lower addresses) as frames stack up, as Chapter 7's address printout showed the stack sitting high and the CPU pushing toward the heap.

The hardware makes this cheap. A CPU register, the **stack pointer**, holds the address of the current top of the stack. Allocating a frame for a new call is, in essence, *subtracting*

the frame's size from the stack pointer — moving it down to reserve that many bytes — and deallocating on return is adding it back. That is the sense in which stack allocation is "free": reserving space for all of a function's locals is a single arithmetic adjustment of one register, not a search or a bookkeeping operation. There is no per-variable cost and nothing to track; the locals are simply the bytes between the old and new stack-pointer positions. CS:APP (Chapter 3) develops the exact frame layout and the call/return instruction sequence; the shape to hold is: **a call pushes a frame by bumping a pointer, a return pops it by bumping the pointer back, and the whole frame's storage appears and vanishes together.**

Two consequences follow immediately. First, automatic variables have a lifetime tied to their block: they come into existence when execution enters the block (or function) and **cease to exist when execution leaves it** — the frame is popped, and those bytes are no longer reserved for them (C standard §6.2.4, automatic storage duration). Second, that is why locals are so cheap and why you do not free them: the pop reclaims the entire frame at once. The flip side — the danger — is that the storage is reclaimed *whether or not any pointer still refers to it*. Hold onto that; it is the whole of §8.4.

QUICK CHECK

Why is allocating a function's local variables often described as "free," in terms of what the CPU actually does to reserve the space? And what event reclaims all of a function's locals at once? (Answer after the next section.)

8.2 Dynamic storage: the heap, and why it must be managed

The stack's discipline — allocate on call, free on return, strict LIFO — is exactly what makes it cheap, and also what makes it insufficient. Plenty of data does not follow call structure: a buffer whose size you only learn at runtime, a data structure that must outlive the function that built it, a result you want to hand back to a caller. For these you need storage whose lifetime you control independently of the call stack. That is the **heap**, and you obtain it through the allocation functions `malloc` and friends (the subject of [Chapter 9](#)): `malloc(n)` asks for n bytes and returns a pointer to them; that block stays alive until *you* explicitly `free` it, across as many function calls and returns as you like.

This freedom is not free. Unlike the stack — where "allocate" is a pointer bump and "free" is a pointer bump back — the heap must track, out of a large pool, which regions are in use and which are available, in whatever arbitrary order you allocate and free them. That tracking is the job of the **allocator**, a nontrivial piece of software (typically inside the C library) that maintains data structures — free lists, size classes — to satisfy requests and reclaim freed blocks. Because allocations and frees can interleave in any order, the allocator must search for a suitable free block, may split or coalesce blocks, and can leave the pool **fragmented** (free space chopped into pieces too small to be useful). So a heap allocation is a genuine function call doing real work, orders of magnitude more expensive than a stack allocation, and the memory it returns must be explicitly returned or it is **leaked**. Chapter 9 opens the allocator up; for now the contrast is the point.


THE SAME IDEA THREE WAYS: STACK VS. HEAP

1. In words. The stack is automatic storage with LIFO lifetime tied to calls, allocated/freed by bumping the stack pointer — nearly free, but a local dies when its function returns. The heap is manually managed storage with a lifetime *you* control, allocated by a real allocator — flexible, but costly and yours to free.

2. As a picture. Figure 8.1: the stack, a tidy pile of frames growing down, each pushed/popped whole; the heap, a pool of scattered blocks (some in use, some free) that an allocator hands out and reclaims in any order.

3. As a table. Stack: allocate = pointer bump, free = automatic on return, lifetime = the block, cost $\approx$ free, failure = overflow. Heap: allocate = malloc (search/split), free = manual free, lifetime = until you free it, cost = real, failures = leak / use-after-free / double-free (Ch 9).


Stack: LIFO frames, grow down, bump the pointer



grows down

return g() → pop its frame (bump SP back); its locals vanish.

Heap: allocator hands out blocks in any order



malloc searches for a free block; free returns one; gaps = fragments

Stack allocate/free = one register adjustment ($\approx$ free). Heap allocate/free = allocator does real work.

Danger: a stack local's storage is reclaimed on return whether or not a pointer still refers to it (§8.4).

Figure 8.1: The stack pushes/pops whole frames by bumping a pointer (nearly free, LIFO lifetime); the heap is a managed pool the allocator hands out and reclaims in any order (flexible, costly, yours to free).

8.3 Stack overflow

The stack is fast but **bounded**: the OS reserves a finite region for it (commonly on the order of a few megabytes per thread by default on Linux, though the exact size is configurable and platform-dependent — check `ulimit -s` rather than assume a number). If the program keeps pushing frames without returning — most often through unbounded or runaway **recursion**, or by declaring a very large local array — the stack pointer eventually runs past the end of that reserved region. This is **stack overflow**: the access falls on a page with no valid mapping (often a deliberately unmapped *guard page*), the hardware faults, and the OS typically terminates the process (on Linux, a segmentation fault). The two everyday causes are worth naming: recursion with no base case or too deep a recursion for the input, and large automatic arrays (declaring `int buf[100000000];` as a local tries to reserve tens of megabytes on a few-megabyte stack). The fixes match the causes: bound the recursion (or convert it to iteration with an explicit heap-backed structure), and put large or runtime-sized buffers on the *heap*, where the space is far larger and the size can be chosen at runtime. Stack overflow is the stack's characteristic failure, just as leaks and use-after-free (Ch 9) are the heap's.

(Answer to the §8.1 quick check: stack allocation is "free" because reserving space for a call's locals is just subtracting the frame size from the stack-pointer register — one arithmetic adjustment, no per-variable work. Returning from the function — popping the frame by moving the stack pointer back — reclaims all of that function's locals at once.)

8.4 Lifetimes and the dangling-pointer bug: why you can't return a pointer to a local

Now the payoff, and one of the most important bugs in this book. Because a stack frame is reclaimed the instant its function returns (§8.1), **a pointer to a local variable is only valid until that function returns**. After the return, the local's storage is no longer reserved for it — the frame has been popped, the stack pointer moved back, and those bytes are free to be reused by the very next function call. A pointer still holding that address is a **dangling pointer**: it names storage whose lifetime has ended. Reading or writing through it is **undefined behavior** (C standard §6.2.4: the lifetime of an object with automatic storage duration ends when execution leaves its block; using a pointer to an object past its lifetime is UB). This is the answer to the opener's prediction — returning `&n` and dereferencing it later does *not* reliably give back `n`, because `n` no longer exists.

Here is the bug and its correct counterpart, side by side:

```
/* BUG: returns the address of a local; the frame is gone on return. */
int *make_counter_bug(void) {
    int n = 0;
    return &n;          /* dangling: n dies when this function returns */
}

/* CORRECT: the caller owns the storage; the callee just fills it in. */
void fill_counter(int *out) {
    *out = 0;           /* writes into storage the caller keeps alive */
}

int main(void) {
    int *p = make_counter_bug();
    printf("%d\n", *p); /* undefined behavior: *p reads a dead local */

    int n;
    fill_counter(&n);   /* n lives in main's frame, valid throughout */
    printf("%d\n", n);  /* well-defined */
}
```

Two things about the buggy version deserve care. First, **the compiler warns you**. Built with `gcc -Wall` on this machine, `make_counter_bug` produces `warning: function returns address of local variable [-Wreturn-local-addr]`; clang emits the analogous `-Wreturn-stack-address`. This is a bug the toolchain actively flags — heed it. Second, **what happens at runtime is not guaranteed**, and this is the crux of undefined behavior. On this machine, compiling the buggy program (at both `-O0` and `-O2`) and running it produced a **segmentation fault** (the process crashed). But that outcome is *not* part of any contract: UB means the standard imposes no requirements at all, so a different compiler, optimization level, or run could instead print `0`, print stale garbage, or appear to "work"

— which is worse, because the latent bug hides until conditions change. Do *not* reason "it segfaulted, so at least it's detectable": you cannot rely on the crash any more than on the value. The correct fix, shown on the right, is the standard pattern — **the caller owns the storage and passes a pointer in** (`fill_counter(&n)`), so the object outlives the callee's frame; or, when the callee must create the object, allocate it on the *heap* and return that pointer (the caller then frees it — Ch 9). Either way the rule is the same: *never hand out a pointer whose target will die before the pointer does*.

TRAP: RETURNING (OR STORING) A POINTER TO A STACK LOCAL

The recurring lifetime bug: a function returns `&local`, or stores `&local` into a longer-lived structure (a global, an out-parameter's target, a caller's list), and something reads through that pointer after the function has returned. The local's frame is gone, so the pointer dangles and the access is undefined behavior — it may crash, may return stale or overwritten bytes, or may *appear* to work until the next call reuses the frame, which is the dangerous case because it hides in testing. The tells: returning the address of a local or parameter; taking `&` of a local and stashing it somewhere that outlives the function; a pointer that "worked yesterday" now returning garbage after unrelated code changed the call pattern. Heed `-Wreturn-local-addr` / `-Wreturn-stack-address` — the compiler catches the direct return. The fixes: return by value for small objects; take an out-pointer the caller owns; or heap-allocate and transfer ownership (documenting who frees). And do not trust a crash to reveal it — UB is not obligated to crash, so the discipline is to reason about lifetime, not to test-and-see.

QUICK CHECK

Two functions each need to give the caller an `int` the caller can keep: one returns `&local`; the other takes `int *out` and writes `*out`. Which one is correct and why? Where does the correct one's storage actually live? (Answer below the next section.)

8.5 Choosing where an object should live

The stack/heap choice, made well, is mostly mechanical once you ask two questions: *how long must this object live*, and *how big is it (and is that known at compile time)*? Prefer the **stack** when the object's lifetime fits within the function that creates it and its size is known and modest — this is the common case, and it is free, automatically cleaned up, and cache-friendly (a frame's locals are hot and contiguous). Reach for the **heap** when the object must *outlive* the function that creates it (you need to return it, or store it in a longer-lived structure), when its size is only known at runtime or is too large for the stack (§8.3), or when its lifetime genuinely does not nest with call structure (shared data structures, caches, anything with a dynamic graph of owners). The cost of the heap is not only the allocator's time (§8.2) but the obligation it creates: *someone must free it, exactly once, after the last use* — and getting that wrong is the entire subject of the next chapter. This is why the guidance leans toward the stack: it removes the question. The senior habit is to make lifetime an explicit, stated design decision — "this object lives on the stack in `f`," or "the heap owns this and `g` frees it" — rather than reaching for `malloc` reflexively or letting ownership stay implicit.

(Answer to the §8.4 quick check: the `int *out` version is correct — it writes into storage the caller already owns and keeps alive (a local in the caller's frame, or wherever the caller got the pointer), so the object outlives the callee's return. The `&local` version returns a dangling pointer to a frame that is popped on return. The correct one's storage lives in the caller's frame, not the callee's.)

SAY IT LIKE A SYSTEMS ENGINEER

Try it first. In code review, someone defends a helper that does `char *label(int id) { char buf[32]; sprintf(buf, "id-%d", id); return buf; }` with "it compiles and it worked in my test." Cover the paragraph and explain the problem in two or three sentences, using *stack frame*, *lifetime*, and *undefined behavior*.

Model answer. "buf is an automatic array in label's stack frame, so its lifetime ends the instant label returns — the returned pointer dangles, and any read through it is undefined behavior. It 'worked in your test' only by luck: UB isn't required to crash, so the frame's bytes happened to still hold the string until something reused them, which is exactly the bug that hides until the call pattern changes. The compiler will even warn with `-Wreturn-local-addr`. Fix it by having the caller pass in a buffer it owns — `label(id, char *out, size_t n)` — or by returning a heap allocation the caller frees." Naming the frame's lifetime and calling the crash *unreliable* — rather than "it segfaults sometimes" — is the senior register, and it points at the real fix: make ownership explicit.

8.6 What to carry forward

Objects have lifetimes, and a pointer is only valid while its target lives. **Automatic (stack) storage** is created and destroyed with function calls: a frame is pushed by bumping the stack pointer down and popped by bumping it back, which is why stack allocation is nearly free and why a local's lifetime ends exactly when its function returns. **Dynamic (heap) storage** is managed by a real allocator, costs real time, lives until you free it, and exists precisely for objects that must outlive their creating call or whose size or lifetime does not fit the stack's LIFO discipline. The stack's characteristic failure is **overflow** (runaway recursion or huge locals); the heap's failures are the next chapter's. And the through-line bug of this chapter — **returning or storing a pointer to a stack local** — is undefined behavior because the storage is reclaimed on return whether or not a pointer still refers to it; the compiler warns, the crash is *not* guaranteed, and the fix is to make the storage outlive the pointer (caller-owned or heap-owned).

Chapter 9 picks up the heap's controls in full: what `malloc`, `calloc`, `realloc`, and `free` actually promise; what the allocator is doing underneath (free lists, size classes, fragmentation); and the precise catalogue of manual-memory bugs — use-after-free, double-free, leak, buffer overflow, uninitialized read — each stated exactly as the undefined (or indeterminate) behavior it is, with the real tools that catch them. The dangling-pointer idea you just learned on the stack reappears there on the heap as *use-after-free*: the same bug — a pointer outliving its storage — in the region where *you* control the lifetime, and therefore *you* can get it wrong.

CAN YOU... WITHOUT LOOKING?

- Explain what a **stack frame** is, why a call pushing one and a return popping one is nearly **free**, and what a local's **lifetime** is tied to.
- Say why the **heap** must be **managed** (an allocator tracking free space in arbitrary allocate/free order) and how that makes it costlier than the stack.
- Define **stack overflow**, name its two common causes, and give the matching fixes.
- State precisely why **returning a pointer to a local** is a bug — in terms of frame lifetime — and why it is **undefined behavior** rather than "returns garbage."
- Explain why a crash from that bug is *not* guaranteed, and give the two correct patterns (caller-owned storage, or heap-allocate-and-transfer-ownership).
- Decide, for a given object, whether it belongs on the **stack** or the **heap**, from its lifetime and size.

RETRIEVAL — CLOSED BOOK

Cover the chapter. Answer from memory; rate confidence first.

1. **R1.** What is a stack frame, and in what sense is allocating one "free"? What reclaims a function's locals? (§8.1)
2. **R2.** Why must the heap be managed by an allocator, and name two costs that follow. (§8.2)
3. **R3.** What is stack overflow, what are its two usual causes, and how do you fix each? (§8.3)
4. **R4.** Precisely why is returning a pointer to a local undefined behavior, and why is a crash *not* guaranteed? (§8.4)
5. **R5.** Give two questions that decide whether an object belongs on the stack or the heap, and the correct fix for a callee that must hand the caller an object it created. (§8.4–8.5)

FURTHER READING

- **Bryant & O'Hallaron, *Computer Systems: A Programmer's Perspective (CS:APP)***, Chapter 3 (procedure calls, the stack frame, the call/return sequence and stack-pointer mechanics). The reference for how frames are pushed and popped; the free CMU 15-213 materials cover it.
- **Arpaci-Dusseau, *Operating Systems: Three Easy Pieces (OSTEP)***, the address-space and memory-API chapters (free full text). The OS view of the stack and heap regions, stack limits and guard pages, and the allocation interface.
- **cppreference.com** and the **C standard** — storage duration (§6.2.4) and object lifetime. The authority for when an automatic object's lifetime ends and why using a pointer past it is undefined.
- **Beej's Guide to C** (free). Practical treatment of the stack, the heap, and the common lifetime mistakes, with runnable examples.

Exercises

- 8.1** **WARM-UP** In one sentence each: (a) what is a stack frame? (b) what CPU operation allocates one? (c) when does an automatic (local) variable's lifetime end?
- 8.2** **WARM-UP** Name one reason to allocate on the heap rather than the stack, and one cost the heap imposes that the stack does not.
- 8.3** **WARM-UP** Give the two most common causes of stack overflow, and the matching fix for each.
- 8.4** **CORE** Explain, in terms of stack frames and the stack pointer, precisely why this is a bug: `int *f(void) { int x = 7; return &x; }`. State what happens to `x`'s storage on return, why using the returned pointer is undefined behavior, and what warning gcc/clang emit.
- 8.5** **CORE** Rewrite `int *f(void) { int x = 7; return &x; }` two correct ways: (a) so the *caller* owns the storage; (b) so the object lives on the heap. For (b), state who is responsible for freeing it and when.
- 8.6** **CORE** A colleague says "my function returns a pointer to a local and it works fine — I tested it." Explain why "it worked" is *not* evidence the code is correct, using the definition of undefined behavior. What is the danger of a bug that happens to appear to work?
- 8.7** **COST** Compare the cost of one stack allocation vs. one heap allocation. (a) In terms of operations, what does the machine do to allocate a stack frame's worth of locals? (b) What must an allocator potentially do to satisfy one `malloc`? (c) Given (a) and (b), state a rule of thumb for a hot loop that needs a small scratch buffer each iteration, and justify it.
- 8.8** **FADED** *Worked, with the last step for you.* Decide where a function's result should live for: a helper that builds a linked list of unknown length and returns it to its caller. Steps 1–3 given; complete step 4.
- Step 1 (given):** The list must *outlive* the helper that builds it — the caller uses it after the helper returns.
- Step 2 (given):** Its size (number of nodes) is only known at runtime.
- Step 3 (given):** Both facts rule out automatic (stack) storage: the nodes would die on return, and the count is not fixed at compile time.
- Step 4 (you):** State where the nodes must therefore be allocated, who becomes responsible for freeing them and when, and what bug (from the next chapter's preview) results if that responsibility is forgotten. Name the design principle: who owns the memory?
- 8.9** **MIXED** For each symptom, decide which mechanism from Chapters 4–8 is the primary cause and justify in a line: (a) a recursive function crashes with a segfault only on large inputs; (b) a function returns a pointer that gives correct values in unit tests but garbage in production after unrelated code was added; (c) a random-access loop over a huge array is slower than DRAM latency predicts and huge pages help; (d) a struct is larger than the sum of its fields.

8.10 **MIXED** Without running anything, classify each as *well-defined*, *undefined behavior*, or *a performance issue only*, and say why, drawing on Chapters 7–8: (a) taking `&local` and using it *within* the same function; (b) returning `&local` and dereferencing it in the caller; (c) declaring `int big[80000000];` as a local and using it; (d) storing `&local` into a global and reading it after the function returns.

8.11 **STRETCH** The chapter draws a parallel: a dangling pointer to a stack local (this chapter) and a use-after-free on the heap ([Chapter 9](#)) are "the same bug in different regions." (a) State the one property both violate. (b) Explain why the stack version is in some ways *easier* to catch (name the compiler warning) while the heap version generally is not. (c) Argue why "the storage is reclaimed whether or not a pointer still refers to it" is the root of *both*.

8.12 **LAB** On your own machine, compile (with `-Wall`) the buggy `int *f(void){ int x=7; return &x; }` and a main that dereferences the result. (a) Record the exact compiler warning. (b) Run it and report what happens on *your* machine (crash, a value, garbage) — and note that this outcome is not guaranteed by the standard. (c) Then write the caller-owned-storage fix and confirm it is warning-free and correct. Report only what you actually observed, and state why (b) must not be relied upon.

Chapter 9 — Manual Allocation: `malloc/free` and What Goes Wrong

Before you start: Chapter 8 (stack vs. heap; the dangling-pointer bug; lifetimes), Chapter 7 (pointers, arrays, bounds), Chapter 5 (locality). · **Closes the core of Part 2** and sets up **Rust** (Part 5), where the compiler enforces by rule what this chapter makes you do by hand. Here are the heap's controls — `malloc`, `calloc`, `realloc`, `free` — what the allocator does underneath, and the precise catalogue of ways manual memory management goes wrong, each stated as the undefined behavior it is.

BEFORE YOU READ ON

From memory, reaching back: (1) Chapter 8 — why is a pointer to a stack local invalid after its function returns, and what one word names a pointer whose target's lifetime has ended? (2) Chapter 7 — pointer arithmetic and indexing are only defined *within* an array (plus one-past-the-end). What is stepping outside those bounds, in the language's terms? (3) *Predict*: you `free(p)` a heap block, then a moment later read `p[0]` anyway. What do you think happens — a guaranteed crash, the old value still there, or something the standard refuses to define? By the end you will be able to name this bug and say precisely why its behavior is not guaranteed.

The heap gives you storage whose lifetime you control — and hands you, in exchange, the entire responsibility for managing it correctly. No garbage collector will reclaim what you forget; no bounds checker will stop you writing past the end; no runtime will notice you using a block after you freed it. C gives you the raw controls and trusts you completely, which is exactly why this chapter matters: the bugs here are not "the program crashes," they are **undefined behavior**, and understanding that distinction precisely is the difference between a mysterious, occasional failure and a controllable one. We will cover the four allocation functions and what each actually promises; sketch what an allocator does underneath (free lists, size classes, fragmentation — at the level this book needs); then catalogue the canonical memory bugs — use-after-free, double-free, memory leak, buffer overflow, uninitialized read — each stated exactly, with why it is dangerous going beyond "it crashes." Finally we meet the real tools that catch them, and run one — AddressSanitizer — on this machine to see what it reports. This is the chapter that most directly motivates Rust: every rule here that you must hold in your head, Rust's type system checks for you.

9.1 The allocation functions and what they promise

Four functions from `<stdlib.h>` make up the manual-allocation interface. Their *documented* behavior (cppreference; the C standard §7.22.3) is precise, and precision matters because the failure modes live in the gaps:

- **`void *malloc(size_t size)`** — allocates `size` bytes and returns a pointer to the block, suitably aligned for any object type. **The contents are uninitialized** — indeterminate bytes, *not* zeroed. Returns `NULL` on failure (and if `size` is 0, returns

either `NULL` or a unique pointer you may not dereference). You must check the return value before using it.

- **`void *calloc(size_t n, size_t size)`** — allocates space for `n` objects of `size` bytes each, **initialized to all-zero bytes**. Because it takes the count and element size separately, mainstream implementations guard the `n * size` multiplication against overflow (returning `NULL` rather than silently under-allocating) — which is a practical reason to prefer it over `malloc(n * size)` when allocating an array. Use it when you want zeroed memory or are allocating an array.
- **`void *realloc(void *ptr, size_t size)`** — resizes the block `ptr` previously returned to `size` bytes, preserving the contents up to the smaller of old and new sizes. It *may* move the block (returning a new address), so you must assign its result — `p = realloc(p, n)` loses the block on failure; the safe idiom uses a temporary. On failure it returns `NULL` and leaves the *original* block valid and unchanged. `realloc(NULL, size)` behaves like `malloc(size)`.
- **`void free(void *ptr)`** — returns a block previously obtained from `malloc/calloc/realloc` to the allocator. `free(NULL)` is explicitly a **no-op** (always safe). Passing anything else — a pointer not from the allocator, or one already freed — is **undefined behavior**.

The contract, stated as rules you must uphold: allocate, check for `NULL`, use only within the requested size, free exactly once when done, and never touch the block after freeing. Every bug in §9.3 below is a violation of one of these — and the language does not enforce a single one of them for you.

QUICK CHECK

Does `malloc` zero the memory it returns? Does `calloc`? And why is `p = realloc(p, n)` a dangerous idiom — what do you lose if the `realloc` fails? (Answer after the next section.)


9.2 What an allocator does, conceptually

Between your `malloc` call and the memory you get back sits the **allocator** — usually part of the C library — managing a large pool of memory it obtained from the OS. You do not need its internals to use it, but a conceptual picture explains the costs and one of the bugs (fragmentation). The allocator's job is to satisfy variable-sized requests, in arbitrary allocate/free order, out of that pool. To do so it tracks which regions are free, typically via **free lists**: linked lists of available blocks it searches to find one big enough for a request. To make that search fast, modern allocators group free blocks into **size classes** — separate lists for "16-byte blocks," "32-byte blocks," and so on — so a request goes straight to a list of right-sized blocks instead of scanning one giant list. When a block is freed, the allocator may **coalesce** it with adjacent free blocks into a larger one, and it records bookkeeping (typically a small header near each block) so `free` knows the block's size.

Two consequences matter for this book. First, **cost**: because a request may involve searching a list, splitting a block, or (when the pool is exhausted) asking the OS for more memory, a heap allocation is far more expensive than the stack's pointer-bump (Ch 8) —

which is why hot loops avoid per-iteration `malloc`. Second, **fragmentation**: after many allocations and frees of varying sizes, the free space can end up chopped into many small, non-adjacent pieces. You may have plenty of total free memory yet be unable to satisfy a large request because no single free block is big enough — this is *external* fragmentation. (There is also *internal* fragmentation: rounding a request up to a size class wastes the slack.) Fragmentation is not undefined behavior — it is a performance and capacity problem — but it is a real cost of manual allocation, and it is why allocation patterns (many small short-lived blocks vs. few large long-lived ones) affect a program's memory footprint. CS:APP (Chapter 9) develops allocator design in depth, including the free-list and coalescing mechanics; this book keeps the model conceptual.

The allocator manages a pool; free lists track available blocks by size class.



Fragmentation: total free space is large, but scattered — a big request may find no single block that fits.

size-class free lists: [16B] → □ → □ [32B] → □ [64B] → □ → □ → □

malloc: pick a size class, take a free block (maybe split). free: return it, maybe coalesce with neighbors.

Cost follows: allocate/free do real work; the stack's pointer-bump (Ch 8) does not. Fragmentation is a capacity cost, not UB.

Figure 9.1: An allocator satisfies variable-sized requests from a pool using size-classed free lists, splitting and coalescing blocks. Interleaved allocate/free of varying sizes leaves external fragmentation — free space too scattered to satisfy a large request.

(Answer to the §9.1 quick check: `malloc` does **not** zero its memory — the contents are indeterminate; `calloc` **does** zero it. `p = realloc(p, n)` is dangerous because if `realloc` returns `NULL` on failure, you have overwritten your only pointer to the still-valid original block — leaking it; assign to a temporary and check it first.)

9.3 The canonical memory bugs — each stated precisely

Now the catalogue. The critical framing, which this whole book has been building toward: most of these are **undefined behavior (UB)**, a specific term from the C standard (§3.4.3) meaning behavior "for which this International Standard imposes *no requirements*." That is stronger and stranger than "it crashes." Under UB the program may crash, may produce wrong results, may appear to work, may corrupt unrelated data, may do different things on different runs or compilers — and, crucially, the optimizer is permitted to assume UB *never happens*, so code near a latent UB bug can be transformed in surprising ways. It is worth distinguishing three standard categories you will meet:

- **Undefined behavior** (§3.4.3): no requirements at all — anything may happen (e.g. use-after-free, out-of-bounds access, double-free).
- **Unspecified behavior** (§3.4.4): the standard offers two or more possibilities and does not say which you get (e.g. the order function arguments are evaluated) — but each possibility is a valid, non-catastrophic outcome.
- **Implementation-defined behavior** (§3.4.1): unspecified behavior that each implementation must *document* (e.g. `sizeof(int)`, byte order from Ch 7).

The bugs below are firmly in the first category (with uninitialized reads a closely related "indeterminate value" case). Do not read any of the realistic outcomes as guarantees — they are what *tends* to happen, offered only so you recognize the symptom.

Use-after-free. Reading or writing a block after you have freed it. The pointer still holds the old address, but the block's lifetime has ended (it may already be handed to another allocation). This is exactly [Chapter 8](#)'s dangling-pointer bug, now on the heap, and it is **undefined behavior**. Realistically it may read stale data, read data that now belongs to an unrelated allocation, corrupt the allocator's bookkeeping, or crash — and it is a serious security vulnerability class, because an attacker who controls what gets allocated into the freed slot can influence what your dangling pointer reads or writes. This is the answer to the opener: reading `p[0]` after `free(p)` is not a guaranteed crash and not a guaranteed old value — it is UB.

Double-free. Calling `free(p)` twice on the same block (without a reallocation in between). Also **undefined behavior**: the second `free` corrupts the allocator's internal structures (it tries to reclaim a block already on a free list), which can crash immediately, silently corrupt the heap so a *later* unrelated allocation misbehaves, or — again — be exploited. Note `free(NULL)` is *not* a double-free and is always safe; the bug is freeing the same non-null block twice.

Memory leak. Losing the last pointer to a heap block without freeing it, so the memory can never be reclaimed. This one is *not* undefined behavior — the program stays well-defined — but it is a real defect: a long-running program that leaks steadily grows its memory footprint until it exhausts memory and is killed or drags the system into swap. Leaks are the mirror of use-after-free: use-after-free frees too *early* (while a pointer still needs it), a leak frees too *late* (never).

Buffer overflow (out-of-bounds access). Reading or writing outside the bounds of an allocation — `a[4]` on a 4-element array, or writing 20 bytes into a 16-byte block. Recall from [Chapter 7](#) that pointer arithmetic and indexing are only defined within the array (plus one-past-the-end); going past that is **undefined behavior**. An out-of-bounds *write* is especially dangerous: it can overwrite adjacent data, allocator metadata, or (on the stack) a return address — the classic memory-corruption vulnerability behind a large fraction of security exploits. An out-of-bounds *read* can leak adjacent memory (the Heartbleed class of bug).

Uninitialized read. Reading memory before anything has been written to it — a `malloc`'d block (whose contents are indeterminate, §9.1) or an uninitialized local (Ch 8) — and using the value. Using such an **indeterminate value** is at best unspecified and, in general (for objects whose address is never taken, or trap representations), undefined; either way the value is meaningless and the behavior unreliable. The realistic symptom is a bug that changes with optimization level, unrelated code, or run-to-run — because the "value" is whatever bytes happened to be there. The fix is trivial and non-negotiable: initialize before you read (or use `calloc` for zeroed memory).

THE SAME IDEA THREE WAYS: THE MEMORY BUGS ARE LIFETIME/BOUNDS VIOLATIONS

1. In words. Every bug here breaks one of §9.1's rules: touch a block outside its *lifetime* (use-after-free, double-free), outside its *bounds* (buffer overflow), before it is *initialized* (uninitialized read), or never end its lifetime (leak). Most are undefined behavior — no guaranteed outcome.

2. As a picture. A timeline of one block: malloc (born, contents garbage) → initialize → use → free (dead). Reading before "initialize" = uninitialized read; touching after "dead" = use-after-free; freeing "dead" again = double-free; never reaching "dead" = leak; touching left/right of the block's extent = overflow.

3. As UB, not "a crash." Each undefined case may crash, corrupt, leak data, or seem fine — and may differ by compiler, optimization, or run. The tools in §9.4 turn these unreliable, sometimes-invisible faults into loud, located errors.

TRAP: ASSUMING UNDEFINED BEHAVIOR IS DETERMINISTIC ("IT DIDN'T CRASH, SO IT'S FINE")

The most dangerous misconception in manual memory management is treating a memory bug as if it had a reliable outcome — "the use-after-free returns the old value," "the overflow just writes the next byte," "it ran fine, so it's correct." **Undefined behavior means the standard imposes no requirements** (C §3.4.3): the outcome is not defined, may vary with compiler/optimization/run, and the optimizer may assume the bug cannot occur and transform your code accordingly — so a UB bug can even change behavior elsewhere in the function. A program that "works in testing" with a latent use-after-free or overflow is not correct; it is *unexploded*. The tells: relying on a specific value or crash from a bug; "it only fails in release builds" or "only under load" or "only on the other machine" — all classic UB signatures. The discipline: reason about lifetimes and bounds to *prevent* these (§9.1's rules), and use the sanitizers/Valgrind of §9.4 to *detect* them — never conclude "no crash = no bug." This is precisely the guarantee Rust's type system will provide by construction in Part 5.

9.4 Tools that catch these bugs — and one run on this machine

Because these bugs are unreliable by nature, you do not hunt them by staring; you use tools that instrument memory accesses and turn a silent UB into a loud, located report. Two real, widely used ones:

- **AddressSanitizer (ASan)** — a compiler feature (in gcc and clang) enabled with `-fsanitize=address`. It instruments the program to check every memory access against "shadow memory" and surrounds allocations with *redzones*, catching heap and stack buffer overflows, use-after-free, double-free, and (with its LeakSanitizer component) leaks — reporting the bug's type, address, and source location, with a modest runtime slowdown. It detects *addressability* bugs (touching memory you shouldn't); it does *not* catch reads of uninitialized-but-valid memory — that is MemorySanitizer's job (clang), or Valgrind's.
- **Valgrind (Memcheck)** — a runtime instrumentation tool that runs your unmodified binary on a synthetic CPU, checking each access. It detects use-after-free, invalid frees, leaks, out-of-bounds on heap blocks, and — unlike ASan — use of uninitialized

values, at a larger slowdown. (Valgrind was not installed on this machine, so the runs below use ASan.)

To make this concrete, four tiny buggy programs were compiled with `gcc -fsanitize=address -g` and run on this machine. ASan's actual reports (trimmed to the key lines):

```
/* use-after-free: read p[0] after free(p) */
==ERROR: AddressSanitizer: heap-use-after-free ... READ of size 4
    0x... is located 0 bytes inside of 16-byte region [freed by thread T0 here ...]
SUMMARY: AddressSanitizer: heap-use-after-free  uaf.c:6 in main

/* buffer overflow: a[4] = 7 on a 4-int (16-byte) block */
==ERROR: AddressSanitizer: heap-buffer-overflow ... WRITE of size 4
    0x... is located 0 bytes to the right of 16-byte region [0x..10,0x..20)
SUMMARY: AddressSanitizer: heap-buffer-overflow  oob.c:4 in main

/* double-free: free(p); free(p); */
==ERROR: AddressSanitizer: attempting double-free ...
    [freed by thread T0 here ...]

/* memory leak: malloc(1024), never freed */
==ERROR: LeakSanitizer: detected memory leaks
Direct leak of 1024 byte(s) in 1 object(s) allocated from: ... leak.c:3
```

Every bug was caught, named, and pinned to a source line — the whole point of the tools. For contrast, the **uninitialized read** (printing an uninitialized `int`) ran under ASan *without* an error, printing a garbage value — confirming ASan does not cover that class (you would reach for Valgrind or MemorySanitizer, or catch the simple case at compile time: `gcc -Wall` flags it with `-Wuninitialized`). The lesson is twofold: these tools convert probabilistic, sometimes-invisible UB into deterministic, located failures — run your tests under a sanitizer — and no single tool catches everything, so know which catches which. **Outputs above are from this specific machine and toolchain (gcc 11 on x86-64); exact addresses and wording vary by version.**

QUICK CHECK

Which of these does AddressSanitizer catch: use-after-free, heap buffer overflow, double-free, memory leak, uninitialized read? Which one is the odd one out, and what tool would you use for it? (Answer below the communicate box.)

SAY IT LIKE A SYSTEMS ENGINEER

Try it first. A teammate closes a bug report with "I can't reproduce the crash — it runs fine for me, so I think it's a flaky test." The code has a use-after-free. Cover the paragraph and respond in two or three sentences, using *undefined behavior* and naming a tool.

Model answer. "A use-after-free is undefined behavior, so 'runs fine for me' is exactly what we'd expect some of the time — the standard makes no promise about the outcome, and it can depend on the build, the timing, or what got allocated into the freed slot, which is why it's intermittent rather than flaky. Instead of trying to reproduce by luck, let's build the tests with `-fsanitize=address` (or run under Valgrind); ASan will turn the use-after-free into a deterministic, located error pointing at the exact line. 'Can't reproduce' isn't evidence the bug is gone — with UB, no crash doesn't mean no bug." Naming the behavior as UB and reaching for a sanitizer — rather than accepting "flaky" — is the senior register, and it converts an intermittent mystery into a caught, fixable defect.

(Answer to the §9.4 quick check: AddressSanitizer catches use-after-free, heap buffer overflow, double-free, and (via LeakSanitizer) memory leaks. The odd one out is the **uninitialized read** — ASan does not detect it; use Valgrind or MemorySanitizer, or catch simple cases with `gcc -Wall (-Wuninitialized)`.)

9.5 What to carry forward — and the bridge to Rust

Manual allocation is a contract you uphold by hand. `malloc` gives uninitialized bytes (check for `NULL`), `calloc` gives zeroed bytes (and, on mainstream implementations, checks its $n \times \text{size}$ multiply for overflow), `realloc` resizes and may move (assign to a temporary), and `free` returns a block exactly once (`free(NULL)` is a safe no-op). Underneath, the allocator manages a pool with free lists and size classes, splitting and coalescing blocks, at real cost, and can leave the pool fragmented. And the ways it goes wrong form a small, precise catalogue: **use-after-free** and **double-free** (lifetime violations), **buffer overflow** (bounds violation), **uninitialized read** (reading before writing), and **memory leak** (never freeing) — the first four undefined behavior, meaning *no guaranteed outcome*, not "a crash," and the last a defined-but-real defect. AddressSanitizer and Valgrind turn these unreliable faults into located errors, and each covers a different subset — which is why you run your tests under them rather than trusting that "it didn't crash."

Step back and notice the shape of the last three chapters. You learned that a pointer is an address into regioned memory (Ch 7), that objects have lifetimes and a pointer can outlive its target (Ch 8), and that on the heap *you own every lifetime and every bound* — and can violate them, silently, with the language offering no net (Ch 9). Every safe program in C is one where the programmer has, in their head, proven that no pointer is used outside its target's lifetime, no access falls outside its bounds, and every allocation is freed exactly once. That proof obligation is real and it is heavy — the bugs above are what happens when it fails. This is the exact setup for **Rust** (Part 5): its ownership, borrowing, and lifetime system is a way to make the compiler *check that proof* — to enforce, as type rules verified before the program ever runs, the very invariants you just had to maintain by hand. When you reach that Part, the motivation will already be in your

hands: you will have felt the weight of the net that Rust removes by building the guarantees into the language.

CAN YOU... WITHOUT LOOKING?

- State what `malloc`, `calloc`, `realloc`, and `free` each do and promise — including which zeroes, which may move, and why `free(NULL)` is safe.
- Sketch what an **allocator** does (free lists, size classes, split/coalesce) and define **fragmentation** — and say why it is a capacity cost, not UB.
- Name and precisely define the five canonical bugs, and say which are **undefined behavior** and which is not.
- Explain why "undefined behavior" is stronger than "it crashes," and why "no crash" is *not* evidence of correctness.
- Say which bugs **AddressSanitizer** catches, which it does not, and what tool covers the gap.
- Explain, in one sentence, how this chapter motivates **Rust** — the compiler enforcing what you had to prove by hand.

RETRIEVAL — CLOSED BOOK

Cover the chapter. Answer from memory; rate confidence first.

1. **R1.** What does each of `malloc`/`calloc`/`realloc`/`free` promise, and what is the safe `realloc` idiom? (§9.1)
2. **R2.** What does an allocator do underneath, and what is external fragmentation? (§9.2)
3. **R3.** Define use-after-free, double-free, buffer overflow, uninitialized read, and leak — which are UB? (§9.3)
4. **R4.** What precisely does "undefined behavior" mean, and why is "it ran fine" not evidence of correctness? (§9.3)
5. **R5.** Which bugs does `AddressSanitizer` catch and which does it miss, and how does this chapter set up Rust? (§9.4–9.5)

FURTHER READING

- **Bryant & O'Hallaron, *Computer Systems: A Programmer's Perspective* (CS:APP)**, Chapter 9 (dynamic memory allocation: allocator design, free lists, coalescing, fragmentation, and the common memory bugs). The reference for §9.2's allocator model; the free CMU 15-213 materials cover it.
- **cppreference.com** and the **C standard** — `<stdlib.h>` memory functions (§7.22.3), and the definitions of undefined (§3.4.3), unspecified (§3.4.4), and implementation-defined (§3.4.1) behavior. The authority for what each function promises and what the language leaves undefined.
- **The AddressSanitizer and Valgrind documentation** — the tools' own manuals for what each detects and how to run them. Primary sources; verify flags against your compiler/tool version.
- **The Rust Programming Language ("the book") and the Rustonomicon** (free). Read ahead to ownership and borrowing to see how the invariants this chapter enforced by hand become compiler-checked rules — the payoff Part 5 develops.

Exercises

9.1 **WARM-UP** In one sentence each: (a) does `malloc` zero its memory? (b) does `calloc`? (c) is `free(NULL)` safe?

9.2 **WARM-UP** Write the *safe* `realloc` idiom (using a temporary), and explain in one sentence what goes wrong with `p = realloc(p, n);` when the allocation fails.

9.3 **WARM-UP** Match each bug to its one-line definition: use-after-free, double-free, memory leak, buffer overflow, uninitialized read. Then mark which are undefined behavior.

9.4 **CORE** Explain precisely why "the use-after-free returns the old value, so it's harmless" is wrong. Use the C standard's definition of undefined behavior, and give two realistic outcomes other than "returns the old value."

9.5 **CORE** For each snippet, name the bug and state whether it is undefined behavior: (a) `free(p); return p[0];` (b) `int *p = malloc(4*sizeof(int)); p[4] = 0;` (c) `free(p); free(p);` (d) `int *p = malloc(sizeof(int)); printf("%d", *p);` (e) `int *p = malloc(1024); p = NULL;`

9.6 **CORE** Explain the difference between a *memory leak* and a *use-after-free* in terms of freeing "too late" vs. "too early," and state why exactly one of the two is undefined behavior and the other is not.

9.7 **COST** (a) Why is a heap allocation much more expensive than a stack allocation (tie to §9.2 and Ch 8)? (b) What is external fragmentation, and why can it make a large `malloc` fail even when total free memory exceeds the request? (c) Is fragmentation undefined behavior? Classify it.

9.8 **FADED** *Worked, with the last step for you.* Diagnose and fix: `char *dup(const char *s) { char *r = malloc(strlen(s)); strcpy(r, s); return r; }`. Steps 1–3 given; complete step 4.

Step 1 (given): `strcpy` copies the string *including* its terminating `'\0'`, so it writes `strlen(s) + 1` bytes.

Step 2 (given): `malloc(strlen(s))` reserves only `strlen(s)` bytes — one too few.

Step 3 (given): So `strcpy` writes one byte past the end of the block: a heap buffer overflow, which is undefined behavior.

Step 4 (you): Give the corrected allocation size, add the one check the function is missing (what if `malloc` returns `NULL`?), and name the tool from §9.4 that would have caught the original overflow and the line it would point to.

9.9 **MIXED** For each observation, decide whether the cause is a *use-after-free*, a *buffer overflow*, an *uninitialized read*, a *leak*, or a *non-memory performance issue*, and justify in a line: (a) a long-running server's memory grows without bound until the OOM killer stops it; (b) a test prints a different number each run and only in release builds; (c) an out-of-bounds write corrupts an unrelated variable's value; (d) the program crashes inside `free` with heap-corruption after the same pointer is released twice; (e) a random-access loop over a huge array is slow and huge pages help (careful — one of these is not a memory *bug*).

9.10 **MIXED** Without running code, decide for each whether it is *undefined behavior*, a *defined-but-real defect*, or *well-defined and fine*, drawing on Chapters 7–9: (a) `free(NULL)`; (b) allocating in a loop and never freeing; (c) `a[n]` where `a` has exactly `n` elements; (d) reading a `calloc`'d block before writing it; (e) reading a `malloc`'d block before writing it.

9.11 **STRETCH** The chapter says this material "motivates Rust." (a) State the three proof obligations a C programmer carries for heap memory (lifetime, bounds, free-exactly-once). (b) For each of use-after-free, double-free, and leak, name which obligation it violates. (c) In one or two sentences, describe what it would mean for a *compiler* to check these obligations before the program runs — and why that is more reliable than a sanitizer that checks at runtime.

9.12 **LAB** On your own machine (if a sanitizer is available), write four tiny programs that each contain exactly one of: a use-after-free, a heap buffer overflow, a double-free, and a leak. Compile each with `-fsanitize=address -g` and run it. (a) Record the exact error type and source line ASan reports for each. (b) Write a fifth program with an uninitialized read and confirm whether ASan flags it; if not, note what tool would. Report only what you actually observed, and label your compiler/tool versions, since the exact output varies.

Chapter 10 — Undefined Behavior & the Optimizing Compiler

Before you start: Chapter 9 (the memory bugs, most of them undefined behavior; the §3.4.3 definition), Chapter 7 (pointers, bounds, and one-past-the-end), Chapter 5 (what the compiler is allowed to reorder). · **Continues Part 2.** Chapter 9 named undefined behavior and told you it was "stronger than a crash." This chapter shows you *why* — how the optimizing compiler actively reasons from the assumption that your program contains no UB, so a latent bug does not merely misbehave where it sits: it lets the compiler transform code elsewhere in surprising, sometimes dangerous ways. We will compile real examples on this machine and read the assembly.

BEFORE YOU READ ON

From memory, reaching back: (1) Chapter 9 — in the C standard's terms, what does **undefined behavior** mean, and how is it different from *unspecified* and *implementation-defined* behavior? (2) Chapter 7 — indexing a pointer is defined only within an array (plus one-past-the-end for the address, not the dereference); what is stepping outside that? (3) *Predict*: you write `int f(int x) { return x + 1 > x; }`. For *signed int*, is `x + 1 > x` always true? What do you think an optimizing compiler emits for this function — a comparison, or the constant 1? By the end you will be able to read the assembly and say exactly why.

Here is the mental shift this chapter demands. It is tempting to think of C as a "portable assembler": you write operations, the compiler emits the corresponding instructions, and undefined behavior is what happens when those instructions hit a bad state at runtime. That model is wrong, and believing it is the source of a whole category of baffling bugs. The correct model is that **undefined behavior is a promise you make to the compiler** — a precondition it is entitled to assume you never violate — and the optimizer reasons *from* that promise. When you write code that can execute UB, you have not described "what the machine does in the bad case"; you have told the compiler the bad case cannot happen, and it will rewrite your program on that basis. We will make this concrete: the standard's definition and its two milder cousins (§10.1); the single assumption the optimizer runs on (§10.2); a gallery of UB that actually changes the compiled code on this machine, read as assembly (§10.3); and why the language is built this way, plus the tools and flags that defend you (§10.4). This is the intellectual core of why manual memory management (Ch 9) is so heavy, and — again — the exact weight Rust will lift in Part 5.

10.1 Undefined behavior, precisely — and its two milder cousins

Chapter 9 gave the definition; we now lean on it fully. The C standard (C11 §3.4.3) defines **undefined behavior** as behavior "for which this International Standard imposes *no requirements*," and notes that permissible responses range from "ignoring the situation completely with unpredictable results" to "behaving during translation or program execution in a documented manner characteristic of the environment" to "terminating a translation or execution." Read that carefully: "no requirements" and "during translation" together mean the license extends to *compile time* — the compiler, not just the running

program, may do anything when UB is possible. Three standard categories, sharpened from §9.3, are worth keeping straight because only the first grants the compiler this license:

- **Undefined behavior** (§3.4.3): no requirements at all. The program has no meaning on the inputs that trigger it. Examples: signed integer overflow, dereferencing a null or dangling pointer, an out-of-bounds access, a data race, reading an object through a pointer of an incompatible type (a strict-aliasing violation).
- **Unspecified behavior** (§3.4.4): the standard gives two or more valid possibilities and does not say which you get — e.g. the order in which a function's arguments are evaluated. Each possibility is a normal, non-catastrophic outcome; the compiler need not document its choice.
- **Implementation-defined behavior** (§3.4.1): unspecified behavior that the implementation must *document* — e.g. the width of `int`, or the byte order (Ch 7). Portable code avoids depending on it; it is not dangerous, merely non-portable.

The distinction that matters for this chapter: unspecified and implementation-defined behavior still produce a *valid program* with *some* defined outcome — you just may not know which. Undefined behavior produces **no valid program at all** on the triggering input. There is no "the value you happen to get"; the standard has stopped describing your program entirely. Everything in §10.2 follows from that single fact.

QUICK CHECK

Which of these is undefined, which unspecified, which implementation-defined: (a) `sizeof(long)`; (b) the evaluation order of `f() + g()`; (c) signed integer overflow? (Answer after the next section.)

10.2 The optimizer's one assumption: UB never happens

An optimizing compiler transforms your code under the **as-if rule**: it may emit any instructions it likes as long as the observable behavior of a *well-defined* program is preserved. The crucial word is *well-defined*. For an input that triggers undefined behavior, the program has no observable behavior the compiler is required to preserve — so the optimizer is free to assume that input never occurs. It does not check; it does not warn (usually); it simply reasons: "this operation would be UB if condition *C* held, and the programmer has promised never to invoke UB, therefore *C* is false — and I may use that fact to simplify everything downstream (and, under the as-if rule, upstream too)."

Two consequences of that reasoning surprise people. First, **the optimizer propagates the assumption**: one operation that is only UB when `p` is null lets the compiler conclude `p` is non-null *everywhere* that value is used, deleting your later null check as provably-dead code. Second, **the effects can appear to "travel" in time**: because a program that executes UB has undefined behavior for the *whole* execution, an optimizer may hoist, sink, or delete work around the UB in ways that make observable effects *before* the buggy line also vanish or change. This is why "it printed the right thing up to the crash" is not something you can rely on. Neither of these is a compiler bug — each is the as-if

rule applied to a program whose meaning you forfeited by writing the UB. Lattner's LLVM series and Regehr's guide (Further Reading) walk through the same reasoning at length; the sections below reproduce it on this machine.

How the optimizer reasons about a possibly-UB operation:




Figure 10.1: The optimizer treats a possibly-UB operation as a promise that the triggering condition never holds, then simplifies on that basis — here deleting a null check because the pointer was already dereferenced. The pattern is the real Linux kernel bug reproduced in §10.3.

(Answer to the §10.1 quick check: (a) `sizeof(long)` is **implementation-defined** — it varies but must be documented; (b) the evaluation order of `f() + g()` is **unspecified** — one of two valid orders, undocumented; (c) signed integer overflow is **undefined behavior** — no requirements at all, the case the optimizer exploits.)

10.3 A gallery of UB that changes the compiled code — read on this machine

Abstract claims about "the optimizer may assume" become believable only when you watch it happen. Each example below was compiled with gcc 11.4.0 at -O2 on this x86-64 machine and the relevant assembly read directly. **Exact output varies by compiler and version; these are from this specific toolchain.**

Signed integer overflow (the opener). In C, signed integer overflow is undefined (unsigned overflow is defined and wraps modulo 2^n). So in `int f(int x){ return x + 1 > x; }` the compiler may assume `x + 1` never overflows — and if it never overflows, `x + 1 > x` is *always* true. It compiles the whole function to a constant:

```

/* gcc -O2 -S : int f(int x){ return x + 1 > x; } */
f:
    movl    $1, %eax    ; return 1, unconditionally
    ret
  
```

No addition, no comparison — just `return 1`. This is correct per the standard and often useful (it lets loops with an `int` counter be optimized as if they cannot wrap). But if you were *relying* on wraparound to detect overflow — the classic `if (x + 1 < x) /* overflowed */` — the check is deleted and your overflow goes undetected. (Compile the same with `-fwrapv`, which *defines* signed overflow as two's-complement wraparound, and the comparison reappears.)

Null check deleted after a dereference (a real CVE). Dereferencing a null pointer is UB, so once a pointer has been dereferenced the compiler may assume it is non-null from then on. Consider the pattern below — a simplified form of the actual Linux kernel function `tun_chr_poll`:

```
struct T { int *sk; };
int poll(struct T *t) {
    int *sk = t->sk;      /* dereference t */
    if (!t) return -1;     /* null check AFTER the dereference */
    return *sk;
}
```

At -O2, gcc on this machine emits (the `if (!t)` is *gone* — no test of the pointer against zero):

```
poll:
    endbr64
    movq    (%rdi), %rax    ; sk = t->sk    (dereference t)
    movl    (%rax), %eax    ; return *sk
    ret                                ; no `testq %rdi,%rdi; je ...` -- check deleted
```

Compiled with `-fno-delete-null-pointer-checks`, the guard survives (`testq %rdi,%rdi; je .L3` reappears). This is not a contrived teaching example: it is the shape of **CVE-2009-1897** in the Linux kernel's `drivers/net/tun.c` (kernel 2.6.30). The pointer was dereferenced before the null check; gcc, entitled to assume the dereference could not have been of a null pointer, deleted the check; an attacker able to map memory at address zero could then reach the code that should have been guarded — a privilege-escalation hole. The kernel's fix moved the check before the dereference and, as belt-and-braces, the build adopted `-fno-delete-null-pointer-checks`. The lesson is exact: **a null check that comes after the dereference it is meant to guard is not a weak check — to the optimizer it is a statement that the pointer was already known non-null, and it may be deleted.**

The others you already met (Ch 9), now as UB the optimizer trusts. Out-of-bounds access, use-after-free, and uninitialized reads are all UB; the optimizer is likewise permitted to assume they never happen and to transform code accordingly (for instance, assuming an index stays in range, or that two pointers of incompatible types never alias — the **strict-aliasing** rule, §6.5, which lets the compiler keep a value in a register across a write through an `int*` when it "knows" a `float*` cannot touch the same object; `-fno-strict-aliasing` disables it). The through- line: every UB in Chapter 9's catalogue is not just a runtime hazard but a compile-time license.

THE SAME IDEA THREE WAYS: UB IS A PROMISE, NOT A BEHAVIOR

1. In words. Undefined behavior is a precondition you promise the compiler you will never violate. The compiler does not define "what happens in the bad case"; it assumes the bad case is impossible and optimizes on that assumption. So UB has no outcome to reason about — only consequences the optimizer derives from "this can't happen."

2. As a decision. For an operation that is UB exactly when condition *C* holds, the compiler branches: "either *C* is false (fine, proceed) or *C* is true (UB, undefined — I owe nothing)." Both branches let it act as if *C* is false. So it deletes anything that only mattered when *C* was true — like a null check after a dereference.

3. As compiled code. `x + 1 > x` becomes `return 1; if (!t)` after `t->sk` vanishes. The transform is visible in the assembly (§10.3) and correct per the standard — the "surprise" is entirely in the reader's wrong mental model of C as a portable assembler.

TRAP: "IT WORKED AT -O0, SO THE OPTIMIZER BROKE MY CODE"

When a program behaves at `-O0` and misbehaves at `-O2`, the instinct is "the optimizer introduced a bug." Almost always the truth is the reverse: **the undefined behavior was already in your program, and the optimizer merely exercised the license that UB granted it.** At `-O0` the compiler emits naive code that happens to do what you expected; at `-O2` it reasons from "no UB" and the latent bug surfaces. The tells: "only fails in release builds," "only with optimizations on," "adding a `printf` makes it go away" (the print perturbs what the optimizer can prove). The wrong fix is to compile at `-O0` or sprinkle `volatile` until it "works" — that hides UB, it does not remove it, and the next compiler version may re-expose it. The right response is to find and eliminate the undefined behavior: build with `-fsanitize=undefined` (UBSan) and `-Wall -Wextra`, and reason about which operation had no defined meaning. "Works at `-O0`" is not evidence of correctness, exactly as "didn't crash" was not in Chapter 9.

10.4 Why the language is built this way — and how to defend

It is fair to ask why the standard grants this license at all. The answer is a deliberate bargain. Leaving these situations undefined — rather than mandating a checked outcome — is what lets a C compiler generate fast code across wildly different hardware without inserting a bounds check on every array access, an overflow check on every addition, or a null check on every dereference. Signed-overflow UB, for example, lets the compiler promote loop counters, strength-reduce, and assume loops terminate; defined wraparound would forbid some of those. The cost of the bargain is precisely this chapter: the responsibility for never triggering UB shifts entirely to you, and the compiler will not (in general) tell you when you have. Different languages strike the bargain differently — Java defines integer overflow and mandates bounds checks (paying runtime cost for safety); Rust checks bounds and, in debug builds, overflow, and confines the UB-bearing operations to explicitly `unsafe` blocks (Part 5). C chose maximum latitude for the compiler, and hands you the bill.

Given that, your defenses are concrete and layered:

- **Sanitizers.** `-fsanitize=undefined` (UndefinedBehaviorSanitizer, in gcc and clang) instruments the program to detect many UB categories at runtime — signed

overflow, out-of-bounds on known-size arrays, null dereference, misaligned access, and more — reporting the file and line. Combined with AddressSanitizer (Ch 9) for memory errors, it turns much silent UB into loud, located failures. Run your tests under both.

- **Warnings.** `-Wall` `-Wextra` catch a meaningful fraction statically (uninitialized reads, some overflows, format-string mismatches). Free; turn them on and treat them as errors (`-Werror`) where you can.
- **Defining-the-undefined flags.** When you deliberately depend on a behavior the standard leaves undefined, tell the compiler to define it: `-fwrapv` (signed overflow wraps), `-fno-strict-aliasing` (allow type-punning through pointers), `-fno-delete-null-pointer-checks` (keep null checks even after a dereference). These trade some optimization for predictability and are how the Linux kernel, among others, tames the sharpest edges.
- **Write UB-free code.** The durable fix. Check pointers before dereferencing, indices before indexing, and sizes before copying (Ch 11); use unsigned types or explicit overflow-checked arithmetic when you need wraparound or overflow detection; never read before initializing. Every rule from §9.1 is really a rule for staying inside the defined language.

QUICK CHECK

Your code does `if (x + 1 < x)` to detect signed overflow and the check keeps getting optimized away. Name two ways to make the detection actually work — one that changes the flag you compile with, one that changes the type or arithmetic you use. (Answer below the communicate box.)

SAY IT LIKE A SYSTEMS ENGINEER

Try it first. A teammate files a compiler bug: "gcc deleted my null check! Look — I clearly write `if (!t) return -1;` and it's not in the assembly. This is a miscompilation." The code dereferences `t` one line above the check. Cover the paragraph and respond in three or four sentences, using *undefined behavior* and the as-if rule.

Model answer. "It's not a miscompilation — it's the as-if rule doing exactly what the standard permits. You dereference `t` with `t->sk` *before* the check, and dereferencing a null pointer is undefined behavior, so the compiler is entitled to assume `t` was non-null at that point — which makes the later `if (!t)` provably dead code it can delete. This is literally the shape of the Linux kernel bug CVE-2009-1897. The fix isn't to report gcc; it's to move the null check *before* the dereference (and, if we want a belt-and-braces guard, build with `-fno-delete-null-pointer-checks`). A check that comes after the dereference it guards isn't a weak check — to the optimizer it's a promise the pointer was already non-null." Naming the as-if rule and the CVE, rather than blaming the compiler, is the senior register — and it turns a "compiler bug" into a one-line code fix.

(Answer to the §10.4 quick check: (1) compile with `-fwrapv`, which defines signed overflow as two's-complement wraparound so the comparison is no longer optimized away; or (2) do the arithmetic in an **unsigned** type (defined to wrap) or use a checked-arithmetic

*builtin such as `__builtin_add_overflow`, which reports overflow without relying on UB. Detecting signed overflow by **letting it happen** and testing afterward is itself UB — that is why the naive check disappears.)*

10.5 What to carry forward

Undefined behavior is not "what the machine does in the bad case." It is a precondition you promise the compiler you will never violate, and the optimizer reasons from that promise: under the as-if rule it may assume every UB-triggering condition is false and rewrite your program accordingly — deleting a null check after a dereference, folding `x + 1 > x` to the constant `1`, keeping a value in a register across a write it "knows" cannot alias. We watched each of these happen in the assembly on this machine. The categories to keep separate are undefined (no requirements — the compiler's license), unspecified (a valid but unstated choice), and implementation-defined (unspecified but documented); only the first grants the license. The language makes this bargain on purpose — it is what lets C compile to fast code across every architecture without safety checks on every operation — and it hands you the whole cost. Your defenses are UBSan and ASan for detection, `-Wall -Wextra` for the static fraction, the define-the-undefined flags (`-fwrapv`, `-fno-strict-aliasing`, `-fno-delete-null-pointer-checks`) when you must depend on an edge, and — above all — writing code that never triggers UB.

This chapter closes the loop opened in Chapter 9. There we cataloged the manual-memory bugs and called most of them UB; here you saw why that word carries so much weight — UB is what makes those bugs unpredictable, build-dependent, and, as CVE-2009-1897 shows, exploitable. It also sets up the next chapter: the buffer overflow of Chapter 9 is UB, and Chapter 11 follows exactly what an out-of-bounds write can do to a program's control flow — the single most consequential UB in the history of security. And it is the sharpest possible motivation for Rust (Part 5): a language whose safe subset has *no* undefined behavior, confining every UB-bearing operation to code you must mark `unsafe`, so the promises this chapter made you keep by hand are enforced by the type system before the program ever runs.

CAN YOU... WITHOUT LOOKING?

- State the C standard's definition of **undefined behavior** and distinguish it from **unspecified** and **implementation-defined** behavior.
- Explain the **as-if rule** and why it lets the optimizer assume UB-triggering conditions are false — including how an assumption "propagates" and can affect code before the buggy line.
- Walk through why `int f(int x){ return x + 1 > x; }` can compile to return `1`, and why a null check *after* a dereference can be deleted (CVE-2009-1897).
- Explain why "it worked at -00" and "adding a `printf` fixed it" are UB signatures, not evidence of correctness.
- Name four defenses (UBSan, warnings, define-the-undefined flags, UB-free code) and say what each buys.
- Say, in one sentence, why the language leaves these behaviors undefined at all — the optimization/safety bargain — and how Rust strikes it differently.

RETRIEVAL — CLOSED BOOK

Cover the chapter. Answer from memory; rate confidence first.

1. **R1.** Define undefined, unspecified, and implementation-defined behavior, and say which grants the compiler its license. (§10.1)
2. **R2.** State the as-if rule and explain how "the optimizer assumes UB never happens" leads to deleting code. (§10.2)
3. **R3.** Why can $x + 1 > x$ become 1, and how does `-fwrapv` change that? (§10.3–10.4)
4. **R4.** Describe CVE-2009-1897: what pattern caused it, why the check was deleted, and the two-part fix. (§10.3)
5. **R5.** Name the bargain the language makes and four defenses against UB. (§10.4)

FURTHER READING

- **Chris Lattner, "What Every C Programmer Should Know About Undefined Behavior" (LLVM Project Blog, 2011, three parts).** The optimizer-writer's own account of how and why the compiler exploits UB, with worked examples (null checks, signed overflow, strict aliasing). Free.
- **John Regehr, "A Guide to Undefined Behavior in C and C++" (Embedded in Academia, 2010, three parts).** A researcher's catalogue of UB and its interaction with aggressive compilers, including the "UB can time-travel" framing. Free (blog.regehr.org).
- **The C standard** — §3.4.1/3.4.3/3.4.4 (the three behavior categories), §6.5 (aliasing), and the integer-overflow rules. The authority for what is undefined; cppreference.com summarizes it accessibly.
- **UndefinedBehaviorSanitizer (UBSan) documentation** (Clang/GCC). The runtime detector for many UB categories; verify the exact checks and flags against your compiler version.

Exercises

- 10.1** WARM-UP In one sentence each, define *undefined*, *unspecified*, and *implementation-defined* behavior, and state which one grants the optimizer license to assume the situation never happens.
- 10.2** WARM-UP Is *unsigned* integer overflow undefined behavior in C? Is *signed* integer overflow? Answer each and, for the defined one, say what value results.
- 10.3** WARM-UP A program prints correct output at `-00` but wrong output at `-02`. State the single most likely explanation, and name the two build flags you would add to hunt for it.
- 10.4** CORE Explain, step by step, why `int f(int x){ return x + 1 > x; }` can be compiled to return 1. Then explain what breaks if a programmer used `if (x + 1 < x)` to detect overflow, and give a correct way to detect it.

10.5 **CORE** Given `int *sk = t->sk; if (!t) return -1; return *sk;` (a) name the UB, (b) explain why an optimizer may delete the `if (!t)` check, (c) give the one-line code fix, and (d) name the flag that would preserve the check without changing the code. This is a real CVE — name the pattern.

10.6 **CORE** "The compiler has a bug — my code is fine, it just miscompiled." A teammate says this about the §10.5 null-check example. Explain, using the as-if rule, why the compiler is behaving correctly and the code is not.

10.7 **COST** The standard *could* have defined all of these behaviors (mandating bounds checks, overflow checks, null checks). (a) What runtime cost would that impose? (b) Name one specific optimization that signed-overflow UB enables. (c) State the bargain in one sentence: what the programmer gives up and what the compiler gains.

10.8 **FADED** *Worked, with the last step for you.* A teammate reports: "My bounds check `if (i < 0 || i >= n)` works, but after I access `a[i]` *first* for logging and check `i` afterward, the check stopped firing." Diagnose and fix.

Step 1 (given): Accessing `a[i]` with `i` out of range is an out-of-bounds access — undefined behavior.

Step 2 (given): Once `a[i]` has executed, the compiler may assume `i` was in range (else UB), i.e. `0 <= i < n`.

Step 3 (given): Therefore the later `if (i < 0 || i >= n)` is provably false and can be deleted — same shape as the null-check CVE.

Step 4 (you): State the fix (what must come before what), name one sanitizer that would have flagged the original out-of-bounds access at runtime, and say why moving the log line does not merely "reorder output" but changes what the compiler is allowed to prove.

10.9 **MIXED** For each, decide whether it is *undefined*, *unspecified*, or *implementation-defined*, and justify in a line: (a) the value of `sizeof(int)`; (b) the order `a()` and `b()` are called in `a() + b()`; (c) `INT_MAX + 1`; (d) right-shifting a negative signed integer; (e) reading a `malloc'd` block before writing it. (One of these is a §9.3 memory bug; recognizing that is the point.)

10.10 **MIXED** A colleague proposes three "fixes" for an intermittent bug that only appears in release builds: (a) always compile at `-O0`; (b) mark the variables `volatile` until it stops; (c) build the tests with `-fsanitize=undefined,address` and fix whatever it reports. Rank them from least to most effective and say, for each, whether it removes the bug or merely hides it.

10.11 **STRETCH** The chapter says a UB assumption can "propagate" and even affect code *before* the buggy line. (a) Explain why, using the definition of UB as applying to the whole program execution. (b) Give a concrete scenario where deleting a downstream null check (because of an upstream dereference) turns a would-be crash into a silent security hole. (c) Argue why a *static* guarantee (Rust's safe subset having no UB) is qualitatively stronger than a runtime sanitizer here.

10.12 **LAB** On your own machine, compile `int f(int x){ return x + 1 > x; }` with `gcc -O2 -S` and read the assembly; then recompile with `-fwrapv` and compare. Do the same for the §10.5 `poll` function with and without `-fno-delete-null-pointer-checks`. (a) Record what changed in each case. (b) Write a tiny program with a signed-overflow bug and confirm whether `-fsanitize=undefined` reports it. Report only what you observed and label your compiler/version, since output varies.

Chapter 11 — Arrays, Strings & Buffer Safety

Before you start: Chapter 10 (a buffer overflow is undefined behavior — and the optimizer trusts it never happens), Chapter 9 (the buffer-overflow bug, stated precisely), Chapter 8 (the stack, frames, and the return address that sits on it), Chapter 7 (arrays, pointers, and array-to-pointer decay). · **Continues Part 2.** This chapter takes the buffer overflow — Chapter 9's most consequential bug — and follows it all the way to its worst outcome: overwriting a return address and seizing a program's control flow. Then it covers the C string model that makes overflows so easy, the string-function minefield, and the layered defenses real systems deploy — measuring each on this machine.

BEFORE YOU READ ON

From memory, reaching back: (1) Chapter 8 — a function's local array and its return address both live in the same stack frame; which direction does the stack grow, and which of those two is at a *higher* address on a typical downward-growing stack? (2) Chapter 9 — an out-of-bounds *write* was called "especially dangerous." Why more dangerous than an out-of-bounds read? (3) *Predict:* you replace an unsafe `strcpy(dst, src)` with `strncpy(dst, src, n)` where `n` is exactly the size of `dst`, and `src` is `n` characters long. Do you think the result is a valid, null-terminated C string? We will compile exactly this and read the bytes.

Almost every catastrophic memory-corruption vulnerability in the history of software traces to the same root: a program wrote past the end of a buffer. Chapter 9 named it and called it undefined behavior; Chapter 10 showed why UB is dangerous in the abstract. This chapter makes it visceral, because the buffer overflow is not merely "corrupt some adjacent bytes" — on the stack, the bytes adjacent to a local array include the saved return address, and overwriting *that* hands an attacker control of what the CPU executes next. We will build up to it deliberately: why the C string is a fragile convention rather than a real type (§11.1); the anatomy of a stack buffer overflow and why it is exploitable (§11.2); the string-function minefield and the sized functions that replace it — including the trap that `strncpy` is *not* the safe `strcpy` (§11.3); the layered defenses (stack canaries, `_FORTIFY_SOURCE`, ASLR, NX/DEP) and precisely what each does and does not guarantee, run on this machine (§11.4); and the disciplines that actually prevent the bug (§11.5). Throughout, the refrain from Chapters 9 and 10 holds: the bug is UB, so a program that "runs fine in testing" may be one crafted input away from disaster.

11.1 The C string is a convention, not a type

C has no string type. What C calls a string is a **convention**: a contiguous array of `char` terminated by a null byte `'\0'` (a zero byte, value 0). The length is not stored anywhere — it is *implied* by the position of the terminator, which is why `strlen` is an $O(n)$ scan that walks bytes until it finds the zero. This one design decision is the source of a startling share of C's danger. Because the length is not carried with the data, every function that touches a string must either be told the buffer's capacity or trust that a terminator exists within bounds — and when that trust is misplaced, the read or write runs off the end. Recall from Chapter 7 that an array decays to a pointer to its first element the moment you pass it to a function: the callee receives an address with *no length information at all*.

A `char *` tells you where a string starts; it tells you nothing about how much room there is. Bounds safety in C is therefore never automatic — it is something the programmer must reconstruct by hand, by tracking capacities alongside pointers.

Two consequences follow immediately. First, **a missing terminator is a bug waiting to happen**: a "string" with no `'\0'` in bounds makes `strlen`, `printf("%s")`, and every other string function read past the end — an out-of-bounds read (UB), which can leak adjacent memory (the Heartbleed class, Ch 9). Second, **a copy into an undersized buffer overflows it**: writing an n -plus-one-byte string (n characters plus the terminator) into an n -byte buffer writes one byte past the end — an out-of-bounds write (UB), the subject of the rest of this chapter. The C string's elegance (a string is just bytes) and its danger (a string is *just* bytes, with no length and no bounds) are the same fact.


QUICK CHECK

Why is `strlen` an $O(n)$ operation rather than $O(1)$? And when you pass a `char buf[64]` to a function, how much does the function know about `buf`'s capacity? (Answer after the next section.)

11.2 The stack buffer overflow — anatomy and why it is exploitable

Recall the stack frame from [Chapter 8](#). When a function runs, its local variables — a `char buf[8]`, say — live in its stack frame, and (on the common downward-growing stack of x86-64 and most architectures) the frame also holds the **saved return address**: the address the CPU will jump to when this function returns. Critically, the return address sits at a *higher* address than the local buffer, while a copy into the buffer writes toward *higher* addresses. So a write that overruns `buf` marches straight toward the saved return address. Overflow it by enough bytes and you overwrite the return address itself. Now the picture is stark: when the function returns, the CPU loads that overwritten value into the instruction pointer and jumps *wherever the overflowing bytes told it to*. An attacker who controls the input controls where the program goes next.

One stack frame (stack grows down; a copy into `buf` writes UP toward the return address):



`strcpy(buf, "AAAAAAA...\\xef\\xbe\\xad\\xde");` 8 bytes fill `buf`, more climb over the saved FP, and the final bytes land on the return address. On ``ret``, execution jumps to the attacker's value.

A stack canary (§11.4) sits between `buf` and the return address; the overflow trips it before ``ret``.

Figure 11.1: A stack buffer overflow. Because the saved return address sits above the local buffer and a copy writes upward, overrunning the buffer can overwrite the return address, redirecting control flow when the function returns. A stack canary is placed in the gap to detect exactly this.

This is not a hypothetical; it is the mechanism behind decades of exploits, laid out for the first time in a widely read form by **Aleph One's "Smashing the Stack for Fun and Profit"** (Phrack 49, 1996) and developed rigorously in CS:APP (§3.10). The classic exploit places the attacker's own machine code ("shellcode") in the buffer and overwrites the return address to point at it; modern mitigations (§11.4) have made that specific technique much harder, but the fundamental hazard — an out-of-bounds write reaching control data — remains, and newer techniques (return-oriented programming) route around the defenses. To make it concrete, this program overflows an 8-byte buffer:

```
#include <string.h>
int main(void){
    char buf[8];
    strcpy(buf, "this string is much too long for the buffer"); /* 44 bytes into 8 */
    return buf[0];
}
```

Compiled with stack protection on this machine (`gcc -O0 -fstack-protector-all`) and run, it does *not* silently corrupt and continue — the canary catches it (§11.4). `gcc` even warns at compile time before you run it:

```
warning: '__builtin_memcpy' writing 44 bytes into a region of size 8 overflows the
destination
      [-Wstringop-overflow=]
$ ./a.out
*** stack smashing detected ***: terminated
Aborted (core dumped)          # exit status 134 (SIGABRT)
```

Without those defenses — an older toolchain, or a copy whose length the compiler cannot see — the same overflow would have corrupted the return address silently. **Outputs are from this machine (gcc 11.4.0, x86-64); wording and status vary by toolchain.**

TRAP: "STRNCPY IS THE SAFE STRCPY"

The most common half-fix in C is to reach for `strncpy(dst, src, n)` believing the `n` makes it safe. It bounds the write, which prevents the overflow — but **`strncpy` does not guarantee a null terminator**. Its actual contract (cppreference; the C standard): it copies at most `n` bytes; if `src` is shorter than `n` it pads the rest with zeros, but if `src` is `n` or more bytes long it writes exactly `n` bytes and stops — leaving `dst` **unterminated**. The next `strlen` or `printf("%s")` then runs off the end: you traded a bounded overflow for an unbounded over-read. On this machine, `char dst[4]; strncpy(dst, "abcd", 4);` leaves the four bytes 97 98 99 100 ('a' 'b' 'c' 'd') with *no* terminating zero — a "string" that is not a string. `strncpy` was designed in the 1970s for fixed-width records (like early UNIX directory entries), *not* as a safe string copy, and using it as one is a classic trap. If you use it, you must terminate by hand (`dst[n-1] = '\0';`) — or, better, use `snprintf` (§11.3), which always terminates.

(Answer to the §11.1 quick check: `strlen` is $O(n)$ because a C string stores no length — it must scan byte by byte until it finds the terminating `'\0'`. And a function receiving `char buf[64]` knows **nothing** about the capacity: the array decayed to a bare `char *`, an address with no length attached — you must pass the size separately.)

11.3 The string-function minefield — and the sized alternatives

The classic C string functions are unsafe by construction because they take no destination size: they write until they hit a terminator in the *source*, with no regard for the *destination's* capacity. The worst offenders and their safer replacements:

- **gets(char *s)** — reads a line into *s* with *no* bound whatsoever. It was so irredeemably unsafe that it was deprecated in C99 and **removed from the language in C11**. Never use it; use `fgets(s, size, stdin)`, which takes the buffer size.
- **strcpy(dst, src) / strcat(dst, src)** — copy/append with no destination bound; overflow *dst* whenever *src* is too long. Prefer `snprintf` (below), or bounded copies where you control termination.
- **sprintf(dst, fmt, ...)** — formats into *dst* with no bound; a long formatted result overflows. Replace with `snprintf(dst, size, fmt, ...)`, which writes at most *size* bytes *including* the terminator, always null-terminates (when *size* > 0), and returns the length it *would* have written — so you can even detect truncation.
- **strncpy / strncat** — bounded but treacherous: `strncpy` may not terminate (the §11.2 trap), and `strncat's` size argument is the number of bytes to append, not the buffer size, which readers routinely get wrong. Usable, but only with care and a manual terminator.

The reliable pattern for building a string safely is `snprintf`: it is bounded, always terminates, and reports truncation. When you truly need a raw copy, carry the capacity and check it — `if (srclen + 1 <= dstcap) memcpy(dst, src, srclen + 1);` — the same "reconstruct the length by hand" discipline §11.1 demanded. (Some platforms offer `strlcpy/strlcat`, which always terminate and return the source length; they are widely available but are *not* in the C standard, so their portability must be checked. The standard's own `_s` "bounds-checking" functions, Annex K, exist but are optional and unevenly implemented — do not assume them.)

QUICK CHECK

Name the one standard string function that was *removed* from C11 for being unsafe, and its safe replacement. And what does `snprintf` guarantee that `strncpy` does not? (Answer below the communicate box.)

11.4 Defense in depth — canaries, FORTIFY, ASLR, NX — and what each buys

Because buffer overflows are so consequential, modern toolchains and operating systems layer several *mitigations*. It is essential to understand what each actually guarantees, because none of them *removes* the bug — they raise the cost of exploiting it, and each defeats a specific technique while leaving others open. Four you will meet, three demonstrated on this machine:

- **Compiler bounds warnings.** With optimization and `-Wall`, gcc's `-Wstringop-overflow` can catch overflows it can see at compile time. On this machine the §11.2 program

drew

warning: `'__builtin_memcpy'` writing 44 bytes into a region of size 8 overflows the destination — a static catch, but only when the sizes are visible to the compiler.

- **`_FORTIFY_SOURCE`**. Building with `-D_FORTIFY_SOURCE=2 -O2` makes glibc substitute checked variants (`__strcpy_chk`, `__memcpy_chk`, ...) that know the destination size and abort on overflow — catching many cases at compile time *or* runtime. On this machine, fortifying `char buf[4]; strcpy(buf, "hello");` produced warning: `'__builtin__strcpy_chk'` writing 6 bytes into a region of size 4. It only helps where the size is known and requires optimization; it is not universal.
- **Stack canaries** (`-fstack-protector` / `-fstack-protector-all`). The compiler places a random "canary" value between the local buffers and the saved return address (Figure 11.1) and checks it just before returning; if an overflow overwrote it, the program aborts instead of returning to corrupted control data. On this machine, the §11.2 overflow tripped it: `*** stack smashing detected ***: terminated, exit status 134`. Canaries defeat the classic "overwrite the return address by linear overflow" attack — but not overwrites that skip the canary, or non-control-data corruption, and they act *after* the corruption (they convert exploitation into a crash, i.e. a denial of service, not into "no bug").
- **ASLR and NX/DEP** (OS-level). **Address Space Layout Randomization** randomizes the base addresses of the stack, heap, and libraries each run, so an attacker cannot easily predict where to jump. **NX/DEP** (the no-execute page permission, tied to virtual memory from Chapter 4) marks the stack and heap non-executable, so injected shellcode cannot run as code. Together they defeat the original "inject shellcode on the stack and jump to it" recipe — which is why real attacks moved to *return-oriented programming*, chaining existing code fragments rather than injecting new ones. These are runtime/OS defenses, not compiler ones, and again they raise cost rather than eliminate the bug.

The load-bearing point: **every one of these is a mitigation, not a fix**. They assume the overflow exists and try to contain its blast radius — turning a silent hijack into a crash, or a reliable exploit into an unreliable one. The bug is still there, still undefined behavior (Ch 10), and a determined attacker or an unlucky input may route around any single layer. The only true fix is to not write out of bounds in the first place (§11.5) — or to use a language that enforces bounds by construction (Part 5).

THE SAME IDEA THREE WAYS: A BUFFER OVERFLOW IS A BOUNDS VIOLATION WITH CONSEQUENCES

1. In words. Writing past a buffer's end (§9.3's bounds violation) is undefined behavior; on the stack the bytes just past a local array include control data — the saved return address — so the overflow can redirect execution, not merely corrupt data.

2. As a picture. Figure 11.1: `buf` low, saved frame pointer above it, return address above that; a copy writes upward, so an overrun climbs over the frame pointer onto the return address. A canary sits in the gap to catch the climb.

3. As UB the compiler and OS fight in layers. The write is UB (Ch 10) — no defined outcome. Compiler warnings and `_FORTIFY_SOURCE` catch some cases; canaries turn a hijack into a crash; ASLR and NX defeat specific exploit recipes. None removes the UB; only bounds-correct code does.

SAY IT LIKE A SYSTEMS ENGINEER

Try it first. In a review, a colleague writes: "I turned on `-fstack-protector`, so our buffer overflows are handled — we don't need to fix the `strcpy` calls." Cover the paragraph and respond in three or four sentences, distinguishing *mitigation* from *fix*.

Model answer. "A stack canary is a mitigation, not a fix: it detects a linear stack overflow just before the function returns and aborts, so it converts a control-flow hijack into a crash — a denial of service, still a real incident, and it does nothing for heap overflows, non-control-data corruption, or overwrites that skip the canary. The overflow is still there, still undefined behavior, so the compiler is free to do surprising things around it, and a canary that fires in production is an outage. We should keep the canary *and* ASLR *and* `_FORTIFY_SOURCE` as defense in depth, but the actual fix is to remove the unbounded copies — replace `strcpy` with `snprintf` or a capacity-checked `memcpy` — so the out-of-bounds write cannot happen." Naming the distinction — mitigation contains blast radius, only bounds-correct code removes the bug — is the senior register, and it keeps the defenses without treating them as a license to leave the bug in.

(Answer to the §11.3 quick check: `gets` was removed from C11; its safe replacement is `fgets`, which takes the buffer size. And `snprintf` always null-terminates its output (when the size is nonzero) and reports truncation, whereas `strncpy` may leave the destination with no terminator at all.)

11.5 What to carry forward

The C string is a convention — a null-terminated byte array with no stored length — and that single fact is why bounds safety in C is manual and why buffer overflows are so easy. An out-of-bounds write is undefined behavior (Ch 9, Ch 10), and on the stack it is the most consequential UB in security: the saved return address lives just past a local buffer, so an overrun can overwrite it and redirect control flow — the mechanism Aleph One documented and CS:APP develops. The classic string functions (`gets`, `strcpy`, `strcat`, `sprintf`) take no destination size and are unsafe by construction; `gets` was removed from C11 outright. The sized replacements — `fgets`, `snprintf`, capacity-checked `memcpy` — are safe *when you carry the capacity and check it*, and `strncpy` is a trap because it can leave a string unterminated. Layered defenses (compiler warnings, `_FORTIFY_SOURCE`, stack

canaries, ASLR, NX/DEP) each raise the cost of exploitation and we watched three of them fire on this machine — but every one is a mitigation that contains the blast radius, not a fix that removes the bug.

The discipline that actually prevents the bug is the same one Chapters 7–10 kept circling: a bare pointer carries no bounds, so *you* must carry them — keep the capacity beside every buffer, check every copy against it, prefer sized APIs, and never trust that a terminator is present. That is a proof obligation, on every string operation, that the programmer discharges by hand and that no amount of mitigation discharges for you. It is also the cleanest possible statement of what a safer language buys: in Rust (Part 5) a string and a slice *carry their length*, indexing is bounds-checked, and an out-of-bounds access is a defined panic rather than undefined memory corruption — the "carry the length by hand" discipline of this chapter, promoted into the type system. Part 2 has now taken you from "a pointer is an address" (Ch 7) to "and you own every lifetime, every bound, and every promise to the compiler" (Ch 8–11). The final chapter turns the last of those responsibilities — allocation — from a hazard you survive into a tool you wield.

CAN YOU... WITHOUT LOOKING?

- Explain why a C string is "a convention, not a type," and why that makes `strlen` $O(n)$ and bounds safety manual.
- Draw a stack frame and explain why a buffer overflow can overwrite the **saved return address** and hijack control flow.
- State why `strncpy` is *not* the safe `strcpy` — the exact contract that bites — and name a genuinely safe alternative.
- Name the four classic unsafe string functions and their sized replacements, and which one C11 removed.
- Describe stack canaries, `_FORTIFY_SOURCE`, ASLR, and NX/DEP, and say for each what it buys and why it is a mitigation, not a fix.
- State the one discipline that actually prevents overflows, and how Rust promotes it into the type system.

RETRIEVAL — CLOSED BOOK

Cover the chapter. Answer from memory; rate confidence first.

1. **R1.** What is a C string, physically, and why does it store no length? (§11.1)
2. **R2.** Why can a stack buffer overflow hijack control flow — what sits next to the buffer, and in which direction does a copy write? (§11.2)
3. **R3.** State `strncpy`'s exact contract and the trap it creates; give a safer function. (§11.2–11.3)
4. **R4.** For canaries, `_FORTIFY_SOURCE`, ASLR, and NX/DEP, what does each defeat, and why is each a mitigation not a fix? (§11.4)
5. **R5.** What single discipline prevents overflows, and how does Rust encode it? (§11.5)

FURTHER READING

- **Aleph One, "Smashing the Stack for Fun and Profit," Phrack 49 (1996).** The first widely read, step-by-step account of stack buffer overflow exploitation; historically pivotal. Free (phrack.org). Read it for the mechanism, not as current exploit practice (the mitigations of §11.4 postdate it).
- **Bryant & O'Hallaron, *Computer Systems: A Programmer's Perspective (CS:APP)*, §3.10.** Buffer overflows, stack smashing, and the mitigations (canaries, ASLR, NX) developed rigorously with the stack-frame model of Chapter 8. The reference for this chapter.
- **cppreference.com** and the **C standard** — the exact contracts of `strcpy`, `strncpy`, `strcat`, `snprintf`, `fgets`, and the removal of `gets` in C11. The authority for what each function promises.
- **Beej's Guide to C.** Free (beej.us). Practical, readable coverage of C strings and the safe idioms.

Exercises

- 11.1** **WARM-UP** In one sentence each: (a) what physically marks the end of a C string? (b) why is `strlen` $O(n)$? (c) when you pass an array to a function, what does the function learn about its capacity?
- 11.2** **WARM-UP** Which standard string function was *removed* from the language in C11 for being unsafe, and what should you use instead?
- 11.3** **WARM-UP** True or false, with a one-line reason each: (a) `strncpy(dst, src, n)` always produces a null-terminated string. (b) A stack canary removes the buffer-overflow bug. (c) `snprintf` always null-terminates when its size argument is nonzero.
- 11.4** **CORE** Explain, using a stack-frame diagram in words, why an out-of-bounds *write* into a local array can hijack a program's control flow but an out-of-bounds *read* (usually) cannot. What lies just past the buffer, and in which direction does a copy write?
- 11.5** **CORE** For each, state whether it can overflow the destination and why, and give a safe rewrite: (a) `strcpy(buf, user_input);` (b) `sprintf(buf, "%s-%s", a, b);` (c) `char d[4]; strncpy(d, "abcd", 4); printf("%s", d);`
- 11.6** **CORE** A colleague argues: "We enabled stack canaries and ASLR, so buffer overflows are handled." Give three distinct reasons this is wrong — each naming something a canary or ASLR does *not* prevent — and state what the actual fix is.
- 11.7** **COST** (a) Why is `strlen` $O(n)$ while a length-carrying string type (as in Rust or C++ `std::string`) makes "get the length" $O(1)$? (b) What runtime cost do stack canaries and `_FORTIFY_SOURCE` add, and what do you get for it? (c) Is a canary check on the hot path free? Classify the trade-off (safety vs. speed).

11.8 **FADED** *Worked, with the last step for you.* Diagnose and fix: `void greet(const char *name){ char msg[16]; strcpy(msg, "Hello, "); strcat(msg, name); puts(msg); }`.

Step 1 (given): "Hello, " is 7 characters plus a terminator, using 8 of the 16 bytes.

Step 2 (given): `strcat` appends `name` with no bound; if `name` is longer than 7 characters the total exceeds 16 bytes.

Step 3 (given): So a long `name` overflows `msg` — an out-of- bounds write, undefined behavior, and (on the stack) a potential control-flow hijack.

Step 4 (you): Rewrite the function using `snprintf` so it can never overflow regardless of `name`'s length; state what `snprintf` guarantees that `strcat` does not; and name one §11.4 defense that would turn the original overflow into a crash rather than a silent hijack.

11.9 **MIXED** For each symptom, name the most likely cause — *unterminated string (over-read)*, *buffer overflow (over-write)*, *use-after-free* (Ch 9), or *a mitigation firing* — and justify in a line: (a) a program prints a corrupted string with garbage bytes appended after the expected text; (b) a service aborts with `*** stack smashing detected ***` on a long request; (c) reading a freed heap block returns another request's data; (d) `printf("%s", s)` reads far past where the data should end and leaks adjacent memory.

11.10 **MIXED** Classify each as *undefined behavior*, *a defined-but-real defect*, or *well-defined and fine*, drawing on Chapters 9–11: (a) `strncpy(d, s, sizeof d)` then reading `d` as a string, where `strlen(s) >= sizeof d`; (b) `snprintf(d, sizeof d, "%s", s)` then reading `d`; (c) writing one byte past a heap block; (d) a canary aborting the process on a detected overflow.

11.11 **STRETCH** The chapter calls the mitigations "defense in depth." (a) Explain what each of stack canaries, NX/DEP, and ASLR defeats, and name the exploit technique (return-oriented programming) that was developed to route around NX. (b) Argue why layering mitigations is still worthwhile even though none is a fix. (c) Explain what it means for Rust to make the bug "not happen" rather than "not exploit" — and why a bounds-checked panic is categorically different from undefined memory corruption.

11.12 **LAB** On your own machine, write the §11.2 8-byte-buffer overflow. (a) Compile with `-fstack-protector-all`, run it, and record the exact message and exit status. (b) Compile the same with `-D_FORTIFY_SOURCE=2 -O2` and record any compile-time diagnostic. (c) Write `char d[4]; strncpy(d, "abcd", 4);` and print the four bytes as integers to confirm whether a terminator is present. Report only what you observed and label your compiler/version, since output varies.

Chapter 12 — Custom Allocators: Arenas, Pools & Free Lists

Before you start: Chapter 9 (what the allocator does — free lists, size classes, fragmentation, cost), Chapter 10 (a use-after-free is undefined behavior, whoever owns the memory), Chapter 5 (mechanical sympathy: cost follows layout and access pattern). · **Closes Part 2.** Chapter 9 treated the general-purpose allocator as a fact of the machine and cataloged the ways manual memory goes wrong. This chapter turns allocation from a hazard you survive into a tool you wield: when your allocation *pattern* is known, a custom allocator — an arena, a pool, a free list — can be dramatically faster than `malloc`, and we will build and measure one on this machine. It also completes the arc of the part: even your own allocator does not relax a single rule from Chapters 9–11.

BEFORE YOU READ ON

From memory, reaching back: (1) Chapter 9 — why is a heap `malloc` far more expensive than the stack's pointer-bump, and what two things (a search structure and a per-block record) does the allocator maintain? (2) Chapter 9 — what is *external fragmentation*, and why can it make a large allocation fail even when total free memory is ample? (3) *Predict:* suppose you allocate a million small objects that all die at the same moment. If you could replace a million individual `free` calls with a *single* operation that reclaims them all at once, roughly how much faster do you think freeing would be? We will measure exactly this.

Chapter 9's allocator is a marvel of generality: it satisfies any size, in any order, at any time, and it is the right default for the overwhelming majority of code. But generality has a price — every allocation may search a free list, split a block, and update bookkeeping; every free may coalesce; and interleaved requests fragment the pool. When you know something specific about *how* your program allocates — that many objects share a lifetime, or that they are all the same size — you can trade the general allocator's flexibility for speed and simplicity by writing an allocator specialized to that pattern. This chapter builds the two most useful specializations: the **arena** (also called a region or bump allocator), which makes allocation a pointer increment and frees everything at once (§12.2), and the **pool** or free-list allocator for fixed-size objects (§12.3). We start by asking precisely when the general allocator loses (§12.1), build and measure an arena on this machine, then close with an honest accounting of when a custom allocator is worth it — and the crucial fact that it does not lift the burden of Chapters 9–11 (§12.4–12.5). This is the capstone of "manual memory": having felt its hazards, you now use it deliberately.

12.1 Why not just `malloc`? — allocation patterns

The general-purpose allocator is optimized for the case where it knows *nothing* about your requests. That ignorance is what costs you. Three patterns where a specialized allocator can do far better, precisely because it exploits knowledge the general allocator lacks:

- **Many objects, one lifetime.** A request handler, a compiler pass, a frame of a game — all allocate a burst of objects that become garbage together at a known moment.

The general allocator makes you `free` each one individually (and track every pointer to do so); an allocator that knows they share a lifetime can free the whole batch in one operation.

- **Many objects, one size.** A pool of network connections, tree nodes, fixed-size buffers. The general allocator's size-class machinery and per-block headers are overhead you do not need if every object is identical; a fixed-size allocator can hand out and reclaim blocks in $O(1)$ with almost no bookkeeping and no fragmentation *within* that size.
- **Hot allocation loops.** Per-iteration `malloc/free` in a tight loop pays the general allocator's cost every iteration. Amortizing that — allocating a big block once and sub-dividing it cheaply — can dominate.

The mechanism behind all three is the same idea as Chapter 5's mechanical sympathy: when you know the access or lifetime pattern, you can lay memory out and manage it to match, and the machine rewards you. But there is a discipline here that this chapter insists on and §12.4 returns to: custom allocators are a *measured* optimization, not a default. Berger, Zorn, and McKinley's study of eight real programs using custom allocators (OOPSLA 2002) found that for most of them a good general-purpose allocator matched or beat the custom one — the big wins came specifically from the region/arena pattern (the "one lifetime" case), where they measured improvements up to 44%. The lesson is not "always write your own"; it is "when the pattern is right, the win is real and large — and when it is not, you have added complexity for nothing."

QUICK CHECK

Name the allocation pattern an *arena* exploits and the one a *fixed-size pool* exploits. Which single piece of knowledge does each have that the general-purpose allocator does not? (Answer after the next section.)

12.2 The arena (region / bump) allocator — build and measure it

An **arena** (or region, or bump allocator) exploits the "many objects, one lifetime" pattern. The idea is minimal: reserve one large block up front, keep a single *offset* into it, and satisfy each allocation by returning the current offset and bumping it forward by the requested size (rounded up for alignment). There is no free list, no size class, no per-block header, and — most importantly — no per-object free: individual objects are *never* freed. Instead the entire arena is reclaimed at once, either by resetting the offset to zero (reuse the memory) or by freeing the backing block. Allocation is a pointer increment; bulk-free is a single assignment or one `free`. Here is the whole allocator:

```

typedef struct { unsigned char *base; size_t cap, off; } Arena;

Arena arena_new(size_t cap){ Arena a; a.base = malloc(cap); a.cap = cap; a.off = 0; return
a; }

void *arena_alloc(Arena *a, size_t n){
    size_t p = (a->off + 15) & ~(size_t)15;    /* round up to 16-byte alignment */
    if (p + n > a->cap) return NULL;           /* arena exhausted */
    a->off = p + n;                             /* bump the offset */
    return a->base + p;
}

void arena_reset(Arena *a){ a->off = 0; }      /* frees EVERYTHING in O(1) */

```

That is a complete, working allocator. Compiled and run on this machine, it hands out correctly aligned, independent blocks ($x=42$, $s="hello"$, $y=7$, each 16-byte aligned), and `arena_reset` returns the offset to 0, reclaiming all of them in one operation. The payoff is speed, and it is worth measuring honestly. The arena wins when many objects are live *simultaneously* and die together — so the fair comparison is: allocate N objects that all stay live, then free them. With the arena that is N pointer-bumps and one bulk free; with `malloc` it is N allocations and N individual frees. For $N = 5,000,000$ small objects on this machine (gcc 11.4.0, -O2, x86-64):


```

arena (N live, 1 bulk free) : ~50-90 ms
malloc (N live, N frees)    : ~240-370 ms
speedup                     : ~4-5x

```

Roughly a **four-to-five-fold** speedup, which lines up with the region wins Berger et al. reported. **These figures are illustrative and from this specific machine/run — they vary run-to-run (the ranges above are across repeated runs) and with size, allocator, and hardware; measure your own.** The win has a clear source: the arena replaced a million bookkeeping-heavy free calls with a single offset reset, and made each allocation a pointer add instead of a free-list search.

Arena: one block, one offset. Each alloc bumps the offset; reset frees everything at once.



`alloc(n)`: $p = \text{round_up}(\text{off})$; $\text{off} = p + n$; return $\text{base} + p$; — $O(1)$, a pointer add. No header, no search.

`reset()`: $\text{off} = 0$; — frees A, B, C (and all others) in ONE step. No per-object free.

Trade-off: you CANNOT free obj B alone. If B outlives A and C, its memory is stuck until the whole arena resets.

Right pattern: objects that share a lifetime (a request, a frame, a compiler pass). Wrong pattern: mixed, long lifetimes.

Berger/Zorn/McKinley (OOPSLA 2002): regions win big for the right pattern, but can inflate memory use when objects outlive the batch.

Figure 12.1: An arena allocates by bumping a single offset and frees everything at once by resetting it — $O(1)$ allocation, $O(1)$ bulk free, no fragmentation, no per-object free. The cost is inflexibility: you cannot free one object early, so memory is held until the whole region is reset.

The trade-off is exactly that inflexibility, and it is not free. Because you cannot free an individual object, any object that outlives the batch pins the *entire* arena's memory until reset — so an arena used with the wrong lifetime pattern can inflate memory

consumption badly (Berger et al. note precisely this downside of regions). And — the theme of §12.5 — resetting an arena while a pointer into it is still in use is a use-after-free by another name: undefined behavior, exactly as in Chapter 10, now with *you* as the allocator.

*(Answer to the §12.1 quick check: an **arena** exploits "many objects share one lifetime" — it knows the batch dies together, so it can skip per-object frees and reclaim all at once. A **fixed-size pool** exploits "many objects, one size" — it knows every block is identical, so it needs no size classes, no split/coalesce, and no per-block size header.)*

12.3 The pool / free-list allocator — fixed-size, O(1) both ways

When objects share a *size* rather than a *lifetime* — a pool of tree nodes, connection structs, fixed buffers — a **pool allocator** (a fixed-size free-list allocator) is the specialization that fits. Carve one large block into an array of equal-size slots and thread the free slots onto a singly linked **free list**: because every slot is the same size and currently unused, you can store the "next free slot" pointer *inside the slot itself*, costing no extra memory. Allocation pops the head of the free list; freeing pushes the slot back onto the head. Both are O(1), a couple of pointer operations, with no search, no splitting, no coalescing, and — because all slots are identical — no external fragmentation within the pool at all. The sketch:

```
/* free_list points at the first free slot; each free slot's first bytes hold the next
pointer */
void *pool_alloc(Pool *p){
    if (!p->free_list) return NULL;           /* pool exhausted */
    void *slot = p->free_list;
    p->free_list = *(void **)slot;           /* pop: next free = this slot's stored next */
    return slot;
}
void pool_free(Pool *p, void *slot){
    *(void **)slot = p->free_list;           /* push: this slot's next = old head */
    p->free_list = slot;                     /* new head = this slot */
}
```

Unlike the arena, a pool *does* support freeing individual objects — that is its point — so it fits the "allocate and free repeatedly, but always the same size" pattern the arena cannot. What it gives up is generality: it serves exactly one size. This is not a toy idea. The Linux kernel's **slab allocator** — introduced in Jeff Bonwick's 1994 paper (USENIX) for SunOS and adopted widely since — is a production-grade elaboration: it maintains per-size caches of pre-initialized objects ("slabs") so that allocating a kernel object is a free-list pop and freeing it a push, retaining the object's initialized state between uses and adding a cache-coloring scheme to spread objects across cache lines (Ch 3). Production general allocators (glibc's `malloc`, jemalloc, tcmalloc) likewise use size-classed free lists internally — the pool is the general allocator's own core trick, exposed for you to specialize.

THE SAME IDEA THREE WAYS: SPECIALIZE THE ALLOCATOR TO THE PATTERN

1. In words. The general allocator is fast *because it knows nothing* is a contradiction — its generality costs search, headers, and fragmentation. If you know the pattern (shared lifetime → arena; shared size → pool), you drop the machinery you do not need and get $O(1)$ with almost no bookkeeping.

2. As two pictures. Arena (Fig 12.1): one offset, bump to allocate, reset to free all. Pool: an array of equal slots with a free list threaded through the unused ones; pop to allocate, push to free — both $O(1)$, no fragmentation within the size.

3. As numbers. Arena on this machine: $\sim 4\text{--}5\times$ faster than `malloc` for the one-lifetime batch (§12.2). Pool: allocation and free are each a couple of pointer writes versus the general allocator's free-list search and coalesce. The win is real *only* when the pattern holds (§12.4).

12.4 Choosing an allocator — an honest accounting

With three allocators in hand — general-purpose, arena, pool — the engineering question is when to reach for each. The honest answer starts with a warning that the research backs up: **the general-purpose allocator is the right default, and custom allocation is a measured optimization, not a reflex.** Berger, Zorn, and McKinley found that for six of eight real programs a state-of-the-art general allocator matched or beat the hand-written custom one — the programmers' custom allocators were often *slower* as well as buggier and harder to maintain. Custom allocation earned its keep in exactly the region/arena case. So the decision procedure is: profile first (Chapter 9's cost model, and the performance-engineering discipline Part 6 develops — measure before you optimize); confirm allocation is actually a bottleneck; then match the allocator to the pattern:

- **General-purpose `malloc`** — the default. Any size, any order, any lifetime, individual frees. Use it unless you have measured a reason not to.
- **Arena** — when a burst of objects shares a lifetime and dies together (a request, a frame, a parse). $O(1)$ alloc, $O(1)$ bulk free, no per-object free. Watch the memory downside: anything that outlives the batch pins the whole region.
- **Pool** — when objects share a fixed size and are allocated/freed repeatedly. $O(1)$ both ways, no fragmentation within the size, individual frees supported. Serves exactly one size.

Each choice is a trade, not a free win. The arena trades the ability to free individually for speed and simplicity — and can waste memory if the lifetime assumption is wrong. The pool trades generality (one size only) for $O(1)$ operations and zero fragmentation. And every custom allocator trades the maturity, debugging support, and hardening of the system allocator for code you now own and must get right. That last cost is the bridge to §12.5.

TRAP: "A CUSTOM ALLOCATOR MAKES THE MEMORY BUGS GO AWAY"

It is tempting to think that once *you* control allocation, the hazards of Chapters 9–11 are behind you. The opposite is true: **writing your own allocator makes you responsible for the very invariants the language and the system allocator's tooling used to help enforce** — and it introduces new ways to violate them. Reset an arena while a pointer into it is still live, and the next allocation reuses that memory under the old pointer: a *use-after-reset*, which is a use-after-free (Ch 9), undefined behavior (Ch 10), with you as the allocator. `pool_free` the same slot twice and you corrupt the free list — a double-free in your own code. Hand out a slot without checking the pool is non-empty and you return a null you may dereference. Worse, sanitizers (Ch 9) understand `malloc/free` but do *not*, by default, understand *your* arena's reset or *your* pool's recycling, so they may not catch a use-after-reset at all — you have moved the bug below the tool's visibility. (ASan offers manual "poisoning" hooks precisely so custom allocators can be made visible to it, but you must add them.) A custom allocator does not repeal lifetimes, bounds, and free-exactly-once; it makes you their implementer. Every rule from Chapters 9–11 still holds — you have just taken over enforcing them.

QUICK CHECK

You reset an arena to reclaim a request's memory, but one object from that request was stored in a cache that outlives the request. What bug have you just created, and what is its Chapter-9 name? (Answer below the communicate box.)

SAY IT LIKE A SYSTEMS ENGINEER

Try it first. A teammate proposes: "Our service is slow — let's replace `malloc` with an arena allocator everywhere for a quick speed win." Cover the paragraph and respond in three or four sentences, drawing on §12.1 and §12.4.

Model answer. "An arena is a big win for a specific pattern — a burst of objects that share a lifetime and die together, like everything allocated while handling one request — but it's the wrong tool 'everywhere,' because it can't free individual objects, so any object that outlives its batch pins the whole region and can inflate our memory badly. The evidence backs the caution: Berger, Zorn, and McKinley found a good general allocator matched or beat hand-written custom allocators for most programs, with the arena/region case being the real exception. Let's profile first to confirm allocation is actually the bottleneck, then apply an arena precisely where the lifetime pattern fits — per-request scratch memory — and keep `malloc` for the mixed-lifetime long-lived state. And whatever we pick, we still own every lifetime and bound: resetting an arena with a live pointer into it is a use-after-free that our sanitizers won't see by default." Naming the pattern, citing the measured result, and insisting on profiling — rather than a blanket swap — is the senior register.

(Answer to the §12.4 quick check: you have created a **use-after-free** — specifically a *use-after-reset*. Resetting the arena ended the lifetime of that object's storage; the cached pointer now dangles, and the next allocation from the arena will reuse those bytes under it. It is undefined behavior, and by default the sanitizers will not flag it, because they do not know your arena reset ended a lifetime.)

12.5 What to carry forward — and out of Part 2

The general-purpose allocator is optimized for knowing nothing about your requests, and that ignorance is its cost — search, headers, coalescing, fragmentation (Ch 9). When you *do* know the pattern, a custom allocator trades generality for speed. An **arena** exploits shared lifetime: allocate by bumping an offset, free everything at once by resetting it — $O(1)$ both ways, no per-object free, measured at roughly four-to-five times faster than `malloc` for the batch pattern on this machine, at the cost of being unable to free one object early. A **pool** exploits shared size: an array of equal slots with a free list threaded through the unused ones, $O(1)$ allocation and free with no fragmentation within the size — the same trick the kernel's slab allocator and production `mallocs` use internally. But custom allocation is a *measured* optimization, not a default: the research (Berger et al.) shows a good general allocator usually matches hand-rolled ones, with the arena/region pattern the clear exception — so profile, confirm the bottleneck, and match the allocator to the pattern. And, decisively, writing your own allocator does not lift one obligation from Chapters 9–11: you become the implementer of lifetimes, bounds, and free-exactly-once, with new ways to violate them (use-after-reset, pool double-free) that your tools may not see.

That last point closes Part 2. Across six chapters you went from "a pointer is an address into regioned memory" (Ch 7), through lifetimes and the stack/heap (Ch 8), manual allocation and its bug catalogue (Ch 9), undefined behavior and the optimizer that trusts it (Ch 10), the buffer overflow that hijacks control flow (Ch 11), to writing the allocator itself (Ch 12) — and the through-line never changed: **in C you hold, by hand, a proof that no pointer is used outside its lifetime, no access outside its bounds, every block freed exactly once, and no undefined behavior ever triggered.** That proof is real, it is heavy, and every bug in this part is what happens when it fails. You now understand the machine to the metal and the net that is absent. Two paths lead out. The operating system (Part 3) is what sits beneath all of this — the processes, threads, and virtual memory that gave you the address space and the heap in the first place; you will see the system-call boundary from the other side. And Rust (Part 5) is the language that makes the compiler *check* the proof you have been carrying — ownership, borrowing, and lifetimes turning the invariants of this entire part into type rules verified before the program runs. You have felt the weight; the rest of the book is about the machine underneath it and the tools that lift it.

CAN YOU... WITHOUT LOOKING?

- State the three allocation patterns where a custom allocator beats general-purpose malloc, and the knowledge each exploits.
- Write an **arena** allocator from memory (base, capacity, offset; bump to allocate, reset to free all) and state its one big trade-off.
- Write a **pool**/free-list allocator's alloc and free (pop/push a free list threaded through the slots) and say why both are $O(1)$ with no fragmentation within the size.
- Give the honest decision rule (general by default; profile; arena for shared lifetime; pool for shared size) and cite what the Berger et al. study found.
- Explain why a custom allocator does *not* remove the Chapter 9–11 obligations — and name a bug (use-after-reset) your sanitizers may miss.
- Summarize Part 2's through-line and how it sets up the OS (Part 3) and Rust (Part 5).

RETRIEVAL — CLOSED BOOK

Cover the chapter. Answer from memory; rate confidence first.

1. **R1.** Name three allocation patterns where a custom allocator wins, and what each exploits. (§12.1)
2. **R2.** Describe an arena's allocate and free, its complexity, and its trade-off. (§12.2)
3. **R3.** Describe a pool/free-list allocator and why it has no fragmentation within its size. (§12.3)
4. **R4.** What did Berger et al. find about custom vs. general allocators, and what is the decision rule? (§12.1, §12.4)
5. **R5.** Why does a custom allocator not remove the Chapter 9–11 obligations? Name a bug it introduces. (§12.4–12.5)

FURTHER READING

- **Bryant & O'Hallaron, *Computer Systems: A Programmer's Perspective (CS:APP)*, Chapter 9.** Dynamic memory allocation: allocator design, free lists, coalescing, and the mechanics an arena and pool specialize. The reference for §12.1's cost model.
- **Berger, Zorn & McKinley, "Reconsidering Custom Memory Allocation," OOPSLA 2002.** The empirical study behind §12.4: a good general allocator matches most hand-written custom allocators, with regions (arenas) the clear exception (improvements up to 44%, but a memory-consumption downside). Free from the authors.
- **Bonwick, "The Slab Allocator: An Object-Caching Kernel Memory Allocator," USENIX 1994.** The production pool/slab allocator behind §12.3, still used in the Linux kernel. Free.
- **The jemalloc / tcmalloc documentation.** Real, widely deployed general allocators whose internals (size-classed free lists, thread caches) are the pool idea at scale. Documented behavior only; verify against the version you use.

Exercises

- 12.1** **WARM-UP** In one sentence each: (a) what does an arena exploit? (b) what does a pool exploit? (c) which of the two can free an individual object?
- 12.2** **WARM-UP** What is the time complexity of an arena's `alloc` and its bulk `reset`? What is the one thing an arena cannot do that `malloc` can?
- 12.3** **WARM-UP** In a fixed-size pool, where is the "next free slot" pointer stored, and why does that cost no extra memory?
- 12.4** **CORE** Write, from memory, an arena allocator (a struct with base/capacity/offset, plus `arena_alloc` with 16-byte alignment and `arena_reset`). Explain why `arena_alloc` is $O(1)$ and why `arena_reset` frees everything without touching individual objects.
- 12.5** **CORE** Write a pool allocator's `pool_alloc` and `pool_free` using a free list threaded through the slots. Explain why both are $O(1)$ and why the pool has no external fragmentation within its size class.
- 12.6** **CORE** A teammate wants to replace `malloc` with an arena "everywhere for speed." Give two concrete reasons this is a bad blanket policy — one about the arena's inability to free individually, one citing the Berger et al. finding — and state the correct decision procedure.
- 12.7** **COST** (a) Why is an arena's per-allocation cost lower than `malloc`'s — name two things the general allocator does that the arena skips (tie to Ch 9). (b) The chapter measured $\sim 4\text{--}5\times$ on this machine for the "N live, then free" pattern; why does that comparison favor the arena, and what different pattern would narrow or erase the gap? (c) What memory-consumption risk does an arena carry, and when does it bite?
- 12.8** **FADED** *Worked, with the last step for you.* A per-request arena is reset at the end of each request. A developer adds a cache that stores, across requests, a pointer to a string that was allocated *in* the arena. Diagnose the bug.
- Step 1 (given):** `arena_reset` ends the lifetime of everything in the arena by returning the offset to 0.
- Step 2 (given):** The cached pointer still holds the string's old address, but that storage is now considered free.
- Step 3 (given):** The next request's allocations will reuse those bytes, so a later read through the cached pointer sees another request's data or garbage.
- Step 4 (you):** Name the bug using Chapter 9's vocabulary, state whether it is undefined behavior, explain why the ASan/Valgrind of §9.4 will likely *not* catch it by default, and give a fix (what must be copied where, so the cache does not point into the arena).

12.9 **MIXED** For each workload, choose *general malloc*, *arena*, or *pool*, and justify in a line: (a) a web handler that allocates many small objects while serving one request, all discardable when the response is sent; (b) a graph engine that constantly creates and destroys same-size nodes throughout its run; (c) a general-purpose library that cannot predict its callers' allocation patterns; (d) a compiler pass that allocates a burst of AST nodes freed together when the pass ends.

12.10 **MIXED** Classify each as *undefined behavior*, *a defined-but-real defect*, or *well-defined and fine*, drawing on Chapters 9–12: (a) reading through a pointer into an arena after `arena_reset`; (b) an arena holding a long-lived object that pins the whole region, growing memory use; (c) `pool_free`-ing the same slot twice; (d) resetting an arena when no live pointers into it remain.

12.11 **STRETCH** Part 2 has argued that in C the programmer carries, by hand, a proof about lifetimes, bounds, and freeing. (a) State that proof obligation in one sentence. (b) Explain how a custom allocator *relocates* rather than *removes* that obligation — who now enforces "free exactly once" for arena and pool memory? (c) In one or two sentences, describe what it would mean for a language (Rust, Part 5) to tie an allocation's lifetime to a scope so that "use after reset" becomes a compile-time error — and why that is stronger than a runtime sanitizer that does not understand your allocator.

12.12 **LAB** On your own machine, implement the arena of §12.2. (a) Verify correctness: allocate three objects of different sizes, confirm they are independent and 16-byte aligned, and that `arena_reset` returns the offset to 0. (b) Time allocating N objects that stay live then freeing them, arena (one bulk free) vs. `malloc` (N frees), for a large N , and report the ratio. (c) Repeat with a pattern where each object is freed immediately after allocation and explain why the arena's advantage shrinks. Report only what you measured and label your compiler/hardware, since numbers vary.

VOL 1 · SYSTEMS PROGRAMMING & PERFORMANCE · PART III

The Operating System

Processes, threads, scheduling, virtual memory, and the system-call boundary — how the OS multiplexes the machine.

Chapter 13 — The System-Call Boundary

Before you start: Chapter 4 (virtual addresses; the hardware translates every pointer and can refuse the access), Chapter 11 (a buffer overflow corrupts *this* process — but only this one), Chapter 12 (you became the allocator; the backing memory came from somewhere). · **Opens Part 3.** Part 2 left you holding, by hand, a proof about the memory *inside* one program. This part is about the operating system underneath all of it — the code that gave you an address space, a heap, and the illusion of owning the machine. We start at the seam where your program meets that code: the **system call**, the single controlled doorway from your program into the kernel.

BEFORE YOU READ ON

From memory, reaching back: (1) Chapter 11 — a stack buffer overflow can overwrite the saved return address and hijack *this* program's control flow. Can it, by the same mechanism, overwrite memory belonging to a *different* running program? What stops it? (2) Chapter 4 — when you dereference a pointer, the address is virtual and the hardware translates it; who set up the translation tables, and could your code just rewrite them to point anywhere in physical RAM? (3) *Predict:* when your C program calls `write()` to print to the screen, does *your* code push the bytes out to the terminal device — or does something else do it on your behalf? Guess what stops an ordinary program from writing straight to the disk controller. By the end you will be able to trace exactly what happens on that call.

Every useful program eventually needs to do something it is not allowed to do. Reading a file, sending a network packet, getting more memory from the OS, creating a process, drawing on the screen — all of these touch hardware or shared state that a single process must not control directly, because the whole point of an operating system is that one buggy or malicious program cannot seize the machine or read another program's secrets. So the machine runs your code in a restricted mode, and the only way to get privileged work done is to *ask*: to make a **system call**, a controlled transfer into the kernel that switches into privileged mode, runs a specific pre-approved handler, and returns. This chapter is about that boundary. We establish the two privilege modes and why they exist (§13.1), trace exactly how a system call crosses the boundary on this machine, register by register (§13.2), separate the C library's friendly wrappers from the kernel underneath them (§13.3), measure what a crossing actually costs and how to avoid paying it (§13.4), and carry the boundary forward into processes and threads (§13.5). This is the seam Chapter 12 promised you would see "from the other side."

13.1 Two modes: user and kernel

The CPU itself enforces the boundary. A modern processor runs in one of (at least) two **privilege modes**: **user mode**, in which most of your code runs, and **kernel mode** (also called supervisor or privileged mode), in which the operating system runs. On x86 these are implemented as *rings* — ring 3 for user code, ring 0 for the kernel. The difference is not speed; it is *authority*. In user mode a set of instructions simply will not execute: you cannot halt the CPU, disable interrupts, load the page tables that define address

translation (Chapter 4), or read and write device registers. Attempt one and the hardware traps into the kernel instead of obeying you.

That single mechanism is what makes the isolation of Chapter 11 real. Your buffer overflow can wreck *this* process because it runs with this process's authority — but it cannot reach another process's memory, because the memory-management unit (Chapter 4) translates every address through *this* process's page tables and faults on anything outside them, and it cannot rewrite those page tables to escape, because installing page tables is a privileged instruction it is not allowed to run. Two hardware mechanisms, working together: the MMU draws the boundary around each process's *memory*, and the user/kernel mode bit draws the boundary around the machine's *instructions and devices*. Everything you have written in this book so far has run in ring 3, with no direct access to any of it. The kernel is the only software that runs privileged, and it is the referee that hands out access under its own rules.

This is the answer to the opener's first two questions. A process cannot corrupt another process's memory because it is confined to its own translated address space and cannot touch the translation machinery; it cannot rewrite the page tables because that instruction traps. The confinement is not politeness — it is enforced by silicon on every instruction.

QUICK CHECK

Name one privileged operation that a program in user mode is *not* allowed to perform directly, and say what the hardware does if it tries anyway. (Answer after the next section.)

13.2 Crossing the boundary: the system-call mechanism

If user code cannot touch devices or memory tables, how does it ever read a file? It asks the kernel to do it, through a deliberately narrow door. The mechanism is a **trap**: a special instruction that simultaneously switches the CPU into kernel mode *and* jumps to a single, fixed entry point that the kernel installed at boot. This is the crucial safety property — user code does not get to switch into kernel mode and then run wherever it likes; the trap lands only at the kernel's chosen handler, which decides what to do. On 64-bit x86 Linux that instruction is `syscall` (the older 32-bit path used the software interrupt `int 0x80`).

Because the boundary is narrow, there is a strict convention — an **application binary interface (ABI)** — for how you name the call and pass its arguments. On x86-64 Linux (verified against the `syscall(2)` manual page): the **system-call number** goes in register `rax`; the first six arguments go in `rdi`, `rsi`, `rdx`, `r10`, `r8`, `r9` (note `r10`, not `rcx` — the `syscall` instruction itself clobbers `rcx` and `r11`); you execute `syscall`; and the kernel returns the result in `rax`. A negative return in the range of a small error code means failure — the raw kernel interface returns `-errno` (§13.3 turns that into the friendly `errno` you know). The kernel also *validates* what you pass: hand it a pointer to write into, and it checks that the pointer really belongs to your address space, so you cannot trick it into scribbling on kernel memory on your behalf.

A system call is a controlled trap: one fixed door from ring 3 into ring 0 and back.




Figure 13.1: A system call. User code loads the call number and arguments into the ABI's registers and executes `syscall`; the CPU switches to kernel mode and jumps to one fixed entry point (never an address of the caller's choosing); the kernel validates, does the privileged work, and returns the result in `rax`. The boundary is narrow by design — that narrowness is the security.

Here is the whole path made concrete. When a program runs `write(1, "hi\n", 3)` — write three bytes to file descriptor 1, standard output — the sequence is: put the number for `write` in `rax`, put 1 in `rdi`, the buffer address in `rsi`, and 3 in `rdx`; execute `syscall`; the CPU traps to the kernel, which writes the bytes and returns 3 (the count written) in `rax`. We can watch this happen with `strace`, which intercepts and prints every system call a program makes. Compiling a three-line C program that does exactly that and running it under `strace` on this machine (gcc 11.4.0, x86-64 Linux) shows the crossing directly:

```
$ strace -e trace=write ./hello
write(1, "hi\n", 3)           = 3
+++ exited with 0 +++
```

That one line is the boundary: the arguments the program marshaled, and the `= 3` the kernel returned. Everything else the program did — computing, moving data in its own memory — never appears, because none of it crossed into the kernel.

(Answer to the §13.1 quick check: a user-mode program may not execute privileged instructions such as halting the CPU, disabling interrupts, loading the page tables, or accessing device registers directly. If it tries, the hardware does not obey — it traps into the kernel, which typically terminates the offending process.)

13.3 The library wrapper vs. the raw system call

You have almost certainly never written `syscall` by hand, and you rarely will. The C library (glibc on most Linux systems) provides an ordinary-looking function for each system call — `write`, `read`, `open`, `fork` — and that **wrapper** does the register marshaling for you: it moves your arguments into `rdi/rsi/rdx`, sets `rax`, executes `syscall`, and inspects the result. On the negative-return failure convention from §13.2, the wrapper does the translation you already know: it stores the error code in the global `errno` and returns `-1`. So the raw kernel interface returns `-EBADF`; the wrapper you call returns `-1` and sets `errno == EBADF`. The wrapper is a thin shell around one `syscall`.

The distinction that trips people up is that *not every library function is a wrapper*, and the ones that matter most are not. `printf` is **not** a system call. It is a substantial C library routine that parses your format string, converts numbers to text, and — critically — *buffers* the result in user-space memory, only calling the `write` system call later, when the buffer fills or is flushed. Similarly `malloc` (Chapters 9 and 12) is a library function that manages a heap in user space; it crosses into the kernel only occasionally, calling the `brk` or `mmap` system calls when it needs to grow the heap — which is exactly where your arena's backing block came from in Chapter 12. This is the "other side" that chapter promised: your allocator handed out bytes that the C library had obtained, in bulk, from the kernel across this very boundary. The layering is real and worth holding precisely: *your code* calls a *library function*, which may buffer and compute in user space, and which crosses into the *kernel* only when it must, through a *system-call wrapper*.

QUICK CHECK

You run a program that prints with `printf`, then crashes before returning from `main` — and the printed line never appears in your terminal. Nothing was wrong with the format string. What most likely happened to those bytes? (Answer below the communicate box.)

13.4 The cost of crossing — and how to avoid paying it

Crossing the boundary is not free, and the cost is the reason a great deal of systems performance work exists. A system call must switch privilege mode, save and later restore the caller's state, enter the kernel, validate arguments, and eventually switch back — and along the way it can disturb the CPU's caches and TLB (Chapters 3-4). Compared with an ordinary function call, which on a modern core is a couple of instructions, that is enormously more expensive. Measured on this machine (gcc 11.4.0, -O2, x86-64), calling a trivial non-inlined function ten million times versus making the `getpid` system call the same number of times:

```
plain function call : ~1-2 ns each
getpid system call  : ~150-190 ns each
ratio                : roughly two orders of magnitude
```

These figures are illustrative, from this specific machine and run — they vary with hardware, kernel, and mitigations (some CPU-vulnerability mitigations make the boundary crossing markedly more expensive), and you should measure your own. But the shape is robust and the lesson is durable: a system call costs on the order of a hundred function calls or more. Three consequences follow, and each is a technique you will use:

- **Batch — do not cross per item.** Reading a megabyte one byte at a time is a million crossings; reading it in large blocks is a handful. This is exactly why the C library buffers: `printf` and `fread` accumulate in user space and cross the boundary rarely. The single most common I/O performance bug is an unbuffered per-item system call in a hot loop.

- **Let the kernel share read-only data without a crossing — the vDSO.** A few "system calls" that only *read* kernel-maintained data, such as `clock_gettime` and `gettimeofday`, are served from the **vDSO** (virtual dynamic shared object): a page the kernel maps read-only into every process, containing the data and the code to read it, so the call runs entirely in user mode with no trap. On this machine, a program that calls `clock_gettime` one million times produces, under `strace -c`, *zero* `clock_gettime` system calls — the crossings simply are not there, because the vDSO answered in user space. (You can see the vDSO mapped into any process: `ldd` on a dynamically linked program lists `linux-vdso.so.1`.)
- **Observe the boundary with `strace`.** Because every crossing is visible, `strace` (and `strace -c` for a counted summary) is the first tool to reach for when a program is mysteriously slow or stuck: it shows you exactly which system calls it is making, how many, and where it is blocked.

THE SAME IDEA THREE WAYS: A SYSTEM CALL IS A CONTROLLED TRAP

1. **In words.** User code cannot touch devices or memory tables, so it *asks* the kernel by executing a trap instruction that switches to kernel mode and jumps to one fixed handler — never an address of the caller's choosing. The kernel validates, does the privileged work, and returns. Narrowness is the security.
2. **As a picture.** Figure 13.1: two stacked bands (ring 3, ring 0) with a single arrow down (`syscall` → fixed entry) and a single arrow up (return in `rax`).
3. **As numbers.** The ABI: number in `rax`; args in `rdi,rsi,rdx,r10,r8,r9`; result in `rax`; `-errno` on failure. The cost: ~150-190 ns for `getpid` versus ~1-2 ns for a function call on this machine — roughly 100×, which is why you batch and why the vDSO exists.

TRAP: "PRINTF IS A SYSTEM CALL" — CONFUSING THE LIBRARY WITH THE KERNEL

It is tempting to treat every function that produces output as a direct line to the operating system. **It is not: `printf` is a C library function that formats and buffers in user-space memory, and calls the write system call only when that buffer flushes.** Conflating the two produces real, confusing bugs. A program that `printfs` a line and then crashes may show *nothing* in the terminal, because the bytes were sitting in the library's buffer and the crash skipped the flush — the boundary was never crossed (this is the §13.3 quick check, and it returns in Chapter 14 when `fork` *duplicates* that unflushed buffer and a line prints twice). The fix when you need output to be real *now* is to force the crossing: `fflush(stdout)`, or configure line buffering, or call `write` directly — which is unbuffered and goes straight across. The mental model that prevents a dozen bugs: *your code* → *a library function (may buffer)* → *a system-call wrapper* → *the kernel*. Know which layer you are on. The library's job is often to *avoid* the kernel; the kernel is only reached deliberately.

SAY IT LIKE A SYSTEMS ENGINEER

Try it first. A teammate reports: "Our service writes one log line per request with a bare `write()` call, and profiling says it is spending huge time in the kernel. Let's just delete the logging." Cover the paragraph and respond in three or four sentences, drawing on §13.3–13.4.

Model answer. "The problem is almost certainly not *that* we log, but *how* — one unbuffered write per line is one system-call crossing per line, and a crossing costs on the order of a hundred function calls or more (on our box `getpid` was ~150 ns versus ~1-2 ns for a call), so at high request rates those crossings dominate. The fix is to batch: buffer log lines in user space and flush in large blocks, which is exactly why the C library buffers `printf` in the first place; we cross the boundary a few times instead of thousands. Let's confirm with `strace -c` that `write` is the hot call, switch to a buffered logger, and re-measure — deleting logging throws away observability to fix a batching bug." Naming the crossing cost, reaching for batching rather than removal, and confirming with `strace` is the senior register.

(Answer to the §13.3 quick check: the bytes were almost certainly still in `printf`'s user-space buffer. Because `printf` buffers and only crosses into the kernel via `write` on a flush, a crash before the flush loses them — they never reached the terminal. Force the crossing with `fflush`, line buffering, or a direct write.)

13.5 What to carry forward

The CPU runs your code in **user mode** (ring 3), where privileged instructions — touching devices, loading page tables, disabling interrupts — simply will not execute, and the operating system in **kernel mode** (ring 0) is the only software with that authority. Together the MMU (Chapter 4) and this mode bit are what make process isolation real: a bug confined to your address space cannot reach another process or the machine. To get privileged work done, user code makes a **system call** — a controlled trap (`syscall` on x86-64) that switches mode and jumps to *one fixed kernel entry point*, passing a call number in `rax` and arguments in `rdi,rsi,rdx,r10,r8,r9`, receiving a result in `rax` (with `-errno` on failure). The C library wraps each call in an ordinary function and translates the error convention into `errno` and `-1` — but many library functions (`printf`, `malloc`) are *not* system calls; they compute and buffer in user space and cross only when they must. And the crossing is expensive — on this machine roughly two orders of magnitude more than a function call — so you batch, you lean on the vDSO for cheap read-only data like the time, and you observe the boundary with `strace`.

This is the seam the rest of Part 3 lives on. A **process** (Chapter 14) is the isolated container this boundary protects — its own address space, created and destroyed through system calls (`fork`, `execve`, `exit`, `wait`) that cross exactly this line. A **thread** (Chapter 15) is a second flow of control that shares one side of the boundary — the same address space — trading isolation for cheap sharing. Every abstraction the OS gives you, from the heap that fed your allocator to the files and sockets you will read later, is reached through the doorway you have just learned to trace.

CAN YOU... WITHOUT LOOKING?

- Name the two CPU privilege modes and give two operations user mode is forbidden to do — and say what enforces the ban.
- Explain why a system call jumps to a *fixed* kernel entry point rather than an address the caller chooses, and why that matters.
- State the x86-64 Linux system-call ABI: where the call number, the arguments, and the return value go.
- Explain how `printf` differs from the `write` system call, and why a crash can lose a `printf`ed line.
- Give the rough cost of a system call versus a function call on a machine you have measured, and name two ways to cross the boundary less (batching, the vDSO).
- Say how process isolation follows from the MMU plus the user/kernel mode bit.

RETRIEVAL — CLOSED BOOK

Cover the chapter. Answer from memory; rate confidence first.

1. **R1.** What is the difference between user mode and kernel mode, and what enforces it? (§13.1)
2. **R2.** Describe the system-call mechanism on x86-64 Linux: instruction, where the number and arguments go, where the result comes back. (§13.2)
3. **R3.** Is `printf` a system call? What actually crosses the boundary, and when? (§13.3)
4. **R4.** Roughly how much more expensive is a system call than a function call, and why? (§13.4)
5. **R5.** What is the vDSO, and why does calling `clock_gettime` often *not* show up as a system call? (§13.4)

FURTHER READING

- **Arpaci-Dusseau & Arpaci-Dusseau, *Operating Systems: Three Easy Pieces (OSTEP)*, "Mechanism: Limited Direct Execution."** The free chapter that develops exactly this boundary — user vs. kernel mode, traps, and the system-call mechanism. The reference for §13.1–13.2.
- **Bryant & O'Hallaron, *Computer Systems: A Programmer's Perspective (CS:APP)*, Chapter 8.** Exceptional control flow: how the hardware and OS handle traps and system calls from the programmer's view.
- **The Linux `syscall(2)` and `vdso(7)` manual pages.** The primary sources for the per-architecture calling convention of §13.2 and the vDSO of §13.4. Verify against your kernel and architecture.
- **`strace` documentation and man page.** The tool used throughout this chapter to make the boundary visible; `strace -c` for a counted summary. Free.

Exercises

13.1 **WARM-UP** Name the two CPU privilege modes, and give one operation a program in user mode is not allowed to perform directly.

13.2 **WARM-UP** On x86-64 Linux, which register holds the system-call number, and which register holds the return value?

13.3 **WARM-UP** Is `printf` a system call? If not, what does it do first, and which system call does it eventually use to produce output?

13.4 **CORE** Trace, step by step, what happens when a program executes `write(1, buf, 5)`: what does the C library put in which registers, what instruction crosses the boundary, and what comes back on success? (Use the ABI from §13.2.)

13.5 **CORE** Explain the relationship between the kernel's raw error convention and the `errno/-1` you see from a C library call. If a raw system call returns `-EBADF`, what does the wrapper return and what does it set?

13.6 **CORE** You run `strace` on a program and see, among other lines: `openat(...) = 3`, `read(3, ..., 4096) = 512`, and then thousands of `read(3, ..., 1) = 1` lines. In one or two sentences, say what the last pattern tells you is wrong and what to change.

13.7 **COST** (a) Name two reasons a system call costs far more than an ordinary function call. (b) A program reads a 1 MB file by calling `read` one byte at a time; explain why this is slow in terms of boundary crossings, and give the fix. (c) Why does calling `clock_gettime` in a tight loop often cost *no* system calls at all?

13.8 **FADED** *Worked, with the last step for you.* A program prints a diagnostic line with `printf` just before dereferencing a bad pointer and crashing. The line never appears in the terminal, even though the program definitely reached the `printf`.

Step 1 (given): `printf` writes into the C library's user-space output buffer, not directly to the terminal.

Step 2 (given): The buffer is flushed (crossed to the kernel via `write`) only when it fills, at a newline for a line-buffered terminal, or at normal program exit.

Step 3 (given): A crash from an invalid dereference terminates the process abnormally, skipping the normal flush.

Step 4 (you): State where the bytes actually were when the program died, explain why the line was lost (referencing the boundary), and give two ways to make sure the line appears before a possible crash.

13.9 **MIXED** For each, decide whether it requires crossing the system-call boundary or is pure user-space computation, and say why in a few words: (a) adding two integers in registers; (b) opening a file; (c) sorting an in-memory array; (d) sending a network packet; (e) formatting a number into a string with `sprintf`.

13.10 **MIXED** Drawing on Chapters 10–13, classify each as *undefined behavior*, a *boundary crossing (system call)*, or *neither*: (a) reading one past the end of a stack array; (b) calling `read` on a valid file descriptor; (c) signed integer overflow in a loop counter; (d) a buffered `printf` that has not yet flushed.

13.11 **STRETCH** The kernel validates every pointer you pass to a system call such as `write`, checking it belongs to your address space. (a) Relate this to Chapter 4's address translation: why can a process not simply pass a pointer to another process's memory or to kernel memory and have the kernel read it? (b) Sketch what could go wrong for security if the kernel skipped this validation. (c) Explain how the two mechanisms — per-process address translation and the user/kernel mode bit — together make one process unable to spy on another.

13.12 **LAB** On your own machine: (a) write a program that prints with `write` and run it under `strace -e trace=write`; confirm you see the `write` call and its return value. (b) With `strace -c`, count the total system calls of a trivial "hello world" and note how many are startup overhead (dynamic linking, etc.) versus your own output. (c) Measure the cost of the `getpid` system call versus a trivial function call in a loop, and (d) call `clock_gettime` a million times and confirm with `strace -c` that it does *not* appear as a million system calls. Report your numbers and label your compiler, kernel, and hardware, since figures vary.

Chapter 14 — Processes: fork, exec, and wait

Before you start: Chapter 13 (the system-call boundary — `fork`, `exec`, and `wait` are all system calls), Chapter 4 (each process has its own virtual address space, translated through its own page tables), Chapter 11 (a memory bug is confined to *one* process's address space). · Pairs with Chapter 15 (threads share an address space instead of copying it). Chapter 13 showed the boundary that protects a process; this chapter is about the process itself — what it is, how the OS makes a new one, how one program becomes another, and how a parent collects a finished child.

BEFORE YOU READ ON

From memory, reaching back: (1) Chapter 13 — to do something privileged, user code must do what, and through which instruction on x86-64? (2) Chapter 11 — a buffer overflow can corrupt the process it runs in; can it corrupt a *different* running process's memory, and what two mechanisms decide the answer? (3) *Predict:* `fork()` makes a new process that is a copy of the calling one. If your program is holding a 1 GB array in memory and calls `fork()`, do you think the operating system immediately copies all 1 GB? Guess, and note why the answer matters for how often you can afford to `fork`. We measure the idea behind it shortly.

A running program is not the same thing as the file on disk. The file is inert bytes; the running instance — with its own memory, its own place in the CPU, its own open files — is a **process**, and the process is the operating system's fundamental unit of isolation and accounting. Everything Chapter 13 protected was protecting a process: its private address space, its authority level, its resources. This chapter opens the process up. We define precisely what a process bundles together (§14.1); we build a new one with the Unix system call `fork`, and see why it returns *twice* and why copying a process can be cheap (§14.2); we turn a copied process into a *different* program with `exec`, and meet the fork-then-exec pattern that every shell uses (§14.3); we collect a finished child with `wait`, and meet the zombie and the orphan (§14.4); and we carry the model into threads (§14.5). By the end you will be able to write, and reason about, the create-run-reap lifecycle that underlies every command you have ever typed.

14.1 What a process is — the unit of isolation

A process is three things bundled and protected together. First, an **address space**: the private virtual memory of Chapters 4, 7, and 8 — text, initialized data, bss, heap, and stack — that the MMU translates through this process's own page tables so that its pointers name only its own memory. Second, one or more **threads of execution**: the actual running context — the program counter, the stack pointer, the register values — that says *where* in the code the process currently is (this chapter's processes have exactly one; Chapter 15 adds more). Third, a bundle of **kernel-managed resources**: the open file descriptors, the process ID (PID) and parent PID, the current working directory, the user and group credentials, and more, all tracked by the kernel on the process's behalf.

The reason these are bundled is **isolation**. The kernel gives each process the illusion of owning the machine: its own address space (so one process's pointers are meaningless in another's — the answer to Chapter 11's confinement), its own turn on the CPU, its own resource handles. A fault in one process — a wild pointer, a crash, even a security compromise — is contained to that process's address space and cannot, by the mechanisms of Chapter 13, corrupt another process or the kernel. This is why the process is the natural boundary for a sandbox, a multi-tenant service, or any place you need one failure not to become everyone's failure. The process is precisely the thing the system-call boundary was built to protect.

QUICK CHECK

Name the three things a process bundles together. Which one makes two processes' identical-looking pointer values refer to completely different memory? (Answer after the next section.)

14.2 Creating a process: `fork` and copy-on-write

Unix creates a new process in a way that surprises everyone the first time: the `fork()` system call is called *once* but returns *twice*. It makes the calling process (the **parent**) into two nearly identical processes — the parent and a new **child** — and then both continue from the point of the return. The child is a copy: same code, same current values of every variable, copies of the open file descriptors. The one difference is the return value, which is how each copy knows who it is: `fork` returns **0 in the child** and **the child's PID in the parent** (or `-1` on failure). You branch on it. Here is the whole pattern, compiled and run on this machine (gcc 11.4.0, x86-64 Linux):

```
int x = 100;
pid_t pid = fork();           /* one call, returns TWICE */
if (pid == 0) {               /* child sees 0 */
    x += 1;
    printf("child : pid=%d ppid=%d x=%d (fork returned %d)\n", getpid(), getppid(), x, pid);
} else {                      /* parent sees the child's pid */
    int st; waitpid(pid, &st, 0);
    x += 10;
    printf("parent: pid=%d x=%d (fork returned %d; child exit=%d)\n", getpid(), x, pid,
WEXITSTATUS(st));
}
```

```
child : pid=1341934 ppid=1341933 x=101 (fork returned 0)
parent: pid=1341933 x=110 (fork returned 1341934; child exit=0)
```


Read that output closely, because it shows every essential fact. The two processes have different PIDs (1341934 and 1341933), and the child's parent PID equals the parent's PID — the family tree is real. Both started with `x == 100`, but the child's `x += 1` and the parent's `x += 10` did not interfere: the child ends at 101 and the parent at 110. Their memory is *independent* — the same variable name, at what is even the same virtual address, but two separate physical copies, because each process has its own address space (§14.1). And `fork` returned 0 down one path and the child's PID down the other, which is the only thing distinguishing the twins.

Now the predict from the opener: if the parent held a 1 GB array, did `fork` copy a gigabyte? No — and it could not afford to, or forking would be unusable. The kernel uses **copy-on-write (COW)**: rather than duplicating the parent's pages, it maps the *same* physical pages into both processes and marks them read-only. Both share the physical memory as long as they only read it. The instant *either* one writes to a page, the hardware faults, the kernel makes a private copy of just that one page for the writer, and both continue. So `fork`'s cost is proportional not to the size of the address space but to the number of pages actually modified afterward — which is why forking a huge process is cheap, and why the child's `x += 1` above triggered a copy of exactly the one page holding `x` and nothing else. Copy-on-write is Chapter 4's page tables doing double duty: the same translation-and-fault machinery that isolates processes also lets them share until the moment sharing becomes incorrect.

THE SAME IDEA THREE WAYS: FORK RETURNS TWICE INTO TWO ADDRESS SPACES

- 1. In words.** One call, two returns: the parent gets the child's PID, the child gets 0, and from then on they are separate processes with separate (copy-on-write) memory. You branch on the return value to run different code in each.
- 2. As a picture.** Figure 14.1: a single arrow entering `fork` and two arrows leaving, into two boxes that start as copies and diverge as each writes its own pages.
- 3. As a trace.** Before: one process, `x = 100`. After `fork`: two processes, both `x = 100`, distinguished by return value (child 0, parent child-PID). After each runs: child `x = 101`, parent `x = 110` — independent, exactly as the measured output shows.

`fork()`: one call, two returns, two independent (copy-on-write) address spaces.



The only difference between the twins is `fork()`'s return value — that is how each knows which one it is.

Memory is independent afterward: the child's `x += 1` and the parent's `x += 10` touch separate physical pages.

Figure 14.1: `fork` returns twice — 0 in the child, the child's PID in the parent — producing two processes that start as copies and, thanks to copy-on-write, share physical pages until one writes. The branch on the return value is how each copy runs different code.

(Answer to the §14.1 quick check: a process bundles an **address space**, one or more **threads of execution**, and a set of **kernel resources** such as open file descriptors and its PID. The **address space** — separate per process and translated through its own page tables — is what makes two processes' identical-looking pointer values refer to different memory.)

14.3 Running a new program: `exec` and the fork-exec pattern

`fork` gives you a copy of the *same* program. To run a *different* program you need `exec` — the `execve` system call and its convenience wrappers. `exec` does something drastic: it **replaces the current process image** — text, data, heap, stack — with a brand-new program loaded from disk, and jumps to its entry point. Everything the process was running is gone; the new program starts fresh. Two things survive the replacement, and they are the important ones: the **PID stays the same** (it is the same process, now running different code), and the **open file descriptors stay open** (unless marked close-on-exec). On success, `exec` *does not return* — there is nothing to return to, because the code that called it no longer exists.

Put the two together and you get the pattern behind every command you run: **fork, then exec, then wait**. To launch a program, a process forks; the child `execs` the new program (replacing itself with it); and the parent `waits` for the child to finish. This is exactly what a shell does for every command you type — `fork` a child, have it `exec /bin/ls`, and wait for it to complete before printing the next prompt — and it is how your terminal launched the demo program in §14.2. The pattern is also why the parent in that demo called `waitpid` before continuing.

Why split creation (`fork`) from loading (`exec`) into two calls, when almost every use runs them back to back? Because the *gap between them* is where the elegance lives. After the child is forked but before it `execs`, it is still running the parent's code and can adjust its inherited resources: it can redirect standard output to a file, wire its input or output to a pipe connected to a sibling, drop privileges, or change directory — and then `exec` the target program, which starts up already plumbed into that environment without knowing anything about it. Unix I/O redirection and pipelines (`ls | grep foo > out.txt`) are built entirely in that fork-to-exec gap. The two-call design is not clumsiness; it is the seam that makes composition possible.

QUICK CHECK

A shell runs the command `ls`. In order, which of `fork`, `exec`, and `wait` does the shell process perform, which does the new child process perform, and which call does *not* return on success? (Answer below the communicate box.)

14.4 Reaping: `wait`, exit status, zombies, and orphans

A process ends by calling `exit` (or returning from `main`, or being killed). But it does not vanish the instant it ends. The kernel must keep a small amount of information about it — most importantly its **exit status**, the number it passed to `exit` — so that the parent can find out how the child fared. A process that has terminated but whose exit status has not yet been collected is a **zombie** (shown as `<defunct>` in `ps`): it holds no memory or CPU, just a slot in the process table and its exit status, waiting to be read.

The parent collects it by calling `wait` or `waitpid`, which blocks until a child finishes (or returns immediately if one already has), hands back the child's PID and its exit status (decode it with macros like `WEXITSTATUS`), and lets the kernel finally free the zombie's last

slot. This is **reaping**. It is the parent's responsibility, and forgetting it is a real defect: a long-running parent that `forks` children and never `waits` accumulates zombies, slowly filling the process table. Note the category carefully — a leaked zombie is a *defined* resource leak, not undefined behavior; the program is not doing anything the language forbids, it is just failing to clean up, exactly as a memory leak is defined-but-bad (Chapter 9).

The mirror-image case is the **orphan**: a child whose parent exits *first*. An orphan is not left dangling — the kernel **reparents** it to `init` (PID 1, or a subreaper), which `waits` for it when it finishes, so orphans are reaped automatically. The lifecycle, then, is: a process is `forked`, runs, `exits` (becoming a zombie holding its status), and is reaped by a waiting parent (or by `init` if orphaned). Create, run, exit, reap.

TRAP: FORK DUPLICATES YOUR UNFLUSHED OUTPUT BUFFER (THE DOUBLE-PRINT)

Chapter 13 warned that `printf` buffers in user-space memory and crosses to the kernel only on a flush. `fork` copies the *entire* address space — **including that buffer with its unflushed contents**. So if you `printf` a line and *then* `fork`, the un-crossed bytes sit in both the parent's and the child's copy of the buffer, and when each process later flushes (at exit), the line is written *twice* — once per process — even though your code printed it once. This is not hypothetical: constructing exactly this (a `printf` before the `fork`, with output going to a pipe so buffering is block-mode) reproduces the doubled line every time. The trap is subtle because it disappears when output goes to a terminal (line-buffered, so the newline already flushed before the `fork`) and reappears when output is redirected to a file or pipe. The fix is the same as Chapter 13's: force the crossing before you branch — `fflush(stdout)` before `fork` (or use unbuffered write). The deeper lesson: `fork` copies *everything*, including the parts of your program's state you had forgotten were state.

SAY IT LIKE A SYSTEMS ENGINEER

Try it first. A teammate proposes: "Let's handle each incoming request by forking a fresh process — that way one bad request can't take down the others." Cover the paragraph and respond in three or four sentences, drawing on §14.1–14.2.

Model answer. "That is a genuinely good instinct for *isolation*: a separate process has its own address space, so a crash or memory corruption in one request is contained and can't touch a sibling — it's the classic robustness pattern, and `fork` is cheaper than it looks because copy-on-write means we don't duplicate the whole address space, only the pages a child actually writes. The costs to weigh are that processes don't share memory, so passing data between them needs explicit IPC (pipes, shared memory) rather than a pointer, and each request pays process-creation and context-switch overhead, which can hurt at very high rates. So process-per-request is the right call when fault-containment or a security boundary is the priority; if instead we need cheap shared-memory concurrency and can enforce discipline, threads (Chapter 15) share the address space and skip the copy — at the cost of exactly that isolation. Let's decide based on whether containment or shared-state throughput is what this service needs, and measure." Naming copy-on-write, the isolation-vs-sharing trade, and deferring to a measurement is the senior register.

(Answer to the §14.3 quick check: the shell process performs `fork` and then `wait`; the new child process performs `exec` (replacing itself with `ls`). `exec` is the one that does not return on success — on success there is no original program left to return to.)

14.5 What to carry forward

A **process** is the operating system's unit of isolation: an **address space** (private virtual memory, Chapter 4), one or more **threads of execution** (the running context), and a bundle of **kernel resources** (open files, PID, credentials), all protected by the boundary of Chapter 13. You create one with `fork`, which is called once and **returns twice** — 0 in the child, the child's PID in the parent — producing two processes that start as copies and diverge; copying is cheap because of **copy-on-write**, which shares physical pages until one side writes, so forking a huge process costs only the pages later modified. You turn a copied process into a different program with `exec`, which **replaces the image** (keeping the PID and open descriptors) and does not return on success; the universal **fork-exec-wait** pattern — with the fork-to-exec gap used to set up redirection and pipes — is how every shell runs every command. And you clean up with `wait`, which reaps a finished child's exit status; an unreaped terminated child is a **zombie** (a defined leak), while an orphaned child is reparented to `init` and reaped for you.

Every process in this chapter had a single thread of execution — one program counter walking through the code. Chapter 15 relaxes that. A **thread** is a second (third, hundredth) flow of control *within one process*, sharing its address space rather than copying it as `fork` does. That one change — share instead of copy — buys cheap concurrency and drops the isolation this chapter spent its length building: the very memory that `fork` kept separate, threads deliberately share, and the proof burden of Part 2 returns, now with multiple threads reaching for the same bytes at once.

CAN YOU... WITHOUT LOOKING?

- State the three things a process bundles together, and which one enforces memory isolation between processes.
- Explain why `fork` returns twice, what it returns to each side, and how a program uses that to run different code in parent and child.
- Explain copy-on-write and why forking a 1 GB process does not copy 1 GB.
- Describe what `exec` replaces and what it keeps, and why it does not return on success.
- Lay out the fork-exec-wait pattern and say what the gap between `fork` and `exec` is used for.
- Define a zombie and an orphan, say who reaps each, and classify a leaked zombie (UB or defined defect?).

RETRIEVAL — CLOSED BOOK

Cover the chapter. Answer from memory; rate confidence first.

1. **R1.** What three things make up a process, and what does the address space give it? (§14.1)
2. **R2.** What does `fork` return in the parent and in the child, and why is that the only difference between them? (§14.2)
3. **R3.** What is copy-on-write, and why does it make `fork` cheap? (§14.2)
4. **R4.** What does `exec` do, what survives it, and does it return on success? (§14.3)
5. **R5.** What is a zombie, what is an orphan, and who reaps each? (§14.4)

FURTHER READING

- **Arpaci-Dusseau & Arpaci-Dusseau, *Operating Systems: Three Easy Pieces (OSTEP)*, "Abstraction: The Process" and "Interlude: Process API."** The free chapters that build exactly this material — the process abstraction and the `fork/exec/wait` API. The reference for §14.1–14.4.
- **Bryant & O'Hallaron, *Computer Systems: A Programmer's Perspective (CS:APP)*, Chapter 8 (Exceptional Control Flow).** Processes, process creation, and reaping children, from the programmer's view.
- **The Linux `fork(2)`, `execve(2)`, and `wait(2)` manual pages.** Primary sources for the exact semantics, including copy-on-write, the fate of file descriptors across `exec`, and zombie/orphan behavior. Verify per system.
- **Beej's Guide to Unix IPC / C.** Free, practical walkthroughs of `fork/exec` and inter-process communication (pipes) — the plumbing that lives in the `fork-to-exec` gap.

Exercises

- 14.1** **WARM-UP** Name the three things a process bundles together.
- 14.2** **WARM-UP** What does `fork` return in the parent process, and what does it return in the child?
- 14.3** **WARM-UP** Does `exec` return on success? What does it replace, and name one thing it keeps.
- 14.4** **CORE** Write, from memory, the `fork/if-else` skeleton that distinguishes child from parent, with the parent calling `waitpid`. Explain in a sentence what each branch is and how each process knows which it is.
- 14.5** **CORE** Explain copy-on-write: what does the kernel actually do to the parent's pages at `fork` time, what triggers a real copy, and why does this make forking a 1 GB process fast?
- 14.6** **CORE** Describe the `fork-exec-wait` pattern a shell uses to run a command. Then explain what useful work can happen in the gap between `fork` and `exec`, and give one concrete example (e.g. I/O redirection).

14.7 **COST** (a) Explain why `fork` is cheaper than the size of the address space would suggest — name the mechanism. (b) Describe a workload for which copy-on-write buys you almost nothing (the child writes to nearly all inherited pages) and say why. (c) Compare process-per-request and thread-per-request on one cost each pays that the other does not.

14.8 **FADED** *Worked, with the last step for you.* A long-running server forks a worker for each connection and never calls `wait`. After hours, `ps` shows hundreds of `<defunct>` entries and the server eventually cannot create new processes.

Step 1 (given): When a child exits, the kernel keeps its exit status until the parent collects it, leaving the child as a zombie.

Step 2 (given): This server never calls `wait/waitpid`, so no child's status is ever collected.

Step 3 (given): Each zombie holds a process-table slot, and the table is finite.

Step 4 (you): Name the accumulating state, explain why it eventually stops new forks from succeeding, say whether this is undefined behavior or a defined resource leak, and give a fix (how the parent should reap, mentioning `wait/waitpid` or handling `SIGCHLD`).

14.9 **MIXED** For each goal, decide whether you need `fork`, `exec`, or `fork` then `exec`, and say why in a line: (a) run `/bin/ls` from inside your program and return to your program afterward; (b) create a second process that runs the *same* code on the second half of an array; (c) a shell running a command the user just typed.

14.10 **MIXED** Drawing on Chapters 9–14, classify each as *undefined behavior*, *a defined-but-real defect*, or *well-defined and fine*: (a) after `fork`, the parent expects to see changes the child made to `x` in the child's memory; (b) a server that never reaps its children; (c) both parent and child printing a line that was `printfed` (unflushed) before the `fork`; (d) a child that `execs` a new program while the parent keeps running.

14.11 **STRETCH** Tie process isolation back to Chapters 11 and 13. (a) Explain why a buffer overflow in process A cannot corrupt process B's memory, naming the two hardware mechanisms (Chapters 4 and 13) that enforce it. (b) Explain why this makes a separate process, not a thread, the right unit for sandboxing untrusted code. (c) In one or two sentences, describe what a multi-tenant service gains by putting each tenant's work in its own process rather than sharing one.

14.12 **LAB** On your own machine: (a) write the `fork` demo of §14.2 and confirm from its output that `fork` returned twice, that the child's `ppid` equals the parent's `pid`, and that the two processes' copies of a variable are independent. (b) Add a `printf` of some line *before* the `fork`, without an `fflush`, and run the program with its output redirected to a file; observe whether the line appears once or twice and explain why. (c) Add `fflush(stdout)` before the `fork` and confirm the fix. Report your output and your explanation; note your compiler and system.

Chapter 15 — Threads: Shared Memory and Its Hazards

Before you start: Chapter 14 (a process copies its address space; threads will share it), Chapter 13 (the system-call boundary — `clone` is the syscall underneath thread creation), Chapter 12 (in C you carry, by hand, a proof about memory — that proof is about to get a new clause). · Chapter 14's processes each had one thread of execution. This chapter adds more threads to a single process — a cheaper concurrency than `fork`, bought by giving up the isolation that made processes safe.

BEFORE YOU READ ON

From memory, reaching back: (1) Chapter 13 — a system call is a controlled trap; name the instruction on x86-64 and say where the system-call number goes. (2) Chapter 14 — `fork` gives the child its own address space (via copy-on-write), so when the parent and child both hold a variable `x`, are they the same storage or two copies? (3) *Predict*: two threads share a single counter and each executes `counter++` one million times, with no locking. When both finish, is the final value exactly two million, less than two million, or possibly more? Write down your guess and your reasoning — we will run this exact program and measure the answer.

Chapter 14 gave you a way to do two things at once — `fork` a second process — but at the price of a second, isolated address space: the processes cannot see each other's memory, so sharing data means explicit inter-process communication. Often that isolation is exactly what you want. But often you want the opposite: several flows of control working on the *same* data structures in the *same* memory, cheaply. That is a **thread**. This chapter defines a thread as a second flow of control inside one address space (§15.1), creates threads with the POSIX `pthread`s API (§15.2), explains how the kernel actually schedules them and why Linux uses a one-to-one model (§15.3), and then confronts the hazard that sharing buys you — the **data race**, which we will reproduce and measure losing more than half of a program's work (§15.4). We close by carrying the new proof burden forward toward the concurrency part that makes shared memory correct (§15.5). Threads make sharing trivial; this chapter is about why that is both the point and the danger.

15.1 A thread is a second flow of control in one address space

A process can contain more than one **thread** of execution, and the defining fact of threads is *what they share and what they don't*. All threads in a process **share** one address space: the same code (text), the same heap, the same global and static variables, and the same open file descriptors. What each thread keeps **private** is only what it needs to be an independent flow of control: its own **stack**, its own **registers** (including the program counter and stack pointer), and its own thread-local storage. That is the whole model. Two threads in one process see the same `malloc'd` object at the same address and can both read and write it — directly, through an ordinary pointer, with no copy and no IPC.

Set this against Chapter 14 precisely, because the contrast *is* the concept. When you `fork`, the child gets a *copy* of the address space (copy-on-write), so the parent's and child's variables are separate storage — writes in one are invisible to the other, and that isolation is what contains a crash. When you create a thread, the new thread gets the *same* address space, so a write by one thread is immediately visible to the others — and a wild pointer or a crash in one thread corrupts or kills the *whole* process, all threads with it. Process is copy-and-isolate; thread is share-and-expose. The power of threads (cheap communication through shared memory) and their danger (no isolation, and the races of §15.4) are the same fact viewed twice.

One process, two threads: shared address space, private stacks and registers.




Figure 15.1: Threads in one process share the address space — code, heap, globals, file descriptors — while each keeps a private stack, registers, and thread-local storage. Compare Chapter 14's `fork`, which *copies* the address space. Sharing is why threads communicate cheaply and why they race.

QUICK CHECK

Name two things all threads in a process share, and two things each thread keeps private. If thread A stores 7 into a global variable, can thread B see it — and how does that differ from a parent and child after `fork`? (Answer after the next section.)

15.2 Creating threads: the `pthread` model

On Unix systems the portable interface is **POSIX threads** (`pthread`). Two calls carry most of the weight. `pthread_create(&tid, attr, fn, arg)` starts a new thread that begins by running `fn(arg)` concurrently with the caller, and writes the new thread's identifier into `tid`. `pthread_join(tid, &ret)` blocks until that thread finishes and collects its return value — it is the thread analog of Chapter 14's `wait`: create is to `fork` as join is to `wait`. The thread function has the fixed signature `void *fn(void *)`; a thread ends by returning from it (or calling `pthread_exit`). You compile such a program with `-pthread`. A minimal two-thread program:

```

void *worker(void *arg){ /* runs concurrently */ return NULL; }

int main(void){
    pthread_t t1, t2;
    pthread_create(&t1, NULL, worker, NULL); /* start thread 1 */
    pthread_create(&t2, NULL, worker, NULL); /* start thread 2 */
    pthread_join(t1, NULL); /* wait for it, like wait() for a child */
    pthread_join(t2, NULL);
    return 0;
}

```

After the two `pthread_create` calls, three flows of control exist in one address space — the main thread and the two workers — all sharing the heap and globals, each on its own stack. The two `pthread_join` calls make the main thread wait for both to finish before returning, so the process does not exit (and tear down the shared memory the workers are using) out from under them. That is the whole lifecycle: create, run concurrently, join.

15.3 Kernel threads, user threads, and the one-to-one model

Who actually schedules a thread onto a CPU? There are two classic models. In a **user-level** (or **M:N**) model, a runtime library multiplexes many user threads onto a smaller number of kernel-scheduled entities, switching between them in user space; this makes creating and switching threads very cheap, but it has a sharp edge — if one user thread makes a blocking system call, it can stall the underlying kernel thread and thus every user thread mapped onto it. In a **kernel-level** (or **one-to-one, 1:1**) model, each user thread corresponds to its own kernel-scheduled thread, so the kernel schedules them independently and a blocking call in one blocks only that one, at the cost of a kernel-managed object per thread.

Linux uses the **one-to-one** model, implemented by the **NPTL** (Native POSIX Thread Library, which replaced the older LinuxThreads around the 2.6 kernel). Under the hood, `pthread_create` calls the `clone` system call with a set of flags that ask for sharing — `CLONE_VM` (share the address space), `CLONE_FILES` (share the file-descriptor table), `CLONE_THREAD`, `CLONE_SIGHAND`, and others — so the new thread is a full kernel-scheduled entity that happens to share memory and resources with its creator. This reveals a clean underlying truth: `clone` is the general primitive, and `fork` and thread-creation are two points on its spectrum. `fork` (Chapter 14) is `clone` asking to *copy* the address space; thread creation is `clone` asking to *share* it. A thread and a process are the same kind of schedulable thing; they differ only in how much they share. (User-space schedulers do still exist above the kernel — language runtimes such as Go schedule many lightweight tasks onto a pool of OS threads — but on Linux the OS threads underneath are 1:1 kernel threads.)

QUICK CHECK

Does Linux map each pthread to its own kernel-scheduled thread (1:1) or multiplex many onto few (M:N)? Which system call does `pthread_create` use underneath, and which flag makes the new thread share the caller's memory? (Answer below the communicate box.)

15.4 The shared-memory hazard: the data race

Now the predict. Threads share memory, so two threads can touch the same variable at the same time — and that is where correctness goes to die. Take the simplest possible shared operation, `counter++` on a shared integer. It looks atomic in C, but it is not: at the machine level it is a **read-modify-write** — load the current value into a register, add one, store it back — three steps that can interleave with another thread's three steps. If thread A loads 5, thread B loads 5, both add one, both store 6, then two increments produced a net change of *one*: an update was lost. Run this at scale and the losses pile up. Here is the exact program from the opener — four threads, each incrementing a shared counter one million times, no locking — compiled and run on this machine (gcc 11.4.0, -O2, x86-64), five times:

```
void *bump(void *a){ for (int i = 0; i < 1000000; i++) counter++; return NULL; }
/* 4 threads all run bump() on the same shared `counter`, then we join and print it */

expected = 4000000, got = 1092939
expected = 4000000, got = 1170537
expected = 4000000, got = 1372703
expected = 4000000, got = 1810815
expected = 4000000, got = 1324249
```

The program should have counted to four million. It counted to between roughly 1.1 and 1.8 million — **more than half of the four million increments simply vanished**, and the answer was different every run. This is a **data race**: concurrent accesses to the same memory location, at least one of them a write, with no synchronization ordering them. In C and C++ a data race is **undefined behavior** — Chapter 10's UB returns, now arising from concurrency rather than from a bad pointer or an overflow — and the nondeterministic, mostly-wrong output above is exactly what UB looks like when it manifests. The lost updates come straight from the interleaved read-modify-writes: whenever two threads read the same value before either stores, one increment is overwritten and lost.

THE SAME IDEA THREE WAYS: A DATA RACE LOSES UPDATES

1. In words. `counter++` is not atomic — it is load, add, store. Two threads running it on the same variable can both load the same value, both add one, and both store, so their two increments net only one. Nothing orders them, so results are nondeterministic. That unordered, conflicting access is a data race, and in C/C++ it is undefined behavior.

2. As a picture. An interleaving timeline: A loads 5 · B loads 5 · A adds → 6 · B adds → 6 · A stores 6 · B stores 6. Two increments, one net change.

3. As numbers. Four threads, one million increments each, no lock: expected 4,000,000; measured on this machine ~1.1–1.8 million across five runs — over half the work lost, and different every time. Adding synchronization (a mutex or an atomic increment) restores exactly 4,000,000.

TRAP: "IT PRINTED THE RIGHT ANSWER ON MY MACHINE, SO THE CODE IS CORRECT"

A data race is undefined behavior *whether or not it visibly misbehaves on a given run*, and it can easily produce the right answer while being thoroughly broken — which is exactly what makes it dangerous. **Correct-looking output is not evidence of correctness for racing code.** Writing this very chapter, an earlier version of the counter program — compiled at `-O2` without the variable marked `volatile` — printed `4000000` on every single run, because the optimizer had collapsed each thread's million-iteration loop into a single add, shrinking the race window to almost nothing so the threads happened to serialize. The bug was fully present (four threads doing unsynchronized read-modify-writes on shared memory); it simply did not *manifest*, just as Chapter 10's undefined behavior need not misbehave to be undefined. A race that surfaces only under a particular compiler, optimization level, CPU load, or core count will pass every test on your laptop and corrupt data in production. The lesson: reason about shared-memory correctness from the *rules* (is there an unsynchronized conflicting access? then it is a race, full stop), never from whether a test run looked fine. The fix is real synchronization — a mutex around the update, or an atomic increment — not a comforting green test.

SAY IT LIKE A SYSTEMS ENGINEER

Try it first. A teammate says: "Threads are just cheaper processes. Let's convert our process-per-request server to threads across the board for the speed win." Cover the paragraph and respond in three or four sentences, drawing on §15.1 and §15.4.

Model answer. "Threads *are* cheaper — they share the address space instead of copying it, so creation skips the copy-on-write setup and they can communicate through shared memory with no IPC — but 'cheaper processes' misses the crucial difference: threads drop the isolation that processes gave us. A memory bug or crash in one thread takes down the whole process — every in-flight request — where a bad process only took itself down, and every piece of mutable state we now share becomes a potential data race, which is undefined behavior. That is not theoretical: an unsynchronized shared counter under four threads lost over half its increments in our measurement and gave a different wrong answer each run, and a race can even pass tests by printing the right number. So threads are right where we want cheap shared-memory concurrency and can enforce synchronization discipline; keep processes where fault containment or a security boundary matters. Let's not convert blindly — decide per component and protect every shared write." Naming the isolation trade, the data-race hazard, and refusing the blanket swap is the senior register.

(Answer to the §15.3 quick check: Linux uses the one-to-one (1:1) model via NPTL — each `pthread` is its own kernel-scheduled thread. `pthread_create` uses the `clone` system call underneath, and `CLONE_VM` is the flag that makes the new thread share the caller's address space.)

15.5 What to carry forward — and toward concurrency

A **thread** is a second flow of control inside one process: all threads **share** the address space — code, heap, globals, file descriptors — while each keeps a **private** stack, registers, and thread-local storage. You create and reap them with `threads` — `pthread_create` starts a thread running a `void *fn(void *)`, `pthread_join` waits for it and collects its result (create/join is the thread's fork/wait). Linux schedules threads **one-to-**

one with kernel threads via NPTL, and `pthread_create` is really the `clone` system call asking to share memory (`CLONE_VM`) — the same primitive `fork` uses to *copy* it, so a thread and a process differ only in how much they share. The price of sharing is the **data race**: unsynchronized conflicting access to the same memory, which is undefined behavior in C/C++ and which we measured losing more than half of a four-million-increment workload, nondeterministically — and which can hide by printing the right answer, so correctness must be reasoned from the rules, not from a passing run.

That last point reopens the theme of Part 2 at a new level. Chapter 12 said that in C you carry, by hand, a proof about lifetimes, bounds, and freeing. Threads add a clause: *no unsynchronized concurrent access to shared mutable memory* — and violating it is undefined behavior just like the rest. Making shared-memory concurrency correct is a large subject of its own: the synchronization primitives that order accesses (locks, condition variables, atomics), the lock-free techniques that avoid them, and the memory model that defines precisely what a multithreaded program means. That is the concurrency-and-parallelism part of this volume, which the data race you just measured is the doorway into. And it is one more place where a language can help: just as Rust's ownership makes use-after-free a compile error, its type system also encodes which data may cross between threads and which must be synchronized — turning "don't race" from a discipline you enforce by hand into an invariant the compiler checks, the payoff this whole part has been building toward.

CAN YOU... WITHOUT LOOKING?

- State what threads in a process share and what each keeps private, and contrast that with parent/child after `fork`.
- Write a minimal two-thread pthreads program and explain `pthread_create`/`pthread_join` as the thread analog of `fork`/`wait`.
- Explain the difference between 1:1 and M:N threading, say which Linux uses, and name the system call and flag underneath `pthread_create`.
- Explain why `counter++` from two threads can lose updates, decomposing it into load/add/store with an interleaving.
- State what a data race is, why it is undefined behavior, and why a passing test does not prove racing code correct.
- Give the new clause threads add to Part 2's hand-carried proof, and name the fix category (synchronization).

RETRIEVAL — CLOSED BOOK

Cover the chapter. Answer from memory; rate confidence first.

1. **R1.** What does a thread share with its process's other threads, and what does it keep private? (§15.1)
2. **R2.** What do `pthread_create` and `pthread_join` do, and what earlier pair are they analogous to? (§15.2)
3. **R3.** What is the 1:1 threading model, which library implements it on Linux, and what system call/flag underlie a thread? (§15.3)
4. **R4.** Why can two threads doing `counter++` lose updates, and what did the measurement show? (§15.4)
5. **R5.** What is a data race, why is it undefined behavior, and why can it pass a test yet be wrong? (§15.4)

FURTHER READING

- **Arpaci-Dusseau & Arpaci-Dusseau, *Operating Systems: Three Easy Pieces (OSTEP)*, "Concurrency: An Introduction" and "Interlude: Thread API."** The free chapters that introduce threads, the shared/private split, the pthreads API, and the race that motivates all of concurrency. The reference for §15.1–15.4.
- **Bryant & O'Hallaron, *Computer Systems: A Programmer's Perspective (CS:APP)*, Chapter 12 (Concurrent Programming).** Threads, shared variables, and races from the programmer's view.
- **The Linux `pthread(7)` and `clone(2)` manual pages.** Primary sources for the NPTL one-to-one model of §15.3 and the exact clone flags (`CLONE_VM`, `CLONE_FILES`, `CLONE_THREAD`, ...) that distinguish a thread from a fork. Verify per system.
- **The C11/C++11 standard's memory model (see cppreference.com).** The authority for why a data race is undefined behavior and what "synchronizes-with" means — the formal ground the concurrency part builds on.

Exercises

- 15.1** **WARM-UP** Name two things all threads in a process share, and two things each thread keeps private.
- 15.2** **WARM-UP** What does `pthread_create` take as arguments, and what does `pthread_join` do?
- 15.3** **WARM-UP** Does Linux use a 1:1 or an M:N threading model, and which system call does `pthread_create` use underneath?
- 15.4** **CORE** Write, from memory, a minimal program that starts two threads running the same function and waits for both. Explain how `pthread_create`/`pthread_join` mirror Chapter 14's `fork/wait`, and why the joins are necessary before `main` returns.

15.5 **CORE** Explain why two threads each doing `counter++` on a shared variable can lose updates. Decompose `counter++` into its machine-level steps and give a concrete interleaving in which two increments produce a net change of one.

15.6 **CORE** Contrast a process and a thread on four axes: address space, isolation, cost to create, and what data sharing costs you. For each, say which is cheaper/safer and why.

15.7 **COST** (a) Why is creating a thread cheaper than forking a process — what work does thread creation skip? (b) What new cost or hazard do you take on in exchange for sharing memory? (c) The measured demo lost more than half of its increments and gave a different total each run; is that loss deterministic, and what does its nondeterminism imply for whether tests can be trusted to catch it?

15.8 **FADED** *Worked, with the last step for you.* A web service keeps a shared `long hits` counter that each request thread increments with `hits++`. Under load, the reported hit count is noticeably lower than the true number of requests — but every unit test passes.

Step 1 (given): `hits++` is a non-atomic read-modify-write on memory shared by all request threads.

Step 2 (given): The unit tests exercise one request at a time, so no two increments ever overlap and nothing races.

Step 3 (given): Under real load, many threads increment concurrently and their read-modify-writes interleave, so some increments overwrite others.

Step 4 (you): Name the bug precisely, say whether it is undefined behavior, explain why the passing tests failed to catch it, and name the category of fix (and one concrete mechanism) that makes the count correct.

15.9 **MIXED** For each need, choose a separate *process* (via `fork`) or a *thread*, and justify in a line: (a) run an untrusted plugin so that if it crashes or corrupts memory it cannot take down the host; (b) compute a parallel sum over a large in-memory array that all workers read; (c) run a command the user typed at a shell.

15.10 **MIXED** Drawing on Chapters 9–15, classify each as *undefined behavior*, a *defined-but-real defect*, or *well-defined and fine*: (a) two threads writing the same non-atomic `long` with no synchronization; (b) two threads each writing a *different* variable with no synchronization; (c) a thread reading through a pointer whose target was already freed; (d) two threads reading (never writing) the same shared value.

15.11 **STRETCH** Chapter 12 said that in C the programmer carries, by hand, a proof about lifetimes, bounds, and freeing. (a) State the additional clause that threads add to that proof. (b) Explain why violating it is undefined behavior and why, unlike a use-after-free, it may never reproduce in testing. (c) In one or two sentences, describe how a language could move this from a runtime hope to a compile-time guarantee — for example, a type system that tracks which data may be shared across threads and which accesses must be synchronized (Rust's `Send/Sync` and borrow checking; forward reference, prose only).

15.12 **LAB** On your own machine: (a) run the shared-counter race — four threads, one million increments each of one shared counter, no lock — five times, and report the wrong totals and how much they vary run to run. (b) Add a mutex around the increment (or use an atomic increment) and confirm the total is now exactly four million every time. (c) Try building the unsynchronized version at `-O0` versus `-O2`, and with the counter marked `volatile` or not; report which combinations make the bug visible and which hide it — and explain why a "correct" run does not mean the code is correct. Label your compiler, core count, and hardware.

Chapter 16 — CPU Scheduling

Before you start: Chapter 15 (Linux threads are 1:1 kernel-scheduled entities — *these* are what the scheduler picks among), Chapter 14 (a process blocks in wait or on I/O — a blocked task is off the CPU and off the scheduler's radar), Chapter 13 (a system call is a trap into the kernel — where the scheduler lives). · Chapters 14 and 15 gave you more runnable things than you have CPUs. This chapter is about the OS's answer to the question that creates: *who runs next, and for how long?*

BEFORE YOU READ ON

From memory, reaching back: (1) Chapter 14 — when a process calls `read` on a socket with no data, it *blocks*; is a blocked process using the CPU while it waits? (2) Chapter 15 — Linux schedules each pthread one-to-one with a kernel thread; so is the unit the scheduler places on a CPU the process or the thread? (3) *Predict*: three jobs arrive at almost the same instant — one needs 100 ms of CPU, the other two need 10 ms each. If the scheduler runs them strictly in arrival order (the 100 ms job first), what is the *average* time from arrival to completion across the three? What if it ran the two short ones first instead? Write both numbers down — we will compute them exactly in §16.2.

A modern machine has a handful of CPUs and, at any moment, dozens or hundreds of runnable threads. Something must decide which runnable thread occupies each CPU and when to take it away and give the CPU to another — that something is the **scheduler**, and the decision is **scheduling**. Chapter 13 gave the mechanism the scheduler rides on (a timer interrupt traps into the kernel, which can switch the running thread); this chapter is about the **policy** — the algorithm that chooses. We start by naming what a scheduler optimizes and the metrics that make "good" precise (§16.1); build up the classic policies from the simplest, FIFO, through shortest-job-first and its convoy trap (§16.2); add preemption and the time slice with round-robin, exposing the response-vs-turnaround trade-off (§16.3); and then meet the schedulers real systems ship — the multi-level feedback queue and Linux's proportional-share scheduler — which approximate the ideal without an oracle (§16.4). We close by carrying forward what "the OS multiplexes the CPU" really costs and buys (§16.5).

16.1 The scheduling problem, and how to measure a scheduler

At any instant each thread is in one of a few states. A **running** thread is on a CPU right now. A **ready** (runnable) thread could run but no CPU is free for it. A **blocked** thread is waiting for something — I/O to complete, a child to exit, a lock to free — and *cannot* run until that event arrives, so it is not the scheduler's concern until it wakes. The scheduler's job is to pick, from the ready threads, which one each CPU runs, and a **context switch** (Chapter 13's mechanism) is how it hands a CPU from one thread to another: save the outgoing thread's registers, load the incoming thread's, resume. The switch is not free — it costs the direct register save/restore plus the indirect cost of a cold cache and TLB (Chapter 3) for the new thread — which is why the scheduler cannot simply switch on every instruction.

"Good" needs a definition, and there are several, often in tension. The two that matter most here are **turnaround time** — completion time minus arrival time, how long the whole job took to finish, the metric a batch workload cares about — and **response time** — first-run time minus arrival time, how long until the job *started* getting served, the metric an interactive user feels as lag. A scheduler also cares about **throughput** (jobs completed per unit time), **fairness** (each job gets a defensible share), and avoiding **starvation** (no ready job waits forever). The reason scheduling is hard is that these goals conflict: as we will see, the policy that minimizes average turnaround is not the one that minimizes response time, and the one that is "fair" may not be either.

QUICK CHECK

Define turnaround time and response time in one line each. A job arrives at $t=0$, first gets the CPU at $t=5$, and finishes at $t=30$. What is its response time and its turnaround time? (Answer after the next section.)


16.2 FIFO, shortest-job-first, and the convoy effect

The simplest possible policy is **FIFO** (first-in, first-out, also called FCFS): run jobs to completion in arrival order, no preemption. It is obviously fair in one sense and trivial to implement — and it has a sharp failure mode. Take the opener's workload: three jobs, A needs 100 ms and B and C need 10 ms each, all arriving at essentially $t=0$ with A first. FIFO runs A to completion, then B, then C. A finishes at 100, B at 110, C at 120, so average turnaround is $(100 + 110 + 120) / 3 = \mathbf{110\ ms}$. The two short jobs sat behind the long one for almost their entire wait. This is the **convoy effect**: short jobs pile up behind a long one like cars behind a truck, and average turnaround balloons because everyone inherits the long job's runtime.


The fix, if you knew the run times, is to run the short jobs first. **Shortest-Job-First (SJF)** schedules the ready jobs in increasing order of length. On the same workload it runs B, then C, then A: B finishes at 10, C at 20, A at 120, and average turnaround drops to $(10 + 20 + 120) / 3 = \mathbf{50\ ms}$ — less than half of FIFO's, for the exact same jobs. That is not a coincidence: when all jobs arrive together, SJF is *provably optimal* for average turnaround (any swap that puts a longer job earlier only adds its length to more jobs' waits). When jobs arrive over time, the preemptive version — **Shortest-Time-to-Completion-First (STCF)**, which preempts the running job if a newly arrived one has less work left — extends the same idea. These are the numbers you predicted; here they are, computed by a scheduling simulator:

```
workload: A=100ms, B=10ms, C=10ms, all arrive ~t=0
FIFO (order A,B,C): finish A=100 B=110 C=120 -> avg turnaround 110.0 ms
SJF (order B,C,A): finish B=10 C=20 A=120 -> avg turnaround 50.0 ms
```

The convoy effect: same three jobs, two orders. Bar length = CPU time; number = completion.

FIFO (A first):  120

avg turnaround = $(100+110+120)/3 = 110$ ms — B and C wait behind the truck

SJF (short first):  120

avg turnaround = $(10+20+120)/3 = 50$ ms — short jobs clear first

Same jobs, same total work (120 ms), same last-completion (120). Only the ORDER changed — and average turnaround more than halve

Figure 16.1: FIFO vs SJF on A=100, B=10, C=10 (arriving together). Running the long job first makes every short job inherit its runtime (the convoy effect); running short jobs first minimizes average turnaround — which SJF provably does when all jobs arrive at once.

THE SAME IDEA THREE WAYS: ORDER DOMINATES AVERAGE TURNAROUND

1. In words. Turnaround is completion minus arrival. When jobs share a CPU, a job's turnaround includes the runtime of everything scheduled before it. So putting a long job first adds its length to every later job's turnaround; putting short jobs first minimizes the total. That is why SJF beats FIFO on average turnaround.

2. As a picture. Figure 16.1: the same three bars, packed truck-first versus short-first. The sum of the completion marks — not the total width — is what average turnaround measures, and reordering shrinks it.

3. As numbers. A=100, B=10, C=10: FIFO averages $(100+110+120)/3 = 110$ ms; SJF averages $(10+20+120)/3 = 50$ ms. Identical work, identical finish of the last job (120 ms), less than half the average turnaround — purely from order.

(Answer to the §16.1 quick check: response time = first-run – arrival = 5 – 0 = 5; turnaround = completion – arrival = 30 – 0 = 30.)

16.3 Preemption, the time slice, and round-robin

SJF minimizes turnaround, but it has two problems. First, it needs an **oracle** — you must know each job's length in advance, which for interactive and server workloads you never do. Second, even with the oracle its *response* time is bad: a job that arrives just after a long one starts must wait for it to finish (or, with STCF, only if it is shorter). Interactive workloads care about response time above all — you want the cursor to move the instant you type — so we need a policy that shares the CPU in *small slices* rather than running each job to completion. That policy is **round-robin (RR)**: run each ready job for a fixed **time slice** (quantum), then preempt it (via the timer interrupt) and move to the next, cycling forever until jobs finish.

Round-robin's whole point is response time. On the A/B/C workload with a 1 ms slice, all three jobs get their first turn within the first few milliseconds, so every job's response time is essentially immediate (about 1 ms on average) rather than SJF's — where C waited 10 ms and A waited 20. But that responsiveness has a price: RR's average

turnaround is worse than SJF's, because it stretches every job out by interleaving. Here are all three policies on the same workload from the simulator:

```
workload: A=100ms, B=10ms, C=10ms
FIFO:      avg turnaround 110.0 ms,  avg response  70.0 ms  (B waits 100, C waits 110)
SJF:       avg turnaround  50.0 ms,  avg response  10.0 ms  (C waits 10, A waits 20)
RR (q=1ms): avg turnaround  59.7 ms,  avg response  ~1.0 ms  (every job starts almost at
once)
```

That is the fundamental trade-off in one table: **SJF wins turnaround, round-robin wins response, and you cannot have both**. The **size** of the time slice is the knob. A tiny slice gives excellent response (jobs cycle fast) but pays a context-switch cost on every single switch — and if the slice approaches the switch cost itself, the CPU spends a large fraction of its time switching rather than working. A large slice amortizes the switch cost (better throughput) but degrades response time toward FIFO's. Real schedulers pick a slice large enough that the switch overhead is a small percentage, and small enough that interactivity stays crisp — typically a few milliseconds.

QUICK CHECK

Round-robin makes the time slice small to improve one metric and worsen another — which is which? And what goes wrong if you make the slice as small as the cost of a single context switch? (Answer below the communicate box.)

16.4 What real schedulers do: MLFQ and proportional share

Real systems cannot use pure SJF (no oracle) or pure round-robin (ignores that jobs differ in importance and behavior). Two ideas dominate real schedulers, and both aim to get SJF-like behavior — favor short, interactive work — *without* knowing job lengths in advance, by **learning from the past**.

The first is the **multi-level feedback queue (MLFQ)**. Keep several ready queues, each a priority level; always run a job from the highest non-empty level, round-robin within a level. The trick is how a job's priority *changes*: a new job enters at the top; if it uses its whole time slice (CPU-bound behavior) it is demoted a level; if it gives up the CPU before the slice ends (blocking on I/O — interactive behavior) it stays high. Over a few slices this *learns* each job's character: interactive jobs, which block often, float near the top and get scheduled quickly (great response time); CPU-bound jobs sink to the bottom and share the leftovers. That approximates "short jobs first" using observed behavior instead of an oracle. Two refinements make it work in practice: to prevent a long job from **starving** at the bottom forever, MLFQ periodically **boosts** every job back to the top; and to stop a job from **gaming** the rule (yielding just before each slice ends to keep its priority), it accounts for total CPU used at a level rather than per-slice behavior.

The second idea is **proportional share** (fair-share): rather than juggle priorities, give each job a **weight** and hand out CPU time in proportion to weight, so a job with twice the weight gets twice the CPU over time. This is the family Linux's normal scheduler belongs to. For over a decade Linux used the **Completely Fair Scheduler (CFS)**, which tracks

each task's accumulated virtual runtime and always runs the task that has had the least, weighted by priority; kernel 6.6 (2023) replaced CFS with **EEVDF** (Earliest Eligible Virtual Deadline First), a proportional-share scheduler with better latency guarantees. In both, a process's **nice** value (−20 to +19, set with `nice/renice`; lower is "less nice," i.e. higher priority) sets its weight: a lower nice value claims a larger share. You can see this directly. Two identical CPU-bound loops pinned to the *same* core so they must compete, one at nice 0 and one at nice 19, run for five wall-clock seconds on this machine (Linux 6.18, which uses EEVDF; numbers illustrative):

```
taskset -c 0 nice -n 0 ./hog -> nice0: cpu-seconds used = 4.66
taskset -c 0 nice -n 19 ./hog -> nice19: cpu-seconds used = 0.09
```

The nice-0 loop got roughly fifty times the CPU of the nice-19 loop over the same wall-clock window — not because nice 19 was "paused," but because the proportional-share scheduler weighted its slice far smaller. That is the whole idea made visible: `nice` does not change what a program computes, only what *fraction of the CPU* it is entitled to when there is contention. (Linux also offers strict real-time classes — `SCHED_FIFO` and `SCHED_RR` — that sit above the normal class and always preempt it; those are for bounded-latency work and are a foot-gun for everything else, as the warning below explains.)

TRAP: REACHING FOR REAL-TIME PRIORITY (`SCHED_FIFO`) TO "MAKE IT FASTER"

When a latency-sensitive service is being starved by a background job, the tempting fix is to promote the service to a real-time scheduling class — `SCHED_FIFO` or `SCHED_RR` — so it always wins. This is *strict fixed-priority* scheduling, and it does not mean "a bit faster": a runnable `SCHED_FIFO` thread preempts **every** normal task and runs until it blocks or a higher real-time thread appears. If that thread is CPU-bound and never blocks — a busy loop, a spin-wait, a bug — it will monopolize its core and **starve everything else on it, including kernel threads**, and can wedge the machine hard enough to require a reboot. (Linux mitigates this with real-time throttling — `sched_rt_runtime_us` caps RT time per period — precisely because the default is so dangerous.) Real-time priority is the right tool only for work that is both genuinely latency-critical *and* bounded in CPU (it blocks regularly). To bias the normal scheduler toward an important service, reach first for nice or cgroup CPU weights, which change the *share* under contention without the power to starve the system. The lesson: "higher priority" is not a speed dial; strict priority is a loaded gun pointed at every other task on the CPU.

SAY IT LIKE A SYSTEMS ENGINEER

Try it first. A teammate says: "Our request service and the nightly compaction job share a box, and compaction is starving the service's tail latency. Let's just give the service `SCHED_FIFO` real-time priority so it always wins." Cover the paragraph and respond in three or four sentences, drawing on §16.4.

Model answer. "The instinct is right — the service should win against a background batch job — but `SCHED_FIFO` is a much bigger hammer than we want: it is strict fixed priority, so if any of the service's threads ever spins or busy-loops without blocking, it will preempt everything on that core, including kernel threads, and can hang the box, which is why the kernel even throttles real-time time by default. What we actually want is a bigger *share* under contention, not absolute priority — so lower the service's nice value or, better, put the two workloads in separate cgroups and give the service a larger CPU weight (and cap or de-prioritize compaction). That fixes the starvation through the fair scheduler, keeps the system stable, and leaves headroom if compaction ever needs to burst. Save real-time scheduling for work that is both latency-critical and provably bounded in CPU." Distinguishing "bigger share" from "strict priority," and naming the starvation and hang risk, is the senior register.

(Answer to the §16.3 quick check: a small time slice improves **response** time (jobs cycle onto the CPU quickly) and worsens **turnaround** and throughput (more interleaving and more context switches). If the slice approaches the cost of one context switch, the CPU spends a large fraction of its time switching rather than doing useful work — overhead dominates.)

16.5 What to carry forward

The **scheduler** decides which ready thread each CPU runs and for how long; a **context switch** is the (non-free) mechanism that hands a CPU from one thread to another. We measure schedulers by **turnaround** (completion – arrival) and **response** (first-run – arrival) time, among others, and the central fact is that these goals conflict. FIFO is simple but suffers the **convoy effect**; **SJF/STCF** minimize average turnaround but need an oracle and give poor response; **round-robin** with a small time slice wins response at the cost of turnaround, with the slice size trading responsiveness against switch overhead. Real schedulers approximate "favor short, interactive work" without an oracle by learning from behavior — the **multi-level feedback queue** (priorities that adapt, with boosting to prevent starvation) — or by handing out CPU in proportion to a **weight** — Linux's proportional-share scheduler (CFS, then EEVDF from 6.6), where the **nice** value sets the share, which we watched split one core roughly fifty to one between a nice-0 and a nice-19 loop.

Step back and see what this part is assembling. Chapter 14 gave you many processes and Chapter 15 many threads; this chapter is how one machine *pretends* to run them all at once — by slicing the CPU in time so fast that each thread perceives a CPU (nearly) to itself. That is the CPU half of the operating system's core illusion. The other half is the same trick played on *memory*: giving each process the illusion of a large private address space on a machine whose physical memory is smaller and shared. Chapter 4 showed the hardware that makes an address virtual; the next chapter is the OS policy that makes the

illusion hold when memory runs short — demand paging, page faults as a mechanism, and the replacement decisions that are this chapter's turnaround-versus-response trade-off played again on pages instead of milliseconds.

CAN YOU... WITHOUT LOOKING?

- Name the three scheduling-relevant states of a thread (running, ready, blocked) and say which ones the scheduler chooses among.
- Define turnaround time and response time, and compute both for a job given its arrival, first-run, and completion times.
- Explain the convoy effect with the $A=100/B=10/C=10$ example, and give FIFO's and SJF's average turnaround (110 vs 50 ms).
- State why SJF is optimal for average turnaround yet unusable in practice, and what round-robin trades to win response time.
- Explain how a multi-level feedback queue approximates SJF without an oracle, and why it needs priority boosting.
- Explain what a nice value does under Linux's proportional-share scheduler, and why `SCHED_FIFO` is dangerous.

RETRIEVAL — CLOSED BOOK

Cover the chapter. Answer from memory; rate confidence first.

1. **R1.** What are the running/ready/blocked states, and which does the scheduler pick among? (§16.1)
2. **R2.** Define turnaround and response time. (§16.1)
3. **R3.** What is the convoy effect, and what were FIFO's and SJF's average turnarounds on A/B/C? (§16.2)
4. **R4.** Why does round-robin improve response time but hurt turnaround, and what does the time-slice size trade? (§16.3)
5. **R5.** How does an MLFQ learn to favor interactive jobs without knowing their lengths, and what does a nice value control? (§16.4)

FURTHER READING

- **Arpaci-Dusseau & Arpaci-Dusseau, *Operating Systems: Three Easy Pieces (OSTEP)*, "Scheduling: Introduction" and "Scheduling: The Multi-Level Feedback Queue."** The free chapters that develop turnaround/response metrics, FIFO/SJF/STCF/round-robin, and MLFQ (including boosting and anti-gaming). The reference for §16.1–16.4.
- **Bryant & O'Hallaron, *Computer Systems: A Programmer's Perspective (CS:APP)*, Chapter 8 (Exceptional Control Flow).** Context switches and the timer interrupt that preemption rides on, from the programmer's view.
- **The Linux `sched(7)` and `nice(2)` / `sched_setscheduler(2)` manual pages.** Primary sources for the scheduling classes (`SCHED_OTHER/normal`, `SCHED_FIFO`, `SCHED_RR`), the nice range, and real-time throttling. Verify per kernel version.
- **The Linux kernel scheduler documentation (docs.kernel.org/scheduler/), including the EEVDF scheduler page.** The authority for the CFS-to-EEVDF change in 6.6 and how proportional share is implemented. Fast-moving; treat details as version-specific.

Exercises

- 16.1** **WARM-UP** Name the three scheduling-relevant states of a thread, and say which state(s) the scheduler chooses the next runner from.
- 16.2** **WARM-UP** Define turnaround time and response time. For a job that arrives at $t=0$, first runs at $t=8$, and completes at $t=40$, give both.
- 16.3** **WARM-UP** In one sentence, what is the convoy effect, and which simple scheduling policy suffers from it?
- 16.4** **CORE** Three jobs arrive at $t=0$: X needs 8 ms, Y needs 4 ms, Z needs 4 ms. Compute the average turnaround time under FIFO (order X, Y, Z) and under SJF, showing each job's completion time. Which is lower, and why is that no accident?
- 16.5** **CORE** Explain why SJF minimizes average turnaround when all jobs arrive together, and give the two reasons it is nonetheless unusable as a general-purpose scheduler.
- 16.6** **CORE** Describe how a multi-level feedback queue decides a job's priority over time, and explain how this makes interactive (I/O-bound) jobs get scheduled quickly and CPU-bound jobs sink — without the scheduler ever being told which is which.
- 16.7** **COST** (a) What does shrinking the round-robin time slice buy you, and what does it cost? (b) A context switch on a machine costs roughly a few microseconds; if you set the time slice equal to that, what fraction of the CPU is wasted on switching, and why? (c) Give the qualitative reason real schedulers use a slice of a few milliseconds rather than a few microseconds or a few seconds.
- 16.8** **FADED** *Worked, with the last step for you.* A batch analytics job and an interactive dashboard API share one server. Users complain the dashboard is laggy whenever the batch job runs, even though the box has spare capacity at night.
- Step 1 (given):** The dashboard is I/O-bound and latency-sensitive (response time matters); the batch job is CPU-bound and only turnaround matters.
- Step 2 (given):** Under a pure round-robin or a naive equal-share policy, the CPU-heavy batch job gets an equal turn each cycle, so the dashboard waits behind it and its response time suffers.
- Step 3 (given):** The two workloads should not get equal CPU: the dashboard should be favored under contention, but the batch job must still make progress (not starve).
- Step 4 (you):** Name a concrete Linux mechanism (not real-time priority) that gives the dashboard a larger share under contention while letting the batch job run when the dashboard is idle, and say in one line why this is safer than `SCHED_FIFO`.
- 16.9** **MIXED** For each situation, name the scheduling metric that matters most (turnaround, response time, or throughput), and justify in a line: (a) a text editor reacting to keystrokes; (b) a nightly batch job that must finish before the business day; (c) a build farm maximizing jobs completed per hour across many machines.

16.10 **MIXED** Drawing on Chapters 13–16, classify each claim as *true* or *false* and fix the false ones: (a) "A process blocked on `read` is consuming CPU while it waits." (b) "On Linux the scheduler picks among processes, not threads." (c) "A context switch is free, so the time-slice size does not matter." (d) "Lowering a process's `nice` value increases the CPU share it gets under contention."

16.11 **STRETCH** Starvation is a scheduler's cardinal sin. (a) Show how strict shortest-job-first can starve a long job indefinitely if short jobs keep arriving. (b) Explain how a multi-level feedback queue's periodic priority *boost* prevents this, and what would go wrong without it. (c) Explain why a proportional-share scheduler (weight-based) does not starve a low-weight job even under load — contrast "smaller share" with "no share," and connect this to why `nice` cannot starve a task but `SCHED_FIFO` can.

16.12 **LAB** On your own machine: (a) write a CPU-bound loop that runs for a fixed number of wall-clock seconds and reports the CPU-seconds it actually accumulated (via `CLOCK_THREAD_CPUTIME_ID`). (b) Run two copies pinned to the *same* core with `taskset -c 0` so they must compete, one at `nice 0` and one at `nice 19`, and report the CPU-seconds each got — you should see the `nice-0` copy dominate by a large ratio. (c) Now run them on *different* cores (or without pinning on a multi-core box) and explain why the `nice` difference nearly disappears. Label your kernel version, core count, and scheduler (CFS vs EEVDF). Explain what `nice` did and did not change.

Chapter 17 — Virtual Memory: Demand Paging and Page Replacement

Before you start: Chapter 4 (the hardware that makes an address virtual — the MMU, page tables, and the TLB translate a virtual page to a physical frame; *this* chapter is the OS policy that fills those tables), Chapter 16 (the OS multiplexes the CPU in time; it multiplexes memory in space the same way), Chapter 13 (a page fault is exceptional control flow — a trap into the kernel). · Chapter 4 showed *how* a virtual address becomes a physical one. This chapter is what the OS does when the answer is "that page isn't in memory."

BEFORE YOU READ ON

From memory, reaching back: (1) Chapter 4 — an address is *virtual*; what hardware structure translates a virtual page number to a physical frame, and what small cache speeds that translation up? (2) Chapter 16 — the scheduler faced conflicting goals (turnaround vs response); name the sin a scheduler must avoid (a job that never runs). (3) *Predict*: a program calls `mmap` to reserve 512 MiB of memory and immediately checks its resident memory (RSS) *before touching any of it*. Is its RSS about 512 MiB, about 0, or somewhere in between? Then it writes to every byte and checks again. Write down both guesses — we will measure them in §17.2.

Chapter 4 gave a process the illusion of a large, private, contiguous address space, and showed the hardware — the MMU walking a page table, the TLB caching recent translations — that turns each virtual address into a physical one on every memory access. But that chapter assumed the page was *there*: that every virtual page a process uses has a physical frame behind it. It usually does not, and cannot: the sum of every process's address space is far larger than physical RAM. The operating system sustains the illusion anyway, and this chapter is how. We reframe memory as the OS sees it — physical frames it hands out and page tables it owns (§17.1); introduce **demand paging**, where a page gets a frame only when first touched, and the **page fault** that makes it happen — measured (§17.2); go beyond physical memory to **swapping** and the replacement question it forces (§17.3); and work through the replacement policies — OPT, FIFO, LRU — and the anomaly that makes FIFO treacherous (§17.4). We close by carrying forward the mechanisms this unlocks — `mmap`, copy-on-write, the page cache — that the rest of the book runs on (§17.5).

17.1 Memory as the OS sees it: frames, page tables, and the illusion

Recall the two vocabularies. A process sees **virtual pages**: its address space, chopped into fixed-size pages (4 KiB on x86-64 by default). Physical memory is chopped into equal-size **frames**. A **page table**, one per process and owned by the OS, maps each virtual page to a frame — or records that it has none. Each page-table entry (PTE) holds the frame number plus control bits, and the one that matters here is the **present bit**: set if the page currently lives in a frame, clear if it does not. On every access the MMU consults the page table (via the TLB); if the translation says present, the access proceeds

at hardware speed. This is the machinery of Chapter 4, now seen from the OS's chair: the OS is the party that *fills in* those tables, hands out frames, and decides what to do when a PTE's present bit is clear.

That last case is the whole subject of this chapter. Because every process has its own page table, each gets its own private virtual address space — the same virtual address in two processes maps to different frames, which is the isolation Chapter 14 relied on. And because the OS controls the present bit, it can lie productively: it can promise a process a page it has not actually backed with a frame yet, and make good on the promise only when the process reaches for it. When the process touches a page whose PTE is not present, the hardware cannot proceed — so it traps.

QUICK CHECK

What does the *present bit* in a page-table entry mean, and who sets it — the hardware or the OS? If a process accesses a virtual page whose PTE has the present bit clear, what happens next? (Answer after the next section.)

17.2 Demand paging and the page fault, measured

When a running instruction accesses a virtual page whose PTE is not present, the MMU raises a **page fault** — a hardware exception, Chapter 13's exceptional control flow, that traps into the kernel's page-fault handler. The handler inspects why the page is absent and chooses: if the access is *legal* but the page simply was never materialized (a freshly `mmap`'d region, the first touch of newly allocated heap), it finds a free frame, fills it appropriately — zeroed for anonymous memory, read from a file for a file mapping — sets the PTE's present bit, and restarts the faulting instruction, which now succeeds. If the access is *illegal* (a wild pointer into an unmapped region), the handler delivers `SIGSEGV` instead — the segmentation fault of Chapters 7–11, now revealed as "the page-fault handler found no valid mapping." This lazy strategy — give a page a frame only when it is first used — is **demand paging**, and it is why allocating memory is cheap: `malloc` or `mmap` mostly just records a promise in the page tables; the physical frames are spent only on the pages you actually touch.

You can watch it directly. A program that `mallocs` 100 MiB and writes one byte to every 4 KiB page should take exactly one page fault per page — $100 \text{ MiB} / 4 \text{ KiB} = 25,600$ of them — and, because anonymous pages are served from zeroed frames with no disk I/O, they should all be **minor** faults (frame available without I/O), never **major** (required disk I/O). Running it under `/usr/bin/time -v (gcc 11.4.0, Linux 6.18, x86-64)`:

Maximum resident set size (kbytes): 103852	# ~100 MiB actually resident after touching
Minor (reclaiming a frame) page faults: 25672	# ~one per 4 KiB page touched (25,600) + a little
Major (requiring I/O) page faults: 0	# no disk I/O — anonymous demand-zero pages

And the reservation-versus-use split from the opener: a program that `mmaps` 512 MiB and checks its resident memory before and after touching it:

```

after mmap 512 MiB: VmRSS = 1364 kB      # the mapping exists, but almost nothing is
resident
after touching all: VmRSS = 525976 kB   # now every page has been faulted in (~512 MiB)

```

The 512 MiB mapping cost essentially *no* physical memory until it was used: RSS was about 1.3 MiB (just the program itself), not 512 MiB, right after the `mmap` succeeded. Touching every page faulted each one in, and only then did RSS climb to ~512 MiB. That is demand paging in two numbers: **allocation is a promise; residency is paid page by page, on first touch.**

THE SAME IDEA THREE WAYS: DEMAND PAGING PAYS PER TOUCHED PAGE


1. In words. Reserving address space (via `mmap/malloc`) only records mappings in the page table with the present bit clear; no frame is spent. The first access to each page faults, the kernel gives it a frame and sets present, and the instruction restarts. So physical memory is consumed in proportion to pages *touched*, not pages *reserved*.

2. As a picture. Figure 17.1: a virtual page with `present=0` → access → page fault trap → handler grabs a free frame, zeroes/loads it, sets `present=1` → instruction restarts and succeeds. Later faults on other pages repeat the path.

3. As numbers. `mmap 512 MiB`: RSS 1,364 kB before touching, 525,976 kB after touching all. Writing one byte per page across 100 MiB took 25,672 minor faults ($\approx$ one per 4 KiB page) and *zero* major faults — anonymous pages are served from zeroed frames, no disk I/O.

(Answer to the §17.1 quick check: the *present* bit means the page currently occupies a physical frame; the **OS** sets it (the hardware only reads it during translation). Accessing a not-present page raises a page fault — a trap into the kernel's fault handler, which either materializes the page or delivers `SIGSEGV`.)

A page fault on first touch (demand paging), then the restart.



The faulting instruction is re-executed from scratch after the handler returns; the program never sees the fault, only a slight pause.

Figure 17.1: The demand-paging fault path. A first touch of a not-present page traps into the kernel, which materializes the page (a minor fault if no disk I/O is needed, a major fault if it must read from a file or swap), sets the present bit, and restarts the instruction. Illegal accesses become `SIGSEGV` instead.

17.3 Beyond physical memory: swapping and the replacement question

Demand paging is easy while free frames last. The hard case is when physical memory fills and a new page must be brought in with no frame to spare. The OS's answer is

swapping (paging to a backing store): it *evicts* some resident page — writing it out to a swap area on disk if it has been modified since it was loaded (a "dirty" page; a clean page backed by a file can just be dropped and re-read later) — freeing that frame for the incoming page. Later, if the evicted page is touched again, it faults back in from disk. This is where the minor/major distinction bites: a fault that can be satisfied from a free or droppable frame is a **minor** fault (fast, no I/O); a fault that must *read the page back from disk or swap* is a **major** fault, and it costs a disk access — Chapter 3's numbers make the gap vivid, since a disk (even an SSD) is orders of magnitude slower than DRAM, and a spinning disk millions of times slower than a cache hit.

Swapping keeps programs running past the size of RAM, but it turns the choice of *which page to evict* into a performance decision of the first order: every eviction of a page that is about to be needed again forces an expensive major fault to bring it back. This is the **page-replacement problem**, and it is Chapter 16's dilemma in a new guise — a policy question about a scarce, multiplexed resource, where a wrong choice is not incorrect, merely slow. The metric here is the **page-fault rate** (equivalently, the miss rate against the frames you have), and the goal is a replacement policy that minimizes it.

QUICK CHECK

What is the difference between a *minor* and a *major* page fault, and which one involves disk I/O? When the OS evicts a page to make room, why does it matter whether that page is "dirty" (modified since it was loaded)? (Answer below the communicate box.)

17.4 Replacement policies: OPT, FIFO, LRU, and Belady's anomaly

What should the OS evict? The unbeatable baseline is **OPT** (Belady's MIN, 1966): evict the page whose *next* use is farthest in the future. It provably minimizes faults — but it needs to see the future, so it is only a yardstick, not an implementable policy. The simplest real policy is **FIFO**: evict the page that has been resident longest, regardless of use. It is trivial but ignores *how* pages are used. The workhorse approximation is **LRU** (least-recently-used): evict the page that has gone longest without being accessed, betting that the recent past predicts the near future — which, thanks to **locality of reference** (Chapter 3), it usually does. On a short reference string, simulated exactly:

```
reference string: 0 1 2 0 1 3 0 3 1 2 1    (3 frames)
FIFO: 7 page faults
LRU : 5 page faults
OPT : 5 page faults      # the lower bound; no policy can do better on this string
```

FIFO takes 7 faults here where LRU and OPT take 5, because FIFO evicts pages by age even when they are about to be reused. But FIFO's real hazard is subtler and worth naming, because it violates an intuition everyone has: that giving a program *more* memory can only help. For FIFO, it can hurt. On the classic reference string of Belady, Nelson & Shedler (1969), FIFO takes *more* faults with four frames than with three:

```
reference string: 1 2 3 4 1 2 5 1 2 3 4 5
FIFO, 3 frames: 9 page faults
FIFO, 4 frames: 10 page faults      # MORE memory, MORE faults
```

This is **Belady's anomaly**: for FIFO, adding a frame can increase the fault count, because a larger FIFO queue changes the eviction order in a way that can evict a soon-to-be-reused page it would have kept with fewer frames. LRU and OPT are immune — they belong to a class (stack algorithms) whose resident set with N frames is always a superset of its resident set with $N-1$, so more memory never adds faults. The anomaly is the sharpest argument against FIFO and for use-based policies like LRU: a policy that ignores usage can be not just suboptimal but *non-monotonic*, punishing you for adding memory. (Real kernels do not implement true LRU — tracking exact recency on every access is too expensive — but approximate it cheaply with a per-page reference bit and a clock/second-chance sweep, keeping LRU's monotonic, locality-exploiting behavior without its bookkeeping cost.)

TRAP: THRASHING — WHEN THE WORKING SET DOESN'T FIT

Replacement policy only helps while the pages a program actively uses — its **working set** — fit in the frames it has. When the sum of all running processes' working sets exceeds physical memory, *no* replacement policy can win: every page the OS evicts is one that is about to be needed, so the system generates major fault after major fault, and the CPU sits idle waiting on disk while throughput collapses. This is **thrashing**, and its signature is a machine that is nearly idle on CPU yet grinding — high major-fault and swap-I/O rates, everything slow. The trap is diagnosing it as a CPU or code problem and "optimizing the hot loop," when the real fix is to reduce the memory footprint (smaller working set, fewer concurrent jobs) or add RAM so the working set fits. Modern Linux often does not let you thrash gently to a halt: when memory (plus swap) is truly exhausted, the **OOM killer** picks a process and kills it outright. Either way the lesson is the same — paging is a graceful extension of memory only up to the working-set boundary; past it, performance does not degrade linearly, it falls off a cliff. Watch major-fault and swap rates, not just CPU.

SAY IT LIKE A SYSTEMS ENGINEER

Try it first. A teammate says: "Our service malloced a 40 GB buffer on a 16 GB box and the allocation *succeeded* — so we clearly have enough memory, and we can fill it up. Great." Cover the paragraph and respond in three or four sentences, drawing on §17.2 and the thrashing trap.

Model answer. "A successful malloc doesn't mean the memory exists — it means the kernel recorded the mapping; under demand paging (and overcommit) no physical frames are spent until you *touch* the pages, so a 40 GB allocation on a 16 GB box can succeed while backing almost nothing. The moment we actually write across all 40 GB, each page faults in, resident memory climbs past 16 GB, and the OS has to swap — if our working set genuinely exceeds RAM we'll thrash (near-idle CPU, endless major faults, everything crawling) or get OOM-killed outright. So allocation success proves nothing about capacity; what matters is the *resident working set*. Let's size the buffer to what fits with headroom, stream instead of buffering everything, and watch RSS and major-fault rates under load — not just whether malloc returned non-NULL." Separating "reserved" from "resident," and naming thrashing and the OOM killer, is the senior register.

*(Answer to the §17.3 quick check: a **minor** fault is satisfied without disk I/O — the page is served from a zeroed or already-cached frame; a **major** fault must read the page from a file or swap, so it costs a disk access. A "dirty" (modified) page must be **written out** to swap before its frame can be reused, whereas a clean page — an unmodified copy of file data — can be dropped and re-read later, so evicting clean pages is cheaper.)*

17.5 What to carry forward

The OS owns physical **frames** and each process's **page table**, and the **present bit** lets it back a virtual page with a frame lazily. **Demand paging** materializes a page only on first touch: the access to a not-present page raises a **page fault** (a trap), and the handler either supplies the page — a **minor** fault if no I/O is needed, a **major** fault if it must read from disk or swap — or delivers SIGSEGV. Allocation is therefore a promise (we watched a 512 MiB mapping cost ~1 MiB of RSS until touched), and residency is paid page by page (25,672 minor faults to touch 100 MiB). When RAM fills, **swapping** evicts pages to disk, making the choice of victim — the **page-replacement problem** — a performance decision: OPT is the unbeatable but unimplementable yardstick, FIFO is simple but suffers **Belady's anomaly** (more memory, more faults), and LRU (approximated in real kernels) exploits locality and is monotonic. Past the working-set boundary, no policy saves you: **thrashing**, then the OOM killer.

With this the operating system's core illusion is complete: Chapter 16 sliced the CPU in time, this chapter backs each address space with a smaller physical memory, and together they let many processes share one machine as if each owned it. The mechanisms this chapter unlocked are the ones the rest of the book leans on. The same page table that holds the present bit also holds protection and sharing bits, which give us **memory-mapped files** (`mmap` a file and read/write it as memory, the page cache faulting data in on demand) and **copy-on-write** — the trick from Chapter 14's `fork`, now seen fully: parent and child share every frame read-only, and only a write faults and copies the one page touched. Copy-on-write and shared mappings are also how separate processes can deliberately share a region of memory — the fast path for the inter-process communication that a later chapter takes up, and the reason the isolation of processes need not mean the cost of copying. Virtual memory is not only protection; it is the substrate for sharing, laziness, and speed.

CAN YOU... WITHOUT LOOKING?

- Explain, using the present bit, how the OS can promise a process a page it has not yet backed with a frame.
- Trace the page-fault path for a first touch of anonymous memory, and say why it is a minor fault, not a major one.
- State what demand paging makes cheap, and give the two measured numbers (RSS before/after touch; minor faults per page).
- Distinguish minor and major faults, and explain why evicting a dirty page costs more than evicting a clean one.
- Rank OPT, FIFO, and LRU, and explain Belady's anomaly and why LRU/OPT are immune to it.
- Define thrashing and the working set, and explain why a successful malloc does not prove you have the memory.

RETRIEVAL — CLOSED BOOK

Cover the chapter. Answer from memory; rate confidence first.

1. **R1.** What does the present bit mean, and what happens on an access to a not-present page? (§17.1, §17.2)
2. **R2.** What is demand paging, and what did the mmap-512-MiB measurement show about reserved vs resident memory? (§17.2)
3. **R3.** Distinguish minor and major page faults; which involves disk I/O? (§17.2, §17.3)
4. **R4.** What is the page-replacement problem, and how do OPT, FIFO, and LRU differ? (§17.4)
5. **R5.** What is Belady's anomaly, and what is thrashing? (§17.4)

FURTHER READING

- **Arpaci-Dusseau & Arpaci-Dusseau, *Operating Systems: Three Easy Pieces (OSTEP)*, "Paging: Introduction," "Beyond Physical Memory: Mechanisms," and "Beyond Physical Memory: Policies."** The free chapters on the page as the unit, the page fault and swapping mechanism, and the replacement policies (OPT/FIFO/LRU, the anomaly, clock). The reference for §17.1–17.4.
- **Bryant & O'Hallaron, *Computer Systems: A Programmer's Perspective (CS:APP)*, Chapter 9 (Virtual Memory).** Address spaces, page tables, demand paging, and memory mapping from the programmer's view; the companion to Chapter 4's hardware treatment.
- **Belady, "A study of replacement algorithms for a virtual-storage computer," *IBM Systems Journal* 5(2):78–101 (1966); and Belady, Nelson & Shedler, "An anomaly in space-time characteristics of certain programs running in a paging machine," *Communications of the ACM* 12(6):349–353 (1969).** The primary sources for the optimal (MIN) policy and for Belady's anomaly, respectively.
- **The Linux `mmap(2)`, `madvise(2)`, and `proc(5)` manual pages.** Primary sources for memory mapping, demand-paging hints, and the `/proc/<pid>/status` fields (`VmRSS`) used to measure residency. Verify per kernel version.

Exercises

- 17.1** **WARM-UP** Define a virtual page and a physical frame, and say what a page-table entry's present bit records.
- 17.2** **WARM-UP** What is a page fault, and what two broad outcomes can the kernel's fault handler produce for it?
- 17.3** **WARM-UP** In one line each, distinguish a minor page fault from a major page fault.
- 17.4** **CORE** A program `mmaps` 1 GiB and its RSS is a few megabytes immediately afterward, then climbs toward 1 GiB as it works. Explain what happened at each stage in terms of demand paging, and why the allocation did not consume 1 GiB of physical memory up front.
- 17.5** **CORE** Simulate FIFO and LRU on the reference string 1 2 3 1 4 2 5 with 3 frames. Give the number of page faults for each and show the resident set after each reference. Which policy does better here, and why?
- 17.6** **CORE** Explain Belady's anomaly: state what it is, which policy exhibits it, and why LRU and OPT cannot. Use the phrase "resident set is a superset with more frames" in your explanation of the immune policies.
- 17.7** **COST** (a) Why is a major page fault dramatically more expensive than a minor one — put rough numbers on it using Chapter 3's memory hierarchy. (b) Why does evicting a *clean* page cost less than evicting a *dirty* one? (c) A service's p99 latency has sudden 50 ms spikes that correlate with high major-fault counts; what is the likely cause, and is the fix in the CPU-bound code?
- 17.8** **FADED** *Worked, with the last step for you.* A batch job processes a 30 GB dataset on a 32 GB machine by `mmap`ing the whole file and iterating over it once, sequentially. It runs fine. A colleague "optimizes" it to make ten concurrent passes over the same mapping from ten threads to "use all the cores," and now the machine crawls with near-zero CPU utilization.
- Step 1 (given):** The single sequential pass has a small working set at any instant (a sliding window of recently touched pages), so it fits comfortably in RAM and pages stream in and out cleanly.
- Step 2 (given):** Ten concurrent passes at different file offsets touch ten far-apart regions at once, so the combined working set is large and scattered.
- Step 3 (given):** That combined working set exceeds physical memory, so every page the OS evicts is soon needed by another thread, generating major fault after major fault.
- Step 4 (you):** Name the phenomenon precisely, explain why CPU utilization is near *zero* during it (what are the cores waiting on?), and give the direction of the fix (not "optimize the loop").

17.9 **MIXED** For each observation, name the most likely cause and one thing to measure to confirm it: (a) allocating a huge array succeeds instantly but the program slows down later as it fills the array; (b) a program that fit in RAM last week now runs 100× slower after the dataset grew slightly; (c) reading a memory-mapped file is fast the second time through but slow the first.

17.10 **MIXED** Drawing on Chapters 4, 14, and 17, classify each statement as *true* or *false* and fix the false ones: (a) "After a successful 8 GB `malloc`, 8 GB of physical RAM is now in use." (b) "A segmentation fault is just a page fault the handler could not satisfy legally." (c) "`fork` copies the parent's entire address space immediately." (d) "Two processes with the same virtual address are reading the same physical frame."

17.11 **STRETCH** Copy-on-write ties this chapter to Chapter 14. (a) Explain how the page table lets `fork` give the child a "copy" of the address space while physically copying no data at first. (b) Describe exactly what happens, page by page, when the child later writes to one page — which fault fires, what the handler does, and how many pages get copied. (c) Explain why this makes the common `fork-then-exec` pattern (Chapter 14) nearly free even for a large parent, and what a single write to a shared page costs.

17.12 **LAB** On your own machine: (a) write a program that `mallocs` (or `mmaps`) a few hundred MiB and touches one byte per 4 KiB page; run it under `/usr/bin/time -v` and report the minor-fault count — confirm it is about one per page. (b) Read `VmRSS` from `/proc/self/status` before and after touching a large mapping, and report both values to show reservation vs residency. (c) Optional and careful: on a machine you can afford to stress (or a VM), allocate and touch more than physical RAM and observe swap I/O / major faults climbing — describe the thrashing behavior and stop before you wedge the box. Label your kernel version, page size, and RAM.

Chapter 18 — Signals: Asynchronous Control Flow

Before you start: Chapter 14 (when a child exits, the kernel notifies the parent with `SIGCHLD` — that notification *is* a signal; and `wait` reaps the child), Chapter 13 (a blocking system call can be interrupted — this chapter is what interrupts it), Chapter 11 (a bad memory access became a segmentation fault — `SIGSEGV` is a signal too). · Every earlier chapter's control flow was *synchronous*: the program ran its own instructions in order. This chapter adds a second, *asynchronous* flow the kernel can force on a process at almost any instant.

BEFORE YOU READ ON

From memory, reaching back: (1) Chapter 14 — after a child process exits but before the parent waits for it, what is the child called, and how did the parent learn the child had finished? (2) Chapter 13 — a process blocked in a read that never completes is off the CPU; name one thing that could wake it besides the data arriving. (3) *Predict*: a program installs a handler for `SIGUSR1`, then blocks that signal, sends it to itself *five* times in a tight loop, and finally unblocks it. How many times does the handler run — five, one, or zero? Write down your guess and why — we will measure it in §18.4.

So far a program's control flow has been its own: it executes its instructions, calls functions, makes system calls, and returns — a single thread of cause and effect (threads in Chapter 15 added more such flows, but each was still running *its own* code). A **signal** breaks that assumption. It is an **asynchronous notification** the kernel delivers to a process to tell it that an event happened — a software analog of a hardware interrupt — and when one arrives, the process's normal flow is suspended, a designated **handler** runs, and then the normal flow resumes where it left off. You have already met signals without naming them: the segmentation fault (`SIGSEGV`) of Chapter 11, the `Ctrl-C` that kills a program (`SIGINT`), the `SIGCHLD` that told a parent its child had exited (Chapter 14). This chapter makes the mechanism precise: what a signal is and its default dispositions (§18.1), how to catch one with a handler (§18.2), how signals interrupt blocking system calls — the `EINTR` you must handle (§18.3) — and the sharp hazards of running code asynchronously, namely `async-signal-safety` and the fact that standard signals do not queue (§18.4). We close by placing signals among the OS's communication mechanisms (§18.5). Signals are the kernel's way of interrupting your program to tell it something; the danger is that "at almost any instant" means the handler can run in the middle of anything.

18.1 What a signal is, and its default disposition

A **signal** is a small integer, identified by name, delivered to a process to indicate an event. Some are sent by the kernel in response to hardware or program conditions — `SIGSEGV` (invalid memory access), `SIGFPE` (arithmetic error like divide-by-zero), `SIGCHLD` (a child changed state). Some are sent by users or programs to control a process — `SIGINT` (from `Ctrl-C`), `SIGTERM` (the polite "please exit," what `kill` sends by default), `SIGKILL` (the un-catchable "die now"), `SIGUSR1/SIGUSR2` (application-defined). A process can send a

signal to another (subject to permissions) with the `kill` system call — poorly named, since it sends *any* signal, not only lethal ones.

Every signal has a **default disposition** — what happens if the process does nothing about it. The common defaults are: **Terminate** the process (`SIGTERM`, `SIGINT`); **Terminate and dump core** (`SIGSEGV`, `SIGABRT` — the core file for a debugger); **Ignore** the signal (`SIGCHLD`'s default is to be ignored — which is why an un-reaped child still becomes a zombie, per Chapter 14; the notification is delivered but discarded); and **Stop / Continue** the process (`SIGSTOP`/`SIGCONT`, job control). A process can *change* the disposition of most signals — catch them with a handler, or ask to ignore them — with two crucial exceptions: **SIGKILL and SIGSTOP can neither be caught nor ignored**. That is by design: the OS must retain an unblockable way to kill or halt any process, so that a program cannot make itself immortal by trapping every signal.

QUICK CHECK

Name the four broad default dispositions a signal can have. Which two signals can a process *not* catch or ignore, and why does the OS insist on that? (Answer after the next section.)

18.2 Catching a signal: handlers and `sigaction`

To catch a signal, a process installs a **signal handler**: a function the kernel will call when that signal is delivered. The portable, well-defined way to install one is the `sigaction` system call (the older `signal()` function has historically ambiguous semantics across systems; prefer `sigaction`). You fill in a `struct sigaction` with your handler function and flags, and register it for a given signal number:

```
void handler(int signum){ /* runs asynchronously when the signal arrives */ }

struct sigaction sa;
memset(&sa, 0, sizeof sa);
sa.sa_handler = handler;
sigaction(SIGINT, &sa, NULL); /* from now on, Ctrl-C runs handler() instead of terminating */
```

Once installed, the handler runs *asynchronously*: when `SIGINT` arrives, the kernel suspends whatever the process was doing, calls `handler`, and — when the handler returns — resumes the interrupted code exactly where it left off. The program did not *call* the handler; the kernel injected it into the control flow. This is the defining strangeness of signals and the source of every hazard in §18.4: the handler can begin executing between any two instructions of your main code, including in the middle of a library call. The handler's signature is fixed (`void handler(int)`, receiving the signal number so one handler can serve several signals), and while it runs, the same signal is (by default) blocked so the handler does not interrupt itself.

QUICK CHECK

Which system call installs a signal handler the well-defined way, and why is the older `signal()` discouraged? In what sense does the program *not* call its own handler? (Answer below the communicate box.)

18.3 Signals interrupt system calls: `EINTR`

Here is the consequence that bites every real program. Suppose a process is blocked in a slow system call — `read` waiting for input that has not arrived, `wait` for a child, `sleep` — and a signal arrives that the process catches. The kernel cannot both run the handler and leave the process cleanly blocked, so it interrupts the system call: it runs the handler, and then the blocked call returns early with the error `EINTR` ("interrupted system call"). The call did not fail for any reason of its own — the data may still be coming — it was simply cut short to deliver the signal. Here is a program that blocks in `read` on an empty pipe, arranges `SIGALRM` to fire after one second (handler installed *without* the restart flag), and reports what `read` returns (gcc 11.4.0, Linux 6.18):

```
alarm(1);                                /* deliver SIGALRM in 1 second */
ssize_t n = read(fd[0], buf, sizeof buf); /* blocks; nobody writes to the pipe */
/* ... one second later the signal fires, the handler runs, and read returns: */

read returned -1, errno=4 (Interrupted system call)
```

The `read` returned `-1` with `errno == EINTR`, not because the pipe failed but because the signal interrupted it. A correct program must expect this: any slow system call that can be interrupted by a caught signal may return `EINTR`, and the usual fix is to **retry the call** (the classic `while ((n = read(...)) < 0 && errno == EINTR); loop`). Alternatively you can install the handler with the `SA_RESTART` flag, which asks the kernel to *automatically restart* many interrupted system calls instead of failing them with `EINTR` — but the coverage is incomplete (some calls still return `EINTR` even with `SA_RESTART`), so robust code handles `EINTR` regardless. The naive assumption that a system call either blocks until it succeeds or fails for a real reason is simply wrong once signals are in play.

THE SAME IDEA THREE WAYS: A CAUGHT SIGNAL CUTS A BLOCKING CALL SHORT


1. In words. A process blocked in a slow system call cannot stay cleanly blocked while a handler runs, so the kernel interrupts the call: it runs the handler, then the call returns `-1` with `errno == EINTR`. The operation did not truly fail; it was cut short to deliver the signal, so the program must retry it (or use `SA_RESTART`, with caveats).

2. As a picture. Figure 18.1: main thread blocked in `read` → signal arrives → handler runs → control returns to `read`, which unwinds and returns `EINTR` rather than resuming the block.

3. As numbers. A `read` on an empty pipe with a `SIGALRM` after one second (no `SA_RESTART`): `read` returned `-1`, `errno=4` (Interrupted system call). With a correct retry loop it would resume waiting; with `SA_RESTART` the kernel would restart it for you.

(Answer to the §18.1 quick check: the four default dispositions are *Terminate*, *Terminate-and-dump-core*, *Ignore*, and *Stop/Continue*. `SIGKILL` and `SIGSTOP` cannot be caught or ignored, so the OS always retains a way to kill or halt any process — a program cannot trap its way to immortality.)

A signal injects a handler into the control flow — here, interrupting a blocked read.



The handler ran between the block and the return; the read did not fail on its own — it was interrupted.

Figure 18.1: A caught signal interrupts a blocking system call. The kernel suspends the block, runs the handler, then the interrupted call returns `-1` with `errno == EINTR` instead of continuing to wait — which the program must handle (retry) or opt out of with `SA_RESTART`.

18.4 The hazards: async-signal-safety and non-queuing

Because a handler can run between any two instructions of your program — including inside a library function that is halfway through updating its own state — the code a handler may safely execute is sharply restricted. A function is **async-signal-safe** if it is safe to call from a signal handler, which essentially means it is reentrant: it does not rely on global state that a signal could catch mid-update. Most of the standard library is *not* async-signal-safe. The classic trap is calling `printf` (or `malloc`) in a handler: if the signal arrives while the main program is already inside `printf` or `malloc` — holding an internal lock or partway through a data structure — the handler's second call to the same function re-enters it and can deadlock (waiting on a lock the interrupted call holds) or corrupt the heap. The kernel documents the exact set of functions guaranteed safe (the `signal-safety(7)` man page); it is a short list — `write`, `_exit`, and a handful of others — and `printf`/`malloc`/most of `libc` are not on it. The safe idiom is to do almost nothing in a handler: set a flag of type `volatile sig_atomic_t` and return, then let the main loop notice the flag and do the real work synchronously.

The second hazard is delivery semantics, and it settles the opener's prediction. Standard signals (numbers 1–31) are **not queued**: a signal is either pending or not, a single bit, so if the same signal is generated several times while it is blocked, the process still sees it *once* when unblocked — the extra instances coalesce and are lost. Here is the opener's exact program — install a `SIGUSR1` handler that counts calls, block `SIGUSR1`, send it to ourselves five times, then unblock:

```
sigprocmask(SIG_BLOCK, ...);      /* block SIGUSR1 */
for (int i=0;i<5;i++) raise(SIGUSR1); /* send it 5 times while blocked */
/* handler count so far = 0 (blocked, so not delivered yet) */
sigprocmask(SIG_UNBLOCK, ...);    /* unblock: the pending signal is delivered */

sent 5 while blocked; handler count so far = 0
after unblock, handler ran 1 time(s)
```

The handler ran **once**, not five times: the five sends collapsed into a single pending bit. This is why you must never use a standard signal as a counter or a reliable event stream — if two children exit at nearly the same moment, the parent may get a single `SIGCHLD`, so a correct `SIGCHLD` handler reaps *all* available children in a loop (`while (waitpid(-1, ..., WNOHANG) > 0)`), not one per signal. (POSIX real-time signals, `SIGRTMIN` through `SIGRTMAX`, *do* queue and can carry a small payload, at the cost of more machinery — the exception that proves the rule.) The robust modern pattern for turning asynchronous signals into something a normal event loop can consume is the **self-pipe trick** (the handler does one `async-signal-safe write` to a pipe the main loop reads) or Linux's `signalfd`, both of which move the real work out of the handler and back into the synchronous main flow.

TRAP: DOING REAL WORK IN A SIGNAL HANDLER

The natural instinct on catching a signal — log a message, update a data structure, free some resources, flush state — is exactly what you must not do, because a handler runs *asynchronously, in the middle of arbitrary code*. Calling anything not on the `async-signal-safe` list (which is nearly all of `libc`, including `printf`, `malloc/free`, and most of the standard library) risks a deadlock or heap corruption if the signal interrupted that very function. The bug is vicious because it is *rare*: it only manifests when the signal happens to land inside the unsafe call, so the program runs fine in testing and hangs or crashes in production under load — the same "passes the test, broken anyway" pathology as the data race of Chapter 15, and for the same reason (a schedule-dependent window). The discipline: a handler should do the minimum `async-signal-safe` thing — set `volatile sig_atomic_t got_signal = 1`; (or write one byte down a self-pipe) and return — and the main loop does the real work when it is safe. If you find yourself wanting to `printf` from a handler to debug it, that itch is the bug.

SAY IT LIKE A SYSTEMS ENGINEER

Try it first. A teammate says: "For graceful shutdown I added a SIGTERM handler that logs 'shutting down,' flushes the in-memory cache to disk, closes the DB connections, and calls exit — all right there in the handler. Clean, right?" Cover the paragraph and respond in three or four sentences, drawing on §18.4.

Model answer. "The intent is right but the handler is doing far too much: it runs asynchronously, possibly in the middle of a malloc or an fprintf, and logging, flushing, and closing all call functions that are not async-signal-safe — so if SIGTERM lands at the wrong instant, the handler deadlocks on a lock the interrupted code holds or corrupts the heap, and your 'graceful' shutdown hangs the process precisely when you're trying to stop it cleanly. The safe pattern is to make the handler trivial — set a volatile sig_atomic_t shutdown_requested = 1; (or write a byte to a self-pipe) and return — and let the main loop, on seeing the flag, do the flush, close, and exit synchronously where all of libc is safe to call. That also lets you bound and order the shutdown steps instead of racing them inside an interrupt. Keep handlers to a flag; do the work in the main flow." Naming async-signal-safety and moving the work out of the handler is the senior register.

(Answer to the §18.2 quick check: *sigaction* is the well-defined way to install a handler; the older *signal()* is discouraged because its semantics — handler persistence, whether the signal is blocked during its own handler, and system-call restart behavior — historically vary across systems. The program does not call its handler: the kernel **injects** it into the control flow at delivery time, suspending whatever was running and resuming it when the handler returns.)

18.5 What to carry forward

A **signal** is an asynchronous notification the kernel delivers to a process — a software interrupt — and each has a **default disposition** (terminate, terminate-and-core, ignore, stop/continue), with SIGKILL and SIGSTOP the two that can never be caught or ignored. A process catches a signal by installing a handler with *sigaction*; the handler runs asynchronously, injected into the control flow, and returns to where the program was interrupted. Two consequences dominate correct use. First, a caught signal **interrupts blocking system calls**, which return -1 with `errno == EINTR` (we measured a `read` cut short to `errno 4`) — you must retry, or use `SA_RESTART` with care. Second, a handler runs in the middle of arbitrary code, so it may call only **async-signal-safe** functions (not `printf`/`malloc`/most of `libc`), and standard signals **do not queue** — five rapid SIGUSR1s coalesced into a single delivery in our measurement — so the safe pattern is a handler that sets a volatile `sig_atomic_t` flag (or writes a self-pipe) and defers the real work to the synchronous main loop.

Signals complete the picture of how the operating system talks to a process: the system-call boundary (Chapter 13) is how a process asks the kernel for things synchronously; signals are how the kernel (or another process) interrupts a process asynchronously to tell it something happened. But a signal carries almost no information — just its number (real-time signals add a small payload) — so it is a doorbell, not a conversation. When processes need to exchange real *data* rather than notifications, they need genuine inter-

process communication: pipes and FIFOs for byte streams, shared memory (built, as Chapter 17 showed, on shared mappings in the page tables) for the fastest sharing, and sockets for talking across machines. That IPC toolkit — the channels that carry data between the isolated address spaces this part has spent five chapters building and protecting — is where the operating system part of this volume finishes, and it is the bridge to the low-level I/O and networking that the data systems in later volumes are made of.

CAN YOU... WITHOUT LOOKING?

- Define a signal as asynchronous control flow, and name its four default dispositions.
- Say which two signals can never be caught or ignored, and why the OS insists on that.
- Install a handler with `sigaction` from memory, and explain in what sense the kernel "injects" the handler.
- Explain `EINTR`: why a caught signal makes a blocking call return it, and the two ways to cope.
- Explain `async-signal-safety` and why `printf` in a handler is a bug, and give the safe handler pattern.
- Explain why standard signals do not queue, and what that means for a `SIGCHLD` handler reaping children.

RETRIEVAL — CLOSED BOOK

Cover the chapter. Answer from memory; rate confidence first.

1. **R1.** What is a signal, and what are the four default dispositions? (§18.1)
2. **R2.** Which two signals cannot be caught or ignored, and why? (§18.1)
3. **R3.** What does `sigaction` do, and in what sense does a handler run asynchronously? (§18.2)
4. **R4.** What is `EINTR`, when does it occur, and how do you handle it? (§18.3)
5. **R5.** What does `async-signal-safe` mean, and why do standard signals not queue? (§18.4)

FURTHER READING

- **Bryant & O'Hallaron, *Computer Systems: A Programmer's Perspective* (CS:APP), Chapter 8 (Exceptional Control Flow), §8.5 (Signals).** Signals as exceptional control flow, handlers, blocking, and the concurrency hazards of handlers — the reference for §18.1–18.4.
- **Stevens & Rago, *Advanced Programming in the UNIX Environment*, 3rd ed. (2013), Chapter 10 (Signals).** The canonical, thorough treatment of signal semantics, `sigaction`, reliable signals, and interrupted system calls. [[△](#) reference — not free]
- **The Linux `signal(7)`, `sigaction(2)`, and `signal-safety(7)` manual pages.** Primary sources for the signal list and default dispositions, the `SA_RESTART/EINTR` behavior, and the exact set of `async-signal-safe` functions. Verify per system.
- **The Linux `signalfd(2)` manual page and the "self-pipe trick" (D. J. Bernstein).** The documented ways to turn asynchronous signal delivery into a synchronous, event-loop-friendly readable stream — the robust alternative to doing work in a handler.

Exercises

- 18.1** **WARM-UP** In one sentence, what is a signal, and how does it differ from a normal function call in how it enters the program?
- 18.2** **WARM-UP** Name the four broad default dispositions a signal can have, with one example signal for each if you can.
- 18.3** **WARM-UP** Which two signals can a process neither catch nor ignore, and what would go wrong if it could?
- 18.4** **CORE** A program blocks in `read` waiting for input and catches `SIGINT` with a handler. Explain, step by step, what happens when the user presses `Ctrl-C`: what the `read` returns, what `errno` is, and why the call did not simply resume waiting. Give the two standard ways to make the program robust to this.
- 18.5** **CORE** Explain why calling `printf` (or `malloc`) inside a signal handler is unsafe. What property must a function have to be callable from a handler, and what is the safe pattern for a handler that needs to trigger real work?
- 18.6** **CORE** Explain why standard signals do not queue, and why this means a `SIGCHLD` handler must reap children in a loop (with `WNOHANG`) rather than reaping exactly one child per delivered signal. Tie this back to Chapter 14's zombies.
- 18.7** **COST** (a) Why is a handler that only sets a volatile `sig_atomic_t` flag both cheap and safe, compared with one that does the work directly? (b) What is the latency cost of the self-pipe / `signalfd` approach versus a direct handler, and what correctness benefit do you buy with it? (c) Why can a signal-based design never be a reliable *counter* of events under load?
- 18.8** **FADED** *Worked, with the last step for you.* A long-running server installs a `SIGHUP` handler to reload its config. The handler re-reads the config file, rebuilds several in-memory structures with `malloc`, and logs progress with `fprintf`. In testing it works; in production, under heavy request load, the server occasionally hangs completely when config is reloaded.
- Step 1 (given):** The handler runs asynchronously and can interrupt the main code at any instant, including while the main code is inside `malloc` or `fprintf`.
- Step 2 (given):** `malloc` and `stdio` hold internal locks / mutable state; they are not async-signal-safe.
- Step 3 (given):** Under heavy load the odds rise that `SIGHUP` lands while a request thread is inside one of those functions, so the handler's own call to it re-enters and deadlocks.
- Step 4 (you):** Name the bug precisely, explain why it appears only under load (not in tests), and describe the redesign that fixes it — what the handler should do and where the reload work should actually run.

18.9 **MIXED** For each need, say whether a *signal* is the right mechanism or whether you want real IPC (a pipe, shared memory, or a socket), and justify in a line: (a) tell a running daemon to re-read its configuration; (b) stream a megabyte of data from one process to another on the same host; (c) notify a parent that a child has exited; (d) send a structured request and get a structured reply between two processes.

18.10 **MIXED** Drawing on Chapters 13–18, classify each statement as *true* or *false* and fix the false ones: (a) "A signal handler is called by the program the way any other function is." (b) "If you install a handler with `SA_RESTART`, you never have to worry about `EINTR`." (c) "Sending a process `SIGKILL` lets it run cleanup code first." (d) "A segmentation fault is the delivery of the `SIGSEGV` signal."

18.11 **STRETCH** A signal handler and the main program both touch a shared `int` `running` flag. (a) Why must that flag be declared `volatile sig_atomic_t` rather than a plain `int` — what could the compiler or the hardware otherwise do? (b) Compare this hazard with the data race of Chapter 15: how is a handler-vs-main-flow conflict similar to a thread-vs-thread race, and how is it different (single thread, asynchronous interruption)? (c) Explain why this makes signals a form of concurrency even in a single-threaded program.

18.12 **LAB** On your own machine: (a) write a program that blocks in `read` on an empty pipe, sets an `alarm`, and prints what `read` returns and `errno` when `SIGALRM` fires — confirm you see `EINTR` (`errno` 4); then reinstall the handler with `SA_RESTART` and observe the difference. (b) Write a program that installs a counting handler for `SIGUSR1`, blocks it, `raises` it five times, unblocks, and prints how many times the handler ran — confirm it is once, not five. (c) Explain both results in terms of interrupted system calls and non-queuing standard signals. Label your compiler and kernel version.

Chapter 19 — Inter-Process Communication: Pipes, Shared Memory, and Sockets

Before you start: Chapter 14 (fork gives each process its own *isolated* address space — this chapter is how those isolated processes talk), Chapter 17 (a shared mapping lets two page tables point at the same physical frame — that is the machinery under shared memory), Chapter 18 (a signal is a doorbell, not a conversation — it notifies but carries no data). · Chapter 14 spent real effort keeping processes apart; this chapter is the controlled set of channels that let them exchange data *across* that isolation, and it is where Part III finishes.

BEFORE YOU READ ON

From memory, reaching back: (1) Chapter 17 — copy-on-write and shared mappings both rest on one fact about page tables; what lets two processes' page tables refer to the *same* physical frame? (2) Chapter 18 — why did we call a signal a "doorbell, not a conversation," and what does that tell you it is bad at? (3) *Predict*: two related processes `mmap` the same page with `MAP_SHARED`; one writes the value 42 into it and exits, the other reads. Does the reader see 42? Now suppose the mapping had instead been `MAP_PRIVATE` — same question. Write down both guesses; we measure them in §19.3.

A process is an island. Chapter 14 built that isolation deliberately: each process gets its own virtual address space, so a wild pointer in one cannot corrupt another, and Chapter 17 showed the page tables that enforce it. Isolation is a safety property you want — and an obstacle the moment two processes must cooperate. A shell pipeline (`sort | uniq`), a web server handing a request to a worker, a database client talking to a server: all need to move data between address spaces that, by construction, cannot see each other's memory. **Inter-process communication** (IPC) is the family of kernel-provided channels that bridge the gap. This chapter covers the three that matter most, along the axis of *what* they move and *how far*: **pipes and FIFOs**, the unidirectional byte stream that powers the shell (§19.2); **shared memory**, the fastest channel because it moves nothing — the processes map the same frames (§19.3); and **sockets**, the bidirectional channel that, unlike the others, also reaches *across machines* (§19.4). We start by placing them on a map (§19.1) and close by carrying the toolkit forward into concurrency and networking (§19.5). The through-line: every channel trades cost against reach against safety, and shared memory in particular gives you speed by handing back a hazard you have met before.

19.1 Why IPC exists, and the shape of the toolkit

Because address spaces are private, one process cannot simply pass a pointer to another: the same virtual address means different physical memory (or nothing) in the other process, so a raw pointer is meaningless across the boundary. Every IPC mechanism is therefore either a **channel the kernel copies bytes through** or an **arrangement to share the same physical memory**. That single distinction organizes the whole toolkit. A *copying* channel — a pipe, a FIFO, a socket — is safe and simple: the sender hands bytes

to the kernel, the kernel buffers them, the receiver reads them out, and neither ever touches the other's memory. A *sharing* arrangement — shared memory — is fast because it skips the copy entirely: both processes map one region and read and write it directly, at memory speed, but they inherit the full burden of coordinating that concurrent access themselves.

Cutting the other way is **reach**. Pipes, FIFOs, shared memory, and Unix-domain sockets are *same-machine* mechanisms: they rely on a shared kernel, so they cannot connect two processes on different computers. Network sockets (TCP/UDP) are the exception — the same socket API that talks to a process next door talks to one across the world, at the price of copies, protocol overhead, and the possibility of failure in between. And cutting a third way is **message shape**: some channels preserve the boundaries between writes (a datagram socket delivers each message whole), while others are pure byte streams that erase them (a pipe or a TCP connection). Keep those three axes — copy vs. share, local vs. remote, stream vs. message — in mind; every mechanism below is a point in that space, and choosing well is mostly a matter of locating your need on the map.

QUICK CHECK

Why can't one process simply pass a pointer to another and have it work? Name the single distinction that splits the IPC toolkit in two, and say which side "moves no data." (Answer after the next section.)

19.2 Pipes and FIFOs: the byte stream

A **pipe** is the oldest and simplest IPC channel: a unidirectional, in-kernel byte buffer with two file descriptors, a read end and a write end. The `pipe(fd)` system call fills `fd[0]` (read) and `fd[1]` (write); bytes written to `fd[1]` come out of `fd[0]` in order. On its own a pipe connects nothing — the trick is `fork`: a parent creates the pipe, forks, and now parent and child share both descriptors (file descriptors survive `fork`, Chapter 14), so one can close the read end and one the write end, leaving a one-way channel between relatives. That is exactly how the shell builds `a | b`: it pipes `a`'s standard output to `b`'s standard input. A **FIFO** (named pipe) is the same buffer given a *filesystem name* with `mkfifo`, so that *unrelated* processes — with no common ancestor to inherit descriptors from — can rendezvous by opening the same path.

The defining property of a pipe is that it is a **byte stream with no message boundaries**. The kernel does not remember how many writes produced the bytes; it just concatenates them. Two writes of three bytes each are indistinguishable from one write of six. Here is that fact measured (gcc 11.4.0, Linux 6.18): write `"ABC"` then `"DEF"` to a pipe, then do a single read:

```
write(fd[1], "ABC", 3);
write(fd[1], "DEF", 3);
ssize_t n = read(fd[0], buf, sizeof buf);

two writes of 3 bytes; one read got 6 bytes: "ABCDEF"
```

One read returned all six bytes at once: the boundary between the writes is gone. This is not a bug — it is what a stream *is* — but it means that if you need messages (records, requests), you must impose them yourself with **framing**: length prefixes, delimiters, or fixed-size records. (Small writes are still atomic up to a limit: POSIX guarantees a write of at most `PIPE_BUF` bytes — 4096 on Linux — will not be interleaved with another writer's, which is why the boundary-erasing above is about *your own* back-to-back writes, not corruption.) A pipe is also **finite**: it holds a bounded amount before a writer must wait. Filling a non-blocking pipe one byte at a time until it refuses shows the capacity:

```
while (write(fd[1], &c, 1) == 1) total++;

pipe full after 65536 bytes (errno=11 EAGAIN)
```

The kernel's pipe buffer is **65536 bytes** (16 pages of 4096) by default on Linux — matching the documented capacity in `pipe(7)`, adjustable with `fcntl`'s `F_SETPIPE_SZ`. This bound gives pipes their flow control: a fast writer *blocks* when the buffer is full until the reader drains it, and a reader *blocks* on an empty pipe until data arrives — so a pipeline naturally paces its stages to the slowest one. Two more behaviors complete the model. When *all* write ends are closed, a `read` returns **0** — end-of-file — which is how the reader learns the writer is done. And writing to a pipe whose *read* end is all closed raises **SIGPIPE** (default: terminate), or, if you ignore that signal, fails the `write` with `errno == EPIPE` — the kernel's way of telling a writer that nobody is listening, so it should stop rather than pour bytes into a void.

QUICK CHECK

A pipe delivered six bytes from two three-byte writes as one read. What property of a pipe does that show, and what must you add on top if you need discrete messages? What does a `read` return when every write end is closed? (Answer below the communicate box.)

19.3 Shared memory: the fastest channel moves nothing

Pipes and sockets copy: bytes go from the sender's memory into a kernel buffer and out again into the receiver's memory — two copies and a system call per exchange. **Shared memory** is the mechanism that removes the copy by removing the boundary. Two processes arrange for a region of their address spaces to be backed by the *same physical frames* — precisely the shared-mapping machinery of Chapter 17, where two page tables point at one frame — and thereafter one process's store to that region is the other's load. Nothing traverses the kernel per access; communication runs at the speed of memory. The arrangement comes in a few forms: an *anonymous* `mmap` with `MAP_SHARED` inherited across `fork` (for related processes); POSIX shared memory (`shm_open` to get a named object, then `mmap` it) for unrelated processes; and the older System V `shmget/shmat` interface. All reach the same end: one region, many mappers.

The opener's prediction pins down what "shared" means and what it does not. Map one page `MAP_SHARED|MAP_ANONYMOUS` and another `MAP_PRIVATE|MAP_ANONYMOUS`, `fork`, let the child write 42 into both, and have the parent read them back:

```
int *shared = mmap(NULL, 4096, ..., MAP_SHARED | MAP_ANONYMOUS, -1, 0);
int *private = mmap(NULL, 4096, ..., MAP_PRIVATE | MAP_ANONYMOUS, -1, 0);
/* child: *shared = 42; *private = 42; exit */

after child wrote 42: parent sees shared=42, private=0
```

The parent sees `shared == 42` but `private == 0`. The `MAP_SHARED` page is one frame both processes map, so the child's write is visible to the parent; the `MAP_PRIVATE` page is copy-on-write, so the child's write faulted a private copy the parent never sees (the same copy-on-write from `fork` in Chapter 14). Shared memory is `MAP_SHARED`: genuine sharing, zero copies. But here is the catch, and it is the whole reason concurrency is the next part of this book. Shared memory gives you the sharing and *nothing else* — no ordering, no mutual exclusion. Two processes writing the same shared counter with no coordination is **exactly the data race of Chapter 15**, now across process boundaries rather than threads. Measured with two processes each incrementing a shared counter a million times, no synchronization (gcc 11.4.0 at -O0, so each ++ is a real load-add-store):

```
two processes each +1000000 on a shared counter, no lock: final=1305815 (expected 2000000)
```

Roughly 700,000 updates vanished — the same lost-update interleaving as Chapter 15's threads, for the same reason: `counter++` is not atomic, and two flows of control raced through its load-add-store. Shared memory hands you raw, unsynchronized concurrency; making it correct needs the locks, condition variables, and atomics of the next part. The map's lesson is sharp: shared memory is the fastest channel and the most dangerous, because it is the only one that does not serialize access through the kernel for you.

THE SAME IDEA THREE WAYS: SHARED MEMORY SHARES FRAMES, NOT SYNCHRONIZATION

1. In words. A `MAP_SHARED` region is backed by physical frames that appear in both processes' page tables, so a write by one is directly a read by the other — no copy, no kernel round-trip. What the kernel does *not* supply is any ordering or exclusion between those accesses; that is the caller's job.

2. As a picture. Figure 19.1: two page tables, both pointing at frame F (shared) — one write, both see it — versus two page tables pointing at private copies (copy-on-write), where a write splits them apart.

3. As numbers. `MAP_SHARED`: parent sees the child's 42. `MAP_PRIVATE`: parent still sees 0. And two processes racing a shared counter a million times each land at 1305815, not 2000000 — sharing without synchronization is the Chapter 15 data race again.

*(Answer to the §19.1 quick check: a pointer is meaningless across the boundary because the same virtual address maps to different physical memory in each process's private address space. The toolkit splits into channels the kernel **copies** bytes through and arrangements that **share** the same physical memory; the sharing side — shared memory — moves no data.)*

MAP_SHARED: two page tables, one frame — a write by either is seen by both.




Figure 19.1: Shared memory (MAP_SHARED) maps one physical frame into both processes, so a store by one is a load by the other with no copy. A private (copy-on-write) mapping instead gives each a separate copy on the first write, so changes do not cross. Sharing frames is not sharing synchronization — coordinating the concurrent access is still the program's job.

19.4 Sockets: bidirectional channels that cross machines

Pipes are unidirectional and local; the moment you need two-way traffic, or need to reach a process on another host, you reach for a **socket**. A socket is a bidirectional endpoint, created with `socket()`, that comes in two dominant flavors distinguished by *message shape*. A **stream socket** (`SOCK_STREAM`, TCP over a network) is a reliable, ordered *byte stream* — like a pipe, it has **no message boundaries**, so the same framing burden applies. A **datagram socket** (`SOCK_DGRAM`, UDP) preserves *message boundaries* — each send is one recv — but drops the reliability and ordering guarantees, so datagrams can be lost, duplicated, or reordered. The same API spans reach: with the `AF_UNIX` address family a socket is a fast **local** channel (a Unix-domain socket, which can even pass open file descriptors between processes); with `AF_INET/AF_INET6` the identical call sequence talks **across machines**. That uniformity is the socket's great virtue: one programming model from "next door" to "across the world."

What you buy with that reach is paid in copies and uncertainty. Every socket exchange, local or remote, passes bytes through kernel buffers — so, like a pipe, a socket *copies* rather than shares, and cannot match shared memory's zero-copy speed for two processes on one machine. And once the channel spans the network, the far end can vanish mid-conversation: a socket write can fail, a read can hang, a connection can reset. Those failure modes — and the `epoll/io_uring` machinery for handling thousands of them efficiently — are the subject of the low-level I/O and networking material later in this volume; here the point is placement. Use a pipe or FIFO for a local, one-way byte stream; shared memory when two local processes must exchange data at memory speed and you are prepared to synchronize; a Unix-domain socket for local bidirectional or descriptor-passing needs; and a network socket when the other process is on another machine and you accept the cost and the failure modes that come with distance.

TRAP: ASSUMING A STREAM PRESERVES YOUR MESSAGE BOUNDARIES

The single most common IPC bug is treating a **byte stream** — a pipe, a FIFO, or a TCP connection — as if it preserved the boundaries of your writes. You send a 200-byte request in one call and assume the peer's one `recv` returns exactly those 200 bytes; it does not have to. The stream may coalesce two sends into one read (as we measured for pipes: two 3-byte writes came back as one 6-byte read) or split one send across two reads. Over TCP this is guaranteed to bite eventually, because the network fragments and recombines the stream at will, so the code that "worked on localhost" corrupts messages under real traffic — the same "passes the test, broken anyway" pathology as a data race, and just as schedule- and load-dependent. The fix is **framing**: agree on a length prefix or a delimiter and loop reading until you have a full message. If you need discrete messages for free, use a *datagram* mechanism (a `SOCK_DGRAM` socket, which preserves boundaries) — but then handle loss and reordering. Never let "one write, one read" be an unstated assumption in stream code.

SAY IT LIKE A SYSTEMS ENGINEER

Try it first. A teammate says: "Our two services run on the same box and swap a lot of data, so I'll skip the socket and just put everything in a big `MAP_SHARED` region — it's basically free, right? Both sides read and write the struct directly." Cover the paragraph and respond in three or four sentences, drawing on §19.3.

Model answer. "Shared memory is the right instinct for speed — it's the only local channel with no per-message copy, because both processes map the same frames, so it can be far faster than a socket for high-volume same-host exchange. But it is *not* free: it gives you the sharing and nothing else — no ordering, no mutual exclusion — so the instant both sides write the same fields concurrently you have the Chapter 15 data race, across processes now, and we measured that losing hundreds of thousands of updates. To use it you must add your own synchronization in the shared region — a mutex or semaphore both processes can see, or lock-free structures with atomics — and think about what happens if one side crashes mid-update and leaves the data half-written. So: shared memory for the throughput, plus explicit coordination, plus a recovery story — not 'basically free.'" Naming that shared memory buys speed by handing back the synchronization problem is the senior register.

(Answer to the §19.2 quick check: the six-bytes-from-two-writes result shows a pipe is a **byte stream with no message boundaries** — the kernel concatenates writes — so to send discrete messages you must add framing (length prefixes, delimiters, or fixed records). A `read` returns **0** (end-of-file) once every write end is closed.)

19.5 What to carry forward

Processes are isolated by design, so **inter-process communication** is the kernel's set of channels across that isolation, and each channel is a point on three axes: *copy* vs. *share*, *local* vs. *remote*, and *stream* vs. *message*. A **pipe** is a unidirectional, in-kernel byte stream shared across `fork` (a **FIFO** is a pipe with a filesystem name for unrelated processes); it has no message boundaries — we watched two 3-byte writes return as one 6-byte read — is bounded (65536 bytes on Linux, giving flow control by blocking a full writer and an empty reader), signals end-of-file with a `read` of 0, and raises `SIGPIPE`/`EPIPE` when no reader remains. **Shared memory** (`MAP_SHARED`) is the fastest channel because it

copies nothing — both processes map the same frames, the Chapter 17 machinery — but it supplies sharing without synchronization, so two processes racing a shared counter lost ~700,000 of two million updates, the Chapter 15 data race again. **Sockets** are bidirectional and, uniquely, cross machines: stream sockets (TCP) are boundary-less byte streams needing framing, datagram sockets (UDP) preserve message boundaries but drop reliability, and the same API spans a local Unix-domain socket to a worldwide connection, paying copies and failure modes for the reach.

With this, Part III is complete. The operating system took one machine and multiplexed it: the system-call boundary (Chapter 13) is how a process asks the kernel for service; `fork/exec/wait` (Chapter 14) create and isolate processes; threads (Chapter 15) add flows of control inside one; scheduling (Chapter 16) shares the CPU in time; virtual memory (Chapter 17) shares physical memory across address spaces; signals (Chapter 18) let the kernel interrupt a process; and IPC (this chapter) opens controlled channels between the isolated processes the OS worked to keep apart. Two threads run out of this part. One points *forward in this volume*: shared memory just handed us unsynchronized concurrency across processes, and threads did the same within one — the data race is now a debt owed twice over. Making shared-memory concurrency *correct* — mutual exclusion with locks, waiting with condition variables and semaphores, ordering with atomics and the memory model, and the lock-free techniques that avoid locks entirely — is the concurrency-and-parallelism part that begins next. The other points toward *later volumes*: sockets, framing, and the cost of the kernel boundary are the substrate under every networked data system, and the efficient handling of many connections at once (`epoll`, `io_uring`, zero-copy) is the low-level I/O material this volume closes with. The channels are open; now we have to make the traffic on them correct.

CAN YOU... WITHOUT LOOKING?

- Explain why a raw pointer cannot be passed between processes, and state the copy-vs-share distinction that organizes the IPC toolkit.
- Describe a pipe as a unidirectional byte stream, explain why it has no message boundaries, and say what framing is for.
- State a pipe's flow-control behavior (block on full/empty), what a read of 0 means, and when `SIGPIPE/EPIPE` fires.
- Explain why shared memory is the fastest channel, and why it is also the most dangerous — tie it to the Chapter 15 data race.
- Contrast `MAP_SHARED` and `MAP_PRIVATE` across `fork` using the measured 42-vs-0 result.
- Distinguish stream and datagram sockets by message shape, and say what sockets uniquely offer that the other channels do not.

RETRIEVAL — CLOSED BOOK

Cover the chapter. Answer from memory; rate confidence first.

1. **R1.** What is IPC for, and what are the three axes a channel sits on? (§19.1)
2. **R2.** What is a pipe, why does it have no message boundaries, and what is a FIFO? (§19.2)
3. **R3.** What is a pipe's capacity and flow-control behavior, and what does a read of 0 mean? (§19.2)
4. **R4.** Why is shared memory the fastest channel, and what does it fail to provide? (§19.3)
5. **R5.** How do stream and datagram sockets differ, and what do sockets offer that pipes and shared memory cannot? (§19.4)

FURTHER READING

- **Stevens & Rago, *Advanced Programming in the UNIX Environment*, 3rd ed. (2013), Chapters 15 (Interprocess Communication), 16 (Network IPC: Sockets), and 17 (Advanced IPC).** The canonical, thorough treatment of pipes, FIFOs, shared memory, and sockets, including Unix-domain sockets and descriptor passing. [[△ reference](#) – not free]
- **Bryant & O'Hallaron, *Computer Systems: A Programmer's Perspective (CS:APP)*, Chapter 10 (System-Level I/O) and Chapter 11 (Network Programming).** Unix I/O and the sockets interface from the programmer's view — the reference for §19.2 and §19.4. [[△ reference](#) – not free]; the CMU 15-213 materials are free.
- **Beej's Guide to Network Programming** (beej.us). Free. A practical, readable introduction to the sockets API, stream vs. datagram, and the framing you must add on top of TCP.
- **The Linux `pipe(2)`, `pipe(7)`, `mmap(2)`, `shm_overview(7)`, `unix(7)`, and `socket(2)` manual pages.** Primary sources for the pipe capacity and `PIPE_BUF` atomicity, `MAP_SHARED` semantics, POSIX shared memory, Unix-domain sockets, and the socket types. Verify per kernel version.

Exercises

- 19.1** **WARM-UP** In one sentence each, name the copy-vs-share distinction and give one IPC mechanism on each side of it.
- 19.2** **WARM-UP** A pipe is unidirectional and only usefully connects processes created how? What does a FIFO add that lets unrelated processes use it?
- 19.3** **WARM-UP** What does a read on a pipe return when all write ends are closed, and what happens when you write to a pipe whose read ends are all closed?
- 19.4** **CORE** Explain why a pipe has no message boundaries, using the measured "two 3-byte writes, one 6-byte read" result. If your protocol needs discrete messages over a pipe or a TCP stream, what must you add, and give two concrete framing schemes.

19.5 **CORE** Explain why shared memory is the fastest local IPC channel *and* why it is the most hazardous. Contrast `MAP_SHARED` and `MAP_PRIVATE` across `fork` using the 42-vs-0 measurement, and tie the hazard back to Chapter 15.

19.6 **CORE** Distinguish a stream socket from a datagram socket by message shape and by guarantees. For each of these, say which you would pick and why: (a) a reliable in-order request/response protocol; (b) low-latency telemetry where an occasional lost sample is fine. What does `AF_UNIX` change versus `AF_INET`?

19.7 **COST** (a) Roughly how many copies does a byte make going from one process to another through a pipe versus through shared memory, and why does that make shared memory faster? (b) The pipe holds 65536 bytes before a writer blocks — explain how that bound provides flow control for a pipeline. (c) What is the cost you pay to get shared memory's zero copies, stated as an engineering burden rather than a number?

19.8 **FADED** *Worked, with the last step for you.* A logging service on one host receives log lines from many local producers. The team uses a single shared-memory ring buffer that all producers and the one consumer read and write directly, with no locks, "because shared memory is fast." Lines occasionally come out garbled or duplicated under load.

Step 1 (given): Shared memory shares the frames but provides no mutual exclusion or ordering.

Step 2 (given): Multiple producers advancing the ring's write index concurrently is a read-modify-write on shared state — a data race.

Step 3 (given): The race only manifests when two producers hit the index at overlapping moments, so light load looks fine and heavy load corrupts.

Step 4 (you): Name the bug precisely, explain why it appears only under load, and describe two fixes — one using a synchronization primitive placed in the shared region, and one that changes the channel choice to avoid the shared write index entirely.

19.9 **MIXED** For each need, choose the best IPC mechanism (pipe/FIFO, shared memory, Unix-domain socket, or network socket) and justify in a line: (a) a shell pipeline `grep | sort`; (b) a video frame passed 60 times a second between two processes on the same GPU host; (c) a client on a laptop calling a service in a datacenter; (d) two local daemons that must exchange requests *and* pass an open file descriptor between them.

19.10 **MIXED** Drawing on Chapters 14–19, classify each statement as *true* or *false* and fix the false ones: (a) "A TCP connection delivers each `send` as exactly one `recv`." (b) "Two processes can communicate by one passing the other a pointer into its heap." (c) "Shared memory needs no synchronization because it is just memory." (d) "A signal can carry a kilobyte of structured data from one process to another." (e) "Closing all write ends of a pipe is how the reader learns the stream is finished."

19.11 **STRETCH** A pipeline `producer | consumer` runs with the producer far faster than the consumer. (a) Explain, using the 65536-byte buffer and blocking semantics, why the producer does not run away and exhaust memory — where does it end up waiting? (b) How is this *backpressure* related to the flow control you would otherwise have to build by hand over a raw shared-memory channel? (c) Now the consumer dies. Trace what the producer observes on its next write, naming the signal and the `errno`, and explain why the OS designed it to notify the writer rather than silently discard the bytes.

19.12 **LAB** On your own machine: (a) write a program that creates a pipe, writes "ABC" then "DEF", and does one `read` — confirm you get six bytes with no boundary. (b) Fill a non-blocking pipe one byte at a time and print the capacity; confirm it near 65536 and try changing it with `fcntl(F_SETPIPE_SZ)`. (c) `mmap` one page `MAP_SHARED` and one `MAP_PRIVATE`, `fork`, have the child write both, and print what the parent sees — confirm shared changes cross and private ones do not. Label your compiler and kernel version.

VOL 1 · SYSTEMS PROGRAMMING & PERFORMANCE · PART IV

Concurrency & Parallelism

Threads, locks and their hazards, lock-free techniques, atomics, and the memory model that makes concurrent code correct (or not).

Chapter 20 — Mutual Exclusion: Locks and the Critical Section

Before you start: Chapter 15 (two threads incrementing a shared counter lost more than half their updates — the data race this chapter fixes), Chapter 19 (shared memory gave two processes raw sharing with no synchronization — the same debt), Chapter 16 (the scheduler can preempt a thread between *any* two instructions — which is why the race window is always open). · Part III showed how the OS creates concurrency; Part IV is how to make it *correct*, and it starts with the most basic tool: excluding all but one thread from a critical section.

BEFORE YOU READ ON

From memory, reaching back: (1) Chapter 15 — decompose `counter++` into the three machine steps it really is, and show one interleaving of two threads that loses an update. (2) Chapter 19 — shared memory gave two processes the sharing; what one thing did it *not* give them, that this chapter supplies? (3) *Predict*: wrapping a one-nanosecond increment in a lock and unlock, on an uncontended lock (no other thread competing), how much time do you think the lock and unlock add — about the same, $\sim 10\times$, $\sim 1000\times$? Write your guess; we measure it in §20.3.

Chapter 15 left us with a debt and Chapter 19 doubled it: whenever two flows of control — threads in a process, or processes over shared memory — read and write the same memory without coordination, the result is a **data race**, which loses updates nondeterministically and is undefined behavior in C and C++. The fix has a name and a shape. The updates were lost because `counter++` is really *load*, *add*, *store*, and two threads interleaved those steps; the cure is to ensure that once a thread begins that sequence, no other thread may begin it until the first finishes — to make the sequence a **critical section** that runs with **mutual exclusion**. This chapter builds that guarantee: what a critical section is and why it needs mutual exclusion (§20.1), the **mutex** that provides it and what has to exist underneath (§20.2), the cost of locking — uncontended and contended (§20.3) — and the sharpest failure mode locks introduce, **deadlock**, with the four conditions that produce it (§20.4). We close by carrying locks forward to the coordination primitives of the next chapter (§20.5). A lock is the simplest correct answer to the data race, and every complication in this part is a reaction to a cost a lock imposes.

20.1 The critical section and mutual exclusion

A **race condition** is a situation where the correctness of a computation depends on the relative timing or interleaving of multiple flows of control. The specific race we keep meeting is on a shared read-modify-write: `counter++` loads the value, adds one, and stores it back, and if a second thread loads the same old value before the first stores, one increment is overwritten. The region of code that must not be interleaved — here, the load-add-store — is a **critical section**: code that touches shared state and produces a wrong answer if two threads execute it concurrently. The property we need over a critical section is **mutual exclusion**: at most one thread is inside it at any instant. Mutual

exclusion turns "load-add-store" from a sequence three other threads can climb into the middle of into an *indivisible* step, restoring the intuition that `counter++` just adds one.

Here is the debt and its repayment measured on the same workload — four threads, each incrementing a shared counter one million times, so the correct total is 4,000,000 (gcc 11.4.0, Linux 6.18):

```
no lock:    counter=1187737 (expected 4000000)
no lock:    counter=1144730 (expected 4000000)
no lock:    counter=2766329 (expected 4000000)
with mutex: counter=4000000 (expected 4000000)
with mutex: counter=4000000 (expected 4000000)
with mutex: counter=4000000 (expected 4000000)
```

Unsynchronized, the count is wrong and *different every run* — over half the updates lost in some runs, nondeterministically, because the interleaving depends on the scheduler. With a mutex guarding the increment, the count is exactly 4,000,000 every time. That determinism is the point of mutual exclusion: it does not make the program faster (it is slower, §20.3) — it makes it *correct*, and correct *reproducibly*, which no amount of re-running an unlocked version can achieve. As Chapter 15 warned, a racing program can even print the right answer by luck; mutual exclusion is how you stop depending on luck.

QUICK CHECK

Define a critical section and the property (mutual exclusion) you need over it. Why does the unsynchronized counter give a *different* wrong answer each run, while the locked one is always exactly right? (Answer after the next section.)

20.2 The mutex: how to lock, and what is underneath

The primitive that provides mutual exclusion is the **mutex** (mutual-exclusion lock). It has two operations: **lock** (acquire) and **unlock** (release). A mutex is either free or held by one thread; `lock` on a free mutex takes it and returns, `lock` on a held mutex *blocks* the caller until the holder unlocks. So the discipline is: lock before the critical section, unlock after, and only the one thread holding the lock runs the protected code. In POSIX threads:

```
pthread_mutex_t m = PTHREAD_MUTEX_INITIALIZER;

pthread_mutex_lock(&m);    /* enter critical section – blocks if another thread holds m */
counter++;                /* protected: at most one thread here at a time */
pthread_mutex_unlock(&m); /* leave; a waiting thread may now proceed */
```

Two disciplines matter beyond the mechanics. First, all threads that touch the shared state must lock the *same* mutex — a lock excludes only the threads that agree to take it; a critical section guarded by a lock one racer ignores is not guarded at all. Second, hold the lock for *as short a critical section as correctness allows*: while you hold it, every other thread that needs it is stalled, so the lock serializes exactly the code it covers (§20.3). There is a bonus the mutex delivers beyond exclusion: the POSIX lock and unlock

operations also **synchronize memory**, establishing an ordering (a happens-before edge) between the unlock in one thread and the next lock in another. That is why the non-atomic `counter++` inside a critical section is not merely excluded but *well-defined* — the lock supplies the ordering the memory model otherwise denies a plain shared access (the memory model itself is a later chapter in this part).

What has to exist under `lock` for it to work? The lock operation must itself be race-free — but "check if the mutex is free and take it" is a read-modify-write, the very thing we cannot do safely without help. The help comes from the **hardware**: CPUs provide atomic read-modify-write instructions — test-and-set, compare-and-swap, atomic exchange — that perform "read the old value and write a new one" as one indivisible step no other core can split. A mutex is built on one of these. (Pure-software mutual exclusion is possible — Dekker's and Peterson's algorithms achieve it with only ordinary loads and stores — but it is subtle and slow, and real locks use the hardware atomic; those atomics are a chapter of their own later in this part.) On Linux, a typical mutex takes the fast path entirely in user space with an atomic when uncontended, and only traps into the kernel — via the `futex` system call — to sleep when it must actually wait, which is why an uncontended lock is cheap but not free, the subject of §20.3.

QUICK CHECK

What are the two operations on a mutex, and what does `lock` do when the mutex is already held? Why must a lock be built on a hardware atomic instruction rather than an ordinary "if free, take it"? (Answer below the communicate box.)

20.3 The cost of a lock: uncontended and contended

A lock is not free even when no one is competing for it. Measured single-threaded — a plain increment versus the same increment wrapped in an uncontended lock and unlock (gcc 11.4.0 -O2, Linux 6.18; illustrative, machine- and run-dependent):

plain increment:	~1.0 ns/op
lock + increment + unlock:	~12 ns/op
ratio:	~12x

The uncontended lock adds roughly an order of magnitude to a trivial operation — a handful of nanoseconds for the atomic and the bookkeeping, answering the opener: not "about the same," and not 1000×, but ~10× on a bare increment. That overhead is the *floor*. The real cost arrives under **contention**, when threads actually compete, and it grows in three ways. (1) **Serialization**: the lock forces the critical section to run one thread at a time, so if a large fraction of the work is inside the lock, adding cores stops helping — this is Amdahl's law applied to a critical section, and it is why "hold the lock briefly" is a performance rule, not just style. (2) **Cache-line bouncing**: the lock word (and the data it guards) ping-pongs between cores' caches as ownership changes hands, and each transfer is a cross-core cache miss — the mechanical-sympathy cost from Part I, now paid per contended acquisition. (3) **Context switches**: when a blocking mutex

makes a waiting thread sleep, acquiring the lock later means waking it, a scheduler round-trip (Chapter 16) far more expensive than the atomic itself.

These costs drive the central design tension of locking: **coarse-grained** versus **fine-grained**. One big lock around a whole data structure is simple and easy to reason about, but it serializes all access, throttling concurrency. Many small locks — one per bucket, per node, per shard — let independent operations proceed in parallel, but they multiply the overhead, complicate the code, and, crucially, open the door to **deadlock** when a thread must hold more than one at a time (§20.4). There is no free lunch: you trade simplicity and safety against concurrency, and the right point depends on how contended the structure actually is — which, per Part I's discipline, you should *measure* before assuming.

THE SAME IDEA THREE WAYS: A LOCK SERIALIZES A CRITICAL SECTION

1. In words. A mutex admits one thread at a time to the critical section; others block at lock until the holder unlocks. That guarantees correctness (mutual exclusion) and imposes a cost (serialization plus per-acquisition overhead), so the lock should cover the smallest region that is still correct.

2. As a picture. Figure 20.1: two threads' timelines; thread B's lock blocks while A is inside the critical section, then proceeds after A's unlock — the protected region never overlaps.

3. As numbers. Uncontended, a lock+unlock turns a ~1 ns increment into ~12 ns (~10×). Contended, the four-thread counter is *correct* (4,000,000 every run) but slower than the racing version — you buy determinism with throughput.

(Answer to the §20.1 quick check: a critical section is code touching shared state that misbehaves if two threads run it concurrently; mutual exclusion means at most one thread is inside it at a time. The unlocked counter varies because the lost updates depend on the scheduler's interleaving, which differs each run; the locked counter is always exactly right because mutual exclusion makes the load-add-store indivisible, removing the interleaving entirely.)

A mutex admits one thread at a time; the other blocks at lock() until unlock().




Figure 20.1: Mutual exclusion by a mutex. While thread A holds the lock, thread B blocks inside `lock()`; only after A's `unlock()` does B enter its critical section. The protected regions are serialized — that is the correctness guarantee and, simultaneously, the throughput cost.

20.4 Deadlock and the four conditions

Fine-grained locking buys concurrency but introduces a new way to fail completely. When a thread must hold more than one lock at a time, two threads acquiring the same locks in *different orders* can each grab one and then wait forever for the other's — a **deadlock** (Dijkstra's "deadly embrace"). Here are two threads and two mutexes A and B: thread 1 locks A then B, thread 2 locks B then A. Run with opposite orders versus a consistent order (gcc 11.4.0, Linux 6.18; the deadlocking run is killed after a 3-second timeout):

opposite lock order (t1: A,B ; t2: B,A):	hangs; killed after 3s (exit 124 = timed out)
consistent lock order (both: A,B):	t1 got both; t2 got both; both threads finished

With opposite orders the program hangs: thread 1 holds A and waits for B while thread 2 holds B and waits for A — neither can proceed, forever. With a consistent order it always completes, because two threads wanting both locks can never form the cycle. That fix generalizes into the theory. Deadlock requires **all four** of the classic conditions (Coffman, Elphick & Shoshani, 1971) to hold simultaneously: **mutual exclusion** (a resource is held exclusively), **hold-and-wait** (a thread holds one resource while requesting another), **no preemption** (a resource cannot be forcibly taken from its holder), and **circular wait** (a cycle of threads each waiting on the next). Break *any one* and deadlock is impossible. The most practical break is the last: impose a **global lock order** — always acquire locks in the same agreed sequence — and no circular wait can form, which is exactly what the consistent-order run did.

Two relatives round out the failure family. **Livelock**: threads are not blocked but keep reacting to each other and make no progress (two people stepping aside in the same direction repeatedly). And **priority inversion**: a low-priority thread holds a lock a high-priority thread needs, and a medium-priority thread — runnable, so it preempts the low one under a strict-priority scheduler (Chapter 16's `SCHED_FIFO`) — keeps the low thread off the CPU, so it never releases the lock, and the high-priority thread is stuck behind both. (This is the mechanism famously diagnosed on the 1997 Mars Pathfinder mission; the standard cure is *priority inheritance*, where the lock holder temporarily inherits the priority of the highest waiter.) All three share a moral: locks compose badly, and the more locks you hold at once, the more carefully you must reason about the order and the priorities.

TRAP: INCONSISTENT LOCK ORDERING (THE DEADLY EMBRACE)

The classic deadlock bug is two code paths that acquire the same two locks in *opposite* orders — one `lock(A); lock(B)`, another `lock(B); lock(A)`. Each thread grabs its first lock, then blocks forever on the second, which the other holds: total, permanent stall. The bug is vicious for the usual reason — it needs the two paths to interleave at just the wrong moment, so it is *rare*: the code passes tests and every demo, then one day under load two requests hit the two paths simultaneously and the service freezes with no crash, no error, no log — just threads stuck in `lock()`. It is especially insidious when the two locks are acquired far apart in the call graph, or hidden inside library calls, so nobody sees both orders in one place. The discipline: define a **global lock ordering** and always acquire multiple locks in that order (break circular wait), hold as few locks as possible for as short a time as possible (attack hold-and-wait), and use `trylock` with a back-off if you truly cannot order them. If you ever write `lock(X)` while already holding `lock(Y)`, you owe yourself a proof that no other path ever takes them the other way.

SAY IT LIKE A SYSTEMS ENGINEER

Try it first. A teammate says: "Our service froze in production — no crash, no error, CPU near zero, it just stopped serving. I restarted it and it was fine for a week, now it's frozen again. Must be a hardware glitch, right?" You have a thread dump showing several threads stuck in `pthread_mutex_lock`. Cover the paragraph and respond in four or five sentences, drawing on §20.4.

Model answer. "That is the signature of a *deadlock*, not hardware: CPU near zero (nobody is spinning, everyone is blocked), no crash or error (deadlock is a permanent wait, not a fault), and threads parked in `pthread_mutex_lock` — which the dump confirms. It reproduces rarely because it needs two code paths to acquire the same locks in opposite orders and to interleave at just the wrong instant, so it hides for a week and then, under the right concurrent load, forms a cycle: each thread holds one lock and waits on the other forever. The fix is to establish a *global lock ordering* so every path takes those locks in the same sequence — that breaks the circular-wait condition and makes the cycle impossible — plus shrink the critical sections so we hold fewer locks at once. Let's diff the two paths in the dump to find the inconsistent order, and add a lint or a runtime lock-order check so the next one fails loudly in testing instead of silently in production." Naming the deadlock from its fingerprint — blocked threads, zero CPU, no error — and prescribing a lock order is the senior register.

(Answer to the §20.2 quick check: the two operations are `lock` (acquire) and `unlock` (release); `lock` on a held mutex **blocks** the caller until the holder releases it. A lock must be built on a hardware atomic read-modify-write because "check if free, then take it" is itself a read-modify-write — without an indivisible instruction, two threads could both see the mutex free and both "take" it, which is the very race the lock is supposed to prevent.)

20.5 What to carry forward

A **race condition** is correctness that depends on interleaving; the region that must not interleave is a **critical section**; and the property that fixes it is **mutual exclusion** — at most one thread inside at a time. A **mutex** supplies it with `lock/unlock`: we watched a four-thread counter go from a nondeterministic ~1.1–2.8 million to exactly 4,000,000

every run. A lock also *synchronizes memory*, giving the ordering that makes the guarded non-atomic accesses well-defined, and it rests on a hardware atomic instruction because "if free, take it" is itself a race. Locks are not free: an uncontended lock is $\sim 10\times$ a bare increment, and under contention the costs are serialization (Amdahl on the critical section), cache-line bouncing between cores, and context switches when waiters sleep — which is why you hold the lock for the smallest correct region and choose coarse vs. fine granularity by measurement. And locks introduce **deadlock**: all four Coffman conditions (mutual exclusion, hold-and-wait, no preemption, circular wait) must hold, so breaking one — most practically by imposing a global lock order — prevents it, as the consistent-order run showed against the hanging opposite-order one.

Locks give us *safety*: they keep threads out of each other's critical sections. What they do not give us is a way to *wait for a condition*. A consumer thread that needs data another thread has not yet produced cannot make progress by locking harder — and it must not sit spinning while holding the lock, since that both burns the CPU and keeps the producer out. What it needs is to *release the lock and sleep until the condition it is waiting for becomes true*, then wake and re-check — coordination, not just exclusion. That is the **condition variable**, and its counting cousin the **semaphore**: the next chapter builds the producer-consumer pattern on them, and shows why the wait must sit in a loop rather than an `if`. Mutual exclusion was the first half of shared-memory correctness; waiting and signalling is the second, and together they are the classical toolkit that everything lock-free and every memory-model rule later in this part is measured against.

CAN YOU... WITHOUT LOOKING?

- Define race condition, critical section, and mutual exclusion, and say how they relate.
- Write a mutex-guarded critical section from memory, and explain why every racing thread must take the *same* lock.
- Explain why a mutex must rest on a hardware atomic, and what "if free, take it" gets wrong without one.
- State the uncontended lock cost ($\sim 10\times$ a bare op) and the three ways contention makes it worse.
- Explain the coarse- vs. fine-grained locking trade-off in terms of concurrency, overhead, and deadlock risk.
- List the four deadlock conditions and give the most practical one to break, with the fix that breaks it.

RETRIEVAL — CLOSED BOOK

Cover the chapter. Answer from memory; rate confidence first.

1. **R1.** What is a critical section, and what property must hold over it? (§20.1)
2. **R2.** What are a mutex's two operations, and why must all racers take the same one? (§20.2)
3. **R3.** Why must a lock be built on a hardware atomic instruction? (§20.2)
4. **R4.** Name the three ways contention makes a lock expensive. (§20.3)
5. **R5.** List the four deadlock conditions and name the most practical one to break. (§20.4)

FURTHER READING

- **Arpaci-Dusseau, *Operating Systems: Three Easy Pieces (OSTEP)*, "Locks" and "Lock-based Concurrent Data Structures."** Free (pages.cs.wisc.edu/~remzi/OSTEP/). Building a lock from hardware atomics (test-and-set, compare-and-swap), spin vs. sleeping locks, and locking real data structures — the reference for §20.1–§20.3.
- **Bryant & O'Hallaron, *Computer Systems: A Programmer's Perspective (CS:APP)*, Chapter 12 (Concurrent Programming).** Threads, shared variables, mutual exclusion, and synchronization from the programmer's view. [[▲](#) reference – not free]; CMU 15-213 materials free.
- **Coffman, Elphick & Shoshani, "System Deadlocks," *ACM Computing Surveys* 3(2): 67-78 (1971), DOI 10.1145/356586.356588.** The primary source for the four necessary conditions for deadlock. Author/venue/year verified via the ACM Digital Library.
- **Dijkstra, "Cooperating Sequential Processes" (EWD123, 1965; published 1968).** The origin of the critical-section and mutual-exclusion problem, the "deadly embrace" (deadlock), and the semaphore. Free via the E. W. Dijkstra Archive (cs.utexas.edu/~EWD/).
- **The Linux `pthread_mutex_lock(3)` and `futex(2)` manual pages.** Primary sources for POSIX mutex semantics and the fast-user-space/kernel-sleep mechanism under a mutex. Verify per system.

Exercises

- 20.1** **WARM-UP** In one sentence each, define a critical section and mutual exclusion, and say why `counter++` needs the latter.
- 20.2** **WARM-UP** What are the two operations on a mutex, and what does the acquire operation do when the mutex is already held by another thread?
- 20.3** **WARM-UP** Name the four conditions that must all hold for deadlock, and give the one most practical to break in application code.
- 20.4** **CORE** Explain why a mutex must be built on a hardware atomic read-modify-write instruction. Walk through what could go wrong if `lock` were implemented as "if the mutex is free, set it to held" with ordinary loads and stores, and connect that failure to the data race of Chapter 15.
- 20.5** **CORE** Two threads each lock mutexes A and B, but thread 1 does A then B and thread 2 does B then A. (a) Show an interleaving that deadlocks. (b) Explain which of the four conditions your fix removes when you make both threads acquire in the order A then B. (c) Why does this bug typically pass tests and only appear in production?
- 20.6** **CORE** A mutex both *excludes* other threads and *synchronizes memory*. Explain each role, and why the second is what makes a non-atomic `counter++` inside a critical section not just serialized but well-defined under the C/C++ memory model. Tie this to why a plain shared access without a lock is undefined behavior (Chapter 15).

20.7 **COST** (a) An uncontended lock+unlock measured ~12 ns against a ~1 ns bare increment — explain where that ~10× goes. (b) Under heavy contention the cost is far higher; name the three sources (serialization, cache-line bouncing, context switches) and say which dominates when the critical section is tiny versus when threads must sleep. (c) Using Amdahl's reasoning, explain why shrinking the critical section can matter more than making the lock itself faster.

20.8 **FADED** *Worked, with the last step for you.* A bank-transfer function locks the source account, then the destination account, then moves money. Under load the service occasionally freezes; a thread dump shows threads blocked in `lock`.

Step 1 (given): Transfer $x \rightarrow y$ locks X then Y; a concurrent transfer $y \rightarrow x$ locks Y then X.

Step 2 (given): If both start together, the first holds X and waits for Y; the second holds Y and waits for X.

Step 3 (given): That is a circular wait — all four Coffman conditions now hold — so both threads block forever.

Step 4 (you): Name the bug, explain why it is load- and timing-dependent, and give a concrete fix that imposes a global lock order on accounts (say, by account id) so no cycle can form — and state which condition your fix removes.

20.9 **MIXED** For each scenario, decide whether the right tool is (i) a single mutex, (ii) fine-grained per-item mutexes, or (iii) no lock at all because there is no sharing — and justify: (a) a global counter incremented by all threads; (b) a hash map where operations on different buckets are independent; (c) a value each thread keeps in its own thread-local storage; (d) a rarely-updated configuration read by everyone on every request.

20.10 **MIXED** Drawing on Chapters 15–20, classify each as *true* or *false* and fix the false ones: (a) "If two threads take different locks, they are safely mutually exclusive." (b) "A deadlocked program spins the CPU at 100%." (c) "A mutex makes the guarded code faster." (d) "Holding a lock longer is always safer." (e) "You can implement a correct lock with only ordinary loads and stores and no hardware atomic."

20.11 **STRETCH** A strict priority scheduler (Chapter 16's `SCHED_FIFO`) runs three threads: low L holds a lock, high H needs it, medium M is CPU-bound and needs no lock. (a) Explain how H can end up blocked indefinitely even though it is the highest priority — name the phenomenon. (b) Why does M, not H, keep L off the CPU? (c) Describe priority inheritance and exactly which link in the chain it breaks. (d) Relate this to why holding a lock while doing unbounded work is dangerous under real-time scheduling.

20.12 **LAB** On your own machine: (a) run four threads incrementing a shared counter one million times each, first with no lock, then with a mutex — confirm the unlocked total is wrong and varies while the locked total is exactly 4,000,000. (b) Time a plain increment versus a lock+increment+unlock in a tight loop and report the ratio (expect roughly an order of magnitude). (c) Write two threads that lock two mutexes in opposite orders and observe the hang; then fix the order and confirm it completes. Label your compiler and kernel version, and note that the timing numbers are illustrative.

Chapter 21 — Coordinating Threads: Condition Variables and Semaphores

Before you start: Chapter 20 (a mutex gives mutual exclusion over a critical section, but no way to *wait* for a condition), Chapter 16 (a blocked thread is off the CPU; a runnable one competes for it — the difference between sleeping and spinning), Chapter 15 (threads share memory, so one can produce what another consumes). · Locks kept threads *out* of each other's way; this chapter is how one thread *waits* for another to make something true — coordination on top of exclusion.

BEFORE YOU READ ON

From memory, reaching back: (1) Chapter 20 — what does a mutex guarantee, and name one of the four conditions required for deadlock. (2) Chapter 16 — a thread *spinning* in a loop checking a flag versus a thread *blocked* waiting: which one competes for the CPU, and what does that cost? (3) *Predict*: a consumer waits ~0.5 seconds for a producer, first by spinning in a loop that re-checks a flag under a lock, then by blocking on a proper wait primitive. How much CPU time do you think each consumes during that half-second? Write both guesses; we measure them in §21.1.

A mutex answers "may I touch the shared state?" It does not answer "is the shared state *ready* yet?" Consider a bounded buffer between a producer and a consumer: the consumer must not dequeue from an empty buffer, and the producer must not enqueue into a full one. A lock protects the buffer's fields, but it cannot make a consumer *wait* for an item to appear — and the naive fix, spinning in a loop re-checking under the lock, is a disaster: it burns a whole CPU doing nothing and, worse, holds the lock the producer needs, so the very condition it waits for can never come true. What a waiting thread needs is to **release the lock and sleep until signalled**, then wake and re-check. This chapter builds that: why busy-waiting is wrong (§21.1); the **condition variable** — wait, signal, broadcast — and the non-negotiable rule that the wait sits in a `while` loop, not an `if` (§21.2); the **semaphore**, Dijkstra's counting primitive that *remembers* signals (§21.3); and the classic ways to get coordination wrong, plus the monitor discipline that packages it (§21.4). We close by carrying the blocking toolkit toward atomics and the memory model (§21.5). Exclusion was half of shared-memory correctness; waiting is the other half.

21.1 Waiting without wasting: why busy-waiting is wrong

The obvious way to wait for a condition is to loop until it holds — a **busy-wait**, or spin. It is almost always wrong for a thread that may wait more than a few instructions, for two reasons. First, a spinning thread is *runnable*, so the scheduler keeps giving it the CPU (Chapter 16) while it does no useful work — it burns a core. Second, and more insidious in the bounded-buffer case, if the spin holds the lock while re-checking, it locks out the very thread that would make the condition true, so the wait can never end: a self-inflicted deadlock. Here is the CPU cost measured — a thread waits ~0.5 s for a signal, first by

busy-polling a flag under a lock, then by blocking on a condition variable (gcc 11.4.0, Linux 6.18; CPU time via `getrusage`):

```
busy-poll (spins):  waited ~0.5s wall, consumed 0.499 s CPU
condvar (blocks):  waited ~0.5s wall, consumed 0.001 s CPU
```

The spinning waiter burned essentially a full core-second of CPU for half a second of wall-clock waiting — 100% of a core doing nothing — while the blocking waiter consumed a thousandth of that, because it was *off the CPU*, asleep, until woken. That five-hundred-fold difference answers the opener and states the whole motivation: a thread that must wait should **sleep**, not spin, so the CPU goes to threads with work to do. (Spinning is not always wrong — for a wait known to be shorter than a context switch, a brief spin, or a spinlock, can beat sleeping; the point is that unbounded or long waits must block.) The primitive that lets a thread sleep until a condition changes, without holding the lock while it sleeps, is the condition variable.

QUICK CHECK

Give the two reasons busy-waiting is wrong for a thread that may wait a while, and say which one turns a spin-under-lock into a deadlock. Roughly how much CPU did the spinning waiter burn versus the blocking one? (Answer after the next section.)

21.2 Condition variables: wait, signal, and the `while` rule

A **condition variable** is a queue of threads waiting for some predicate over shared state to become true. It is always paired with a mutex (the one guarding that state) and has three operations. `wait(cv, m)` must be called with the mutex held; it *atomically* releases the mutex and puts the thread to sleep on the condition variable, and when the thread is later woken it *re-acquires the mutex* before returning — so the thread sleeps without holding the lock but wakes holding it again. `signal(cv)` wakes one waiting thread; `broadcast(cv)` wakes all of them. The atomic "release-and-sleep" is the crucial part: it closes the window in which a thread could check the predicate, decide to wait, and miss a signal that arrives before it actually sleeps.

The rule that trips up everyone learning this is: **always wait in a while loop that re-checks the predicate, never a bare if**. Here is the producer-consumer bounded buffer done correctly (gcc 11.4.0, Linux 6.18) — note the `while`:

```
/* consumer */
pthread_mutex_lock(&m);
while (count == 0)           /* while, NOT if */
    pthread_cond_wait(&not_empty, &m);
item = buf[out]; out = (out+1) % CAP; count--;
pthread_cond_signal(&not_full);
pthread_mutex_unlock(&m);

produced sum=210 consumed sum=210 (1..20 sum=210); buffer never exceeded CAP=4
```

Producer and consumer agreed exactly (both sums 210, the sum of 1..20) and the buffer never exceeded its capacity of four — the coordination is correct. The `while` is mandatory for two independent reasons. (1) **Spurious wakeups**: POSIX explicitly permits `cond_wait` to return without any signal, so on waking you must re-verify the predicate actually holds. (2) **Stolen conditions**: even with a real signal, between the signal and this thread re-acquiring the mutex, *another* thread may have run and consumed the item, so the predicate you were signalled about may already be false again. In both cases the `while` simply re-checks and, if the predicate is still false, waits again — turning "I was woken" into "I woke *and* confirmed the condition," which is what correctness requires. An `if` assumes the wakeup means the predicate holds; under either hazard, that assumption is a bug.

THE SAME IDEA THREE WAYS: WAIT RELEASES THE LOCK, SLEEPS, RE-CHECKS IN A LOOP


1. In words. `cond_wait` atomically releases the mutex and sleeps; a signal wakes a waiter, which re-acquires the mutex and — because the wakeup may be spurious or the condition already re-consumed — re-tests the predicate in a `while` loop, waiting again if it is still false.

2. As a picture. Figure 21.1: the consumer holds the mutex, finds `count == 0`, calls `wait` (drops the lock, sleeps); the producer locks, enqueues, signals; the consumer wakes, re-acquires the lock, re-checks `count`, and proceeds.

3. As numbers. The bounded buffer with `while`-loops: produced and consumed sums both 210 over items 1..20, buffer never above `CAP = 4`. Blocking on the condition variable cost 0.001 s CPU versus the spinner's 0.499 s.

(Answer to the §21.1 quick check: busy-waiting is wrong because a spinning thread stays runnable and burns a CPU doing no work, and — the deadlock case — if it spins while holding the lock, it locks out the thread that would make the condition true, so the wait never ends. The spinning waiter burned ~0.499 s of CPU for a ~0.5 s wait, roughly 500× the ~0.001 s the blocking waiter used.)

`cond_wait` releases the lock and sleeps; signal wakes it; it re-acquires and re-checks.



The `while` loop guards against spurious wakeups and a condition another thread already consumed.

Figure 21.1: A condition variable. The consumer, holding the mutex, finds the buffer empty and `wait`s — atomically releasing the lock and sleeping. The producer enqueues and `signal`s; the consumer wakes, re-acquires the mutex, and re-checks the predicate in its `while` loop before proceeding.

21.3 Semaphores: counting permits that remember

A **semaphore** (Dijkstra, 1965) is an integer with two atomic operations and a waiting queue: **wait** (historically *P*; POSIX `sem_wait`) decrements the count, blocking if it is already zero, and **post** (historically *V*; POSIX `sem_post`) increments the count, waking a waiter if any. A *binary* semaphore (count 0 or 1) behaves much like a mutex; a *counting* semaphore initialized to *N* models a pool of *N* interchangeable permits. Ten threads competing for a semaphore initialized to three never exceed three inside the guarded section (gcc 11.4.0, Linux 6.18):

```
sem_init(&slots, 0, 3); /* three permits */
/* each worker: sem_wait(&slots) ... work ... sem_post(&slots) */

10 threads, semaphore initialized to 3: max concurrently in section = 3
```

The semaphore capped concurrency at exactly its initial count — the natural tool for "at most *N* at once" (a connection pool, a bounded worker set, *N* copies of a resource). The producer-consumer buffer has an elegant semaphore form too: a semaphore `empty` initialized to the buffer size counts free slots (the producer *waits* it before enqueueing), and a semaphore `full` initialized to zero counts filled slots (the consumer *waits* it before dequeuing), with a mutex protecting the buffer indices. The crucial property that distinguishes a semaphore from a condition variable is **memory**: a semaphore's count *remembers* posts. A `sem_post` with no waiter increments the count, so a later `sem_wait` succeeds immediately — the signal is not lost. A `cond_signal` with no waiter, by contrast, *vanishes*: it wakes no one and leaves no trace. That single difference is why a condition variable must always be paired with a shared-state predicate checked under the mutex (the state is the memory; the condition variable is only the wakeup), whereas a semaphore *is* its own state.

QUICK CHECK

What do a semaphore's two operations do, and what does a counting semaphore initialized to *N* give you? State the one key difference between a semaphore and a condition variable regarding a signal sent with no one waiting. (Answer below the communicate box.)

21.4 Getting it wrong: lost wakeups, *if-not-while*, and the monitor

The coordination primitives are small, and the bugs cluster around a few classic mistakes — all versions of confusing "I was woken" with "the condition holds," or of signalling without the shared state to back it. The most common is using *if* where the rule demands *while*: a single check on waking, which a spurious wakeup or a condition consumed by another thread turns into acting on a false predicate — dequeuing from a now-empty buffer, proceeding on data that is not there. Its cousin is the **lost wakeup**: signalling a condition variable *without* updating (and holding the lock over) the shared state, so a waiter that has not yet slept misses the signal and a waiter that has slept is never woken. Because a condition variable does not remember signals (§21.3), the discipline is ironclad: change the shared state under the mutex, *then* signal; and always

wait in a loop on a predicate over that state. A semaphore, which remembers, forgives the timing — but a semaphore misused as a mutex you forget to post, or posted the wrong number of times, deadlocks or over-admits just as badly.

These patterns recur so reliably that they have a structuring discipline: the **monitor** (Hoare, 1974; Brinch Hansen). A monitor bundles shared state with the mutex that protects it and the condition variables threads wait on, so that all access goes through procedures that hold the lock and coordinate through the conditions — the programmer stops hand-wiring locks and signals and instead works with an object whose methods are automatically mutually exclusive. This is exactly what a Java object's `synchronized` methods with `wait/notify` package, and what higher-level constructs like a thread-safe blocking queue or a Go channel give you: the mutex, the condition variables, and the predicate discipline, wrapped so you cannot forget the `while` loop or the "signal under the lock" rule. The lesson to carry: prefer the highest-level correct abstraction available — a concurrent queue over hand-rolled condition variables, a monitor over raw locks and signals — and drop to the primitives only when you must, because the primitives are where the subtle bugs live.

TRAP: IF INSTEAD OF WHILE AROUND COND_WAIT (AND SIGNALLING WITHOUT THE STATE)

Two coupled mistakes cause most condition-variable bugs. The first is guarding `cond_wait` with an `if` instead of a `while`: you check the predicate once, wait, and on waking assume it holds — but POSIX permits **spurious wakeups** (waking with no signal), and even a real signal can be "stolen" by another thread that re-acquires the lock first and consumes the condition, so on waking the predicate may be false. Acting on it — dequeuing from an empty buffer, using data that is not there — corrupts state or crashes, and it is rare and load-dependent, so it survives testing. The second is the **lost wakeup**: because a condition variable does *not* remember a signal sent with no waiter (unlike a semaphore's count), if you signal without holding the lock and updating the shared state, a thread that is about-to-wait can miss the signal and sleep forever. The discipline that kills both: (1) always `while (!predicate) cond_wait(...)`, never `if`; (2) always update the shared state under the mutex and then signal. If you find yourself reasoning "the wakeup means the item is there," you have the `if-not-while` bug.

SAY IT LIKE A SYSTEMS ENGINEER

Try it first. A teammate says: "My worker waits for jobs like this: `if (queue_empty) pthread_cond_wait(&cv, &m);` then it pops a job. And to avoid missing wakeups I call `pthread_cond_signal` right before I take the lock to enqueue. Simple and fast — what could go wrong?" Cover the paragraph and respond in four or five sentences, drawing on §21.2 and §21.4.

Model answer. "Two bugs, both timing-dependent, so they'll pass your tests and fail in production. First, the `if` must be a `while`: `cond_wait` can return spuriously, and even on a real signal another worker may have popped the job before you re-acquire the lock, so on waking you must re-check `queue_empty` — with `if`, you'll pop from an empty queue. Second, signalling *before* taking the lock and enqueueing is a lost-wakeup race: a condition variable doesn't remember a signal sent with no waiter, so if the worker is between its check and its `cond_wait` when you signal, it misses it and may sleep forever with a job sitting in the queue. The fix is the standard discipline — enqueue under the mutex, *then* signal while still holding it (or right after), and have the worker loop `while (queue_empty) cond_wait(...)`. Honestly, I'd wrap this in a thread-safe blocking queue so nobody has to remember these rules per call site." Naming both the `if-not-while` and the lost-wakeup, and reaching for a higher-level abstraction, is the senior register.

(Answer to the §21.3 quick check: *wait/P decrements the count, blocking if it is zero; post/V increments it, waking a waiter if any. A counting semaphore initialized to N admits at most N threads at once — a pool of N permits. The key difference from a condition variable: a semaphore **remembers** a post made with no waiter (the count rises, so a later wait succeeds), whereas a condition-variable signal with no waiter is lost.*)

21.5 What to carry forward

Locks (Chapter 20) exclude; this chapter's primitives **coordinate**. A thread that must wait for a condition should **sleep, not spin** — we measured a busy-poll burning 0.499 s of CPU against a blocking waiter's 0.001 s for the same half-second wait. A **condition variable** lets a thread atomically release its mutex and sleep until `signal/broadcast` wakes it, whereupon it re-acquires the mutex; the wait *must* sit in a `while` loop re-checking the predicate, because wakeups can be spurious and the condition can be consumed by another thread first — our bounded buffer with `while`-loops kept producer and consumer in exact agreement (sums 210) and the buffer within capacity. A **semaphore** is a counting primitive whose `wait/post` (P/V) admit at most its initial count of threads — three permits capped ten threads at three — and, unlike a condition variable, it *remembers* posts sent with no waiter, which is why a condition variable must always pair with a predicate over shared state while a semaphore is its own state. The bugs cluster on `if-not-while` and lost wakeups; the **monitor** (and its descendants — Java `synchronized/wait/notify`, blocking queues, channels) packages the discipline so you need not rebuild it per call site.

With locks, condition variables, and semaphores, you have the classical *blocking* synchronization toolkit — enough to write correct shared-memory concurrency. But it rests on assumptions this part must now examine directly. Every primitive here blocks a thread (a scheduler cost), and every one relies on the guarantee that memory operations become visible in a sensible order across cores — a guarantee that real multicore

hardware does *not* provide for plain loads and stores, and that these primitives quietly supply through the same atomics that a lock is built on. The rest of this part opens that box: the hardware **atomics** (compare-and-swap and friends) that let you build synchronization — and lock-free data structures that avoid blocking entirely — and the **memory model** that states precisely what a concurrent program means, why a data race is undefined behavior, and what ordering you are and are not entitled to assume. The primitives you now know are the correct, comprehensible baseline; the chapters ahead show what they cost, how they work underneath, and when you can do without them.

CAN YOU... WITHOUT LOOKING?

- Give the two reasons busy-waiting is wrong, and the CPU numbers that show a blocking wait beating a spin.
- Explain what `cond_wait` does atomically, and why that atomicity closes a missed-signal window.
- State the while-not-if rule and give both reasons it is mandatory (spurious wakeups; stolen condition).
- Describe a semaphore's wait/post, what a counting semaphore of N gives you, and that it remembers posts.
- State the one difference between a semaphore and a condition variable regarding a signal with no waiter.
- Explain the lost-wakeup bug and the two-part discipline (update-under-lock-then-signal; wait in a loop) that prevents it.

RETRIEVAL — CLOSED BOOK

Cover the chapter. Answer from memory; rate confidence first.

1. **R1.** Why is busy-waiting wrong, and which failure makes a spin-under-lock a deadlock? (§21.1)
2. **R2.** What does `cond_wait` do, and why must the wait be in a while loop? (§21.2)
3. **R3.** What do a semaphore's two operations do, and what does an initial count of N mean? (§21.3)
4. **R4.** How does a semaphore differ from a condition variable when a signal is sent with no waiter? (§21.3)
5. **R5.** What is a lost wakeup, and what discipline prevents it? (§21.4)

FURTHER READING

- **Arpaci-Dusseau, *Operating Systems: Three Easy Pieces (OSTEP)*, "Condition Variables," "Semaphores," and "Common Concurrency Problems."** Free (pages.cs.wisc.edu/~remzi/OSTEP/). The while-loop rule, the producer-consumer and reader-writer problems, semaphores, and the taxonomy of concurrency bugs — the reference for §21.1–§21.4.
- **Bryant & O'Hallaron, *Computer Systems: A Programmer's Perspective (CS:APP)*, Chapter 12 (Concurrent Programming), §12.5 (synchronizing threads with semaphores).** Semaphores, mutual exclusion, and producer-consumer from the programmer's view. [[▲](#) reference – not free]; CMU 15-213 materials free.
- **Dijkstra, "Cooperating Sequential Processes" (EWD123, 1965; published 1968).** The origin of the semaphore and its P/V operations. Free via the E. W. Dijkstra Archive (cs.utexas.edu/~EWD/).
- **Hoare, "Monitors: An Operating System Structuring Concept," *Communications of the ACM* 17(10):549-557 (1974), DOI 10.1145/355620.361161.** The primary source for monitors, with the bounded buffer and reader-writer examples. Author/venue/year verified via the ACM Digital Library.
- **The Linux `pthread_cond_wait(3)` and `sem_overview(7)` manual pages.** Primary sources for POSIX condition-variable semantics (including permitted spurious wakeups) and POSIX semaphores. Verify per system.

Exercises

- 21.1** **WARM-UP** In one sentence, why should a thread that may wait a while *sleep* rather than *spin*?
- 21.2** **WARM-UP** What three operations does a condition variable have, and what does `wait` do to the mutex when it is called?
- 21.3** **WARM-UP** What does a counting semaphore initialized to N guarantee about how many threads can be inside the section it guards?
- 21.4** **CORE** Explain why `cond_wait` must be called inside a `while` loop that re-checks the predicate, not an `if`. Give the two independent reasons (spurious wakeups and a condition consumed by another thread) and show, with a two-consumer example, how an `if` leads to acting on a false predicate.
- 21.5** **CORE** A condition variable does not remember a signal sent with no waiter, but a semaphore's count does. (a) Explain the lost-wakeup bug this creates for condition variables. (b) State the discipline (update shared state under the lock, then signal; wait in a loop) that prevents it. (c) Why does a semaphore not need this discipline for the "signal arrives before the waiter sleeps" case?

21.6 **CORE** Implement (in pseudocode) a bounded buffer with a mutex and two condition variables (`not_full`, `not_empty`). Show the producer and consumer, mark where each waits and signals, and explain why the buffer can never overflow or underflow and why the loops are necessary.

21.7 **COST** (a) A busy-poll waiter burned ~ 0.499 s of CPU for a ~ 0.5 s wait while a blocking waiter used ~ 0.001 s — explain the $\sim 500\times$ gap in terms of runnable-vs-blocked (Chapter 16). (b) When is a short spin actually *cheaper* than blocking, and what is the crossover quantity? (c) What does `cond_wait` cost that a spin does not (name the two scheduler events on the sleep/wake path), and why is it still the right default for a long wait?

21.8 **FADED** *Worked, with the last step for you.* A thread pool's workers wait for tasks. Each worker does `if (queue.empty()) cond_wait(&cv, &m);` then pops. Occasionally, under load, a worker pops from an empty queue and crashes; occasionally a submitted task is never picked up though workers are idle.

Step 1 (given): The `if` checks the predicate only once; on waking (spuriously, or after another worker took the task) the queue may be empty.

Step 2 (given): If the producer signals without holding the lock and updating the queue, a not-yet-sleeping worker can miss the signal.

Step 3 (given): A condition variable does not remember a signal with no waiter, so the missed signal is gone and the task waits with no one woken.

Step 4 (you): Name both bugs precisely, explain why each is intermittent and load-dependent, and give the corrected worker and producer (the `while` loop and the enqueue-under-lock-then-signal order), stating which bug each change fixes.

21.9 **MIXED** For each need, choose the best primitive (a plain mutex, a condition variable + mutex, or a counting semaphore) and justify in a line: (a) protect a single shared counter; (b) let at most 5 requests use a resource pool at once; (c) have a consumer block until a producer makes an item available; (d) signal "shutdown requested" to many waiting threads at once.

21.10 **MIXED** Drawing on Chapters 18–21, classify each as *true* or *false* and fix the false ones: (a) "A `cond_signal` with no thread waiting is remembered and delivered to the next waiter." (b) "`cond_wait` keeps holding the mutex while it sleeps." (c) "A binary semaphore behaves much like a mutex." (d) "You can safely replace the `while` around `cond_wait` with an `if` if your platform has no spurious wakeups." (e) "A blocking mutex, like a signal handler flag, is a way to coordinate two flows of control over shared memory."

21.11 **STRETCH** A monitor bundles shared state, its mutex, and its condition variables. (a) Explain what discipline a monitor enforces that raw locks-and-signals leave to the programmer, and how that eliminates the `if-not-while` and lost-wakeup traps at the call site. (b) A Java object's `synchronized` methods with `wait/notify`, and a thread-safe blocking queue, are both monitors — explain what each hides. (c) Argue for "prefer the highest-level correct abstraction" using the fact that the primitives are where the subtle bugs live, and give one case where you would still drop to raw condition variables.

21.12 **LAB** On your own machine: (a) write a waiter that blocks ~ 0.5 s two ways — busy-polling a flag under a lock, and blocking on a condition variable — and measure CPU time with `getrusage`; confirm the spinner burns $\sim 100\%$ of a core and the blocker ~ 0 . (b) Build a producer-consumer bounded buffer with a mutex and two condition variables (using `while` loops); confirm produced and consumed totals match and the buffer never exceeds its capacity. (c) Run ten threads against a semaphore initialized to three and confirm at most three are ever in the section at once. Label your compiler and kernel version, and note the CPU numbers are illustrative.

Chapter 22 — Atomics and Compare-and-Swap: Building Blocks Below the Lock

Before you start: Chapter 20 (a lock gives mutual exclusion — but what is a lock *built from*?), Chapter 21 (blocking primitives coordinate threads; all of them rest on an indivisible hardware operation), Chapter 15 (threads share memory, so an unsynchronized read-modify-write races), Chapter 3 (a core sees memory through its cache; "atomic" is enforced at the cache line). · The last two chapters used locks and condition variables as given; this one opens the box and shows the hardware primitive — the atomic read-modify-write — that every one of them is built on.

BEFORE YOU READ ON

From memory, reaching back: (1) Chapter 15 — why is `counter++` across two threads a race, and what are the three machine steps hiding inside that one line of C? (2) Chapter 20 — a mutex `lock()` has to test-and-set some word indivisibly; what would go wrong if the test and the set were two separate instructions? (3) *Predict*: four threads each add 1 to a shared counter one million times. With a plain `count++`, what final value do you expect (the ideal is 4,000,000)? With an atomic increment? Write both guesses; we measure them in §22.1.

A lock is not a primitive of the machine; it is a *program*, and like every program it needs some smaller indivisible operation to stand on. That operation is the **atomic read-modify-write** (RMW): a single hardware instruction that reads a memory location, computes a new value, and writes it back with no other core able to observe or interpose a write in the middle. This chapter builds up from the problem that motivates it. First, why an ordinary `x++` on shared memory silently loses updates — it is three steps, not one, and threads interleave inside it (§22.1). Then the atomic RMW that fixes it, and how the hardware actually provides indivisibility — the x86 `LOCK` prefix and the cache-line lock beneath it (§22.2). Then the one RMW that is more than a fixed arithmetic operation, the **compare-and-swap** (CAS), and the retry loop that turns it into any update you like (§22.3). Finally, why CAS in particular is the *universal* primitive — Herlihy's consensus hierarchy — and what atomics cost, since they are far from free (§22.4). We close by carrying CAS toward the lock-free data structures it makes possible (§22.5). Locks were the abstraction; atomics are the mechanism underneath.

22.1 The lost update: `x++` is three steps, not one

A single line, `count++`, compiles to a **read-modify-write**: load `count` into a register, add one, store it back. Those are three separate instructions, and a thread can be preempted — or simply run in parallel on another core — between any two of them. When two threads read the same old value before either has stored, both compute the same "new" value and both store it: two increments, one net effect. The update is *lost*. It is not a rare corner case; under contention it is the common case. Here are four threads each incrementing a shared `long` one million times, compiled at `-O0` so each iteration is a genuine load-add-store (gcc 11.4.0, Linux 6.18, x86-64):

```

expected      = 4000000
plain  x++    = 1355281  (lost 2644719 updates)
atomic fetch_add = 4000000

```

The plain counter lost roughly two-thirds of its four million increments — and the exact figure changes every run, because it is a function of how the threads happened to interleave, which is precisely the hallmark of a data race. The atomic version, using a single indivisible increment, landed on 4,000,000 exactly, every run. (At -O2 the plain loop can appear to "work": the optimizer collapses the million iterations into one `count += 1000000`, shrinking the race window to a single load-add-store per thread that rarely collides — a reminder that a data race is undefined behavior, so a passing test proves nothing.) The lesson is blunt: **a read-modify-write on shared memory must be performed as one indivisible operation**, or updates vanish. The tool that makes it one operation is the atomic.

QUICK CHECK

Name the three machine steps hiding inside `count++`, and describe the interleaving of two threads that loses one increment. Why does the plain counter give a *different* wrong answer each run? (Answer after the next section.)

22.2 Atomic read-modify-write and how the hardware makes it indivisible

C11 (and C++11) added `<stdatomic.h>`: types like `atomic_long` and operations like `atomic_fetch_add`, `atomic_fetch_or`, `atomic_exchange` that the compiler lowers to a single atomic machine instruction. Replacing `count++` with `atomic_fetch_add(&count, 1)` is the whole fix in §22.1. What does "single atomic instruction" mean at the hardware level? On x86, an ordinary increment of a memory operand and an atomic one differ by one prefix byte. Compare the compiler's output for a plain post-increment versus `atomic_fetch_add` (gcc 11.4.0, -O2, `objdump -d`):

```

plain_add:                                atomic_add:
  mov    (%rdi),%rax                      mov    $0x1,%eax
  lea    0x1(%rax),%rdx                   lock xadd %rax,(%rdi)
  mov    %rdx,(%rdi)                      ret
  ret

```

The plain version is exactly the three steps of §22.1: `mov` (load), `lea` (add one), `mov` (store) — three instructions, interruptible between any pair. The atomic version is *one* instruction, `lock xadd` ("exchange-and-add"), which loads, adds, and stores as a single indivisible unit. The magic is the `LOCK` prefix. It tells the CPU to hold exclusive ownership of the affected cache line for the duration of the read-modify-write, so no other core can read or write that line in between — historically by asserting a bus lock, and on modern chips by the cache-coherence protocol (Chapter 3) refusing to yield the line while the locked operation completes. The result is that the whole load-add-store is atomic with respect to every other core. Note the cost this implies, which §22.4 measures: the core

must own the line exclusively, so a `lock`-prefixed op serializes at the cache line and cannot be reordered around freely — it is much dearer than a plain store.

THE SAME IDEA THREE WAYS: AN ATOMIC RMW IS ONE INDIVISIBLE READ-MODIFY-WRITE

1. In words. A plain `x++` is load, add, store — three instructions another core can interleave into, losing updates. An atomic RMW (`lock xadd`) does the same read-modify-write as a single instruction that owns the cache line throughout, so no other core can observe or interpose a write in the middle.

2. As a picture. Figure 22.1: two threads each load the same old value 41, each add one, each store 42 — one increment lost — versus the atomic path where the second thread cannot touch the line until the first's read-modify-write retires, so it reads 42 and stores 43.

3. As numbers. Four threads $\times$ one million: plain `++` lost 2,644,719 of 4,000,000 (and a different count each run); atomic `fetch_add` lost none — exactly 4,000,000, every run. In the disassembly, three instructions collapse to one `lock xadd`.

(Answer to the §22.1 quick check: the three steps are load `count` into a register, add one, store it back. If thread A loads 41 and, before it stores, thread B also loads 41, both compute 42 and both store 42 — two increments, one net change, so one is lost. The wrong answer differs each run because it depends on the exact interleaving of loads and stores across cores, which the scheduler and hardware vary run to run — the signature of a data race.)

Plain `x++`: two threads interleave inside the read-modify-write and lose an increment.




Figure 22.1: The lost update. With a plain `x++`, T2 reads the counter before T1 has stored, so both write 42 and one increment is lost. With an atomic `lock xadd`, T2's read-modify-write cannot begin until T1's has completed, because the `LOCK` prefix holds the cache line exclusively for the whole operation.

22.3 Compare-and-swap: the conditional RMW and the retry loop

`fetch_add` and its siblings each hard-wire one operation (add, or, exchange). But most interesting updates are not a fixed arithmetic function of the old value — "install this new head only if the head is still what I read," "raise the maximum only if my value is larger." The primitive that expresses *any* such conditional update is **compare-and-swap**. In C11 it is `atomic_compare_exchange_weak(obj, &expected, desired)`: it atomically compares the value in `obj` with `expected`, and *only if they are equal* writes `desired` and returns true; otherwise

it writes the current value back into `expected` and returns false. In one indivisible step it asks "is it still what I think?" and, if so, updates it. That conditional is what lets you build a lock-free update out of a plain read, an arbitrary computation, and a CAS in a **retry loop**: read the current value, compute the new one, attempt to CAS it in; if the CAS fails (someone else changed it first), the failure has already reloaded the current value, so loop and try again. Here is an atomic *maximum* — an operation with no dedicated instruction — built from a CAS loop, run by eight threads (gcc 11.4.0, Linux 6.18):

```
void atomic_max(atomic_long *p, long v) {
    long cur = atomic_load(p);
    while (v > cur)
        if (atomic_compare_exchange_weak(p, &cur, v)) /* only if we'd raise it */
            return; /* cur is reloaded on failure */
    /* installed */
}
/* 8 threads hammering atomic_max: */
max seen = 7499 (true max = 7499)
```

Eight threads racing to raise a shared maximum agreed exactly on the true maximum, 7499, with no lock. The structure is the template for essentially every lock-free algorithm in the next chapter: **read a snapshot, compute off it, and commit with a CAS that succeeds only if nothing changed underneath you** — retrying if it did. Two details matter. The `expected` argument is updated by a failed CAS to the value it actually found, so the loop always retries against fresh state without a separate reload. And the `_weak` form is allowed to fail *spuriously* (report mismatch even when the values were equal) in exchange for cheaper code on architectures that implement CAS with load-linked/store-conditional; because it already lives in a retry loop, a spurious failure just means one more iteration, so `_weak` is the right choice inside a loop and `_strong` (no spurious failures) for a one-shot check.

QUICK CHECK

State exactly what `compare_exchange` does on success and on failure, including what happens to the `expected` argument when it fails. Why does the CAS-loop idiom not need a separate reload of the current value on retry? When would you pick `_weak` over `_strong`? (Answer below the communicate box.)

22.4 Why CAS is universal — and what atomics cost

Why single out compare-and-swap rather than, say, atomic add? Because CAS is *universal* in a precise, proven sense. Herlihy's "Wait-Free Synchronization" (1991) ranks synchronization primitives by their **consensus number**: the maximum number of threads that can use the primitive (plus plain reads and writes) to solve *consensus* — all agreeing irrevocably on one value despite arbitrary delays. Atomic reads and writes alone have consensus number **1**: they cannot solve consensus for even two threads. Atomic `fetch_add`, `test-and-set`, and their kin sit higher but still at a finite number. Compare-and-swap has consensus number **infinity**: with CAS, any number of threads can reach consensus, and from it you can build a wait-free implementation of *any* object with a sequential specification. That is why hardware designers provide CAS (or load-linked/

store-conditional, its equivalent) and why every lock-free structure in Chapter 23 is a CAS loop: CAS is the primitive from which all the others, and all lock-free objects, can be constructed. Plain loads and stores cannot substitute — no arrangement of them solves two-thread consensus.

Universality does not mean free. An atomic RMW must own its cache line exclusively and cannot be reordered around like a plain access, so it costs far more than an ordinary operation — and under contention, more still, because the cache line ping-pongs between cores (Chapter 3's coherence traffic). Measured uncontended, one operation each, 100 million times (gcc 11.4.0, -O2):

```
plain  ++      : 0.9 ns/op
atomic fetch_add : 15.3 ns/op      (~15× a plain op)
mutex lock/unlock : 11.9 ns/op      (itself two atomic ops)
```

A plain increment is under a nanosecond; the atomic is an order of magnitude dearer, and a full mutex lock/unlock — which is itself built from atomics — is in the same league (a lock is not magically slower than an atomic; it *is* atomics plus a slow path for when it must block). These are illustrative single-machine numbers, but the shape is universal: atomics are cheap enough to use freely for a shared counter and expensive enough that a hot, highly contended atomic can dominate a profile. The other cost is *correctness*: atomic loads and stores are individually indivisible, but they do **not compose** into an atomic compound. Two separate atomic operations have a gap between them, and the whole point of §22.1 was that gaps lose updates — which is exactly the trap that sends you to CAS.

TRAP: COMPOSING ATOMIC OPERATIONS (AN ATOMIC LOAD THEN AN ATOMIC STORE IS NOT ATOMIC)

Because `atomic_load` and `atomic_store` are each atomic, it is tempting to write a "conditional update" as two of them: `if (atomic_load(&p) == expected) atomic_store(&p, desired);`. Each call is indivisible, but *together* they are not: between the load and the store, another thread can change `p`, and your store clobbers its update — the lost-update bug of §22.1 wearing atomic clothing. This is a check-then-act race, and it is the reason compare-and-swap exists: CAS folds the compare **and** the swap into one indivisible operation, so nothing can slip between them. Any time you find yourself reading an atomic, deciding based on the value, and writing back a value that depends on that decision, you need a single CAS (usually in a retry loop), not a load followed by a store. The tell is a sentence of the form "if it's still X, make it Y" spanning two atomic calls — that gap is the bug.

SAY IT LIKE A SYSTEMS ENGINEER

Try it first. A teammate says: "I don't need a mutex for this flag. I'll just do `if (atomic_load(&busy) == 0) atomic_store(&busy, 1);` to claim it, and `atomic_store(&busy, 0)` to release. Both are atomic, so it's a lock-free lock — right?" Cover the paragraph and respond in four or five sentences, drawing on §22.3 and §22.4.

Model answer. "The two operations are each atomic, but the sequence isn't, and that gap is the whole problem: two threads can both load `busy == 0`, both decide it's free, and both store 1 — you've handed the 'lock' to two threads at once, which is the lost-update race from the counter example. To claim the flag indivisibly you need a single *compare-and-swap*: `expected = 0; atomic_compare_exchange_strong(&busy, &expected, 1)` succeeds for exactly one thread and fails for the other, because the compare and the set happen as one operation with no window between them. That is, in fact, roughly how a spinlock's fast path is built. But before you build your own, use a real mutex or a tested lock-free type: hand-rolled synchronization is where the subtle, load-dependent bugs live, and the atomic `compare_exchange` is the sharp tool you reach for only when you must." Naming the check-then-act gap and reaching for CAS — not two atomic accesses — is the senior register.

(Answer to the §22.3 quick check: on success, `compare_exchange` finds `*obj == expected`, writes `desired`, and returns `true`; on failure, it leaves `*obj` unchanged, writes the current value of `*obj` into `expected`, and returns `false`. The loop needs no separate reload because a failed CAS has already refreshed `expected` with what it found. Prefer `_weak` inside a retry loop, where a rare spurious failure just costs one more iteration and the code is cheaper on load-linked/store-conditional machines; prefer `_strong` for a single, non-looping check.)

22.5 What to carry forward

Beneath every lock and condition variable is an **atomic read-modify-write**: a single hardware instruction that reads, computes, and writes with no other core able to interpose. An ordinary `x++` is not atomic — it is load, add, store, and threads interleave inside it, losing updates (our four-threaded counter lost 2.6 million of 4 million; the atomic lost none). The hardware provides indivisibility through a locked operation — on x86 the `LOCK` prefix, which holds the cache line exclusively for the whole read-modify-write (`lock xadd` is one instruction where the plain version is three). The one RMW that expresses *arbitrary* conditional updates is **compare-and-swap**: "write this only if the value is still what I read," folded into one indivisible step, and wrapped in a **retry loop** — read, compute, CAS, retry on failure — it implements any update lock-free (our eight-threaded atomic maximum agreed on 7499). CAS is not merely convenient but *universal*: Herlihy's consensus hierarchy places plain reads/writes at consensus number 1 and CAS at infinity, so any lock-free object can be built from it. Atomics are far from free (roughly fifteen times a plain op, more under contention), and — the trap — they do not compose: an atomic load then an atomic store leaves a gap another thread slips through, which is precisely why CAS exists.

With compare-and-swap in hand, a new possibility opens: synchronization *without locks at all*. If a single CAS can commit an update only when nothing changed underneath it, then a data structure can let many threads race, each preparing an update and

committing it with a CAS, retrying the few that lose — no thread ever blocked, no lock ever held. The next chapter builds exactly these **lock-free data structures**: the Treiber stack and the Michael-Scott queue, both nothing but a CAS loop over a shared pointer. It also confronts what CAS alone does not solve — the notorious *ABA problem*, where a value returns to what you read while its meaning has changed underneath, defeating the compare — and how real systems reclaim memory safely without locks. And underneath all of it sits an assumption we have quietly leaned on and must finally make explicit: that these atomic operations become visible to other cores in a sensible order. The chapter after next, on the memory model, shows that plain loads and stores carry no such guarantee, and that the ordering a lock or a lock-free structure needs is something the atomic must be asked for by name.

CAN YOU... WITHOUT LOOKING?

- Explain why `x++` on shared memory loses updates, naming the three steps and the racing interleaving.
- Say what the x86 `LOCK` prefix does, and why `lock xadd` is atomic where three plain instructions are not.
- State precisely what `compare_exchange` does on success and failure, and write the CAS retry-loop idiom.
- Explain what "CAS is universal" means via consensus numbers (plain reads/writes = 1, CAS = infinity).
- Give the rough cost of an atomic RMW versus a plain op, and why contention makes it worse.
- State why two atomic operations do not compose, and what to use instead for a conditional update.

RETRIEVAL — CLOSED BOOK

Cover the chapter. Answer from memory; rate confidence first.

1. **R1.** Why is `count++` across threads a race, and what does an atomic RMW change? (§22.1)
2. **R2.** What does the `LOCK` prefix do at the cache-line level, and how many instructions is `lock xadd` versus a plain increment? (§22.2)
3. **R3.** What does `compare_exchange` return and write on success and on failure, and what is the CAS-loop idiom? (§22.3)
4. **R4.** What is a consensus number, and what are the consensus numbers of plain reads/writes and of CAS? (§22.4)
5. **R5.** Why do two atomic operations fail to compose into an atomic compound, and what fixes it? (§22.4)

FURTHER READING

- **Herlihy, "Wait-Free Synchronization," *ACM Transactions on Programming Languages and Systems* 13(1):124-149 (1991), DOI 10.1145/114005.102808.** The consensus hierarchy: plain reads/writes have consensus number 1, compare-and-swap has infinity, and CAS is universal for wait-free constructions. The source for §22.4; author, venue, and year verified via dblp and the ACM Digital Library.
- **Herlihy & Shavit, *The Art of Multiprocessor Programming* (Morgan Kaufmann; 2nd ed. 2020).** The standard text on atomics, the consensus hierarchy, and lock-free programming, with the primitives of this chapter developed rigorously. [[▲](#) reference – not free].
- **cppreference, `<stdatomic.h>` / `std::atomic` and `atomic_compare_exchange` (en.cppreference.com).** The precise contract for `fetch_add`, the `_weak` vs `_strong` compare-exchange, and spurious failure. Primary reference for §22.2–§22.3; verify per standard version.
- **Intel 64 and IA-32 Architectures Software Developer's Manual, Vol. 3A, "Locked Atomic Operations" (the `LOCK` prefix and `CMPXCHG/XADD`).** Primary source for how x86 makes a read-modify-write atomic at the cache line. Free from Intel.
- **Drepper, "What Every Programmer Should Know About Memory" (Red Hat, 2007).** Free (akkadia.org/drepper/cpumemory.pdf). Cache coherence and why a lock-prefixed operation costs what it does — background for the numbers in §22.4.

Exercises

- 22.1** **WARM-UP** Name the three machine steps that `count++` compiles to, and in one sentence say why that makes it non-atomic across threads.
- 22.2** **WARM-UP** What does the x86 `LOCK` prefix guarantee about the cache line during a read-modify-write, and how many instructions is `lock xadd`?
- 22.3** **WARM-UP** In one sentence each: what does `atomic_compare_exchange` return on success, and what does it write into its `expected` argument on failure?
- 22.4** **CORE** Two threads run `count++` on a shared counter starting at 0, each once. (a) Enumerate an interleaving of the six machine steps (three per thread) that ends with `count == 1` instead of 2. (b) Explain why replacing the body with `atomic_fetch_add(&count, 1)` makes that interleaving impossible. (c) Why can the same program appear correct at 02 yet be lost-update-prone at 00?
- 22.5** **CORE** Write the compare-and-swap retry loop for an atomic operation with no dedicated instruction — say, atomically setting a counter to the larger of its current value and `v` (an atomic max). Explain what the `expected` argument holds after a failed CAS and why the loop therefore needs no explicit reload, and state one reason to use `_weak` rather than `_strong` inside the loop.

22.6 **CORE** A colleague implements "claim a resource" as `if (atomic_load(&owner) == 0) atomic_store(&owner, my_id);`. (a) Show the interleaving that lets two threads both believe they own the resource. (b) Rewrite it with a single compare-and-swap so exactly one thread succeeds. (c) Explain the general rule: which shape of update always requires CAS rather than a separate atomic load and store?

22.7 **COST** Using the measured figures (plain `++` ~ 0.9 ns, atomic `fetch_add` ~ 15 ns, uncontended): (a) explain where the $\sim 15\times$ comes from in terms of the cache line and instruction reordering. (b) Predict qualitatively what happens to the atomic's cost when eight cores hammer the *same* counter, and why (invoke Chapter 3's coherence traffic). (c) Given that, when is a per-thread counter summed at the end preferable to one shared atomic, and what does that cost instead?

22.8 **FADED** *Worked, with the last step for you.* You must maintain a running maximum latency observed across many worker threads, lock-free.

Step 1 (given): A plain `if (v > hi) hi = v;` is a check-then-act race — two threads can both pass the test and the smaller write can land last.

Step 2 (given): There is no `atomic_fetch_max` in C11, so you build it from `compare_exchange`.

Step 3 (given): The loop reads `cur = atomic_load(&hi)`, and while `v > cur` tries `compare_exchange_weak(&hi, &cur, v)`; a failed CAS reloads `cur`.

Step 4 (you): Write the complete function, explain why it terminates (each failed retry means someone else raised `hi`, so `cur` only grows), and state what invariant it guarantees about `hi` versus every `v` ever submitted.

22.9 **MIXED** For each need, choose the cheapest correct tool — a plain access, an atomic `fetch_*`, a single CAS, or a CAS retry loop — and justify in a line: (a) increment a shared hit counter; (b) set a flag to 1 exactly once and detect who did it first; (c) push onto a shared linked-list stack; (d) read a value no other thread ever writes.

22.10 **MIXED** Drawing on Chapters 20–22, classify each as *true* or *false* and fix the false ones: (a) "A mutex is a fundamental hardware primitive, unlike an atomic add." (b) "`lock_xadd` and a plain `add` to memory differ by a one-byte prefix." (c) "Because each is atomic, an `atomic_load` followed by an `atomic_store` forms an atomic compound operation." (d) "Compare-and-swap can solve consensus for any number of threads, but plain reads and writes cannot solve it for two." (e) "An uncontended atomic increment costs about the same as a plain increment."

22.11 **STRETCH** Herlihy's consensus hierarchy ranks primitives by consensus number. (a) Explain what it means that atomic read/write has consensus number 1, and give the intuition for why two threads cannot reach consensus with reads and writes alone. (b) Explain what "CAS has consensus number infinity, and is universal" buys you — what class of objects can you build from CAS? (c) Connect this to the previous two chapters: in what sense is every lock and every lock-free structure "just" CAS (or an equivalent primitive), and why does the hardware bother to provide CAS specifically?

22.12 **LAB** On your own machine: (a) write four threads each doing one million `count++` on a shared `long`, compile at `-O0`, and confirm the total falls short of 4,000,000 and varies run to run; switch to `atomic_fetch_add` and confirm it is exactly 4,000,000 every run; then recompile the plain version at `-O2` and observe it may now "pass" (explain why). (b) Disassemble a plain memory increment and an `atomic_fetch_add` with `objdump -d` and find the `lock` prefix. (c) Build an atomic max from a `compare_exchange_weak` loop, run it under many threads, and confirm it agrees with the true maximum. Label your compiler and CPU; note the timing figures are illustrative.

Chapter 23 — Lock-Free Data Structures

Before you start: Chapter 22 (compare-and-swap and the retry loop — the one primitive this whole chapter is built from), Chapter 20 (a lock serializes access; what if a delayed lock-holder must not stall everyone else?), Chapter 7 (pointers and nodes — a stack and a queue are pointer surgery), Chapter 9 (free and use-after-free — reclaiming a node another thread still holds is a disaster). Chapter 22 gave you CAS; this chapter uses nothing but CAS in a loop to build shared stacks and queues that no thread ever locks.

BEFORE YOU READ ON

From memory, reaching back: (1) Chapter 22 — write the compare-and-swap retry-loop idiom in three steps (read, compute, commit-or-retry), and say what a failed CAS does to your local copy of the value. (2) Chapter 20 — name one thing that goes wrong when the thread holding a lock is preempted, delayed, or crashes. (3) *Predict*: eight threads push and pop on a *shared* stack with no lock — just a CAS on the head pointer. Do you expect any pushed item to be lost or duplicated? And if a node can be freed and its address immediately reused, what could break? Write your guesses; §23.2 and §23.4 settle them.

A lock makes concurrency correct by making it *sequential*: one thread at a time inside the critical section. That is simple and often right, but it has a structural weakness — the fate of every thread is tied to the one holding the lock. If that thread is preempted (Chapter 16), page-faults (Chapter 17), or is simply slow, everyone waiting is stuck, and if it dies holding the lock, they are stuck forever. **Lock-free** programming removes that coupling: many threads operate on the shared structure at once, each committing its change with a single compare-and-swap that succeeds only if nothing moved underneath it, and retrying if it did. No thread ever holds a lock, so no thread's delay can stall the others. This chapter builds the idea concretely. First, what "lock-free" actually guarantees — the ladder of progress conditions from blocking up to wait-free, and why "lock-free" is a promise about the *system*, not the absence of locks (§23.1). Then the two canonical structures, each a CAS loop over a pointer: the **Treiber stack** (§23.2) and the **Michael-Scott queue** (§23.3). Then the hazard that makes lock-free memory management genuinely hard — the **ABA problem** and safe reclamation (§23.4). We close by carrying the open question of *ordering* — which lock-free code depends on and which we have not yet justified — into the memory model (§23.5).

23.1 What "lock-free" means: the ladder of progress

"Lock-free" is a precise progress guarantee, not a synonym for "uses no mutex." The guarantees form a ladder. A **blocking** (lock-based) algorithm offers no non-blocking guarantee at all: a thread delayed inside the critical section blocks every other thread indefinitely. **Obstruction-free**, the weakest non-blocking level, promises that a thread will finish if it runs *in isolation* for long enough (no guarantee under contention). **Lock-free** promises that the *system as a whole* always makes progress: out of all contending threads, *some* thread completes its operation in a bounded number of steps, so the

structure can never be globally stuck — though an unlucky individual thread might retry forever (starve) while others succeed. **Wait-free**, the strongest, promises that *every* thread completes in a bounded number of *its own* steps, with no starvation. These are Herlihy's conditions, and they trade strength for cost: wait-free is ideal but often intricate and slower, so most practical "lock-free" structures — including the two in this chapter — give the lock-free guarantee, which is enough to eliminate the lock-holder problem.

Why does the guarantee matter in practice? Because a lock-free structure is immune to the pathologies that come from a thread being stopped at the wrong moment. There is no lock to hold, so there is no **deadlock** and no **priority inversion** (a low-priority thread cannot block a high-priority one by holding a lock, Chapter 16); a thread preempted mid-operation delays only itself, never the structure; and a thread that dies leaves the structure usable by everyone else. The price is subtlety: every operation must be designed so that its single committing CAS is a **linearization point** — the instant the operation appears to take effect atomically, in Herlihy and Wing's sense (1990), so that concurrent operations still look like *some* valid sequential order. Get the linearization point right and the structure is correct under any interleaving; get it wrong and you have a race that no test reliably catches.

QUICK CHECK

Put the four progress conditions in order from weakest to strongest, and state the difference between the lock-free and wait-free guarantees (whose progress is promised?). Name one pathology that lock-free code cannot suffer, and why. (Answer after the next section.)

23.2 The Treiber stack: a CAS loop on the head pointer

The simplest lock-free structure is a stack, published by Treiber (IBM, 1986). It is a singly linked list with an atomic `top` pointer, and both operations are the CAS retry loop of Chapter 22 applied to that pointer. To **push**: set the new node's `next` to the current `top`, then CAS `top` from that old value to the new node; if the CAS fails (someone else pushed or popped first), the failure reloaded `top`, so re-point `next` and retry. To **pop**: read `top`, and CAS it to `top->next`; retry if it moved. That is the whole algorithm — no lock, no blocking. Eight threads pushing 100,000 items each and then popping every node concurrently agree exactly (gcc 11.4.0, Linux 6.18, x86-64):

```

void push(long v) {
    node *n = malloc(sizeof *n); n->val = v;
    node *old = atomic_load(&top);
    do { n->next = old; }
    while (!atomic_compare_exchange_weak(&top, &old, n)); /* old reloaded on failure */
}
node *pop(void) {
    node *old = atomic_load(&top);
    while (old) { if (atomic_compare_exchange_weak(&top, &old, old->next)) return old; }
    return NULL;
}
/* 8 threads push 100000 each, then 8 threads pop concurrently: */
pushed total = 800000
popped count = 800000, popped sum = 800000    (no node lost or duplicated)

```

Every one of the 800,000 nodes was popped exactly once — none lost, none returned twice — under eight-way contention with no lock. The linearization point is the successful CAS: a push "happens" the instant its CAS installs the new top, a pop the instant its CAS unlinks the old top, and any two concurrent operations therefore order by which CAS won. Under contention the losers simply retry, which is why the guarantee is *lock-free* (the system always advances — some CAS always wins) and not *wait-free* (a persistently unlucky thread could keep losing). One honest caveat, deferred to §23.4: this code *leaks* — it never frees popped nodes. Freeing them safely while other threads may still be reading them is the hard part of lock-free programming, and it is where the ABA problem lives.

THE SAME IDEA THREE WAYS: PUSH IS A CAS LOOP THAT INSTALLS A NEW HEAD

1. In words. Point the new node at the current top, then compare-and-swap the top from that old value to the new node. If another thread changed the top first, the CAS fails and hands you the new top; re-point and try again.

2. As a picture. Figure 23.1: thread A reads `top = X`, sets `new->next = X`; thread B pushes first, so `top` is now `Y`; A's `CAS(top, X, new)` fails (`top` is `Y`, not `X`), reloads `old = Y`, re-points `new->next = Y`, and its second CAS succeeds.

3. As numbers. Eight threads $\times$ 100,000 pushes then concurrent pops: 800,000 pushed, 800,000 popped, sum 800,000 — no node lost or duplicated. Under an 8-thread push/pop benchmark the lock-free stack ran at ~176 ns/op against a mutex-guarded stack's ~359 ns/op (illustrative, contention-dependent).

(Answer to the §23.1 quick check: *weakest to strongest — blocking, obstruction-free, lock-free, wait-free. Lock-free guarantees that the **system** makes progress (some thread always completes), allowing an individual thread to starve; wait-free guarantees that **every** thread completes in a bounded number of its own steps. Lock-free code cannot deadlock or suffer priority inversion, because there is no lock to hold — a delayed thread stalls only itself, never the structure.*)

Treiber push: CAS the head; retry if another thread moved it first.




Figure 23.1: The Treiber stack's push. Thread A prepares a node pointing at the top it read (X), but thread B pushes Y first, so A's compare-and-swap fails — and the failure reloads the current top (Y) into A's local variable. A re-points its node at Y and retries; the second CAS succeeds. The successful CAS is the push's linearization point.

23.3 The Michael-Scott queue: two pointers, and helping

A FIFO queue is harder than a stack because two ends change independently: enqueue appends at the `Tail`, dequeue removes at the `Head`. The classic lock-free solution is the **Michael-Scott queue** (PODC 1996), the algorithm behind Java's `ConcurrentLinkedQueue` and countless runtimes. It keeps a permanent **dummy node** so the list is never empty (`Head` and `Tail` start pointing at it), which removes the awkward special cases at the boundaries. Enqueue is a two-step CAS: first CAS the current tail node's `next` from `NULL` to the new node (this is the linearization point — the item is now in the queue); then CAS `Tail` forward to the new node. Because those two CASes are not one atomic step, another thread can arrive between them and find `Tail` "lagging" — pointing at a node whose `next` is no longer `NULL`. The elegant fix is **helping**: any thread that observes a lagging tail advances it (CAS `Tail` to `tail->next`) before proceeding, so no operation ever waits for the thread that fell behind. Four producers and four consumers hammering the queue agree exactly (gcc 11.4.0, Linux 6.18):

```
enqueue: link new node after tail (CAS tail->next: NULL -> n) // linearization point
        then swing Tail forward  (CAS Tail: tail -> n)      // may be done by a helper
dequeue: read value from Head->next, then CAS Head forward to it

4 producers x 200000, 4 consumers:
expected count=800000 sum=1279999600000
dequeued  count=800000 sum=1279999600000 OK
```

All 800,000 items came out exactly once and the checksums matched, with no lock anywhere. The two lessons generalize beyond the queue. First, when a logical update spans more than one atomic step, you cannot hold a lock across them — so you design the structure to be *correct in every intermediate state* and let any thread **help** complete a half-finished operation it stumbles upon; this "help the laggard" pattern recurs throughout lock-free design. Second, the linearization point is a single, identifiable CAS (here, linking the new node), which is what lets you reason that the concurrent execution is equivalent to some sequential one. As with the stack, this code as written leaks nodes; making dequeue free the node it removed — safely, when no other thread is still looking at it — is the reclamation problem of the next section.

QUICK CHECK

Why does the Michael-Scott queue keep a permanent dummy node? Enqueue takes two CAS steps — what can another thread observe between them, and what does "helping" do about it? Which single step is the enqueue's linearization point? (Answer below the communicate box.)

23.4 The ABA problem and safe memory reclamation

Everything so far leaked memory on purpose, because reclaiming it exposes the subtlest bug in lock-free programming. A CAS checks whether a pointer *still equals* the value you read — but "equal" means the same bits, not the same object. If between your read and your CAS a value changes from **A** to **B** and back to **A**, your CAS sees **A**, concludes "nothing changed," and succeeds — even though the world moved on underneath it. This is the **ABA problem**. It bites the Treiber stack precisely when nodes are freed and reused: a thread reads `top = A` and plans to CAS `top` to `A->next`; meanwhile other threads pop **A** and **B**, free **A**, and a fresh `malloc` hands back **A**'s address for a new node that is pushed, so `top` is **A** again — but `A->next` now points somewhere entirely different (or into freed memory). The stalled thread's CAS succeeds on the recycled address and installs a stale, dangling `next`, corrupting the stack. ABA is intermittent by nature — it needs a specific interleaving plus address reuse — which is exactly what makes it dangerous: it survives testing and strikes in production.

There are three standard defenses, and every real lock-free system uses one. A **tagged (versioned) pointer** pairs the pointer with a counter that is incremented on every update, and CASes the pair; **A-to-B-to-A** now differs in the counter, so the CAS correctly fails. This needs a double-width CAS — on x86-64 the `lock cmpxchg16b` instruction (compile with `-mcx16`) — and the versioned Treiber stack passes the same 800,000-node test the plain one did. **Hazard pointers** (Michael, 2004) take a different tack: before dereferencing a node, a thread publishes it in a per-thread "hazard" slot; a thread that wants to free a node first checks that no hazard pointer references it, deferring reclamation until it is unreferenced — so a node in use is never freed, and ABA-via-reuse cannot happen. **Epoch-based reclamation** and **RCU** (read-copy-update) generalize this: readers announce they are active, and a retired node is freed only after every reader that could have seen it has departed. All three exist for one reason — a lock-free structure cannot `free` a node the moment it is unlinked, because another thread may still be about to read it, and freeing it is a use-after-free (Chapter 9) waiting to happen.

TRAP: THE ABA PROBLEM (A CAS SEES THE SAME VALUE, NOT THE SAME WORLD)

A compare-and-swap only checks that a location still holds the *bits* you read, so a value that goes $A \rightarrow B \rightarrow A$ slips through: your CAS succeeds as if nothing happened, though the object those bits refer to may have been popped, freed, and replaced. In the Treiber stack this corrupts the list — the CAS installs a next pointer read from a node that has since been recycled, silently losing or duplicating items, or dangling into freed memory. It is intermittent (it needs a preemption at the wrong instant plus address reuse), so it passes tests and fails under production load. The tell is any lock-free code that frees nodes and reuses their memory while a CAS elsewhere depends on a pointer's identity. The fixes are the three of §23.4: a tagged pointer (version the value so A-to-A differs), hazard pointers, or epoch/RCU reclamation (never free a node another thread might still hold). If you are writing a lock-free structure that frees nodes and you have not chosen one of these, you have an ABA bug you simply have not hit yet.

SAY IT LIKE A SYSTEMS ENGINEER

Try it first. A teammate says: "I wrote a lock-free stack — push and pop are just a CAS on the head, straight out of the textbook. In the unit test it's perfect. But under load it occasionally segfaults or loses an item, and I can't reproduce it in the debugger. The CAS is obviously correct — it only swaps the head if it hasn't changed. What's going on?" Cover the paragraph and respond in four or five sentences, drawing on §23.2 and §23.4.

Model answer. "The CAS logic is right; the bug is memory reclamation, and the symptom is the ABA problem. If pop frees the node it removed, that address can be handed back by malloc and pushed again while another thread is mid-pop holding the old head value — its CAS sees the same pointer bits, concludes nothing changed, and succeeds on a node whose next now points into recycled or freed memory, which is your segfault-or-lost-item. It's intermittent because it needs a preemption at exactly the wrong moment plus address reuse, so it never shows up in a single-stepping debugger. The fix isn't the CAS — it's safe reclamation: a tagged/versioned pointer with a double-width CAS so A-to-A differs by version, or hazard pointers, or epoch-based reclamation so you never free a node another thread might still be reading. Honestly, unless we need this specific structure, I'd use a vetted concurrent queue from the standard library — lock-free reclamation is exactly where hand-rolled code goes wrong." Naming ABA, tying it to reclamation rather than the CAS, and reaching for a vetted implementation is the senior register.

(Answer to the §23.3 quick check: the dummy node keeps the list non-empty so Head and Tail always point at a real node, eliminating null-boundary special cases. Between enqueue's two CASes another thread can see Tail lagging — pointing at a node whose next is no longer NULL; "helping" means that thread advances Tail itself (CAS it to tail->next) rather than waiting. The linearization point is the first CAS, linking the new node after the tail — that is the instant the item is in the queue.)

23.5 What to carry forward

Lock-free programming trades the simplicity of a lock for immunity to the lock-holder's fate: many threads operate at once, each committing with a single compare-and-swap that wins only if nothing moved underneath, retrying if it lost, so no thread ever blocks another. "Lock-free" is a precise guarantee — the *system* always makes progress (some thread completes), which is weaker than **wait-free** (every thread completes) but strong

enough to rule out deadlock and priority inversion. The two canonical structures are each a CAS loop over a pointer: the **Treiber stack** (CAS the head; our eight-threaded run pushed and popped 800,000 nodes with none lost) and the **Michael-Scott queue** (a dummy node, a two-step enqueue, and *helping* to advance a lagging tail; four producers and four consumers moved 800,000 items exactly). Each operation has a single CAS as its **linearization point**, which is what makes the concurrent execution equivalent to a sequential one. The genuinely hard part is memory: a CAS checks bits, not identity, so a value that cycles A-B-A defeats it — the **ABA problem** — and you cannot free an unlinked node while another thread may still read it. The three defenses are tagged pointers (double-width CAS), hazard pointers, and epoch/RCU reclamation.

Every algorithm in this chapter quietly assumed something we have not earned: that when one thread's CAS publishes a node, another thread that reads the updated pointer also sees the node's *contents* — its `val` and `next` — fully written. On the abstract machine that is obvious; on real multicore hardware it is not. Cores have store buffers and caches, and both the compiler and the processor reorder memory operations, so a thread can observe a new pointer while the writes that initialized what it points to are still in flight — a reader following a freshly published pointer into *garbage*. What saves the code is that the atomic operations carry **ordering** guarantees, and those guarantees are not automatic: they must be requested, by name, through the atomic's memory-ordering argument. The final chapter of this part makes that machinery explicit — the **memory model**: why a data race is undefined behavior, why plain loads and stores may be reordered across cores, and what `memory_order_relaxed`, `acquire/release`, and `seq_cst` actually promise. It is the specification that says precisely when the node you just published is safe for another thread to read.

CAN YOU... WITHOUT LOOKING?

- Order the four progress conditions and state what lock-free versus wait-free each guarantee.
- Write the Treiber stack's push and pop as CAS loops, and name each operation's linearization point.
- Explain the Michael-Scott queue's dummy node, two-step enqueue, and the "helping" that advances a lagging tail.
- Explain the ABA problem: why a CAS on the same bits can be wrong, and where it bites the Treiber stack.
- Name the three safe-reclamation techniques (tagged pointers, hazard pointers, epoch/RCU) and what each prevents.
- Say why lock-free code cannot deadlock or suffer priority inversion.

RETRIEVAL — CLOSED BOOK

Cover the chapter. Answer from memory; rate confidence first.

1. **R1.** What does the lock-free guarantee promise, and how does it differ from wait-free? (§23.1)
2. **R2.** Write the Treiber stack's push and pop, and name the linearization point of each. (§23.2)
3. **R3.** Why does the Michael-Scott queue use a dummy node and "helping"? (§23.3)
4. **R4.** What is the ABA problem, and how does it corrupt the Treiber stack when nodes are freed? (§23.4)
5. **R5.** Name the three safe-reclamation techniques and what each does. (§23.4)

FURTHER READING

- **Michael & Scott, "Simple, Fast, and Practical Non-Blocking and Blocking Concurrent Queue Algorithms," *Proc. 15th ACM PODC* (1996), pp. 267-275, DOI 10.1145/248052.248106.** The lock-free queue of §23.3, with the dummy node and the tail-helping mechanism. Author, venue, and year verified via dblp and the ACM Digital Library.
- **Treiber, "Systems Programming: Coping with Parallelism," *IBM Almaden Research Report RJ 5118* (April 1986).** The compare-and-swap stack of §23.2 (the "Treiber stack"). Report number and date verified via the IBM Research report archive.
- **Herlihy & Wing, "Linearizability: A Correctness Condition for Concurrent Objects," *ACM TOPLAS* 12(3):463-492 (1990), DOI 10.1145/78969.78972.** The definition of the linearization point used throughout this chapter. Verified via dblp and the ACM Digital Library.
- **Michael, "Hazard Pointers: Safe Memory Reclamation for Lock-Free Objects," *IEEE Transactions on Parallel and Distributed Systems* 15(6):491-504 (2004), DOI 10.1109/TPDS.2004.8.** The hazard-pointer technique of §23.4 and its explicit treatment of the ABA problem. Verified via the IEEE/ACM Digital Library.
- **Herlihy & Shavit, *The Art of Multiprocessor Programming* (Morgan Kaufmann; 2nd ed. 2020), chapters on concurrent stacks, queues, and memory reclamation.** The textbook development of everything here, with proofs of linearizability. [△ reference – not free].

Exercises

- 23.1** WARM-UP In one sentence, what does the *lock-free* progress guarantee promise about the system as a whole?
- 23.2** WARM-UP In the Treiber stack, what single machine operation is the linearization point of a push? Of a pop?
- 23.3** WARM-UP State the ABA problem in one sentence: what does a CAS check, and why can "the same value" be misleading?

23.4 **CORE** Implement the Treiber stack's push and pop as CAS loops (pseudocode is fine). For each, explain (a) what a failed CAS tells you and how the loop recovers, (b) why the successful CAS is the linearization point, and (c) why the structure gives the lock-free guarantee but not the wait-free one.

23.5 **CORE** The Michael-Scott queue enqueues in two CAS steps. (a) Why can it not be done in one? (b) Describe what a second thread can observe about `tail` after the first CAS but before the second, and how "helping" keeps that from stalling anyone. (c) Explain why the dummy node removes the empty-queue special cases at both ends.

23.6 **CORE** Walk through a concrete ABA interleaving on the Treiber stack with nodes freed and reused: thread A reads `top = A` and is about to CAS to `A->next`; describe the sequence of pops, frees, and a push that returns address A, and show exactly how A's CAS then corrupts the stack. Then say which of the three defenses would prevent it and how.

23.7 **COST** (a) Under low contention, why can a lock-free stack beat a mutex-guarded one (recall the ~176 ns vs ~359 ns/op figure), and what is the lock-free structure spending its time on instead of blocking? (b) Under *very high* contention on a single hot pointer, why might the lock-free version's throughput *collapse* — what does a storm of failing CASes cost (Chapter 3's coherence traffic)? (c) Name one structural change (e.g. elimination, or per-thread sub-stacks) that reduces contention, and what it trades away.

23.8 **FADED** *Worked, with the last step for you.* You must make the leaking Treiber stack free popped nodes safely.

Step 1 (given): Freeing a node the instant pop unlinks it is unsafe: another thread mid-pop may still hold that node as its `old` and be about to read `old->next`.

Step 2 (given): Immediate free also enables ABA: the freed address can be reused and pushed, so a stale CAS succeeds on recycled memory.

Step 3 (given): Hazard pointers have each thread publish the node it is about to dereference; a would-be reclaimer frees a retired node only if no hazard pointer references it.

Step 4 (you): Describe how a hazard-pointer-protected pop works step by step (publish `top` as hazardous, re-verify it is still `top`, then read `next` and CAS), and explain why this makes both the use-after-free and the ABA impossible.

23.9 **MIXED** For each requirement, choose between a mutex-guarded structure and a lock-free one, and justify in a line: (a) a shared work queue in a soft-real-time audio callback that must never block; (b) a rarely-contended configuration map read at startup; (c) a structure accessed from a signal handler (Chapter 18); (d) a structure whose operations are large and complex, with modest contention.

23.10 **MIXED** Drawing on Chapters 20–23, classify each as *true* or *false* and fix the false ones: (a) "Lock-free means the code uses no atomic instructions." (b) "A lock-free structure can deadlock if two threads retry forever." (c) "The Michael-Scott queue's dummy node is a performance optimization that can be omitted." (d) "A CAS that finds the pointer equal to what you read proves nothing was popped and freed in between." (e) "Wait-free is a strictly stronger guarantee than lock-free."

23.11 **STRETCH** "Helping" is a recurring lock-free pattern. (a) Explain what problem helping solves in the Michael-Scott queue and why a thread advances *another* thread's half-finished operation rather than waiting. (b) Argue how helping relates to the lock-free guarantee: why does a design where every thread can complete any pending operation avoid a stall when one thread is preempted mid-update? (c) Contrast this with a lock, where a preempted holder stalls everyone — and explain the cost helping pays for that immunity (extra CASes, more complex invariants).

23.12 **LAB** On your own machine: (a) implement the Treiber stack with CAS, run eight threads pushing a known number of items and then popping concurrently, and confirm no item is lost or duplicated (leak the nodes for now). (b) Benchmark it against a mutex-guarded stack under contention and report ns/op for both (note the result is workload-dependent). (c) Implement the tagged-pointer version using a double-width CAS (`-mcx16` on x86-64, linking `-latomic`), confirm it passes the same correctness test, and find the `lock cmpxchg16b` in the disassembly. Label your compiler and CPU; note the throughput figures are illustrative.

Chapter 24 — The Memory Model: `memory_order` and Happens-Before

Before you start: Chapter 22 (atomics — every one takes a memory-ordering argument we have so far ignored), Chapter 23 (a lock-free structure publishes a pointer and trusts the reader to see what it points to — why is that safe?), Chapter 10 (undefined behavior and the optimizing compiler — a data race is UB, and the compiler exploits it), Chapter 3 (each core sees memory through its own cache and store buffer). The last two chapters assumed memory operations become visible in a sensible order across cores; this chapter shows they do not, and specifies exactly what ordering you may demand and how to ask for it.

BEFORE YOU READ ON

From memory, reaching back: (1) Chapter 10 — what does it mean that something is *undefined behavior*, and what license does that give the optimizer? (2) Chapter 3 — a store does not go straight to shared memory; where does it sit first, and why does that let one core lag another's view? (3) *Predict*: two threads run `x = 1; r0 = y;` and `y = 1; r1 = x;` from `x = y = 0`. On paper, can both `r0` and `r1` end up 0? On a real x86 machine, out of two million trials, how many times do you think both read 0? Write your guesses; §24.1 measures it.

Every chapter of this part has leaned on an assumption so natural it was invisible: that there is *one* memory, and that all cores see writes to it in a single consistent order — the model Lamport named **sequential consistency** (1979). It is the mental picture behind "thread A sets the flag, thread B sees the flag, so thread B sees the data A wrote first." That picture is false. Real hardware has per-core store buffers and caches, and both the CPU and the optimizing compiler reorder memory operations for speed, so a program can observe outcomes no sequentially consistent execution allows. This chapter makes the true model precise. First, the demonstration that reordering is real, on x86 no less — a litmus test whose "impossible" outcome happens millions of times (§24.1). Then the framework the C11/C++11 standard built to tame it: a **data race is undefined behavior**, and correctness is defined through the **happens-before** relation (§24.2). Then the menu of orderings an atomic can request — `relaxed`, `acquire/release`, and `seq_cst` — and the *synchronizes-with* edge that makes one thread's writes visible to another (§24.3). Then the workhorse pattern, message passing with `acquire/release`, what it guarantees, and what it costs on real hardware (§24.4). We close the part by carrying this specification forward as the ground truth beneath every concurrent program (§24.5).

24.1 The single-memory illusion: reordering is real

Consider the **store-buffering** litmus test (it underlies Dekker's mutual-exclusion algorithm). Two threads, starting from `x = y = 0`: thread 0 does `x = 1` then reads `r0 = y`; thread 1 does `y = 1` then reads `r1 = x`. Reason it through under a single shared memory: whichever store happens first, the later thread's load must see a 1, so `r0 == 0 && r1 == 0` is *impossible*. Yet it happens — because each core's store sits in a private **store buffer** (Chapter 3) before reaching cache, so each thread reads the other's not-yet-published

value as 0. Here it is with relaxed atomics (to avoid the undefined behavior of a plain data race), two million trials on x86-64 (gcc 11.4.0, Linux 6.18):

```
relaxed store-buffering, 2000000 trials:  
r0==0 && r1==0 observed 349165 times    (about 1 in 5; sequential consistency forbids it)
```

The "impossible" outcome occurred in roughly one trial in five. This is not a bug in the program or the hardware: x86's memory model (formalized as **x86-TSO**, total store order, by Sewell et al., 2010) explicitly allows a store to be reordered after a later load to a *different* address, precisely because of the store buffer. And x86 is one of the *stronger* models — ARM and POWER reorder far more aggressively. Two forces reorder your memory operations: the **hardware**, as here, and the **compiler**, which freely reorders and elides plain loads and stores it cannot prove another thread observes. The consequence is stark: the intuitive single-memory, everyone-agrees-on-one-order model you have been using is not what the machine provides. To write correct concurrent code you need a precise model of what *is* guaranteed — and a way to buy back the ordering you need where you need it.

QUICK CHECK

In the store-buffering test, explain why `r0 == 0 && r1 == 0` is impossible under sequential consistency but common on real x86. Name the hardware structure responsible, and name the *other* agent (besides the CPU) that reorders memory operations. (Answer after the next section.)

24.2 Data races are undefined behavior; happens-before

Faced with hardware and compilers that reorder, the C11/C++11 standard did not promise sequential consistency for ordinary code — that would forbid the optimizations that make programs fast. Instead it drew a bright line. A **data race** — two accesses to the same location from different threads, at least one a write, not ordered by synchronization, with at least one non-atomic — is **undefined behavior** (Chapter 10). Not "unspecified value," not "torn read": undefined, the compiler may assume it never happens and optimize accordingly, so a racy program has no meaning at all. The flip side is the contract: if your program is **data-race-free** — every conflicting pair of accesses ordered by synchronization through atomics — the standard guarantees it behaves as some sequentially consistent execution. This is the DRF-SC guarantee that Boehm and Adve's "Foundations of the C++ Concurrency Memory Model" (2008) established for C and C++: race-free code gets the simple model; racy code gets nothing.

"Ordered by synchronization" is made precise by **happens-before**. Within a single thread, operations are **sequenced-before** one another in program order. Across threads, an atomic store and the atomic load that reads its value can form a **synchronizes-with** edge (the details in §24.3). Happens-before is, roughly, the transitive closure of sequenced-before and synchronizes-with: if A happens-before B, then A's effects are guaranteed visible to B. Two conflicting accesses *not* ordered by happens-before are a data race. And the default atomic ordering, `memory_order_seq_cst`, additionally forces all such atomic operations into a single global total order — restoring the sequentially

consistent illusion for code that uses it. Rerun the store-buffering test with `seq_cst` instead of `relaxed` (gcc 11.4.0):

```
seq_cst store-buffering, 2000000 trials:
r0==0 && r1==0 observed 0 times    (sequential consistency restored)
```

THE SAME IDEA THREE WAYS: MEMORY OPERATIONS REORDER UNLESS AN ORDERING FORBIDS IT


1. In words. Plain and relaxed memory operations may be reordered by the CPU and compiler, so one thread can see another's operations out of program order. Sequential consistency (the default `seq_cst` atomics) forces a single global order that forbids such reorderings; a data race — conflicting accesses not ordered by happens-before — is undefined behavior.

2. As a picture. Figure 24.1: each thread's store sits in a private store buffer before reaching shared memory, so thread 0 reads `y` (still 0) before thread 1's `y = 1` drains — and symmetrically — yielding `r0 == r1 == 0`.

3. As numbers. Store-buffering, two million trials: with relaxed atomics the SC-forbidden outcome happened 349,165 times (~1 in 5); with `seq_cst` it happened 0 times.

(Answer to the §24.1 quick check: under sequential consistency the two stores occur in some global order, and the thread that stores second must, when it later loads, see the first thread's already-committed store — so both loads reading 0 cannot happen. On x86 each store first enters a private **store buffer** and is not yet visible to the other core, so each thread's load reads the other's stale 0. Besides the CPU, the **compiler** also reorders memory operations.)

Store buffering: each store waits in a private buffer, so each load reads a stale 0.



Both loads see the pre-store value: `r0 == r1 == 0`, which sequential consistency forbids.

Figure 24.1: Store buffering. Each core's store to `x` (resp. `y`) sits in a private store buffer before it reaches shared memory, so each core's subsequent load of the other variable reads the old value 0. The outcome `r0 == r1 == 0` is impossible under sequential consistency yet common on real x86 — measured here about one trial in five.

24.3 The ordering menu: relaxed, acquire/release, `seq_cst`

Every atomic operation takes a `memory_order` argument choosing how much ordering it enforces — a dial from "atomic but unordered" to "globally ordered," trading cost for guarantees. There are three levels worth knowing. `memory_order_relaxed` guarantees only atomicity (no torn or lost updates) and coherence on the *same* object; it imposes *no* ordering on other memory operations and creates no synchronizes-with edge. It is right

for a statistics counter whose exact interleaving with other data does not matter — the store-buffering test used it, and indeed it let the reordering show. `memory_order_release` (on a store) paired with `memory_order_acquire` (on a load) is the key construct: if an acquire-load reads the value written by a release-store to the same atomic, the store **synchronizes-with** the load, and everything sequenced-before the release in the storing thread happens-before everything sequenced-after the acquire in the loading thread. That is the machinery that makes a published pointer safe: all the writes that set up the object are visible to any thread that acquires the pointer. `memory_order_seq_cst`, the default, adds to acquire/release a single global total order over all `seq_cst` operations — the strongest and most expensive, and the only one that forbids the store-buffering reordering of §24.1.

The relationships to keep straight: *sequenced-before* is ordering within one thread (program order); *synchronizes-with* is the cross-thread edge a release/acquire pair (or a lock release/acquire, or a thread create/join) creates; and *happens-before* is their transitive closure, the relation that actually determines visibility and defines a data race. Relaxed atomics participate in none of the cross-thread ordering — they are atomic dots with no edges — which is exactly why relaxed is cheap and why it cannot, by itself, publish anything but the atomic's own value. When you need one thread's ordinary writes to become visible to another, you need a release/acquire (or stronger) pair, not relaxed.

QUICK CHECK

What does `memory_order_relaxed` guarantee, and what does it *not*? State precisely what a release-store/acquire-load pair establishes when the acquire reads the release's value. What does `seq_cst` add on top of acquire/release? (Answer below the communicate box.)

24.4 Acquire/release in practice: message passing, and what it costs

The canonical use of acquire/release is **message passing**: one thread fills a payload with ordinary writes, then publishes a flag with a release-store; another thread waits on the flag with an acquire-load and, once it sees the flag, reads the payload. The release/acquire pair guarantees the payload writes happen-before the reads — so the reader never sees a published flag but stale payload. This is precisely the guarantee the lock-free structures of Chapter 23 needed when they published a node pointer. A producer and consumer handing off a four-word payload a million times, verifying every one (gcc 11.4.0, Linux 6.18):

```
/* producer */                /* consumer */
payload[0..3] = ...;          while (load_acquire(&flag) != r) {}
store_release(&flag, r); /* publish */ read payload[0..3]; /* guaranteed visible */

message passing via release/acquire, 1000000 rounds: payload mismatches = 0
```

Not one of a million handoffs saw a stale payload: the acquire that observed the flag was guaranteed, by synchronizes-with, to see every write sequenced before the release. Now

the cost, which is where the memory model gets concrete and surprising. On x86 the ordering you pay for depends entirely on which order you ask for. Compare the compiler's code for stores and loads at each level (gcc 11.4.0, `-O2`, `objdump`):

```
store_relaxed : mov    %edi, f(%rip)    ← plain store
store_release : mov    %edi, f(%rip)    ← plain store (identical!)
store_seqcst  : xchg   %edi, f(%rip)    ← locked exchange – drains the store buffer
load_acquire  : mov    f(%rip), %eax    ← plain load
```

On x86-TSO, release-stores and acquire-loads compile to *plain* `mov` instructions — identical to relaxed — because the hardware already guarantees the ordering they need (it never reorders store-store or load-load, only store-load). The ordering is enforced only by forbidding the *compiler* from reordering across them, at no runtime cost. Only `seq_cst`'s store costs a real instruction — the `xchg` (a locked operation) that drains the store buffer and thereby forbids the store-load reordering that produced §24.1's result. This is the honest picture of the memory model: acquire/release is often free on strong hardware and lets the compiler still forbid its own reorderings, while `seq_cst` buys a global order at the price of a fence. On weaker architectures (ARM, POWER) acquire and release require real barrier instructions too — which is why using the weakest ordering that is correct matters for portable performance. But the danger runs the other way: using an ordering that is too *weak* is not slow, it is *wrong*.

TRAP: USING RELAXED WHERE YOU NEED RELEASE/ACQUIRE (RELAXED ORDERS NOTHING BUT ITSELF)

Because `memory_order_relaxed` is the cheapest ordering and "the flag is atomic anyway," it is tempting to publish data with a relaxed store and read it with a relaxed load. But relaxed creates *no* synchronizes-with edge and therefore *no* happens-before: it guarantees the atomic's own value is not torn, and nothing whatsoever about the ordinary writes around it. So a consumer can see the relaxed flag set to "ready" while the payload writes that preceded it are still invisible — it reads the flag as published and the data as garbage. Worse, it may work on x86 (where the hardware happens to order store-store) and fail on ARM, so it passes your tests and corrupts on a phone. The rule: to publish data, the store that publishes must be `release` and the load that observes it must be `acquire` (or use `seq_cst`); relaxed is only for atomics whose *ordering relative to other memory does not matter*, such as a statistics counter or a stop flag whose late arrival is harmless. If you are using an atomic to make *other* writes visible, relaxed is a correctness bug, not an optimization.

SAY IT LIKE A SYSTEMS ENGINEER

Try it first. A teammate says: "I made the handoff faster. The producer fills the buffer and sets `ready` with a relaxed atomic store; the consumer spins on a relaxed load of `ready` and then reads the buffer. It's all atomic and it passes every test on our CI machines — why do you keep saying it's broken?" Cover the paragraph and respond in four or five sentences, drawing on §24.3 and §24.4.

Model answer. "It passes because your CI machines are x86, whose hardware happens to keep stores in order, so the buffer writes reach memory before the flag — but the C memory model gives you no such guarantee, and the bug is real. A relaxed store creates no synchronizes-with edge and thus no happens-before, so nothing orders the buffer writes relative to the flag; the consumer is entitled to see `ready` set and the buffer still empty, which is exactly what will happen on ARM. The fix costs nothing on x86: make the publishing store `release` and the observing load `acquire`, and now every write before the release is guaranteed visible after the acquire — that pairing is the whole point of the ordering argument. Relaxed is correct only when the atomic's ordering relative to other data doesn't matter, like a counter; here you're using it to *publish* data, so it's a correctness bug, not a speedup." Explaining that x86 hides the bug, naming synchronizes-with/happens-before, and prescribing release/acquire is the senior register.

(Answer to the §24.3 quick check: `relaxed` guarantees atomicity and same-object coherence but imposes no ordering on other memory and creates no synchronizes-with edge. A release-store paired with an acquire-load that reads its value establishes that the store synchronizes-with the load, so everything sequenced-before the release happens-before everything sequenced-after the acquire — making the storing thread's prior writes visible to the loading thread. `seq_cst` adds a single global total order over all `seq_cst` operations, restoring sequential consistency and forbidding the store-buffering reordering.)

24.5 What to carry forward

The single-memory, everyone-agrees model — **sequential consistency** — is an illusion real hardware and compilers break for speed. Our store-buffering litmus test produced the SC-forbidden outcome `r0 == r1 == 0` about one time in five on x86 (349,165 of two million), because each core's store waits in a private buffer before it is visible; the CPU and the compiler both reorder plain memory operations. The C11/C++11 memory model tames this with a bright line: a **data race is undefined behavior**, but a **data-race-free** program behaves as some sequentially consistent execution (Boehm and Adve's DRF-SC guarantee). Correctness is defined through **happens-before** — the transitive closure of *sequenced-before* (program order in a thread) and *synchronizes-with* (the cross-thread edge from a release-store to the acquire-load that reads it). The `memory_order` menu picks how much ordering an atomic buys: `relaxed` (atomic but unordered — no synchronizes-with), `acquire/release` (the publish/observe pair that makes prior writes visible — our message-passing handoff verified a million payloads with zero stale reads), and `seq_cst` (a global total order, the only one that forbade the store-buffering reordering, and on x86 the only one that costs a real fence). The trap: relaxed publishes nothing but its own value, so using it to hand off data is a correctness bug that x86 hides and ARM exposes.

This closes the concurrency part, and with it the last piece of shared-memory correctness. You now hold the full toolkit and, more importantly, the model beneath it: threads share memory and race (Chapter 15); locks and condition variables give blocking mutual exclusion and coordination (Chapters 20–21); atomics and compare-and-swap are the hardware primitives they rest on (Chapter 22); lock-free structures use those primitives to avoid blocking entirely (Chapter 23); and the memory model is the specification that says, precisely, when a write by one thread is guaranteed visible to another (this chapter). The recurring lesson has been the same at every level — the abstract machine's simple, sequential memory is not what runs, and the gap between the two is where concurrency bugs live. What comes next is a language built to close that gap by construction. Everything in this part you had to prove by hand and hold in your head — which data is shared, which access needs a lock, which ordering an atomic requires, whether a pointer outlives the thread reading it — a type system can instead enforce at compile time, turning "I reasoned about it carefully" into "it does not compile otherwise." The next part turns to Rust: ownership, borrowing, and lifetimes as the machinery that encodes these very invariants — data-race freedom not as a discipline you maintain, but as a property the compiler checks.

CAN YOU... WITHOUT LOOKING?

- Describe the store-buffering test and why `r0 == r1 == 0` is SC-forbidden yet common on x86, and name the store buffer.
- State the DRF-SC contract: what a data race is (and that it is UB), and what a race-free program is guaranteed.
- Define sequenced-before, synchronizes-with, and happens-before, and how they relate.
- Say what relaxed, acquire/release, and `seq_cst` each guarantee.
- Explain the message-passing pattern and why a release/acquire pair makes the payload visible.
- Explain why publishing data with a relaxed store is a bug, and why it can pass on x86 and fail on ARM.

RETRIEVAL — CLOSED BOOK

Cover the chapter. Answer from memory; rate confidence first.

1. **R1.** Why does the store-buffering outcome `r0 == r1 == 0` occur on x86, and what forbids it? (§24.1)
2. **R2.** What is a data race, what does the standard say it is, and what does DRF-SC guarantee? (§24.2)
3. **R3.** Define sequenced-before, synchronizes-with, and happens-before. (§24.2–§24.3)
4. **R4.** What do relaxed, acquire/release, and `seq_cst` guarantee, and which is free on x86? (§24.3–§24.4)
5. **R5.** Why is publishing data with a relaxed store a correctness bug, and why might it pass on x86? (§24.4)

FURTHER READING

- **Boehm & Adve, "Foundations of the C++ Concurrency Memory Model," *Proc. ACM PLDI* (2008), pp. 68–78, DOI 10.1145/1375581.1375591.** The DRF-SC design of §24.2: race-free programs get sequential consistency, races are undefined. Author, venue, and year verified via dblp and the ACM Digital Library.
- **Lamport, "How to Make a Multiprocessor Computer That Correctly Executes Multiprocess Programs," *IEEE Transactions on Computers* C-28(9):690–691 (1979), DOI 10.1109/TC.1979.1675439.** The definition of sequential consistency — the model §24.1 shows hardware violating. Verified via IEEE Xplore and dblp.
- **Sewell, Sarkar, Owens, Zappa Nardelli & Myreen, "x86-TSO: A Rigorous and Usable Programmer's Model for x86 Multiprocessors," *Communications of the ACM* 53(7): 89–97 (2010), DOI 10.1145/1785414.1785443.** The precise model of x86 memory ordering behind the store-buffering result of §24.1. Verified via the ACM Digital Library.
- **`cppreference, std::memory_order (en.cppreference.com)`.** The exact semantics of relaxed, acquire/release, and seq_cst, and of synchronizes-with and happens-before. Primary reference for §24.3–§24.4; verify per standard version.
- **Preshing, "Memory Reordering Caught in the Act" (blog, preshing.com, 2012).** A hands-on walkthrough of the store-buffering experiment reproduced in §24.1, with the reordering measured directly. Accessible background; a blog, not a peer-reviewed source.

Exercises

- 24.1** **WARM-UP** In one sentence, what is sequential consistency, and name one hardware structure that causes real machines to violate it.
- 24.2** **WARM-UP** What is a data race, and what does the C11/C++11 standard say a program containing one means?
- 24.3** **WARM-UP** What does `memory_order_relaxed` guarantee, and what does it not guarantee about other memory operations?
- 24.4** **CORE** For the store-buffering test (`x = 1; r0 = y; and y = 1; r1 = x; from x = y = 0`): (a) prove `r0 == r1 == 0` is impossible under sequential consistency. (b) Explain, using the store buffer, why it happens on x86 anyway. (c) Which single change — and to which operations — forbids the outcome, and what does it cost on x86 (name the instruction)?
- 24.5** **CORE** Define *sequenced-before*, *synchronizes-with*, and *happens-before*, and state how happens-before is built from the other two. Then explain, in those terms, why two conflicting accesses not ordered by happens-before are a data race.
- 24.6** **CORE** Write the message-passing pattern with a payload and a flag. (a) Mark which store must be `release` and which load must be `acquire`. (b) Explain, via *synchronizes-with* and *happens-before*, why the consumer is guaranteed to see the payload once it sees the flag. (c) Show what breaks if both are `relaxed`, and why it may still pass on x86.

24.7 **COST** Using the disassembly (relaxed store = `mov`, release store = `mov`, `seq_cst` store = `xchg`, acquire load = `mov`): (a) explain why acquire/release are free on x86 but `seq_cst`'s store is not. (b) What does the `xchg` do that the plain `mov` does not, and how does that forbid store-buffering? (c) Why might the same acquire/release require real barrier instructions on ARM, and what does that imply for "use the weakest correct ordering"?

24.8 **FADED** *Worked, with the last step for you.* A lock-free structure publishes a freshly built node by CAS-ing a pointer, and readers follow the pointer and read the node's fields. Occasionally on an ARM device a reader sees the new pointer but garbage fields.

Step 1 (given): The node's fields are written with plain stores, then the pointer is published.

Step 2 (given): The publish and the read use `relaxed` atomics, which create no synchronizes-with edge.

Step 3 (given): With no happens-before between the field writes and the reader's field reads, the reader may observe the pointer before the field writes are visible — undefined behavior in general, and visible reordering on ARM.

Step 4 (you): State the fix precisely (which operation becomes `release`, which becomes `acquire`), explain why that pairing makes the field writes visible, and say why the bug hid on x86.

24.9 **MIXED** For each atomic use, choose the weakest correct ordering — `relaxed`, `acquire/release`, or `seq_cst` — and justify in a line: (a) a hit counter incremented by many threads, read only at shutdown; (b) a flag that publishes a fully initialized buffer to a consumer; (c) a "please stop" flag a worker polls, where a slightly late stop is harmless; (d) two flags in a Dekker-style algorithm needing store-load ordering.

24.10 **MIXED** Drawing on Chapters 22–24, classify each as *true* or *false* and fix the false ones: (a) "A data race produces an unspecified value but defined behavior." (b) "On x86, a release-store compiles to the same instruction as a relaxed store." (c) "`memory_order_relaxed` makes the writes around it visible to other threads in order." (d) "A race-free program using `seq_cst` atomics behaves as some sequentially consistent execution." (e) "`seq_cst` is free on x86, just like `acquire/release`."

24.11 **STRETCH** The DRF-SC guarantee says race-free programs get sequential consistency and racy ones get undefined behavior. (a) Explain why this bargain lets the compiler and hardware keep their reordering optimizations while still giving well-behaved programs a simple model. (b) Argue why "undefined," not merely "unspecified," is the right (if harsh) definition for a data race — what would the compiler be forbidden from doing under a weaker definition? (c) Connect this to Chapter 10: how is a data race like other undefined behavior the optimizer is allowed to assume never happens, and what does that mean for "but it worked when I tested it"?

24.12 **LAB** On your own machine: (a) implement the store-buffering test with two threads and relaxed atomics, run a couple million trials, and count how often `r0 == r1 == 0` — confirm it is nonzero on x86; switch the atomics to `seq_cst` and confirm the count drops to zero. (b) Disassemble a relaxed store, a release store, a `seq_cst` store, and an acquire load with `objdump`, and identify which one differs (the `seq_cst` store's locked instruction). (c) Build the message-passing handoff with release/acquire and verify a large number of payloads with zero stale reads. Label your compiler and CPU architecture; note that the reordering counts are hardware-dependent.

Chapter 25 — Thread Pools and Task Parallelism

Before you start: Chapter 15 (threads share an address space — the substrate every worker here runs on), Chapter 21 (condition variables and the producer-consumer pattern — a work queue *is* a producer-consumer), Chapter 16 (the scheduler multiplexes runnable threads onto cores — more threads is not more parallelism), Chapter 6 (the cost model — a thread is not free). · The last five chapters were about making concurrent code *correct*; this one is about making parallel code *fast to structure* — decoupling the work you have from the threads that run it.

BEFORE YOU READ ON

From memory, reaching back: (1) Chapter 21 — sketch the producer-consumer pattern: how does a consumer wait for an item without busy-spinning, and what wakes it? (2) Chapter 6 — a thread has a stack (often ~8 MiB of address space) and must be set up and torn down by the kernel; is creating one closer to a function call or to a system call in cost? (3) *Predict*: on a laptop, roughly how many microseconds does it take to `pthread_create` a thread and `pthread_join` it? If you have 100,000 tiny tasks, how much faster might a pool of 8 reused workers be than spawning one thread per task? Write your guesses; §25.1 and §25.2 measure both.

You have a machine with many cores and a pile of work. The naive move is to spawn a thread for each unit of work and let the scheduler sort it out. It is a trap on both ends: a thread is expensive to create and destroy, and thousands of runnable threads oversubscribe the cores and drown the scheduler in context switches. The discipline that fixes both is to *separate the work from the workers*. A fixed pool of long-lived threads pulls units of work — **tasks** — off a shared queue; you submit tasks, the pool runs them. This chapter builds that idea up from the cost of a bare thread. First, why a thread-per-task does not scale — the measured cost of creation and the problem of oversubscription (§25.1). Then the **thread pool**: workers draining a queue, built from the very condition-variable machinery of Chapter 21 (§25.2). Then **task parallelism** and the **fork-join** model — decomposing a computation into tasks whose results you wait on, with **futures** as handles to not-yet-computed values (§25.3). Then the scalability problem a single shared queue creates, and its classic answer, **work-stealing** (§25.4). We close by carrying the task abstraction forward as the unit in which parallel work is expressed (§25.5).

25.1 The cost of a thread: why not one per task

A thread is not a function call. Creating one asks the kernel to allocate a stack, set up a task structure, and make the thread schedulable; joining it tears all that down. That is real work on the system-call boundary of Chapter 13, not the handful of cycles a call costs. Measure it directly — create and immediately join an empty thread, 200,000 times (gcc 11.4.0, `-O2 -pthread`, Linux 6.18 under WSL2; this environment has higher thread-setup overhead than bare metal, so treat the absolute number as illustrative):

```
pthread_create + pthread_join, 200000 iters: 44.07 s (220.4 us each)
```

Roughly *220 microseconds* per thread on this machine. In that time a modern core executes on the order of a million instructions — so if your task is smaller than that, you spend more time managing the thread than doing the work. And there is a second, subtler cost: **oversubscription**. Spawning 100,000 threads for 100,000 tasks does not give you 100,000-way parallelism; you have a fixed number of cores (Chapter 16), so the scheduler must time-slice all those threads onto them, paying a context switch and cache disruption at every slice, plus the memory for every thread's stack. More runnable threads than cores is not more parallelism — it is more overhead. The fix follows from stating the problem precisely: you want as much *parallelism* as you have cores, and as many *tasks* as the problem naturally has, and those two numbers should be decoupled. That decoupling is exactly what a thread pool provides.

QUICK CHECK

Name the two distinct costs that make thread-per-task a bad structure for many small tasks — one paid at creation/teardown, one paid while they run. Why does spawning 100,000 threads on an 8-core machine not give 100,000-way parallelism? (Answer after the next section.)

25.2 The thread pool: workers draining a work queue

A **thread pool** creates a fixed set of worker threads once, up front, and keeps them alive. Work is submitted as **tasks** onto a shared queue; each idle worker loops: take a task, run it, repeat. This is the producer-consumer pattern of Chapter 21 wearing a new name — the queue is guarded by a mutex, and workers block on a condition variable when it is empty rather than busy-waiting, so an idle pool costs no CPU. The core of a worker is exactly the Chapter 21 consumer:

```
void *worker(void *) {
    for (;;) {
        pthread_mutex_lock(&mtx);
        while (queue_empty() && !shutdown) /* predicate loop, never an if */
            pthread_cond_wait(&not_empty, &mtx);
        if (queue_empty() && shutdown) { pthread_mutex_unlock(&mtx); return NULL; }
        task_t t = dequeue();
        pthread_mutex_unlock(&mtx);          /* run the task OUTSIDE the lock */
        t.fn(t.arg);                          /* do the work */
    }
}
```

The pool pays the ~220 us thread-creation cost exactly *w* times — once per worker at startup — and then *amortizes* it across every task ever submitted. Run 100,000 tiny tasks (each a ~200-iteration arithmetic loop) two ways on this machine: one thread per task versus a pool of 8 workers (gcc 11.4.0, -O2 -pthread):

```
thread-per-task : 100000 tasks, 21.73 s (217.2 us/task)
thread pool (8) : 100000 tasks, 0.10 s ( 1.04 us/task)
```


The pool ran the same work about **200× faster** per task, because it stopped paying the thread tax on every task and paid it only eight times. Two design points matter. First, run the task *outside* the lock — the mutex guards the queue, not the work, so holding it during `t.fn()` would serialize the whole pool back down to one worker (that is the contention lesson of the next chapter, previewed). Second, the pool size is a knob: too few workers and cores sit idle; too many and you are back to oversubscription. A common default is one worker per hardware thread for CPU-bound work, more if tasks block on I/O. The pool has bought us the decoupling we wanted: submit as many tasks as the problem has, run them on as many workers as the machine has.

THE SAME IDEA THREE WAYS: A POOL AMORTIZES THREAD COST ACROSS MANY TASKS


1. In words. Creating a thread is expensive and having more runnable threads than cores just adds overhead. A thread pool creates a fixed set of workers once and feeds them tasks from a shared queue, so the creation cost is paid a handful of times and reused for every task, and parallelism is capped at the number of workers rather than the number of tasks.

2. As a picture. Figure 25.1: producers enqueue tasks; a fixed row of workers each loop “lock → wait if empty → dequeue → unlock → run”, blocking on a condition variable when the queue drains.

3. As numbers. 100,000 tiny tasks: thread-per-task took 21.73 s (217 us/task, paying ~220 us creation every time); a pool of 8 took 0.10 s (1.04 us/task) — about 200× faster.

(Answer to the §25.1 quick check: the two costs are (1) **creation/teardown** — the kernel allocates a stack and task structure and makes the thread schedulable, ~220 us measured here, and (2) **oversubscription** — once runnable threads exceed cores, the scheduler time-slices them, paying context switches and cache disruption. 100,000 threads on 8 cores gives at most 8-way parallelism because only 8 threads run at once; the other ~99,992 just wait, so you get the overhead of many threads with the parallelism of eight.)

Thread pool: producers enqueue tasks; a fixed set of workers drain the queue.



Creation cost is paid once per worker at startup, then reused for every task.

Parallelism is capped at the number of workers; task count is whatever the problem needs.

Figure 25.1: A thread pool. A fixed set of workers repeatedly lock the queue, wait on a condition variable while it is empty, dequeue a task, unlock, and run the task outside the lock. Thread-creation cost is paid once per worker; every submitted task reuses those workers.

25.3 Task parallelism and fork-join

A pool runs whatever tasks you give it; **task parallelism** is the discipline of *decomposing a computation* into independent tasks that can run concurrently, as opposed to **data parallelism**, which applies the same operation across many data elements. The canonical structure is **fork-join**: split the work into parts (*fork* subtasks), run them in parallel, then *join* — wait for all of them and combine results. Summing a large array is the textbook case: partition it into *P* ranges, hand each to a thread, and add the partial sums. Measured on this machine (64M doubles, each element passed through `sqrt`; gcc 11.4.0 -O2 -pthread -lm; warmed, averaged over 3 runs):

```
P= 1 : 0.369 s (baseline)
P= 2 : 0.186 s speedup 1.98x
P= 4 : 0.134 s speedup 2.75x
P= 8 : 0.084 s speedup 4.38x
P=16 : 0.076 s speedup 4.86x
```

Two threads nearly halve the time (1.98×) — near-perfect for a problem that splits cleanly — but the speedup curve bends below linear as *P* grows: 4.38× at eight threads, not 8×. Why it bends (memory bandwidth, a serial remainder, this CPU's mix of performance and efficiency cores) is the whole subject of the next chapter; here the point is the *structure*. Fork-join needs a way to hand back a subtask's result, and the abstraction for that is a **future** (sometimes a *promise* on the producing side): a handle to a value that is not computed yet. Submitting a task returns a future immediately; calling `.get()` on it blocks until the task completes and yields the result — the join. Futures let you express a dependency graph directly: fork several tasks, get their futures, and a task that needs their outputs simply gets them, blocking only if they are not ready. The pool of §25.2 plus futures is enough to express most fork-join parallelism: decompose, submit, collect.

QUICK CHECK

Distinguish *task* parallelism from *data* parallelism in one line each. In a fork-join computation, what does “join” do, and what does a *future* represent? Why does calling `future.get()` sometimes block and sometimes return immediately? (Answer below the communicate box.)

25.4 Load balancing and work-stealing

The single shared queue of §25.2 has a scaling flaw: every worker locks the *same* mutex to get its next task, so as you add workers they contend on that one lock, and the queue itself becomes a serial bottleneck — precisely the contention the next chapter measures. It also balances load only coarsely. The standard answer, introduced by the Cilk system (Frigo, Leiserson & Randall, 1998) and analyzed by Blumofe and Leiserson, is **work-stealing**. Give *each* worker its own double-ended queue (a *deque*) of tasks. A worker pushes newly forked subtasks onto, and pops them from, the *bottom* of its own deque — a purely local, uncontended operation in the common case. When a worker's own deque runs empty, it becomes a **thief**: it picks another worker at random and steals a task from the *top* of that victim's deque. Pushing/popping one end locally while thieves take from

the other end keeps the fast path contention-free and distributes load automatically — idle workers go find work rather than waiting to be fed.

This is not just an engineering nicety; it comes with a proof. Blumofe and Leiserson (“Scheduling Multithreaded Computations by Work Stealing,” 1999) showed that for a fully-strict (well-structured fork-join) computation with work T_1 (total operations) and span T_∞ (the critical-path length, the time on infinitely many processors), a randomized work-stealing scheduler runs in expected time $O(T_1/P + T_\infty)$ on P processors — asymptotically optimal, with provably bounded space and communication. In words: you get near-linear speedup as long as there is enough parallelism ($T_1/T_\infty > P$) to keep the processors fed. Work-stealing is why fork-join frameworks scale: Cilk, Intel's TBB, and Doug Lea's Java Fork/Join framework (2000, the engine under Java's parallel streams) are all work-stealing pools. When you write “fork these subtasks and join,” a work-stealing runtime is what turns that into balanced parallelism without your managing which core does what.

TRAP: SUBMITTING TASKS THAT BLOCK ON OTHER TASKS IN THE SAME BOUNDED POOL (POOL DEADLOCK)

A thread pool has a fixed number of workers, and that number is a resource that can be exhausted. The classic deadlock: a task running on the pool submits a *subtask* to the same pool and then *blocks waiting for that subtask's result*. If every worker is simultaneously blocked waiting on a subtask that is still sitting in the queue — because there is no free worker to run it — the pool is wedged forever: the waiters hold the only threads that could make progress. With a pool of W workers, W tasks that each block on an un-started subtask is all it takes. This bites people who treat a pool as infinitely deep, and it is especially nasty because it depends on timing and pool size, so it passes tests and hangs in production under load. The disciplines that avoid it: never block a pool worker waiting on work queued to the *same* pool; use a work-stealing runtime whose join can execute the pending subtask on the current thread instead of blocking (this is exactly what fork-join frameworks do); or separate blocking and CPU-bound work into different pools. If your tasks have dependencies, express them as futures the runtime understands — do not have a worker sit on a lock waiting for a task that has nowhere to run.

SAY IT LIKE A SYSTEMS ENGINEER

Try it first. A teammate says: “Our request handler spawns a fresh thread for every incoming job so they run in parallel. Under load the box is at 100% CPU but throughput actually *drops*, and latency is all over the place. Should we just spawn the threads at higher priority?” Cover the paragraph and respond in four or five sentences, drawing on §25.1–§25.2.

Model answer. “Higher priority won't help — the problem is that you have far more runnable threads than cores, so the CPU is busy context-switching and thrashing caches instead of doing work, and you're also paying ~hundreds of microseconds to create and destroy a thread per job. Bound the concurrency: use a fixed thread pool sized to about the number of hardware threads for CPU-bound jobs, and submit jobs as tasks onto its queue. Now creation cost is paid once per worker and amortized over every job, and you never oversubscribe the cores, so throughput goes up and latency stabilizes. If jobs block on I/O rather than CPU, size that pool larger or give I/O its own pool — but the fix is bounding and reusing threads, not spawning more of them faster.” Naming oversubscription and creation cost, and prescribing a bounded reused pool, is the senior register.

(Answer to the §25.3 quick check: **task** parallelism decomposes a computation into different tasks that run concurrently; **data** parallelism applies the same operation across many data elements. “Join” waits for all forked subtasks to finish and combines their results. A **future** is a handle to a value that may not be computed yet; `future.get()` blocks if the task is still running and returns immediately if it has already completed — it is how a dependent task waits exactly as long as necessary and no longer.)

25.5 What to carry forward

The organizing idea of this chapter is to **separate the work from the workers**. A thread is expensive — about 220 microseconds to create and join on the machine here — and more runnable threads than cores is not more parallelism, just more context-switching overhead, so spawning a thread per task fails on both cost and oversubscription. A **thread pool** fixes both: a fixed set of long-lived workers drains a shared work queue (the producer-consumer pattern of Chapter 21, built from a mutex and a condition variable so an idle pool burns no CPU), paying the creation cost a handful of times and amortizing it across every task — 100,000 tiny tasks ran ~200× faster on a pool of 8 than one-thread-per-task. **Task parallelism** decomposes a computation into tasks; the **fork-join** model forks subtasks, runs them in parallel, and joins their results, with **futures** as handles to not-yet-computed values (our parallel sum hit 1.98× on two threads, then bent below linear). A single shared queue contends, so real fork-join runtimes use **work-stealing** — per-worker dequeues with idle workers stealing from busy ones — which Blumofe and Leiserson proved runs a well-structured computation in near-optimal $O(T_1/P + T_\infty)$ time. The trap: a bounded pool is a finite resource, so a worker that blocks waiting on a subtask queued to the same pool can deadlock it.

You now have the vocabulary for the rest of the part and much of the book: *task* as the unit of work, *pool* as the bounded set of workers, *fork/join* and *future* as the way to express and collect parallel work, and *work-stealing* as the runtime that balances it. But every speedup number in this chapter came with a caveat — two threads gave 1.98×,

eight gave only 4.38×. That gap between the parallelism you asked for and the speedup you got is not sloppiness; it is fundamental, and it has names. The next chapter makes it precise: **Amdahl's law** and its serial fraction, **Gustafson's** rejoinder about growing the problem, and the two silent killers of scalability — **lock contention** and **false sharing** — that turn adding cores into a slowdown. Having learned to express parallel work, we turn to the question of how much of it actually pays off.

CAN YOU... WITHOUT LOOKING?

- State the two costs (creation/teardown and oversubscription) that make thread-per-task a poor structure.
- Describe a thread pool and how it is the Chapter 21 producer-consumer pattern, and why the task runs outside the lock.
- Explain why a pool amortizes thread-creation cost, and roughly what that bought in the 100,000-task measurement.
- Define task vs data parallelism, the fork-join model, and what a future represents.
- Explain work-stealing: per-worker dequeues, local push/pop vs stealing, and what it buys over one shared queue.
- Explain how submitting a blocking, dependent task to a bounded pool can deadlock it, and how to avoid that.

RETRIEVAL — CLOSED BOOK

Cover the chapter. Answer from memory; rate confidence first.

1. **R1.** Why is thread-per-task a bad structure for many small tasks — name both costs. (§25.1)
2. **R2.** What is a thread pool, which Chapter 21 pattern is it, and why is the creation cost amortized? (§25.2)
3. **R3.** Define task parallelism, the fork-join model, and a future. (§25.3)
4. **R4.** What is work-stealing, and what did Blumofe & Leiserson prove about its running time? (§25.4)
5. **R5.** How can a bounded pool deadlock, and how do you prevent it? (§25.4)

FURTHER READING

- **Blumofe & Leiserson, “Scheduling Multithreaded Computations by Work Stealing,” *Journal of the ACM* 46(5):720-748 (1999), DOI 10.1145/324133.324234.** The provably-good work-stealing scheduler of §25.4, with the $O(T_1/P + T_\infty)$ time and space bounds. First author Robert D. Blumofe; verified via dblp and the ACM Digital Library.
- **Frigo, Leiserson & Randall, “The Implementation of the Cilk-5 Multithreaded Language,” *Proc. ACM PLDI* (1998), pp. 212-223, DOI 10.1145/277650.277725.** The work-first principle and the runtime that popularized work-stealing fork-join. First author Matteo Frigo; verified via dblp and the ACM Digital Library.
- **Lea, “A Java Fork/Join Framework,” *Proc. ACM 2000 Java Grande Conference*, pp. 36-43, DOI 10.1145/337449.337465.** A practical work-stealing pool — the design under Java's parallel streams. First author Doug Lea; verified via dblp and the ACM Digital Library.
- **Arpaci-Dusseau & Arpaci-Dusseau, *Operating Systems: Three Easy Pieces*, “Concurrency” chapters (free online, ostep.org).** Threads, locks, and the producer-consumer / condition-variable foundation the pool of §25.2 is built on. Accessible background; verify chapter numbers per edition.

Exercises

- 25.1** **WARM-UP** In one sentence each, name the two costs that make spawning one thread per task a poor structure for many small tasks.
- 25.2** **WARM-UP** A thread pool is an instance of which pattern from Chapter 21? Name the two synchronization primitives its work queue uses and what each is for.
- 25.3** **WARM-UP** What is a *future*, and what happens when you call `.get()` on one whose task has not finished versus one whose task has?
- 25.4** **CORE** Explain why a pool of w workers running N tasks pays the thread-creation cost only about w times, not N times. Using the measured ~ 220 us creation cost and the 1.04 us/task pool result, estimate what fraction of the thread-per-task time (217 us/task) was creation overhead rather than useful work, and why the pool avoids it.
- 25.5** **CORE** In the worker loop of §25.2, the task is run *after* unlocking the queue mutex. (a) What would go wrong for throughput if you held the lock while running `t.fn()`? (b) Why must the empty-queue wait be a `while` loop around `cond_wait`, not an `if` (recall Chapter 21)? (c) How does the pool ensure an idle worker uses no CPU?
- 25.6** **CORE** Describe work-stealing precisely: what data structure does each worker own, which end does the owner use versus a thief, and why does that split keep the common case contention-free? Then state, in words, the Blumofe–Leiserson running-time bound and what condition on the computation is needed for near-linear speedup.

25.7 **COST** Consider a single-shared-queue pool with P workers, each task taking time τ , where dequeuing requires the one global lock held for time c . (a) Argue that when c is not negligible relative to τ , throughput stops improving past some number of workers, and identify what caps it. (b) Explain qualitatively how per-worker dequeues with work-stealing change this. (c) Why does making tasks *coarser* (more work per task) improve a shared-queue pool's scaling, and what is the downside of tasks that are too coarse?

25.8 **FADED** *Worked, with the last step for you.* A service uses a pool of 4 workers. Each request task, midway through, submits a second task to the *same* pool and blocks on its future before finishing. Under light load it is fine; under load it hangs.

Step 1 (given): The pool has exactly 4 worker threads; parallelism is capped at 4 running tasks.

Step 2 (given): Under load, 4 request tasks each reach the point of submitting a subtask and blocking on its future — all 4 workers are now blocked.

Step 3 (given): The 4 subtasks sit in the queue with no free worker to run them, so none of the futures will ever complete.

Step 4 (you): Name this failure, explain precisely why it is a deadlock (what holds what, what waits on what), and give two concrete fixes — one about pool structure and one about how dependent work should be expressed.

25.9 **MIXED** For each workload, say whether a fixed pool sized to the hardware-thread count is right, and if not what to change: (a) CPU-bound numeric tasks that never block; (b) tasks that spend most of their time blocked on network I/O; (c) a recursive divide-and-conquer computation with millions of tiny subtasks and deep nesting; (d) a mix of short CPU tasks and occasional long-blocking database calls. Justify each in a line.

25.10 **MIXED** Drawing on Chapters 21 and 25, classify each as *true* or *false* and fix the false ones: (a) “Spawning one thread per task gives you as much parallelism as you have tasks.” (b) “A thread pool's work queue is a producer-consumer buffer guarded by a mutex and a condition variable.” (c) “An idle thread pool spins, burning CPU while it waits for work.” (d) “Work-stealing eliminates the single shared queue by giving each worker its own deque.” (e) “Calling `future.get()` always blocks until the whole pool is idle.”

25.11 **STRETCH** Work-stealing is described as “idle workers steal from busy ones.” (a) Explain why stealing from the *top* (opposite end from the owner's push/pop) both reduces contention with the owner and tends to steal the *oldest*, largest-grained task — and why stealing a large task is good for load balance. (b) The Blumofe–Leiserson bound is $O(T_1/P + T_\infty)$. Interpret each term, and explain why a computation with too little parallelism (span T_∞ close to work T_1) cannot be sped up much no matter how many workers you add — and connect this to the next chapter's Amdahl's law.

25.12 **LAB** On your own machine: (a) measure the cost of `pthread_create` + `pthread_join` by creating and joining an empty thread in a loop; report microseconds per thread and your CPU/OS. (b) Build a small thread pool (fixed workers, a mutex+condvar work queue) and run a large number of tiny tasks through it, then run the same tasks one-thread-per-task; report the per-task times and the speedup. (c) Implement a fork-join parallel sum of a large array for $P = 1, 2, 4, 8$ threads and plot speedup; note where it bends below linear. Warm up and average your runs, and state that absolute numbers are hardware- and OS-dependent.

Chapter 26 — Scalability: Amdahl's Law, Contention, and False Sharing

Before you start: Chapter 25 (fork-join gave 1.98× on two threads but only 4.38× on eight — this chapter explains the gap), Chapter 6 (the cost model — speedup is a ratio you measure, not assume), Chapter 3 (caches and the 64-byte line — false sharing lives here), Chapter 20 (a critical section is serialized — that is a serial fraction). · The last chapter taught you to *express* parallel work; this one asks the harder question — how much of it actually pays off, and what silently steals the rest.

BEFORE YOU READ ON

From memory, reaching back: (1) Chapter 20 — a critical section runs one thread at a time; if every thread must pass through the same lock, what happens to parallel speedup? (2) Chapter 3 — a cache coherence protocol keeps one value per line consistent across cores; what must happen when one core writes a line another core holds? (3) *Predict*: eight threads each increment their *own* separate counter, a billion times. If the eight counters sit in one array (adjacent in memory) versus each padded onto its own 64-byte cache line, how much slower do you think the adjacent version runs — 1.1×, 2×, or 10×? Write your guess; §26.4 measures it.

Adding cores is supposed to make a program faster. Chapter 25 showed it does — up to a point, and then the curve bends. This chapter is about that ceiling and the two forces that push you well below it. The first ceiling is arithmetic and unavoidable: any part of the program that must run serially caps the whole thing, no matter how many cores you throw at it — that is **Amdahl's law** (§26.1). Its optimistic counterpart, **Gustafson's law**, observes that in practice we grow the problem to fit the machine, which changes the accounting (§26.2). Then the two silent killers, both of which turn “add a core” into “go slower.” **Lock contention**: when threads pile up on the same lock, the critical section becomes a serial bottleneck and scaling goes *negative* (§26.3). And **false sharing**: two threads touching *different* variables that happen to share a cache line ping-pong that line between cores, paying coherence traffic for data they never actually share (§26.4). We close the concurrency part by carrying forward a discipline: measure speedup honestly, and hunt the serial fraction (§26.5).

26.1 Amdahl's law: the serial fraction is the ceiling

Suppose a fraction s of a program's runtime is inherently serial — it cannot be parallelized — and the remaining $1 - s$ parallelizes perfectly across P processors. The parallel part shrinks to $(1 - s)/P$, but the serial part s does not move. So the speedup is

$$\text{speedup}(P) = 1 / (s + (1 - s)/P)$$

This is **Amdahl's law** (Gene Amdahl, 1967). Its brutal consequence is the limit as $P \rightarrow \infty$: the $(1 - s)/P$ term vanishes, and speedup tops out at **1/s**. If 10% of your program is serial, the most you can ever get — with infinite cores — is 10×. Not 100×, not 1000×: ten. The serial fraction, not the core count, is the ceiling. Measure it: a compute-bound workload

(no shared memory, so we isolate the serial fraction from the contention effects below) with a serial section of about 9% and the rest split across P threads (gcc 11.4.0, -O2 -pthread -lm, this machine; illustrative):

```
measured serial fraction s = 0.089    (ceiling 1/s = 11.2x)
P= 2 : speedup 2.01x    Amdahl predicts 1.84x
P= 4 : speedup 3.20x    Amdahl predicts 3.16x
P= 8 : speedup 4.58x    Amdahl predicts 4.93x
P=16 : speedup 6.37x    Amdahl predicts 6.85x
```

The measured speedups track Amdahl's prediction closely and are visibly bending toward the 11.2× ceiling set by the 9% serial part — sixteen threads bought only 6.4×, and no number of threads would break 11.2×. This reframes optimization: to make a parallel program scale, attack the *serial fraction* — the sequential setup, the single lock every thread passes through, the one-threaded merge at the end. Doubling cores helps less and less; shrinking s raises the ceiling for everyone. Amdahl's law is why “we'll just add more cores” is rarely the answer, and why profiling to find what is *not* parallel is the highest-leverage move.

QUICK CHECK

Write Amdahl's law for speedup on P processors with serial fraction s . What is the limit as $P \rightarrow \infty$? If a program is 5% serial, what is the best possible speedup, and does buying a 128-core machine change that ceiling? (Answer after the next section.)

26.2 Gustafson's law: grow the problem with the machine

Amdahl's law fixes the *problem size* and asks how much faster a bigger machine runs it — a pessimistic framing, and for a fixed problem, an honest one. But John Gustafson (1988), working on a 1024-processor machine at Sandia, observed that people do not buy bigger machines to run the *same* problem faster; they buy them to run *bigger* problems in the same time. When you scale the problem up with the processor count, the serial part (setup, I/O) tends to stay roughly constant while the parallel part grows with the data. Under that assumption, **Gustafson's law** gives the *scaled* speedup:

```
scaled_speedup(P) = P - s*(P - 1)    (s = serial fraction of the scaled work)
```

This grows nearly linearly with P — no fixed ceiling — because as the problem grows, the parallel fraction grows too, and the constant serial part becomes a smaller share of a larger job. Gustafson reported over 1000-fold speedups on 1024 processors precisely this way. The two laws are not contradictory; they answer different questions. **Amdahl** is *strong scaling*: fixed problem, more cores — the serial fraction caps you at $1/s$. **Gustafson** is *weak scaling*: problem grows with cores — scaling stays near-linear. Which one governs depends on what you actually do when you get a bigger machine. If your problem size is fixed (render *this* frame, sort *this* file), Amdahl rules and you must attack the serial fraction. If it grows (train on more data, simulate a finer grid), Gustafson says throwing cores at a proportionally bigger problem keeps paying off. The practical lesson

is to know which regime you are in before you reason about whether more cores will help.

THE SAME IDEA THREE WAYS: THE SERIAL FRACTION GOVERNS SCALING


1. In words. For a fixed problem, speedup is capped by the serial fraction at $1/s$ (Amdahl) — the part that can't parallelize dominates as cores grow. If instead you grow the problem with the machine, the parallel part grows too and the constant serial part shrinks in relative terms, so speedup stays near-linear (Gustafson).

2. As a picture. Figure 26.1: the serial slice (dark) stays fixed while the parallel slice (light) shrinks as you add cores — Amdahl; or the parallel slice grows with the problem while the serial slice stays put — Gustafson.

3. As numbers. With $s = 0.089$, Amdahl caps speedup at $11.2\times$ (16 threads got $6.4\times$, tracking the prediction). Gustafson: $P - s(P-1)$ at $P = 16$ with the same s is $14.7\times$ — near-linear, because the work grew with the cores.

(Answer to the §26.1 quick check: $\text{speedup}(P) = 1/(s + (1-s)/P)$; as $P \rightarrow \infty$ it approaches $1/s$. At 5% serial the ceiling is $1/0.05 = 20\times$, and a 128-core machine cannot beat it — the serial fraction, not the core count, sets the limit.)

Amdahl (fixed problem): the serial slice stays fixed; only the parallel slice shrinks with more cores.



Gustafson (grow the problem): the parallel slice grows with P while the serial slice stays put — near-linear scaling.



Figure 26.1: Amdahl vs Gustafson. Under Amdahl (fixed problem) the serial slice is a fixed floor, so speedup approaches $1/s$ no matter how many cores divide the parallel slice. Under Gustafson (problem grows with the machine) the parallel slice grows while the serial slice stays constant, so scaled speedup stays near-linear.

26.3 Contention: the critical section is a serial fraction

Amdahl's serial fraction is often not an explicit sequential block — it is *hidden inside a lock*. A critical section runs one thread at a time (Chapter 20), so if every thread must pass through the same lock on every operation, that lock *is* the serial fraction, and Amdahl's ceiling applies to it. Worse, contention does not merely fail to speed up — it can make the program *slower* as you add threads, because the lock itself becomes a contended cache line that bounces between cores, and threads spend their time waiting and handing off rather than working. Measure the extreme case: many threads incrementing one shared counter, each increment under one global mutex (gcc 11.4.0, `-O2 -pthread`, this machine):

```

one global mutex, 40M increments total:
P= 1 : 1.80 s    speedup 1.00x
P= 2 : 5.42 s    speedup 0.33x
P= 4 : 5.52 s    speedup 0.33x
P= 8 : 7.06 s    speedup 0.26x

```

Adding threads made it *nearly three times slower*, not faster. This is **negative scaling**: the entire work is inside the critical section, so it is 100% serial (Amdahl ceiling 1×), and on top of that the lock's cache line ping-pongs among the contending cores, adding coherence overhead that grows with P . The lesson is that a lock on a hot path is a scalability tax, and the fixes are all about *shrinking or removing the serial fraction*: make the critical section smaller (hold the lock for less), shard the state so threads touch different locks (per-thread or per-bucket counters combined at the end), or use a lock-free structure (Chapter 23) or per-thread accumulation to remove the shared write entirely. “Just add threads” applied to a contended lock is how you make a program slower and call it optimization.

QUICK CHECK

Why is a lock on a hot path a “serial fraction” in Amdahl's sense? Explain how adding threads to a single-mutex counter can make it *slower*, not just fail to speed up — name what physically bounces between cores. Give one fix that shrinks the serial fraction. (Answer below the communicate box.)

26.4 False sharing: contention you did not write

The most insidious scalability bug involves no shared variable at all. Recall from Chapter 3 that cache coherence operates at the granularity of a **cache line** — 64 bytes on this machine, not one variable. When one core writes any byte of a line, the coherence protocol invalidates that line in every other core's cache; a core that then reads or writes its own, *different* variable in the same line must re-fetch the whole line. So if two threads update two *independent* variables that happen to land in the same 64-byte line, every update by one invalidates the other's cache, and the line ping-pongs between cores on every write — the threads pay the full cost of sharing a variable while sharing nothing. This is **false sharing**. Measure it: eight threads, each incrementing its own counter 200 million times, with the eight counters packed adjacently in an array (they fall in the same one or two lines) versus each padded to its own 64-byte line (gcc 11.4.0, -O2 -pthread, this machine, 64-byte lines):

```

packed (false sharing) : 2.56 s
padded (own cache line): 0.25 s
false sharing slowdown : ~10x

```

The *same* arithmetic ran about **ten times slower** purely because the counters shared cache lines. No variable is shared between threads; the bug is entirely in the memory layout. The fix is to **pad and align** hot per-thread data so each thread's data owns its own cache line — the padded version's per-counter struct is exactly 64 bytes. This is the same mechanical-sympathy discipline as Chapter 3 and 5: layout is performance. False

sharing is especially dangerous because it is invisible in the source — the code looks embarrassingly parallel and correct, and it *is* correct; it is just slow, and profilers that don't attribute cache-coherence stalls can hide it. Whenever independent threads write to nearby memory — adjacent array slots, fields of a shared struct, a counters table — suspect false sharing and check the line layout.

TRAP: FALSE SHARING — INDEPENDENT DATA COLLIDING ON ONE CACHE LINE (CORRECT BUT 10× SLOW)

You split a counter into one slot per thread precisely to avoid contention: `long counts[NUM_THREADS];`, each thread touches only `counts[my_id]`. There is no shared variable, no data race, the code is correct — and it is several times slower than you expect and gets worse as you add threads. The reason is that `counts` is a dense array, so eight longs fit in one or two 64-byte cache lines, and every thread's write to its own slot invalidates the whole line in every other thread's cache — the line ping-pongs between cores on every increment. The threads share a cache line even though they share no data, so they pay the full coherence cost of true sharing. The tell is that padding each slot to its own line — `struct { long v; char pad[56]; } counts[NUM_THREADS];`, or C11 `alignas(64) / __attribute__((aligned(64)))` — makes the slowdown vanish (the §26.4 measurement went from 2.56 s to 0.25 s). The discipline: hot, frequently-written per-thread data must not share a cache line with another thread's hot data; align it to the line size, and be suspicious of any “per-thread” array of small elements. It is a correctness-preserving performance bug, which is why it survives testing and hides in production.

SAY IT LIKE A SYSTEMS ENGINEER

Try it first. A teammate says: “I parallelized the histogram: each of the 8 threads owns one entry of a `long hist[8]` array, so there's no lock and no shared variable — but with 8 threads it's barely faster than 1, and with 16 it's actually slower. The threads are totally independent, so I don't get it.” Cover the paragraph and respond in four or five sentences, drawing on §26.3–§26.4.

Model answer. “The threads are independent in the source but not in the cache: those eight longs sit in one 64-byte cache line, so every increment by any thread invalidates that line in every other core, and the line ping-pongs between cores on every write — that's false sharing, and it costs as much as truly sharing the variable. Nothing is wrong with correctness; the bug is the memory layout. Pad each counter onto its own cache line — an `alignas(64)` struct per thread, or 64 bytes of stride — and the slowdown disappears; in a similar test that took a per-counter workload from ~2.5 s to ~0.25 s. Adding a sixteenth thread made it worse because you added another core fighting over the same line, more coherence traffic, not more work done. The rule is that hot per-thread data must own its cache line.” Naming false sharing, the 64-byte line, and the padding fix — and distinguishing it from a correctness bug — is the senior register.

(Answer to the §26.3 quick check: a lock serializes its critical section — one thread at a time — so a lock every thread passes through on the hot path is a serial fraction, and Amdahl caps speedup at 1 over that fraction. Adding threads makes it slower because the lock's own **cache line** bounces between the contending cores (coherence traffic) and threads spend time waiting and handing off rather than working, so overhead grows with **P**. A fix that shrinks the serial fraction: shard the state into per-thread or per-bucket counters combined at the end, or shrink/remove the critical section with a lock-free or per-thread approach.)

26.5 What to carry forward

Parallel speedup has a ceiling, and it is set by the part that does *not* parallelize. **Amdahl's law** — $\text{speedup}(P) = 1/(s + (1-s)/P)$ — says a fixed problem with serial fraction s tops out at $1/s$ no matter how many cores you add (our 9%-serial workload bent toward $11.2\times$ and 16 threads got only $6.4\times$). **Gustafson's law** — $P = s(P-1)$ — answers the other question: grow the problem with the machine and scaling stays near-linear, because the parallel work grows while the serial part stays constant. Strong scaling (fixed problem) is Amdahl; weak scaling (growing problem) is Gustafson; know which you are in. Two forces push you below even Amdahl's ceiling. **Lock contention** turns a hot lock into a serial fraction *and* a bouncing cache line, so adding threads can go *negative* — our single-mutex counter went from 1.8 s at one thread to 7.1 s at eight. **False sharing** makes independent threads collide on a shared 64-byte cache line, paying the full cost of sharing while sharing nothing — our packed counters ran $\sim 10\times$ slower than the padded ones, fixed by aligning hot per-thread data to its own line.

This closes the concurrency and parallelism part, and with it Volume 1's tour of how the machine really behaves. The through-line of the whole part has been a single, unforgiving fact: the abstract machine's simple, sequential, single-memory model is not what runs, and every gap between the model and the metal — reordering, the store buffer, the coherence protocol, the cost of a thread, the serial fraction hiding in a lock — is where concurrency's bugs and slowdowns live. You now hold the full shared-memory toolkit: threads and their hazards (Chapter 15), locks and condition variables (Chapters 20-21), atomics and lock-free structures (Chapters 22-23), the memory model that says when a write is visible (Chapter 24), the pool-and-task machinery that structures parallel work (Chapter 25), and now the scalability laws that tell you how much of it pays off. Every one of these correctness properties you have had to prove by hand and hold in your head — which data is shared, which access needs a lock, which ordering an atomic requires, whether padding avoids false sharing. The next part turns to a language built to encode those very invariants in its type system, so that data-race freedom and memory safety become properties the compiler checks rather than disciplines you maintain: Rust, and it begins with the idea beneath all of it — ownership.

CAN YOU... WITHOUT LOOKING?

- State Amdahl's law and its $1/s$ limit, and explain why the serial fraction — not the core count — is the ceiling.
- State Gustafson's law and explain how it answers a different question (weak vs strong scaling).
- Explain why a hot lock is a serial fraction and how contention can make a program slower as threads are added.
- Define false sharing, name the granularity (the 64-byte cache line) that causes it, and give the padding/alignment fix.
- Given a scaling problem, say whether to suspect the serial fraction, lock contention, or false sharing, and why.
- Explain why false sharing is a correctness-preserving performance bug that hides in testing.

RETRIEVAL — CLOSED BOOK

Cover the chapter. Answer from memory; rate confidence first.

1. **R1.** State Amdahl's law and its limit as $P \rightarrow \infty$; what does it imply about a 5%-serial program? (§26.1)
2. **R2.** State Gustafson's law and explain how it differs from Amdahl (weak vs strong scaling). (§26.2)
3. **R3.** Why is a hot lock a serial fraction, and how can contention cause negative scaling? (§26.3)
4. **R4.** What is false sharing, what granularity causes it, and what is the fix? (§26.4)
5. **R5.** Contrast lock contention and false sharing — what is actually shared in each? (§26.3–§26.4)

FURTHER READING

- **Amdahl, “Validity of the Single Processor Approach to Achieving Large Scale Computing Capabilities,”** *AFIPS Conference Proceedings*, vol. 30 (1967), pp. 483–485, DOI 10.1145/1465482.1465560. The original statement of the serial-fraction ceiling of §26.1. First author Gene M. Amdahl; verified via dblp and the ACM Digital Library.
- **Gustafson, “Reevaluating Amdahl's Law,”** *Communications of the ACM* 31(5):532–533 (1988), DOI 10.1145/42411.42415. The scaled-speedup (weak-scaling) rejoinder of §26.2, from the 1024-processor experiments at Sandia. First author John L. Gustafson; verified via the ACM Digital Library.
- **Hill & Marty, “Amdahl's Law in the Multicore Era,”** *IEEE Computer* 41(7):33–38 (2008), DOI 10.1109/MC.2008.209. Amdahl's law applied to multicore chip design — how the serial fraction shapes core-count trade-offs. First author Mark D. Hill; verified via IEEE Xplore.
- **Drepper, “What Every Programmer Should Know About Memory” (2007, Red Hat; freely available online).** Section 6 covers false sharing and cache-line effects behind §26.4 in depth. A technical whitepaper (no DOI); widely cited reference, verify the current PDF.

Exercises

26.1 **WARM-UP** Write Amdahl's law for speedup on P processors with serial fraction s , and state its limit as $P \rightarrow \infty$.

26.2 **WARM-UP** Define *false sharing* in one sentence, and name the memory granularity (and its size on a typical x86 machine) that causes it.

26.3 **WARM-UP** In one sentence, why is a lock that every thread passes through on the hot path a “serial fraction” in Amdahl's sense?

26.4 **CORE** A program is 20% serial. (a) Compute the best possible speedup with infinite cores. (b) Compute the actual speedup with $P = 4$ and $P = 16$ from Amdahl's law. (c) Going from 4 to 16 cores quadruples the hardware — by what factor does the speedup actually improve, and what does that say about “just add cores”?

26.5 **CORE** Explain, in Amdahl's terms, why the single-mutex counter of §26.3 exhibited *negative* scaling (1.80 s at one thread, 7.06 s at eight). (a) What is the serial fraction of that workload? (b) Beyond the serial fraction, what extra cost grows with the number of threads, and what physically causes it? (c) Give two distinct fixes and say what each does to the serial fraction.

26.6 **CORE** Two threads each increment their own `long` counter a billion times. Version A stores them as `long c[2];`; version B pads each onto its own 64-byte cache line. (a) Which is faster and by roughly how much (cite the §26.4 result), and why — what happens to the cache line on each write? (b) Confirm there is no data race or shared variable in either version. (c) Show the padding fix in code (a struct with padding, or an alignment attribute).

26.7 **COST** For a workload with serial fraction s and perfectly parallel remainder: (a) Derive the efficiency $\text{speedup}(P)/P$ as a function of s and P , and explain why efficiency falls as P grows for any $s > 0$. (b) At what point does buying more cores stop being cost-effective, in terms of efficiency? (c) Contrast this with Gustafson's regime: why can efficiency stay high under weak scaling even though Amdahl's efficiency falls?

26.8 **FADED** *Worked, with the last step for you.* A parallel word-count gives each of 8 threads a slot in `long counts[8]` to tally its chunk. It is correct but barely faster than single-threaded, and adding threads makes it worse.

Step 1 (given): There is no lock and no shared variable — each thread writes only `counts[my_id]`, so it is not lock contention.

Step 2 (given): The eight `long`s occupy 64 bytes — one cache line — so all eight slots live on the same line.

Step 3 (given): Each thread's write invalidates that line in every other core's cache, so the line ping-pongs between cores on every increment.

Step 4 (you): Name the phenomenon, explain why more threads make it worse, and give the precise fix (including how you would lay out or align the counters), citing what the §26.4 measurement showed the fix is worth.

26.9 **MIXED** For each scaling problem, name the most likely cause — *serial fraction* (Amdahl), *lock contention*, or *false sharing* — and one diagnostic that would confirm it: (a) speedup plateaus at $\sim 10\times$ regardless of cores, with no lock on the hot path; (b) throughput *drops* as threads are added, and a profiler shows most time in `pthread_mutex_lock`; (c) an embarrassingly parallel loop over a per-thread results array barely scales, with no locks anywhere; (d) a single-threaded merge at the end takes a fixed 2 seconds no matter how fast the parallel phase runs.

26.10 **MIXED** Drawing on Chapters 20, 23, and 26, classify each as *true* or *false* and fix the false ones: (a) “With enough cores, any parallel program can reach arbitrarily high speedup.” (b) “Gustafson's law contradicts Amdahl's law.” (c) “False sharing is a data race.” (d) “A lock on the hot path can make a program slower as you add threads.” (e) “Padding per-thread counters to separate cache lines changes the program's results.”

26.11 **STRETCH** Amdahl and Gustafson give different answers because they hold different things fixed. (a) Show algebraically that if you fix the problem size and let P grow, Amdahl's $1/s$ ceiling follows; if instead you fix the *time* and grow the problem with P , Gustafson's near-linear scaled speedup follows. (b) A team reports “linear speedup to 1000 cores” on their simulation. Explain how that can be true under Gustafson yet not contradict Amdahl, and what question you would ask to tell which regime they measured. (c) Connect this to Chapter 25's span bound $O(T_1/P + T_\infty)$: how is the span like Amdahl's serial fraction?

26.12 **LAB** On your own machine: (a) build a workload with a tunable serial fraction (a fixed serial section plus a parallel section split across P threads), measure speedup for $P = 1, 2, 4, 8, 16$, and compare to Amdahl's prediction for your measured s . (b) Increment one shared counter under a single mutex with 1, 2, 4, 8 threads and observe whether scaling goes negative. (c) Have N threads each increment their own counter a billion times, packed in an array versus padded to 64-byte lines, and report the false-sharing slowdown. Label your CPU (and its cache-line size) and note that the numbers are hardware-dependent.

VOL 1 · SYSTEMS PROGRAMMING & PERFORMANCE · PART V

Rust for Systems

Ownership, borrowing, and lifetimes — memory safety without a GC, and how the type system encodes the invariants the previous parts made you prove by hand.

Chapter 27 — Ownership and Moves

Before you start: Chapter 8 (the stack and the heap — ownership is about who is responsible for heap memory), Chapter 9 (malloc/free and what goes wrong — use-after-free, double-free, leaks are exactly what ownership prevents), Chapter 24 (a data race is undefined behavior — the memory model you had to obey by hand, a type system can enforce), Chapter 10 (undefined behavior and the optimizing compiler — the failure mode ownership rules out). · This chapter opens the Rust part; every hazard the C parts made you track by hand, Rust encodes as a rule the compiler checks — and it all rests on one idea: ownership.

BEFORE YOU READ ON

From memory, reaching back: (1) Chapter 9 — what are the three classic bugs of manual malloc/free memory management, and which one does calling free twice on the same pointer cause? (2) Chapter 8 — when does a stack variable's storage get reclaimed, and does the compiler need you to say so? (3) *Predict*: in C, you write `char *t = s;` then `free(s)` then use `t` and `free(t)`. Does that *compile*? Does it *run* correctly? In Rust, the closest equivalent — `let t = s;` then using `s` — will the compiler accept it? Write your guesses; §27.1 and §27.3 answer all three.

A note on this part before we begin. The C parts of this book gave you no safety net on purpose: you allocated and freed by hand, reasoned about lifetimes yourself, and proved data-race freedom by careful thought (Chapter 24). Rust keeps the same low-level control — no garbage collector, the same machine model, predictable performance — but moves those proofs into the type system, so the compiler rejects the programs you would otherwise have to keep correct by discipline. The engine of that is **ownership**, and this chapter is only about ownership and its most surprising consequence, **moves**. First, the problem ownership solves — the manual-memory bugs of Chapter 9, shown in real C that compiles and then crashes (§27.1). Then the three **ownership rules** and how “dropped when the owner goes out of scope” gives automatic, deterministic cleanup with no GC (§27.2). Then **move semantics**: assigning or passing a value transfers ownership and *invalidates the source*, so use-after-move is a compile error, not a runtime crash (§27.3). Then ownership across function boundaries, and a first glimpse of why this same rule makes data races impossible (§27.4). We close by carrying ownership forward toward borrowing — the way to use a value without taking ownership (§27.5).

On the code in this part. The book's environment has a C toolchain (gcc 11.4.0) but *no Rust compiler installed*, so the C programs below were compiled and run and their output is real, whereas the Rust snippets were *not* compiled here; their behavior and the quoted compiler diagnostics are drawn from *The Rust Programming Language* (Klabnik & Nichols, 2nd ed., 2023) and the official Rust documentation, and are labeled as such. Nothing about Rust's behavior below is invented — it is cited.

27.1 The problem ownership solves

Recall the three failure modes of manual memory management from Chapter 9: **use-after-free** (touching memory you already returned), **double-free** (returning the same block twice), and **leaks** (never returning it). All three come from the same root: in C,

freeing memory is decoupled from any notion of *who owns it*, so nothing stops two pointers from aliasing one block and both trying to manage it. Here is the bug in real C — it compiles cleanly with no warning, and then aborts at run time (gcc 11.4.0, -02, Linux 6.18):

```
char *s = malloc(16);
strcpy(s, "hello");
char *t = s;          /* alias: t and s point at the SAME buffer */
free(s);              /* one "owner" frees it */
printf("%s\n", t);    /* use-after-free through t - undefined behavior */
free(t);              /* double free */

$ gcc -02 usefree.c -o usefree      # compiles, exit 0, no warning
$ ./usefree
free(): double free detected in tcache 2
Aborted (core dumped)              # exit 134 (SIGABRT)
```

The compiler accepted it — the aliasing and the second `free` are invisible to the C type system — and the program died at run time (a lucky outcome; use-after-free often *silently* corrupts instead). The root cause is that `s` and `t` both claim the right to free the same buffer, and C has no way to express that only one of them should. Rust's answer is to make *ownership* a first-class, compiler-enforced concept: every value has exactly one owner, responsibility for freeing follows the owner, and the type system tracks it. The promise is **memory safety without a garbage collector** — the same manual, predictable, no-runtime layout as C, but with use-after-free, double-free, and leaks of owned memory ruled out at compile time. That the guarantee is real, not just a convention, was established formally by the RustBelt project (Jung et al., 2018), which gave a machine-checked safety proof for a realistic subset of the language. The rest of this chapter is how the rule works.

QUICK CHECK

In the C snippet, name the two distinct bugs that occur once `t` aliases `s` and `s` is freed. Why can the C compiler not catch either? State Rust's one-sentence promise about memory safety and the garbage collector. (Answer after the next section.)

27.2 The ownership rules and drop

Rust's model is three rules, stated in *The Rust Programming Language* (Klabnik & Nichols, 2023, ch. 4):

1. Each value in Rust has an owner.
2. There can only be one owner at a time.
3. When the owner goes out of scope, the value is dropped.

“Dropped” means Rust automatically runs cleanup — for a heap-owning type like `String`, its `drop` frees the heap buffer — at the exact point the owning variable's scope ends. This is **RAII** (resource acquisition is initialization, the C++ idea from Chapter 12's arena discussion): cleanup is tied to scope, not to a manual call. Because there is exactly *one* owner (rule 2), the buffer is freed exactly *once* (rule 3) — no double-free — and because

drop is automatic at end of scope, it is never forgotten — no leak of owned memory. Consider (Rust; not compiled here — behavior per the cited book):

```
{
    let s = String::from("hello"); // s owns a heap buffer
    // ... use s ...
}                                // s goes out of scope here: Rust calls drop(s),
                                // which frees the buffer — automatically, exactly once
```

No `free` call appears, yet the memory is released deterministically at the closing brace — not by a garbage collector scanning the heap later, but by a `drop` the compiler inserts at compile time, at a place you can point to in the source. This is the crucial difference from a GC language: the cleanup is *deterministic* and has the same performance characteristics as the manual `free` you would have written in C, without your having to write it or risk forgetting it. Ownership turns “remember to free exactly once, at the right time” — a discipline you maintained by hand and sometimes got wrong — into a property of the program's structure that the compiler guarantees. The single-owner rule is what makes it sound; the next section shows the price it extracts, which is that ordinary assignment does not mean what it means in C.

THE SAME IDEA THREE WAYS: ONE OWNER FREES EXACTLY ONCE, AUTOMATICALLY


1. In words. Every value has exactly one owner; when the owner's scope ends, the value is dropped (its resources freed). One owner means one free — no double-free — and automatic drop means no forgotten free — no leak — all decided at compile time, with no garbage collector.

2. As a picture. Figure 27.1: a `String` value is a small stack record (pointer, length, capacity) owning a heap buffer; when the owning variable leaves scope, `drop` frees the buffer — one arrow, one free.

3. As numbers/contrast. The C snippet of §27.1 has two pointers to one buffer and calls `free` twice → abort. The Rust equivalent has one owner and one automatic drop → freed once; the aliasing that caused the double-free cannot even be written (see §27.3).

(Answer to the §27.1 quick check: once `t` aliases `s` and `s` is freed, using `t` is a **use-after-free** and the second `free(t)` is a **double-free**. The C compiler cannot catch either because its type system does not track ownership or lifetimes — a pointer carries no information about whether the memory it points to is still valid or who is responsible for freeing it. Rust's promise: memory safety — no use-after-free, double-free, or leak of owned memory — with no garbage collector, enforced at compile time.)

A String owner: a small stack record owning one heap buffer; drop frees it once at end of scope.



End of s's scope → Rust inserts drop(s) → heap buffer freed exactly once (no GC, deterministic).

One owner (rule 2) → one free (rule 3). A second owner would mean a second free — which moves forbid.

Figure 27.1: A String owner. The stack record (pointer, length, capacity) owns a heap buffer; when the owning variable leaves scope, the compiler-inserted drop frees the buffer exactly once. One owner means one free — the structural reason double-frees and leaks of owned memory cannot occur.

27.3 Moves: assignment transfers ownership

Here is where Rust departs sharply from C and even C++. In C, `char *t = s;` copies the pointer, giving two pointers to one buffer — the aliasing that caused §27.1's double-free. Rust cannot allow that, because two owners would mean two drops. So for a heap-owning type, assignment is a **move**: it transfers ownership to the destination and *invalidates the source*. After the move the source is no longer a valid owner, and any use of it is a compile error. Consider (Rust; not compiled here — diagnostic quoted from the cited documentation):

```
let s1 = String::from("hello");
let s2 = s1;           // MOVE: ownership transfers from s1 to s2; s1 is now invalid
println!("{s1}");      // error: use of moved value

error[E0382]: borrow of moved value: `s1`
  = note: move occurs because `s1` has type `String`, which does not implement the `Copy` trait
```

The program does not compile: after `let s2 = s1;`, the value has moved, so `s1` is dead and touching it is rejected at compile time with error E0382. This is the whole game. The C double-free of §27.1 is *unwritable* in Rust: the moment you make a second binding, the first stops being an owner, so there is never a moment when two owners could both free. The bug that compiled-and-crashed in C is caught before the program ever runs — and note that this is a *compile-time* rejection, costing nothing at run time. Two refinements complete the picture. If you genuinely want two independent copies, you ask for one explicitly with `.clone()`, which allocates a new buffer and deep-copies — now there are two owners of two *separate* buffers, each dropped once. And small values that live entirely on the stack — integers, `bool`, `char`, and tuples of them — implement the `Copy` trait, so `let y = x;` *copies* them instead of moving (there is no heap buffer to double-free, so copying is safe and cheap); `x` stays usable. The rule of thumb: values that own a resource move; plain stack data copies.

QUICK CHECK

In Rust, after `let s2 = s1;` where `s1` is a `String`, what is the status of `s1`, and what happens if you use it? Why does the same pattern work fine for `let y = x;` when `x` is an `i32`? How do you get two independent `Strings` from one? (Answer below the communicate box.)

27.4 Ownership across functions, and a first look at safety

Ownership follows the same move rule across function boundaries: passing a value to a function *moves* it into the function (unless the type is `Copy`), so the caller can no longer use it; returning a value moves ownership back out. This makes the flow of responsibility explicit in the signatures (Rust; not compiled here — behavior per the cited book):

```
fn consume(s: String) {           // s is moved IN; consume now owns it
    println!("{s}");
}                                 // s dropped here – buffer freed

let a = String::from("hi");
consume(a);                       // a is MOVED into consume
// println!("{a}");               // error[E0382]: a was moved, no longer valid
```

If a function needs a value but should not take ownership, moving it in and out by value is clumsy — which is exactly the motivation for *borrowing*, the subject of the next chapters: a way to lend access without transferring ownership. But ownership already buys something profound beyond memory safety. Because a moved-from value is dead, and because there is only ever one owner, ownership controls *aliasing* — and aliasing plus mutation is the root of both memory-corruption bugs and data races (Chapter 24). A taste: when you hand a value to another thread, Rust *moves* it there, so the original thread can no longer touch it — two threads cannot both own and mutate the same value by accident, which is the beginning of how Rust makes data races a compile error rather than the undefined behavior you had to prevent by hand in Chapter 24. The full story needs borrowing and the `Send/Sync` traits, developed in the chapters ahead. The point to hold now is that ownership is not merely a memory-management scheme; it is a discipline on aliasing, and aliasing is where the hazards of the whole first volume lived.

TRAP: EXPECTING ASSIGNMENT OR A FUNCTION CALL TO COPY (IT MOVES, AND THE SOURCE GOES DEAD)

Coming from C or a garbage-collected language, the deepest habit to unlearn is that `let b = a;` or `f(a)` leaves `a` usable. For an owning type in Rust it does not: the value *moves*, and `a` is invalidated — using it afterward is a compile error (E0382, “borrow of moved value”). Beginners read this as the compiler being obstinate and reach for the wrong fixes: sprinkling `.clone()` everywhere to silence it (correct but often needlessly copies heap data, throwing away the performance you came to Rust for), or fighting the borrow checker without understanding that the error is *telling them something true* — that they tried to use a value whose ownership they gave away. The right response is to decide, at each move, what you actually mean: if the callee should take over the value, move is correct and you simply must not use the source again; if the caller still needs it, you want to *borrow* it (lend a reference — the next chapters) rather than move it; and only if you truly need a second independent value should you `clone`. The move error is not noise; it is the ownership rule catching an aliasing decision you have to make consciously. Reaching for `clone` to make the compiler quiet is the analog of casting away a warning in C — it works and hides the point.

SAY IT LIKE A SYSTEMS ENGINEER

Try it first. A teammate porting C to Rust says: “This is ridiculous — I wrote `let b = a; use(a);` and it won't compile, saying `a` was moved. In C this just copies the pointer and works. I've been adding `.clone()` everywhere to get it to build. Isn't the borrow checker just getting in the way?” Cover the paragraph and respond in four or five sentences, drawing on §27.1–§27.3.

Model answer. “In C, `let b = a;` copying the pointer is exactly the bug that bites you later — now two pointers own one buffer, and whoever frees second gets a double-free or use-after-free, which is what that C program did at run time. Rust makes assignment of an owning type a *move*: ownership transfers to `b` and `a` goes dead, so there is never a moment when two owners could both free — the crash is turned into a compile error. So the checker isn't in the way; it's reporting that you used a value after giving its ownership away. `clone` silences it by making a real second buffer, which is sometimes right but often wasteful — if you only need to *look at* the value in `use()`, you want to borrow a reference, not clone. Decide per case: move if you're handing it off, borrow if the caller still needs it, clone only if you truly need two independent values.” Framing the move as preventing the C double-free, and clone-vs-borrow as a real decision, is the senior register.

(Answer to the §27.3 quick check: after `let s2 = s1;` with a `String`, `s1` has been **moved from** and is invalid; using it is a compile error (E0382). It works for `let y = x;` with an `i32` because integers implement the `Copy` trait — they live entirely on the stack with no heap resource to double-free, so assignment copies the bits and leaves `x` valid. To get two independent `Strings`, call `s1.clone()`, which allocates and deep-copies a separate buffer.)

27.5 What to carry forward

Ownership is the idea beneath everything in this part. The manual-memory bugs of Chapter 9 — use-after-free, double-free, leaks — all stem from C's having no notion of who is responsible for a piece of memory: our C snippet aliased one buffer with two pointers, compiled without a warning, and aborted at run time on the double-free. Rust replaces the discipline with three compiler-enforced rules: every value has exactly **one**

owner, and when the owner goes out of scope the value is **dropped** — cleanup tied to scope (RAII), deterministic, with no garbage collector, so owned memory is freed exactly once and never forgotten. The consequence is **move semantics**: assigning or passing an owning value *transfers* ownership and *invalidates the source*, so the two-owners situation that caused the C double-free literally cannot be written — using a moved-from value is a compile error (E0382), caught before the program runs and at zero runtime cost. Explicit `.clone()` makes a genuine second copy when you want one; small stack types (`Copy`) copy instead of move. And ownership is more than memory management: because a moved-from value is dead and there is only ever one owner, ownership is a discipline on *aliasing* — the root of both corruption and the data races of Chapter 24.

You have now seen the core move, but two things are still missing, and they are what the rest of the part builds. First, constantly moving values in and out of functions just to read them is untenable — you need a way to *use* a value without *owning* it. That is **borrowing**: taking a reference, shared or exclusive, under rules that keep aliasing safe — the next chapter. Second, borrowing raises the question of *how long* a reference may live relative to the value it points at, which is what **lifetimes** make precise — and together, ownership, borrowing, and lifetimes are how Rust's type system encodes, and checks at compile time, the very invariants you spent the C parts proving by hand: that memory is freed once and used only while valid, and that shared, mutable state is not raced upon. Everything you learned about the machine in Volume 1 still holds underneath — Rust compiles to the same instructions, the same cache and memory model — but the proofs have moved from your head into the compiler. We begin, in the next chapter, with borrowing.

CAN YOU... WITHOUT LOOKING?

- Name the three manual-memory bugs ownership prevents, and why C's type system cannot catch them.
- State the three ownership rules and what “dropped when the owner goes out of scope” means (RAII, no GC).
- Explain why single ownership means owned memory is freed exactly once and never leaked.
- Define a *move*: what happens to the source, and why using it afterward is a compile error.
- Distinguish *move*, *clone*, and *Copy*, and say which applies to a `String` vs an `i32`.
- Explain how ownership controls aliasing, and connect that to preventing data races (Chapter 24).

RETRIEVAL — CLOSED BOOK

Cover the chapter. Answer from memory; rate confidence first.

1. **R1.** What three manual-memory bugs does ownership prevent, and why can't the C compiler catch them? (§27.1)
2. **R2.** State the three ownership rules and what “drop” does at end of scope. (§27.2)
3. **R3.** What is a move — what happens to the source — and why is using a moved-from value a compile error? (§27.3)
4. **R4.** Distinguish move, clone, and Copy; which applies to String vs i32? (§27.3)
5. **R5.** How does ownership control aliasing, and why does that matter beyond memory management? (§27.4)

FURTHER READING

- **Klabnik & Nichols, *The Rust Programming Language*, 2nd ed. (No Starch Press, 2023), ch. 4 “Understanding Ownership” (also free at doc.rust-lang.org/book).** The primary source for the ownership rules, moves, clone, and Copy of §27.2–§27.3; the quoted rules and E0382 behavior are from here. Authors and edition verified via the publisher; the book was not compiled in this environment.
- **Jung, Jourdan, Krebbers & Dreyer, “RustBelt: Securing the Foundations of the Rust Programming Language,” *Proc. ACM on Programming Languages* 2(POPL), Article 66 (2018), DOI 10.1145/3158154.** The first machine-checked safety proof for a realistic subset of Rust — the formal basis for §27.1's “safety without a GC” claim. First author Ralf Jung; verified via dblp and the ACM Digital Library.
- ***The Rustonomicon* (official, doc.rust-lang.org/nomicon).** The advanced guide to ownership, moves, and the memory model underneath unsafe Rust — the “why” behind the rules of this chapter. Reference for later chapters; verify against the current edition.
- **Stroustrup, *The C++ Programming Language*, 4th ed. (Addison-Wesley, 2013), on RAII and move semantics.** The C++ lineage of the drop/RAII idea of §27.2 and of moves — useful context for where Rust's model came from.

Exercises

- 27.1** WARM-UP State the three ownership rules of Rust in your own words.
- 27.2** WARM-UP In one sentence, what does it mean that a value is “dropped” when its owner goes out of scope, and what C operation does drop correspond to for a heap-owning value?
- 27.3** WARM-UP After `let s2 = s1;` where `s1` is a `String`, is `s1` still usable? What does the compiler do if you try to use it?
- 27.4** CORE Take the C snippet of §27.1 (alias, free, use, free). (a) Identify each of the two bugs and the line it occurs on. (b) Explain why it compiles in C despite being wrong. (c) Explain precisely why the analogous Rust code cannot be written — what happens at the second binding that removes the possibility of a double-free?

27.5 **CORE** Explain how the three ownership rules together guarantee that owned heap memory is freed *exactly once* and is *never leaked*, and why this needs no garbage collector. Which rule prevents the double-free, and which prevents the leak?

27.6 **CORE** Distinguish *move*, *clone*, and *copy*. (a) For each, say what happens to the source value and whether a heap buffer is duplicated. (b) Why does `let y = x;` move when `x: String` but copy when `x: i32`? (c) When is reaching for `.clone()` the wrong instinct, and what should you often use instead?

27.7 **COST** Compare the runtime cost of Rust's ownership checks with a garbage collector. (a) At what time — compile or run — is the move/drop machinery resolved, and what runtime cost does the ownership checking itself add? (b) Contrast the memory-reclamation behavior of drop-at-end-of-scope with a tracing GC: which is deterministic, and what does that imply for pause times and peak memory? (c) Why is `.clone()` a real, measurable cost while a move is essentially free — what does each do at the machine level?

27.8 **FADED** *Worked, with the last step for you.* A Rust newcomer writes `let cfg = load_config();` then `process(cfg);` then `log(cfg);` and gets `error[E0382]: use of moved value: cfg` on the `log` line.

Step 1 (given): `cfg` has an owning type (say `Config`, not `Copy`).

Step 2 (given): `process(cfg)` takes its argument by value, so it *moves* `cfg` into `process`; after that call `cfg` is invalid.

Step 3 (given): `log(cfg)` then tries to use a moved-from value, which the compiler rejects.

Step 4 (you): Give two different correct fixes — one that lets both functions use the value without cloning, and one where cloning is genuinely appropriate — and say how you would decide between them and why blindly cloning is the wrong default.

27.9 **MIXED** For each, say whether assignment/passing *moves* or *copies*, and whether the source stays usable: (a) `let b = a;` with `a: String`; (b) `let b = a;` with `a: i32`; (c) passing a `Vec<u8>` by value to a function; (d) `let b = a;` with `a: (i32, bool)`. Justify each in a line, naming `Copy` where relevant.

27.10 **MIXED** Drawing on Chapters 9, 24, and 27, classify each as *true* or *false* and fix the false ones: (a) “Rust achieves memory safety with a garbage collector.” (b) “In Rust, `let s2 = s1;` for a `String` leaves both `s1` and `s2` valid.” (c) “A move error is caught at compile time and costs nothing at run time.” (d) “`.clone()` and a move do the same thing.” (e) “Ownership only matters for memory management, not for concurrency.”

27.11 **STRETCH** Ownership controls aliasing, and aliasing is the root of both memory-corruption bugs and data races. (a) Explain how the single-owner rule, applied across threads by *moving* a value into a thread, prevents two threads from both owning and mutating the same value. (b) Argue why “aliasing + mutation” is the common cause of both a use-after-free (Chapter 9) and a data race (Chapter 24), and how ownership attacks the shared root rather than the two symptoms separately. (c) What remaining capability does ownership-by-move alone *not* give you — that is, why do you still need borrowing to write useful programs, and what will borrowing have to guarantee to stay safe?

27.12 **LAB** On your own machine (a C toolchain suffices; a Rust toolchain is ideal if you can install one): (a) In C, compile and run the alias/free/use/free snippet of §27.1; confirm it compiles without warning and observe what happens at run time (a double-free abort or, with a sanitizer such as `-fsanitize=address`, a diagnosed use-after-free). (b) If you have `rustc`, write the corresponding `let s2 = s1; println!("{}", s1);` and confirm it fails to compile with `E0382`; if you do not, explain from this chapter exactly which line would be rejected and why. (c) Write down, for three variables in a small program, whether each assignment is a move or a copy, and justify from the type. Note your toolchain versions.

Chapter 28 — Borrowing and References

Before you start: Chapter 27 (ownership and moves — borrowing is the way to use a value *without* taking ownership, so the source stays valid), Chapter 9 (malloc/free and use-after-free — a reference that outlives or aliases badly is the same bug), Chapter 11 (iterator/pointer invalidation when a buffer is reallocated), Chapter 24 (a data race is two accesses to one location, one a write, unsynchronized — the borrow rule is exactly the aliasing discipline that rules it out). · Chapter 27 left us moving values in and out of functions just to read them; this chapter is how to lend access instead, and the one rule the compiler enforces to keep lending safe.

BEFORE YOU READ ON

From memory, reaching back: (1) Chapter 27 — after `let s2 = s1;` for a `String`, why is `s1` dead, and what problem would arise if it were *not*? (2) Chapter 11 — in C, you hold a pointer into a growable array and then the array is reallocated larger; what can happen to that pointer? (3) *Predict:* in Rust you take `let r = &v;` (a shared reference to a vector) and then, while `r` is still in use, call `v.push(x)` (which needs to mutate `v`). Does the compiler accept it? Write your guesses; §28.2 and §28.3 answer all three.

Ownership alone is not enough to write real programs. If the only way to give a function access to a value were to *move* it in (Chapter 27), every function that merely reads a value would have to move it back out, and two functions could never look at the same value at once. What you want is to **lend** access without transferring ownership — a **reference** — and the discipline that keeps lending safe is **borrowing**. This chapter is only about references and the borrow rule. First, why references are needed and what a shared reference (`&T`) is: read-only access, any number at once (§28.1–§28.2). Then the mutable reference (`&mut T`): exclusive write access, and why there may be only one, with no other reference alive alongside it (§28.3). Then the borrow checker's single rule — *one writer XOR many readers* — and how it turns the iterator-invalidation and data-race hazards of the earlier volumes into compile errors (§28.4). We close by carrying borrowing forward to the question it forces open: how long may a reference live relative to the value it points at — lifetimes (§28.5).

On the code in this part. The book's environment has a C toolchain (gcc 11.4.0) but *no Rust compiler installed*, so the C programs below were compiled and run and their output is real, whereas the Rust snippets were *not* compiled here; their behavior and the quoted compiler diagnostics are drawn from *The Rust Programming Language* (Klabnik & Nichols, 2nd ed., 2023) and the official Rust error index, and are labeled as such. Nothing about Rust's behavior below is invented — it is cited.

28.1 Why references: using without owning

Recall the awkwardness at the end of Chapter 27: a function that only wants to *read* a `String` — say, to compute its length — must take it by value, which *moves* it in and leaves the caller unable to use it afterward. To keep the value, the function would have to return it back out, threading ownership through every call. That is untenable. A **reference**

solves it: `&s` creates a value that *points at* `s` without owning it, so the caller keeps ownership and the function borrows access for the duration of the call (Rust; not compiled here — behavior per the cited book):

```
fn length(s: &String) -> usize {    // s borrows: it does NOT own the String
    s.len()
}                                     // s goes out of scope, but nothing is dropped –
                                     // the reference did not own the buffer

let owner = String::from("hello");
let n = length(&owner);              // lend a reference; ownership stays with `owner`
println!("{owner} has len {n}");     // `owner` is STILL valid – it was never moved
```

The `&` in `&owner` is “take a reference to”; the `&String` in the signature is “a reference to a `String`.” The reference is a borrow: the callee may use the value through it, but it does not own the value and so does not drop it when the reference goes out of scope — only the owner drops. Because no move happened, `owner` is still usable after the call. This is the everyday shape of Rust code: pass owning values by reference to read them, and reserve moves for when you genuinely mean to hand off ownership. But a reference is a pointer into someone else's value, and Chapters 9 and 11 are a catalogue of how pointers into other people's memory go wrong — dangling after a free, aliasing a buffer that is then mutated out from under you. The rest of the chapter is the two kinds of reference and the one rule that makes them safe where the C pointer was not.

QUICK CHECK

Why does passing `&owner` to `length` leave `owner` usable afterward, when passing `owner` by value would not? When the reference parameter `s` goes out of scope at the end of `length`, is the `String`'s buffer freed? Why or why not? (Answer after the next section.)

28.2 Shared references: many readers, no writers

The reference of §28.1 is a **shared reference**, written `&T`. Its defining property: you may have *any number* of shared references to a value at the same time, and through a shared reference you may *read* the value but not mutate it. That is safe for the same reason many threads reading (but not writing) shared memory is safe (Chapter 24): if nobody writes, no reader can observe a torn or stale value, and no two accesses conflict. Shared references are `Copy` — cheap to duplicate, since they are just pointers — so handing the same value to several readers is free (Rust; not compiled here — behavior per the cited book):

```
let v = vec![1, 2, 3];
let a = &v;           // shared reference #1
let b = &v;           // shared reference #2 – fine, many readers allowed
println!("{a} {b}", a.len(), b.len()); // both read v; no conflict
```

Two immutable views of one vector coexist without trouble because neither can change it. The read-only-ness is not a convention you must remember — it is in the type: a `&T` simply has no method that mutates, so an attempt to write through one does not compile.

This is the “many readers” half of Rust's borrowing rule, and it is exactly as permissive as it can safely be: sharing is unlimited precisely while nothing mutates. The moment you need to *change* the value, you need a different, stricter kind of reference — because mutation plus aliasing is where every hazard lived.

THE SAME IDEA THREE WAYS: ONE WRITER XOR MANY READERS

1. In words. At any point a value may be borrowed either by any number of shared references (`&T`, read-only) *or* by exactly one mutable reference (`&mut T`, read-write) — never both, never two mutable ones. Sharing is safe while nothing writes; writing is safe only when nothing else can observe the value meanwhile.

2. As a picture. Figure 28.1: a value with several thin arrows labelled `&` pointing in (all readers, allowed together), versus the same value with one bold arrow labelled `&mut` and no other arrow permitted (the exclusive writer).

3. As numbers/contrast. The C program of §28.4 holds one pointer into a buffer and then reallocates the buffer → the pointer dangles → heap-use-after-free (ASan-confirmed). The Rust equivalent cannot compile: the reallocation needs `&mut`, which the outstanding shared reference forbids (E0502) → the bug is rejected before it runs.

(Answer to the §28.1 quick check: passing `&owner` lends a reference and does not move the value, so `owner` keeps ownership and stays valid; passing by value would move the `String` into the function and invalidate `owner`. When the reference parameter `s` goes out of scope, nothing is freed, because a reference does not own the buffer — only the owner drops it, and the owner is still `owner` back in the caller.)

Borrowing rule: many shared readers together, OR one exclusive writer alone — never both.




Figure 28.1: The borrowing rule. Top: a value may have any number of shared references (`&T`, read-only) at once. Bottom: or exactly one mutable reference (`&mut T`, read-write), with no other reference of either kind alive alongside it. Never both. This is aliasing XOR mutation, enforced by the type system.

28.3 Mutable references: exactly one writer, no aliases

To mutate a value through a reference you take a **mutable reference**, `&mut T`. Its rule is the opposite of the shared reference's permissiveness: while a `&mut T` to a value exists, it must be the *only* live reference to that value — no second mutable reference, and no shared reference either. A mutable borrow is *exclusive*. The reason is precisely the aliasing-plus-mutation hazard: if two references could reach one value and at least one

could write, you would have, in a single thread, the same conflict a data race is across threads (Chapter 24) — a write that another accessor can observe half-done or unexpectedly, or that invalidates a pointer the other holds. Rust rules it out by construction (Rust; not compiled here — diagnostics quoted from the official Rust error index):

```
let mut v = vec![1, 2, 3];
let m = &mut v;           // exclusive mutable borrow of v
let n = &mut v;           // ERROR: a second mutable borrow while m is live
m.push(4);

error[E0499]: cannot borrow `v` as mutable more than once at a time
```

```
let mut v = vec![1, 2, 3];
let r = &v;                // shared (read) borrow
let m = &mut v;            // ERROR: mutable borrow while a shared borrow is live
println!("{}", r[0]);

error[E0502]: cannot borrow `v` as mutable because it is also borrowed as immutable
```

The first is rejected with E0499 — “*cannot borrow as mutable more than once at a time*” — because two mutable references would be two writers to one value. The second is rejected with E0502 — “*cannot borrow as mutable because it is also borrowed as immutable*” — because a writer alongside a reader is exactly aliasing-plus-mutation. Both diagnostics are from the official Rust error index. Note what the exclusivity buys: through a `&mut T` the compiler *knows* nothing else can see the value, so mutation through it is as safe and as optimizable as if the value were not shared at all — the reference is effectively a compiler-checked exclusive lock, but one with zero runtime cost, resolved entirely at compile time. The borrow of a value ends when the reference is last used (this is “non-lexical lifetimes”: the borrow’s span is how long the reference is actually used, not the enclosing braces), so you can take a shared borrow, finish with it, and then take a mutable one — just never overlapping.

QUICK CHECK

State the two-part borrowing rule in one sentence. Which error code fires when you take two `&mut` to one value, and which when you take a `&mut` while a `&` is still live? Why is a `&mut T` being exclusive the single-threaded analog of preventing a data race? (Answer below the communicate box.)

28.4 The borrow checker, and how it kills iterator invalidation and data races

The two rules — many shared XOR one exclusive, and (next chapter) references must not outlive their value — are enforced by the **borrow checker**, a compile-time pass. To see it earning its keep, take a bug from Chapter 11 that C compiles without complaint. Here is real C: hold a pointer into a heap buffer, then grow the buffer with `realloc`, which is free

to move it — after which the old pointer dangles. Compiled and run in this environment (gcc 11.4.0):

```
int *v = malloc(4 * sizeof(int));
for (int i = 0; i < 4; i++) v[i] = i * 10;
int *p = &v[0];                /* alias into the buffer */
v = realloc(v, 4096 * sizeof(int)); /* may move the block, invalidating p */
printf("via p: %d\n", *p);        /* use-after-free if the block moved */
free(v);

$ gcc -O2 -Wall borrow.c -o borrow          # compiles, NO warning
$ gcc -O1 -fsanitize=address borrow.c -o borrow && ./borrow
==ERROR: AddressSanitizer: heap-use-after-free on address 0x502000000010
... in main
SUMMARY: AddressSanitizer: heap-use-after-free
```

The plain build gives no warning at all; only AddressSanitizer, at run time, catches the **heap-use-after-free**. This is the iterator-invalidation hazard of Chapter 11 in its rawest form — a live pointer into a container that is then reallocated. Now the Rust equivalent. A `Vec`'s `push` may reallocate, so its signature takes `&mut self`; holding a reference into the vector means holding a `&` borrow of it. The two cannot overlap (Rust; not compiled here — diagnostic per the official Rust error index):

```
let mut v = vec![1, 2, 3];
let p = &v[0];           // shared borrow: a reference INTO the vector
v.push(4);                // push needs &mut v – conflicts with the live borrow p
println!("{}", p);        // use of p after the (attempted) reallocation

error[E0502]: cannot borrow `v` as mutable because it is also borrowed as immutable
```

The exact bug the C program hid until run time — a reference into a buffer that is then reallocated — is a compile error in Rust, `E0502`, because `push`'s mutable borrow cannot coexist with the shared borrow `p`. The borrow checker does not need to know that `push` reallocates; the rule “no mutable borrow while a shared borrow is live” is enough. And this is the same rule that forbids data races: a data race (Chapter 24) is two accesses to one location, at least one a write, without synchronization — which is aliasing (two accessors) plus mutation (a write). Rust's borrow rule bans exactly that combination for a single value, and, extended across threads by the `Send/Sync` traits (a later chapter), it is why safe Rust cannot have a data race. The undefined behavior you proved absent by hand in Chapters 11 and 24 is, here, a property the compiler checks.

TRAP: MUTATING A COLLECTION WHILE A REFERENCE INTO IT IS ALIVE (THE “FIGHTING THE BORROW CHECKER” REFLEX)

The most common early-Rust frustration is writing a loop that reads from a collection and mutates it in the same breath — `for x in &v { if cond(x) { v.push(...) } }` — and being stopped by E0502: *cannot borrow v as mutable because it is also borrowed as immutable*. Beginners read this as the compiler being pedantic about a program that “obviously works,” and reach for the wrong escapes: cloning the whole collection to get an owned copy to iterate, or wrapping everything in `RefCell` to push the check to run time, or dropping into `unsafe`. But the error is reporting a *real* hazard — the exact iterator-invalidation bug the C program above hit at run time, where mutating the container can reallocate its buffer and dangle the reference you are iterating through. The right responses match what you actually mean: if you only need to *read* while deciding what to change, collect the decisions first (e.g. gather the indices or new items into a separate vector) and then apply the mutations after the read-borrow ends; if you need to mutate in place without growing, iterate with `&mut (iter_mut)`, which is a single exclusive borrow and is allowed; if you genuinely need shared mutable access with runtime-checked rules, `RefCell` is a deliberate, documented choice, not a way to silence the compiler. Fighting the borrow checker by cloning or reaching for `unsafe` to make the message go away is the analog of casting away a warning in C — it works and discards the thing the warning was protecting you from.

SAY IT LIKE A SYSTEMS ENGINEER

Try it first. A teammate says: “The borrow checker is being ridiculous. I’m iterating over a vector and pushing to it inside the loop, and it won’t compile — E0502, can’t borrow as mutable because it’s borrowed as immutable. This is a five-line loop that would just work in C++ or Python. Why is Rust making me clone the whole thing?” Cover the paragraph and respond in four or five sentences, drawing on §28.3–§28.4.

Model answer. “That loop in C++ is the classic iterator-invalidation bug: pushing can reallocate the vector’s buffer, and the reference you’re iterating through then dangles — in C the same pattern is a silent heap-use-after-free that only a sanitizer catches at run time. Rust’s rule is that a mutable borrow (push needs one) can’t coexist with the shared borrow your loop is holding, which is exactly why it flags it. So the checker isn’t being pedantic; it caught the reallocation hazard before the program ran. Don’t clone the whole vector to dodge it — that pays for a full copy to hide a real bug. Collect the items you want to add into a separate vector during the read pass, then push them after the loop, or use `iter_mut` if you’re mutating in place; the fix is to separate the read from the write, which is what the rule is telling you to do.” Framing E0502 as the compile-time form of the C use-after-free, and offering the collect-then-mutate fix instead of a clone, is the senior register.

(Answer to the §28.3 quick check: the rule is “at any time a value may have either any number of shared references `&T` or exactly one mutable reference `&mut T`, never both.” Two `&mut` to one value is E0499 (“cannot borrow as mutable more than once at a time”); a `&mut` while a `&` is live is E0502 (“cannot borrow as mutable because it is also borrowed as immutable”). A `&mut T` being exclusive is the single-threaded analog of preventing a data race because a data race is aliasing plus a write; exclusivity guarantees no other accessor can observe the value while it is being written.)

28.5 What to carry forward

Borrowing is how you use a value without owning it. A **reference** points at a value it does not own, so lending one leaves the owner intact and drops nothing when the reference's scope ends. There are two kinds, and the whole of borrowing is one rule dividing them: a value may be borrowed by *any number of shared references* (`&T`, read-only) or by *exactly one mutable reference* (`&mut T`, read-write), never both at once and never two mutable ones. That is **aliasing XOR mutation**: unlimited sharing is safe precisely while nothing writes, and writing is safe precisely when nothing else can see the value meanwhile. The **borrow checker** enforces this at compile time, at zero runtime cost, and in doing so it turns concrete hazards from the earlier volumes into compile errors: the iterator-invalidation use-after-free that the C program of §28.4 hid until a sanitizer caught it is, in Rust, `E0502` at compile time (push's mutable borrow cannot coexist with a live shared borrow into the vector), and two writers to one value are `E0499`. The same “no aliasing with mutation” rule, lifted across threads later, is why safe Rust has no data races — the Chapter 24 hazard made structural.

One question is now unavoidable, and it is the whole of the next chapter. A reference is a pointer into another value; the second borrow rule is that a reference must never be used after the value it points at is gone — a dangling reference (the C return-a-pointer-to-a-local bug) must be impossible. But how does the compiler *know*, at the point you use a reference, that the value behind it is still alive — especially when references are passed into and out of functions, stored in structs, and returned? The answer is **lifetimes**: a way to name and relate the spans over which references are valid, so the borrow checker can prove that no reference outlives its referent. Everything the machine does underneath is unchanged — a Rust reference compiles to exactly the pointer C would use, with no extra runtime data — but the compiler now carries a proof, expressed in lifetimes, that the pointer is always valid when dereferenced. That proof, and the small annotations it sometimes needs from you, is where we go next.

CAN YOU... WITHOUT LOOKING?

- Say what a reference is, and why lending one leaves the owner valid where moving would not.
- State the two-part borrowing rule (many shared XOR one exclusive) and why each half is safe.
- Explain why a `&mut T` must be the only live reference, in terms of aliasing plus mutation.
- Name the error codes for two mutable borrows (`E0499`) and a mutable-while-shared borrow (`E0502`).
- Show how the borrow rule turns the C iterator-invalidation use-after-free into a compile error.
- Connect the borrow rule to data-race prevention (aliasing XOR mutation, Chapter 24).

RETRIEVAL — CLOSED BOOK

Cover the chapter. Answer from memory; rate confidence first.

1. **R1.** What is a reference, and why does passing `&owner` leave `owner` usable while passing it by value does not? (§28.1)
2. **R2.** How many shared references (`&T`) to one value may exist at once, and what may you do through them? (§28.2)
3. **R3.** State the exclusivity rule for `&mut T`, and give the two error codes it is enforced by. (§28.3)
4. **R4.** Take the C realloc/dangling-pointer bug; explain why the analogous Rust code is E0502 at compile time. (§28.4)
5. **R5.** Why is “aliasing XOR mutation” the same rule that prevents data races? (§28.4)

FURTHER READING

- **Klabnik & Nichols, *The Rust Programming Language*, 2nd ed. (No Starch Press, 2023), ch. 4.2 “References and Borrowing” (also free at doc.rust-lang.org/book).** The primary source for shared vs. mutable references and the borrowing rule of §28.2–§28.3. Authors and edition verified via the publisher; the book was not compiled in this environment.
- **The Rust Error Index, entries E0499, E0502, and E0505 (doc.rust-lang.org/error_codes).** The official one-line meanings quoted in §28.3–§28.4: E0499 “borrowed as mutable more than once,” E0502 “borrowed as mutable because also borrowed as immutable,” E0505 “a value was moved out while it was still borrowed.” Each verified against the current error index.
- **Jung, Jourdan, Krebbers & Dreyer, “RustBelt: Securing the Foundations of the Rust Programming Language,” *Proc. ACM on Programming Languages* 2(POPL), Article 66 (2018), DOI 10.1145/3158154.** The machine-checked proof that the borrowing discipline of this chapter is sound. First author Ralf Jung; verified via the ACM Digital Library.
- **The Rust reference and the “non-lexical lifetimes” RFC 2094 (rust-lang.github.io/rfcs).** Why a borrow ends at last use rather than at the closing brace — the refinement noted in §28.3. Reference for the precise model; verify against the current edition.

Exercises

- 28.1** WARM-UP In one sentence each, define a shared reference (`&T`) and a mutable reference (`&mut T`), saying what you may do through each and how many of each may exist at once.
- 28.2** WARM-UP Why does `fn length(s: &String)` leave the caller's `String` usable after the call, whereas `fn length(s: String)` would not?
- 28.3** WARM-UP State the two-part borrowing rule in one sentence, and name the two situations it forbids together with their error codes.

28.4 **CORE** Take the C snippet of §28.4 (pointer into a buffer, `realloc`, use the old pointer). (a) Identify the bug and the line it occurs on. (b) Explain why the plain `gcc -O2 -Wall` build gives no warning, and what AddressSanitizer reports at run time. (c) Explain precisely why the analogous Rust code is `E0502` at compile time — which borrow conflicts with which, and why the checker does not need to know that `push` reallocates.

28.5 **CORE** Explain why a mutable reference must be exclusive. (a) What single-threaded hazard would two references to one value, at least one mutable, reproduce? (b) Connect this to the definition of a data race from Chapter 24. (c) What does the exclusivity of `&mut T` let the compiler assume, and why is that good for optimization?

28.6 **CORE** For each, say whether it compiles, and if not, which error code fires and why: (a) `let a = &v; let b = &v;` (both shared); (b) `let m = &mut v; let n = &mut v;;` (c) `let r = &v; let m = &mut v;;` (d) a shared borrow used and finished, then a mutable borrow taken afterward (non-overlapping).

28.7 **COST** References are “zero-cost.” (a) At the machine level, what is a `&T` — how much runtime data does it carry beyond the pointer C would use? (b) At what time is the borrow checking done, and what runtime cost does it add? (c) Contrast this with `RefCell`, which enforces the same borrow rule at run time: what does `RefCell` cost that a plain `&mut` does not, and when is paying it justified?

28.8 **FADED** *Worked, with the last step for you.* A newcomer writes a loop that scans a `Vec<i32>` and pushes each value's double back onto the same vector, and gets `error[E0502]: cannot borrow `v` as mutable because it is also borrowed as immutable`.

Step 1 (given): `for x in &v { ... }` holds a shared borrow of `v` for the whole loop body.

Step 2 (given): `v.push(2 * x)` inside the body needs a mutable borrow of `v`, which cannot coexist with the shared borrow.

Step 3 (given): this is the iterator-invalidation hazard: a `push` could reallocate `v` and dangle the iterator.

Step 4 (you): Give two correct fixes — one that separates the read pass from the write pass without cloning the whole vector, and one using `iter_mut` if the mutation were in place — and say why cloning the entire vector to dodge the error is the wrong default.

28.9 **MIXED** For each, state whether it is allowed by the borrow rule and justify in a line: (a) three `&v` at once; (b) one `&mut v` and one `&v` at once; (c) two `&mut v` at once; (d) a `&mut v` after all shared borrows of `v` have ended.

28.10 **MIXED** Classify each as *true* or *false* and fix the false ones: (a) “You may have many mutable references to a value as long as they don't write at the same time.” (b) “A shared reference lets you mutate the value it points at.” (c) “The borrow checker adds runtime overhead to every dereference.” (d) “Holding a reference into a vector and then pushing to it is a compile error.” (e) “A reference drops (frees) the value when it goes out of scope.”

28.11 **STRETCH** The borrow rule is “aliasing XOR mutation.” (a) Show that both a use-after-free through an aliased pointer (Chapter 9/11) and a data race (Chapter 24) are instances of “two paths to one value, at least one writing,” and how the single rule attacks their shared root. (b) Explain why unlimited *shared* references are nonetheless safe, and unlimited access is only a problem once mutation enters. (c) The rule as stated is for a single value in one thread; name what still has to be added to extend “no data races” across threads, and why the borrow rule alone is not yet enough there.

28.12 **LAB** On your own machine (a C toolchain suffices; a Rust toolchain is ideal): (a) Compile and run the §28.4 C snippet (pointer into a buffer, `realloc` to a much larger size, use the old pointer). Confirm it builds without warning under `-O2 -Wall`, then rebuild with `-fsanitize=address` and observe the heap-use-after-free. (b) If you have `rustc`, write the `let p = &v[0]; v.push(4); println!("{p}");` version and confirm it fails with `E0502`; if not, name the exact conflicting borrows from this chapter. (c) Write a version that compiles by ending the shared borrow before the `push`, and explain why it is now allowed. Note your toolchain versions.

Chapter 29 — Lifetimes

Before you start: Chapter 28 (borrowing — a lifetime is how the compiler proves the second borrow rule, that a reference never outlives its value), Chapter 27 (ownership and drop — a value is dropped at end of scope, so a reference to it must not be used past that point), Chapter 8 (the stack: a local's storage is reclaimed when its scope ends), Chapter 9 (dangling pointers and use-after-free — exactly what lifetimes rule out). · Chapter 28 left one borrow rule unproven: a reference must never be used after the value it points at is gone. This chapter is how the compiler proves it — by tracking, and sometimes making you name, the spans over which references are valid.

BEFORE YOU READ ON

From memory, reaching back: (1) Chapter 8 — when is a stack local's storage reclaimed, and what becomes of a pointer to it afterward? (2) Chapter 28 — a reference does not own its value; who is responsible for dropping the value, and when? (3) *Predict*: in C you write a function that returns `&local` (the address of one of its own stack variables) and the caller dereferences it. Does `gcc -Wall` warn? In Rust, a function that returns a reference to one of its own locals — will the compiler accept it? Write your guesses; §29.1 and §29.3 answer all three.

Chapter 28 gave one borrow rule in full — aliasing XOR mutation — and named a second only in passing: a reference must never be used after the value it refers to is gone. That second rule is what this chapter is about, and the machinery that enforces it is **lifetimes**. A lifetime is not a runtime thing and it changes nothing about the code the machine runs; it is a *compile-time* name for a region of the program over which a reference is valid, and the borrow checker uses lifetimes to prove that every reference is dereferenced only while its referent is still alive. First, the hazard: the dangling reference — the C return-a-pointer-to-a-local bug — and why the compiler must forbid it (§29.1). Then how references usually need no annotation, and when they do: naming a lifetime `'a` in a function signature to relate an output reference to its inputs (§29.2). Then what the borrow checker actually checks, and the `E0597` error it produces when a value does not live long enough (§29.3). Then lifetime **elision** — why most signatures omit lifetimes — and structs that hold references, where a lifetime is mandatory (`E0106`) (§29.4). We close by carrying lifetimes forward: ownership, borrowing, and lifetimes together are the whole memory-safety story, and the next chapters build reusable abstractions on top of it (§29.5).

On the code in this part. The book's environment has a C toolchain (gcc 11.4.0) but *no Rust compiler installed*, so the C programs below were compiled and run and their output is real, whereas the Rust snippets were *not* compiled here; their behavior and the quoted compiler diagnostics are drawn from *The Rust Programming Language* (Klabnik & Nichols, 2nd ed., 2023) and the official Rust error index, and are labeled as such. Nothing about Rust's behavior below is invented — it is cited.

29.1 The dangling reference, and why it must be forbidden

A reference is a pointer into a value it does not own. The value is dropped by its owner at end of scope (Chapters 27, 28); if a reference is used after that point, it points at reclaimed storage — a **dangling reference**, the use-after-free of Chapter 9 in reference form. The archetype is returning a pointer to a stack local. Here is real C; the trivial form is caught by a best-effort warning, but that warning is a heuristic, not a guarantee. Compiled in this environment (gcc 11.4.0):

```
int *make(void) {
    int local = 42;
    return &local;      /* address of a stack local; reclaimed when make returns */
}
/* caller: */ int *r = make(); printf("%d\n", *r);    /* dangling: use-after-scope */

$ gcc -O2 -Wall dangle.c -o dangle
dangle.c: In function 'make':
dangle.c:3:12: warning: function returns address of local variable [-Wreturn-local-addr]
```

gcc catches this direct form. But launder the same pointer through a struct and the heuristic misses it — the program compiles clean and is only caught at run time by a sanitizer (gcc 11.4.0, this environment):

```
struct Ref { int *p; };
struct Ref borrow_local(void) {
    int local = 42;
    struct Ref r; r.p = &local;    /* store address of a local in a struct */
    return r;                      /* r.p now dangles */
}
/* caller: */ struct Ref r = borrow_local(); printf("%d\n", *r.p);

$ gcc -O2 -Wall dangle2.c -o dangle2      # compiles, NO warning
$ gcc -O1 -fsanitize=address dangle2.c -o d && ./d
==ERROR: AddressSanitizer: stack-use-after-scope ... in main
```

The point is not that C never warns — it is that C's warnings are a best-effort heuristic that follows only simple data flow and gives up as soon as the pointer passes through a struct, a second function, or a container. The general problem — does this reference outlive the value it points at? — is not something C's type system tracks at all. Rust makes it a rule the type system enforces in *every* case: a reference may not be used after its referent is dropped, no exceptions, checked at compile time. To do that in the general case — references flowing into and out of functions, stored in structs, returned — the compiler needs a way to *name* and *compare* the spans over which references and values are valid. That name is a lifetime.

QUICK CHECK

Why is returning `&local` from a function a bug — what has happened to `local`'s storage by the time the caller dereferences the result? Why does `gcc -Wall` catch the direct form but not the version laundered through a struct? State the general rule Rust enforces that the C heuristic only approximates. (Answer after the next section.)

29.2 Lifetimes named: relating an output reference to its inputs

Most of the time you never write a lifetime — the compiler infers the valid regions on its own (§29.4). You must name one when a function returns a reference and the compiler cannot tell, from the signature alone, *which* input the returned reference borrows from — because that determines how long the result stays valid. The canonical case is a function returning whichever of two string slices is longer (Rust; not compiled here — behavior per the cited book):

```
fn longest<'a>(x: &'a str, y: &'a str) -> &'a str {
    if x.len() > y.len() { x } else { y }
}
```

The `<'a>` declares a **lifetime parameter** named `'a` (lifetimes are written with a leading apostrophe and are usually short and lowercase). The signature now says: for some region `'a`, both inputs are references valid for at least `'a`, and the returned reference is valid for `'a` too. This is a *constraint*, not a runtime value — it does not change what the function does or make anything live longer; it tells the borrow checker how the output's validity relates to the inputs'. The payoff is at the call sites: the compiler now knows the returned reference cannot outlive the shorter-lived of the two arguments, and it will reject any caller that tries to use the result after either input has gone. Without the annotation the compiler could not know whether the result borrows from `x` or `y`, so it could not check the callers — which is exactly why the annotation is required here. Lifetime parameters, then, are how a function's signature carries a promise about the validity of the references it returns, so that borrow-checking composes across function boundaries.

THE SAME IDEA THREE WAYS: A REFERENCE MAY NOT OUTLIVE ITS REFERENT

1. In words. Every reference is valid only for a region — its lifetime — bounded by the scope of the value it points at. The borrow checker requires that a reference is used only within that region; using it after the referent is dropped is rejected. A lifetime parameter like `'a` lets a signature state how an output reference's region relates to its inputs'.

2. As a picture. Figure 29.1: a value's lifetime drawn as a bar from its creation to its drop; a reference's use drawn as a shorter bar that must fit *inside* the value's bar. A reference bar that extends past the value's drop is the error.


3. As numbers/contrast. The C function of §29.1 returns a pointer whose referent's storage is reclaimed the instant the function returns — the reference bar extends past the value bar → use-after-scope (ASan-confirmed). The Rust equivalent is E0597 at compile time: the borrowed value does not live long enough, so the reference cannot be returned.

(Answer to the §29.1 quick check: `local` lives on the function's stack frame, which is reclaimed the moment the function returns, so by the time the caller dereferences the returned pointer the storage is gone — a use-after-scope. gcc's `-Wreturn-local-addr` follows simple direct data flow and catches `return &local;`, but it gives up once the address is stored in a struct and returned indirectly, so that version compiles clean. Rust's rule,

enforced in all cases at compile time, is that a reference may never be used after the value it refers to has been dropped.)

A reference's use must fit inside its referent's lifetime. Extending past the drop is the error.

OK: reference used only while value is alive



BAD: reference used after value dropped → E0597




Figure 29.1: Lifetimes as fitting. A reference is valid only within its referent's lifetime (top, allowed). A reference used after the value is dropped (bottom) is a dangling reference — the C use-after-scope of §29.1, and in Rust a compile error, E0597, “borrowed value does not live long enough.”

29.3 What the borrow checker proves, and E0597

With lifetimes as its vocabulary, the borrow checker's second rule is a containment check: for every use of a reference, the region over which the reference is used must be contained in the region over which its referent is alive. When it is not, the compiler reports that the borrowed value does not live long enough. Here is the dangling-reference bug of §29.1 in Rust — a function trying to return a reference to its own local (Rust; not compiled here — diagnostic quoted from the official Rust error index):

```
fn dangle<'a>() -> &'a str {
    let s = String::from("hello"); // s is a local; it will be dropped at the `}`
    &s                               // return a reference to s
}                                   // s is dropped here – the reference would dangle

error[E0597]: `s` does not live long enough
  = note: values in a scope are dropped in the opposite order they are defined
```

The compiler rejects it with E0597 — *the borrowed value does not live long enough*, per the official Rust error index — because `s` is dropped at the closing brace, so the returned reference would point at freed storage; its required lifetime (as long as the caller keeps it) exceeds `s`'s actual lifetime (the function body). This is the C use-after-scope of §29.1, but caught at compile time in every case, not by a heuristic that a struct can defeat. The fix is never to make the reference live longer — you cannot extend a dropped value — but to return an *owned* value (return `String`, moving ownership out, Chapter 27) or to borrow from something that outlives the function (an input reference, §29.2). Note again that lifetimes are inert at run time: they are erased before code generation, so a correct program pays nothing for them, and the reference compiles to the same bare pointer C would use. The entire cost of the guarantee is borne by the compiler, once.

QUICK CHECK

What does E0597 report, and what condition on regions is the borrow checker verifying when it fires? Why can't you fix a “does not live long enough” error by making the reference live longer, and what are the two real fixes? At run time, what does a lifetime cost? (Answer below the communicate box.)

29.4 Elision, and structs that hold references

If every reference needed an explicit lifetime, Rust would be unbearable. It does not, because of **lifetime elision**: a set of deterministic rules by which the compiler fills in the common cases so you write no annotation at all. The three elision rules (Klabnik & Nichols, 2023, ch. 10) are: (1) each elided input reference gets its own distinct lifetime; (2) if there is exactly one input lifetime, it is assigned to all elided output lifetimes; and (3) if there are multiple input lifetimes but one of them is `&self` or `&mut self` (a method), the lifetime of `self` is assigned to all elided outputs. When these rules determine every lifetime, you write none; the `longest` function of §29.2 needed an explicit `'a` only because it has two input references and neither is `self`, so rule 2 and rule 3 do not apply and the compiler cannot guess which input the output borrows. Elision is not a special case in the type system — it is pure syntactic sugar for lifetimes you could always write out.

One place a lifetime is never optional: a struct that *holds* a reference. Such a struct is only valid as long as the value it borrows, and that relationship must be named. Omitting it is E0106 (Rust; not compiled here — diagnostic per the official Rust error index):

```
struct Excerpt { part: &str }           // ERROR: reference field with no lifetime
error[E0106]: missing lifetime specifier

struct Excerpt<'a> { part: &'a str } // OK: the struct borrows for 'a and may not
//      outlive the &str it holds
```

The first is rejected with E0106 — *missing lifetime specifier* — because a struct with a reference field carries an implicit promise (“I do not outlive what I borrow”) that must be written into the type as `<'a>`. Once annotated, the borrow checker enforces that no `Excerpt` value is used after the `str` it points into has been dropped — the same containment rule as §29.3, now for a reference living inside a struct. This is why data structures that own their data (holding `String`, `Vec`) are simpler than ones that borrow it (holding `&str`, `&[T]`): owning avoids lifetime parameters entirely, at the cost of an allocation; borrowing avoids the allocation, at the cost of threading a lifetime through the type. That trade — own vs. borrow — is a recurring design decision, and lifetimes are the language in which the borrowing side of it is expressed.

TRAP: READING A LIFETIME AS “HOW LONG THE VALUE LIVES” AND TRYING TO MAKE REFERENCES LIVE LONGER

The word “lifetime” misleads newcomers into thinking `'a` is a duration you set — that annotating a reference with a longer lifetime makes the underlying value live longer, the way a smart pointer or a garbage collector would keep it alive. It does not. A lifetime is purely *descriptive*: it names a region that the reference is valid for, bounded by facts already fixed by ownership — when the value is created and when its owner drops it. Writing `'a` does not extend anything; it only lets the compiler *check* a relationship. So when you hit E0597, “`s` does not live long enough,” the wrong instinct is to search for a lifetime annotation that will make the error disappear — adding `'static`, or inventing longer-named lifetimes — because no annotation can make a value outlive its owner's scope. The error is telling you a structural truth: you are trying to use a reference past the life of what it points at. The real fixes are to change ownership, not the annotation: return an owned value (move it out, Chapter 27), borrow from something that genuinely outlives the use (an input reference, or a value in an outer scope), or restructure so the reference is never needed past the value's death. Reaching for `'static` to silence E0597 is the analog of casting a pointer to make a type error vanish — it either does not compile anyway or forces the value to be truly static, which is rarely what you meant.

SAY IT LIKE A SYSTEMS ENGINEER

Try it first. A teammate says: “I have a function that builds a string and returns a `&str` to it, and I keep getting E0597, `does not live long enough`. I tried annotating it `'a`, then `'static`, and it still won't compile. Lifetimes are just confusing syntax — why can't I just tell it the reference lives as long as I need?” Cover the paragraph and respond in four or five sentences, drawing on §29.3–§29.4.

Model answer. “The annotation can't fix it because a lifetime doesn't *make* anything live longer — it only names how long a value already lives, which ownership fixed: your local string is dropped when the function returns, so any reference to it would dangle, which is exactly the C ‘return a pointer to a local’ bug. That's what E0597 is reporting, and no lifetime name, not even `'static`, can extend a value past its owner's scope. The fix is about ownership, not syntax: return the `String` by value so ownership moves out to the caller, or borrow from an input that outlives the call. Adding `'static` only makes it compile if the data really is static, which a freshly built string isn't. So don't hunt for the magic annotation — decide who should own the result.” Framing the error as the C return-a-local bug and the fix as an ownership decision, not an annotation search, is the senior register.

(Answer to the §29.3 quick check: E0597 reports that a borrowed value does not live long enough; the borrow checker is verifying that the region over which a reference is used is contained within the region its referent is alive. You cannot fix it by making the reference “live longer” because a lifetime only describes existing facts — it cannot extend a value past its owner's drop; the two real fixes are to return an owned value (move ownership out) or to borrow from something that genuinely outlives the use. At run time a lifetime costs nothing — lifetimes are erased before code generation.)

29.5 What to carry forward

Lifetimes complete the memory-safety story. Ownership (Chapter 27) says every value has one owner and is dropped at end of scope; borrowing (Chapter 28) lets you use a

value without owning it, under aliasing-XOR-mutation; and lifetimes enforce the second borrow rule — that a reference is never used after its referent is dropped. A **lifetime** is a compile-time name for the region a reference is valid over; the borrow checker proves, for every use, that the reference's region is contained in its referent's, and reports E0597 — “borrowed value does not live long enough” — when it is not, catching the C dangling-pointer/use-after-scope bug in every case, where C's warning is only a best-effort heuristic a struct can defeat. Most lifetimes are inferred by **elision**, so you rarely write one; you must name a lifetime ('a) when a function returns a reference whose origin among several inputs is ambiguous, and you must give a struct a lifetime parameter when it holds a reference (E0106, “missing lifetime specifier”). Crucially, lifetimes are inert at run time — erased before code generation — so the whole guarantee is paid for once, at compile time, and a Rust reference is the same bare pointer C would emit.

You now have the three pillars — ownership, borrowing, lifetimes — that let Rust guarantee memory safety without a garbage collector, encoding in the type system the very invariants you proved by hand across the C volumes: memory freed exactly once, used only while valid, and never shared-and-mutated in a way that races. What is missing is not safety but *reuse*. The `longest` function worked on `&str`; what about a function that finds the largest element of any slice, of any type that can be compared? Writing one version per type is the C-with-copy-paste world these volumes have been leaving behind. The next chapter is how Rust expresses code that is generic over types while staying fully checked and monomorphized to concrete, zero-overhead machine code: **traits and generics**. And traits, it will turn out, are also how the concurrency guarantee is finally stated — the `Send` and `Sync` markers that lift borrowing's aliasing discipline across threads. The machine underneath is unchanged; what grows is the vocabulary in which you state, and the compiler checks, what your code requires of the types it works with.

CAN YOU... WITHOUT LOOKING?

- Define a lifetime as a compile-time region, and say what the second borrow rule is.
- Explain why returning `&local` is a dangling reference, and why C's warning is only a heuristic.
- Read the `longest<'a>` signature and say what the 'a asserts and why it is needed there.
- State what E0597 reports and the two real fixes (own vs. borrow-from-something-longer-lived).
- Give the elision intuition and say why a struct holding a reference needs a lifetime (E0106).
- Say what a lifetime costs at run time, and why.

RETRIEVAL — CLOSED BOOK

Cover the chapter. Answer from memory; rate confidence first.

1. **R1.** What is a dangling reference, and why does returning `&local` create one? Why is C's `Wreturn-local-addr` only best-effort? (§29.1)
2. **R2.** In `fn longest<'a>(x: &'a str, y: &'a str) -> &'a str`, what does `'a` assert, and why is an annotation required here? (§29.2)
3. **R3.** What does E0597 report, and what containment condition is the borrow checker verifying? (§29.3)
4. **R4.** State the intuition for lifetime elision, and why a struct with a reference field needs a lifetime (E0106). (§29.4)
5. **R5.** What do lifetimes cost at run time, and how do the three pillars (ownership, borrowing, lifetimes) fit together? (§29.3, §29.5)

FURTHER READING

- **Klabnik & Nichols, *The Rust Programming Language*, 2nd ed. (No Starch Press, 2023), ch. 10.3 “Validating References with Lifetimes” (also free at doc.rust-lang.org/book).** The primary source for lifetime syntax, the `longest` example, and the three elision rules of §29.2 and §29.4. Authors and edition verified via the publisher; the book was not compiled in this environment.
- **The Rust Error Index, entries E0597 and E0106 (doc.rust-lang.org/error_codes).** The official meanings quoted in §29.3–§29.4: E0597 “a value was dropped while it was still borrowed” (borrowed value does not live long enough) and E0106 “a lifetime is missing from a type” (missing lifetime specifier). Each verified against the current error index.
- **Jung, Jourdan, Krebbers & Dreyer, “RustBelt: Securing the Foundations of the Rust Programming Language,” *Proc. ACM on Programming Languages* 2(POPL), Article 66 (2018), DOI 10.1145/3158154.** The formal model in which lifetimes-as-regions and the “references never outlive their referent” guarantee are proved sound. First author Ralf Jung; verified via the ACM Digital Library.
- **The “non-lexical lifetimes” RFC 2094 and the Polonius project notes (rust-lang.github.io).** How the borrow checker computes lifetime regions precisely (last-use rather than lexical scope) — the modern engine behind §29.3. Reference for the checker's internals; verify against the current edition.

Exercises

29.1 WARM-UP In one sentence, what is a lifetime, and what is the second borrow rule that lifetimes enforce?

29.2 WARM-UP Why is `fn make() -> &i32 { let x = 5; &x }` a bug? What has happened to `x` by the time the caller would dereference the result?

29.3 WARM-UP In `fn longest<'a>(x: &'a str, y: &'a str) -> &'a str`, state in words what the three occurrences of `'a` jointly assert.

29.4 **CORE** Take the two C snippets of §29.1 (direct `return &local;` and the version laundered through a struct). (a) Why is each a use-after-scope? (b) Why does `gcc -Wall` warn on the first but not the second? (c) Explain how Rust's rule differs from a best-effort warning, and what error code the analogous Rust function produces.

29.5 **CORE** Explain what the borrow checker verifies when it emits `E0597`. (a) State the containment condition on regions. (b) Why can't the error be fixed by annotating a longer lifetime? (c) Give the two structural fixes and say which each is appropriate for.

29.6 **CORE** State the three lifetime-elision rules and apply them: (a) `fn first(s: &str) -> &str` — why is no annotation needed? (b) `fn get(&self, i: usize) -> &T` — which rule supplies the output lifetime? (c) `fn longest(x: &str, y: &str) -> &str` — why do the rules fail to determine the output, forcing you to write `'a`?

29.7 **COST** Lifetimes are “zero-cost.” (a) At what phase are lifetimes checked, and what happens to them before code generation? (b) What is the runtime representation of a `&'a T` compared with the C pointer it replaces? (c) Contrast this with a garbage-collected or reference-counted language's approach to the same “is this reference still valid?” question: what does each pay, and when?

29.8 **FADED** *Worked, with the last step for you.* A newcomer writes `fn greeting() -> &str { let s = String::from("hi"); &s }` and gets error `[E0597]: `s` does not live long enough`.

Step 1 (given): `s` is a local `String`, dropped at the function's closing brace.

Step 2 (given): `&s` would be returned to the caller, who could use it after `s` is dropped — a dangling reference.

Step 3 (given): no lifetime annotation (including `'static`) can make `s` outlive its scope, so the error is structural, not syntactic.

Step 4 (you): Give the correct fix (change the return type so ownership moves out) and one alternative (borrow from an input that outlives the call), and explain why reaching for `'static` is the wrong instinct here.

29.9 **MIXED** For each, say whether a lifetime annotation is *required*, *elided*, or the code is a *dangling-reference error*, and justify: (a) `fn id(x: &i32) -> &i32`; (b) `struct Wrapper { r: &u8 };` (c) a function returning a reference to its own local; (d) `fn pick(a: &str, b: &str) -> &str` returning one of the two.

29.10 **MIXED** Classify each as *true* or *false* and fix the false ones: (a) “Annotating a reference `'static` makes the value it points at live forever.” (b) “Lifetimes exist at run time and cost a pointer of extra data.” (c) “A struct with a `&str` field can omit its lifetime.” (d) “`E0597` means the reference outlives its referent.” (e) “Most function signatures need explicit lifetimes.”

29.11 **STRETCH** Owning vs. borrowing in data structures. (a) Contrast a struct that stores a `String` with one that stores a `&'a str`: what does each cost (allocation vs. lifetime parameter), and what constraint does the borrowing version place on how the struct may be used? (b) Explain why a self-referential struct (a field that borrows from another field of the same struct) is hard to express with lifetimes, and why. (c) Relate the own-vs-borrow choice back to the C question of “who owns this buffer” from Chapters 9 and 27, and argue that lifetimes are just that question made explicit and checked.

29.12 **LAB** On your own machine (a C toolchain suffices; a Rust toolchain is ideal): (a) Compile the two C snippets of §29.1 with `-O2 -Wall`; confirm the direct `return &local;` warns (`-Wreturn-local-addr`) and the struct-laundered version does not, then catch the latter with `-fsanitize=address` (stack-use-after-scope). (b) If you have `rustc`, write the `fn dangle() -> &str` version and confirm `E0597`; if not, explain from this chapter which line is rejected and why. (c) Write the `longest` function with its lifetime annotation, then a struct with a `&str` field, and note what happens if you omit the struct's lifetime (`E0106`). Note your toolchain versions.

Chapter 30 — Traits and Generics

Before you start: Chapter 29 (lifetimes — with ownership and borrowing, the memory-safety story is complete; this chapter is about *reuse* on top of it), Chapter 27 (ownership; Copy, Clone, and Drop are *traits*, so this chapter names machinery you have already met), Chapter 10 (the optimizing compiler — monomorphization is a compile-time specialization that makes generics free), Chapter 5 (mechanical sympathy — static vs. dynamic dispatch is a real cost you can now name). · The longest function of Chapter 29 worked only on string slices; this chapter is how to write code once that works over many types, fully checked and with no runtime cost.

BEFORE YOU READ ON

From memory, reaching back: (1) Chapter 27 — Copy decided whether `let y = x;` copies or moves; what kind of thing is Copy (a keyword? a type? something a type can implement)? (2) Chapter 10 — what does it mean to specialize code at compile time, and why can a compile-time transformation be free at run time? (3) *Predict:* in C you write a generic sort with `qsort`, passing a comparison *function pointer*; how is the comparison call dispatched — is the target known at compile time or chosen at run time? In Rust, a function generic over a type `T` that you compare with `<` — does the compiler need to know something about `T` to allow the comparison? Write your guesses; §30.2 and §30.3 answer all three.

Ownership, borrowing, and lifetimes made Rust memory-safe; they did nothing for *reuse*. The `longest` function worked on `&str`; a function to find the largest element of a slice would, in C, be written once per element type or smeared over `void*` and function pointers that throw away compile-time type checking. Rust's answer is two intertwined features. **Generics** let you write code parameterized over a type `T`; **traits** let you say what a `T` must be able to *do*, and are how shared behavior is named and required. First, the reuse problem in C, concretely, with a real `qsort` run (§30.1). Then generics and **monomorphization**: the compiler stamps out a specialized copy per concrete type, so generic code is as fast as hand-written type-specific code — zero-cost static dispatch (§30.2). Then traits and **trait bounds**: a generic can only use capabilities its bounds guarantee, and a missing bound is a compile error, `E0277` (§30.3). Then the other dispatch strategy — **trait objects** (`dyn Trait`), one shared version that dispatches through a vtable at run time, exactly the C function-pointer mechanism — and the cost trade between the two (§30.4). We close by carrying traits forward: they are also how error handling and, later, thread-safety are expressed (§30.5).

On the code in this part. The book's environment has a C toolchain (gcc 11.4.0) but *no Rust compiler installed*, so the C programs below were compiled and run and their output is real, whereas the Rust snippets were *not* compiled here; their behavior and the quoted compiler diagnostics are drawn from *The Rust Programming Language* (Klabnik & Nichols, 2nd ed., 2023) and the official Rust error index, and are labeled as such. Nothing about Rust's behavior below is invented — it is cited.

30.1 The reuse problem in C

C gives you two ways to write code that works over many types, and both are unsatisfying. The first is duplication: write `max_int`, `max_float`, `max_long`, one per type, differing only in a type name. The second is type erasure through `void*` plus a function pointer — how the standard library's `qsort` sorts anything. Here is a real `qsort`, compiled and run in this environment (gcc 11.4.0):

```
int cmp_int(const void *a, const void *b) {
    return (*(const int*)a) - (*(const int*)b); /* caller-supplied comparison */
}
int a[] = {5, 2, 9, 1, 7};
qsort(a, 5, sizeof(int), cmp_int);           /* generic sort via void* + fn ptr */

$ gcc -O2 -Wall gen.c -o gen && ./gen
1 2 5 7 9
```

This works, and it is genuinely generic — `qsort` sorts arrays of any element type. But look at the two costs. First, *type safety is gone*: `qsort` sees `void*`, so if you pass a comparator written for `double` while sorting `ints`, it compiles cleanly and corrupts silently — the compiler cannot check that the comparator matches the array, because the types were erased. Second, every comparison is an **indirect call through a function pointer**: `cmp_int` is chosen at run time, so the CPU cannot inline it and pays an indirect branch per comparison (Chapter 5). Rust's generics fix the first cost entirely and let you *choose* whether to pay the second. The `qsort` pattern — a function pointer resolved at run time — is exactly what Rust calls a trait object (§30.4); the duplication-free, fully-checked, fully-inlined alternative is generics with monomorphization (§30.2). Both start from a single idea: name the behavior a type must provide, and let the type system check it.

QUICK CHECK

Name the two ways C lets you write code over many types, and the cost of each. In the `qsort` call, why can the compiler not check that `cmp_int` matches the array being sorted? How is each comparison dispatched — is the target known at compile time? (Answer after the next section.)

30.2 Generics and monomorphization

A **generic** function is parameterized over a type. You write it once, with a type parameter T , and the compiler generates a specialized copy for each concrete type you actually use it with — a process called **monomorphization** (Rust; not compiled here — behavior per the cited book):

```
fn largest<T: PartialOrd + Copy>(list: &[T]) -> T {
    let mut biggest = list[0];
    for &item in list {
        if item > biggest { biggest = item; } // needs T: PartialOrd to use `>`
    }
    biggest
}
// called with &[i32] and &[f64]; the compiler emits a separate specialized
// copy of `largest` for i32 and for f64 – as if you'd written each by hand.
```

At compile time the compiler sees that you called `largest` with an `&[i32]` and an `&[f64]`, and it stamps out two concrete functions — one operating on `i32`, one on `f64` — each with the type filled in. There is no type parameter left at run time: each call goes to an ordinary, type-specific function whose comparison is a direct, inlinable operation, not an indirect call. This is **static dispatch**, and it is *zero-cost*: generic code runs exactly as fast as the hand-duplicated `max_int/max_float` of §30.1, because after monomorphization that is literally what it is — but you wrote it once and the compiler checked it once. The price, paid only in binary size and compile time, is that each concrete instantiation is a separate copy of the code (code bloat if a large generic is used at many types). The abstraction is free where it matters — at run time — which is the recurring promise of this whole part: expressive source, no runtime tax.

THE SAME IDEA THREE WAYS: WRITE ONCE, SPECIALIZE PER TYPE, PAY NOTHING AT RUN TIME

1. In words. A generic function is written once over a type parameter `T`. The compiler monomorphizes it — emits a separate specialized copy for each concrete type used — so at run time there is no type parameter and no dispatch: each call is an ordinary direct call, as fast as hand-written type-specific code.

2. As a picture. Figure 30.1: one generic `largest<T>` box fanning out at compile time into `largest_i32` and `largest_f64` (static dispatch, direct calls), contrasted with one `largest(&dyn ..)` box that stays single and dispatches through a vtable at run time (dynamic dispatch).

3. As numbers/contrast. The C `qsort` of §30.1 pays one indirect function-pointer call per comparison (chosen at run time, not inlinable). Monomorphized Rust generics pay zero: the comparison is a direct, inlinable operation in each specialized copy — the same machine code you'd get from writing the type-specific function by hand.

(Answer to the §30.1 quick check: C lets you write over many types by (1) duplicating code per type — a maintenance cost — or (2) erasing types through `void*` plus a function pointer, as `qsort` does — which costs type safety and an indirect call. The compiler cannot check that `cmp_int` matches the array because `qsort`'s parameters are `void*`: the element type is erased, so a mismatched comparator compiles anyway. Each comparison is dispatched through the function pointer at run time, so the target is not known at compile time and cannot be inlined.)

Two dispatch strategies from one generic idea: monomorphized (static) vs. trait object (dynamic).




Figure 30.1: Static vs. dynamic dispatch. A generic monomorphizes into a direct, inlinable copy per concrete type (top, zero runtime cost, larger binary). A trait object keeps one copy and dispatches each method through a vtable at run time (bottom, an indirect call, smaller binary) — the same mechanism as C's `qsort` function pointer.

30.3 Traits and trait bounds

The `largest` function used `item > biggest`. For that to typecheck, the compiler must know that `T` supports comparison — otherwise a caller could instantiate `largest` with a type that has no ordering, and `>` would be meaningless. A **trait** names a set of behaviors a type can implement; a **trait bound** on a generic (`T: PartialOrd`) requires that `T` implement the trait, and in return lets the generic body use that trait's operations. You have already met traits: `Copy`, `Clone`, and `Drop` from Chapter 27 are traits, which is why “`i32` is `Copy`” was phrased as implementing something. A trait definition and an implementation look like this (Rust; not compiled here — behavior per the cited book):

```
trait Summary {                               // a named set of behavior
    fn summarize(&self) -> String;
}
impl Summary for Article {                    // Article promises to provide it
    fn summarize(&self) -> String { format!("{}", self.title) }
}

fn notify<T: Summary>(item: &T) {             // bound: T must implement Summary
    println!("Breaking! {}", item.summarize()); // allowed BECAUSE of the bound
}
```

The bound `T: Summary` is a contract in both directions: the caller must supply a type implementing `Summary`, and in exchange the body of `notify` may call `summarize`. If you drop the bound and try to call `summarize` anyway, or call `notify` with a type that does not implement `Summary`, the compiler rejects it with `E0277` — *the trait bound is not satisfied* (per the official Rust error index):

```
notify(&5); // i32 does not implement Summary
error[E0277]: the trait bound `i32: Summary` is not satisfied
```

This is the decisive difference from C's `void*` genericity of §30.1. There, a type mismatch was invisible and corrupted silently at run time; here, “this type does not provide the required behavior” is a compile error naming the exact unmet bound. It is also, notably, different from C++ templates, where an unsatisfied requirement surfaces as an error

deep inside the instantiated body; Rust checks the bound at the *definition*, so the generic is verified once, against its stated requirements, independent of any caller. Trait bounds are how a generic states precisely what it needs of its type parameters — no more (so it stays maximally reusable) and no less (so the body is fully checked) — and E0277 is the compiler holding both sides to the contract.

QUICK CHECK

What is a trait, and what does a trait bound like `T: PartialOrd` give the caller and the body of a generic respectively? Which error code fires when a required bound is unmet, and what does it say? Why is that better than C's silent `void*` mismatch and than a C++ template error deep in the body? (Answer below the communicate box.)

30.4 Static vs. dynamic dispatch: trait objects

Monomorphization needs to know every concrete type at compile time. Sometimes you cannot — you want a single `Vec` holding a mix of types that all implement one trait, or a plugin resolved at run time. For that Rust offers a **trait object**, written `dyn Trait` behind a pointer (`&dyn Summary` or `Box<dyn Summary>`). A trait object is a pair of pointers: one to the value, one to a **vtable** — a table of the trait's method implementations for that value's concrete type — and each method call is dispatched indirectly through the vtable at run time (Klabnik & Nichols, 2023, ch. 17.2). This is **dynamic dispatch**, and it is precisely the C `qsort` mechanism of §30.1: an indirect call through a pointer chosen at run time (Rust; not compiled here — behavior per the cited book):

```
// static dispatch: one specialized copy per type, direct/inlinable calls, larger binary
fn notify<T: Summary>(item: &T) { println!("{}", item.summarize()); }

// dynamic dispatch: one copy, vtable call per invocation, smaller binary, any Summary type
fn notify_dyn(item: &dyn Summary) { println!("{}", item.summarize()); }

let items: Vec<Box<dyn Summary>> = vec![Box::new(article), Box::new(tweet)];
// a heterogeneous collection — impossible with a single monomorphized type
```

The trade is explicit and it is a cost trade you can now name from Chapter 5. Static dispatch (generics) resolves the call at compile time: direct, inlinable, zero dispatch cost, but a separate code copy per type (larger binary) and every type known at compile time. Dynamic dispatch (trait objects) resolves at run time through the vtable: one copy of the code and the ability to mix types or choose at run time, but an indirect call per method (not inlinable, an indirect branch — the same cost as `qsort`'s function pointer) and the values must sit behind a pointer. Neither is “better”; the default in Rust is generics (pay nothing, specialize) and you reach for `dyn` deliberately when you need heterogeneity or run-time choice, or to bound code size. The important thing is that unlike C — where `void*` dynamic dispatch also threw away type safety — a Rust trait object is still fully type-checked: `&dyn Summary` can only hold a type that actually implements `Summary`, so you get the flexibility of the function-pointer approach without the silent-corruption risk of §30.1.

TRAP: REACHING FOR `Box<dyn Trait>` BY DEFAULT, AND MISREADING E0277 AS “MY TYPE IS BROKEN”

Two related early mistakes. The first: coming from an object-oriented language, newcomers make everything a `Box<dyn Trait>` because it feels like the familiar interface/subclass pattern — and thereby pay an indirect call on every method and give up inlining, when a generic with a trait bound would have cost nothing at run time and inlined fully. Dynamic dispatch is a real tool (heterogeneous collections, run-time-chosen implementations, keeping binary size down), but it is not the default; generics with bounds are, and reaching for `dyn` reflexively is the analog of making every C call go through a function pointer for uniformity's sake. The second: reading error[E0277]: the trait bound `T: Foo` is not satisfied as “my type is wrong” and casting or wrapping to make it quiet. The error is not saying your type is broken; it is saying the type does not (yet) provide the behavior the generic requires — and the two correct responses are to *implement the trait* for your type (give it the behavior) or, if the generic asked for more than it needs, to *relax the bound* to what the body actually uses. Neither is a cast. As with a C warning, E0277 is naming a real gap between what the code requires and what the type provides; the fix is to close the gap, not to silence the message.

SAY IT LIKE A SYSTEMS ENGINEER

Try it first. A teammate says: “My generic function won't compile until I add `T: PartialOrd` — Rust makes me spell out every requirement, when in C++ a template just works and in Python I'd never write any of this. Isn't this bound stuff just ceremony?” Cover the paragraph and respond in four or five sentences, drawing on §30.2–§30.4.

Model answer. “The bound isn't ceremony — it's what lets the compiler check the generic *once*, at its definition, instead of C++'s way where a missing requirement blows up deep inside the instantiated template with a wall of errors, or Python's way where it blows up at run time on some input you didn't test. `T: PartialOrd` is you telling the compiler exactly what the body needs (an ordering to use `>`), and in exchange it guarantees every caller supplies a type that has it — that's the E0277 you'd get otherwise, caught at compile time and naming the missing trait. It also costs nothing at run time: the generic monomorphizes into a direct, inlinable copy per type, so it's as fast as a hand-written type-specific function. So the bound buys you C's performance with checking C++ and Python can't give you at that stage. Write the minimal bound the body actually uses and you get maximal reuse with full checking.” Framing the bound as check-once-at-the-definition and zero-cost after monomorphization is the senior register.

(Answer to the §30.3 quick check: a trait names a set of behaviors a type can implement; a bound `T: PartialOrd` requires the caller to supply a `T` that implements `PartialOrd`, and in return lets the generic's body use that trait's operations (here `>`). An unmet bound is E0277, “the trait bound is not satisfied.” It beats C's silent `void*` mismatch — which corrupts at run time — and a C++ template error deep in the instantiated body, because Rust checks the bound at the generic's definition and names the exact missing trait.)

30.5 What to carry forward

Generics and traits are how Rust gets reuse without giving up the checking or the performance of the earlier volumes. A **generic** function is written once over a type parameter and **monomorphized** — the compiler emits a specialized copy per concrete type — so it runs as fast as hand-written type-specific code with no runtime dispatch:

zero-cost **static dispatch**. A **trait** names behavior a type can implement (you already knew `Copy`, `Clone`, `Drop`), and a **trait bound** (`T: PartialOrd`) states exactly what a generic requires of its type — checked once, at the definition, with an unmet bound reported as `E0277`, “the trait bound is not satisfied.” When you need heterogeneity or a run-time-chosen implementation, a **trait object** (`dyn Trait`) gives **dynamic dispatch** through a `vtable` — the same indirect-call mechanism as C's `qsort` function pointer, but still fully type-checked, unlike C's `void*`. The choice between them is the Chapter 5 cost trade made explicit: static is the default (free, specialized, larger binary), dynamic is deliberate (one copy, flexible, an indirect call per method).

With ownership, borrowing, lifetimes, generics, and traits, you now have the whole conceptual core of Rust — and the last two chapters of this part put it to work. Traits are not only for behavior like `summarize`; they are how Rust expresses its most important guarantees as ordinary, checkable requirements. The next chapter is **error handling**: `Option<T>` and `Result<T, E>` are generic types (built on exactly the generics of this chapter), and they replace C's `error-code-and-errno` conventions and null pointers with types the compiler forces you to handle — turning “forgot to check the return value” into a thing that does not compile. And the chapter after it, on fearless concurrency, rests on two traits — `Send` and `Sync` — that lift the aliasing discipline of borrowing across threads, making the data races of Chapter 24 a compile error rather than undefined behavior. The pattern to carry forward is that in Rust, a guarantee is usually a trait: name the behavior or the capability, require it in the type, and let the compiler check it — the same move, now applied to errors and to threads.

CAN YOU... WITHOUT LOOKING?

- State the two ways C writes code over many types and the cost of each (duplication; `void*`+function pointer).
- Define monomorphization and explain why generic code is zero-cost at run time.
- Say what a trait and a trait bound are, and what each side of a bound gets.
- Name the error code for an unmet bound (`E0277`) and why compile-time checking beats C's silent mismatch.
- Contrast static and dynamic dispatch: mechanism, cost, and when to choose each.
- Explain how a trait object is like C's `qsort` function pointer but still type-safe.

RETRIEVAL — CLOSED BOOK

Cover the chapter. Answer from memory; rate confidence first.

1. **R1.** What are C's two options for generic code, and what does each cost? Why can't `qsort`'s comparator be type-checked? (§30.1)
2. **R2.** What is monomorphization, and why is a Rust generic zero-cost at run time? (§30.2)
3. **R3.** What is a trait bound, what does each side gain, and which error reports an unmet one? (§30.3)
4. **R4.** What is a trait object, how is a method call on it dispatched, and how does that compare to `qsort`? (§30.4)
5. **R5.** When do you choose static dispatch vs. dynamic dispatch, and what is the cost of each? (§30.4)

FURTHER READING

- **Klabnik & Nichols, *The Rust Programming Language*, 2nd ed. (No Starch Press, 2023), ch. 10.1 “Generic Data Types” and ch. 10.2 “Traits” (also free at doc.rust-lang.org/book).** The primary source for generics, monomorphization, traits, and trait bounds of §30.2–§30.3, including the `largest` and `Summary/notify` examples. Authors and edition verified via the publisher; not compiled here.
- **Klabnik & Nichols, *The Rust Programming Language*, 2nd ed. (No Starch Press, 2023), ch. 17.2 “Using Trait Objects That Allow for Values of Different Types.”** The source for trait objects, the value-plus-vtable representation, and dynamic dispatch of §30.4. Chapter number and title verified via the publisher's listing and doc.rust-lang.org/book.
- **The Rust Error Index, entry E0277 (doc.rust-lang.org/error_codes).** The official meaning quoted in §30.3: “you tried to use a type which doesn't implement some trait in a place which expected that trait” — the `unmet-trait-bound` error. Verified against the current error index.
- **Stroustrup, *The C++ Programming Language*, 4th ed. (Addison-Wesley, 2013), on templates.** The C++ template model — instantiation-time errors deep in the body — that Rust's define-time bound checking of §30.3 improves on; useful contrast for where trait bounds come from and what they fix.

Exercises

- 30.1** WARM-UP In one sentence each, define a generic function, a trait, and a trait bound.
- 30.2** WARM-UP What is monomorphization, and why does it mean a generic function has no runtime dispatch cost?
- 30.3** WARM-UP What error code does the compiler produce when a generic is used with a type that does not satisfy its trait bound, and what does the message say?

30.4 **CORE** Take the C `qsort` of §30.1. (a) In what sense is it generic, and what are the two costs it pays? (b) Show concretely how passing a comparator written for the wrong element type would compile yet misbehave, and why the compiler cannot catch it. (c) Explain which Rust feature (generics or trait objects) corresponds to the `qsort` mechanism, and which gives the duplication-free, fully-checked, fully-inlined alternative.

30.5 **CORE** Explain the two-directional contract of a trait bound `T: Summary` on `fn notify<T: Summary>(item: &T)`. (a) What must the caller supply? (b) What may the body do that it could not without the bound? (c) Contrast Rust's checking the bound at the definition with C++'s reporting an unmet requirement inside the instantiated template body.

30.6 **CORE** Contrast static dispatch (generics) with dynamic dispatch (trait objects). (a) How is a call resolved in each — at compile time or run time — and what is the machine-level cost of each? (b) What is a trait object's runtime representation? (c) Give one situation that requires dynamic dispatch and one where static dispatch is clearly preferable.

30.7 **COST** Weigh the costs of monomorphization vs. trait objects. (a) What does monomorphization cost that dynamic dispatch does not, and in what resource? (b) What does dynamic dispatch cost per call that a monomorphized generic does not? (c) For a small trait method called in a tight loop over a homogeneous collection, which dispatch would you choose and why; for a large plugin interface with many implementations, which — and what is the deciding factor each time?

30.8 **FADED** *Worked, with the last step for you.* A newcomer writes `fn largest<T>(list: &[T]) -> T { let mut b = list[0]; for &x in list { if x > b { b = x; } } b }` and gets `error[E0277]: the trait bound `T: PartialOrd` is not satisfied on the x > b line (and a second on the Copy requirement).`

Step 1 (given): the body uses `>`, which requires `T` to have an ordering — the `PartialOrd` trait.

Step 2 (given): the body also copies elements out of the slice (`b = x, list[0]`), which requires `T: Copy`.

Step 3 (given): without those bounds, the compiler cannot know `T` supports the operations, so it rejects the definition — before any caller exists.

Step 4 (you): Give the corrected signature with the minimal bounds, explain why bounding at the definition (not the call site) is what lets the generic be checked once, and say why casting or wrapping to silence `E0277` is the wrong instinct.

30.9 **MIXED** For each, say whether static dispatch (generic + bound) or dynamic dispatch (trait object) is the better fit, and why: (a) summing a slice of any numeric type in a hot loop; (b) a `Vec` holding several different shape types that all implement `Draw`; (c) a function called at exactly one concrete type; (d) a rendering pipeline that loads drawable plugins chosen at run time.

30.10 **MIXED** Classify each as *true* or *false* and fix the false ones: (a) “A generic function carries a type parameter at run time and dispatches on it.” (b) “A trait object is dispatched at run time through a vtable.” (c) “E0277 means your type has a bug.” (d) “Monomorphization makes binaries smaller.” (e) “Like C's `void*` generics, Rust trait objects give up type safety.”

30.11 **STRETCH** “A guarantee in Rust is usually a trait.” (a) Explain how `Copy` (Chapter 27) is a trait that governs whether assignment moves or copies, and how `T: PartialOrd` governs whether `>` is allowed — both the same mechanism. (b) Preview how `Option<T>` and `Result<T, E>` (next chapter) use generics to make error handling checkable. (c) Preview how thread-safety could be expressed as a trait requirement (`Send/Sync`): what would a bound like “`T: Send`” have to mean, and why would requiring it at a thread-spawn boundary turn a data race into a compile error?

30.12 **LAB** On your own machine (a C toolchain suffices; a Rust toolchain is ideal): (a) Compile and run the §30.1 `qsort` demo; then deliberately pass a comparator written for a different element type and confirm it still compiles (type erasure). (b) If you have `rustc`, write `largest<T: PartialOrd + Copy>`, call it at two types, and confirm it compiles; then remove the bounds and observe E0277; if not, explain from this chapter which line is rejected and which bound is missing. (c) Write a trait `Draw` with two implementors and put them in a `Vec<Box<dyn Draw>>`; explain why a monomorphized generic could not hold both in one vector. Note your toolchain versions.

Chapter 31 — Error Handling

Before you start: Chapter 30 (traits and generics — `Option` and `Result` are generic enums, and `?` leans on the `From` trait, so this chapter is those tools put to work), Chapter 9 (manual allocation and what goes wrong — C signals failure with return codes, `errno`, and `NULL`, none of which the compiler makes you check), Chapter 27 (ownership — an error value is owned and moved like any other, which is why `?` can hand it back to the caller), Chapter 10 (undefined behavior — dereferencing a `NULL` you forgot to check is exactly the UB a type can rule out). · Chapter 30 promised that in Rust a guarantee is usually a type; this chapter is that promise applied to failure — “did you handle the error?” becomes a thing that does not compile.

BEFORE YOU READ ON

From memory, reaching back: (1) Chapter 30 — what is a generic type, and what does a trait bound like `T: From<E>` let a body assume? (2) Chapter 9 — name the ways C reports that an operation failed (a sentinel return value, `errno`, a `NULL` pointer); does the C compiler force the caller to inspect any of them? (3) *Predict:* in C, `FILE *fp = fopen(path, "r");` returns `NULL` when the file is missing; if you forget to check and call `fgets(buf, n, fp)`, does that *compile*? does it *run* correctly? In Rust, `File::open` returns a `Result`; can you use the file handle without first confronting the failure case? Write your guesses; §31.1 and §31.3 answer all three.

Every real program meets conditions it cannot control: a file is missing, a socket closes, a parse fails, an allocation returns nothing. The question is not whether errors happen but whether the language *makes you deal with them*. C's answer, from Chapter 9, is a set of conventions — a sentinel return value (`-1`, `NULL`), the global `errno`, sometimes a status code — all of which the compiler is happy to let you ignore. Rust's answer is to make failure a value with a type the compiler forces you to handle: `Option<T>` for “a value that might be absent” (there is no null), and `Result<T, E>` for “an operation that might fail.” Both are the ordinary generic enums of Chapter 30, and handling them is ordinary pattern matching. First, how C signals failure and why every signal is silently ignorable, in real C that compiles and then segfaults (§31.1). Then `Option<T>`: the type that abolishes the null pointer and forces the absent case into the open (§31.2). Then `Result<T, E>` for recoverable failure, marked `#[must_use]` so ignoring it warns (§31.3). Then the `?` **operator**, which propagates an error to the caller in one character, and `panic!` for the unrecoverable — with the compile error `E0277` that fires when `?` is used in a function that returns neither (§31.4). We close by carrying error values forward into concurrency, where a panic on one thread and a `Result` crossing a channel are the next chapter's concern (§31.5).

On the code in this part. The book's environment has a C toolchain (gcc 11.4.0) but *no Rust compiler installed*, so the C programs below were compiled and run and their output is real, whereas the Rust snippets were *not* compiled here; their behavior and the quoted compiler diagnostics are drawn from *The Rust Programming Language* (Klabnik & Nichols, 2nd ed., 2023) and the official Rust error index, and are labeled as such. Nothing about Rust's behavior below is invented — it is cited.

31.1 How C signals failure — and why it is ignorable

C has no single error mechanism; it has conventions. A function returns a *sentinel* — `NULL` from `fopen` and `malloc`, `-1` from `read` and `open` — and often also sets the global `errno` to say why. The caller is supposed to check. Nothing enforces it. Here is real C that forgets one check, compiled and run in this environment (gcc 11.4.0):

```
FILE *fp = fopen("does-not-exist.txt", "r"); /* returns NULL on failure */
char buf[64];
fgets(buf, sizeof buf, fp);                /* fp never checked for NULL */
printf("read: %s\n", buf);

$ gcc -O2 -Wall cerr.c -o cerr
cerr.c:5: warning: ignoring return value of 'fgets' ... [-Wunused-result]
$ ./cerr
Segmentation fault (core dumped)           # exit status 139
```

Look closely at what the compiler did and did not say. It emitted a warning — but about the wrong thing: `fgets` happens to be marked `warn_unused_result`, so ignoring *its* return drew a note. The actual bug — that `fp` is `NULL` and is dereferenced inside `fgets` — drew nothing, because checking `fopen`'s return is a convention, not a rule. The program compiles cleanly enough to ship and then segfaults on the missing file: the `NULL` dereference is the undefined behavior of Chapter 10, reached because the failure was expressible but not enforced. This is the recurring shape of C error handling: the information is there (the sentinel, `errno`), and the discipline to check it is entirely on you. Rust's move is to make the failure part of the return *type*, so there is no unchecked value to forget — the success value is only reachable by first accounting for the failure. That type is an enum, and we meet the simpler one first.

QUICK CHECK

Name three ways C reports failure. In the snippet above, which check was skipped, and why did the compiler not flag the real bug even with `-Wall`? What runtime event does the skipped check ultimately cause, and which chapter's concept is that? (Answer after the next section.)

31.2 `Option<T>`: there is no null

Rust has no null pointer. A value that may be absent has type `Option<T>`, a generic enum with two variants — `Some(T)` carrying a value, or `None`. Because absence is a distinct variant rather than a special pointer bit, the type system knows a `T` might not be there and refuses to let you use it as if it always is; you must pattern-match both cases (Rust; not compiled here — behavior per the cited book):

```
enum Option<T> { Some(T), None }           // in the standard library, generic over T

fn first_char(s: &str) -> Option<char> {
    s.chars().next()                      // Some(c) if non-empty, None if empty
}

match first_char(name) {
    Some(c) => println!("starts with {c}"),
    None    => println!("empty name"),    // the compiler forces this arm to exist
}
```

The point is what you *cannot* do: you cannot treat an `Option<char>` as a `char` and call methods on it, because it is a different type. To get at the inner value you must go through the `None` case — by `match`, or by a combinator like `unwrap_or`, or by `?` (§31.4). This is precisely the check C omitted in §31.1, now made unskippable: the “might be missing” is in the type, so the missing case cannot be silently dereferenced. It is the language-level fix for what C. A. R. Hoare, who introduced the null reference in 1965, later called his “billion-dollar mistake” — a value that inhabits every reference type yet supports no operation, so every use is a latent crash (Hoare, QCon London, 2009). `Option` puts the absent case back into the type where the compiler can see it, turning a whole family of C's NULL-dereference segfaults into a match arm you are made to write.

31.3 `Result<T, E>`: recoverable failure as a value

Absence is one thing; a failed operation that carries a *reason* is another. For that Rust has `Result<T, E>`, a generic enum over the success type `T` and the error type `E`, with variants `Ok(T)` and `Err(E)`. Fallible standard-library operations return it — `File::open` returns `Result<File, io::Error>` — and it is annotated `#[must_use]`, so ignoring one is a warning, not a shrug (Rust; not compiled here — behavior per the cited book):

```
enum Result<T, E> { Ok(T), Err(E) }       // generic over success T and error E

let f = File::open("data.txt");           // Result<File, io::Error>
let file = match f {
    Ok(handle) => handle,                  // the File, only reachable through Ok
    Err(e)     => return Err(e),          // must confront the failure to proceed
};
// File::open("x"); // warning: unused `Result` that must be used (#[must_use])
```

Compare this line by line with the C of §31.1. There, `fopen` returned a `FILE*` that was `NULL` on failure, and reaching the “file” without checking was legal and segfaulted. Here, `File::open` returns a `Result` whose `File` is *inside* the `Ok` variant: the only way to a usable `File` is through code that also names what to do on `Err`. If you throw the whole `Result` away, the `#[must_use]` attribute makes the compiler warn that a result that must be used was not — the diagnostic C's `warn_unused_result` gestured at, here applied to every fallible operation by construction. You still have escape hatches — `unwrap()` returns the `Ok` value or **panics** on `Err`, and `expect("msg")` does the same with your message — but they are visible, greppable choices to convert a recoverable error into a crash, not the silent default. Handling the error is the path of least resistance; ignoring it takes a deliberate keystroke.

QUICK CHECK

What are the two variants of `Result<T, E>`, and where does the success value live? What does `#[must_use]` buy you over C's conventions, and what do `unwrap()` / `expect()` do on an `Err`? (Answer below the communicate box.)

31.4 Propagation with `?`, and `panic!` for the unrecoverable

Matching every `Result` by hand would bury real logic under error plumbing. The `?` **operator** is the shorthand: on an expression of type `Result`, `x?` evaluates to the inner value if it is `Ok`, and otherwise *returns the `Err` from the enclosing function* — an early return that hands the error to the caller. When the caller's error type differs, `?` converts through the `From` trait of Chapter 30 (Rust; not compiled here — behavior per the cited book):

```
fn read_config(path: &str) -> Result<String, io::Error> {
    let mut f = File::open(path)?;           // if Err, return it now; else bind the File
    let mut s = String::new();
    f.read_to_string(&mut s)?;               // same: propagate on Err, continue on Ok
    Ok(s)                                    // reached only if every step succeeded
}
```

Each `?` is the whole match-and-early-return of §31.3 compressed to one character, and the function reads as the happy path with the failures factored out. The catch — and it is a helpful one — is that `?` only makes sense in a function that itself returns a `Result` (or an `Option`): it needs somewhere to return the `Err` to. Use `?` in a function returning `()` and the compiler stops you with **E0277** — “the `?` operator can only be used in a function that returns `Result` or `Option`” (per the official Rust error index) — the same error code as an unmet trait bound in Chapter 30, because “this function can absorb a propagated error” is itself a trait requirement. The fix is to give the function a `Result` return type, not to reach for `unwrap`.

Not every failure is recoverable. When a program hits a bug or a violated invariant it cannot sensibly continue past — an index out of bounds, an assertion, a “this should be impossible” branch — it calls `panic!`, which unwinds the stack (running destructors) and aborts the thread. The distinction is the whole design: `Result` is for *expected, recoverable* conditions the caller should handle (a missing file, a bad request); `panic!` is for *unexpected, unrecoverable* ones that signal a bug (Klabnik & Nichols, 2023, ch. 9). Choosing `unwrap()` is choosing “if this is `Err`, that is a bug, so panic” — sometimes right (a hard-coded string you know parses), often a shortcut that turns a routine error into a crash. The rule of thumb: return a `Result` when the caller can do something about it; `panic!` only when continuing would mean running on broken assumptions.

THE SAME IDEA THREE WAYS: AN ERROR IS A VALUE THE TYPE SYSTEM FORCES YOU TO HANDLE

1. In words. Fallibility is encoded in the return type — `Option<T>` for absence, `Result<T, E>` for failure-with-reason — so the success value is reachable only through code that also accounts for the failure. `?` propagates an `Err` to the caller; `panic!` is reserved for unrecoverable bugs. Ignoring a `Result` warns (`#[must_use]`); using `?` without a place to return the error is `E0277`.

2. As a picture. Figure 31.1: one call fanning into two arms — `Ok(T)/Some(T)` to the happy path, `Err(E)/None` to a handler — with `?` drawn as an arrow from the `Err` arm straight back out to the caller, versus C's single return value with a dashed “did you remember to check?” branch that falls through to a `NULL` dereference.

3. As numbers/contrast. The C of §31.1 compiled and reached exit status 139 (`SIGSEGV`) on a missing file because the `NULL` check was optional. The Rust equivalent cannot reach the `File` without a matching `Err` handler; the same “forgot to check” is a compile-time warning (`#[must_use]`) or, for `?` in the wrong context, a compile error (`E0277`) — a runtime segfault turned into a message before the program runs.

(Answer to the §31.1 quick check: C reports failure by sentinel return values (`NULL`, `-1`), by setting `errno`, and by explicit status codes. The skipped check was `fopen`'s return for `NULL`; `-Wall` did not flag it because checking a return sentinel is a convention the compiler does not enforce — it warned only about ignoring `fgets`'s own `warn_unused_result`. The unchecked `NULL` is dereferenced inside `fgets`, causing a segmentation fault — the undefined behavior of Chapter 10.)

An error is a value with two arms; C's single return has an optional, skippable check.




Figure 31.1: A `Result/Option` forces both arms — the success value lives inside `Ok/Some`, and `?` routes an `Err` back to the caller. C's single return value makes the failure check optional, and skipping it falls through to the `NULL` dereference of §31.1.

TRAP: REFLEXIVE .UNWRAP(), AND SILENCING #[MUST_USE] WITH LET _ =

Two related shortcuts undo everything this chapter buys. The first is reaching for `.unwrap()` or `.expect()` on every `Result` to make the types line up, because matching feels like ceremony. Each `unwrap` is a decision that “an `Err` here is an unrecoverable bug” — fine for a literal you know parses, but on a file open, a network read, or user input it converts a routine, recoverable error into a panic! that aborts the thread. The mechanical fix when the compiler complains that `?` needs a `Result`-returning function (E0277) is *not* to slap on `unwrap`; it is to give the function a `Result` return type and propagate. The second shortcut is silencing the `#[must_use]` warning with `let _ = fallible();` — the exact analog of ignoring C's return sentinel, now with a paper trail. Both are the §31.1 mistake wearing Rust syntax: the error is still there, you have just re-hidden it. Handle it, propagate it with `?`, or — if a failure truly is a bug — panic! *deliberately* with a message that says why; do not launder a recoverable error into a crash to quiet the compiler.

SAY IT LIKE A SYSTEMS ENGINEER

Try it first. A teammate says: “`Result` and `?` everywhere are just checked exceptions with extra typing — in C I return `-1` and check it, which is simpler and I can see the control flow. Isn't all this enum stuff ceremony?” Cover the paragraph and respond in four or five sentences, drawing on §31.1–§31.4.

Model answer. “The difference isn't typing, it's what the compiler enforces. In C the `-1-and-check` is a convention — nothing stops you from skipping it, and when you do you don't get an exception, you get the §31.1 segfault at run time on a missing file. `Result<T, E>` puts the success value inside `Ok`, so you literally can't reach it without naming the `Err` case, and `#[must_use]` warns if you drop the whole thing — that's the check you keep meaning to write in C, made non-optional. And `?` is not hidden control flow like an exception: it's a visible, one-character early return that only compiles in a function that returns `Result`, so the error path is explicit and type-checked all the way up. It's also not exceptions — there's no invisible unwinding through ten stack frames; a panic is the separate, deliberate tool for *unrecoverable* bugs. So you get C's explicit control flow plus the compiler making sure you didn't forget the check.” Framing `Result` as “the check you meant to write, enforced” and separating it cleanly from panics is the senior register.

(Answer to the §31.3 quick check: `Result<T, E>` has variants `Ok(T)` and `Err(E)`, and the success value lives inside `Ok` — reachable only through code that also names the `Err` case (a `match`, a combinator, or `?`). `#[must_use]` makes the compiler warn if you discard the whole `Result`, an enforcement C's return-value convention lacks. `unwrap()` returns the `Ok` value or `panic!`s on `Err`, and `expect("msg")` does the same with your message — visible, greppable ways to convert a recoverable error into a crash.)

31.5 What to carry forward

Error handling in Rust is the Chapter 30 pattern once more: name the possibility in a type, and let the compiler force the handling. `Option<T>` (`Some/None`) replaces the null pointer, so a value that might be absent cannot be used as if it were present — abolishing a whole class of C's NULL-dereference crashes. `Result<T, E>` (`Ok/Err`) makes a fallible operation return its success value *inside* the success variant and is `#[must_use]`, so the

failure cannot be reached without being named and cannot be silently discarded. The `?` **operator** propagates an `Err` to the caller in one character, converting error types through `From`, and is only valid where there is a `Result` or `Option` to return to (E0277 otherwise). `panic!` is the separate tool for the unrecoverable — a bug or violated invariant — and `unwrap()/expect()` are the visible, greppable way to opt into it, never the silent default. The whole design turns §31.1's optional-and-forgotten check, whose penalty was a runtime segfault, into a compile-time obligation.

The next chapter, **fearless concurrency**, is where these threads (literally) come together. Errors do not stop at a thread boundary: a worker thread can `panic!`, and its `Result` — of joining it, of the work it did — has to travel back to whoever spawned it, often across a channel that itself yields `Result` when the other end has hung up. And the deeper link is the one Chapter 30 previewed: just as `Result` encodes “this might fail” in a type, the traits `Send` and `Sync` encode “this is safe to move to, or share across, threads,” lifting the ownership and borrowing discipline of Chapters 27–28 over the thread boundary — so the data races you had to reason about by hand in Chapter 24 become, like the unchecked errors here, a compile error rather than undefined behavior. Same move, applied to time and threads instead of to failure.

CAN YOU... WITHOUT LOOKING?

- Name three ways C signals failure and explain why each is silently ignorable (§31.1's segfault).
- State the variants of `Option<T>` and why it abolishes the null pointer.
- State the variants of `Result<T, E>`, where the success value lives, and what `#[must_use]` does.
- Explain what `?` does, how it converts error types, and why E0277 fires when it is misused.
- Distinguish `Result` (recoverable) from `panic!` (unrecoverable), and say what `unwrap()` chooses.
- Explain why reflexive `unwrap` and `let _ =` re-create the C mistake.

RETRIEVAL — CLOSED BOOK

Cover the chapter. Answer from memory; rate confidence first.

1. **R1.** How does C signal failure, and why did the §31.1 program compile with `-Wall` yet segfault? (§31.1)
2. **R2.** What is `Option<T>`, and how does it eliminate a class of NULL-dereference bugs? (§31.2)
3. **R3.** What is `Result<T, E>`, where is the success value, and what does `#[must_use]` add over C? (§31.3)
4. **R4.** What does `?` do, what trait does it use to convert errors, and which error code reports its misuse? (§31.4)
5. **R5.** When do you return a `Result` vs. call `panic!`, and what does `unwrap()` really choose? (§31.4)

FURTHER READING

- **Klabnik & Nichols, *The Rust Programming Language*, 2nd ed. (No Starch Press, 2023), ch. 9 “Error Handling” (also free at doc.rust-lang.org/book).** The primary source for `Result`, `panic!`, the `?` operator, and the recoverable-vs-unrecoverable guidance of §31.3–§31.4; ch. 6 covers `Option` and enums. Authors, edition, and chapter numbers verified via the publisher and the online book; not compiled here.
- **The Rust Error Index, entries E0277 and E0308 (doc.rust-lang.org/error_codes).** The official meanings used in §31.4: E0277 covers “the `?` operator can only be used in a function that returns `Result` or `Option`” (an unsatisfied `FromResidual`/trait bound), and E0308 is “mismatched types.” Both verified against the current error index.
- **C. A. R. Hoare, “Null References: The Billion Dollar Mistake,” QCon London, 2009 (recorded talk, InfoQ).** Hoare's account of introducing the null reference in ALGOL W (1965) and why it was a mistake — the motivation for `Option<T>` in §31.2. A talk, not a paper; speaker, venue, and year verified via the InfoQ recording.
- **The Rust standard-library documentation for `std::option::Option` and `std::result::Result` (doc.rust-lang.org/std).** The definitions, the `#[must_use]` annotation on `Result`, and the combinators (`unwrap`, `expect`, `unwrap_or`, `map`, and `then`) referenced throughout. Reference for the exact API; verify against the current stable docs.

Exercises

- 31.1** **WARM-UP** In one sentence each, say what `Option<T>` and `Result<T, E>` represent, and name the two variants of each.
- 31.2** **WARM-UP** What does the `?` operator do to an expression of type `Result<T, E>` — what happens on `Ok`, and what on `Err`?
- 31.3** **WARM-UP** What is the difference between returning a `Result` and calling `panic!`, and which is `unwrap()` a shortcut for?
- 31.4** **CORE** Take the C of §31.1. (a) Which failure did the program not check, and how does C signal that failure? (b) Why did `gcc -Wall` not warn about the real bug, and what was the warning it *did* emit about? (c) Explain, in terms of Chapter 10, what runtime event the missing check causes and why it is undefined behavior.
- 31.5** **CORE** Explain how `Result<T, E>` makes the §31.1 mistake hard to repeat. (a) Where does the success value live, and what must you write to reach it? (b) What does `#[must_use]` add if you discard the whole `Result`? (c) In what sense is `unwrap()` the closest Rust analog to skipping C's check, and why is it at least visible?
- 31.6** **CORE** A function `fn load(path: &str)` with return type `()` contains `let f = File::open(path)?;` and fails to compile with `error[E0277]: the ? operator can only be used in a function that returns Result or Option.` (a) Why is `?` illegal here specifically? (b) Give the corrected signature. (c) Why is adding `.unwrap()` instead the wrong fix in general?

31.7 **COST** Consider the cost of the two error styles, not in CPU cycles but in *failure modes*. (a) What is the worst-case runtime outcome of a forgotten check in C (as in §31.1)? (b) What is the worst-case outcome of the corresponding omission in Rust — distinguishing “dropped the `Result`” from “reflexively `unwrap()`ed it”? (c) Argue which class of failure (compile-time diagnostic, warning, or runtime UB) each lands in, and why that ordering matters for a system you cannot easily debug in production.

31.8 **FADED** *Worked, with the last step for you.* A newcomer writes `fn first(path: &str) -> String { let s = std::fs::read_to_string(path)?; s }` in a function returning `String` and gets error[E0277] on the `?`.

Step 1 (given): `read_to_string` returns `Result<String, io::Error>`, so `?` may early-return an `io::Error`.

Step 2 (given): the function's declared return type is `String`, which has no place for that propagated `Err` to go — hence E0277.

Step 3 (given): to keep `?`, the function must return a `Result` whose error type the `io::Error` can become.

Step 4 (you): Give the corrected signature and body (returning `Result<String, io::Error>`), and explain why this is preferable to changing `?` to `.unwrap()` for a file that might legitimately be missing.

31.9 **MIXED** For each situation, say whether you would return `Result`, return `Option`, or `panic!`, and why: (a) looking up a key that may or may not be in a map; (b) parsing a network request whose bytes may be malformed; (c) indexing a slice with an integer your own code just computed and proved in-range; (d) opening a user-specified config file.

31.10 **MIXED** Classify each as *true* or *false* and fix the false ones: (a) “Rust's null pointer is written `None`.” (b) “`?` works in any function.” (c) “Dropping a `Result` is a compile error.” (d) “`unwrap()` on an `Err` returns a default value.” (e) “`?` can convert one error type into another on the way out.”

31.11 **STRETCH** “An error is a value the type system forces you to handle.” (a) Explain how this is the same mechanism as Chapter 30's trait bounds — what “requirement” is `?-` in-a-function stating, and why is its violation the *same* error code (E0277) as an unmet bound? (b) Show how `Option` and `Result` being *generic* (Chapter 30) is what lets one `?` work uniformly across every fallible operation. (c) Preview: if a worker thread's result must return to its spawner, what type would carry “the work, or why it failed,” and how might a thread `panic!` be observed by the parent?

31.12 **LAB** On your own machine (a C toolchain suffices; a Rust toolchain is ideal): (a) Compile and run the §31.1 `cerr.c`; confirm it builds and observe the exit status on a missing file (expect a SIGSEGV, status 139). Add the missing `if (fp == NULL)` check and confirm the crash becomes a clean error message. (b) If you have `rustc`, write `read_config` from §31.4 returning `Result<String, io::Error>`, call it on a missing path, and observe that the error is returned rather than crashing; then change the return type to `()` and observe E0277. Without `rustc`, explain from this chapter which line each change affects. (c) Replace one `?` with `.unwrap()` and describe how the failure mode changes. Note your toolchain versions.

Chapter 32 — Fearless Concurrency

Before you start: Chapter 24 (the memory model — a data race is undefined behavior you had to rule out by hand; this chapter makes it a compile error), Chapter 28 (borrowing — aliasing XOR mutation, the rule that Send and Sync lift across threads), Chapter 20 (mutual-exclusion locks — Rust's Mutex bundles the lock *with* the data it guards), Chapter 31 (error handling — joining a thread and receiving on a channel both yield Result). · This chapter finishes the Rust part: the aliasing discipline you proved by hand in the concurrency volume becomes two traits the compiler checks, so a data race stops compiling.

BEFORE YOU READ ON

From memory, reaching back: (1) Chapter 24 — what exactly is a *data race*, and what is its status in C and C++ (defined behavior? undefined?); (2) Chapter 28 — state the borrow rule in one line (how many mutable references may coexist with how many shared ones?). (3) *Predict*: in C, two threads each increment a shared long a million times with no lock; does the program *compile*? does it print the expected total? In Rust, can you hand the same &mut to two threads at once? Write your guesses; §32.1 and §32.3 answer all three.

The concurrency volume (Chapters 20–26) taught the hazards — data races, torn reads, lost updates, deadlock — and the tools to avoid them: locks, atomics, the memory model. What it could not give you was a *guarantee*. In C, whether your locking is correct is a property you argue for by inspection, and a single unlocked access to shared mutable state is undefined behavior (Chapter 24) that may work in testing and corrupt in production. Rust calls its alternative **fearless concurrency**: the same aliasing rule that made single-threaded code safe — one mutable reference or many shared ones, never both (Chapter 28) — is extended across threads by two marker traits, **Send** and **Sync**, so that sharing mutable state without synchronization does not type-check. First, the C data race, compiled and run so you can see it compile cleanly and print the wrong answer (§32.1). Then Send and Sync: what each means and how `thread::spawn` uses them to police the thread boundary (§32.2). Then the compiler stopping the §32.1 race with two named diagnostics — **E0373** (a closure would outlive its borrow) and **E0277** (a non-thread-safe type crossing threads) (§32.3). Then the safe way to share: **Arc<Mutex<T>>** and message-passing **channels**, where the Mutex holds the data it guards and joining or receiving yields the `Result` of Chapter 31 (§32.4). We close by completing the Rust part and turning toward performance (§32.5).

On the code in this part. The book's environment has a C toolchain (gcc 11.4.0) but *no Rust compiler installed*, so the C programs below were compiled and run and their output is real, whereas the Rust snippets were *not* compiled here; their behavior and the quoted compiler diagnostics are drawn from *The Rust Programming Language* (Klabnik & Nichols, 2nd ed., 2023), the Rustonomicon, and the official Rust error index, and are labeled as such. Nothing about Rust's behavior below is invented — it is cited.

32.1 The C data race, compiled and run

Here is the smallest interesting concurrency bug: two threads increment one shared counter, with no synchronization. Each `counter++` is a read-modify-write — load, add one, store — and if the two threads interleave between a load and its store, one update is lost (Chapter 24). Compiled and run in this environment (gcc 11.4.0), without optimization so each `++` is a real load/add/store:

```
static long counter = 0;           /* shared, no lock */
static void *bump(void *arg) {
    for (int i = 0; i < 2000000; i++) counter++; /* a data race */
    return NULL;
}
// main: create two threads running bump, join both, print counter

$ gcc -O0 -Wall -pthread race.c -o race
$ ./race
expected 4000000, got 2453233
$ ./race
expected 4000000, got 2284135
```

and a different number each run:
2284135, 2799577, 2159965, ...

Two things to notice. First, it *compiled cleanly* — the data race is invisible to the C compiler; there is no rule in the language against two threads touching `counter` at once, only your discipline to add a lock. Second, the answer is *wrong and nondeterministic*: roughly two million of the four million increments vanished, and a different number vanishes each run, because the loss depends on the exact interleaving. This is the undefined behavior of Chapter 24 made concrete — the program has no defined meaning, so “works in testing” tells you nothing. The C fix is a mutex around `counter++`, which you must remember to place on *every* access; forget one and the race returns, silently. Rust's move is to make the sharing itself illegal unless it is synchronized — and to do it with the ordinary trait mechanism of Chapter 30.

QUICK CHECK

Why did the §32.1 program compile despite the data race? Why is the printed total both wrong and different each run, and which chapter's concept is that? What is the C fix, and what is its fragility? (Answer after the next section.)

32.2 `Send` and `Sync`: thread-safety as a trait

Rust encodes thread-safety in two **marker traits** — traits with no methods, which simply classify types. **`Send`** means a value of the type is safe to *transfer* (move) to another thread. **`Sync`** means it is safe to *share* a reference to the type across threads — precisely, `T` is `Sync` when `&T` is `Send` (Rustonomicon, “Send and Sync”). They are *auto traits*: the compiler derives them structurally, so a struct is `Send` if all its fields are, and you almost never write them by hand. Most types are both; the important exceptions are the ones that would make a race possible. Crucially, `thread::spawn` requires its closure to be `Send` and `'static`, which is where the borrow rule of Chapter 28 crosses the thread boundary (Rust; not compiled here — behavior per the cited sources):


```
// thread::spawn's bound (simplified): the closure F must be Send and 'static
pub fn spawn<F>(f: F) -> JoinHandle<...> where F: FnOnce() + Send + 'static { ... }

let data = vec![1, 2, 3];
let handle = thread::spawn(move || {      // `move` gives the closure ownership of data
    println!("{:?}", data);              // data is Send (Vec<i32> is Send), so this is
allowed
});
handle.join().unwrap();                    // wait for the thread; join yields a Result
```

The bound `F: Send + 'static` is the whole guarantee in one line. `Send` ensures every value captured by the closure is safe to move to the new thread; `'static` ensures the closure holds no borrow that could dangle when the parent frame ends (Chapter 29). What the type system will not let you do is give two threads simultaneous mutable access to the same data without synchronization — because that would require sharing something that is not `Sync`, or aliasing a `&mut`, both of which the borrow rules forbid. The §32.1 C race required exactly that unsynchronized shared mutation; in Rust it cannot be expressed. The reference-counting pointer `Rc<T>` from Chapter 30's world of smart pointers is the canonical *not-Send* type: its reference count is a plain non-atomic integer, so sending an `Rc` to another thread could race on that count — and the compiler knows, because `Rc` is marked not `Send`.

QUICK CHECK

Define `Send` and `Sync` in one line each, and state the relationship between them. What two bounds does `thread::spawn` put on its closure, and what does each prevent? Why is `Rc` not `Send`? (Answer below the next section.)

32.3 The compiler stops the race: two named errors

Turn the §32.1 race into Rust and it does not compile — in one of two ways, each a diagnostic you should learn to read. The first arises when a closure *borrow*s local data and is handed to a thread that may outlive the borrow. The compiler reports **E0373**, “closure may outlive the current function, but it borrows a value owned by the current function,” and suggests the fix: `move`, to give the closure ownership instead of a borrow (Rust; not compiled here — behavior per the cited sources):

```
let v = vec![1, 2, 3];
let h = thread::spawn(|| println!("{:?}", v)); // borrows v
error[E0373]: closure may outlive the current function, but it borrows `v`,
               which is owned by the current function
help: to force the closure to take ownership, use the `move` keyword
// fix: thread::spawn(move || println!("{:?}", v));
```

The second arises when you try to *share* a value that is not thread-safe. Suppose you reach for `Rc` to share a counter between threads: because `Rc` is not `Send`, the `Send` bound on `spawn`'s closure is unsatisfied, and the compiler reports **E0277** — the same trait-bound error as Chapters 30 and 31 — with a message naming the exact missing capability (per the official Rust error index):


```
let count = Rc::new(Mutex::new(0));
let h = thread::spawn(move || { *count.lock().unwrap() += 1; }); // Rc is !Send
error[E0277]: `Rc<Mutex<i32>>` cannot be sent between threads safely
note: the trait `Send` is not implemented for `Rc<Mutex<i32>>`
// fix: use Arc (atomic reference count), which IS Send + Sync
```

This is what “fearless” means. The §32.1 program that compiled and silently lost updates is, in Rust, a compile error that names the problem before the program runs — either “this borrow could dangle, take ownership” (E0373) or “this type is not safe to send/share across threads” (E0277). And both are ordinary consequences of machinery you already have: E0373 is the lifetime rule of Chapter 29 applied to a closure that escapes to a thread; E0277 is the trait-bound check of Chapter 30 applied to `Send`. The data race did not become easier to avoid by convention — it became impossible to write, because the type of the shared value refuses to cross the boundary unsynchronized.

THE SAME IDEA THREE WAYS: THREAD-SAFETY IS A TRAIT THE COMPILER CHECKS, NOT A DISCIPLINE YOU KEEP

1. In words. `Send` (safe to move to a thread) and `Sync` (safe to share across threads, i.e. `&T` is `Send`) are auto traits the compiler derives structurally. `thread::spawn` requires its closure to be `Send + 'static`, so unsynchronized shared mutation — the §32.1 race — cannot be expressed; attempting it is E0373 (dangling borrow) or E0277 (a non-`Send` type crossing threads).

2. As a picture. Figure 32.1: two threads reaching for one box of data; a “Send/Sync gate” on the boundary passes an `Arc<Mutex<T>>` (thread-safe) and rejects a bare `&mut` or an `Rc` with a red E0277/E0373 stamp — versus C’s boundary, which is wide open and lets the unlocked `counter++` through to a lost update.

3. As numbers/contrast. The C race of §32.1 compiled and printed 2453233 of an expected 4000000 — a different wrong number each run. The Rust equivalent produces *zero* runs, wrong or right: it stops at compile time with E0373 or E0277. A nondeterministic runtime corruption becomes a deterministic message before execution.

(Answer to the §32.1 quick check: it compiled because C has no rule against two threads touching shared mutable state — the data race is invisible to the compiler, only to your discipline. The total is wrong because increments interleave and lose updates, and different each run because the loss depends on the nondeterministic interleaving — the undefined behavior of Chapter 24. The C fix is a mutex around every access to `counter`; its fragility is that you must remember it on every access, and one omission silently reintroduces the race.)

Send/Sync gate the thread boundary; only thread-safe sharing gets through.




Figure 32.1: The Send/Sync gate on the thread boundary. A thread-safe `Arc<Mutex<T>>` passes; a bare `&mut` or a non-Send `Rc` is rejected at compile time with E0373 or E0277. C's boundary has no gate, so the unsynchronized `counter++` of §32.1 goes through to a lost update.

32.4 Sharing safely: `Arc<Mutex<T>>` and channels

Two patterns cover most shared state. The first is **shared mutable state behind a lock**. To share ownership of a value across threads you wrap it in `Arc<T>` — an *atomic* reference-counted pointer, the `Send + Sync` sibling of `Rc` whose count is updated with the atomics of Chapter 22 — and to mutate it you wrap the inner value in `Mutex<T>`. The decisive difference from Chapter 20's C mutex is that Rust's `Mutex` *owns the data it guards*: there is no way to reach the value without calling `lock()`, which returns a guard granting access and releases the lock automatically when the guard is dropped (Rust; not compiled here — behavior per the cited book):

```
let counter = Arc::new(Mutex::new(0));           // shared, lock bundled with the data
let mut handles = vec![];
for _ in 0..10 {
    let c = Arc::clone(&counter);                // bump the atomic refcount; move the clone
    handles.push(thread::spawn(move || {
        let mut n = c.lock().unwrap();           // acquire lock; `n` is a guard, not an int
        *n += 1;                                 // the ONLY way to the data is through the
    }));
}                                                  // guard dropped here -> lock released
for h in handles { h.join().unwrap(); }
// *counter.lock().unwrap() == 10, deterministically — no lost updates
```

Because the integer lives *inside* the `Mutex`, the Chapter 20 mistake of accessing shared data without holding its lock is not a discipline you maintain — it is unrepresentable, since the only handle to the data is the guard `lock()` hands back. Two threads of Chapter 31 also reappear here: `lock()` returns a `Result` because a `Mutex` becomes *poisoned* if a thread panics while holding it (so the `unwrap()` is you deciding a poisoned lock is unrecoverable), and `join()` returns a `Result` whose `Err` carries a panicked thread's payload. The second pattern is **message passing**: rather than share memory, threads communicate over a **channel** (`std::sync::mpsc`), where sending a value *moves* its ownership down the channel (Chapter 27), so the sender can no longer touch it and there is nothing to race on. Receiving returns a `Result` — `Err` when the sending end has hung

up. Both patterns turn Chapter 24's shared-memory hazards into ownership and type problems the compiler checks.

TRAP: UNSAFE IMPL SEND TO SILENCE E0277, AND TREATING MUTEX AS DEADLOCK-PROOF

Two failure modes, one at each end of caution. The first is over-riding the compiler: seeing `error[E0277]: Rc<...> cannot be sent between threads safely` and writing `unsafe impl Send for MyType {}` to make it compile. That does not make the type thread-safe; it *asserts* to the compiler that it is, switching off the exact check that was protecting you and reintroducing the §32.1 race with a signed confession (`unsafe`) attached. The right fix is almost always to use the thread-safe type — `Arc` for `Rc`, an `atomic` for a plain counter — not to claim `Send` you have not earned. The second failure is the opposite complacency: assuming `Arc<Mutex<T>>` makes concurrency “safe” in every sense. It abolishes *data races* — unsynchronized access is impossible — but it does *not* abolish the higher-level hazards of Chapter 20: two locks acquired in opposite orders by two threads still **deadlock**, and holding a `Mutex` across a long or blocking operation still serializes everything (Chapter 26's contention). Rust guarantees you locked the data; it does not guarantee your *lock ordering* or *granularity* is sane. Fearless means no data races, not no concurrency bugs.

SAY IT LIKE A SYSTEMS ENGINEER

Try it first. A teammate says: “‘Fearless concurrency’ is marketing. A mutex is a mutex; Rust just makes me fight the borrow checker across threads, and I can deadlock in Rust same as in C. What does it actually buy me?” Cover the paragraph and respond in four or five sentences, drawing on §32.1–§32.4.

Model answer. “It buys you exactly one thing, and it's the thing that ruins your week: it makes *data races* — unsynchronized access to shared mutable state — a compile error instead of undefined behavior. The C version in §32.1 compiled fine and printed a different wrong number every run; the Rust version won't compile until the shared thing is `Send/Sync`, which in practice means you wrapped it in `Arc<Mutex>` or an `atomic` — so ‘I forgot to lock this one access’ is unrepresentable, because the data lives inside the lock. You're right that it does *not* save you from deadlock or from holding a lock too long — those are lock-ordering and granularity bugs, not data races, and Rust doesn't check them. So the honest claim is narrow but huge: the entire category of nondeterministic memory-corruption concurrency bugs, the ones that pass tests and page you at 3 a.m., is gone at compile time; the logical concurrency bugs are still yours to design around.” Naming precisely what Rust does and does not guarantee — races yes, deadlock no — is the senior register.

(Answer to the §32.2 quick check: *Send* means a value is safe to move to another thread; *Sync* means it is safe to share a reference across threads, i.e. `T: Sync` exactly when `&T: Send`. `thread::spawn` bounds its closure `Send + 'static: Send` so every captured value can move to the new thread, `'static` so the closure holds no borrow that could dangle after the parent returns. `Rc` is not *Send* because its reference count is a non-atomic integer, so sending it across threads could race on that count.)

32.5 What to carry forward

Fearless concurrency is the last application of the pattern this whole part has repeated: name a property in the type system and let the compiler enforce it. A **data race** — the undefined behavior of Chapter 24 — is prevented not by discipline but by two marker traits: **Send** (safe to move to a thread) and **Sync** (safe to share across threads, `&T: Send`), auto-derived structurally, required by `thread::spawn` as `Send + 'static`. Attempting the unsynchronized sharing of §32.1 is a compile error — **E0373** (a borrow that would outlive its owner, fixed by `move`) or **E0277** (a non-`Send` type crossing threads, fixed by the thread-safe type). Safe sharing uses `Arc<Mutex<T>>`, where the `Mutex` owns the data so it cannot be touched without the lock, or **channels**, where sending moves ownership so there is nothing to race on — with poisoning and hang-ups surfacing as the Results of Chapter 31. And the guarantee is honest about its edge: it kills data races, not deadlock or contention.

That completes the Rust part, and with it the argument the last six chapters have been making. Across the C volumes you proved invariants by hand: memory freed exactly once and used only while valid, references that never dangle, shared state accessed only under a lock, no data races. Rust encodes each as something the compiler checks — **ownership and moves** (Chapter 27), **borrowing** (Chapter 28), **lifetimes** (Chapter 29), **traits and generics** (Chapter 30), **error handling** (Chapter 31), and now **data-race freedom** — all with the same machine underneath and, after monomorphization and lifetime erasure, the same generated code. What Rust does *not* give you is speed: a memory-safe, race-free program can still be slow, cache-hostile, and blocked on I/O. The next part turns to that. It opens with the discipline the whole book has implied and never stated outright — **measure before you optimize** — because the cost model of Chapter 5 is only useful if you can see where the time actually goes. Safety was the compiler's job; performance is yours, and it starts with a profiler, not a guess.

CAN YOU... WITHOUT LOOKING?

- Explain why the §32.1 C program compiled yet printed a wrong, different-each-run total.
- Define `Send` and `Sync` and state the relationship between them.
- Say what bounds `thread::spawn` puts on its closure and what each prevents.
- Name the two compile errors that stop the race (E0373, E0277) and the fix for each.
- Explain how `Mutex<T>` owning its data makes “access without the lock” unrepresentable.
- State precisely what fearless concurrency does and does not guarantee (races vs. deadlock).

RETRIEVAL — CLOSED BOOK

Cover the chapter. Answer from memory; rate confidence first.

1. **R1.** Why did the §32.1 program compile, and why is its output wrong and nondeterministic? (§32.1)
2. **R2.** Define `Send` and `Sync`; what does `thread::spawn` require, and why is `Rc` not `Send`? (§32.2)
3. **R3.** What do `E0373` and `E0277` each report in the concurrency setting, and what fixes each? (§32.3)
4. **R4.** How does `Arc<Mutex<T>>` make unsynchronized access impossible, and where does `Result` appear? (§32.4)
5. **R5.** What does fearless concurrency guarantee, and what class of concurrency bug does it *not* prevent? (§32.4)

FURTHER READING

- **Klabnik & Nichols, *The Rust Programming Language*, 2nd ed. (No Starch Press, 2023), ch. 16 “Fearless Concurrency” (also free at doc.rust-lang.org/book).** The primary source for `thread::spawn`, the move closure, `Arc<Mutex<T>>`, channels, and `Send/Sync` of §32.2–§32.4; ch. 15 covers `Rc`/smart pointers. Authors, edition, and chapter numbers verified via the publisher and the online book; not compiled here.
- **The Rustonomicon, “Send and Sync” (doc.rust-lang.org/nomicon/send-and-sync.html).** The authoritative account of the two marker traits used in §32.2: what each means, why they are auto traits derived structurally, and why `Rc` opts out. Reference for the exact semantics; verify against the current edition.
- **The Rust Error Index, entries `E0373` and `E0277` (doc.rust-lang.org/error_codes).** The official meanings used in §32.3: `E0373` “closure may outlive the current function... use the move keyword,” and `E0277` for an unsatisfied `Send` bound (“cannot be sent between threads safely”). Both verified against the current error index.
- **Jung, Jourdan, Krebbers & Dreyer, “RustBelt: Securing the Foundations of the Rust Programming Language,” *Proc. ACM on Programming Languages* 2(POPL), Article 66 (2018), DOI 10.1145/3158154.** The formal proof that Rust's type system — including `Send/Sync` and the aliasing discipline — actually guarantees data-race freedom. First author Ralf Jung; verified via the ACM Digital Library.

Exercises

32.1 WARM-UP In one sentence each, define `Send` and `Sync`, and state the relationship between them.

32.2 WARM-UP What is a data race, and what is its status in C (per Chapter 24)? What is its status in safe Rust?

32.3 WARM-UP What does the `move` keyword do to a closure passed to `thread::spawn`, and which error does its absence often cause?

32.4 **CORE** Take the C of §32.1. (a) Why did it compile despite the data race? (b) Explain, in read-modify-write terms, why the total is both too low and different each run. (c) State the C fix and one way it is fragile that Rust's approach is not.

32.5 **CORE** Explain how `thread::spawn`'s bound `F: Send + 'static` prevents the §32.1 race. (a) What does the `Send` part rule out? (b) What does the `'static` part rule out, and which chapter's concept is that? (c) Why does trying to share an `Rc` across threads violate the `Send` part specifically?

32.6 **CORE** Contrast E0373 and E0277 as they appear in concurrent code. (a) What does each report, and what borrowing/ trait concept from earlier chapters is each an instance of? (b) Give the standard fix for each. (c) Why is `unsafe impl Send` the wrong fix for E0277 here?

32.7 **COST** Weigh `Arc<Mutex<T>>` against message passing for a shared counter updated by many threads. (a) What does the `Mutex` approach cost per update, and what Chapter 26 hazard grows as thread count rises? (b) What does the channel approach cost, and what does moving ownership down the channel buy you? (c) For a high-contention counter specifically, name a third option from Chapter 22 that avoids a lock entirely, and say when you would still prefer the `Mutex`.

32.8 **FADED** *Worked, with the last step for you.* A newcomer shares a counter with `let c = Rc::new(Mutex::new(0));` and spawns threads that each do `let cc = c.clone();`
`thread::spawn(move || *cc.lock().unwrap() += 1);`, getting error[E0277]: ``Rc<Mutex<i32>>`` cannot be sent between threads safely.

Step 1 (given): `thread::spawn` requires the closure — and thus everything it captures — to be `Send`.

Step 2 (given): `Rc` is not `Send` because its reference count is non-atomic, so the bound fails: E0277.

Step 3 (given): the fix is a shared pointer whose count is atomic and which is therefore `Send + Sync`.

Step 4 (you): Name the replacement type, rewrite the two changed lines, and explain why this is a real fix whereas `unsafe impl Send` for `...` would only silence the check.

32.9 **MIXED** For each, say whether it is prevented by safe Rust's type system, and if so how (or if not, why not): (a) two threads incrementing a shared `i32` with no synchronization; (b) two threads acquiring locks A then B vs. B then A (deadlock); (c) sending an `Rc` to another thread; (d) holding a `Mutex` guard across a slow network call (serializing everyone).

32.10 **MIXED** Classify each as *true* or *false* and fix the false ones: (a) “`Send` and `Sync` have methods you implement.” (b) “A Rust `Mutex` can be unlocked-and-accessed like C's, with the data separate from the lock.” (c) “`Arc` is just `Rc` with an atomic reference count.” (d) “Fearless concurrency prevents deadlock.” (e) “Sending a value down a channel copies it and leaves the sender a usable copy.”

32.11 **STRETCH** “A guarantee in Rust is usually a trait.” Tie the whole part together. (a) Show how `Send/Sync` lift the Chapter 28 borrow rule (aliasing XOR mutation) across the thread boundary — what would sharing a `&mut` between two threads violate? (b) Explain why the concurrency errors reuse the *same* error code (`E0277`) as an unmet bound in Chapter 30 and a misused `?` in Chapter 31. (c) In two sentences, state the through-line of Chapters 27–32: what did the C volumes make you prove by hand, and what does Rust do with each such proof?

32.12 **LAB** On your own machine (a C toolchain suffices; a Rust toolchain is ideal): (a) Compile the §32.1 `race.c` at `-O0` and run it several times; record the totals and confirm they are below `4000000` and vary. Then compile at `-O2` and note whether the race still manifests, and hypothesize why the optimizer changed the outcome. (b) Add a `pthread_mutex_t` around `counter++` and confirm the total becomes exactly `4000000`. (c) If you have `rustc`, write the `Arc<Mutex<i32>>` version from §32.4 and confirm it is deterministic; then try the `Rc` version and observe `E0277`. Without `rustc`, explain from this chapter which line each error points at. Note your toolchain versions.

Performance Engineering

Measure before you optimize: profiling, benchmarking, cache-aware and data-oriented design, SIMD, and reasoning about where time actually goes.

Chapter 33 — Measure First: Profiling and Benchmarking

Before you start: Chapter 5 (mechanical sympathy and the cost model — this part makes that model operational, by measuring where the costs land), Chapter 26 (Amdahl's law — the ceiling on speeding up a part, and why the dominant cost is the only one worth attacking), Chapter 10 (the optimizing compiler — dead-code elimination will silently wreck a naive benchmark), Chapter 3 (memory hierarchy and caches — one of the main places the time actually goes). · This chapter opens the performance part with the discipline the whole book has implied: *measure before you optimize*, because a guess about where time goes is almost always wrong.

BEFORE YOU READ ON

From memory, reaching back: (1) Chapter 26 — state Amdahl's law in one line; if a function is 5% of total runtime and you make it infinitely fast, what is the most you can speed up the whole program? (2) Chapter 10 — what is the optimizing compiler permitted to do to a computation whose result is never used? (3) *Predict*: you write a loop that sums a billion products into a variable, never read the variable, and time it at -O2; what elapsed time do you expect? And if you run that timing exactly once, is the number you get representative? Write your guesses; §33.2 and §33.3 answer all three.

A note on this part before we begin. The previous five parts built up, from the hardware to the type system, a *model* of what a program costs to run. This part is about the other half of engineering with that model: finding out where the cost actually falls, which is almost never where intuition says. The cardinal rule is **measure before you optimize**. Programmers are notoriously bad at guessing bottlenecks — you spend an afternoon hand-tuning a function that turns out to be 2% of runtime while the real cost sits in an allocation or a cache-missing loop you never suspected. Donald Knuth's much-abused line — “premature optimization is the root of all evil” — is precisely this point: he argued that we should forgo optimization in the non-critical 97% of code and instead find, by measurement, the critical few percent where effort pays (Knuth, 1974). So this chapter is the toolset for that finding. First, why measuring is not optional and how Amdahl's law decides *what* to measure (§33.1). Then the single most common way a benchmark lies to you — the optimizer deleting the work you meant to time — shown in real C where the same loop reports 0.0000 s or 1.56 s depending only on whether its result is used (§33.2). Then how to time correctly: which clock, and why you run many trials and report the minimum rather than one number (§33.3). Then **profiling** — sampling vs. instrumentation, and reading a profile as a map of where the time went, so Amdahl points your effort (§33.4). We close by making the measure-optimize loop explicit and turning toward the specific costs the rest of the part attacks (§33.5).

On the code in this part. Unlike the Rust part, the programs here are C, and they were compiled with gcc 11.4.0 and run in this environment, so their *output is real*. Absolute times, however, are hardware- and load-specific: the numbers below are illustrative of the *effect* (a deleted loop, run-to-run variance, a lopsided phase split), not benchmarks of any

particular machine, and your own numbers will differ. Where a tool is named that is *not* installed here (`perf`, `gprof`, `valgrind`), it is discussed from its documentation and cited, not run.

33.1 Measure before you optimize

The discipline has two halves, and Amdahl's law (Chapter 26) supplies the second. The first half is: do not optimize on a hunch, because hunches about performance are wrong at a rate that would be comical if it were not so expensive. The compiler, the cache, the allocator, and the branch predictor all move the real cost somewhere other than the arithmetic you can see in the source. The second half is: even once you know a function is slow, speeding it up is only worth it in proportion to its *share of total runtime*. This is Amdahl's law read as a budgeting rule. If a function is 5% of runtime, then making it infinitely fast — not twice as fast, *infinitely* — speeds the whole program by at most 5% (from 100 units to 95). If another is 90% of runtime, a mere 2× speedup on it saves 45 units, nine times the entire budget of the first. So the order is fixed: *find the dominant cost first, then optimize it*, and measurement is the only way to know which is which. Optimizing before measuring is not just wasted effort; it is effort spent, by the arithmetic of Amdahl, on the part that mathematically cannot matter. Everything else in this chapter serves this: benchmarking to compare two implementations honestly (§33.2–§33.3), and profiling to locate the dominant cost in the first place (§33.4).

QUICK CHECK

State the two halves of “measure before you optimize.” Using Amdahl's law, how much can you speed up a program by perfectly optimizing a part that is 5% of its runtime? Why does that make locating the *dominant* cost the first job? (Answer after the next section.)

33.2 The benchmark that lies: dead-code elimination

The most common benchmarking bug is that the compiler deletes the very work you are trying to time. From Chapter 10: if a computation's result is never observed, the optimizer is free to remove the computation entirely — it is not required to run code whose output no program can see. A microbenchmark that computes into a variable and never uses it hands the optimizer exactly that license. Here is the same loop compiled two ways, differing only in whether the result is consumed — compiled and run in this environment (gcc 11.4.0, `-O2`):

```
double s = 0.0;
for (long i = 0; i < 1000000000L; i++) s += (double)(i & 1023) * 0.5;
// version A: never use s.    version B: consume s (print it / store to volatile).

$ gcc -O2 bench.c -o nosink && ./nosink
no sink    : 0.0000 s                                # the loop was deleted entirely
$ gcc -O2 -DSINK bench.c -o sink && ./sink
with sink  : 1.5592 s (checksum 255749934464.0)      # the loop actually ran
```

Version A reports `0.0000 s` not because a billion iterations are free but because they never happened: with `s` unused, the compiler proved the loop had no observable effect and removed it, exactly as Chapter 10 said it may. A newcomer reads “0.0000 s” as “my code is blazing fast” and ships a benchmark that measures nothing. Version B prints the checksum, so `s` is observed, so the loop must run, and now you see its real cost — 1.56 s here. The fix is a rule: **make the result observable**, by printing a checksum, writing to a `volatile`, or feeding it to a compiler-opaque sink, so the optimizer cannot elide the work. Two companion rules follow. Benchmark the build you will *ship* — measure at `-O2`, not `-O0`, since the debug build's costs are not the release build's. And be sure the input is not a compile-time constant the optimizer can fold away. The through-line is that a benchmark is only honest if the work it claims to time is work the compiler was actually forced to do.

QUICK CHECK

Why did version A report `0.0000 s` — what did the compiler do, and by which Chapter 10 rule? What single change turned the loop back on? Give two further rules for a benchmark that is not lying to you. (Answer below the communicate box.)

33.3 Timing correctly: which clock, and how many runs

Suppose the loop is safe from the optimizer. Two questions remain: *which clock* you read, and *how many times*. For elapsed time you want a monotonic wall clock that never jumps — on Linux, `clock_gettime(CLOCK_MONOTONIC, &t)`, which is immune to wall-time adjustments (NTP, daylight saving) that can make a naive timer run backwards. (Wall-clock time and CPU time differ: a program blocked on I/O consumes little CPU while wall time passes, so choose the one that matches the question you are asking.) The second question matters more than beginners expect: *a single measurement is nearly meaningless*, because runtime varies from run to run due to the OS scheduler, cache and TLB state, frequency scaling, and other tenants on the machine. You must run the workload several times. Compiled and run here (gcc 11.4.0, the same kernel run five times):

```
for (int r = 0; r < 5; r++) { t0 = now(); sink += work(3000000000L); t1 = now();
                             printf("trial %d: %.4f s\n", r+1, t1 - t0); }

$ gcc -O2 trials.c -o trials && ./trials
trial 1: 0.6824 s
trial 2: 0.7081 s
trial 3: 0.6694 s
trial 4: 0.8198 s      <- an outlier: the OS scheduled something else here
trial 5: 0.6743 s
min over 5 trials: 0.6694 s
```

The five runs of identical work span 0.6694 s to 0.8198 s — a 22% spread, from nothing the program did. Reporting any single one of them, especially the first (cold caches) or the fourth (an interruption), would misrepresent the code. The usual practice is to run many trials and report the **minimum** (the run least disturbed by outside interference, closest to the code's intrinsic cost) or the median, and to state the spread rather than hide it behind one number; a first, warm-up run is often discarded because it pays cold-cache and page-fault costs the steady state will not. The mean is a poor summary here

because it is dragged upward by one-directional outliers like trial 4 — interference only ever makes a run *slower*, never faster, so the minimum is the more principled floor. The discipline, then: monotonic clock, warm up, many trials, report the minimum and the spread. A benchmark that reports one number from one run is a benchmark you cannot trust.

THE SAME IDEA THREE WAYS: A MEASUREMENT IS ONLY AS HONEST AS ITS METHOD

1. In words. Optimize only after measuring, and only the dominant cost (Amdahl). Make the timed work observable so the optimizer cannot delete it; time it at the shipping optimization level with a monotonic clock; run many trials and report the minimum and spread, not one number. Then profile to find the dominant cost before you touch anything.

2. As a picture. Figure 33.1: a pipeline — *profile* (find the hot spot) → *benchmark* the hot spot honestly (observable result, -02, many trials) → *optimize* → *re-measure* → stop when good enough — with two red “lies” branching off: a deleted loop reporting 0.0000 s, and a single run reporting an outlier.

3. As numbers/contrast. The same billion-iteration loop reported 0.0000 s (deleted) vs. 1.5592 s (real) depending only on whether its result was used; the same kernel run five times reported 0.6694–0.8198 s, a 22% spread. Two different methods, two different “answers” for identical code — which is why the method, not the stopwatch, is the measurement.

(Answer to the §33.1 quick check: the two halves are (1) do not optimize on a hunch — measure, because bottlenecks are rarely where you guess — and (2) optimize only in proportion to a part's share of runtime, per Amdahl's law. Perfectly optimizing a part that is 5% of runtime speeds the whole program by at most 5% (100 units become 95), even if that part becomes infinitely fast. That is why finding the dominant cost is the first job: effort on a small part cannot pay off no matter how good the optimization, so you must locate the part that dominates before optimizing anything.)

The measure-optimize loop, and the two ways a benchmark lies.




Figure 33.1: The measure-optimize loop — profile to find the hot spot, benchmark it honestly, optimize, re-measure, and stop when good enough — with the two classic benchmark lies: a deleted loop reporting 0.0000 s (§33.2) and a single run reporting an outlier (§33.3).

33.4 Profiling: where the time actually goes

Benchmarking compares two versions of a known piece of code; **profiling** answers the prior question — *which* piece of code to look at. A profiler attributes runtime to the

functions (and lines) that spend it, so Amdahl's dominant cost stops being a guess. There are two families. An **instrumenting** profiler (the classic `gprof`) inserts counting and timing code at function entry and exit; it is exact about call counts but perturbs the very timing it measures and can distort the picture for tiny, hot functions (Graham, Kessler & McKusick, 1982). A **sampling** profiler (Linux `perf`, and most modern tools) instead interrupts the running program many times per second and records where the program counter is; over enough samples this builds a statistical map of where time is spent, with negligible overhead and no recompilation (Gregg, 2020). The output is often drawn as a **flame graph** — stacked bars whose width is time spent — so the widest box is, quite literally, where to optimize. To make the payoff concrete, here is a two-phase program timed by hand, the poor engineer's profile, compiled and run here (gcc 11.4.0, `-O2`):

```
t0 = now(); parse(...); t1 = now(); // a setup phase
      compute(...); t2 = now(); // the main work

$ gcc -O2 profile.c -o profile && ./profile
parse   : 0.1577 s ( 6.4% of runtime)
compute : 2.2904 s (93.6% of runtime)
```

The profile is unambiguous: `compute` is 93.6% of runtime, `parse` only 6.4%. Amdahl now dictates strategy. Every hour spent optimizing `parse` buys at most a 6.4% whole-program improvement even if you make it vanish; a modest 30% cut to `compute` saves $0.30 \times 93.6\% = 28\%$ of the whole. A programmer who *guessed* — who found `parse`'s string handling ugly and spent the afternoon there — would have optimized the part that cannot matter, which is exactly the failure §33.1 warned of. This is why the loop is *profile first*: the profiler replaces the guess with a measurement, points you at the 93.6% box, and only then do you benchmark candidate improvements to it honestly (§33.2– §33.3). A real profiler does this across hundreds of functions you would never think to time by hand — but the logic is the one above, scaled up: measure where the time is, attack the widest box, re-measure, stop.

TRAP: OPTIMIZING WITHOUT A PROFILE, AND TRUSTING A BENCHMARK THAT MEASURED NOTHING

Two mistakes, both the “measure first” rule violated from opposite ends. The first is optimizing on intuition: rewriting the function that *looks* expensive — a clever-looking loop, some string formatting — without a profile showing it is expensive. Nine times in ten it is not the bottleneck, and you have added complexity and risk to code that was 3% of runtime while the real 90% sits untouched; worse, you now believe the problem is “handled.” Always get a profile first, then optimize the box the profile makes wide. The second mistake is trusting a benchmark that never ran the work: the 0.0000 s of §33.2 (dead-code elimination), a loop over a constant the compiler folded, a measurement taken at `-O0` that has nothing to do with the shipped build, or a single run that happened to be an outlier. Both mistakes feel like measurement — you ran something, you got a number — which is what makes them dangerous. The number is only meaningful if the method behind it (a profile to choose the target, an observable result, the shipping optimization level, many trials) is sound. A confident optimization built on a lying benchmark is worse than no optimization, because you will defend it.

SAY IT LIKE A SYSTEMS ENGINEER

Try it first. A teammate says: “I profiled nothing, but I’m sure the JSON parsing is our bottleneck — it just *looks* slow. I rewrote it with a hand-tuned loop and my microbenchmark says it’s 100× faster. Ship it?” Cover the paragraph and respond in four or five sentences, drawing on §33.1–§33.4.

Model answer. “Two problems, and both are ‘measure first.’ First, ‘looks slow’ isn’t a measurement — we need a profile of the real workload showing parsing is actually a big share of runtime, because if it’s 5% of the total, Amdahl says even your 100× win is at most a 5% whole-program improvement, and you’ve added a hand-tuned loop we now have to maintain. Second, a 100× microbenchmark speedup is exactly the number I distrust: it usually means the compiler deleted the work because the result was unused, or you benchmarked a constant, or you compared a -O0 build to a -O2 one — make the result observable, benchmark at the shipping optimization level, and run it many times. So before we ship: profile the real path, confirm parsing is the dominant cost, then re-run your benchmark honestly and show the whole-program time moved. If the profile says parsing is 3%, the right call is to revert and go look at the 90% box instead.” Refusing to ship an optimization without a profile that justifies the target and a benchmark that survives scrutiny is the senior register.

(Answer to the §33.2 quick check: version A reported 0.0000 s because the compiler deleted the loop — with s never used, the loop had no observable effect, and by the Chapter 10 rule the optimizer may remove computations whose results are never observed. Making s observable — printing it as a checksum, or writing it to a volatile/opaque sink — turns the loop back on. Two further rules: benchmark at the optimization level you will ship (-O2, not -O0), and ensure the input is not a compile-time constant the optimizer can fold away.)

33.5 What to carry forward

Performance engineering starts with a refusal: do not optimize what you have not measured. The governing rule is **measure before you optimize**, and Amdahl’s law (Chapter 26) makes it quantitative — a part that is 5% of runtime caps its own whole-program payoff at 5%, so the *dominant* cost is the only target worth your effort, and only a **profile** tells you which part that is. Profiling comes in two forms — **instrumenting** (exact counts, perturbs timing, e.g. gprof) and **sampling** (statistical, low overhead, e.g. perf, drawn as flame graphs) — and its job is to replace a guess with the widest box. **Benchmarking** then compares candidate fixes, but only honestly: make the timed result observable so the optimizer cannot delete it (the 0.0000 s vs. 1.5592 s of §33.2), time at the shipping optimization level, use a monotonic clock, and run many trials reporting the minimum and the spread (the 0.6694–0.8198 s of §33.3), not one number. Put together, these form the loop: profile, benchmark, optimize, re-measure, and stop when good enough.

What the profile does *not* tell you is *why* the widest box is wide — and that is the rest of this part. When `compute` is 93.6% of runtime, the question becomes where *its* time goes, and the answer is usually not the instruction count you can read in the source but the costs the earlier volumes named: cache misses (Chapter 3), branch mispredictions,

memory bandwidth, allocation, the layout of your data. The next chapters take the dominant cost the profiler hands you and open it up — **cache-aware and data-oriented design**, then wider techniques like **SIMD** — always returning to this chapter's loop to prove that a change actually moved the number. Safety was the compiler's job in the last part; speed is yours, and it begins not with a clever rewrite but with a measurement honest enough to bet on.

CAN YOU... WITHOUT LOOKING?

- State “measure before you optimize” and use Amdahl's law to say why the dominant cost is the only target.
- Explain why a loop whose result is unused can report 0.0000 s, and the Chapter 10 rule behind it.
- Give the fix (observable result) and two more rules for an honest benchmark.
- Say which clock to use and why one run is untrustworthy; explain why you report the minimum.
- Distinguish sampling from instrumenting profilers, with a cost of each.
- State the measure-optimize loop end to end.

RETRIEVAL — CLOSED BOOK

Cover the chapter. Answer from memory; rate confidence first.

1. **R1.** What are the two halves of “measure before you optimize,” and how does Amdahl's law make the second quantitative? (§33.1)
2. **R2.** Why can a microbenchmark report 0.0000 s, and what makes the work run again? (§33.2)
3. **R3.** Why is a single timing untrustworthy, and why report the minimum of many trials rather than the mean? (§33.3)
4. **R4.** What is the difference between sampling and instrumenting profilers, and what does each cost? (§33.4)
5. **R5.** Given a profile showing one function at 90%+ of runtime, what does Amdahl say to do, and why? (§33.1, §33.4)

FURTHER READING

- **Brendan Gregg, *Systems Performance: Enterprise and the Cloud*, 2nd ed. (Addison-Wesley, 2020), ISBN 978-0-13-682015-4.** The standard modern reference for performance methodology, sampling profilers (perf, BPF-based tools), and flame graphs of §33.4; the source for the measure-first methodology of this chapter. Author, edition, and year verified via the publisher.
- **Donald E. Knuth, “Structured Programming with go to Statements,” *ACM Computing Surveys* 6(4), 261-301 (1974), DOI 10.1145/356635.356640.** The origin (p. 268) of “premature optimization is the root of all evil” quoted in §33.1 — and the fuller argument that measurement should find the critical few percent of code worth tuning. First author Donald Knuth; DOI verified via the ACM Digital Library.
- **Susan L. Graham, Peter B. Kessler & Marshall K. McKusick, “gprof: a Call Graph Execution Profiler,” *ACM SIGPLAN Notices* 17(6), 120-126 (1982), DOI 10.1145/872726.806987.** The classic instrumenting profiler of §33.4 — call-graph attribution of runtime, and the source of its counts-vs-perturbation trade-off. First author Susan Graham; DOI verified via the ACM Digital Library.
- **The Linux perf tool (perf_events) documentation (perfwiki.github.io, and man perf).** Reference for the sampling profiler named in §33.4 — how it counts hardware events and samples the program counter with low overhead. Not installed in this environment, so discussed from its documentation, not run; verify against your distribution's version.

Exercises

- 33.1** **WARM-UP** State the rule “measure before you optimize” in one sentence, and say what a profiler is for versus what a benchmark is for.
- 33.2** **WARM-UP** Why can a microbenchmark of a loop report 0.0000 s even though the loop appears to do real work?
- 33.3** **WARM-UP** Why do you run a benchmark many times and report the minimum (or median) rather than a single run's time?
- 33.4** **CORE** Take the two-phase profile of §33.4 (parse 6.4%, compute 93.6%). (a) Using Amdahl's law, what is the most you can improve total runtime by perfectly optimizing parse? (b) What whole-program improvement does a 30% cut to compute give? (c) A colleague, without profiling, spends a day optimizing parse because its code “looks slow.” Explain precisely why this violates §33.1.
- 33.5** **CORE** Explain the dead-code-elimination benchmark bug of §33.2. (a) Under what condition is the compiler allowed to delete the timed loop, and which Chapter 10 principle is that? (b) Name two distinct ways to make the result observable so the work runs. (c) Why must you also benchmark at the optimization level you ship rather than at -O0?

33.6 **CORE** Contrast sampling and instrumenting profilers. (a) How does each obtain its data? (b) What is the cost or distortion of each, especially for many tiny, hot functions? (c) For a first look at a large program you do not want to recompile, which would you reach for, and why?

33.7 **COST** Reason about measurement error. (a) The five trials of §33.3 spanned 0.6694–0.8198 s for identical work; list three sources of that run-to-run variation. (b) Why is the *minimum* a more principled estimate of intrinsic cost than the mean, given that interference only ever slows a run? (c) When would the median be a better summary than the minimum, and what should you always report alongside either?

33.8 **FADED** *Worked, with the last step for you.* A developer benchmarks a hashing routine: `t0 = now(); for (i=0;i<N;i++) hash(key); t1 = now();` and reports 0.0000 s, concluding the hash is “essentially free.”

Step 1 (given): the return value of `hash(key)` is never used, so the calls have no observable effect.

Step 2 (given): at -O2 the compiler may therefore remove the entire loop (Chapter 10), so nothing was timed.

Step 3 (given): the input `key` is also loop-invariant, which the optimizer can hoist or fold, compounding the problem.

Step 4 (you): Rewrite the benchmark so it honestly times the hash — make the result observable and vary the input across iterations — and state what you would additionally do (clock, trials) before trusting the number.

33.9 **MIXED** For each situation, say whether you would *profile*, *benchmark*, or *neither yet*, and why: (a) a service is “slow” but you have no idea which part; (b) you have two candidate implementations of the confirmed hot function and must pick one; (c) a teammate proposes rewriting a module that is 2% of runtime; (d) you changed the hot function and want to know if total runtime actually dropped.

33.10 **MIXED** Classify each as *true* or *false* and fix the false ones: (a) “A benchmark reporting 0.0000 s means the code is extremely fast.” (b) “Benchmarking at -O0 tells you the shipped build's performance.” (c) “A sampling profiler requires recompiling the program.” (d) “Optimizing a part that is 4% of runtime can, at best, speed the program by 4%.” (e) “One timing run is enough if the machine is otherwise idle.”

33.11 **STRETCH** Tie this chapter to the cost model. (a) The profiler in §33.4 says `compute` is 93.6% of runtime but not *why*; list three costs from Chapters 3–5 that could be the real reason a loop is slow even at a low instruction count. (b) Explain how you would use the §33.3 timing method to test whether a proposed cache-friendlier layout actually helps — what exactly must you hold constant, and what must you make observable? (c) Argue why “profile → hypothesize a cause → change one thing → re-measure” is the scientific method applied to performance.

33.12 **LAB** On your own machine (a C toolchain suffices): (a) Reproduce §33.2 — write a loop that sums into a variable, compile at `-O2` without using the variable and time it, then print the variable and time it again; report both times and explain the difference. (b) Reproduce §33.3 — run an honest version five times with `clock_gettime(CLOCK_MONOTONIC, ...)`, report the five times and the minimum, and compute the spread. (c) If your platform has `perf` (or `gprof`), profile a small two-function program like §33.4's and confirm the profiler agrees with hand timing about which function dominates; if not, do the hand timing and state which tool you would use and what it would add. Note your toolchain versions and CPU.

Chapter 34 — Cache-Aware Design: Programming for Locality

Before you start: Chapter 3 (the memory hierarchy and caches — this chapter turns that model into code you can time), Chapter 5 (mechanical sympathy — writing with the grain of the machine), Chapter 4 (virtual memory and the TLB — the other table a strided access thrashes), Chapter 33 (measure first — every number here is measured, and you should distrust the ones that are not). · The profiler of Chapter 33 tells you *where* the time goes; this chapter opens the most common answer — the loop is waiting on memory — and shows the code changes that make it wait less.

BEFORE YOU READ ON

From memory, reaching back: (1) Chapter 3 — roughly how many times slower is a main-memory access than an L1 cache hit, order of magnitude? (2) Chapter 3 — what is a *cache line*, and what is the unit the hardware moves between memory and cache? (3) *Predict*: you have an 8192×8192 matrix of double stored row by row, and you sum every element two ways — walking along each row, then walking down each column. Same arithmetic, same element count. How many times slower do you expect the column walk to be: about the same, 2×, or something far larger? Write your guess; §34.2 and §34.3 answer all three.

Chapter 3 gave you the memory hierarchy as a set of facts: registers, then L1, L2, L3 caches, then DRAM, each larger and an order of magnitude slower than the last. This chapter makes those facts *actionable*. The RAM model of Chapter 1 says every memory access costs the same; the real machine says an access that hits L1 is perhaps a hundred times cheaper than one that goes to DRAM, and which one you get is decided by **locality** — whether your accesses reuse recent data (temporal) and touch neighbors of recent data (spatial). The striking consequence is that two loops doing the *identical* arithmetic on the *identical* data can differ in wall-clock time by more than an order of magnitude, decided entirely by the order in which they touch memory. We build this up in five steps: the hierarchy as a cost model, with this machine's real cache sizes (§34.1); spatial locality and why the *cache line*, not the byte, is the unit you pay for (§34.2); the canonical demonstration — row-major versus column-major traversal, measured at 19× here (§34.3); temporal locality, working sets, and blocking to keep reuse in cache (§34.4); and what to carry into the layout and vectorization chapters that follow (§34.5).

On the code in this chapter. As in Chapter 33 the programs are C, compiled with gcc 11.4.0 and run in this environment on an Intel Core i7-12700H, timed with `clock_gettime(CLOCK_MONOTONIC)` and reported as the minimum of several trials (Chapter 33's method). The cache sizes quoted are read from this machine's `/sys/devices/system/cpu/.../cache`, so they are real for this CPU; the *times* are illustrative of the effect and hardware-specific — your machine will differ. A hardware-counter profiler (`perf`) is *not* installed here, so cache-miss counts are not quoted; the effects below are shown through wall-clock timing, which is enough to see an order-of-magnitude gap, and this is stated wherever it matters.

34.1 The hierarchy as a cost model, in real numbers

The reason locality governs performance is that the levels of the hierarchy differ not by a little but by orders of magnitude, and the fast levels are small. On the machine these programs run on — an Intel Core i7-12700H, values read from `/sys` — each core has a 48 KiB L1 data cache and a 1.25 MiB L2, and all cores share a 24 MiB L3; the line size at every level is 64 bytes. Representative published access latencies (order-of-magnitude teaching figures, not measured here; see Drepper 2007 and CS:APP) run about 4 cycles to hit L1, ~12 to L2, ~40 to L3, and *hundreds* of cycles — often quoted near 100 ns — to reach DRAM. Two facts follow that decide everything else. First, the gap is enormous: a DRAM access can cost more than a hundred L1 hits, so a loop that misses cache on most accesses is not a little slower; it is *stalled*, the CPU idling while data crawls in. Second, the fast memory is *small* — 48 KiB of L1, not megabytes — so you cannot keep everything hot; you must arrange for the data you are about to use to already be near the core. That arrangement is what “cache-aware” means, and the hardware gives you two levers, spatial and temporal locality, which the next sections make concrete.

QUICK CHECK

Name the levels of the hierarchy from fastest to slowest and say, to an order of magnitude, how much slower a DRAM access is than an L1 hit. Why does the smallness of L1 (tens of KiB) mean you cannot simply “keep the data in cache” and must instead control locality? (Answer after the next section.)

34.2 Spatial locality: you pay by the cache line, not the byte

The hardware never fetches a single byte or a single `int` from memory. It fetches a whole **cache line** — 64 bytes on this machine — and caches that. So when you read one `int`, the fifteen `ints` next to it in the same line come along for free; if your next access is one of them, it hits cache. This is **spatial locality**: code that touches memory in contiguous, small strides gets many accesses per line fetched, while code that jumps far between accesses pays a full line fetch to use just one element and throws the other 60 bytes away. You can see the line size directly by sweeping the stride of a walk over an array far larger than L3 and measuring time per access — compiled and run here (gcc 11.4.0, `-O2`, a 256 MiB `int` array, a fixed 2^{24} accesses at each stride):

```
for (long k=0, idx=0; k < touches; k++) { s += a[idx]; idx += stride; if (idx >= N) idx = 0; }

$ gcc -O2 stride.c -o stride && ./stride
stride 1 ints ( 4 B): 2.93 ns/access
stride 2 ints ( 8 B): 2.71 ns/access
stride 4 ints (16 B): 3.53 ns/access
stride 8 ints (32 B): 5.82 ns/access
stride 16 ints (64 B): 10.61 ns/access  <- one access per 64-byte line
stride 32 ints (128 B): 18.03 ns/access
stride 64 ints (256 B): 20.91 ns/access
```

Read the shape, not the exact figures. At stride 1 each 64-byte line serves 16 accesses, so the line-fetch cost is amortized 16 ways and each access is cheap (~ 2.9 ns here, with the hardware prefetcher helping). As the stride grows the number of useful accesses per line falls — 8, then 4, then 2 — and the per-access cost climbs, because the same line fetch now serves fewer of them. At a stride of 16 ints — exactly 64 bytes, one cache line — every access lands in a fresh line and uses only 4 of its 64 bytes; the cost has jumped to ~ 10.6 ns, roughly $3.6\times$ the stride-1 cost. The knee at 64 bytes is the cache line making itself visible in wall-clock time. Beyond it the cost keeps rising as larger strides also defeat the prefetcher and stress the TLB (Chapter 4), but the lesson is already fixed: **the cache line is the unit you pay for**, so packing the data a loop uses into as few lines as possible — contiguous, small-stride, no wasted bytes per line — is the first and cheapest performance lever you have.

QUICK CHECK

Why does a stride-1 walk cost far less per access than a stride-16-int walk over the same array, when both read the same number of elements? What is special about a stride of 64 bytes on this machine? (Answer below the communicate box.)

34.3 The canonical case: row-major versus column-major

Now the demonstration that every performance course reaches for, because it isolates spatial locality with nothing else changed. C stores a two-dimensional array in **row-major** order: `A[i][j]` and `A[i][j+1]` are adjacent in memory, while `A[i][j]` and `A[i+1][j]` are a full row apart. Summing the matrix by walking along rows therefore strides by one element (stay in the line, spatial locality is perfect); walking down columns strides by a whole row — 8192 doubles, 64 KiB — so every single access lands in a new line, uses 8 of its 64 bytes, and evicts a line you will need again later. Same elements, same additions, opposite locality. Compiled and run here (gcc 11.4.0, `-O2`, an 8192×8192 double matrix, 512 MiB, minimum of 5 trials):

```
for (i) for (j) s += A[i*N + j]; // row-major: stride 1 element
for (j) for (i) s += A[i*N + j]; // col-major: stride N elements (one row apart)

$ gcc -O2 rowcol.c -o rowcol && ./rowcol
row-major (stride 1): 0.1269 s
col-major (stride N): 2.4744 s
ratio col/row      : 19.50x
```

The column walk is **19.5× slower** for arithmetic that is byte-for-byte identical. Nothing about the computation changed — the same 67 million additions run in both — only the order of memory access, and that order alone cost a factor of nineteen. This is the whole chapter in one measurement: at the RAM-model level (Chapter 1) the two loops are the same program, and a beginner reasoning from instruction counts predicts they run in the same time; the real machine charges an order of magnitude for the difference in locality, exactly the “where does the time go” question Chapter 33 said the instruction count cannot answer. The fix is almost always to reorder the loops or transpose the data so the innermost loop strides by one — and note that the compiler generally may *not* do this for

you, because swapping loop nesting can change results under floating-point rounding and is not always legal, so the choice of traversal order is yours to make and to measure.

THE SAME IDEA THREE WAYS: LOCALITY DECIDES THE COST OF IDENTICAL WORK

1. In words. Memory is a hierarchy whose levels differ by orders of magnitude, and the hardware moves whole cache lines. Touch memory contiguously (spatial locality) and reuse it before it is evicted (temporal locality) and most accesses hit a fast level; stride far or revisit after eviction and most accesses stall on DRAM. The arithmetic is the same either way; the access order sets the price.

2. As a picture. Figure 34.1: the hierarchy as a pyramid (registers → L1 48 KiB → L2 1.25 MiB → L3 24 MiB → DRAM), each level wider and slower; beside it, a 64-byte cache line with a row-major walk using all of it and a column-major walk using one element and discarding the rest.

3. As numbers/contrast. Same array, cost per access rose from ~2.9 ns at stride 1 to ~10.6 ns at a 64-byte stride as line utilization fell (§34.2); the same 8192×8192 sum ran in 0.1269 s row-major and 2.4744 s column-major — **19.5×** — for byte-identical arithmetic (§34.3). Identical work, an order-of-magnitude spread, decided only by access order.

(Answer to the §34.1 quick check: fastest to slowest — registers, L1, L2, L3, DRAM — and a DRAM access is roughly two orders of magnitude slower than an L1 hit (tens of nanoseconds versus a fraction of one). L1 is only tens of KiB, far too small to hold a working set of any size, so you cannot “just keep it in cache”; you must arrange your accesses so the small hot data you need next is already near the core — that arrangement is locality, spatial and temporal.)

The hierarchy (left) and the cache line as the unit of transfer (right).




Figure 34.1: The memory hierarchy (left) charges orders of magnitude more for a DRAM access than an L1 hit, and the hardware moves whole 64-byte lines (right). A row-major walk uses every element of each line it fetches; a column-major walk uses one and discards the rest, evicting lines before their reuse — measured here at a 19.5× gap for identical arithmetic (§34.3).

34.4 Temporal locality, working sets, and blocking

Spatial locality is about the *neighbors* of what you touch; **temporal locality** is about revisiting the *same* data soon enough that it is still cached. The governing quantity is the **working set** — the amount of data a phase of the program touches repeatedly. If the working set fits in a cache level, the second and later passes over it hit that level; if it

overflows, each pass evicts what the next pass needs, and you pay DRAM latency on every access even though you are reusing data. The classic example is dense matrix multiplication: the naive triple loop streams an entire $N \times N$ matrix for each output element, so once N is large enough that a row plus a column no longer fit in cache, nearly every access misses. The standard cure is **blocking** (tiling): restructure the loops to work on small sub-blocks that *do* fit in cache, compute everything that block contributes before moving on, so each loaded line is reused many times before eviction (CS:APP, ch.6; Hennessy & Patterson, ch.2). The block size is chosen so the active tiles fit a target cache level, turning a memory-bound computation into a compute-bound one. The general principle generalizes far past matrices: **size your hot working set to a cache level and reuse each loaded line as much as possible before it is evicted**. When a profile (Chapter 33) shows a loop stalled on memory, the two questions are always the same — am I striding contiguously (spatial), and does my reused data fit in cache (temporal) — and the fixes are reorder the accesses, or shrink the working set until it does.

TRAP: THE INNOCENT TRANSPOSE — STRIDING ACROSS CACHE LINES BY ACCIDENT

The column-major sum of §34.3 was $19\times$ slower than the row-major one, and nobody wrote a “slow” loop — they wrote for (j) for (i) `A[i][j]`, which looks as reasonable as the other order and is a one-token difference. This is the trap: **an access pattern that strides across cache lines instead of along them is invisible in the source**. It hides in iterating a row-major array in column order; in an array of large structs where you touch one field per element (Chapter 35); in chasing pointers through a linked list or tree whose nodes were allocated all over the heap; in a hash table walked in key order. Each looks like ordinary code and each turns a fast level into DRAM latency on nearly every access. Two defenses. First, *know your data's memory layout* and make the innermost, most frequent loop stride by one along it — transpose the data, or swap the loops, so the hot access is contiguous. Second, *measure*: the gap does not show up in a code review or an instruction count, only in a timed run (Chapter 33) or a cache-miss counter (`perf`, where available). The order-of-magnitude cost of getting this wrong is precisely why it is worth a deliberate look at any hot loop's access pattern.

SAY IT LIKE A SYSTEMS ENGINEER

Try it first. A teammate says: “Our image filter is too slow. The inner loop is just `out[x][y] = f(in[x][y])` over a big 2-D array — it's the simplest possible code, so I don't think we can do better without rewriting it in assembly.” Cover the paragraph and respond in four or five sentences, drawing on §34.1–§34.4.

Model answer. “Before assembly, let's check the access pattern, because ‘simplest code’ and ‘fastest code’ aren't the same thing here. If the arrays are stored row-major and the loops are nested for `x { for y }` with `x` the row index, then `[x][y]` is striding by a whole row on the inner loop — every access lands in a new cache line and uses a fraction of it, which on the matrix case cost us about $19\times$. Swapping the loop nesting so the inner index runs along the contiguous dimension, or transposing one array, could give most of that back for a one-line change, no assembly. Let's measure it the Chapter 33 way — same input, `-O2`, many trials, minimum — with the loops each way and compare; if the profile then still shows us memory-bound, we look at blocking the working set to fit L2, and only after that consider vectorizing the kernel.” Reaching for the access pattern and a measurement before reaching for assembly is the senior register: the cheap, high-leverage fix first.

(Answer to the §34.2 quick check: a stride-1 walk reuses each fetched 64-byte line for all 16 ints in it, so the line-fetch cost is amortized across many accesses and each is cheap; a stride-16-int walk strides exactly one 64-byte line per access, so every access triggers a fresh line fetch and uses only 4 of its 64 bytes, and the cost climbs to the per-line fetch cost. A 64-byte stride is special because it equals the cache line: it is the smallest stride at which no two consecutive accesses ever share a line, so spatial locality is completely defeated.)

34.5 What to carry forward

The memory hierarchy is not a detail under the RAM model; on real hardware it is often *the* cost. Its levels differ by orders of magnitude and the fast ones are small (48 KiB of L1 on this machine, 24 MiB of shared L3), so performance is decided by **locality**: touch memory contiguously so each 64-byte line you fetch serves many accesses (**spatial locality**, §34.2), and reuse data while it is still cached by sizing your hot working set to a cache level (**temporal locality** and **blocking**, §34.4). The cost of ignoring this is not marginal: identical arithmetic ran **19.5×** slower in column-major than row-major order (§34.3), a factor that no instruction-count reasoning predicts and that a profile (Chapter 33) surfaces only as “stalled on memory.” The cardinal rule of the chapter: the cache line, not the byte, is the unit you pay for, so pack what a loop uses into as few lines as possible and stride along them, not across them.

This chapter took the data layout as given — a matrix, an array — and controlled the *order* of access. The next chapter, [Chapter 35](#), takes the complementary lever: controlling the *layout itself*, so that the bytes a loop needs are contiguous by construction. That is **data-oriented design** — array-of-structs versus struct-of-arrays, hot/cold splitting, packing — the operational form of the mechanical sympathy Chapter 5 introduced, and it turns the “stride along, not across” rule from something you arrange at the loop into something guaranteed by the data structure. After that, [Chapter 36](#) shows how a cache-friendly, contiguous layout is also the precondition for the CPU's other lever of throughput, **SIMD**: the vector unit can only stream data it can reach quickly. Locality first, because a vector unit starved for memory is no faster than a scalar one.

CAN YOU... WITHOUT LOOKING?

- State the levels of the memory hierarchy and, to an order of magnitude, the L1-to-DRAM latency gap.
- Define spatial and temporal locality, and say why the cache line (not the byte) is the unit you pay for.
- Explain why row-major and column-major traversal of the same matrix can differ by an order of magnitude.
- Explain what a working set is and how blocking keeps reuse in cache.
- Name three ordinary-looking code patterns that secretly stride across cache lines.
- Describe how you would measure whether a proposed locality fix actually helped.

RETRIEVAL — CLOSED BOOK

Cover the chapter. Answer from memory; rate confidence first.

1. **R1.** Why do the orders-of-magnitude latency gaps between cache levels make locality, not instruction count, the thing to optimize? (§34.1)
2. **R2.** What is a cache line, and why does striding by exactly one line size defeat spatial locality? (§34.2)
3. **R3.** Why is column-major traversal of a row-major matrix an order of magnitude slower than row-major, for identical arithmetic? (§34.3)
4. **R4.** What is a working set, and how does blocking (tiling) turn a memory-bound loop into a compute-bound one? (§34.4)
5. **R5.** A profile says a loop is stalled on memory. What two locality questions do you ask, and what are the corresponding fixes? (§34.4)

FURTHER READING

- **Ulrich Drepper, “What Every Programmer Should Know About Memory” (Red Hat, Inc., 2007).** The definitive long-form treatment of caches, cache lines, prefetching, and the locality techniques of this chapter, published as a Red Hat whitepaper and as an LWN.net series (lwn.net/Articles/250967). Author Ulrich Drepper; a whitepaper, not a journal article, so no DOI — cited as such. The 19× effect of §34.3 is exactly what its cache chapters predict.
- **Randal E. Bryant & David R. O'Hallaron, *Computer Systems: A Programmer's Perspective*, 3rd ed. (Pearson, 2015), ISBN 978-0-13-409266-9.** Chapter 6 (the memory hierarchy) is the source for the working-set and blocking discussion of §34.4 and for the representative latency figures of §34.1. First author Randal Bryant; edition, publisher, and ISBN verified via the publisher.
- **John L. Hennessy & David A. Patterson, *Computer Architecture: A Quantitative Approach*, 6th ed. (Morgan Kaufmann, 2017), ISBN 978-0-12-811905-1.** Chapter 2 develops the memory hierarchy quantitatively — cache organization, miss penalties, and the arithmetic behind the cost model of §34.1. First author John Hennessy; edition and ISBN verified via the publisher.
- **Linux kernel sysfs cache topology (`/sys/devices/system/cpu/cpu*/cache/`).** The source of the real cache sizes and 64-byte line size quoted for this machine in §34.1–§34.2 (`size`, `coherency_line_size`, `level`). Read directly from this environment; your machine's values will differ, so read your own before quoting them.

Exercises

34.1 WARM-UP Define spatial locality and temporal locality in one sentence each, and give a small code example that has good spatial locality and one that has poor spatial locality.

34.2 WARM-UP The cache line on this machine is 64 bytes. If you read one `double` (8 bytes) and then read the next 7 `doubles` in memory, how many cache-line fetches does that cost in the best case, and why?

34.3 **WARM-UP** In one sentence, why is a DRAM access being $\sim 100\times$ slower than an L1 hit the reason we care about locality at all, rather than just counting instructions?

34.4 **CORE** The §34.3 measurement: row-major 0.1269 s, column-major 2.4744 s, on an 8192×8192 `double` matrix. (a) Both loops do the same number of additions — so what, precisely, differs between them to produce a $19.5\times$ gap? (b) A `double` is 8 bytes and the line is 64 bytes; how many elements of each fetched line does the column-major walk actually use before the line is (much later) needed again? (c) Name two ways to fix the column-major loop.

34.5 **CORE** The stride sweep of §34.2 rose from ~ 2.9 ns/access at stride 1 to ~ 10.6 ns at a 64-byte stride. (a) Explain the climb in terms of accesses-served-per-line. (b) Why is there a recognizable knee at 64 bytes specifically? (c) Costs kept rising past 64 bytes; name two mechanisms other than the L1 line that explain the further increase.

34.6 **CORE** A colleague's hot loop chases a linked list: each node is `malloc`'d separately and holds one `int` plus a `next` pointer. Profiling shows it stalled on memory. (a) Why does this pattern have poor spatial locality even though each node is small? (b) What layout change would improve it, and why? (c) What does this share with the column-major case of §34.3?

34.7 **COST** Reason about working sets and blocking. A program makes repeated passes over an array of size S , and the last-level cache holds C bytes. (a) Qualitatively, what happens to the miss rate as S grows from below C to well above C ? (b) Explain how blocking the computation into chunks of size $\leq C$ changes the picture. (c) Why can blocking turn a “memory-bound” loop into a “compute-bound” one without reducing the number of arithmetic operations at all?

34.8 **FADED** *Worked, with the last step for you.* You must sum a 4096×4096 `float` matrix stored row-major.

Step 1 (given): row-major storage means `A[i][j]` and `A[i][j+1]` are adjacent; `A[i][j]` and `A[i+1][j]` are one row (16 KiB) apart.

Step 2 (given): the loop `for (j) for (i) sum += A[i][j]` strides by a full row on its inner loop, so nearly every access misses.

Step 3 (given): the loop `for (i) for (j) sum += A[i][j]` strides by one element on its inner loop, so each 64-byte line serves 16 accesses.

Step 4 (you): State which loop to ship and why, and describe exactly how you would measure that the reordered loop is faster on this matrix (clock, optimization level, trials, what to hold constant).

34.9 **MIXED** For each, say whether the problem is primarily *spatial* locality, *temporal* locality, or *neither*, and why: (a) iterating a row-major 2-D array in column order; (b) repeatedly scanning a 100 MiB table that does not fit in a 24 MiB L3; (c) a tight loop summing a 1 KiB array a million times; (d) chasing pointers through a tree whose nodes are scattered across the heap.

34.10 **MIXED** Classify each as *true* or *false* and fix the false ones: (a) “Reading one byte from memory fetches exactly one byte into cache.” (b) “Two loops with the same instruction count always take the same time.” (c) “The compiler will always reorder loops to fix a bad traversal order for you.” (d) “If the working set fits in L2, repeated passes over it mostly hit L2.” (e) “Blocking reduces the number of arithmetic operations.”

34.11 **STRETCH** Tie locality back to the profiler. (a) A profile (Chapter 33) shows one loop at 85% of runtime but a low instruction count per iteration; give the locality hypothesis and how you would test it. (b) Without `perf` or a cache-miss counter, describe a purely wall-clock experiment that distinguishes “memory-bound because of poor locality” from “genuinely compute-bound.” (c) Why is “transpose the data / swap the loops, then re-measure” a stronger move than “rewrite the loop in assembly” as a first response to a memory-bound loop?

34.12 **LAB** On your own machine (a C toolchain suffices): (a) Read your real cache sizes and line size from `/sys/devices/system/cpu/cpu0/cache/index*/` (or your platform's equivalent) and record them. (b) Reproduce §34.3 — sum a large row-major matrix both ways at `-O2`, time each as the minimum of several `clock_gettime(CLOCK_MONOTONIC)` trials, and report the ratio; explain it using your line size. (c) Reproduce the stride sweep of §34.2 over an array larger than your L3 and find the knee; check that it falls at your line size. Note your toolchain version and CPU, and state that these are wall-clock measurements (no hardware counters) if you lack `perf`.

Chapter 35 — Data-Oriented Design: Laying Out Data for the Cache

Before you start: Chapter 5 (mechanical sympathy and data-oriented design — this chapter is that idea made operational and measured), Chapter 34 (locality and the cache line — here we control the layout, not just the access order), Chapter 26 (false sharing — its cause is cache-line layout, and we measure it here), Chapter 33 (measure first). · Chapter 34 took the data layout as given and reordered the accesses; this chapter changes the *layout itself*, so the bytes a loop needs are contiguous by construction.

BEFORE YOU READ ON

From memory, reaching back: (1) Chapter 34 — what is the unit the hardware moves between memory and cache, and how big is it on typical hardware? (2) Chapter 26 — what is *false sharing*, in one line? (3) *Predict*: you have 16 million particles, each a struct of eight doubles (position, velocity, mass, charge), and a loop that sums just the *x* field. Compare that to storing all the *xs* in their own separate array and summing those. Same 16 million additions. How much faster do you expect the separate-array version to be: the same, a little, or several times? Write your guess; §35.2 answers it.

Object-oriented habit says: model a thing as a struct with all its attributes, make an array of those structs, and write loops over them. **Data-oriented design** inverts the priority. It asks first what *transformations* the hot loops perform on the data — which fields they read, in what order, how often — and then lays the data out so those loops touch memory contiguously, with no wasted bytes per cache line. The object is a convenience for the programmer; the layout is a contract with the cache, and Chapter 34 showed that contract can be worth an order of magnitude. Mike Acton put the discipline bluntly in his 2014 keynote: the purpose of the program is to transform data, so “if you don't understand the data, you don't understand the problem,” and the hardware you run on has one overriding cost — memory (Acton, 2014; Fabian, 2018). This chapter makes that concrete in four steps: the core idea and why the object is not the natural unit (§35.1); array-of-structs versus struct-of-arrays, measured at 4× here (§35.2); hot/cold splitting, field ordering, and packing to fit more useful data per line (§35.3); and the place where layout collides with concurrency — **false sharing**, measured at nearly 8× (§35.4). Then what to carry into the vectorization chapter, which needs exactly the layout this one produces (§35.5).

On the code in this chapter. As before the programs are C, compiled with gcc 11.4.0 and run in this environment on an Intel Core i7-12700H, timed with `clock_gettime(CLOCK_MONOTONIC)` and reported as the minimum of several trials (Chapter 33). The times are illustrative of the effect and hardware-specific; a hardware-counter profiler (`perf`) is not installed here, so the effects are shown through wall-clock timing rather than cache-miss counts, and this is stated where it matters.

35.1 The object is not the natural unit

Start from the mismatch. A natural object-oriented model of a particle bundles everything about it — position, velocity, mass, charge — into one struct, and a simulation holds an array of them. But a physics update loop that only needs the `x` coordinate now strides through memory in jumps the size of the *whole struct*, dragging velocity, mass, and charge into cache on every step only to ignore them. The cache line is spent on bytes the loop never reads. Data-oriented design starts from the loop, not the object: it asks which fields the hot transformation actually touches, and arranges memory so *those* bytes are contiguous and the rest are elsewhere. The design question flips from “what is a particle?” to “what does the update loop read, and how do I give it that with no wasted line?” This is the same “stride along, not across” rule as Chapter 34, but achieved by *choosing the layout* rather than by reordering accesses over a fixed one — and because you own the layout, you can guarantee the good pattern instead of hoping the access order preserves it. The rest of the chapter is the concrete moves this principle licenses: splitting one array of fat structs into several arrays of fields (§35.2), separating hot fields from cold (§35.3), and being deliberate about which fields share a cache line when threads are involved (§35.4).

QUICK CHECK

Why does a loop that reads only the `x` field of an array of large structs waste most of every cache line it fetches? State the data-oriented reframing: which question replaces “what is a particle?” (Answer after the next section.)

35.2 Array-of-structs versus struct-of-arrays

The canonical data-oriented transform has a name. **Array-of-structs (AoS)** is the object-oriented default: `Particle a[N]`, each element holding all its fields together. **Struct-of-arrays (SoA)** stores each field in its own contiguous array: `double x[N], y[N], vx[N], ...`. For a loop that reads one field across all elements, SoA is pure spatial locality — every fetched line is full of the field you want — while AoS strides by the struct size and wastes the rest of each line. Here is the exact case from the opener: 16 million particles, an 8-field struct that happens to be 64 bytes (one whole cache line), summing only `x`. Compiled and run here (gcc 11.4.0, -O2, minimum of 7 trials):

```
struct Particle { double x,y,z,vx,vy,vz,mass,charge; }; // 64 bytes = one cache line
for (i) s += aos[i].x;    // AoS: stride 64 B, uses 8 of every 64 bytes fetched
for (i) s += soa_x[i];    // SoA: stride 8 B, every line full of x's

$ gcc -O2 aossoa.c -o aossoa && ./aossoa
AoS  sum of .x (64 B stride): 0.1514 s
SoA  sum of x ( 8 B packed): 0.0371 s
ratio AoS/SoA                : 4.08x
```

SoA is **4.08× faster** for the same 16 million additions. The reason is the same accounting as Chapter 34's cache line: the AoS loop fetches a 64-byte line and uses 8 bytes of it — one `double` — before moving on, so 7 of every 8 bytes it pulls from memory are thrown

away; the SoA loop uses all 64 bytes of every line, and the ratio of useful-to-fetched bytes ($8\times$ better) is most of the speedup. Notice this is not a made-up struct: 64 bytes is an entirely ordinary object size, and the “sum one field” loop is what half of data processing looks like. The tradeoff is real and worth stating plainly: SoA wins decisively when loops sweep *one or a few fields across many elements*; AoS can win when a loop touches *all* fields of *one* element at a time (then the whole struct in one line is exactly what you want). Data-oriented design is choosing per the dominant access pattern the profile shows — and when the hot loop reads a column, you store a column.

QUICK CHECK

Define AoS and SoA. For a loop summing one field across all elements, why is SoA faster, and roughly by what factor did the byte accounting predict here? When would AoS be the better layout instead? (Answer below the communicate box.)

35.3 Hot/cold splitting, field order, and packing

SoA is the extreme; between it and AoS lies a spectrum of cheaper wins from the same principle: *get more useful bytes into each fetched line*. Three moves recur. First, **hot/cold splitting**: if a struct has a few fields the hot loop reads every time (hot) and many it touches rarely (cold), split them into two structs — a dense array of hot fields the loop streams efficiently, and a parallel array of cold fields it visits only when needed. This keeps AoS convenience for the cold data while giving the hot loop SoA-like density. Second, **field ordering and padding**: the compiler inserts padding to satisfy each field's alignment, so a careless field order can bloat a struct with holes — ordering fields large-to-small typically minimizes that padding and shrinks the struct, so more elements fit per line (CS:APP, ch.3). A 24-byte struct that packs to 16 bytes with reordered fields is 50% more elements per line for free. Third, **shrinking the data itself**: using a 32-bit index instead of a 64-bit pointer, or a float instead of a double where precision allows, or a smaller enum, all raise the number of useful elements per cache line and so the effective bandwidth. None of these change the algorithm or the operation count; they change how many of the bytes you drag through the cache are ones you actually use. The unifying question, again, is the one Chapter 34 posed at the loop and this chapter poses at the type: **of every 64-byte line this code fetches, how many bytes does it use?** Every layout move here raises that fraction.

THE SAME IDEA THREE WAYS: RAISE THE USEFUL BYTES PER FETCHED LINE

1. In words. The hardware fetches whole cache lines; design the data so the bytes a hot loop reads are packed into as few lines as possible and the bytes it ignores are elsewhere. Store the column the loop sweeps (SoA), split hot fields from cold, order and shrink fields to cut padding. The operation count is unchanged; the wasted-bytes-per-line falls.

2. As a picture. Figure 35.1: AoS as an array of fat boxes with one thin stripe (the read field) per box, the rest shaded “fetched but unused”; SoA as one tight array of just those stripes, every line full. A cache line drawn over each shows AoS using 1 slot of 8, SoA using 8 of 8.

3. As numbers/contrast. Summing one field of 16M 64-byte structs: AoS 0.1514 s, SoA 0.0371 s — **4.08×**, tracking the 8× better useful-bytes-per-line (§35.2). And the same layout lever, mishandled across threads, produced a **7.9×** slowdown from false sharing (§35.4). Layout gives and layout takes away.


(Answer to the §35.1 quick check: a loop reading only x from an array of large structs fetches a full 64-byte line per element but uses only the few bytes of x in it, so most of every line — and most of the memory bandwidth it consumes — is wasted on fields it never reads. The data-oriented reframing replaces “what is a particle (and what are all its attributes)?” with “what does the hot loop read, and how do I lay memory out so it reads only that, contiguously?”)

Array-of-structs vs struct-of-arrays: summing one field (x) of many records.

AoS: one fetched line (64 B) holds one whole record; the loop uses only x (1 of 8 fields).



SoA: the x 's are their own contiguous array; every fetched line is full of x (8 of 8 used).



Same 16M additions: AoS 0.1514 s vs SoA 0.0371 s = 4.08× faster

Figure 35.1: Summing one field across many records. In AoS each fetched 64-byte line holds one whole record and the loop uses a single field, discarding the rest; in SoA the field has its own contiguous array and every fetched line is fully used. Measured here at 4.08× for identical arithmetic (§35.2).

35.4 When layout meets concurrency: false sharing

Layout can also cost you a factor when threads are involved, and the mechanism is again the 64-byte line. Chapter 26 named **false sharing**: two threads write to two *different* variables that happen to sit in the *same* cache line, and the cache-coherence protocol, which tracks ownership at line granularity, ping-pongs that line between the cores' caches on every write — even though the threads never touch each other's data. The variables are logically independent; the hardware sees one shared line and serializes access to it. This is a pure layout bug: nothing is logically contended, yet the line is. The fix is layout too — **pad or align each thread's data to its own cache line** so no two threads' hot variables share one. Here are four threads each incrementing their own

counter 200 million times, first with the counters packed into one line, then padded to a line each. Compiled and run here (gcc 11.4.0, -O2, four threads on this machine, minimum of 3 trials):

```
struct { volatile long v; } shared[4];           // 4 counters in ONE 64-byte line
struct { volatile long v; char pad[56]; } padded[4]; // each counter on its OWN line

$ gcc -O2 falseshare.c -o falseshare -lpthread && ./falseshare
false-shared (one line) : 1.3925 s
padded (one line each)  : 0.1757 s
slowdown from sharing   : 7.92x
```

Packing the counters into one line made the program **7.9× slower**, and not one byte of data was actually shared between threads — each thread only ever touched its own counter. All 7.9× is the coherence protocol bouncing the single contested line between four cores. The 56 bytes of padding that fixed it look like waste, and in single-threaded terms they are; across threads they buy back a factor of eight, because they guarantee each thread's line is its own. This is the same lesson as SoA from the other direction: in §35.2 *packing* data together helped a single loop; here *separating* data onto distinct lines helped concurrent threads. Both are the same principle — control which bytes share a cache line — applied to opposite access patterns. The general rule for concurrent code: **per-thread hot data should live on its own cache line**, achieved by padding or by `alignas(64)`, and you confirm it the Chapter 33 way, by measuring the padded and unpadded versions.

TRAP: THE ARRAY OF FAT OBJECTS (AND ITS MIRROR, THE SHARED COUNTER LINE)

The default reflex — model a thing as a struct, make an array of them, loop over the array — is invisible as a performance bug because it is *good object-oriented style*. Yet a hot loop that reads one field of a 64-byte struct wastes 7 of every 8 bytes it fetches, which cost **4×** in §35.2, and no code review flags “array of structs” as slow because it is textbook design. The mirror-image trap is its concurrent form: put each thread's counter in an adjacent array slot — the obvious, tidy layout — and false sharing costs **8×** (§35.4), again with nothing in the source that looks wrong. Both are layout bugs hiding behind reasonable-looking code, and both are invisible to instruction counting; they show up only in a timed run (Chapter 33) or a cache/coherence counter. Two defenses. First, at any hot loop, ask the useful-bytes-per-line question of §35.3 — is this loop reading a column out of an array of rows? If so, store the column (SoA), or split hot from cold. Second, for per-thread hot data, pad or align to a cache line and measure both ways. The connecting insight is that *the layout is a performance decision the type system lets you make silently* — so make it deliberately, and measure it, rather than inheriting whatever the object model happened to produce.

SAY IT LIKE A SYSTEMS ENGINEER

Try it first. A teammate says: “Our entity system has an Entity struct with 20 fields, and we store them in one big Entity[] array. The movement system, which runs every frame over all entities, only reads position and velocity, and it's our hottest loop. It's clean OOP — should we leave it?” Cover the paragraph and respond in four or five sentences, drawing on §35.1–§35.4.

Model answer. “The clean-OOP layout is probably costing us several times on that loop, and it won't show up in a review. The movement system reads maybe 4 of 20 fields, so every cache line it fetches is mostly the other 16 fields it never touches — in a comparable case, summing one field of a 64-byte struct, storing it as its own array was about 4× faster. I'd hot/cold split: put position and velocity in their own tightly packed arrays (struct-of-arrays for the hot fields) that the movement system streams, and leave the cold fields in a parallel array visited only when needed. Let's measure it the Chapter 33 way — same entity count, -02, many trials, minimum — with the current layout and the split, and confirm the frame time actually dropped. And while we're at it, if any per-entity data is written by multiple worker threads, we should pad it to a cache line so we don't add false sharing on top.” Choosing the layout from the hot loop's access pattern, and proving it with a measurement, is the senior register — not defending a layout because it is tidy.

(Answer to the §35.2 quick check: AoS is Particle a[N] with all fields of each element together; SoA is a separate contiguous array per field. For a loop summing one field across all elements, SoA is faster because every fetched 64-byte line is full of that field while AoS wastes the bytes of the other fields — here the useful-bytes-per-line was 8× better and the measured speedup 4.08×. AoS is better when a loop touches all fields of one element at a time, since then the whole struct in one line is exactly what it uses.)

35.5 What to carry forward

Data-oriented design is the operational form of the mechanical sympathy Chapter 5 named: decide the memory layout from what the hot loops actually read, not from what makes a tidy object. Because the hardware fetches whole 64-byte lines (Chapter 34), the governing metric is **useful bytes per fetched line**, and every move in this chapter raises it — **struct-of-arrays** to store the column a loop sweeps (4.08× here, §35.2), **hot/cold splitting** and **field ordering and packing** to cut wasted bytes and padding (§35.3), and, across threads, cache-line **padding to avoid false sharing** (7.9× here, §35.4). None of these change the algorithm or the operation count; they change how much of the memory traffic is useful. And all of them are invisible to instruction counting and to code review — an array of fat structs and a row of adjacent counters both look like good code — so they are found and fixed only by measurement (Chapter 33).

The layout this chapter produces is not only cache-friendly; it is the precondition for the next lever of throughput. A vector unit (Chapter 36, **SIMD**) processes 4, 8, or 16 elements per instruction, but only if those elements are contiguous in memory and identically laid out — which is exactly what struct-of-arrays gives it and what array-of-structs denies it. The auto-vectorizer that Chapter 36 leans on typically cannot vectorize a strided AoS field access but happily vectorizes a packed SoA array. So the two chapters compose: Chapter 34 got the access order right, this chapter got the layout right, and

Chapter 36 turns that contiguous, well-laid-out data into vector throughput. Lay the data out for the cache first, because the vector unit, like the cache, rewards contiguity and punishes strides.

CAN YOU... WITHOUT LOOKING?

- State the data-oriented reframing: which question about the data replaces “what is this object?”
- Define AoS and SoA, and say for which access pattern each wins.
- Explain the “useful bytes per fetched line” metric and how SoA, hot/cold splitting, and packing each raise it.
- Explain false sharing as a layout bug and give its fix.
- Say why an array of fat structs and a row of adjacent per-thread counters are both invisible performance bugs.
- Describe how you would measure whether a layout change actually helped.

RETRIEVAL — CLOSED BOOK

Cover the chapter. Answer from memory; rate confidence first.

1. **R1.** What question does data-oriented design ask first, and why is the object not the natural unit of layout? (§35.1)
2. **R2.** What are AoS and SoA, why is SoA faster for a one-field sweep, and when does AoS win? (§35.2)
3. **R3.** Name three layout moves that raise useful-bytes-per-line without changing the algorithm. (§35.3)
4. **R4.** What is false sharing, why did it cost $\sim 8\times$ here, and what is the fix? (§35.4)
5. **R5.** Why are these layout issues invisible to code review and instruction counting, and what makes them visible? (§35.3, §35.4)

FURTHER READING

- **Richard Fabian, *Data-Oriented Design: Software Engineering for Limited Resources and Short Schedules* (2018), ISBN 978-1-916478-70-1.** The book-length treatment of the discipline of this chapter — designing around data transformations and cache-friendly layout, with SoA, hot/cold splitting, and existence-based processing. Author Richard Fabian; year and ISBN verified via the publisher listing. Also available online at dataorienteddesign.com/dodbook.
- **Mike Acton, “Data-Oriented Design and C++,” keynote, CppCon 2014.** The influential talk (Acton was then engine director at Insomniac Games) that popularized the “understand the data, then design the code” framing of §35.1 and the AoS-vs-SoA argument of §35.2. A recorded conference talk, not a paper, so cited as such; slides in the CppCon2014 GitHub archive.
- **Randal E. Bryant & David R. O'Hallaron, *Computer Systems: A Programmer's Perspective*, 3rd ed. (Pearson, 2015), ISBN 978-0-13-409266-9.** Chapter 3 covers struct layout, alignment, and padding — the basis for the field-ordering and packing of §35.3. First author Randal Bryant; edition and ISBN verified via the publisher.
- **Ulrich Drepper, “What Every Programmer Should Know About Memory” (Red Hat, Inc., 2007).** Its cache-coherency sections explain the line-granularity ownership that makes false sharing (§35.4) cost what it does, and the layout techniques that avoid it. A Red Hat whitepaper / LWN.net series (lwn.net/Articles/250967), so no DOI — cited as such.

Exercises

- 35.1** WARM-UP Define array-of-structs (AoS) and struct-of-arrays (SoA) in one sentence each, with a small code sketch of each for a particle with fields `x`, `y`, `vx`, `vy`.
- 35.2** WARM-UP State the “useful bytes per fetched line” metric in one sentence, and use it to explain why summing one field of a 64-byte AoS struct wastes most of each cache line.
- 35.3** WARM-UP What is false sharing, in one or two sentences, and what one-word layout change fixes it?
- 35.4** CORE The §35.2 measurement: AoS 0.1514 s, SoA 0.0371 s (4.08×) summing one field of 16M 64-byte structs. (a) Compute the useful-bytes-per-line ratio between the two layouts and explain why it roughly predicts the speedup. (b) Why is the measured 4.08× a little less than the 8× byte ratio? Give one plausible reason. (c) Describe a loop for which AoS would be the *faster* layout, and why.
- 35.5** CORE Hot/cold splitting. A `Player` struct has `x`, `y` (read every frame) and `name`, `inventory`, `achievements` (read rarely). (a) Sketch the hot/cold split. (b) Why does this help the per-frame loop even though you did not remove any fields? (c) What is the cost of the split — what becomes slower or more awkward, and when is it not worth it?

35.6 **CORE** Field ordering and packing. A struct is declared

`{ char a; double b; char c; }.` (a) Explain why alignment makes this larger than the 10 bytes of its fields, and estimate its actual size on a typical 64-bit target. (b) Reorder the fields to minimize padding and give the new size. (c) In an array of a million of these, why does the smaller struct improve a loop that sweeps them?

35.7 **COST** Reason about the false-sharing measurement (§35.4): packed 1.3925 s, padded 0.1757 s, 7.92×. (a) Given no bytes are logically shared between threads, where does the entire 7.9× come from? (b) Why does the padding (56 wasted bytes per counter) buy back the factor, and what is it trading? (c) Would the gap grow or shrink if you ran *one* thread instead of four, and why?

35.8 **FADED** *Worked, with the last step for you.* A ray tracer stores scene objects as `Sphere{ double cx,cy,cz,r; Material m; char name[48]; }` in a `Sphere[]`, and its hottest loop is an intersection test that reads only `cx,cy,cz,r`.

Step 1 (given): the loop reads 32 bytes (four doubles) of a struct that is well over 100 bytes, so it wastes most of every fetched line.

Step 2 (given): the profile (Chapter 33) confirms the intersection loop dominates runtime and is stalled on memory.

Step 3 (given): two candidate layouts — full SoA (separate `cx[],cy[],cz[],r[]` arrays) and hot/cold split (a packed `{cx,cy,cz,r}` array plus a parallel array of the cold fields).

Step 4 (you): Choose a layout, justify it from the access pattern, and describe exactly how you would measure that it beats the original (clock, optimization level, trials, what to hold constant).

35.9 **MIXED** For each, say whether the better layout is *AoS*, *SoA*, or *hot/cold split*, and why: (a) a loop that sweeps one field across millions of records; (b) a loop that processes one record fully (all its fields) then moves to the next; (c) a struct with 3 hot fields and 30 cold ones, swept every frame on the hot fields; (d) two threads each updating their own per-thread counter in a tight loop.

35.10 **MIXED** Classify each as *true* or *false* and fix the false ones: (a) “SoA is always faster than AoS.” (b) “Reordering struct fields can shrink the struct without removing any field.” (c) “False sharing means two threads are writing the same variable.” (d) “Padding a struct to a cache line is pure waste with no benefit.” (e) “A layout change that keeps the same operation count cannot change runtime.”

35.11 **STRETCH** Tie layout to the whole part. (a) Chapter 34 reordered accesses over a fixed layout; this chapter changed the layout. Give one problem each is the right tool for, and one where you would use both together. (b) Explain why SoA is also the layout a vectorizer (Chapter 36) wants, referencing contiguity. (c) A colleague proposes converting the entire codebase from AoS to SoA up front, before profiling. Argue, from Chapters 33–35, why that is the wrong order.

35.12 **LAB** On your own machine (a C toolchain suffices; a threads library for part c): (a) Reproduce §35.2 — a struct of several `doubles` in an array, versus a separate array of one field; sum that field both ways at `-O2`, time each as the minimum of several `clock_gettime(CLOCK_MONOTONIC)` trials, and report the ratio, explaining it from your struct size and line size. (b) Vary the struct size (16, 32, 64, 128 bytes) and observe how the AoS/SoA ratio changes; explain the trend. (c) If you have threads, reproduce the false-sharing experiment — counters packed into one line versus padded to a line each — and report the slowdown. State that these are wall-clock measurements (no hardware counters) if you lack `perf`, and note your toolchain and CPU.

Chapter 36 — SIMD and Vectorization

Before you start: Chapter 2 (the CPU up close — the vector unit is another execution resource on the same core), Chapter 35 (struct-of-arrays — the contiguous layout the vectorizer needs), Chapter 34 (locality and memory bandwidth — the ceiling that decides whether SIMD helps), Chapter 10 (the optimizing compiler — the auto-vectorizer is one of its transforms), Chapter 33 (measure first). · Chapters 34 and 35 got the data close to the core and contiguous; this chapter turns that data into wide, parallel arithmetic — when the memory system can keep up.

BEFORE YOU READ ON

From memory, reaching back: (1) Chapter 35 — why does a struct-of-arrays layout stream one field faster than an array-of-structs? (2) Chapter 34 — what does it mean for a loop to be “memory-bandwidth bound” rather than compute bound? (3) *Predict*: a modern vector register holds 8 floats (256-bit AVX), so one vector instruction does 8 multiplies at once. If you take a scalar arithmetic-heavy loop over an array that fits in L1 and let the compiler vectorize it, what speedup do you expect: about 8×, less, or possibly more? And would a loop that streams a 64 MiB array instead see the same? Write your guesses; §36.2 and §36.3 answer them.

A scalar instruction operates on one pair of operands: one add, one result. A **SIMD** instruction — Single Instruction, Multiple Data — operates on a whole vector of them at once: with 256-bit AVX registers, eight floats added to eight floats in one instruction, on the same core, in about the same time as the scalar add. This is **data-level parallelism** (Hennessy & Patterson, ch.4): not more cores, but wider instructions on one core, and it is free throughput *if* your data is laid out so the CPU can load a vector's worth at once and your loop has no dependency forcing the elements to be done one at a time. Both conditions point back: contiguity is Chapter 35's struct-of-arrays, and independence is a property of the loop. This chapter proceeds in five steps: what SIMD is and where the lanes come from (§36.1); letting the compiler vectorize for you, measured here from 3.4 to 45 Gflop/s (§36.2); what blocks the vectorizer — dependencies, reductions and floating-point associativity, aliasing (§36.3); intrinsics, and the memory-bandwidth ceiling that decides whether any of this helps (§36.4); and what to carry forward as the performance part closes (§36.5).

On the code in this chapter. As before the programs are C, compiled with gcc 11.4.0 and run in this environment on an Intel Core i7-12700H (which supports AVX2, i.e. 256-bit / 8-wide float vectors, not AVX-512), timed with `clock_gettime(CLOCK_MONOTONIC)` and reported as the minimum of several trials (Chapter 33). Throughputs are in Gflop/s (billions of floating-point operations per second) and are illustrative of the effect, not a benchmark of the machine. Where a compiler decision is claimed — that a loop did or did not vectorize — it is confirmed with gcc's own `-fopt-info-vec` report, quoted inline, not asserted. A hardware-counter profiler (`perf`) is not installed here, so the evidence is wall-clock throughput plus the compiler's vectorization report.

36.1 One instruction, many lanes

The vector unit is extra width on the ordinary core, not a separate processor. A modern x86 core has vector registers — 128-bit (SSE), 256-bit (AVX/AVX2), and on some parts 512-bit (AVX-512) — each holding several elements: a 256-bit register holds 8 floats or 4 doubles, and a single AVX instruction like `vmulps` multiplies all 8 lanes at once. The elements travel through the pipeline (Chapter 2) together, so eight multiplies retire in roughly the time of one, giving up to an 8× throughput gain on 32-bit data for arithmetic that can be done in parallel across the lanes. Two requirements make that gain real, and both were set up by the previous chapters. First, **contiguity**: to fill a vector register in one load, the 8 elements must be adjacent in memory — precisely the struct-of-arrays layout of Chapter 35; a strided array-of-structs field forces the CPU to gather elements one at a time and throws the advantage away. Second, **independence**: the 8 lanes are computed simultaneously, so no lane may depend on another's result within the vector — a loop where each iteration needs the previous one's output cannot be vectorized as written. When both hold, the same loop that issued one scalar operation per element now issues one vector operation per 8 elements, and the arithmetic throughput multiplies. The next section shows how little you often have to do to get this, and how much it is worth.

QUICK CHECK

How many floats fit in a 256-bit AVX register, and how many does one AVX arithmetic instruction process at once? Name the two properties a loop's data and computation must have for the vector unit to deliver its throughput. (Answer after the next section.)

36.2 Auto-vectorization: let the compiler do it

You rarely write vector instructions by hand first. The compiler's **auto-vectorizer** (Chapter 10's optimizer, with the loop transforms of §36.3) recognizes a vectorizable loop and emits vector instructions for it — but only when you ask for the right optimization level and target. gcc does not vectorize at `-O2` in this version; `-O3` enables it against the conservative baseline instruction set (SSE2, 128-bit), and `-march=native` lets it use the widest vectors the actual CPU supports (AVX2, 256-bit here). Here is one compute-heavy, per-element-independent kernel — a Horner-form polynomial over an L1-resident 32 KiB array — compiled three ways and run here (gcc 11.4.0, minimum of 7 trials):

```
for (i) y[i] += (((v*0.11f+0.22f)*v+0.33f)*v+0.44f)*v+0.55f)*v+0.66f;    // v = x[i], ~11
flops/elem

$ gcc -O2 -fno-tree-vectorize simd.c ...    scalar (no vectors) :  3.43 Gflop/s    (1.049 s)
$ gcc -O3 simd.c ...                      -O3 (SSE2, 4-wide)  : 24.89 Gflop/s    (0.145 s)
$ gcc -O3 -march=native simd.c ...         -O3 native (AVX2)  : 45.14 Gflop/s    (0.080 s)

$ gcc -O3 -march=native -fopt-info-vec-optimized simd.c
simd.c:10: optimized: loop vectorized using 32 byte vectors      # AVX2: 8 floats/vector
```

The measured jump is **13×** from scalar to AVX2 (3.43 to 45.14 Gflop/s), and the compiler's own report confirms it vectorized “using 32 byte vectors” — 256-bit AVX2, 8

floats per instruction. Two things deserve care in reading this. First, the speedup *exceeds* the 8-lane width, and honestly so: vectorizing also lets the compiler use fused multiply-add (one instruction for `a*b+c`), unroll the loop, and cut per-iteration overhead, so the win is lanes-times-other-improvements, not lanes alone — which is exactly why you *measure* the speedup rather than assume it equals the vector width. Second, the flags matter enormously: the identical source ran at 3.4, 25, or 45 Gflop/s depending only on `-O2` versus `-O3` versus `-O3 -march=native`, so “did it vectorize?” is a question about the build, and Chapter 33's rule to benchmark the build you ship is what keeps you honest. The practical lesson: reach for auto-vectorization first — right flags, contiguous data, an independent loop — and confirm it happened with `-fopt-info-vec` rather than hoping.

QUICK CHECK

Why did the same source run at 3.4, 25, and 45 Gflop/s? Why is the 13× scalar-to-AVX2 speedup *more* than the 8× you would get from lane count alone? How would you confirm a loop actually vectorized rather than assume it did? (Answer below the communicate box.)

36.3 What blocks the vectorizer

Auto-vectorization is not magic; a few common properties defeat it, and knowing them lets you either fix the loop or understand why the report says “not vectorized.” The first is a **loop-carried dependency**: if iteration `i` reads a value iteration `i-1` wrote (`a[i] = a[i-1] + x[i]`, a running prefix), the lanes cannot be computed in parallel because each needs the one before, and the loop is inherently serial as written. The second is subtle and important — a **reduction under floating-point non-associativity**. Summing an array into one accumulator (`s += a[i]`) is a reduction, and vectorizing it means adding the elements in a different grouping (8 partial sums, combined at the end), which under floating-point rounding can give a *different result* than the strict left-to-right order the source specifies. Because the standard forbids the compiler from silently changing your numerical result, gcc will *not* vectorize a floating-point reduction at `-O3` unless you permit reassociation with `-ffast-math` (or `-funsafe-math-optimizations`). You can see this in a dot product streamed over a 64 MiB array pair here (gcc 11.4.0):

```
float dot = 0; for (i) dot += a[i]*b[i];      // a reduction: reassociating changes FP
rounding

$ gcc -O2 -fno-tree-vectorize dot.c ...      scalar                : 0.96 Gflop/s
$ gcc -O3 -march=native dot.c ...            -O3 native, no fast : 1.12 Gflop/s  # reduction
NOT vectorized
$ gcc -O3 -march=native -ffast-math ...      with -ffast-math    : 2.30 Gflop/s  # now
vectorized
```

Without `-ffast-math` the reduction stays essentially scalar (1.12 vs 0.96 Gflop/s — the small gain is other optimization, not vectorization); allowing reassociation unlocks the vector reduction and roughly doubles throughput. Two lessons follow. One, a **floating-point reduction needs your explicit permission to vectorize**, because the vector version is not bit-identical — so decide deliberately whether your problem tolerates the rounding difference rather than sprinkling `-ffast-math` blindly (it also enables other

unsafe assumptions). Two, the dot-product speedup here is only $\sim 2.4\times$, far short of the $13\times$ of §36.2 on the same machine, and the reason is the whole point of the next section: this loop is limited by memory bandwidth, not arithmetic. Other blockers worth naming: **pointer aliasing** (the compiler must assume two pointers might overlap and serialize to be safe — `restrict`, used in the §36.2 kernel, promises they do not), data-dependent branches inside the loop, and non-contiguous (strided or gathered) access — the array-of-structs pattern of Chapter 35, which is why layout came first.

THE SAME IDEA THREE WAYS: SIMD MULTIPLIES ARITHMETIC, NOT MEMORY

1. In words. A vector instruction does 8 (or 4, or 16) operations at once on contiguous, independent data, so an arithmetic-bound loop can speed up by roughly the lane count (often more, with FMA and unrolling). But if the loop is limited by how fast memory can feed it, widening the arithmetic does little — you were waiting on bytes, not on the adder.

2. As a picture. Figure 36.1: on the left, a scalar lane doing one multiply per step; on the right, 8 lanes doing 8 at once from one contiguous load. Below, two bars: a compute-bound kernel rising from 3.4 to 45 Gflop/s ($13\times$), and a bandwidth-bound dot product barely moving (0.96 to 2.3), capped by a “memory bandwidth” line.

3. As numbers/contrast. Same machine, same compiler: an L1-resident compute kernel went $3.43 \rightarrow 24.89 \rightarrow 45.14$ Gflop/s across scalar, SSE2, and AVX2 — $13\times$ (§36.2); a 64 MiB-streaming dot product went $0.96 \rightarrow 2.30$ Gflop/s even vectorized — $2.4\times$, and only after `-ffast-math` let the reduction vectorize at all (§36.3). Vector width is a throughput ceiling on arithmetic; memory bandwidth is a different ceiling entirely.

(Answer to the §36.1 quick check: a 256-bit AVX register holds 8 floats (or 4 doubles), and one AVX arithmetic instruction processes all 8 lanes at once. The two required properties are contiguity — the 8 elements must be adjacent so one load fills the register, i.e. Chapter 35's struct-of-arrays — and independence — no lane may depend on another's result within the vector, so a loop-carried dependency defeats it.)

Scalar vs SIMD lanes, and the two ceilings: vector width vs memory bandwidth.




Figure 36.1: SIMD does many operations per instruction on contiguous, independent lanes. On an L1-resident, arithmetic-heavy kernel this multiplied throughput $13\times$ (scalar $\rightarrow$ AVX2); on a 64 MiB-streaming dot product it gained only $\sim 2.4\times$, because that loop is limited by memory bandwidth, not by the width of the adder (§36.2–§36.4).

36.4 Intrinsics, and the memory-bandwidth ceiling

When the auto-vectorizer cannot see the parallelism — an awkward data layout, a shuffle it will not synthesize, an algorithm that needs hand-mapping to lanes — you can write the vector instructions yourself with **intrinsics**: C functions like `_mm256_mul_ps` that map one-to-one to AVX instructions, giving you the lanes explicitly at the cost of portability and readability (Intel Intrinsics Guide; Agner Fog's optimization manuals). Intrinsics are the right tool when the payoff is proven and the compiler will not deliver it, but they are a last resort after auto-vectorization, because they are verbose, CPU-specific, and easy to get wrong — and, crucially, they cannot beat the ceiling the previous section exposed. That ceiling is **memory bandwidth**. The dot product of §36.3 gained only 2.4× from vectorization not because the vector unit was idle but because the loop streams 64 MiB pair of arrays that do not fit in any cache, so it is limited by how fast DRAM can deliver bytes, not by how fast the core can multiply them. This is the **roofline** idea (Williams, Waterman & Patterson, 2009): a loop's achievable performance is capped by the lower of two ceilings — peak compute (raised by SIMD) and peak memory bandwidth (raised only by better locality) — and which one binds depends on the loop's *arithmetic intensity*, its ratio of flops to bytes moved. A dot product does ~2 flops per 8 bytes read: low intensity, bandwidth-bound, so SIMD barely helps and the real lever is Chapters 34–35 (locality and layout, to move fewer bytes). The L1 kernel of §36.2 does ~11 flops per element already in cache: high intensity, compute-bound, so SIMD delivered 13×. The discipline: **vectorize the compute-bound loops; for the bandwidth-bound ones, fix locality first** — and let a measurement, not the lane count, tell you which kind you have.

TRAP: COUNTING LANES INSTEAD OF MEASURING — THE SIMD THAT MEMORY BANDWIDTH EATS

The seductive arithmetic is “8 lanes, so 8× faster,” and it fails in both directions. It *over*-predicts on bandwidth-bound loops: the dot product of §36.3 has 8 lanes available and gained only 2.4×, because it waits on DRAM, not on the adder — hand-writing AVX intrinsics for it would add complexity and CPU-specific code for a win the memory system has already capped. And it *under*-predicts on compute-bound loops: the §36.2 kernel gained 13×, *more* than 8×, because vectorization also unlocked FMA and unrolling. Both errors come from reasoning about the hardware instead of measuring it. Two defenses. First, decide whether the loop is compute- or bandwidth-bound before reaching for SIMD — a rough arithmetic-intensity estimate (flops per byte moved) or, decisively, the Chapter 33 experiment of timing scalar versus vectorized; if widening the arithmetic does not move the number, you are bandwidth-bound and the fix is locality and layout (Chapters 34–35), not more lanes. Second, if you do vectorize, confirm it happened (`-fopt-info-vec`) and measure the actual speedup rather than assuming the lane count — because the true number, high or low, is the only one you can ship a decision on. A hand-written SIMD kernel that ignores the bandwidth ceiling is effort spent past the point where it can pay, the vector-unit version of optimizing the 5% part.

SAY IT LIKE A SYSTEMS ENGINEER

Try it first. A teammate says: “I hand-wrote AVX2 intrinsics for our vector-add loop `c[i] = a[i] + b[i]` over big arrays — 8 lanes, so it should be 8× faster — but my benchmark only shows about 1.2×. The intrinsics must be wrong; help me debug them.” Cover the paragraph and respond in four or five sentences, drawing on §36.1–§36.4.

Model answer. “The intrinsics are probably fine — 1.2× is what I'd expect, because that loop is bandwidth-bound, not compute-bound. Vector-add does one add per two loads and a store, so its arithmetic intensity is tiny: you're limited by how fast DRAM can stream the arrays, not by how fast the core adds, and widening the add from 1 lane to 8 doesn't help when the bottleneck is the memory bus. The 8× rule only applies to loops that are actually compute-bound and cache-resident — a math-heavy kernel on data that fits in L1 can hit or exceed 8×, which we've measured. So before debugging the intrinsics, let's confirm the diagnosis: time the plain `-O3 -march=native` auto-vectorized version and check `-fopt-info-vec` — I bet it matches your 1.2×, which means the memory system is the ceiling and the real win is in layout and locality (fewer bytes moved), not more lanes. If we truly need this faster, the lever is Chapters 34–35, not more SIMD.” Diagnosing the ceiling before optimizing against it — and preferring the auto-vectorizer plus a measurement to hand intrinsics — is the senior register.

(Answer to the §36.2 quick check: the same source ran at 3.4/25/45 Gflop/s because vectorization depends on the build — `-O2` did not vectorize, `-O3` used 128-bit SSE2, and `-O3 -march=native` used 256-bit AVX2. The scalar-to-AVX2 gain exceeds 8× because vectorization also enabled fused multiply-add, loop unrolling, and lower per-iteration overhead, so the win is lane count times those other improvements. You confirm a loop vectorized with the compiler's own report, `-fopt-info-vec-optimized`, which named “32 byte vectors,” rather than assuming it did.)

36.5 What to carry forward

SIMD is data-level parallelism on a single core: one instruction operating on a whole vector — 8 floats in 256-bit AVX2 — for up to lane-count throughput, and often more once vectorization unlocks fused multiply-add and unrolling (13× on the compute-bound kernel of §36.2). It is not something you usually write by hand: reach for **auto-vectorization** first, with the right flags (`-O3 -march=native`), contiguous data (Chapter 35's struct-of-arrays), and an independent loop, and *confirm* it with `-fopt-info-vec`. Know what blocks it — loop-carried dependencies, floating-point reductions that need `-ffast-math` to reassociate, pointer aliasing that `restrict` resolves, strided or branchy access — and drop to **intrinsics** only when the payoff is proven and the compiler will not deliver it. Above all, respect the ceiling: SIMD raises the *compute* roof, so it multiplies arithmetic-bound loops but does almost nothing for bandwidth-bound ones (2.4× on the streaming dot product, §36.4). Which ceiling binds is a question of arithmetic intensity, and the honest way to answer it is a measurement, not a lane count.

This chapter closes the arc the part has been building. Chapter 33 said measure first and let a profile find the dominant cost. Chapter 34 opened the most common cause of that cost — the memory hierarchy — and got the access order right. Chapter 35 got the data layout right, contiguous and dense. And this chapter turned that well-placed, well-laid-out

data into wide parallel arithmetic, while insisting that the memory system, not the vector unit, often sets the limit — which is why locality and layout came first, and why the whole part keeps returning to a timed run to tell truth from hope. That is the through-line to carry out of Performance Engineering: the machine has several ceilings — instruction count, cache, bandwidth, vector width — and engineering performance is finding, by measurement, which one you are actually under, and then raising *that* one. The next part turns from the CPU and its memory to the boundary where programs meet the outside world — **I/O and networking** — where a new set of costs, and a new set of “where does the time go” questions, take over.

CAN YOU... WITHOUT LOOKING?

- Explain what SIMD is and where the lanes come from (register width), with the count for 256-bit AVX2.
- Name the two properties a loop needs to auto-vectorize, and the flags that enable and confirm it.
- Explain why a floating-point reduction will not vectorize without `-ffast-math`, and the risk of that flag.
- Explain why the scalar-to-AVX2 speedup can exceed the lane count.
- Distinguish a compute-bound from a bandwidth-bound loop, and say which SIMD helps.
- State when you would drop to intrinsics, and why it is a last resort.

RETRIEVAL — CLOSED BOOK

Cover the chapter. Answer from memory; rate confidence first.

1. **R1.** What is SIMD, how many floats does a 256-bit AVX2 instruction process, and what two loop properties does it require? (§36.1)
2. **R2.** What flags enable and confirm auto-vectorization in gcc, and why did the same source run at 3.4/25/45 Gflop/s? (§36.2)
3. **R3.** Why won't gcc vectorize a floating-point reduction without `-ffast-math`, and what does that flag risk? (§36.3)
4. **R4.** Why did the dot product gain only $\sim 2.4\times$ from SIMD while the L1 kernel gained $13\times$? Use arithmetic intensity. (§36.4)
5. **R5.** When do you reach for intrinsics rather than the auto-vectorizer, and why is it a last resort? (§36.4)

FURTHER READING

- **John L. Hennessy & David A. Patterson, *Computer Architecture: A Quantitative Approach*, 6th ed. (Morgan Kaufmann, 2017), ISBN 978-0-12-811905-1.** Chapter 4 (“Data-Level Parallelism in Vector, SIMD, and GPU Architectures”) is the architectural source for §36.1 — vector registers, lanes, and why width is throughput. First author John Hennessy; edition and ISBN verified via the publisher.
- **Samuel Williams, Andrew Waterman & David Patterson, “Roofline: An Insightful Visual Performance Model for Multicore Architectures,” *Communications of the ACM* 52(4), 65-76 (2009), DOI 10.1145/1498765.1498785.** The roofline model of §36.4 — arithmetic intensity, the compute and bandwidth ceilings, and why a low-intensity loop like the dot product is bandwidth-bound. First author Samuel Williams; DOI verified via the ACM Digital Library.
- **Agner Fog, optimization manuals, *Optimizing software in C++ and Instruction tables* (agner.org/optimize, continuously updated).** The practical reference for vectorization, intrinsics, and per-instruction latencies/throughputs behind §36.3–§36.4. Author Agner Fog; living documents, not a dated publication, so cited as such — consult the current versions.
- **Randal E. Bryant & David R. O'Hallaron, *Computer Systems: A Programmer's Perspective*, 3rd ed. (Pearson, 2015), ISBN 978-0-13-409266-9.** Chapter 5 covers program optimization including vectorization, loop-carried dependencies, and the reduction/associativity issue of §36.3 from the programmer's side. First author Randal Bryant; edition and ISBN verified via the publisher.

Exercises

- 36.1** **WARM-UP** Define SIMD in one sentence. How many `floats` and how many `doubles` fit in a 256-bit AVX register, and how many does one such arithmetic instruction process at once?
- 36.2** **WARM-UP** Name the two properties a loop's data and computation must have for the vector unit to deliver its throughput, and say which chapter's layout supplies the first.
- 36.3** **WARM-UP** Which gcc flags (a) enable vectorization at all, (b) let it use the widest vectors your CPU supports, and (c) let you confirm that a given loop was vectorized?
- 36.4** **CORE** The §36.2 kernel went $3.43 \rightarrow 24.89 \rightarrow 45.14$ Gflop/s across scalar, `-O3` (SSE2), and `-O3 -march=native` (AVX2). (a) Why does the same source produce three different throughputs? (b) The scalar-to-AVX2 gain is $13\times$, more than the 8-lane width — give two reasons vectorization can beat the raw lane count. (c) Why is “confirm with `-fopt-info-vec`” better engineering than assuming the loop vectorized?
- 36.5** **CORE** Floating-point reductions. (a) Why does vectorizing `s += a[i]` change the numerical result, and why does that stop gcc from doing it at `-O3`? (b) What flag permits it, and what else does that flag allow that might be unsafe? (c) The dot product gained only $\sim 2.4\times$ even after vectorizing — name the reason, and say why it is *not* the reduction issue.

36.6 **CORE** For each loop, say whether it can be auto-vectorized as written, and if not, why: (a) `c[i] = a[i] * b[i];` (b) `a[i] = a[i-1] + x[i];` (c) `if (x[i] > 0) y[i] = f(x[i]);` with a heavy data-dependent branch; (d) `s += a[i];` summed into one `float` at plain `-O3`.

36.7 **COST** Arithmetic intensity and the roofline. (a) Define arithmetic intensity in one sentence. (b) Estimate it (flops per byte moved) for the dot product (2 flops per two 4-byte loads) and argue why it is bandwidth-bound. (c) For the §36.2 kernel (~11 flops per 4-byte element, data in L1), argue why it is compute-bound. (d) State the rule for which loops SIMD helps.

36.8 **FADED** *Worked, with the last step for you.* A colleague reports that hand-written AVX2 intrinsics for `c[i]=a[i]+b[i]` over 100 MiB arrays give only 1.2× over scalar, and concludes the intrinsics are buggy.

Step 1 (given): vector-add does 1 add per 2 loads + 1 store — very few flops per byte moved (low arithmetic intensity).

Step 2 (given): 100 MiB arrays do not fit in any cache, so the loop streams from DRAM every pass.

Step 3 (given): therefore the loop is bandwidth-bound; the memory bus, not the adder, sets the ceiling.

Step 4 (you): Explain why 1.2× is expected (not a bug), how you would confirm the diagnosis with the auto-vectorizer and a measurement, and where the real speedup would have to come from instead.

36.9 **MIXED** For each, say whether SIMD is likely to give a large speedup, a small one, or none, and why: (a) a math-heavy per-pixel shader on an image tile that fits in L1; (b) copying a 1 GiB buffer; (c) summing a cache-resident array of 10,000 floats in a tight loop; (d) a loop with a loop-carried dependency computing a running total sequentially.

36.10 **MIXED** Classify each as *true* or *false* and fix the false ones: (a) “8 vector lanes always means 8× faster.” (b) “gcc auto-vectorizes floating-point reductions at `-O3` by default.” (c) “`-march=native` lets the compiler use the widest vectors the build machine supports.” (d) “SIMD helps a memory-bandwidth-bound loop as much as a compute-bound one.” (e) “An array-of-structs field is as easy to vectorize as a packed struct-of-arrays field.”

36.11 **STRETCH** Tie the whole part together. (a) Explain why Chapters 34 and 35 (locality and layout) are prerequisites for Chapter 36 (SIMD), not independent tricks. (b) A profile (Chapter 33) shows one kernel at 80% of runtime; describe the sequence of questions — bound by what? fixable how? — you would ask, in order, drawing on all four chapters. (c) Argue why “the machine has several ceilings; find the one you're under and raise that one” is a better mental model than “make the code do fewer instructions.”

36.12 **LAB** On your own machine (a C toolchain with a vectorizing compiler): (a) Write a compute-heavy, per-element-independent kernel over an array that fits in L1; compile it at `-O2 -fno-tree-vectorize`, `-O3`, and `-O3 -march=native`, time each as the minimum of several `clock_gettime(CLOCK_MONOTONIC)` trials, and report the throughputs and the speedup; confirm vectorization with `-fopt-info-vec-optimized`. (b) Write a dot product over arrays far larger than your last-level cache; show it barely vectorizes without `-ffast-math` and gains little even with it, and explain via arithmetic intensity. (c) State your toolchain version, CPU, and widest vector width, and note that these are wall-clock throughputs (no hardware counters) if you lack `perf`.

VOL 1 · SYSTEMS PROGRAMMING & PERFORMANCE · PART VII

I/O & Networking, Low-Level

Files, sockets, blocking vs async, epoll/io_uring, zero-copy — the systems substrate under every data pipeline and service.

Storage & the Disk

Pages, files, the buffer pool, and the storage manager — the foundation you build a database on.

Storage Engines

B-trees and LSM-trees implemented and analyzed — the two families and when each wins.

Durability & Recovery

Write-ahead logging, ARIES, checkpoints, and crash recovery — how a database keeps its promises across failures.

VOL 2 · DATABASE INTERNALS (BUILD A DATABASE) · PART XI

Transactions & Concurrency Control

ACID, isolation levels and their anomalies, 2PL, MVCC, and OCC — correctness under concurrency, implemented.

Query Processing & Execution

Parser, planner, physical operators, and execution models (Volcano/vectorized/codegen) — turning SQL into work.

Query Optimization

Statistics, cardinality estimation, cost models, and join ordering — the optimizer, honestly.

Indexing & Access Methods

Secondary/clustered/covering indexes, and specialized structures — the access-method layer.

VOL 2 · DATABASE INTERNALS (BUILD A DATABASE) · PART XV

[BUILD] A Working Database

Assemble the pieces into a small but real relational engine —
storage → txn → query → execute — with tests.

VOL 2 · DATABASE INTERNALS (BUILD A DATABASE) · PART XVI

Toward Distributed Databases

Sharding, replication, and the distributed-transaction problem — the bridge to Volume 3.

VOL 3 · DISTRIBUTED SYSTEMS, DEEPLY · PART XVII

Models & Foundations

System and failure models, the FLP impossibility, time and clocks, causality and logical time — the formal ground.

VOL 3 · DISTRIBUTED SYSTEMS, DEEPLY · PART XVIII

Replication & Consistency

Consistency models, linearizability, CAP/PACELC, and quorums — what 'consistent' actually means.

Consensus

Paxos, Raft, and view-stamped replication (with a note on Byzantine fault tolerance) — agreement under failure.

VOL 3 · DISTRIBUTED SYSTEMS, DEEPLY · PART XX

Coordination & Transactions

Leader election, distributed transactions, 2PC/3PC, and sagas — coordinating across nodes.

VOL 3 · DISTRIBUTED SYSTEMS, DEEPLY · PART XXI

Replicated Data & Convergence

Eventual consistency, anti-entropy, gossip, and CRDTs — converging without coordination.

VOL 3 · DISTRIBUTED SYSTEMS, DEEPLY · PART XXII

Partitioning & Scale

Consistent hashing, rebalancing, and hot-spot handling — spreading state and load.

VOL 3 · DISTRIBUTED SYSTEMS, DEEPLY · PART XXIII

Distributed Algorithms & Patterns

Broadcast/ordering, failure detectors, distributed snapshots, and the patterns real systems use.

VOL 3 · DISTRIBUTED SYSTEMS, DEEPLY · PART XXIV

Formal Methods

Specifying and model-checking a protocol with TLA+ — finding the bug before production does.

VOL 3 · DISTRIBUTED SYSTEMS, DEEPLY · PART XXV

Building & Operating Distributed Systems

From spec to a real, observable, testable distributed system — and operating it.

VOL 4 · SECURITY & PRIVACY ENGINEERING · PART XXVI

Security Foundations & Threat Modeling

The security mindset, threat modeling, trust boundaries, and reasoning about adversaries.

VOL 4 · SECURITY & PRIVACY ENGINEERING · PART XXVII

Applied Cryptography

Symmetric/asymmetric crypto, hashing, MACs, signatures, TLS, and key management — used correctly, not invented.

VOL 4 · SECURITY & PRIVACY ENGINEERING · PART XXVIII

Authentication & Authorization

Identity, sessions, OAuth2/OIDC, access-control models, and zero-trust — who can do what.

VOL 4 · SECURITY & PRIVACY ENGINEERING · PART XXIX

Software & Systems Security

Memory-safety bugs, injection, the supply chain, and secure-by-design coding — the common failure modes and their defenses.

Data Security

Encryption at rest and in transit, secrets management, tokenization, and secure key handling for data platforms.

Privacy Engineering

Differential privacy, anonymization vs pseudonymization, GDPR/CCPA, and data governance — privacy as an engineering discipline.

Secure Data Pipelines & Platforms

Applying the above to real data systems: lineage, access, multi-tenancy, and compliance-by-design.

Detection, Response & Operations

Logging, monitoring, detection, and incident response — assuming breach and limiting blast radius.

The Staff+ Role & Scope

Archetypes (tech lead, architect, solver, right-hand), what 'impact' means above senior, and the shift from output to leverage.

Technical Leadership & Strategy

Technical vision, strategy, and roadmaps — setting direction others can follow.

VOL 5 · THE STAFF/PRINCIPAL ENGINEER & TECH LEAD · PART XXXVI

Architecture & Design at Scale

Design review, ADRs, trade-off analysis, and evolving systems without a rewrite — the architect's judgment, applied.

VOL 5 · THE STAFF/PRINCIPAL ENGINEER & TECH LEAD · PART XXXVII

Influence Without Authority

Building alignment and buy-in, driving decisions across teams, and disagree-and-commit.

Communication & Writing

The written word as a staff engineer's main tool: design docs, RFCs, and making complex ideas land.

Mentoring & Growing Engineers

Mentorship, sponsorship, code/design review as teaching, and multiplying the people around you.

VOL 5 · THE STAFF/PRINCIPAL ENGINEER & TECH LEAD · PART XXXX

Execution & Judgment

Leading projects and incidents, making decisions under uncertainty, and knowing what NOT to do.

Career & the Organization

Navigating the org, the promotion case, sustainability, and a durable definition of technical leadership.

BACK MATTER

Solutions to Exercises

Worked solutions to the end-of-chapter exercises. Try each before reading its solution.

Chapter 1 — Solutions

1.1 Uniform cost: every basic operation — arithmetic, comparison, and any memory read or write — costs the same fixed unit of time, independent of which address is accessed. The three false assertions about real hardware: (1) all memory accesses cost the same (they vary by orders of magnitude with which tier of the hierarchy holds the data); (2) one instruction takes one unit of time (real CPUs overlap, reorder, and speculate, retiring several per cycle or stalling for hundreds — [Chapter 2](#)); (3) operations are independent and equally priced (some ops cost more than others, and a dependent instruction must wait for its input). *Method note:* the trap is listing symptoms ("caches exist") without naming the underlying assumption they violate — always tie the fact back to *uniform cost*.

1.2 Fastest to slowest: **register** (<1 ns) < **L1** (~1 ns) < **L2** (a few ns) < **L3** (~10 ns) < **DRAM** (~100 ns) < **NVMe SSD** (tens of μ s) < **same-datacenter network round trip** (hundreds of μ s to ~1 ms). All figures are orders of magnitude, machine- dependent (§1.3).

1.3 The memory wall is the large and historically widening gap between how fast a CPU can execute instructions and how long DRAM takes to return a requested byte, because processor throughput improved much faster than DRAM latency for decades (Wulf & McKee, 1995). It makes uniform cost worse over time because the ratio between "one arithmetic op" and "one DRAM access" grew, so pricing a DRAM access at the same "one unit" as an add gets more wrong every hardware generation, not less.

1.4 (a) A few billion ops/second is $\sim 3 \times 10^9$ /s, so in $100 \text{ ns} = 10^{-7} \text{ s}$ the core could do on the order of **a few hundred operations** ($3 \times 10^9 \times 10^{-7} = 300$). (b) A program dominated by random pointer-following spends most of its time *waiting* for DRAM, leaving the arithmetic units idle for hundreds of potential operations per miss — so it can run at a small fraction of the core's peak throughput even though its big-O looks fine. *Order-of-magnitude flags:* both the 100 ns and the few-billion-ops/s are approximate and machine-dependent; the point is the ratio ($\sim 10^2$), not the exact count. **Feed-forward:** if this surprised you, revisit §1.2.

1.5 Two explanations: (1) **Constant factors dominate at this size.** Big-O ignores the constant, and the $O(n)$ bucketing pass may do far more *slow* memory traffic per element (e.g. scattering writes across buckets spread through DRAM) than the $O(n \log n)$ sort, which may stream memory in a cache-friendly order (§1.4). (2) **Wrong regime.** The asymptotic advantage only shows up once n is large enough; on this dataset n may be small enough that the $\log n$ factor is tiny and the bucketing's larger constant wins for the wrong reasons (§1.4). The single settling action: **measure both on the real data** — big-O cannot answer a "how fast" question. *Tempting wrong answer:* "the $O(n)$ version must be faster, the benchmark is flawed" — this is trusting the RAM model for a speed question, exactly the §1.4 trap. **Feed-forward:** §1.4 and the warning box.

1.6 Step 4: The binary search over a small contiguous sorted array is likely to make *fewer trips to slow memory* per lookup — its handful of accesses are clustered and cache-friendly and the small array may live largely in cache, whereas the hash map's single lookup typically lands on a random heap address and pays a full cache/DRAM miss. So the asymptotically worse structure can win on real hardware because *the metric that matters here is costly-memory-accesses per lookup, not operation count*. The one thing to do before trusting the argument: **benchmark both on the real data and access pattern**. Why this final step and not "pick the lower big-O": big-O measures scaling and is silent about the constant factor — where the data lives — which is the entire basis of the argument; only measurement can confirm the memory-access reasoning holds for your actual sizes and hardware. **Feed-forward:** §1.4; the array-vs-pointer effect is measured in [Chapter 3](#).

1.7 (a) **RAM/big-O** — the tables grow without bound, so the deciding factor is scaling (quadratic vs linear), which is exactly what big-O captures. (b) **Real cost model** — same size, same $O(n)$; the $8\times$ gap can only come from constant factors, i.e. memory-access patterns. (c) **RAM/big-O** — "1000 $\times$ bigger" is a scaling question. (d) **Real cost model** — identical length and $O(n)$; the difference is contiguous streaming vs random pointer-chasing through the hierarchy. *Method-selection cue:* if the question mentions "grows / bigger / scales," reach for big-O; if it fixes the size and asks "why is one faster," reach for the real cost model.

1.8 The claims are **not in contradiction**. Claim (i) is about the RAM/big-O model (optimal *scaling*); claim (ii) is about the real cost model (*speed* at a given size, set by constant factors the memory hierarchy dominates). Both can be true simultaneously: two $O(n)$ algorithms are equally optimal asymptotically yet differ $10\times$ in wall-clock time. To choose for production you still need: the actual input sizes you will run at, and a **measurement** of both on representative data and hardware. *Tempting wrong answer:* "'optimal' means fastest, so (i) and (ii) conflict" — conflating asymptotic optimality with speed, the core confusion of §1.4.

1.9 Mechanical sympathy is writing code that works *with* the grain of the hardware — its memory hierarchy, pipeline, and I/O costs — rather than against it, by understanding enough of how the machine works to get the most from it. The ordering rule: choose a good *algorithm* first (a bad big-O cannot be rescued by micro-optimization — no cache trick fixes an accidental quadratic), *then* apply mechanical sympathy to win the constant factor. Failure modes: knowing only big-O $\rightarrow$ asymptotically optimal code that needlessly runs $10\times$ slow; knowing only micro-optimization $\rightarrow$ a lovingly hand-tuned $O(n^2)$ that dies at scale.

1.10 Three reasons the absolute numbers can't be trusted: (1) they are **microarchitecture-specific** — cache sizes, latencies, and cycle counts differ across CPU models and generations; (2) they **drift with hardware** over time, so any printed figure is dated; (3) measured latency depends on **context** — contention, frequency scaling, what else is resident in cache, and access pattern — so there is no single "the" latency. The ratios are more stable because they reflect the *structural* reason each tier exists (each is a deliberate size/speed tradeoff a step down from the last), and that layering persists even as every absolute number moves. To get trustworthy numbers: **measure on the actual target machine** under representative load (a microbenchmark or a profiler), which is the discipline the performance-engineering chapters build. **Feed-forward:** §1.3; measurement technique in [Chapter 3](#).

1.11 Answers are machine-specific by design; a correct write-up (i) reports real cache-line size (commonly 64 bytes on x86-64) and your actual L1/L2/L3 sizes from `getconf/lscpu`, (ii) reports two measured per-element times — the in-L1 loop should be markedly faster per element than the exceeds-L3 loop, which is bottlenecked on DRAM bandwidth — and (iii) states the ratio you observed. The learning goal is the habit: on a "how fast" question you *generated your own numbers* rather than quoting the book's orders of magnitude. Exact values depend entirely on your hardware; the *shape* (in-cache faster than out-of-cache) is what must appear. **Feed-forward:** the mechanism behind the ratio is §1.2-1.3 and all of [Chapter 3](#).

Chapter 2 — Solutions

2.1 In order: (1) **pipelining** — overlap the stages so different instructions occupy different units in the same cycle; (2) **superscalar** — duplicate execution units and issue several instructions per cycle; (3) **out-of-order** — reorder execution to run later independent instructions while an earlier one is stalled; (4) **speculation / branch prediction** — guess a branch's direction and execute down the predicted path so the pipeline need not wait for the branch to resolve.

2.2 Latency is the number of cycles for one instruction to pass through all stages; **throughput** is how often a new instruction *completes*. Because the stages overlap — while one instruction is in "execute," the next is in "decode," and so on — a new instruction can finish every cycle even though each individually spends several cycles traversing the pipeline. The two numbers measure different things: time-through-one vs. completions-per-time.

2.3 A data hazard: an instruction needs a result that an earlier, not-yet-finished instruction will produce, so it cannot proceed until that value is ready. **A control hazard:** the CPU does not yet know which instruction comes next because an unresolved branch determines it. **Branch prediction addresses the control hazard** — it guesses the branch direction so fetching can continue instead of stalling.

2.4 In loop A each add depends on the previous add's result (they all target one accumulator), forming a single dependency chain: the out-of-order engine can never have more than one add in flight, so the superscalar units sit mostly idle and the loop runs at roughly one add per dependency latency. In loop B the four partial sums are *independent*, so up to four adds can be in flight at once, keeping multiple execution units busy and finishing in a fraction of the time — same number of additions, more ILP. The enabling condition: floating-point addition is not associative, so B only computes the "same" result if you accept that the grouping (and thus rounding) differs, or the compiler is allowed to reassociate — B's parallelism is possible precisely because its adds have no data dependency between them. *Method-selection note:* the tell that ILP is the lever is "same operation count, one version serializes them." **Feed-forward:** §2.3 and the threeways box.

2.5 On P the data-dependent branch resolves in long predictable runs, so the branch predictor is almost always right, speculation is nearly free, and the pipeline stays full. On Q the branch is a patternless coin flip: the predictor is wrong about half the time, and each misprediction flushes the speculatively executed work and refills the pipeline — a penalty on the order of tens of cycles (machine-specific) paid on ~50% of iterations. The extra time on Q is spent flushing and refilling the pipeline, not doing arithmetic. *Tempting wrong answer:* "Q must be doing more comparisons" — no, the arithmetic is identical; the cost is misprediction, not extra work. **Feed-forward:** §2.4.

2.6 (a) Average cost per iteration $\approx b + p \cdot k$ cycles (caveat: k is an order-of-magnitude, microarchitecture-specific figure). (b) The misprediction term dominates when $p \cdot k > b$, i.e. roughly $p > b/k$; for a cheap body and a $\sim$ tens-of-cycles penalty, even a modest p (a branch that misses a nontrivial fraction of the time) makes misprediction the largest term. (c) Making the branch predictable lowers p (the predictor is right more often); a branchless rewrite removes the branch so there is no k to pay at all. You measure first because on an already-predictable branch p is already tiny, so a branchless rewrite adds complexity for no gain — and may even be slower if it does more arithmetic unconditionally. **Feed-forward:** §2.4 warning box.

2.7 Step 4: The miss stalls the CPU *only if there is no independent work to do while it is serviced*. The deciding property is the amount of **instruction-level parallelism in the surrounding code** — whether there are nearby instructions whose inputs are ready and do not depend on the missing load. If there is enough such work, the out-of-order engine hides the miss latency behind it; if the code is a tight dependency chain that all waits on the loaded value, the machine has nothing else to do and stalls for the full miss. It is that property, not the miss latency alone, that decides the outcome, because the latency is a fixed cost of reaching DRAM — what varies, and what you control by how you write the code, is whether the CPU can *overlap* useful work with it. **Feed-forward:** §2.3 (out-of-order) and Chapter 1 §1.2 (miss latency).

2.8 (a) **Branch prediction** — sorting changed only predictability, not arithmetic. (b) **Dependency chain / ILP** — splitting an accumulator exposes independent work. (c) **Memory hierarchy** — a linked list scatters accesses to random addresses (cache misses, Chapter 1) where an array streams contiguously. (d) None of the three is a bottleneck — no branches, no memory pressure, presumably enough ILP — so the loop is already near peak; the lesson is that when none of the mechanisms is being violated, there is little constant-factor slack left to recover. *Method-selection cue:* match the symptom — "sorting helped" $\rightarrow$ branches; "extra accumulators helped" $\rightarrow$ ILP; "layout/pointers hurt" $\rightarrow$ memory.

2.9 The **real cost model** (Chapter 1 §1.4) resolves it: equal big-O guarantees equal *scaling*, never equal *speed*; a $5\times$ gap lives entirely in the constant factor big-O discards. Three independent mechanisms that could each cause it: (1) **memory access pattern** — one version streams cache-friendly memory, the other chases pointers into DRAM (Chapter 1, [Chapter 3](#)); (2) **branch misprediction** — one has an unpredictable data-dependent branch flushing the pipeline (§2.4); (3) **dependency chains / low ILP** — one serializes its work and starves the superscalar units (§2.3). The single identifying action: **profile / measure** (e.g. with a profiler reporting cache misses and branch mispredictions) — the mechanism is an empirical question, not one big-O can answer. *Tempting wrong answer:* "it's impossible / the benchmark is wrong," which is the Chapter 1 trap of trusting big-O for a speed question. **Feed-forward:** §2.5 and Chapter 1 §1.4.

2.10 (a) On typical code branches are highly predictable, so speculation is almost always correct and the CPU keeps working instead of stalling several cycles at every branch — the occasional wasted work is far cheaper than stalling on every branch. (b) On an unpredictable branch the predictor is wrong a large fraction of the time, so much speculative work is discarded and the misprediction penalty dominates — a net loss versus not having to speculate at all. (c) At a high level: because speculatively executed instructions can affect microarchitectural state (such as what ends up in the cache) even when their results are later discarded, that residual state can, in principle, be observed and used to infer data that should not have been accessible — this is the basis of the real speculative-execution vulnerability class. The precise mechanisms and mitigations are a genuine, separate topic (touched on in Volume 4); describing them accurately requires primary sources, and this answer deliberately does not invent specifics. *No-fabrication note*: full credit requires the high-level reasoning *and* the acknowledgement that the details need real sources, not invented ones.

2.11 Machine- and compiler-specific by design. A correct write-up reports two measured times (single vs. four accumulators) and their ratio, with the four-accumulator version typically faster when the compiler has not already reassociated or vectorized the single-accumulator loop. It must state which numbers are the reader's own and note the ratio depends on CPU and compiler. Crucially, if optimization collapses the difference (the compiler reassociates or auto-vectorizes both), the correct response is to *report that* and describe what was observed — never to quote a ratio that was not actually measured. *Method note*: the exercise is as much about honest measurement discipline (no fabricated numbers) as about ILP. **Feed-forward**: §2.3 threeways box.

Chapter 3 — Solutions

3.1 $64 / 8 = 8$ **doubles** per line. Reading one from a cold line brings in the whole line, so the **7 neighbors** become cheap hits. This rewards **spatial locality** — you get the nearby elements for free because memory arrives a full line at a time.

3.2 Temporal: a loop accumulator `sum += x[i]` — `sum` is reused every iteration, so after its first access it stays resident (retention of recently used lines). **Spatial:** walking `x[0]`, `x[1]`, `x[2]`, ... — consecutive elements share lines (line fill) and the sequential pattern lets the *prefetcher* fetch ahead. Line fill and prefetching exploit spatial locality; keeping recent lines exploits temporal.

3.3 (a) **Compulsory (cold)** — first-ever reference to those lines. (b) **Capacity** — the 4 GB working set vastly exceeds the 8 MB cache, so lines are evicted before reuse. (c) **Conflict** — a specific power-of-two row length aligns many accesses onto the same cache sets, so they evict each other despite free space elsewhere.

3.4 (a) **Inner loop over columns (row-major order) is faster:** consecutive accesses are adjacent in memory (stride 1), so each 64-byte line is fully used and the prefetcher runs ahead. Inner loop over rows strides by N elements, touching a new line almost every access. (b) Same big-O and same additions govern *scaling*, not *speed*; the constant factor — memory moved and misses incurred — differs by roughly an order of magnitude because only one version has spatial locality (this is the Chapter 1 §1.4 point, now concrete). (c) The slow version suffers mostly **capacity misses** (the whole matrix far exceeds cache, and its poor locality means it re-fetches lines it effectively wasted) — with a real **conflict** component too, since the power-of-two dimension ($N = 8192$) gives a power-of-two stride that aliases into a few cache sets (the very symptom §3.4 described); its lack of spatial reuse is what turns each fetch into wasted bandwidth. *Tempting wrong answer:* "they're both $O(N^2)$ so the difference must be a measurement error" — the §3.6 trap. **Feed-forward:** §3.6.

3.5 Replace the array-of-pointers with a contiguous **array of the objects themselves** (array of structs, or a struct-of-arrays if you only touch some fields), laid out in traversal order. Why it reduces misses: the objects now share cache lines, so walking them in order has spatial locality — each fetched line yields several objects instead of one random-heap object per pointer (fewer misses). Why it helps hide the rest: sequential, contiguous access is prefetchable and, unlike pointer-chasing, does not force a dependency chain where each object's address depends on reading the previous one — so the out-of-order engine (Chapter 2 §2.3) can overlap independent accesses and hide the misses that remain. *Method note:* the two wins are distinct — locality (fewer misses) and independence (hide the misses) — and the pointer layout loses on both. **Feed-forward:** §3.5.

3.6 (a) 8 bytes used of 64 fetched = **1/8** of each line. (b) A stride-1 traversal uses all 8 values per line, so the stride-8 traversal moves roughly **8× more memory** for the same values. (c) That $\sim 8\times$ is a floor on the slowdown from wasted bandwidth alone; the measured row-major/column-major ratio was larger ($\sim 24\text{--}26\times$) because real effects compound it — two that push the measured number above the naive estimate: (i) the hardware **prefetcher** helps the sequential (row-major) loop even more, widening the gap; (ii) **loss of memory-level parallelism / stalling** on the strided loop, and possibly **TLB misses** from touching many pages, add cost the simple bytes-moved model ignores. *No-fabrication note:* the $8\times$ is arithmetic; the $\sim 25\times$ is measured — do not present the estimate as the measurement. **Feed-forward:** §3.6.

3.7 Step 4: In this particular experiment the dominant lever is the **number of misses (locality)**, not the inability to hide each one. The column-major loop touches a new line almost every iteration and uses only 1 of 8 values per line, so it issues roughly $8\times$ as many line fetches and moves many times more memory than the row-major loop; that sheer volume of DRAM traffic is the main cost. The dependency argument is secondary here because both versions are simple predictable summations with similar per-iteration dependency structure — what differs overwhelmingly is how many expensive lines each one demands. The number of misses is primary *because the two loops differ in locality but not meaningfully in their dependency chains* — so the variable that changed is misses. To test the claim: measure cache misses and memory traffic for both (e.g. with a profiler's cache-miss counters) and confirm the column-major loop's miss count is the multiple that tracks its slowdown. **Feed-forward:** §3.5–3.6, Chapter 2 §2.3.

3.8 (a) **Capacity miss fixed by blocking** — tiling keeps each block's working set inside cache so it is reused before eviction. (b) **Branch prediction** — sorting makes the filter's data-dependent branch predictable (Chapter 2 §2.4). (c) **Spatial locality / cache lines** — consuming each line fully raises useful-bytes-per-line. (d) **ILP / dependency chain** — independent accumulators expose parallelism (Chapter 2 §2.3). *Method-selection cue:* "tiling/ blocking" → capacity; "sorting helped a branch" → prediction; "layout/stride" → spatial locality; "extra accumulators" → ILP.

3.9 Candidate causes across Chapters 1–3 (at least three): **poor spatial locality / high miss rate** (§3.1–3.5); **capacity misses** from a working set exceeding cache, fixable by blocking (§3.3); **branch mispredictions** from a data-dependent branch (Ch 2 §2.4); **a serial dependency chain / low ILP** (Ch 2 §2.3); pointer-chasing combining bad locality and dependencies (§3.5). The first measurement: **run a profiler that reports cache-miss rate and branch-misprediction rate** — it is most informative because it directly distinguishes the two biggest mechanism families (memory vs. branches) in one shot, telling you which chapter's toolkit to reach for before you change any code. *Tempting wrong answer:* "start rewriting the layout" — that guesses the cause; measure first (the discipline of this whole Part). **Feed-forward:** §3.6 and Chapter 1 §1.4.

3.10 (a) Three properties and their effect on the ratio: **cache/last-level-cache size** — a much larger cache (or a smaller matrix) shrinks the gap by keeping more of the working set resident; **memory bandwidth and prefetcher quality** — a stronger prefetcher and higher bandwidth help the sequential loop more, tending to *raise* the ratio; **compiler behavior** — if the compiler vectorizes or reorders one loop, the ratio can move either way. (b) The effect's existence and rough size are robust because they follow from a structural fact true of essentially all cache-based machines: memory moves in lines, and a strided traversal wastes most of each line — so a large (not small) multiplier is guaranteed even though its exact value tracks the specific hardware. (c) Reporting "about 25× on this specific CPU, reproduce it yourself" is honest because the number was measured on one machine and genuinely varies; it is more useful because it teaches the reader the *method* (measure on your own hardware) rather than a false universal constant they might quote as fact — which is exactly the no-fabrication discipline of this book. **Feed-forward:** §3.6 and the warning box.

3.11 Machine-specific by design. A correct write-up records the reader's real cache-line size (commonly 64 bytes) and L1/L2/L3 sizes; confirms both loops produce the identical sum (proving equal work); reports two measured times and the observed ratio; and compares that ratio to the book's ~25× and to the ~8× stride estimate from §3.6, expecting it to sit somewhere around or above the estimate and to vary with hardware. The essential discipline: report only measured numbers, and if the compiler elides a loop, change the code (consume the result / use a `volatile` sink) and say what was done — never quote a number you did not observe. **Feed-forward:** §3.6; deeper measurement in the performance-engineering chapters.

Chapter 4 — Solutions

4.1 (a) **12 bits** ($2^{12} = 4096 = 4 \text{ KiB}$). (b) **One translation**: `a[0]` and `a[1]` sit on the same 4 KiB page (unless `a[0]` happens to be the last element of a page), so the same page number translates both. (c) A 4 KiB page holds $4096/8 = 512$ **longs**, so up to ~512 consecutive elements share one translation before you cross into the next page. This is the TLB analogue of the cache-line reuse from §3.1 — one translation amortized over many accesses.

4.2 Isolation — each process names only its own frames, so it cannot read or corrupt another's memory (this is the security/isolation boundary). **A simple, uniform programming/layout model** — every process sees the same clean virtual layout regardless of where it lands in physical RAM. **OS flexibility** — a virtual page can map to a frame, a file, nothing yet (lazy allocation), or be shared read-only. The isolation one is the protection boundary.

4.3 A data cache stores recently used **cache lines** so a repeated data access is a cheap hit; the TLB stores recently used **translations (virtual-page→physical-frame mappings)** so a repeated access to the same **page** skips the **page-table walk**. Same hit/miss/eviction logic (§3.3), different key and value: address→line for the data cache, page→frame for the TLB.

4.4 A multi-level walk reads one table to learn the address of the next table, reads that to learn the next, and so on — **each read's address depends on the value returned by the previous read**, so the reads cannot be issued in parallel or overlapped; they are strictly serial. That matters because a dependency chain defeats the CPU's latency-hiding (Chapter 2 §2.3): there is no independent work to run during each memory access, so the walk pays the full latency of every step in sequence. This is exactly the **pointer-chasing** structure of a linked-list traversal (§3.5): in both, you cannot compute the next address until the current access completes, so both are serial-latency-bound rather than bandwidth-bound. *Method note*: recognizing "the next address depends on this read" is the general tell for a latency-bound pattern. **Feed-forward**: §3.5, Chapter 2 §2.3.

4.5 (a) A random probe pays (i) a **data-cache miss** to fetch the line holding the bucket (~100 ns to DRAM) and (ii) a **TLB miss / page-table walk** to translate the bucket's address before the fetch — several dependent memory accesses. (b) With 4 KiB pages, TLB reach is (entries × 4 KiB), typically only a few megabytes; a 400 MiB table spans far more pages than the TLB can hold, so random probes keep landing on pages whose translations were evicted — heavy TLB misses (a capacity miss at the translation level, §4.4). (c) To cut translation cost: map the table with **huge pages** (raises TLB reach ~512×), assuming a profiler confirms page-walk cost. To cut the data-miss cost: improve **locality** of the buckets (e.g. keys and values stored together, open addressing with good probe locality, or a more cache-friendly layout) so more probes hit. *Tempting wrong answer*: "it's all just DRAM latency, nothing to do" — the §4.4 trap; the surplus over ~100 ns is the translation cost. **Feed-forward**: §4.4–4.5.

4.6 (a) **reach = $E \times 4 \text{ KiB}$** . (b) Expect heavy thrashing when the randomly-touched footprint exceeds reach: **$W > E \times 4 \text{ KiB}$** (more precisely, when the number of distinct pages touched, $W / 4 \text{ KiB}$, far exceeds E), so translations are evicted before reuse. (c) Switching to 2 MiB pages multiplies reach by **$2 \text{ MiB} / 4 \text{ KiB} = 512\times$** , so the threshold in (b) becomes **$W > E \times 2 \text{ MiB}$** — a working set $512\times$ larger now fits within reach, which is exactly why huge pages relieve TLB thrashing. *No-fabrication note*: the $512\times$ is arithmetic from real, documented page sizes; the actual E is a chip-specific number this book does not invent. **Feed-forward**: §4.5.

4.7 Step 4: The random-by-page walker needs a *new* translation on essentially **every jump**, because each jump targets a different page and it does not stay on any page long enough to reuse a translation. Since the TLB's reach (a few MiB) is far smaller than the 512 MiB footprint, the translation for a page it jumps to was almost certainly evicted since the last visit — so nearly every random access **misses the TLB and pays a page-table walk**, on top of the data-cache miss. The sequential walker, by contrast, amortizes one translation over ~ 512 accesses (Step 3), so its TLB-miss *rate* is ~ 1 per 512 accesses versus ~ 1 per access for the random walker. To confirm the TLB — not just the data cache — is the culprit, measure the **dTLB-load-miss (or page-walk) count/rate** for both runs (e.g. `perf stat -e dTLB-load-misses`): the random run should show dramatically more, isolating translation cost from data-fetch cost. **Feed-forward**: §4.4, and the Chapter 6 capstone.

4.8 (a) **TLB miss / page-walk** — huge pages helping is the signature; they raise reach, not data-cache capacity. (b) **Data-cache capacity miss** — 64 MiB exceeds the 24 MiB L3 so it lives in DRAM, while 1 MiB fits in L2; same accesses, different tier (Ch 3 §3.3). (c) **Dependency chain / low ILP** combined with poor locality — pointer-chasing a list is serial and scattered (§3.5); the standout distinguishing feature here is the serial dependency. (d) **Branch misprediction** — sorting makes the data-dependent branch predictable (Ch 2 §2.4). *Method-selection cue*: "huge pages helped" → TLB; "bigger set, re-read" → capacity; "linked list" → dependency/locality; "sorting fixed a branch" → prediction.

4.9 (a) Huge pages help specifically when **TLB misses / page-walk cost are a measured bottleneck** — i.e. a large, scattered working set whose page footprint exceeds TLB reach. (b) Downsides: internal **fragmentation** / wasted memory (a 2 MiB page backing a few live KiB); difficulty **finding contiguous physical memory** once RAM is fragmented (allocation can fail or stall); and no benefit — pure overhead — for workloads that were never TLB-bound. (c) First measurement: the **TLB-miss / page-walk rate** (e.g. `perf stat -e dTLB-load-misses`); if it is low, huge pages will not help. *Discipline*: this is the "measure before you optimize" rule of Part 1 — don't flip huge pages on faith. **Feed-forward**: §4.5 and the Chapter 6 lab.

4.10 (a) **Compulsory**: yes — the first access to a page's translation always misses the TLB (it was never cached), just as the first touch of a line is a cold miss. **Capacity**: yes — when the page footprint exceeds TLB reach, translations are evicted before reuse; this is TLB thrashing (§4.4). **Conflict**: an analogue exists for set-associative TLBs (translations colliding on the same TLB set), though it is secondary and hardware-specific; the dominant effect for large working sets is capacity. (b) Huge pages attack the *capacity* analogue: they do not add TLB entries but make each entry cover far more memory, so the same footprint needs fewer translations and fits within reach. (c) Larger base pages are not a free system-wide win because programs with small, sparse footprints would waste memory to internal fragmentation (rounding every allocation up to a huge page) and the OS would find it harder to place them; the 4 KiB default is a compromise, and huge pages are opt-in for the workloads that benefit. *Tempting wrong answer*: "just make all pages 2 MiB" — ignores fragmentation and allocation cost. **Feed-forward**: §4.5.

4.11 Machine-specific by design. A correct write-up records the real page size (commonly 4 KiB) and cites the architecture's documented page-table depth and huge-page sizes (do not guess — for x86-64, a four-level table with 2 MiB and 1 GiB huge pages); confirms both loops touch the same data the same number of times; and reports two measured ns-per-access figures showing the random-by-page run slower. Where TLB counters are available, the random run should show far more dTLB-load-misses / page-walks than the sequential run, isolating translation cost. The essential discipline: report only measured numbers, attribute the architecture facts to real documentation, and if the compiler elides a loop, consume the result (a `volatile` sink) and say so. **Feed-forward**: §4.4, and the Chapter 6 capstone which measures exactly this latency-bound gap.

Chapter 5 — Solutions

5.1 Array-of-structs (AoS): one array whose elements are whole records, so a record's own fields are adjacent in memory. **Struct-of-arrays (SoA):** one array per field, so the same field across different records is contiguous. **SoA** stores "the same field across different records" contiguously.

5.2 (a) 8 of 64 bytes — one double per fetched line. (b) SoA uses all 8 values per line, so AoS moves roughly **8× more memory** for the same field. (c) SoA restores **spatial locality**: consecutive values of the scanned field are adjacent, so each line is full of values the loop uses (§3.1–3.2).

5.3 A hot loop reads only a couple of a record's many fields, so most of every fetched cache line is fields the loop never touches — dead weight that lowers useful-bytes-per-line. (Split the few hot fields into their own compact record; move the rest to a cold side-structure.)

5.4 (a) AoS — reads/writes all fields of one particle before moving on, so a record's adjacent fields are all used from one or two fetched lines. (b) **SoA** — summing one column over ten million rows is a stride-1 scan of that field; contiguous storage uses every byte of every line. (c) **SoA** — a SIMD kernel over one field wants that field contiguous so a vector load grabs consecutive values. (d) **AoS** — position, color, and size of one object are used together, so keeping a record's fields adjacent brings them in one fetch. *Method-selection cue*: "one field across records / SIMD one channel" → SoA; "whole record at a time" → AoS.

5.5 (a) The `double` needs 8-byte alignment, so the compiler inserts **padding** after `flag` to place `value` at an 8-byte offset, and adds tail padding so an array keeps every element aligned — pushing the size above 13 bytes (commonly to 24). (b) Order large/strictly-aligned first: `struct S { double value; int id; char flag; };` — `value` at offset 0, `id` at 8, `flag` at 12, then a little tail padding to the struct's alignment; this minimizes internal gaps. (c) A smaller struct means **more records per cache line and per page**, so a scan over ten million of them fetches fewer lines and touches fewer pages — better spatial locality and less TLB pressure, for free, from field ordering alone. *Tempting wrong answer*: "field order doesn't affect layout in C" — it does, through alignment padding. **Feed-forward**: §5.4.

5.6 (a) The scan uses **8 useful bytes of each 64-byte line** (assuming one field per record on a line), a fraction of $8/64 = 1/8$ when $S \geq 64$; for smaller S a line may hold more than one record's copies but you still touch one field's worth per record. (b) As S grows past the line size, each record's scanned field sits on its own line, so the AoS scan moves $\sim S/8$ lines' worth where SoA moves 1 — the bandwidth ratio grows with S (bounded by one line fetched per record). (c) The measured speedup can be smaller than the bandwidth ratio because the SoA loop may be limited by something other than bandwidth — e.g. a **single-accumulator dependency chain** (Ch 2 §2.3) that caps how fast it can consume values, so it cannot fully cash in the saved bandwidth. That is exactly why this chapter's demo measured $\sim 4\times$, not the naive $8\times$; removing the accumulator ceiling widens it. *No-fabrication note:* $8\times$ is the arithmetic ceiling; $\sim 4\times$ is what was measured — keep them distinct. **Feed-forward:** §5.2 and Chapter 6.

5.7 Step 4: Decide empirically. **Measure** where the runtime actually goes — profile both loops (or time them separately) to learn which dominates and by how much; the layout should favor whichever loop owns the most wall-clock time. If loop A (scan) dominates, lean SoA; if loop B (whole-record) dominates, lean AoS. A **hybrid** can serve both partly: hot/cold-split so the field loop A scans lives in its own contiguous array (SoA for that field) while the fields loop B updates together stay grouped — or duplicate/derive the scanned field into a packed side-array refreshed when needed. The posture "pick the layout that helps the hotter loop, then re-measure" is right because the interaction of layout with the prefetcher, the accumulator dependency, and the two loops' real proportions is hard to predict on paper (the demo's $\sim 4\times$ -not- $8\times$ is proof), so the paper argument only picks the *candidate* and the measurement decides. This reflects the **measure-before-you-optimize** discipline of Part 1. **Feed-forward:** §5.2–5.3, Chapter 6.

5.8 (a) **Hot/cold splitting** — moving a cold field out raises useful-bytes-per-line for the hot loop. (b) **SoA / whole-line use** — column arrays make the column scan stride-1. (c) **Capacity miss fixed by blocking** — tiling keeps each block resident (Ch 3 §3.3). (d) **TLB / huge-page effect** — 2 MiB pages raise TLB reach for the random index (Ch 4 §4.5). (e) **ILP / multiple accumulators** — independent accumulators break the dependency chain (Ch 2 §2.3). *Method-selection cue:* "cold field moved" → hot/cold; "rows→columns" → SoA; "tiling" → capacity/blocking; "huge pages" → TLB; "more accumulators" → ILP.

5.9 Candidate causes across Chapters 2–5 (at least three, incl. a layout cause and an earlier one): **AoS-when-scanning-one-field / poor useful-bytes-per-line** (§5.2); **cold fields bloating hot records**, fixable by hot/cold splitting (§5.3); **a capacity miss** from a working set exceeding cache (Ch 3 §3.3); **TLB misses** on a large scattered footprint (Ch 4); **a single-accumulator dependency chain / low ILP** (Ch 2 §2.3); a **data-dependent branch** misprediction (Ch 2 §2.4). First measurement: run a **profiler reporting cache-miss rate (and, if available, TLB-miss and branch-miss rates)** — it distinguishes the big mechanism families in one shot, telling you whether to reach for a layout change, a blocking change, or a branch/ILP fix before touching code. *Tempting wrong answer:* "rewrite the layout first" — that guesses; measure first. **Feed-forward:** §5.2 and Chapter 6.

5.10 (a) Both follow from one fact: **a whole cache line moves as a unit**. Single-threaded, you want the bytes you use together on one line (packing) so one fetch serves them all. Multi-threaded, if two cores write different variables on one line, that same all-or-nothing line movement forces the line to bounce between caches on every write (false sharing). Same mechanism, opposite consequence depending on who writes. (b) Deliberately pad a struct up to a full line for **per-thread counters/state written by different threads** — pad so each thread's variable sits on its own line, eliminating false sharing; you trade away memory (and some single-thread locality) to stop the coherence ping-pong. (c) The choice depends on core count because false sharing only occurs when *multiple cores write* the shared line; with one writer, packing is pure win — so "how many cores write this?" decides whether to pack or spread. *Tempting wrong answer*: "packing is always good for cache performance" — true single-threaded, false under concurrent writers. **Feed-forward**: §5.5 and the concurrency Part.

5.11 Machine-specific by design. A correct write-up records the cache-line size, defines a line-sized record, confirms both sums are identical (equal work), reports two measured times and the observed SoA/AoS ratio, and compares it to the book's $\sim 3.7\text{--}4\times$ and to the $\sim 8\times$ bandwidth ceiling — expecting a several-fold win that varies with hardware. The second part should show the ratio widening toward $8\times$ once the accumulator dependency is removed (multiple accumulators / vectorization), confirming the demo's explanation for why it measured $\sim 4\times$. Essential discipline: report only measured numbers, and if a loop is elided, consume the result and say what was done. **Feed-forward**: §5.2, §5.6, and the Chapter 6 capstone.

Chapter 6 — Solutions

6.1 Latency is the time to complete *one* access you must wait for; **bandwidth** is the throughput when streaming *many* contiguous items. Use **latency** for (a) fetching one random element (a single wait) and **bandwidth** for (b) streaming a huge array in order (the prefetcher keeps many accesses in flight).

6.2 (a) **Memory-bound** — `memcpy` does no arithmetic per byte; it is pure data movement at memory bandwidth. (b) **Compute-bound** — SHA-256 does much arithmetic per byte and the data is already in L1, so the execution units are the limit. (c) **I/O-bound** — the time is dominated by per-row network/SSD latency, not CPU or memory. *Method-selection cue*: "work per byte" — near zero → memory; high → compute; waiting on storage/network → I/O.

6.3 Estimate (predict the class and rough cost from the cost model), **measure** (run it and observe the real time / counters), **reconcile** (explain any gap and update the model). A gap is useful because it usually reveals a wrong assumption — most often using latency where bandwidth applies (or vice versa) — so the discrepancy points directly at the true bottleneck.

6.4 (a) 512 MiB of `long` is 2^{26} = ~67 million elements. Latency estimate: $67\text{M} \times 100\text{ ns} \approx \sim\mathbf{6.7\text{ s}}$. Bandwidth estimate: $512\text{ MiB} / 6\text{ GB/s} \approx \sim\mathbf{0.09\text{ s}}$. (b) The **bandwidth** estimate is right for a sequential sum: the prefetcher streams the data, so accesses overlap rather than each incurring the full latency (this matches the ~0.08–0.09 s measured in §6.3). (c) They differ by roughly **~70–100×** — the cost of using the latency model for a bandwidth-bound workload is being wrong by about two orders of magnitude, exactly the §6.3 surprise. *Tempting wrong answer*: the ~6.7 s latency figure — tempting because ~100 ns is the "known" DRAM number, but it is the wrong number for streaming. **Feed-forward**: §6.3.

6.5 (a) The sum is stride-1, so the hardware prefetcher recognizes the pattern and fetches lines ahead of the loop, keeping many independent accesses in flight; the loop is then limited by how fast memory can *stream* (bandwidth), not by any single access's latency. (b) Even the L2 sum is capped by the **single-accumulator dependency chain** (§2.3): each `sum +=` depends on the previous, so the additions serialize at roughly one add-latency apart, limiting throughput below what the cache could supply. (c) Sum into **several independent accumulators** (or let the compiler vectorize), breaking the dependency chain so more values are consumed per cycle and the loop can approach memory bandwidth. *Method note*: the two limiters — prefetchable bandwidth and the add dependency — are distinct; naming both is the full explanation. **Feed-forward**: §6.3 and Chapter 2 §2.3.

6.6 (a) **compute-time** $\approx F / 10^{10}$ s; **memory-time** $\approx B / 10^{10}$ s; the runtime floor is the larger of the two. (b) Compute-bound when compute-time > memory-time, i.e. $F / B > 1$ op per byte for these rates (more generally, when work per byte exceeds the machine's ops-per-byte balance point). (c) A sum does 1 add per 8 bytes, so $F/B = 1/8 < 1 \rightarrow$ **memory-bound** by this test. §6.3 agrees: the sum's speed tracked memory (GB/s), not the add rate — though the single-accumulator dependency kept it from reaching full bandwidth, which is why the measured GB/s was modest. *No-fabrication note:* the 10^{10} rates are round orders of magnitude for estimation, not a spec of any particular chip. **Feed-forward:** §6.2–6.3.

6.7 Step 4: ~ 150 ns/step exceeds the ~ 100 ns bare DRAM latency because the random chase over 512 MiB also **thrashes the TLB** (Chapter 4): each jump lands on a new page whose translation has been evicted, so most steps pay a **page-table walk** — extra dependent memory accesses to translate the address — *on top of* the data miss. So the ~ 150 ns is roughly the DRAM data latency plus the translation cost. To confirm the surplus is translation and not data, measure the **dTLB-load-miss / page-walk rate** (e.g. `perf stat -e dTLB-load-misses`) and check it is high for the random run and negligible for the sequential one. The pair teaches the rule: estimate with **bandwidth** for prefetchable sequential access (§6.3's $\sim 1.7\times$) and with **latency** — plus translation cost — for scattered, dependent access (this $\sim 25\times$); the access pattern, not just the data's tier, chooses the model. **Feed-forward:** §6.3–6.4, Chapter 4 §4.4.

6.8 (a) **Memory-bound**, memory bandwidth saturated — a single core's sequential scan already uses much of the available bandwidth, so more threads add little (the bottleneck is the memory pipe, not cores). (b) **Compute-bound-ish / control**, branch misprediction — sorting made the data-dependent branch predictable (Ch 2 §2.4), cutting flushes. (c) **Memory-bound**, TLB miss / page-walk — the ~ 160 ns over ~ 100 ns and the huge-page fix are the TLB signature (Ch 4). (d) **Compute-bound**, execution throughput — the multiply-add issue rate limits an in-cache matmul (Ch 2). *Method-selection cue:* "adding threads doesn't help a scan" $\rightarrow$ bandwidth-bound; "sorting fixed it" $\rightarrow$ branch; "huge pages fixed it" $\rightarrow$ TLB; "issue rate limited" $\rightarrow$ compute.

6.9 (a) The $100\times$ is a *latency* ratio, but a sequential scan is prefetchable and therefore **bandwidth-bound**, so latency is the wrong model — the actual slowdown vs. cache is a small factor, not $100\times$ (§6.3). (b) Rough class: **memory-bound (streaming)**; rough slowdown vs. an in-cache version: on the order of a $\sim 2\times$, since DRAM streaming bandwidth is within a small factor of cache (and it may be further gated by the file/I-O layer for a file-backed array). (c) Confirm by measuring (i) the scan's **effective GB/s** (is it near memory bandwidth?) and (ii) its **cache-miss / prefetch behavior** (or simply compare to the same scan on an in-cache slice) before optimizing. *Tempting wrong answer:* accept the $100\times$ and build a cache — that optimizes a bottleneck that isn't there. **Feed-forward:** §6.3 and the warning box.

6.10 (a) The **random** ratio is more portable because it is bounded by DRAM *latency* (plus translation), which is set by physics and the memory subsystem and varies modestly across machines; the **sequential** ratio is bounded by the *bandwidth* gap between L2 and DRAM and by loop details (accumulators, vectorization, prefetcher quality), which vary a lot — so the sequential number is the fragile one. (b) Raise the sequential ratio: a stronger prefetcher / higher L2 bandwidth relative to DRAM (widens the streaming gap), or removing the accumulator dependency so the L2 loop speeds up more than the DRAM loop. Lower the random ns/step: faster DRAM / lower memory latency, a larger TLB or huge pages (cuts page-walk cost), or a bigger last-level cache that keeps more of the footprint on-chip. (c) Reporting both ratios with their access patterns is honest because there is no single "DRAM is $N\times$ slower" — the answer depends entirely on whether access is streaming or scattered; a lone headline number invites readers to apply it to the wrong pattern and mis-estimate by orders of magnitude, which is exactly the §6.3 trap. *Tempting wrong answer:* "just quote the bigger, scarier ratio" — that mispredicts every bandwidth-bound workload. **Feed-forward:** §6.3–6.4.

6.11 Machine-specific by design, and the write-up matters as much as the numbers. A correct submission (1) records real L2/L3 and cache-line sizes; (2) states a *written-first* prediction with the expected bound (sequential → bandwidth-bound, small ratio; DRAM chase → latency-bound, large ratio); (3) reports measured ns/element and GB/s for both sums (identical sums verified); (4) reports ns/step for both pointer-chases; and (5) reconciles each measurement against the prediction — expecting a modest sequential ratio (bandwidth) and a large random one (latency plus TLB), ideally backed by `perf` cache-miss and dTLB-miss counters showing the random DRAM run dominated by misses and page-walks. The essential discipline: predict before measuring, report only observed numbers, explain surprises with latency-vs-bandwidth and the TLB rather than hand-waving, and if a loop is elided, consume the result and say so. *Tempting wrong answer:* skipping the written prediction and only reporting measurements — you lose the reconciliation that is the whole point. **Feed-forward:** §6.3–6.4 and the performance-engineering Part.

Chapter 7 — Solutions

7.1 (a) A pointer variable stores a **memory address** (a virtual address, per Ch 4). (b) ***p dereferences** — goes to the address in `p` and reads or writes the object there, using `p`'s type to know how many bytes and how to interpret them. (c) `&x` produces the **address of** `x` (a pointer to `x`).

7.2 As a double `*`: `q + 4 = 2000 + 4 × sizeof(double) = 2000 + 4 × 8 = 2032`. As an int `*`: `2000 + 4 × sizeof(int) = 2000 + 4 × 4 = 2016`. The rule: `p + n` advances by `n × sizeof(*p)` bytes, so the byte step is the element size. *Method note*: always multiply the integer by the pointee's `sizeof`, never treat it as a byte offset. **Feed-forward**: §7.2.

7.3 `sizeof(a) = 40` (10 ints × 4 bytes — the whole array); `sizeof(p) = 8` (just the pointer, on this x86-64 machine). Difference: `a` names the array object, so `sizeof` measures all its storage; `p` is a separate pointer variable, so `sizeof` measures the pointer, not what it points at. *Tempting wrong answer*: "both are 40 because `p` points at the array" — no, `sizeof` looks at the operand's type, and `p`'s type is `int *`. **Feed-forward**: §7.3.

7.4 Array-to-pointer decay: in most expression contexts an array name is automatically converted to a pointer to its first element (C §6.3.2.1), which is why `arr` can be assigned to an `int *`, indexed, or passed to a function taking `int *`. In `size_t len(int a[]) { return sizeof(a)/sizeof(a[0]); }` the parameter `a` is *already a pointer* (a function parameter of "array" type is adjusted to pointer type; the array decayed at the call), so `sizeof(a)` is the **pointer size** (8 here), not the array size. The returned value is `8 / sizeof(int) = 2` on this machine, regardless of the caller's real length. Correct approach: pass the length explicitly. *Tempting wrong answer*: "it returns the element count" — it cannot; the length was lost at the call boundary. **Feed-forward**: §7.3.

7.5 `struct S { char a; short b; char c; int d; };` with `align = size (char 1, short 2, int 4)`: `a` at offset 0 (byte 0); `b` (align 2) needs an even offset, next is 2, so 1 byte padding at 1, `b` occupies 2-3; `c` at offset 4 (byte 4); `d` (align 4) needs a multiple of 4, next after byte 4 is 8, so 3 bytes padding at 5-7, `d` occupies 8-11. Struct alignment = `max(1,2,1,4) = 4`; running size 12 is already a multiple of 4, so no tail padding. **sizeof(struct S) = 12**; real data = `1+2+1+4 = 8`, so **4 bytes padding**. Verify with `offsetof(struct S, member)` for each member and `sizeof(struct S)`. *Tempting wrong answer*: "`1+2+1+4 = 8`" — ignores alignment padding. **Feed-forward**: §7.5.

7.6 Order largest-alignment first: `struct S { int d; short b; char a; char c; };` — `d` at 0 (0-3), `b` at 4 (4-5), `a` at 6, `c` at 7; running size 8, alignment 4, 8 is a multiple of 4, so no tail padding. **sizeof = 8** — the size of the real data, zero padding. Rule applied: placing members in decreasing alignment order minimizes the internal gaps needed to re-align each successive member (and packs the small fields into what would otherwise be padding). *Tempting wrong answer*: "field order can't matter in C" — it does, via alignment padding. **Feed-forward**: §7.5, and §5.4.

7.7 (a) $10\text{ M} \times 24\text{ bytes} = \mathbf{240\text{ MB}}$. (b) $10\text{ M} \times 16 = \mathbf{160\text{ MB}}$. A 64-byte line holds $64/24 = 2$ whole 24-byte records vs. $64/16 = 4$ whole 16-byte records — **2 more per line**; a 4 KiB page holds $4096/24 = 170$ vs. $4096/16 = 256$ — **86 more per page**. (c) The smaller record raises useful-bytes-per-line and records-per-page, so a scan fetches fewer cache lines and touches fewer pages — better spatial locality and less TLB pressure (Ch 3–5), for free, from field ordering. *Tempting wrong answer*: ignoring that a line/page holds a *whole* number of records (partial records straddle boundaries), which is why we floor. **Feed-forward**: §7.5, §5.4.

7.8 Step 4: The struct's alignment is $\max(8, 1, 4) = \mathbf{8}$. After `b` the running size is 16, which is a multiple of 8, so **no tail padding** is needed. Final `sizeof(struct T) = 16`; real data = $8 + 1 + 4 = 13$, so **3 padding bytes** (the ones at offsets 9–11). Confirm each member's offset with `offsetof(struct T, member)` (and `sizeof` for the total). *Method note*: always finish by rounding the size up to the struct's alignment — here it happened to already be aligned. **Feed-forward**: §7.5.

7.9 (a) **Array-to-pointer decay** (Ch 7) — the parameter is a pointer, so `sizeof` is 8; length was lost at the call. (b) **TLB miss / page-walk** (Ch 4) — huge pages helping is the signature; they raise TLB reach. (c) **Struct padding/alignment** (Ch 7/5) — the compiler inserts padding to align members. (d) **Endianness** (Ch 7) — the other machine's byte order differs, so integers appear byte-swapped. *Method-selection cue*: "sizeof of a parameter is 8" → decay; "huge pages helped" → TLB; "bigger than the fields" → padding; "byte-swapped across machines" → endianness. **Feed-forward**: §7.3–7.6.

7.10 (a) A **virtual** address (Ch 4) — the hardware translates it to physical on access. (b) `&arr[5] - &arr[2] = 3` (an element count of type `ptrdiff_t`; pointer subtraction divides out the element size). (c) **Not stable** — a PIE build is relocated by ASLR each run, so absolute addresses change; the relative layout is stable. (d) **No** — the C standard guarantees only `sizeof(char) == 1` and minimum ranges (`int ≥ 16` bits); 4 bytes is this ABI's implementation-defined value, verified, not guaranteed. *Tempting wrong answer*: on (b), answering "12 bytes" — subtraction yields elements, not bytes. **Feed-forward**: §7.2, §7.4–7.5, and Ch 4.

7.11 (a) For `int arr[10]`, `&arr[10]` is the special "one past the last element" address, which the standard explicitly permits you to *form* and compare against (useful as a loop end), but the object there does not exist, so *dereferencing* it is undefined (C §6.5.6). (b) Making it undefined (rather than "wraps" or "reads adjacent memory") lets the optimizer assume in-bounds access — enabling optimizations and diagnostics — and reflects that beyond the array there may be padding, another object, an unmapped page, or a segment boundary (Ch 4), so no single "reasonable" behavior exists to define. (c) This is exactly the **buffer overflow** bug of Ch 9: an out-of-bounds access is UB, and an out-of-bounds write is the classic memory-corruption vulnerability. *Tempting wrong answer*: "one-past-the-end is illegal to compute" — it is legal to compute and compare, only illegal to dereference. **Feed-forward**: §7.2 and §9.3.

7.12 Machine-specific by design. A correct write-up reports measured `sizeof/_Alignof` for the scalar types (on x86-64 LP64: char 1, short 2, int 4, long 8, double 8, pointer 8, alignment = size), a struct whose hand-computed offsets and total `sizeof` match `offsetof/sizeof` exactly, and an endianness check (x86-64 prints the low byte first — little-endian). The essential discipline: label the machine/ABI, report only measured numbers, and note that a different platform (16-bit, or big-endian, or a different ABI) would legitimately differ — which is the whole reason to measure rather than assume. *Tempting wrong answer*: hard-coding "int is 4 bytes" as universal — it is implementation-defined. **Feed-forward**: §7.5–7.6.

Chapter 8 — Solutions

8.1 (a) A **stack frame** (activation record) is the block of storage for one function call — its locals, arguments, saved return address, and bookkeeping. (b) Subtracting the frame's size from the **stack-pointer register** (moving it down) allocates it — one arithmetic operation. (c) A local's lifetime ends when execution **leaves its block/function** and the frame is popped.

8.2 Heap reason (any one): the object must **outlive the function that creates it**, its size is only known at runtime, or it is too large for the stack. Cost the heap imposes: the allocator does real work (searching/splitting free lists), and — the big one — **you must free it exactly once**, an obligation the stack does not have (frames free automatically on return). **Feed-forward:** §8.2, §8.5.

8.3 (1) **Runaway/unbounded recursion** (missing base case or too deep for the input) — fix by bounding the recursion or converting to iteration with an explicit (heap-backed) stack. (2) **Very large automatic arrays/locals** — fix by allocating large or runtime-sized buffers on the **heap** instead. *Tempting wrong answer:* "increase the stack limit" — a band-aid that does not fix unbounded recursion. **Feed-forward:** §8.3.

8.4 `int *f(void){ int x=7; return &x; }` returns the address of the local `x`, which lives in `f`'s stack frame. On `return`, that frame is **popped** — the stack pointer moves back — so `x`'s storage is no longer reserved for it and can be reused by the next call. The returned pointer therefore **dangles**, and reading/writing through it is **undefined behavior** (C §6.2.4: using a pointer to an automatic object past its lifetime). gcc emits warning: `function returns address of local variable [-Wreturn-local-addr]`; clang emits `-Wreturn-stack-address`. *Tempting wrong answer:* "it returns 7 because the value is still in memory" — not guaranteed; UB imposes no such promise. **Feed-forward:** §8.4.

8.5 (a) **Caller owns the storage:** `void f(int *out){ *out = 7; }`, called `int x; f(&x);` — `x` lives in the caller's frame and outlives `f`. (b) **Heap:** `int *f(void){ int *p = malloc(sizeof(int)); if(!p) return NULL; *p = 7; return p; }` — the block outlives the call; **the caller must free it** exactly once, after its last use (ownership transfers to the caller). *Method note:* the choice is "who owns the memory?" — make it explicit. **Feed-forward:** §8.4–8.5.

8.6 "It worked" is not evidence of correctness because the bug is **undefined behavior**: the C standard imposes no requirements, so the program is free to *appear* to work — the dead frame's bytes may simply still hold the old value until something reuses them. Correctness is a property of the code against the standard, not of one run's observed output. The danger of a bug that appears to work is that it is **latent**: it hides through testing and surfaces later when the optimization level, compiler, or call pattern changes — often in production, and possibly as data corruption rather than an obvious crash. *Tempting wrong answer:* "if it passes tests it's fine" — UB can pass tests and still be a defect. **Feed-forward:** §8.4.

8.7 (a) A stack frame's locals are reserved by **subtracting the frame size from the stack pointer** — a single register adjustment, with no per-variable work and nothing to track. (b) One `malloc` may **search a free list** for a fitting block, **split** it, update bookkeeping, and, if the pool is exhausted, **ask the OS for more memory** — orders of magnitude more work. (c) Rule of thumb: for a small scratch buffer needed each iteration, use a **stack local (or one buffer reused across iterations)**, not a per-iteration `malloc/free` — the stack version is nearly free and avoids allocator traffic and fragmentation in the hot loop. *Tempting wrong answer*: "malloc a fresh buffer each iteration for cleanliness" — needless cost in a hot loop. **Feed-forward**: §8.1–8.2.

8.8 Step 4: The nodes must be allocated on the **heap** (they outlive the helper and their count is dynamic). The **caller becomes responsible for freeing** every node — exactly once each, after it is done traversing the list (walking and freeing node by node). If that responsibility is forgotten, the result is a **memory leak** (Ch 9 preview): the nodes are never reclaimed. The design principle is **ownership** — decide and document who owns the memory and therefore who frees it. *Method note*: "outlives its creator" + "size known only at runtime" is the standard signal for heap allocation with transferred ownership. **Feed-forward**: §8.5 and §9.1, §9.3.

8.9 (a) **Stack overflow** (Ch 8) — deep recursion only on large inputs exhausts the bounded stack. (b) **Dangling pointer to a local** (Ch 8) — correct in tests, garbage after unrelated code changes the call pattern that reuses the frame; classic UB signature. (c) **TLB miss / page-walk** (Ch 4) — huge pages helping points at translation reach. (d) **Struct padding/alignment** (Ch 7) — the compiler pads to align members. *Method-selection cue*: "deep recursion, large input" → overflow; "works then garbage after unrelated change" → dangling/UB; "huge pages helped" → TLB; "bigger than its fields" → padding. **Feed-forward**: §8.3–8.4.

8.10 (a) **Well-defined** — using `&local` within the same function is fine; the object is alive. (b) **Undefined behavior** — the local's frame is popped on return, so the pointer dangles (§8.4). (c) **Performance/robustness issue tending to undefined behavior** — `int big[8000000]` is ~32 MB, likely overflowing a few-MB stack; if it overflows, the access faults (typically a crash). (d) **Undefined behavior** — storing `&local` in a global and reading it after return is the same dangling-pointer bug as (b), just via a longer-lived stash. *Tempting wrong answer*: calling (d) safe "because a global is permanent" — the global outlives the object, which is exactly the problem. **Feed-forward**: §8.3–8.4.

8.11 (a) Both violate the property that **a pointer must not be used outside its target's lifetime**. (b) The stack version is easier to catch because the direct pattern — returning `&local` — is statically visible, and compilers warn (`-Wreturn-local-addr` / `-Wreturn-stack-address`); the heap version (use-after-free) generally is not statically detectable, since the `free` and the later use can be far apart and flow-dependent, so it usually needs a runtime tool (Ch 9's sanitizers). (c) Both stem from "the storage is reclaimed whether or not a pointer still refers to it" — the frame pop (stack) and `free` (heap) both end the object's lifetime without invalidating the pointers that name it, leaving a dangling handle. *Tempting wrong answer*: "they're unrelated bugs" — they are the same lifetime violation in different regions. **Feed-forward**: §8.4 and §9.3.

8.12 Machine-specific by design. A correct write-up records the exact warning (on gcc, `-Wreturn-local-addr`; on clang, `-Wreturn-stack-address`); reports what the run actually did on that machine (on the book's machine the plain build **segfaulted**, but the write-up must note this is *not* guaranteed by the standard — a different build could print a value or garbage); and shows the caller-owned fix (`void f(int *out){ *out = 7; }`) compiling warning-free and behaving correctly. Essential discipline: never rely on the observed (b) outcome — UB is not obligated to crash — and reason about lifetime instead. *Tempting wrong answer*: "it segfaults, so it's at least safe to detect" — the crash is not reliable. **Feed-forward**: §8.4, and Ch 9's tools.

Chapter 9 — Solutions

9.1 (a) **No** — `malloc` returns uninitialized (indeterminate) bytes. (b) **Yes** — `calloc` returns all-zero bytes. (c) **Yes** — `free(NULL)` is an explicit no-op, always safe.

9.2 Safe idiom: `void *tmp = realloc(p, n); if (tmp) p = tmp; else { /* handle failure; p still valid */ }`. With `p = realloc(p, n);`, if `realloc` fails it returns `NULL` and leaves the original block *valid but unreferenced* — you have overwritten your only pointer to it, so it is **leaked** (and you have lost the data). *Tempting wrong answer*: "assign directly, it's cleaner" — cleaner and leaky. **Feed-forward**: §9.1.

9.3 **Use-after-free**: accessing a block after `free` — *UB*. **Double-free**: freeing the same block twice — *UB*. **Memory leak**: losing the last pointer without freeing — *not UB* (a defined defect). **Buffer overflow**: accessing outside an allocation's bounds — *UB*. **Uninitialized read**: using memory before writing it — *indeterminate value; unreliable, UB in general*. Four of five are undefined behavior; the leak is the exception. **Feed-forward**: §9.3.

9.4 "Returns the old value, so it's harmless" is wrong because use-after-free is **undefined behavior** (C §3.4.3): the standard imposes *no requirements*, so there is no guaranteed value and no guarantee it is harmless. Two realistic other outcomes: (1) the freed slot has been handed to another allocation, so you read (or clobber) **unrelated live data** — a correctness and security hole; (2) the access **corrupts allocator metadata or crashes**, possibly far from the buggy line. It can also differ by build or run, and be exploitable. *Tempting wrong answer*: the premise itself — treating UB as having a defined outcome. **Feed-forward**: §9.3.

9.5 (a) **Use-after-free** — reading `p[0]` after `free(p)`; *UB*. (b) **Buffer overflow** — `p[4]` on a 4-element (indices 0–3) block; *UB*. (c) **Double-free** — freeing the same block twice; *UB*. (d) **Uninitialized read** — reading a `malloc'd` `int` before writing it; indeterminate value, unreliable (*UB in general*). (e) **Memory leak** — overwriting the only pointer to a 1024-byte block without freeing; a defect, but *not UB*. *Method-selection cue*: "touched after free" → UAF; "index past the end" → overflow; "freed twice" → double-free; "read before write" → uninitialized; "pointer lost, not freed" → leak. **Feed-forward**: §9.3.

9.6 A **leak** frees *too late* — never; the block stays allocated but unreachable, so memory is wasted but no invalid access occurs, hence **not undefined behavior** (the program stays well-defined, just wasteful). A **use-after-free** frees *too early* — while a pointer still needs the block; the later access touches storage whose lifetime has ended, which **is undefined behavior**. The asymmetry: accessing outside a lifetime is undefined; merely retaining allocated memory forever is defined (if undesirable). *Tempting wrong answer*: "both are UB because both are memory bugs" — a leak is a defect but not UB. **Feed-forward**: §9.3.

9.7 (a) A heap allocation may search a free list, split a block, update bookkeeping, and sometimes call the OS for more memory — real work — whereas a stack allocation is a single stack-pointer adjustment (Ch 8). (b) **External fragmentation** is free space chopped into many small, non-adjacent pieces by interleaved allocate/free of varying sizes; a large request can fail because *no single free block* is big enough, even though the *total* free memory exceeds the request. (c) **Not undefined behavior** — fragmentation is a *performance/capacity* problem, not a language-rule violation. *Tempting wrong answer*: "fragmentation corrupts memory" — it does not; it just wastes/scatters it. **Feed-forward**: §9.2.

9.8 Step 4: Correct allocation size: `malloc(strlen(s) + 1)` — room for the string *and* its terminating `'\0'`. Missing check: after `malloc`, verify it did not return `NULL` before writing (`if (!r) return NULL;`). Corrected function: `char *dup(const char *s){ char *r = malloc(strlen(s)+1); if(!r) return NULL; strcpy(r,s); return r; }`. The tool from §9.4 that catches the original one-byte overflow is **AddressSanitizer** (`-fsanitize=address`), which would report a heap-buffer-overflow on the `strcpy` line (Valgrind would too). *Tempting wrong answer*: "strlen already counts the null terminator" — it does not; it counts the characters before it. **Feed-forward**: §9.3–9.4.

9.9 (a) **Leak** — steadily growing memory until the OOM killer acts is the leak signature. (b) **Uninitialized read** — a different value each run, only in release builds, is the classic "garbage bytes" symptom (build-dependent). (c) **Buffer overflow** — an out-of-bounds write corrupting an adjacent variable. (d) **Use-after-free / double-free** — crashing inside `free` with heap corruption after releasing the same pointer twice is a double-free. (e) **Non-memory performance issue** (a TLB/huge-page effect, Ch 4) — *not* a memory bug; the trap in this item. *Method-selection cue*: "unbounded growth" → leak; "differs per run/release" → uninitialized; "adjacent corruption" → overflow; "crash in free after two frees" → double-free; "huge pages help" → TLB, not a bug. **Feed-forward**: §9.3, Ch 4.

9.10 (a) **Well-defined and fine** — `free(NULL)` is a guaranteed no-op. (b) **A defined-but-real defect** — a memory leak; wasteful but not UB. (c) **Undefined behavior** — `a[n]` on an `n`-element array is one past the last valid index (valid indices `0..n-1`); dereferencing it is out of bounds. (d) **Well-defined and fine** — a `calloc`'d block is zeroed, so reading it yields 0, a defined value. (e) **Reads an indeterminate value (at best unspecified, in general undefined behavior)** — a `malloc`'d block is uninitialized, so reading before writing uses an *indeterminate value*: on this ABI an `int` has no trap representation, so you get an unspecified (garbage) value rather than a crash, but the standard permits it to be full UB in general (matching §9.3), so never rely on it. *Tempting wrong answer*: conflating (d) and (e) — `calloc` zeroes, `malloc` does not. **Feed-forward**: §9.1, §9.3.

9.11 (a) The three obligations: use a pointer only within its target's **lifetime**; access only within the allocation's **bounds**; and **free each block exactly once** (not zero, not twice). (b) **Use-after-free** violates *lifetime*; **double-free** violates *free-exactly-once*; a **leak** also violates *free-exactly-once* (frees zero times) — while use-after-free additionally implies the free happened too early. (c) A compiler that checks these before the program runs would **reject**, at compile time, any program that could use a pointer past its target's lifetime, index out of bounds, or mismanage a free — turning a whole class of runtime UB into type errors. That is more reliable than a runtime sanitizer because a sanitizer only catches a bug on *inputs/paths you actually execute* during testing, whereas a static guarantee holds for *all* executions — which is exactly Rust's ownership/borrowing promise. *Tempting wrong answer*: "a sanitizer is just as good" — it is path-dependent, not a proof. **Feed-forward**: §9.5 and Part 5 (Rust).

9.12 Machine-specific by design. A correct write-up reports, for each of the four bugs, ASan's exact error type and the source line — on the book's machine (gcc 11, x86-64): heap-use-after-free (a READ on the line reading the freed block), heap-buffer-overflow (a WRITE "0 bytes to the right of" the allocated region), attempting double-free, and (via LeakSanitizer) Direct leak of N byte(s) pinned to the allocation line. The fifth program (uninitialized read) should show that ASan does **not** flag it — reach for Valgrind or MemorySanitizer, or catch the simple case with gcc -Wall (-Wuninitialized). Essential discipline: label compiler/tool versions, report only observed output (addresses and wording vary), and note that ASan detects addressability bugs, not uninitialized-value bugs. *Tempting wrong answer*: "ASan catches everything" — it misses uninitialized reads. **Feed-forward**: §9.4.

Chapter 10 — Solutions

10.1 Undefined behavior: the standard imposes *no requirements* (C §3.4.3) — the program has no meaning on the triggering input. **Unspecified behavior:** the standard offers two or more valid possibilities and does not say which you get (e.g. argument evaluation order). **Implementation-defined:** unspecified but the implementation must *document* its choice (e.g. `sizeof(int)`). Only **undefined** behavior grants the optimizer license to assume the situation never happens. **Feed-forward:** §10.1.

10.2 Unsigned integer overflow is *not* UB — it is defined to wrap modulo 2^n . **Signed** integer overflow *is* undefined behavior. The defined case (unsigned) yields the mathematically correct value reduced modulo 2^n (i.e. the low n bits). *Tempting wrong answer:* "both wrap, it's two's complement" — the hardware may wrap, but the *language* leaves signed overflow undefined, which is what the optimizer exploits. **Feed-forward:** §10.3.

10.3 Most likely explanation: the program contains **undefined behavior** that was harmless-looking at -O0 but that the optimizer exploited at -O2 (it assumed the UB never happens and transformed the code). The two flags to hunt it: `-fsanitize=undefined` (UBSan) and `-fsanitize=address` (ASan), ideally with `-Wall -Wextra`. *Tempting wrong answer:* "the optimizer has a bug" — far more often the UB was already there. **Feed-forward:** §10.2, §10.4.

10.4 Signed overflow is UB, so the compiler may assume `x + 1` never overflows; if it never overflows, `x + 1 > x` is always true, so the function folds to `return 1` (on this machine, `movl $1, %eax; ret`). What breaks: a programmer using `if (x + 1 < x)` to detect overflow relies on the overflow *happening* and wrapping — but that is UB, so the compiler assumes it cannot happen and deletes the check. Correct detection: compile with `-fwrapv` (defines signed overflow as wraparound), or do the arithmetic in an unsigned type, or use `__builtin_add_overflow`, which reports overflow without invoking UB. *Tempting wrong answer:* "test after the add" — testing *after* signed overflow is testing after UB. **Feed-forward:** §10.3-10.4.

10.5 (a) The UB is a **null-pointer dereference**: `t->sk` dereferences `t` before it is checked. (b) Because dereferencing null is UB, the compiler may assume `t` was non-null at the dereference, so `if (!t)` is provably false and is deleted as dead code. (c) Fix: move the check before the dereference — `if (!t) return -1; int *sk = t->sk; return *sk;`. (d) The flag that preserves the check without moving code: `-fno-delete-null-pointer-checks`. The pattern is the real Linux kernel bug **CVE-2009-1897** (`tun_chr_poll`). *Tempting wrong answer:* "the check is there, so it's fine" — it is there in the source, not in the binary. **Feed-forward:** §10.3.

10.6 The compiler is behaving correctly under the **as-if rule**: it must preserve the observable behavior of a *well-defined* program, but a program that dereferences a possibly-null pointer has undefined behavior on the null input, so there is no observable behavior it must preserve there. Given the dereference `t->sk`, the compiler infers `t != NULL` and legally removes the redundant `if (!t)`. The code, not the compiler, is wrong: the check comes after the dereference it is meant to guard. *Tempting wrong answer*: "the source clearly has the check, so the compiler dropped correct code" — to the optimizer, a check after a dereference is a statement the pointer was already non-null. **Feed-forward**: §10.2.

10.7 (a) Mandating checks would insert a bounds check on every array access, an overflow check on every signed add, and a null check on every dereference — runtime cost on the hot path of essentially every program. (b) One optimization signed-overflow UB enables: treating a loop counter as unable to wrap, so the compiler can prove loop termination, promote the counter to a wider type, or strength-reduce — none guaranteed if overflow were defined to wrap. (c) The bargain: the programmer gives up "defined behavior in the bad case" (and takes on the duty never to trigger UB); the compiler gains the freedom to generate fast code across all hardware without pervasive safety checks. *Tempting wrong answer*: "UB is just a mistake in the standard" — it is a deliberate performance/portability choice. **Feed-forward**: §10.4.

10.8 Step 4: Fix: the bounds check must come *before* the access — check `if (i < 0 || i >= n) return;` (or handle the error) *then* touch `a[i]`, including for logging. A sanitizer that would have flagged the original out-of-bounds access at runtime: **AddressSanitizer** (`-fsanitize=address`) for a heap/stack array, or UBSan's bounds check for a known-size array. Moving the log line is not "just reordering output": once `a[i]` executes, the compiler may *prove* `0 <= i < n` (else UB) and delete the later check — so the position of the access changes what the compiler is allowed to assume about `i`, not merely when a byte is printed. *Tempting wrong answer*: "it's just a logging line, order doesn't matter" — the access is what licenses the deletion. **Feed-forward**: §10.2–10.3.

10.9 (a) **Implementation-defined** — `sizeof(int)` varies but must be documented. (b) **Unspecified** — the order of evaluating `a()` and `b()` is one of two valid, undocumented choices. (c) **Undefined** — `INT_MAX + 1` is signed overflow. (d) **Implementation-defined** — right-shift of a negative signed integer is implementation-defined (the standard lets the implementation choose arithmetic vs. logical shift and document it). (e) **Undefined behavior** (in general) — reading a `malloc`'d block before writing it uses an indeterminate value (the §9.3 memory bug). *Method-selection cue*: "varies but documented" → implementation-defined; "two valid orders" → unspecified; "no meaning at all" → undefined. **Feed-forward**: §10.1, §9.3.

10.10 Least to most effective: (a) `-O0` — *hides* the bug (the UB remains; a future build or a different path re-exposes it); (b) `volatile` — also *hides* it (it defeats some optimizer assumptions but does not remove the UB, and is fragile and slow); (c) `-fsanitize=undefined,address` plus fixing what it reports — *removes* the bug by finding and eliminating the actual UB. So the order is (a) < (b) < (c), and only (c) removes rather than hides. *Tempting wrong answer*: "compile at `-O0` in production to be safe" — you keep the bug and lose performance. **Feed-forward**: §10.4.

10.11 (a) The standard says a program that executes UB has undefined behavior for that *execution* — not just from the buggy line onward — so the compiler need not preserve observable effects that occur before the UB either; it may hoist, sink, or delete surrounding work ("UB can time-travel," in Regehr's phrasing). (b) Concrete scenario: an upstream `t->sk` lets the compiler delete a downstream `if (!t) return ERR;;`; a caller that would have safely returned an error on a null `t` now falls through into code that trusts `t`, and an attacker who maps memory at address 0 reaches it — a would-be crash becomes a silent privilege escalation (the CVE-2009-1897 shape). (c) A static guarantee is stronger because a runtime sanitizer only catches UB on the *inputs and paths you actually execute* during testing, whereas "the safe subset has no UB" holds for *all* executions — a proof, not a spot check. *Tempting wrong answer*: "sanitizers make UB safe" — they detect it on exercised paths only. **Feed-forward**: §10.2, §10.5, Part 5.

10.12 Machine-specific by design. A correct write-up reports: at `-O2`, `x + 1 > x` folds to `return 1` (on the book's machine, `movl $1, %eax; ret`), and adding `-fwrapv` makes the addition and comparison reappear; for `poll`, the default `-O2` omits the `testq %rdi,%rdi` null check while `-fno-delete-null-pointer-checks` restores it. The signed-overflow program should be flagged by `-fsanitize=undefined` as a runtime error (signed integer overflow). Essential discipline: report only observed assembly/output and label the compiler/version (gcc 11.4.0, x86-64 here), since exact instructions and wording vary. *Tempting wrong answer*: "the assembly will be the same as mine" — it depends on compiler and version. **Feed-forward**: §10.3–10.4.

Chapter 11 — Solutions

11.1 (a) A single **null byte** `'\0'` (value 0) marks the end. (b) `strlen` is $O(n)$ because the length is *not* stored — it must scan byte by byte until it finds the terminator. (c) The function learns **nothing** about capacity: the array decays to a bare `char *`, an address with no length attached. **Feed-forward:** §11.1.

11.2 `gets` — it read a line with no bound at all and was **removed from the language in C11** (deprecated in C99). Use `fgets(s, size, stdin)`, which takes the buffer size. *Tempting wrong answer:* "`scanf("%s")` is fine" — unbounded `%s` has the same overflow problem; use a width, e.g. `%3s`. **Feed-forward:** §11.3.

11.3 (a) **False** — `strncpy` does not terminate when the source is at least `n` bytes long; it leaves the destination unterminated. (b) **False** — a canary *detects* a linear stack overflow and aborts; the bug (the out-of-bounds write) is still there. (c) **True** — `snprintf` always null-terminates when its size argument is nonzero. **Feed-forward:** §11.2–11.4.

11.4 On a downward-growing stack, a local array sits *below* the saved frame pointer and saved **return address**, while a copy into the array writes toward *higher* addresses — so an overrun climbs over the frame pointer onto the return address. Overwriting it means that when the function executes `ret`, the CPU jumps to the attacker-supplied value: control-flow hijack. An out-of-bounds *read* (usually) cannot redirect execution — it leaks adjacent bytes but does not write control data. *Tempting wrong answer:* "reads are as dangerous as writes" — reads leak (serious, e.g. Heartbleed) but do not seize control. **Feed-forward:** §11.2.

11.5 (a) **Overflows** whenever `user_input` is longer than `buf` — no destination bound; rewrite `snprintf(buf, sizeof buf, "%s", user_input);`. (b) **Overflows** if the formatted result exceeds `buf` — `sprintf` is unbounded; rewrite `snprintf(buf, sizeof buf, "%s-%s", a, b);`. (c) **Over-reads** — `strncpy(d, "abcd", 4)` fills all 4 bytes with no terminator, so `printf("%s", d)` runs off the end; fix by terminating (`d[3] = '\0';`, losing a char) or, better, `snprintf(d, sizeof d, "%s", "abcd")`. *Method-selection cue:* "no size argument" → can overflow; "size argument but no guaranteed terminator" → can over-read. **Feed-forward:** §11.2–11.3.

11.6 Three reasons a canary + ASLR do not "handle" overflows: (1) a **canary only catches a linear stack overflow and only at return** — it does nothing for heap overflows or overwrites of non-control data, and it acts after the corruption (turning a hijack into a crash/DoS); (2) **ASLR only randomizes addresses** — it raises the difficulty of a working exploit but is defeated by info-leaks or techniques that do not need fixed addresses; (3) both are **mitigations, not fixes** — the out-of-bounds write is still undefined behavior, still present. The actual fix: remove the unbounded writes (use `snprintf`/bounded copies) so the overflow cannot occur. *Tempting wrong answer:* "mitigations on = bug fixed" — they contain blast radius, they do not remove the bug. **Feed-forward:** §11.4.

11.7 (a) `strlen` is $O(n)$ because a C string stores no length — it scans to the terminator — whereas a length-carrying string stores the count, so reading the length is $O(1)$. (b) Stack canaries add a random-value write on entry and a compare on return per protected frame; `_FORTIFY_SOURCE` adds size-checked variants of the string/memory calls — small overhead in exchange for turning many overflows into detected aborts or compile errors. (c) Not free: the canary compare is on the return path of every protected function, a real (if usually small) cost. Trade-off: **safety vs. speed** — you pay a little runtime to detect a class of memory corruption. *Tempting wrong answer*: "canaries are free" — they add work on every protected return. **Feed-forward**: §11.1, §11.4.

11.8 Step 4: Rewrite with a single bounded format: `void greet(const char *name){ char msg[16]; snprintf(msg, sizeof msg, "Hello, %s", name); puts(msg); }`. `snprintf` writes at most `sizeof msg` bytes *including* the terminator and always null-terminates (when the size is nonzero), so no `name`, however long, can overflow `msg` — it truncates instead; `strcat` guarantees neither a bound nor (relative to the destination capacity) safety. A §11.4 defense that turns the original overflow into a crash rather than a silent hijack: a **stack canary** (`-fstack-protector`), which aborts with "stack smashing detected" instead of returning through a corrupted return address (`_FORTIFY_SOURCE` would also catch it, at compile or run time). *Tempting wrong answer*: "use `strncat` with `sizeof msg`" — `strncat`'s bound is the bytes to *append*, not the buffer size, and it is easy to miscompute. **Feed-forward**: §11.3–11.4.

11.9 (a) **Buffer overflow (over-write)** is possible, but "garbage appended *after* the expected text" most cleanly matches an **unterminated string (over-read)** — the print continues past the intended end until it hits a stray zero. (b) **A mitigation firing** — `*** stack smashing detected ***` is the stack canary aborting on a detected overflow. (c) **Use-after-free** — reading a freed block that was reallocated to another request (Ch 9). (d) **Unterminated string (over-read)** — `printf("%s")` reading far past the data and leaking adjacent memory (Heartbleed class). *Method-selection cue*: "garbage after expected text / reads too far" → missing terminator; "stack smashing detected" → canary; "freed block reused" → UAF. **Feed-forward**: §11.1–11.4, §9.3.

11.10 (a) **Undefined behavior** — `strncpy` with `strlen(s) >= sizeof d` leaves `d` unterminated, so reading it as a string over-reads out of bounds. (b) **Well-defined and fine** — `snprintf` always terminates, so reading `d` is safe (possibly truncated, but valid). (c) **Undefined behavior** — a one-byte heap overflow is an out-of-bounds write. (d) **Well-defined** (a defined defensive action) — a canary deliberately aborting the process on a detected overflow is defined behavior, not UB (it is the mitigation working). *Tempting wrong answer*: conflating (a) and (b) — the difference is the guaranteed terminator. **Feed-forward**: §11.2–11.4.

11.11 (a) A **stack canary** defeats a linear overflow that reaches the return address (detected at return); **NX/DEP** marks the stack/heap non-executable, defeating "inject shellcode and jump to it"; **ASLR** randomizes segment base addresses, defeating attacks that need to predict them. **Return-oriented programming (ROP)** was developed to route around NX by chaining fragments of existing executable code ("gadgets") instead of injecting new code. (b) Layering is worthwhile because each mitigation raises the cost and reliability barrier for a different technique; an attacker must defeat all of them, and many real exploits fail against the stack. (c) Rust makes the bug "not happen" by carrying the length with every slice/string and bounds-checking indexing: an out-of-bounds access is a defined **panic**, not undefined memory corruption — categorically different because a panic has a defined, contained outcome (the program stops safely) whereas UB has no defined outcome and can corrupt control flow. *Tempting wrong answer*: "a panic is just a crash, same as UB" — a panic is defined and cannot be leveraged into memory corruption. **Feed-forward**: §11.4–11.5, Part 5.

11.12 Machine-specific by design. A correct write-up reports: (a) with `-fstack-protector-all` the 8-byte overflow aborts with `*** stack smashing detected ***` (on the book's machine, exit status 134, SIGABRT); (b) with `-D_FORTIFY_SOURCE=2 -O2` a compile-time diagnostic like

`__builtin___strcpy_chk / __memcpy_chk` writing N bytes into a region of size M (and gcc's `-Wstringop-overflow` may also fire); (c) `strncpy(d, "abcd", 4)` leaves the four bytes 97 98 99 100 with *no* terminating zero. Essential discipline: report only observed output and label compiler/version (gcc 11.4.0, x86-64 here), since exact messages and status vary.

Tempting wrong answer: "the canary means the overflow was harmless" — it converted a hijack into a crash; the bug is the same. **Feed-forward**: §11.2–11.4.

Chapter 12 — Solutions

12.1 (a) An arena exploits **shared lifetime** — many objects that die together. (b) A pool exploits **shared size** — many objects of one fixed size. (c) The **pool** can free an individual object (push its slot back on the free list); the arena cannot. **Feed-forward:** §12.1–12.3.

12.2 `arena_alloc` is **O(1)** (round up the offset, bump it, return the address); `arena_reset` is **O(1)** (set the offset to 0). The one thing an arena cannot do that `malloc` can: **free an individual object** — it frees only the whole region at once. *Tempting wrong answer:* "reset is O(n) because it frees n objects" — it frees them all with a single offset assignment, touching none of them. **Feed-forward:** §12.2.

12.3 The "next free slot" pointer is stored **inside the free slot itself** (its first bytes). It costs no extra memory because a free slot holds no live object — its space is unused — so the free list is threaded through the slots rather than kept in a separate structure. *Tempting wrong answer:* "you need a separate array of next-pointers" — unnecessary; the free slots' own bytes hold them. **Feed-forward:** §12.3.

12.4 `typedef struct { unsigned char *base; size_t cap, off; } Arena;` with `void *arena_alloc(Arena *a, size_t n){ size_t p = (a->off + 15) & ~(size_t)15; if (p + n > a->cap) return NULL; a->off = p + n; return a->base + p; }` and `void arena_reset(Arena *a){ a->off = 0; }`. `arena_alloc` is O(1) because it does a fixed amount of work — align, bounds-check, bump — with no search or coalescing. `arena_reset` frees everything without touching individual objects because "free" here just means "the offset no longer points past them," so resetting the offset makes all of that space available again in one step. *Tempting wrong answer:* forgetting the alignment round-up — unaligned returns are UB for some types (Ch 7). **Feed-forward:** §12.2.

12.5 `void *pool_alloc(Pool *p){ if (!p->free_list) return NULL; void *slot = p->free_list; p->free_list = *(void **)slot; return slot; }` and `void pool_free(Pool *p, void *slot){ *(void **)slot = p->free_list; p->free_list = slot; }`. Both are O(1): allocation pops the free-list head and freeing pushes onto it — a couple of pointer writes, no search. There is no external fragmentation within the size class because every slot is identical and interchangeable, so any free slot satisfies any request — free space is never "too scattered to fit." *Tempting wrong answer:* "you still need to coalesce" — with one fixed size there is nothing to coalesce. **Feed-forward:** §12.3.

12.6 Reason 1 (inflexibility): an arena cannot free individual objects, so any object that outlives its batch pins the *whole* region until reset — used "everywhere," including for long-lived mixed-lifetime state, it can inflate memory badly. Reason 2 (evidence): Berger, Zorn & McKinley (OOPSLA 2002) found a good general-purpose allocator matched or beat hand-written custom allocators for most programs; the arena/region case was the exception, not the rule. Correct procedure: **profile first** to confirm allocation is the bottleneck, then apply an arena only where the shared-lifetime pattern fits (e.g. per-request scratch), keeping `malloc` elsewhere. *Tempting wrong answer*: "arenas are always faster, so use them everywhere" — faster only for the right pattern, and costly in memory otherwise. **Feed-forward**: §12.1, §12.4.

12.7 (a) The arena skips, per allocation: the **free-list search** (it just bumps an offset) and the **per-block header / coalescing bookkeeping** (it keeps none) — both of which the general allocator does (Ch 9). (b) The $\sim 4\text{--}5\times$ measurement used "N objects live at once, then freed": the arena frees all N with one reset while `malloc` pays N individual free calls, so the comparison rewards the arena's bulk free. A pattern that narrows or erases the gap: freeing each object immediately after allocating it, where glibc's `malloc` reuses one hot block and the per-object cost nearly vanishes (as the chapter's first, discarded measurement showed). (c) Memory risk: because it never frees individual objects, an arena holding a long-lived object pins the entire region — it bites when object lifetimes are mixed rather than shared. *Tempting wrong answer*: "the arena is always $\sim 5\times$ faster" — only for its pattern. **Feed-forward**: §12.2, §12.4.

12.8 Step 4: The bug is a **use-after-free** (a use-after-reset): the cached pointer dangles into arena storage whose lifetime `arena_reset` ended, and the next request reuses those bytes under it. It **is undefined behavior** (Ch 10). ASan/Valgrind will likely *not* catch it by default because they instrument `malloc/free`, not *your* arena's reset — to them the arena's backing block is still one big live allocation, so a read into it looks legal (you would have to add ASan's manual poisoning hooks for the tool to see the reset). Fix: do not let the cache point into the arena — **copy** the string into memory with a lifetime that matches the cache (e.g. `strdup/malloc` owned by the cache), so nothing the cache holds lives in per-request arena storage. *Tempting wrong answer*: "the sanitizer would have caught it" — not for a custom allocator's reset, by default. **Feed-forward**: §12.4–12.5, §9.3–9.4.

12.9 (a) **Arena** — many small objects sharing one lifetime (the request), freed together at response time. (b) **Pool** — constant create/destroy of same-size nodes throughout the run; $O(1)$ both ways, no fragmentation, individual frees. (c) **General malloc** — a library cannot predict its callers' patterns, so the general allocator is the safe default. (d) **Arena** — a burst of AST nodes freed together when the pass ends (the classic region/compiler-pass case Berger et al. highlighted). *Method-selection cue*: "share a lifetime, die together" → arena; "same size, repeated alloc/free" → pool; "unknown pattern" → general. **Feed-forward**: §12.1, §12.4.

12.10 (a) **Undefined behavior** — reading through a pointer into an arena after `arena_reset` is a use-after-reset (a use-after-free). (b) **A defined-but-real defect** — a long-lived object pinning the whole region wastes memory but is not a language-rule violation (analogous to a leak: defined, undesirable). (c) **Undefined behavior** — `pool_free`-ing the same slot twice corrupts the free list (a double-free in your allocator). (d) **Well-defined and fine** — resetting an arena with no live pointers into it is exactly how it is meant to be used. *Tempting wrong answer:* calling (b) UB — over-retaining memory is a defect, not undefined behavior. **Feed-forward:** §12.2-12.5, §9.3.

12.11 (a) In C the programmer must guarantee, by hand, that no pointer is used outside its target's lifetime, no access falls outside its bounds, and every block is freed exactly once — with no undefined behavior triggered. (b) A custom allocator *relocates* the obligation: it does not free individual objects for you (arena) or it recycles slots you must not double-free (pool), so *you* — the allocator's author and its caller — now enforce "free exactly once" and "no use after the lifetime ends" for arena/pool memory, and the system allocator's tooling no longer helps. (c) A language that ties an allocation's lifetime to a scope (Rust's ownership/borrowing) can make "use after reset" a **compile-time error** — the borrow checker refuses to let a reference outlive the region it points into — which is stronger than a runtime sanitizer because it holds for all executions and needs no test input to trigger, and because the sanitizer does not even understand your custom allocator's reset. *Tempting wrong answer:* "custom allocators remove the obligation" — they move it onto you. **Feed-forward:** §12.5, Part 5.

12.12 Machine-specific by design. A correct write-up reports: (a) three allocations return independent, 16-byte-aligned pointers with the expected values, and `arena_reset` sets the offset back to 0; (b) for large N with all objects live then freed, the arena (one bulk free) beats `malloc` (N frees) by a several-fold ratio — on the book's machine roughly 4–5×, varying run-to-run; (c) when each object is freed immediately after allocation, the arena's advantage shrinks or disappears, because the general allocator reuses one hot block and per-object cost nearly vanishes — showing the win depends entirely on the lifetime pattern. Essential discipline: report only measured numbers and label compiler/hardware (gcc 11.4.0, x86-64 here), since figures vary. *Tempting wrong answer:* "the arena is faster in every test" — not for the free-immediately pattern. **Feed-forward:** §12.2, §12.4.

Chapter 13 — Solutions

13.1 The two modes are **user mode** (ring 3, where your code runs) and **kernel mode** (ring 0, where the OS runs). A user-mode program may not, for example, execute privileged instructions such as halting the CPU, disabling interrupts, loading the page tables (Ch 4), or accessing device registers directly. *Tempting wrong answer*: "user mode is just slower than kernel mode" — the difference is authority, not speed; the same instruction runs at the same rate, but privileged ones simply trap in user mode. **Feed-forward**: §13.1.

13.2 On x86-64 Linux the system-call number goes in **rax**, and the return value comes back in **rax** as well (a negative small value being `-errno`). *Tempting wrong answer*: "the fourth argument goes in `rcx`" — it goes in `r10`, because the `syscall` instruction clobbers `rcx` (and `r11`). **Feed-forward**: §13.2.

13.3 No — `printf` is **not** a system call; it is a C library function that parses the format string, converts values to text, and **buffers** the result in user-space memory. It eventually produces output through the **write** system call, but only when the buffer flushes (fills, hits a newline for a line-buffered terminal, or the program exits normally). *Tempting wrong answer*: "every output function is a system call" — most output is buffered in user space precisely to avoid crossing the boundary per call. **Feed-forward**: §13.3.

13.4 The C library's `write` wrapper puts the `write` system-call number in `rax`, the file descriptor 1 in `rdi`, the buffer's address in `rsi`, and the length 5 in `rdx`; it executes the `syscall` instruction, which switches to kernel mode and jumps to the fixed kernel entry point; the kernel validates the arguments, does the I/O, and returns the number of bytes written (5 on success) in `rax`. *Tempting wrong answer*: "the program jumps directly into the kernel function" — it cannot choose the target; the trap lands only at the kernel's fixed entry point. **Feed-forward**: §13.2.

13.5 The raw kernel interface signals failure by returning a **negative error code** (e.g. `-EBADF`) in `rax`. The C library wrapper detects that, stores the code in the global `errno`, and returns `-1` to the caller. So a raw `-EBADF` becomes: the wrapper returns `-1` and sets `errno == EBADF`. *Tempting wrong answer*: "errno is set by the kernel" — the kernel returns a negative number; the library *translates* it into `errno` and `-1`. **Feed-forward**: §13.3.

13.6 The thousands of `read(3, ..., 1) = 1` lines mean the program is reading the file **one byte per system call** — one boundary crossing per byte, which is the most expensive way possible. The fix is to **read in large blocks** (a big buffer, or buffered I/O such as `fread/a FILE*`), turning thousands of crossings into a handful. *Tempting wrong answer*: "nothing is wrong; it's reading the file" — it is, but at $\sim 100\times$ the necessary cost. **Feed-forward**: §13.4.

13.7 (a) A system call must **switch privilege mode** (save/restore state, enter and leave the kernel) and it **disturbs caches/TLB** and runs argument validation — none of which a plain function call does. (b) Reading 1 MB one byte at a time is $\sim 1,000,000$ boundary crossings, and each costs $\sim 100\times$ a function call; buffering (read a large block, or use a `FILE*`) cuts it to a handful of crossings. (c) `clock_gettime` is served by the **vDSO** — a page the kernel maps read-only into every process — so it runs entirely in user mode with no trap, and the "call" makes no system call at all. *Tempting wrong answer*: "reading byte-by-byte is slow because the disk is slow" — the dominant cost here is the boundary crossings, not the device (the data may already be cached). **Feed-forward**: §13.4.

13.8 Step 4: When the program died, the diagnostic bytes were still in the **C library's user-space output buffer** — they had never crossed the boundary, because `printf` buffers and only calls `write` on a flush. The abnormal termination skipped the normal exit-time flush, so the line was lost: it never reached the terminal. Two ways to guarantee it appears before a possible crash: call `fflush(stdout)` immediately after the `printf` (force the crossing), or use **unbuffered output** — write with the `write` system call directly, or set the stream unbuffered / line-buffered. *Tempting wrong answer*: "the `printf` never executed" — it did; its output was simply stranded in the buffer. **Feed-forward**: §13.3, §13.4.

13.9 (a) **Neither / user space** — an integer add is a plain instruction. (b) **Boundary crossing** — opening a file needs the kernel (`openat`). (c) **User space** — sorting an in-memory array is pure computation. (d) **Boundary crossing** — sending a packet touches the network device, so it needs the kernel. (e) **User space** — `sprintf` formats into a memory buffer; no I/O, no crossing (contrast `printf`, which will eventually `write`). *Method-selection cue*: touch a device, another process, or kernel-managed state $\rightarrow$ crossing; compute in your own memory $\rightarrow$ user space. **Feed-forward**: §13.2–13.4.

13.10 (a) **Undefined behavior** — reading one past a stack array is an out-of-bounds access (Ch 11). (b) **A boundary crossing** — read on a valid descriptor is a system call. (c) **Undefined behavior** — signed integer overflow is UB (Ch 10). (d) **Neither** — a buffered `printf` that has not flushed is well-defined; the bytes are simply still in user space and no crossing has happened yet. *Tempting wrong answer*: calling (d) a system call — the crossing happens only at the flush, not at the `printf`. **Feed-forward**: §13.3, and Ch 10–11.

13.11 (a) A process's pointers are **virtual addresses translated through its own page tables** (Ch 4); a pointer value that names another process's or the kernel's memory does not map to anything in *this* process's address space, so the kernel's validation rejects it (or it simply is not this process's memory to read). (b) Without validation, a process could pass the kernel a pointer into *kernel* memory as the "buffer" for a `read`, tricking privileged code into copying secrets out, or into *another* process's memory — a confused-deputy attack; validation is what keeps the kernel from acting on addresses the caller has no right to. (c) Per-process address translation confines each process to its own memory (it cannot even name another's), and the user/kernel mode bit prevents it from rewriting the translation tables or reading devices to get around that — together, one process can neither address nor reach another's memory. *Tempting wrong answer*: "the kernel trusts the pointer because the process passed it" — the kernel must never trust a user pointer; validating it is essential to isolation. **Feed-forward**: §13.1–13.2, Ch 4.

13.12 Machine-specific by design. A correct write-up reports: (a) `strace` shows the `write` call with its arguments and return value (e.g. `write(1, "...", n) = n`); (b) a trivial hello-world makes a couple dozen system calls, most of them process startup (dynamic linker `mmap/openat/mprotect`, `brk`, etc.) with only one or two being the program's own output; (c) the `getpid` system call costs on the order of $\sim 100\times$ a trivial function call (on the book's machine $\sim 150\text{--}190$ ns vs. $\sim 1\text{--}2$ ns), varying run to run and with mitigations; (d) one million `clock_gettime` calls produce *zero* `clock_gettime` system calls under `strace -c`, because the vDSO answers in user space. Essential discipline: report only measured numbers and label compiler/kernel/hardware, since figures vary. *Tempting wrong answer*: "the syscall and the function call cost about the same" — the whole point is that they differ by roughly two orders of magnitude. **Feed-forward**: §13.2, §13.4.

Chapter 14 — Solutions

14.1 A process bundles (1) an **address space** (its private virtual memory — text, data, bss, heap, stack), (2) one or more **threads of execution** (the running context: program counter, stack pointer, registers), and (3) a set of **kernel resources** (open file descriptors, PID/parent PID, credentials, working directory). *Tempting wrong answer:* "a process is just the program's code" — the code on disk is inert; the process is the running instance with memory, execution state, and resources. **Feed-forward:** §14.1.

14.2 `fork` returns the **child's PID in the parent** and **0 in the child** (or `-1` on failure). That difference is how each of the two otherwise-identical processes learns which one it is, so it can branch to different code. *Tempting wrong answer:* "`fork` returns 0 to whichever called it" — both continue from the return; the caller becomes the parent and gets the child's PID, while the new child gets 0. **Feed-forward:** §14.2.

14.3 No — on success `exec` does **not return** (there is no original program to return to). It **replaces the entire process image** — text, data, heap, stack — with a new program, while **keeping the PID** (and the open file descriptors, unless marked close-on-exec). *Tempting wrong answer:* "`exec` starts a new process" — it does not; it transforms the *current* process into a different program, same PID. **Feed-forward:** §14.3.

14.4

```
pid_t pid = fork(); if (pid == 0) { /* child */ } else if (pid > 0) { int st;
waitpid(pid, &st, 0); /* parent, after child exits */ } else { /* fork failed */ }. The pid == 0 branch is the child; the pid > 0 branch is the parent, which waits for the child; each knows which it is solely from fork's return value. Tempting wrong answer: forgetting the parent must wait — without it the finished child becomes a zombie (§14.4). Feed-forward: §14.2, §14.4.
```

14.5 At `fork`, the kernel does **not** copy the parent's pages; it maps the same physical pages into both processes and marks them **read-only**. A real copy is triggered only when *either* process **writes** a page: the hardware faults, the kernel makes a private copy of that one page for the writer, and both continue. So the cost is proportional to the number of pages actually modified, not to the size of the address space — forking a 1 GB process is fast because almost none of that gigabyte is written before use. *Tempting wrong answer:* "`fork` duplicates all of memory immediately" — that would make forking large processes prohibitively slow; copy-on-write avoids it. **Feed-forward:** §14.2, Ch 4.

14.6 The shell `forks` a child; the child `execs` the requested program (replacing itself with it); the parent (shell) `waits` for the child to finish before prompting again. The gap between `fork` and `exec` is where the child, still running the shell's code, adjusts its inherited resources — e.g. **redirecting standard output to a file** or wiring its input/output to a **pipe** — so that the program it then `execs` starts up already plumbed into that environment. *Tempting wrong answer:* "the shell `execs` the command directly" — then the shell would replace *itself* and never return to the prompt; it `forks` first so a child does the `exec`. **Feed-forward:** §14.3.

14.7 (a) `fork` is cheap because of **copy-on-write**: it shares the parent's physical pages read-only rather than copying them, so it does work proportional to pages later modified, not to address-space size. (b) A workload where the child promptly **writes to nearly every inherited page** (e.g. it mutates a large data structure it inherited) gets little from COW — almost every page faults and is copied anyway, so you pay close to a full copy. (c) Process-per-request pays **process-creation and context-switch overhead** and needs IPC to share data; thread-per-request pays the **loss of isolation** (a crash or memory bug in one thread takes down all requests in the process) and the burden of synchronizing shared state. *Tempting wrong answer*: "COW makes `fork` free regardless of workload" — a write-heavy child erodes the benefit. **Feed-forward**: §14.2, and Ch 15.

14.8 Step 4: The accumulating state is **zombies** — terminated children whose exit status the kernel is still holding because the parent never collected it. They eventually stop new `forks` because each zombie occupies a **process-table slot**, and once the table (a finite resource) fills, the kernel cannot allocate a slot for a new process, so `fork` fails. This is a **defined resource leak, not undefined behavior** — the program violates no language rule; it simply fails to clean up (like a memory leak). Fix: the parent must **reap** its children — call `wait/waitpid` (e.g. from a `SIGCHLD` handler, or a reaping loop) so each child's status is collected and its slot freed. *Tempting wrong answer*: "the zombies are using up memory/CPU" — a zombie holds essentially no memory or CPU, only a table slot and its exit status; the leak is of slots. **Feed-forward**: §14.4.

14.9 (a) **`fork then exec`** — `fork` a child, have it `exec /bin/ls`, and the parent waits and continues (a lone `exec` would replace your program permanently). (b) **`fork` alone** — you want a second process running the *same* code, which is exactly what `fork` gives (no `exec` needed). (c) **`fork then exec`** — the shell `forks` and the child `execs` the typed command, while the shell waits. *Method-selection cue*: same program, new process → `fork`; different program, replace this one → `exec`; new process running a different program → both. **Feed-forward**: §14.2–14.3.

14.10 (a) **A defined-but-real defect** — the expectation is simply wrong: after `fork` the two processes have independent address spaces, so the child's changes are invisible to the parent (not UB, just a mistaken mental model that will produce bugs). (b) **A defined-but-real defect** — never reaping children leaks zombies (a resource leak, defined). (c) **Well-defined and fine (though usually unintended)** — both flushing a buffer copied across `fork` is defined behavior; it prints the line twice, which is a correctness surprise to fix with `fflush` before `fork`, but not UB. (d) **Well-defined and fine** — a child `execing` a new program while the parent runs is the normal pattern. *Tempting wrong answer*: calling (a) or (c) undefined behavior — both are defined; the issues are a wrong assumption and a duplicated buffer, not language-rule violations. **Feed-forward**: §14.2–14.4.

14.11 (a) A buffer overflow in process A writes through A's pointers, which are virtual addresses translated by **A's page tables** (Ch 4); they cannot name B's physical memory, and A cannot rewrite the translation tables to escape because installing page tables is a **privileged instruction** that traps in user mode (Ch 13). So the two mechanisms — per-process address translation and the user/kernel mode bit — confine the corruption to A. (b) Because a thread shares its process's address space, a compromised or crashing thread can corrupt the whole process; a separate *process* has its own address space and authority, so untrusted code in it cannot reach the host's memory — making the process the right sandbox unit. (c) Putting each tenant in its own process means one tenant's crash, memory corruption, or compromise cannot read or damage another tenant's data — fault and security isolation come for free from the process boundary. *Tempting wrong answer*: "a thread is isolated enough because it has its own stack" — a private stack does not isolate memory; all threads share the heap and globals and one bug corrupts them all. **Feed-forward**: §14.1, Ch 11, Ch 13, and Ch 15.

14.12 Machine-specific by design. A correct write-up reports: (a) the demo prints two lines with different PIDs, the child's `ppid` equal to the parent's `pid`, and independent values of the variable (e.g. child 101, parent 110) — confirming `fork` returned twice into two separate address spaces; (b) with a `printf`d line *before* the `fork` and output redirected to a file (block buffering), the line appears **twice**, because `fork` copied the unflushed buffer into both processes and both flushed it at exit; (c) adding `fflush(stdout)` before the `fork` makes the line appear once, because the bytes crossed to the kernel before the buffer was duplicated. Essential discipline: report exactly what you observed and label your compiler/system. *Tempting wrong answer*: "the double line means `fork` ran the `printf` twice" — it ran once; `fork` duplicated the *buffer*, not the print. **Feed-forward**: §14.2, §14.4, and Ch 13.

Chapter 15 — Solutions

15.1 Threads in a process **share** the address space — code (text), the heap, global/static data, and open file descriptors. Each thread keeps **private** its own stack, its own registers (program counter and stack pointer), and its thread-local storage. *Tempting wrong answer*: "each thread has its own heap" — no; the heap is shared, which is exactly why threads can pass data by pointer and also why they race. **Feed-forward**: §15.1.

15.2 `pthread_create(&tid, attr, fn, arg)` takes a pointer to a `pthread_t` to fill in, optional attributes, the function `fn` (of type `void (*)(void *)`) the new thread will run, and the `arg` passed to it; it starts `fn(arg)` running concurrently. `pthread_join(tid, &ret)` blocks until that thread finishes and collects its return value. *Tempting wrong answer*: "`pthread_create` runs the function and returns its result" — it starts the function on a *new* thread and returns immediately; `pthread_join` is what waits. **Feed-forward**: §15.2.

15.3 Linux uses the **one-to-one (1:1)** model — each pthread is its own kernel-scheduled thread (implemented by NPTL). `pthread_create` uses the `clone` system call underneath, with flags including `CLONE_VM` to share the address space. *Tempting wrong answer*: "Linux threads are M:N green threads scheduled in user space" — that was never NPTL; Linux threads are 1:1 kernel threads (language runtimes may add their own user-space scheduling above them). **Feed-forward**: §15.3.

15.4 `pthread_t t1, t2; pthread_create(&t1, NULL, worker, NULL); pthread_create(&t2, NULL, worker, NULL); pthread_join(t1, NULL); pthread_join(t2, NULL);` with `void *worker(void *a) { ... return NULL; }`. `pthread_create` is the thread analog of `fork` (start a new flow of control) and `pthread_join` of `wait` (wait for it and collect its result). The joins are necessary because if `main` returned first, the process would exit and tear down the shared address space the workers are still using. *Tempting wrong answer*: skipping the joins "because the threads will finish anyway" — `main` returning ends the process and can kill still-running threads. **Feed-forward**: §15.2, and Ch 14.

15.5 `counter++` compiles to a **read-modify-write**: (1) load `counter` into a register, (2) add one, (3) store it back. With two threads and no synchronization, an interleaving such as — A loads 5, B loads 5, A computes 6, B computes 6, A stores 6, B stores 6 — has both threads increment from the same starting value, so two increments produce a net change of only **one**; one update is lost. Repeated across millions of iterations, many updates vanish. *Tempting wrong answer*: "`++` is a single atomic operation" — in C it is not; it is three separate steps that can interleave. **Feed-forward**: §15.4.

15.6 Address space: processes have separate address spaces; threads share one.

Isolation: processes are isolated (a crash is contained); threads are not (a crash or memory bug kills the whole process). **Cost to create:** a thread is cheaper (no address-space copy / copy-on-write setup); a process is heavier. **Cost of sharing data:** processes need explicit IPC (pipes, shared memory); threads share by pointer for free — but that "free" sharing costs you the obligation to synchronize (data races). *Tempting wrong answer:* "threads are strictly better because they're cheaper" — they trade away isolation, which is sometimes exactly what you need. **Feed-forward:** §15.1, and Ch 14.

15.7 (a) Thread creation skips duplicating the address space — no copy-on-write page setup, no new page tables — because the new thread *shares* the existing address space; it just needs a stack and a kernel-scheduled context. (b) In exchange you take on the **data race** hazard (and the loss of isolation): every piece of shared mutable state must now be synchronized or it is undefined behavior. (c) The loss is **not deterministic** — it depends on how the threads' read-modify-writes happen to interleave, which varies with scheduling, load, core count, and compiler, so the demo gave a different wrong total each run. Its nondeterminism means **tests cannot be trusted to catch it:** a racing program can pass repeatedly (even print the exactly-right answer) and still be broken. *Tempting wrong answer:* "if the test passes, there's no race" — a race is defined by the rules (unsynchronized conflicting access), not by whether a run misbehaved. **Feed-forward:** §15.4.

15.8 Step 4: The bug is a **data race** on `hits` — multiple threads performing unsynchronized read-modify-writes (`hits++`) on shared memory — and in C/C++ it is **undefined behavior**. The unit tests missed it because they exercised requests **one at a time**, so no two increments ever overlapped and nothing raced; the bug only manifests under genuine concurrency, which the tests never created. The fix category is **synchronization:** make the update atomic — e.g. protect it with a **mutex** (`pthread_mutex_lock/unlock` around `hits++`) or use an **atomic** increment (a `_Atomic/ std::atomic` counter, or an atomic fetch-add). *Tempting wrong answer:* "add more tests" — single-threaded tests, however many, cannot exercise a race; you need concurrency plus a race detector, and the real fix is synchronization. **Feed-forward:** §15.4.

15.9 (a) **Separate process** — isolation is the whole point; an untrusted plugin in its own address space cannot corrupt or crash the host, whereas a thread shares memory and could take everything down. (b) **Thread** — the workers all read one large in-memory array, so shared-address-space threads let them access it directly with no copy or IPC (and reads alone do not race). (c) **Separate process** (via `fork + exec`) — running a typed command means launching a different program, which is a process, not a thread. *Method-selection cue:* need containment/a different program → process; need cheap shared-memory work → thread. **Feed-forward:** §15.1, and Ch 14.

15.10 (a) **Undefined behavior** — two threads writing the same non-atomic variable with no synchronization is a data race. (b) **Well-defined and fine** — writing two *different* variables with no shared location is not a race (barring false-sharing performance effects, which are a cost, not UB). (c) **Undefined behavior** — reading through a freed pointer is a use-after-free (Ch 9), regardless of threads. (d) **Well-defined and fine** — concurrent *reads* of a shared value, with no writer, do not race. *Tempting wrong answer*: calling (b) or (d) a race — a data race requires a conflicting access (at least one write) to the *same* location; disjoint writes and read-only sharing are fine. **Feed-forward**: §15.4, and Ch 9.

15.11 (a) The added clause: **no unsynchronized concurrent access to shared mutable memory** (any two threads' accesses to the same location, at least one a write, must be ordered by synchronization). (b) Violating it is a data race, which is **undefined behavior** in C/C++; unlike a use-after-free — which reliably touches freed memory every run — a race depends on interleaving, so it **may never reproduce** under testing even though the program is incorrect for some schedule. (c) A language can move this to compile time with a type system that tracks sharing and mutation: Rust's `Send/ Sync` traits mark which data may cross or be shared between threads, and the borrow checker forbids aliasing a mutable reference — so an unsynchronized shared mutation is a *compile error*, not a runtime gamble. *Tempting wrong answer*: "a thorough test suite proves there's no race" — no finite set of runs can, because the offending interleaving may simply never occur; only reasoning from the rules (or a checker/ detector) settles it. **Feed-forward**: §15.4, §15.5, and Part 5 (Rust).

15.12 Machine-specific by design. A correct write-up reports: (a) the four-thread, one-million-each unsynchronized counter prints totals well below 4,000,000 that differ every run (on the book's machine ~1.1-1.8 million across five runs) — lost updates from interleaved read-modify-writes; (b) adding a mutex around the increment (or using an atomic increment) restores exactly 4,000,000 every run; (c) the bug's visibility depends on build and variable: at `-O2` without `volatile` the optimizer may collapse each loop into a single add, shrinking the race window so the "wrong" answer hides (it printed 4,000,000 for the author), while `-O0` or a `volatile` counter forces per-iteration memory writes and exposes the loss — demonstrating that a correct-looking run does not mean the code is correct. Essential discipline: report exactly what you saw and label compiler/core count/hardware. *Tempting wrong answer*: "the version that printed 4,000,000 is the correct one" — it is the same racy code; it merely failed to manifest. **Feed-forward**: §15.4.

Chapter 16 — Solutions

16.1 The states are **running** (on a CPU now), **ready** (runnable but no free CPU), and **blocked** (waiting for an event — I/O, a child, a lock — and unable to run until it arrives). The scheduler chooses the next runner from the **ready** threads only; a blocked thread is not a candidate until the event wakes it. *Tempting wrong answer*: "the scheduler picks among all threads" — it cannot run a blocked thread; picking one would accomplish nothing until its event fires. **Feed-forward**: §16.1.

16.2 Turnaround time = completion – arrival; **response time** = first-run – arrival. For the job: response = 8 – 0 = **8**; turnaround = 40 – 0 = **32**. *Tempting wrong answer*: reporting response as the time it *finished* (40) — response measures when the job first got the CPU, not when it completed. **Feed-forward**: §16.1.

16.3 The **convoy effect**: short jobs pile up behind a long one (like cars behind a truck), so they inherit its runtime and average turnaround balloons. It afflicts **FIFO** (first-come-first-served, no preemption). *Tempting wrong answer*: attributing it to round-robin — RR's small slices are exactly what break the convoy. **Feed-forward**: §16.2.

16.4 FIFO (X, Y, Z): X completes at 8, Y at 12, Z at 16 → average turnaround $(8+12+16)/3 = 12$ ms. **SJF (Y, Z, X)**: Y at 4, Z at 8, X at 16 → average $(4+8+16)/3 = 9.33$ ms. SJF is lower, and it is no accident: when all jobs arrive together, running shorter jobs first minimizes the sum of completion times, which is exactly what average turnaround measures — SJF is provably optimal for it. *Tempting wrong answer*: assuming the last-completion time (16 ms in both) decides the average — it is the same, yet the averages differ because ordering changes the *earlier* completions. **Feed-forward**: §16.2.

16.5 SJF minimizes average turnaround because, with simultaneous arrivals, a job's turnaround includes the runtime of everything before it; ordering shortest-first makes each long job's runtime fall on the fewest later jobs (any swap putting a longer job earlier only adds its length to more waits). It is unusable in general because (1) it needs an **oracle** — you must know each job's length in advance, which for interactive and server work you never do — and (2) even given the oracle it gives poor **response time**: a job arriving after a long one starts must wait (or, with STCF, only if shorter). *Tempting wrong answer*: "just estimate the length" — MLFQ does something like this by observing behavior, but a pure SJF that mis-estimates loses its optimality and can starve. **Feed-forward**: §16.2, §16.3.

16.6 An MLFQ keeps several priority queues and always runs from the highest non-empty one (round-robin within a level). A job's priority *adapts*: it enters at the top; if it uses its entire time slice (CPU-bound), it is demoted a level; if it yields before the slice ends (blocks on I/O — interactive), it keeps its level. So interactive jobs, which block frequently, stay near the top and are scheduled quickly (low response time), while CPU-bound jobs sink and share the leftovers — the scheduler *learns* each job's character from behavior rather than being told it. *Tempting wrong answer*: "the programmer sets the priority" — MLFQ derives it from observed CPU/blocking behavior; a boost and anti-gaming accounting keep it honest. **Feed-forward**: §16.4.

16.7 (a) Shrinking the slice improves **response time** (jobs cycle onto the CPU sooner) and hurts **turnaround/throughput** (more interleaving, more context switches). (b) If the slice equals the switch cost, then for every unit of switching the CPU does one unit of useful work — roughly **half** the CPU is wasted on switching (and it only gets worse as the slice drops below the switch cost). (c) A few milliseconds is large enough that a few-microsecond switch is a tiny percentage overhead, yet small enough that interactive jobs still feel instant; seconds would wreck response time, microseconds would drown in switch overhead. *Tempting wrong answer*: "smaller is always more responsive, so make it tiny" — below a threshold, switch overhead dominates and total throughput collapses. **Feed-forward**: §16.3.

16.8 Step 4: Lower the dashboard's `nice` value (or, better, put the two workloads in separate **cgroups** and give the dashboard a larger **CPU weight**). Under Linux's proportional-share scheduler this grants the dashboard a bigger *share* of the CPU when the two contend, while the batch job still runs freely whenever the dashboard is idle. It is safer than `SCHED_FIFO` because it changes the *proportion*, not the absolute priority: a low-weight job still gets *some* CPU, so nothing starves and a runaway loop cannot freeze the machine, whereas a CPU-bound `SCHED_FIFO` thread can starve everything on its core, kernel threads included. *Tempting wrong answer*: "give the dashboard real-time priority" — that fixes the starvation by creating a worse one and risks wedging the box. **Feed-forward**: §16.4.

16.9 (a) **Response time** — the editor must react to each keystroke immediately; the user feels lag directly. (b) **Turnaround time** — the batch job's only requirement is that the whole thing finishes by a deadline; no user is watching it start. (c) **Throughput** — a build farm is judged by jobs completed per hour in aggregate, not any one job's latency. *Method-selection cue*: a human waiting on each action → response; a job that must complete → turnaround; a fleet maximizing aggregate work → throughput. **Feed-forward**: §16.1.

16.10 (a) **False** — a process blocked on `read` is off the CPU entirely; it consumes no CPU until the data arrives and it becomes ready. (b) **False** — Linux threads are 1:1 kernel-scheduled entities (Chapter 15), so the scheduler picks among *threads*. (c) **False** — a context switch costs register save/restore plus cold cache/TLB; the slice size trades responsiveness against that overhead. (d) **True** — a lower `nice` value means a larger weight and thus a larger CPU share under contention. *Tempting wrong answer*: calling (d) false because "nice" sounds like it should lower priority — the naming is inverted: lower `nice` is *higher* priority. **Feed-forward**: §16.1, §16.4, and Ch 15.

16.11 (a) Strict SJF always runs the shortest ready job; if short jobs keep arriving, a long job is never the shortest and **never runs** — indefinite starvation. (b) An MLFQ periodically **boosts** every job back to the top priority, guaranteeing even a long, demoted job eventually gets CPU; without the boost, CPU-bound jobs stuck at the bottom could starve behind a steady stream of interactive work. (c) A proportional-share scheduler gives every job CPU in proportion to its weight, so a low-weight job gets a *small* share, never *zero* — "smaller share" is not "no share." That is exactly why `nice` cannot starve a task (it only shrinks the proportion) while `SCHED_FIFO`, being strict priority, can (a higher-priority runnable thread gets *all* the CPU). *Tempting wrong answer*: "a low `nice` value can starve a process" — no; only strict-priority classes can starve, which is the whole safety argument for using `nice`/weights instead. **Feed-forward**: §16.4.

16.12 Machine-specific by design. A correct write-up reports: (a) a loop that spins for `N` wall-seconds and prints the CPU-seconds it accrued; (b) pinned to one shared core, the `nice 0` copy gets far more CPU than the `nice 19` copy — on the book's machine (Linux 6.18, EEVDF) roughly 4.66 vs 0.09 CPU-seconds over a five-second window, about fifty to one; (c) on separate cores the two no longer compete for the same CPU, so each gets a full core and the `nice` difference nearly vanishes — `nice` only matters *under contention for the same CPU*. The discipline: report exactly what you saw and label kernel/core count/scheduler. *Tempting wrong answer*: "`nice 19` was paused/slept" — it was not; it ran continuously but was granted a much smaller *share* of the contended core. **Feed-forward**: §16.4.

Chapter 17 — Solutions

17.1 A **virtual page** is a fixed-size chunk (4 KiB on x86-64) of a process's virtual address space; a **physical frame** is an equal-size chunk of physical RAM. A page-table entry's **present bit** records whether that virtual page currently occupies a frame (present) or does not (not present). *Tempting wrong answer*: "pages and frames are the same thing" — they are the same size, but a page is virtual and a frame is physical; the page table is exactly the mapping between them. **Feed-forward**: §17.1.

17.2 A **page fault** is a hardware exception raised when a process accesses a virtual page whose PTE is not present; it traps into the kernel's fault handler. The two broad outcomes: (1) the access is **legal**, so the handler materializes the page — finds a frame, zeroes it (anonymous) or reads it from a file/swap, sets present, and restarts the instruction; or (2) the access is **illegal**, so the handler delivers SIGSEGV. *Tempting wrong answer*: "a page fault means a crash" — most page faults are the normal, successful mechanism of demand paging; only the illegal case becomes a fault the program sees. **Feed-forward**: §17.2.

17.3 A **minor** fault is satisfied without disk I/O — the page comes from a zeroed frame or one already in memory (e.g. the page cache). A **major** fault must read the page in from a file or swap, costing a disk access. *Tempting wrong answer*: "all page faults hit the disk" — anonymous demand-zero and cache-hit faults are minor and never touch the disk. **Feed-forward**: §17.2, §17.3.

17.4 The `mmap` only recorded a 1 GiB mapping in the page table with the present bits clear — no frames spent — so RSS was just the program itself immediately afterward. As the program touched pages, each first touch raised a page fault, the handler gave that page a frame and set present, and RSS climbed one page at a time toward 1 GiB. The allocation did not consume 1 GiB up front because **demand paging** makes allocation a promise; *residency* is paid per touched page. *Tempting wrong answer*: "the OS was slow to account for the memory" — nothing was resident to account for; the frames genuinely did not exist until the touches. **Feed-forward**: §17.2.

17.5 Reference 1 2 3 1 4 2 5, 3 frames. **FIFO** = 5 faults: [1]→[1,2]→[1,2,3]→hit 1→evict 1 for 4 [2,3,4]→hit 2→evict 2 for 5 [3,4,5]. **LRU** = 6 faults: [1]→[1,2]→[1,2,3]→hit 1 (1 now MRU) [2,3,1]→evict LRU 2 for 4 [3,1,4]→evict LRU 3 for 2 [1,4,2]→evict LRU 1 for 5 [4,2,5]. Here **FIFO does better** (5 vs 6) — a deliberate reminder that LRU is a *heuristic* that pays off when the reference stream has locality; on a short string engineered without it, LRU can lose. Only OPT is guaranteed optimal. *Tempting wrong answer*: "LRU always beats FIFO" — LRU usually wins on real (local) workloads, but not on every string; this one is a counterexample. **Feed-forward**: §17.4.

17.6 Belady's anomaly is that giving a program more page frames can *increase* its fault count. It is exhibited by **FIFO** (on 1 2 3 4 1 2 5 1 2 3 4 5, FIFO takes 9 faults with 3 frames but 10 with 4). LRU and OPT cannot exhibit it because they are *stack algorithms*: their **resident set with more frames is a superset of the resident set with fewer**, so any page that would have been a hit with N frames is still resident with $N+1$ — adding frames can only remove faults, never add them. FIFO lacks this property because a larger queue reorders evictions and can eject a soon-to-be-reused page it would have kept. *Tempting wrong answer*: "more memory always means fewer faults" — true for LRU/OPT, false for FIFO; that surprise is the whole point of the anomaly. **Feed-forward**: §17.4.

17.7 (a) A minor fault costs a handful of instructions plus a zeroed or already-resident frame — sub-microsecond; a major fault waits on a disk. Using Chapter 3's hierarchy, DRAM access is tens of nanoseconds while an SSD read is tens of microseconds and a spinning disk seek is milliseconds — so a major fault is roughly a thousand to a million times slower than a minor one. (b) A *clean* page is an unmodified copy of data still on disk (or a file), so the OS can drop it and re-read later; a *dirty* page has been modified, so it must be **written out** to swap before its frame is reused — an extra disk write. (c) The likely cause is paging/swap pressure: the working set no longer fits, so the service takes major faults that stall it for disk-access times. The fix is *not* in the CPU-bound code — it is to shrink the memory footprint or add RAM so the working set fits. *Tempting wrong answer*: "profile and optimize the hot function" — the CPU is not the bottleneck; the cores are idle waiting on disk. **Feed-forward**: §17.3, §17.4, and Ch 3.

17.8 Step 4: The phenomenon is **thrashing** — the combined working set exceeds physical memory, so nearly every eviction removes a page another thread is about to need, producing a storm of **major** page faults. CPU utilization is near *zero* because the cores are almost always *blocked waiting on disk I/O* (paging pages in and out), not computing — the machine is I/O-bound on the swap device, not CPU-bound. The fix is to shrink the working set, not "optimize the loop": go back to one (or a few) sequential passes so the active window fits in RAM, or add memory. *Tempting wrong answer*: "add more threads / cores to speed it up" — more concurrent passes enlarge the working set and make the thrashing worse; the original single sequential pass was faster precisely because its working set fit. **Feed-forward**: §17.4.

17.9 (a) **Demand paging / overcommit** — the allocation only reserved mappings; the slowdown is the pages faulting in as they are first touched. Measure: RSS climbing over time (vs the reserved VSZ). (b) **Thrashing** — the dataset just crossed the size where the working set stopped fitting in RAM, so the system is now swapping. Measure: major-fault rate and swap I/O. (c) **Page cache / demand paging of a file mapping** — the first pass takes major faults reading the file from disk; the second pass hits the now-cached pages (minor or no faults). Measure: major faults on first pass vs near-zero on the second. *Method-selection cue*: "allocate huge/fast then slow" → demand paging; "small growth, huge slowdown" → crossed the working-set cliff (thrashing); "slow first, fast second" → cache warming. **Feed-forward**: §17.2, §17.3, §17.4.

17.10 (a) **False** — a successful `malloc` records mappings; physical RAM is consumed only as pages are touched (demand paging/overcommit). (b) **True** — a segmentation fault is a page fault the handler found no legal mapping for, so it delivered `SIGSEGV`. (c) **False** — `fork` uses *copy-on-write*: it copies page tables, not data; frames are shared read-only until a write faults and copies one page. (d) **False** — each process has its own page table, so the same virtual address in two processes normally maps to *different* frames (that is the isolation); they share a frame only through explicit shared memory or a shared file mapping. *Tempting wrong answer*: calling (d) true by analogy to threads — threads share one page table; distinct processes do not. **Feed-forward**: §17.1, §17.2, and Ch 4, Ch 14.

17.11 (a) `fork` copies the parent's *page tables*, mapping the child's pages to the same physical frames as the parent's, but marks every shared writable page **read-only** — so the "copy" is instantaneous and shares all data. (b) When the child (or parent) writes to a read-only shared page, the write raises a **protection page fault**; the handler sees a copy-on-write page, allocates a fresh frame, copies *that one page's* contents into it, remaps the writer's PTE to the new frame with write permission, and restarts the instruction — exactly **one** page is copied per first write. (c) This makes `fork-then-exec` nearly free even for a large parent: `exec` immediately replaces the address space, so almost none of the shared pages are ever written before being discarded — copying them eagerly would be wasted work. A single write to a shared page costs one page fault plus one page copy. *Tempting wrong answer*: "the first write copies the whole address space" — copy-on-write copies only the single page touched, on demand. **Feed-forward**: §17.5, and Ch 14.

17.12 Machine-specific by design. A correct write-up reports: (a) touching one byte per 4 KiB page across the region yields about one minor fault per page (on the book's machine, 25,672 minor faults and 0 major faults to touch 100 MiB), shown via `/usr/bin/time -v`; (b) `VmRSS` is tiny right after a large `mmap` and grows to about the mapping size after touching it all (on the book's machine, 1,364 kB → 525,976 kB for 512 MiB) — reservation vs residency; (c) if you dare to exceed RAM, major-fault and swap-I/O rates climb while CPU utilization drops, the signature of thrashing — stop before the box wedges or the OOM killer fires. Label kernel/page size/RAM. *Tempting wrong answer*: "RSS should equal the mapping size right after `mmap`" — it should not; that is exactly what demand paging prevents. **Feed-forward**: §17.2, §17.4.

Chapter 18 — Solutions

18.1 A **signal** is an asynchronous notification the kernel delivers to a process to indicate an event (a software interrupt). It differs from a normal function call in that the program never *calls* the handler: the kernel *injects* it into the control flow at an arbitrary instant, suspending whatever was running, then resumes that code when the handler returns. *Tempting wrong answer*: "a signal is just a callback you invoke" — you register it, but delivery is forced by the kernel, not called by your code. **Feed-forward**: §18.1, §18.2.

18.2 **Terminate** (SIGTERM, SIGINT); **Terminate and dump core** (SIGSEGV, SIGABRT); **Ignore** (SIGCHLD's default); and **Stop/Continue** (SIGSTOP/SIGCONT). *Tempting wrong answer*: "the default is always to terminate" — several signals default to ignore or to stop/continue, not termination. **Feed-forward**: §18.1.

18.3 **SIGKILL** and **SIGSTOP** cannot be caught or ignored. If a process could catch them, it could refuse to die or to be stopped — a runaway or malicious program would be unkillable, so the OS would lose its last-resort control. Keeping these two unblockable guarantees the kernel can always terminate or halt any process. *Tempting wrong answer*: "you can catch SIGKILL to clean up first" — you cannot; that is exactly why SIGTERM (catchable) exists for graceful shutdown and SIGKILL for the forced kill. **Feed-forward**: §18.1.

18.4 When Ctrl-C sends SIGINT, the kernel suspends the blocked `read`, runs the handler, and then the `read` returns `-1` with `errno == EINTR` ("interrupted system call"). It did not resume waiting because the kernel had to unblock the process to deliver the signal, and the default is to report the interruption rather than silently re-block. The two standard fixes: (1) **retry** the call when `errno == EINTR` (a loop), or (2) install the handler with **SA_RESTART** so the kernel auto-restarts many interrupted calls — though not all, so robust code still handles `EINTR`. *Tempting wrong answer*: "treat `EINTR` as a real read error and abort" — it is not an error about the data; the correct response is to retry. **Feed-forward**: §18.3.

18.5 It is unsafe because a handler can run in the middle of arbitrary code, including inside `printf`/`malloc` themselves; those functions hold internal locks / mutable state and are not reentrant, so the handler's re-entry can deadlock or corrupt the heap. A function callable from a handler must be **async-signal-safe** (essentially reentrant — the `signal-safety(7)` list, e.g. `write`, `_exit`). The safe pattern: the handler sets a volatile `sig_atomic_t` flag (or does one `write` to a self-pipe) and returns; the main loop notices and does the real work synchronously. *Tempting wrong answer*: "it's fine if you only `printf` occasionally" — it only has to interrupt `printf` once, under load, to hang; rarity is what makes the bug dangerous, not safe. **Feed-forward**: §18.4.

18.6 Standard signals (1–31) are represented by a single pending bit each, so they do **not queue**: if the same signal is generated multiple times while pending/blocked, the process sees it once when it is delivered; the extra instances coalesce. So if several children exit close together, the parent may receive a single `SIGCHLD` for all of them — reaping exactly one child per signal would leave the others as zombies (Chapter 14). A correct handler therefore loops: `while (waitpid(-1, &status, WNOHANG) > 0) ;` reaps *all* currently-exited children per delivery. *Tempting wrong answer*: "one `SIGCHLD` per child, so reap one" — signals coalesce, so the count of deliveries does not match the count of events. **Feed-forward**: §18.4, and Ch 14.

18.7 (a) Setting a volatile `sig_atomic_t` flag is a single, atomic, async-signal-safe store — it cannot deadlock or corrupt anything — whereas doing the work directly risks re-entering unsafe libc; and it is cheap (one write). (b) The self-pipe / `signalfd` approach adds a little latency (a write then the main loop's next read) but buys correctness: the real work runs in the synchronous main flow where all functions are safe, and it integrates signals into an ordinary event loop. (c) It can never be a reliable counter because standard signals do not queue — multiple occurrences while pending collapse into one delivery, so the number of handler runs under-counts the number of events. *Tempting wrong answer*: "increment a counter in the handler to count events" — coalescing makes that count wrong under load. **Feed-forward**: §18.4.

18.8 Step 4: The bug is an **async-signal-safety violation**: the `SIGHUP` handler calls non-reentrant functions (`malloc`, `stdio fprintf`), so when the signal interrupts a request thread that is already inside one of them, the handler re-enters it and **deadlocks** (or corrupts the heap). It appears only under load because it requires the signal to land *during* an unsafe call; light testing rarely hits that window, but heavy request traffic keeps threads inside `malloc/stdio` a large fraction of the time, so the odds rise — the same schedule-dependent, "passes tests / breaks in production" pathology as a data race. The fix: make the handler trivial (set volatile `sig_atomic_t` `reload_requested = 1`; or write a self-pipe byte and return), and perform the config re-read, `malloc`, and logging in the **main loop** when it observes the flag, where libc is safe to call. *Tempting wrong answer*: "add a lock around the handler's `malloc`" — the handler may already hold or be blocked on that very lock via the interrupted call; locking cannot make a handler reentrant. **Feed-forward**: §18.4.

18.9 (a) **Signal** (`SIGHUP` by convention) — a content-free "reload" notification is exactly what a signal is for. (b) **Real IPC — a pipe** — a signal carries no data; a byte stream needs a pipe (or socket). (c) **Signal** — this *is* a signal (`SIGCHLD`); the parent then waits. (d) **Real IPC — a socket** (or a pipe pair) — structured request/reply is a conversation, and a signal is only a doorbell. *Method-selection cue*: a content-free notification → signal; moving data or a request/reply → pipe/socket/shared memory. **Feed-forward**: §18.1, §18.5.

18.10 (a) **False** — the kernel injects the handler asynchronously; the program does not call it. (b) **False** — `SA_RESTART` restarts *many* but not all interrupted calls, so robust code still handles `EINTR`. (c) **False** — `SIGKILL` cannot be caught, so no cleanup runs; use `SIGTERM` for graceful shutdown. (d) **True** — an invalid memory access causes the kernel to deliver `SIGSEGV`, whose default disposition terminates (and cores) the process. *Tempting wrong answer:* calling (d) false because "a segfault is a hardware thing" — the hardware fault is *translated* by the kernel into the `SIGSEGV` signal. **Feed-forward:** §18.1, §18.2, §18.3, and Ch 11.

18.11 (a) `volatile` stops the compiler from caching the flag in a register or optimizing the read away, so the main loop actually re-reads memory the handler wrote; `sig_atomic_t` guarantees the read/write is a single uninterruptible operation, so the handler cannot catch a half-written value. Without both, the compiler might never observe the change, or the main loop might see a torn value. (b) It is *similar* to a data race: two flows of control touch the same memory with no ordering, so the shared access needs care. It is *different* in that there is one thread — the handler cannot run truly simultaneously with the main flow, only interrupt it — so the concern is atomicity of a single access and compiler visibility, not simultaneous multi-core writes (which is why `sig_atomic_t` suffices here where full synchronization is needed between threads). (c) Because the handler is a second, asynchronously-scheduled flow of control sharing memory with the main flow, a single-threaded program with signals already has concurrency — the interruption can happen between any two instructions. *Tempting wrong answer:* "a single-threaded program can't have concurrency bugs" — signals introduce asynchronous control flow, which is concurrency. **Feed-forward:** §18.4, and Ch 15.

18.12 Machine-specific by design. A correct write-up reports: (a) with a plain handler (no `SA_RESTART`), the `read` interrupted by `SIGALRM` returns `-1` with `errno == EINTR` (on the book's machine, "errno=4 (Interrupted system call)"); reinstalling with `SA_RESTART` makes the kernel restart the `read` so it keeps blocking (no `EINTR` surfaced). (b) The counting handler runs **once**, not five times, because the five blocked `SIGUSR1`s coalesce into one pending bit ("handler count so far = 0" while blocked, then "ran 1 time(s)" after unblock). (c) The first shows a caught signal interrupts a blocking call; the second shows standard signals do not queue. Label compiler/kernel. *Tempting wrong answer:* "the handler should run five times" — non-queuing standard signals collapse the repeats. **Feed- forward:** §18.3, §18.4.

Chapter 19 — Solutions

19.1 The toolkit splits into channels the kernel **copies** bytes through and arrangements that **share** the same physical memory. Copy side: a **pipe** (or FIFO, or socket) — bytes go sender → kernel buffer → receiver. Share side: **shared memory** (`MAP_SHARED`) — both processes map the same frames and touch them directly, moving no data. *Tempting wrong answer*: "sockets are the sharing mechanism" — a socket copies through the kernel like a pipe; only shared memory truly shares. **Feed-forward**: §19.1.

19.2 A bare pipe usefully connects processes that share its descriptors, which happens because file descriptors survive `fork` — so a parent creates the pipe and forks, and parent and child inherit the ends. A **FIFO** (named pipe, made with `mkfifo`) adds a *filesystem name*, so *unrelated* processes with no common ancestor can meet by opening the same path. *Tempting wrong answer*: "two unrelated processes can just call `pipe()` and connect" — a plain pipe's descriptors are private to the creator and its descendants; unrelated processes need a FIFO (or a socket). **Feed-forward**: §19.2.

19.3 When all write ends are closed, a `read` returns **0** — end-of-file — which is how the reader learns the writer is done. Writing to a pipe whose read ends are all closed raises **SIGPIPE** (default: terminate the writer), or, if that signal is ignored, fails the `write` with `errno == EPIPE`. *Tempting wrong answer*: "read blocks forever when the writer is gone" — it returns 0; it blocks only while a writer still holds the write end but has sent nothing. **Feed-forward**: §19.2.

19.4 A pipe is a byte stream: the kernel concatenates all writes and does not record how many writes produced the bytes, so two 3-byte writes are indistinguishable from one 6-byte write — measured as one `read` returning "ABCDEF". For discrete messages you must add **framing**. Two schemes: (1) a fixed-size **length prefix** before each message (read the prefix, then read exactly that many bytes); (2) a **delimiter** (e.g. newline-terminated records), reading until the delimiter. *Tempting wrong answer*: "just do one `write` per message and one `read` per message" — the stream may coalesce or split regardless of how you wrote it, so the boundary is not preserved. **Feed-forward**: §19.2, §19.4.

19.5 Fastest because it **copies nothing**: both processes map the same physical frames (Chapter 17's shared mapping), so a store by one is a load by the other with no kernel round-trip. Most hazardous because it provides sharing and *nothing else* — no ordering, no mutual exclusion. `MAP_SHARED`: the child's write of 42 is visible to the parent (`shared == 42`). `MAP_PRIVATE`: copy-on-write splits the page on the child's write, so the parent still sees 0 (`private == 0`). Two processes racing a shared counter lost ~700,000 of two million updates — the Chapter 15 data race across processes. *Tempting wrong answer*: "shared memory is safe because there is only one copy of the data" — one copy is exactly why concurrent unsynchronized writes corrupt it. **Feed-forward**: §19.3, and Ch 15.

19.6 A **stream socket** (`SOCK_STREAM`, TCP) is a reliable, ordered byte stream with no message boundaries; a **datagram socket** (`SOCK_DGRAM`, UDP) preserves message boundaries but may lose, duplicate, or reorder datagrams. (a) reliable in-order request/response → **stream** (TCP), which gives you reliability and ordering for free (add framing for message boundaries). (b) low-latency telemetry tolerating loss → **datagram** (UDP), whose per-message delivery and lack of retransmission fit lossy, latest-value data. `AF_UNIX` makes the socket a *local* (same-host) channel — faster, and able to pass file descriptors — while `AF_INET` uses the same API to cross machines. *Tempting wrong answer*: "UDP is just faster TCP, always prefer it" — UDP drops reliability and ordering; choosing it means signing up to handle loss and reordering yourself. **Feed-forward**: §19.4.

19.7 (a) Through a pipe a byte is copied twice — sender → kernel buffer → receiver — plus system-call overhead; through shared memory it is copied **zero** times, because both processes read and write the same frames. Fewer copies and no per-message trap into the kernel make shared memory faster. (b) The 65536-byte buffer bounds how far ahead a writer can get: when it fills, the next write *blocks* until the reader drains it, so a fast producer is paced to the slow consumer — flow control (backpressure) for free. (c) The cost is that you must supply your own synchronization and a crash-recovery story: shared memory hands you unsynchronized concurrency, so you carry, by hand, the mutual exclusion and ordering the copying channels gave you automatically. *Tempting wrong answer*: "shared memory is strictly better because it is faster" — it trades the kernel's built-in serialization and flow control for engineering burden you now own. **Feed-forward**: §19.2, §19.3.

19.8 Step 4: The bug is a **data race** on the ring buffer's shared write index (and the slots): multiple producers perform an unsynchronized read-modify-write on the same index, so two can grab the same slot or advance the index inconsistently, garbling and duplicating lines. It appears only under load because the corruption requires two producers to touch the index at *overlapping* instants; light traffic rarely overlaps, heavy traffic overlaps constantly — the same schedule-dependent, "passes tests / breaks in production" pathology as any race. Fix 1 (synchronize in the shared region): place a mutex or semaphore *inside* the shared memory (a process-shared `pthread_mutex` or a POSIX semaphore both processes map) and hold it while advancing the index — or make the index a lock-free atomic. Fix 2 (change the channel): give each producer its own pipe or socket to the consumer, so there is no shared write index to race on at all; the kernel serializes each channel for you. *Tempting wrong answer*: "add a lock in each producer's private memory" — a lock only excludes if all racers take the *same* lock, which must therefore live in the shared region. **Feed-forward**: §19.3, and Ch 15.

19.9 (a) **Pipe/FIFO** — a one-way local byte stream is exactly a pipe; the shell builds it with `fork`. (b) **Shared memory** — 60 large frames a second on one host demands zero-copy; a socket's per-frame copies would waste bandwidth. (c) **Network socket** — different machines rules out every local mechanism; only a socket reaches across the network. (d) **Unix-domain socket** — bidirectional local exchange plus descriptor passing (`SCM_RIGHTS`) is a capability only Unix-domain sockets offer. *Method-selection cue*: local one-way stream → pipe; local high-volume data → shared memory; crosses machines → network socket; local bidirectional or fd-passing → Unix-domain socket. **Feed-forward**: §19.1–§19.4.

19.10 (a) **False** — TCP is a byte stream with no message boundaries; a send may be split or coalesced, so you must frame. (b) **False** — the pointer is meaningless in the other private address space; pass data through IPC, not a pointer. (c) **False** — shared memory provides no synchronization; concurrent writes race. (d) **False** — a standard signal carries only its number (real-time signals add a small payload); it is a doorbell, not a data channel (Chapter 18). (e) **True** — a read returns 0 once all write ends close, signalling end-of-file. *Tempting wrong answer*: calling (a) true because "TCP is reliable" — reliability (no loss, in order) is orthogonal to message framing; TCP is reliable *and* boundary-less. **Feed-forward**: §19.2, §19.4, and Ch 18.

19.11 (a) The producer fills the pipe's 65536-byte buffer, and its next write **blocks** until the consumer reads some out; it cannot run ahead or exhaust memory because the kernel buffer is bounded and a full pipe suspends the writer. It ends up waiting inside `write`. (b) This is **backpressure**: the slow stage throttles the fast one automatically. Over a raw shared-memory channel you get no such thing — you would have to build the bounded buffer and the block-when-full/block-when-empty logic yourself (which is precisely the producer-consumer problem of the next chapters). (c) When the consumer dies, all read ends close; the producer's next `write` raises **SIGPIPE** (default terminate) or, if ignored, returns `-1` with `errno == EPIPE`. The OS notifies rather than discards so the writer stops promptly instead of doing useless work forever pouring bytes to nobody. *Tempting wrong answer*: "the producer keeps buffering until it runs out of memory" — the bounded pipe blocks it long before that. **Feed-forward**: §19.2.

19.12 Machine-specific by design. A correct write-up reports: (a) one read after "ABC" and "DEF" returns six bytes, "ABCDEF", showing the erased boundary. (b) The non-blocking pipe fills near **65536** bytes (on the book's machine, exactly 65536) before `write` returns `EAGAIN`; `fcntl(fd, F_SETPIPE_SZ, n)` changes the capacity. (c) With `MAP_SHARED` the parent sees the child's write (42); with `MAP_PRIVATE` it does not (0), because the private page is copy-on-write. Label compiler/kernel. *Tempting wrong answer*: "the two 3-byte writes should come back as two reads" — a stream erases the boundary; expecting message preservation is the classic framing bug. **Feed-forward**: §19.2, §19.3.

Chapter 20 — Solutions

20.1 A **critical section** is a region of code that touches shared state and produces a wrong result if two threads execute it concurrently. **Mutual exclusion** is the property that at most one thread is inside such a region at any instant. `counter++` needs it because it is really load-add-store: without exclusion, two threads can load the same old value and each store back a value one greater, losing an update. *Tempting wrong answer:* "`counter++` is a single instruction, so it's atomic" — it compiles to several steps, and even a single instruction is not atomic across cores without special guarantees. **Feed-forward:** §20.1.

20.2 **Lock** (acquire) and **unlock** (release). Acquiring a mutex already held by another thread **blocks** the caller until the holder releases it, after which one waiter proceeds. *Tempting wrong answer:* "lock returns an error if the mutex is held" — a plain lock blocks; it is `trylock` that returns immediately without acquiring. **Feed-forward:** §20.2.

20.3 **Mutual exclusion**, **hold-and-wait**, **no preemption**, and **circular wait** — all four must hold. The most practical to break in application code is **circular wait**, by imposing a *global lock ordering* so every thread acquires multiple locks in the same sequence. *Tempting wrong answer:* "break mutual exclusion by not using locks" — you usually cannot drop exclusion without a race; ordering is the practical lever. **Feed-forward:** §20.4.

20.4 "If free, take it" is itself a read-modify-write: read the mutex's state, decide, write "held." Two threads can both read "free" before either writes "held," so both believe they acquired the lock and both enter the critical section — the exact data race of Chapter 15, now on the lock word itself. A hardware atomic (test-and-set, compare-and-swap, exchange) performs the read-and-conditional-write as one indivisible step no other core can split, so exactly one thread wins. Without it, the lock cannot be race-free, and a racy lock provides no exclusion. *Tempting wrong answer:* "disabling interrupts is enough" — that works on a uniprocessor for the kernel, but not across multiple cores in user space; you need the atomic. **Feed-forward:** §20.2, and Ch 15.

20.5 (a) Thread 1 locks A; thread 2 locks B; thread 1 tries to lock B (held by 2) and blocks; thread 2 tries to lock A (held by 1) and blocks — permanent circular wait. (b) Making both acquire A then B removes **circular wait**: whoever gets A first will get B before releasing A, so no cycle of waiters can form. (c) It requires the two paths to interleave at just the wrong instant, which is rare under light test load but common under production concurrency — the schedule-dependent "passes tests / breaks in production" pathology. *Tempting wrong answer:* "add a timeout and retry" — `trylock/back-off` can help but risks livelock and masks the design flaw; a global lock order fixes it outright. **Feed-forward:** §20.4.

20.6 Exclusion: at most one thread runs the guarded code, so the load-add-store cannot interleave. **Memory synchronization:** the POSIX unlock in one thread and the next lock in another establish a happens-before ordering, so writes made under the lock are guaranteed visible, in order, to the next holder. The second role is what makes the non-atomic `counter++` *well-defined*: the C/C++ memory model calls an unsynchronized conflicting access a data race, which is undefined behavior; the lock's ordering removes the "conflicting, unsynchronized" property, so the access is legal. Without a lock (or atomics), a plain shared read-write is UB — the compiler and hardware may reorder and cache it freely. *Tempting wrong answer:* "exclusion alone makes it correct" — exclusion without the memory ordering still leaves visibility/reordering undefined; POSIX locks supply both, which is why they work. **Feed-forward:** §20.2, and Ch 15.

20.7 (a) The $\sim 10\times$ is the atomic read-modify-write, the surrounding bookkeeping, and the memory-ordering fence the lock and unlock impose — each far more expensive than a bare register add. (b) Serialization (only one thread in the critical section), cache-line bouncing (the lock word and guarded data migrate between cores' caches, a cross-core miss per handoff), and context switches (a blocking mutex sleeps and wakes waiters via the scheduler). Cache-line bouncing tends to dominate when the critical section is tiny and threads spin/retry; context-switch cost dominates when threads actually sleep. (c) By Amdahl, the serial (locked) fraction caps speedup regardless of core count, so reducing the fraction of work under the lock lifts the ceiling more than shaving nanoseconds off each acquisition. *Tempting wrong answer:* "a faster lock implementation is the main win" — if a large fraction of work is serialized, no lock speed helps; shrink the critical section. **Feed-forward:** §20.3.

20.8 Step 4: The bug is a **deadlock** from inconsistent lock ordering: transfer $X \rightarrow Y$ and $Y \rightarrow X$ acquire the two account locks in opposite orders, so concurrent opposite-direction transfers form a circular wait and both threads block in lock forever. It is load- and timing-dependent because it needs the two transfers to start close enough in time to each grab one lock before the other grabs the second — rare in tests, common under production traffic. Fix: impose a **global lock order** on accounts — e.g. always lock the account with the smaller id first, regardless of transfer direction — so both a $X \rightarrow Y$ and a $Y \rightarrow X$ transfer lock the same account first and no cycle can form. This removes the **circular-wait** condition. *Tempting wrong answer:* "lock a single global bank mutex for every transfer" — correct but serializes all transfers, killing throughput; ordering per-account locks keeps concurrency. **Feed-forward:** §20.4.

20.9 (a) **Single mutex** — one shared counter, all threads contend on the same word; one lock suffices (or an atomic). (b) **Fine-grained per-bucket mutexes** — operations on different buckets are independent, so per-bucket locks let them proceed in parallel; one big lock would needlessly serialize them. (c) **No lock** — thread-local storage is not shared, so there is no race to prevent. (d) **No lock (or a read-optimized scheme)** — a rarely-written, frequently-read value can use no lock if updates are done safely (e.g. atomic pointer swap / RCU-style publish); a plain mutex on every read would serialize all requests. *Method-selection cue:* is the state shared and mutated? If not shared → no lock; if shared and independent partitions → fine-grained; if a single hot datum → one lock (or atomic). **Feed-forward:** §20.3.

20.10 (a) **False** — mutual exclusion requires racers to take the *same* lock; different locks do not exclude each other. (b) **False** — deadlocked threads are *blocked*, so CPU is near zero; 100% CPU is a spin/livelock, not a classic deadlock. (c) **False** — a mutex adds overhead and serializes; it makes code *correct*, not faster. (d) **False** — holding a lock longer increases contention and deadlock exposure; hold it for the smallest correct region. (e) **False** — Peterson/Dekker use only loads and stores but assume ordering guarantees; real, portable locks need a hardware atomic. *Tempting wrong answer*: calling (e) true by citing Peterson's algorithm — it is correct only under a memory model with sufficient ordering, which on real hardware you get via atomics/fences, so "no atomic at all" is misleading. **Feed-forward**: §20.2, §20.3, §20.4.

20.11 (a) **Priority inversion**: H is blocked on the lock L holds, so H can proceed only when L runs and releases it. (b) Under strict priority scheduling, medium M is runnable and higher priority than L, so it preempts L and never yields; L therefore never gets the CPU to release the lock, and H — though highest priority — waits behind both. (c) **Priority inheritance**: while L holds a lock a higher-priority thread (H) is waiting on, L temporarily inherits H's priority, so L now outranks M, runs, releases the lock, and H proceeds. It breaks the link where M starves L. (d) Holding a lock while doing unbounded work under real-time scheduling can stall higher-priority threads for that whole duration — bounded, minimal critical sections are essential when priorities are strict. *Tempting wrong answer*: "just raise H's priority" — H is already highest; the problem is L is too low to run, which inheritance, not H's priority, fixes. **Feed-forward**: §20.4, and Ch 16.

20.12 Machine-specific by design. A correct write-up reports: (a) the unlocked four-thread total is below 4,000,000 and differs each run (on the book's machine, e.g. 1,187,737 / 1,144,730 / 2,766,329), while the mutex version is exactly 4,000,000 every run. (b) A lock+increment+unlock is roughly an order of magnitude slower than a bare increment (the book measured ~1 ns vs ~12 ns, ~12×) — illustrative and machine-dependent. (c) Opposite lock orders hang (a deadlock; kill it after a timeout), while a consistent order always completes. Label compiler/kernel. *Tempting wrong answer*: "re-running the unlocked version until it prints 4,000,000 proves it's fine" — a lucky correct run does not make a racing program correct; correctness must come from the rules. **Feed-forward**: §20.1, §20.3, §20.4.

Chapter 21 — Solutions

21.1 Because a spinning thread stays *runnable*, so the scheduler keeps giving it a CPU while it does no useful work — it burns a core (and, if it spins holding the lock, prevents the very progress it waits for). A blocking wait takes the thread *off* the CPU until it is woken, freeing the core for threads with work. *Tempting wrong answer*: "spinning is faster because there's no context switch" — true only for waits shorter than a context switch; for any appreciable wait, spinning wastes a whole core. **Feed-forward**: §21.1.

21.2 `wait`, `signal`, and `broadcast`. `wait` (called with the mutex held) **atomically releases the mutex and sleeps**; when later woken it re-acquires the mutex before returning. `signal` wakes one waiter; `broadcast` wakes all. *Tempting wrong answer*: "`wait` keeps the mutex so no one else can interfere" — it must release it, or the signalling thread could never acquire the lock to change the state. **Feed-forward**: §21.2.

21.3 At most **N** threads can be inside the guarded section at once: the semaphore starts with N permits, each `wait` takes one (blocking at zero), each `post` returns one. *Tempting wrong answer*: "it lets exactly N threads run and blocks the rest forever" — blocked threads proceed as permits are posted back; N is a concurrency cap, not a hard total. **Feed-forward**: §21.3.

21.4 Two reasons. (1) **Spurious wakeups**: POSIX permits `cond_wait` to return with no signal, so the predicate may not hold on waking. (2) **Stolen condition**: even after a real signal, between the signal and this thread re-acquiring the mutex, another thread can run and consume the item, so the predicate is false again. Two-consumer example: buffer has one item; producer signals; consumer C1 and C2 both were waiting; C1 wakes, re-acquires, pops the item; C2 wakes next, re-acquires — with `if` it does not re-check and pops from the now-empty buffer, reading garbage or crashing. With `while`, C2 re-tests `count == 0` and waits again. *Tempting wrong answer*: "one signal per item means one wakeup per item, so `if` is fine" — spurious wakeups and stolen conditions both break that count. **Feed-forward**: §21.2, §21.4.

21.5 (a) A condition variable's `signal` with no waiter is lost — it wakes no one and leaves no trace. So if a producer signals in the window after a consumer checked the predicate but before it called `cond_wait`, the consumer misses the signal and sleeps forever though the item is present. (b) Discipline: update the shared state *under the mutex* and then `signal` (still holding, or immediately after, the lock), and have the waiter loop `while (!predicate) cond_wait(...)`; holding the lock means the waiter cannot be mid-check-and-not-yet-asleep when you signal, and the loop re-checks the state on waking. (c) A semaphore's count *remembers* the post: a `sem_post` with no waiter increments the count, so a later `sem_wait` succeeds immediately — the "signal before wait" ordering is harmless. *Tempting wrong answer*: "signal before locking to be fast" — that is exactly the lost-wakeup race. **Feed-forward**: §21.3, §21.4.

21.6

```

producer(x):
    lock(m)
    while (count == CAP)
        wait(not_full, m)
    buf[in]=x; in=(in+1)%CAP; count++
    signal(not_empty)
    unlock(m)

consumer():
    lock(m)
    while (count == 0)
        wait(not_empty, m)
    x=buf[out]; out=(out+1)%CAP; count--
    signal(not_full)
    unlock(m)

```

It cannot overflow because the producer waits on `not_full` while `count == CAP`, so it never enqueues into a full buffer; it cannot underflow because the consumer waits on `not_empty` while `count == 0`. The loops are necessary because a wakeup may be spurious or the slot/item may have been taken by another producer/consumer before this thread re-acquires the mutex, so the predicate must be re-checked. *Method-selection cue*: "wait until a condition over shared state, then act" is precisely mutex + condition variable in a loop.

Feed-forward: §21.2.

21.7 (a) The spinner stays runnable, so the scheduler keeps it on a CPU for the whole ~ 0.5 s ($\approx 100\%$ of a core = ~ 0.5 s CPU), while the blocker is off the CPU (blocked state) the entire time, waking only at the end (~ 0.001 s CPU). (b) A short spin beats blocking when the expected wait is *shorter than the cost of a context switch pair* (sleep + wake); the crossover quantity is that context-switch cost — below it, spinning avoids two scheduler round-trips. (c) `cond_wait` costs a context switch to sleep and another to wake (plus scheduler bookkeeping); it is still right for a long wait because it frees the core for real work, and the two switches are negligible against a long wait. *Tempting wrong answer*: "blocking is always cheaper" — for sub-microsecond waits the two context switches can cost more than a brief spin. **Feed-forward:** §21.1, and Ch 16.

21.8 Step 4: Bug 1 is `if-not-while`: a worker checks the queue once, waits, and on a spurious wakeup (or after another worker took the task) pops from an empty queue and crashes. Bug 2 is a **lost wakeup**: the producer signals without holding the lock and enqueueing under it, so a worker between its check and its `cond_wait` misses the signal, and since the condition variable does not remember it, the task sits with no one woken. Both are intermittent and load-dependent because each needs a specific interleaving that is rare under light load and common under heavy concurrency. Corrected:

```

worker:
    lock(m)
    while (queue.empty())
        cond_wait(cv, m)
    t = queue.pop()
    unlock(m); run(t)

producer(task):
    lock(m)
    queue.push(task)
    cond_signal(cv)    // under the lock, after enqueue
    unlock(m)

```

The `while` fixes bug 1 (re-check on wake); enqueue-under-lock-then-signal fixes bug 2 (no window to miss the signal). *Tempting wrong answer*: "signal before locking so the worker wakes sooner" — that reintroduces the lost wakeup. **Feed-forward:** §21.2, §21.4.

21.9 (a) **Plain mutex** — a single shared counter needs exclusion, not waiting. (b) **Counting semaphore** initialized to 5 — "at most 5 at once" is exactly a permit pool. (c) **Condition variable + mutex** — "block until a predicate over shared state (an item exists) becomes true" is the condition-variable pattern. (d) **Condition variable + mutex, using broadcast** — a one-shot event to *many* waiters wants a broadcast (or a latch); a single signal would wake only one. *Method-selection cue*: just exclude → mutex; cap concurrency at N → counting semaphore; wait for a state predicate → condition variable (broadcast to wake all). **Feed-forward**: §21.2, §21.3.

21.10 (a) **False** — a `cond_signal` with no waiter is lost, not remembered (that is a semaphore's property). (b) **False** — `cond_wait` releases the mutex while it sleeps and re-acquires it on waking. (c) **True** — a binary semaphore (0/1) behaves much like a mutex (with caveats about ownership). (d) **False** — even without spurious wakeups, a stolen condition (another thread consumes it first) still requires the `while`; keep the loop. (e) **True** — a mutex, like a signal-handler flag, is a mechanism for coordinating two flows of control over shared memory (Chapter 18's single-threaded concurrency generalizes). *Tempting wrong answer*: calling (d) true because "spurious wakeups are the only reason for the loop" — the stolen condition is a second, independent reason. **Feed-forward**: §21.2, §21.3, and Ch 18.

21.11 (a) A monitor enforces that all access to the shared state goes through procedures that automatically hold the monitor's mutex and coordinate via its condition variables, so the "lock before touching state," "signal under the lock," and "wait in a loop" rules are built into the structure rather than re-implemented (and forgotten) at each call site. (b) Java `synchronized/wait/notify` hides the mutex and condition variable behind object methods; a thread-safe blocking queue hides the mutex, the two condition variables, and the predicate discipline behind `put/take`. (c) Because the primitives are where the subtle, load-dependent bugs live (`if-not-while`, lost wakeups, lock ordering), using a vetted higher-level abstraction removes whole bug classes; you would still drop to raw condition variables for a custom waiting policy the high-level type does not express (e.g. priority-ordered wakeups, or waiting on a compound predicate over several structures). *Tempting wrong answer*: "always hand-roll for performance" — the abstraction's overhead is usually negligible against the cost of a concurrency bug in production. **Feed-forward**: §21.4.

21.12 Machine-specific by design. A correct write-up reports: (a) the busy-poll waiter consumes ~100% of a core for the wait (on the book's machine ~0.499 s CPU for a ~0.5 s wait) while the condition-variable waiter consumes almost none (~0.001 s), because it is blocked, not runnable. (b) The bounded buffer's produced and consumed totals match (e.g. both 210 for items 1..20) and the buffer never exceeds its capacity, confirming correct coordination via the `while`-loops. (c) With a semaphore initialized to three, at most three of the ten threads are ever in the guarded section simultaneously. Label compiler/kernel; the CPU figures are illustrative. *Tempting wrong answer*: "the spinner is fine, it got the same result" — it wasted a core to do so; correctness plus efficiency both matter. **Feed-forward**: §21.1, §21.2, §21.3.

Chapter 22 — Solutions

22.1 Load `count` from memory into a register, add one, store the register back to memory. It is non-atomic across threads because another thread can read the same old value (or run in parallel) between the load and the store, so two increments produce one net change. *Tempting wrong answer:* "it's one C statement, so it's one operation" — one line of C is three machine instructions, and atomicity is a property of the machine operation, not the source. **Feed-forward:** §22.1.

22.2 The `LOCK` prefix makes the CPU hold the affected cache line exclusively for the entire duration of the read-modify-write, so no other core can read or write that line in between — the whole load-add-store is indivisible. `lock xadd` is **one** instruction, where a plain memory increment is three (`mov, lea/add, mov`). *Tempting wrong answer:* "it locks the whole bus so nothing else runs" — modern CPUs lock at the cache line via the coherence protocol, not the whole bus, except in rare unaligned cases. **Feed-forward:** §22.2.

22.3 On success it returns **true** (having written `desired` because `*obj` equalled `expected`). On failure it returns false and writes the *current* value of `*obj` into `expected`, leaving `*obj` unchanged. *Tempting wrong answer:* "on failure it leaves `expected` alone" — it overwrites `expected` with what it found, which is exactly what makes the retry loop reload for free. **Feed-forward:** §22.3.

22.4 (a) T1 load (0) → T2 load (0) → T1 add (1) → T1 store (1) → T2 add (1) → T2 store (1): `final count == 1`, one increment lost. (b) `atomic_fetch_add` compiles to a single `lock xadd` that holds the cache line exclusively for the whole read-modify-write, so T2's increment cannot begin until T1's has completed — the interleaving in (a) cannot occur. (c) At `-O2` the optimizer can collapse the million-iteration loop into a single `count += N`, shrinking each thread's race window to one load-add-store that rarely collides; a data race is undefined behavior, so "passes at `-O2`" proves nothing. *Tempting wrong answer:* "the `-O2` version is correct because it gives the right answer" — it is still a data race; it merely hides. **Feed-forward:** §22.1, §22.2.

22.5

```
void atomic_max(atomic_long *p, long v) {
    long cur = atomic_load(p);
    while (v > cur)
        if (atomic_compare_exchange_weak(p, &cur, v))
            return;
    /* installed v */
}
```

After a failed CAS, `cur` holds the value the CAS actually found in `*p`, so the `while` re-tests against fresh state and no explicit reload is needed. Use `_weak` in the loop because it is allowed to fail spuriously in exchange for cheaper code on load-linked/store-conditional machines, and a spurious failure merely costs one extra iteration. *Method-selection cue:* "update only if the old value satisfies a condition" → single CAS in a retry loop. **Feed-forward:** §22.3.

22.6 (a) `T1 atomic_load(&owner) == 0 (true) → T2 atomic_load(&owner) == 0 (true) → T1 atomic_store(&owner, 1) → T2 atomic_store(&owner, 2)`: both passed the test, both stored, and now T2 "owns" but T1 also believed it did. (b) `long expected = 0; if (atomic_compare_exchange_strong(&owner, &expected, my_id)) { /* I claimed it */ }` — exactly one thread finds `owner == 0` and swaps in its id; the other's CAS fails. (c) Any update that reads a value, decides based on it, and writes a value depending on that decision ("if it's still X, make it Y") must be a single CAS, because a separate atomic load and store leave a gap. *Tempting wrong answer*: "put the load and store close together so nothing runs between" — proximity in source is irrelevant; only a single indivisible operation closes the gap. **Feed-forward**: §22.3, §22.4.

22.7 (a) The atomic must own the cache line exclusively and acts as a reordering barrier, so it cannot overlap with surrounding work the way a plain store can; that serialization is the $\sim 15\times$. (b) Under eight-core contention the same line ping-pongs between cores' caches — each atomic must first pull the line to its own core in exclusive state (coherence traffic, Chapter 3) — so per-op cost rises well beyond the uncontended 15 ns, sometimes by an order of magnitude. (c) When updates vastly outnumber reads, give each thread its own counter (no sharing, no coherence traffic) and sum them only when a total is needed; the cost moves to the occasional summation and to cache footprint, not to every increment. *Tempting wrong answer*: "atomics have a fixed cost, so contention doesn't matter" — contention can dominate, which is why sharding a counter helps. **Feed-forward**: §22.4, and Ch 3.

22.8 Step 4:

```
void record_max(atomic_long *hi, long v) {
    long cur = atomic_load(hi);
    while (v > cur)
        if (atomic_compare_exchange_weak(hi, &cur, v))
            return;
}
```

It terminates because each failed CAS means another thread installed a value at least as large, which is written back into `cur`; so `cur` is non-decreasing, and once `cur >= v` the loop exits without another attempt. The invariant it guarantees: at all times `hi` is greater than or equal to every `v` that has been submitted, and eventually equals the maximum submitted. *Tempting wrong answer*: "it could loop forever under contention" — every retry is caused by a real increase in `hi`, which is bounded by the largest submitted value, so retries are finite. **Feed-forward**: §22.3.

22.9 (a) **Atomic fetch_add** — a fixed arithmetic RMW with a dedicated instruction. (b) **Single CAS** (`expected = 0`, swap in 1) — the success/failure result tells you who won. (c) **CAS retry loop** — set the new node's next to the current head, then CAS the head; retry if it moved (the Treiber stack of Chapter 23). (d) **Plain access** — a value no thread writes concurrently needs no atomicity. *Method-selection cue*: fixed arithmetic → `fetch_*`; one-shot "claim" → single CAS; update depending on current shared state → CAS loop; no concurrent writer → plain. **Feed-forward**: §22.2, §22.3.

22.10 (a) **False** — a mutex is a program built *from* atomics (plus a blocking slow path); the atomic is the hardware primitive. (b) **True** — on x86 the atomic and plain forms differ by the one-byte `LOCK` prefix. (c) **False** — the two atomics leave a gap another thread can slip through; only a single CAS composes compare-and-set. (d) **True** — that is exactly the consensus hierarchy: CAS has consensus number infinity, plain reads/writes number 1. (e) **False** — an uncontended atomic increment is roughly an order of magnitude dearer than a plain one (~ 15 ns vs ~ 0.9 ns here), and worse under contention. *Tempting wrong answer*: calling (c) true because "each part is atomic" — atomicity of the parts is not atomicity of the whole. **Feed-forward**: §22.2, §22.3, §22.4.

22.11 (a) Consensus number 1 means reads and writes can solve consensus only for a single thread (trivially) — for two, no protocol of plain reads and writes lets both irrevocably agree on one value regardless of timing; whichever writes "last" is timing-dependent and a reader cannot atomically observe-and-decide, so the two can diverge. (b) CAS at infinity means any number of threads can reach consensus, and from CAS you can build a wait-free implementation of *any* object with a sequential specification (a universal construction) — CAS is a universal primitive. (c) Every lock's fast path is a CAS (or test-and-set) claiming the lock word, and every lock-free structure is a CAS loop committing an update only if nothing changed; the hardware provides CAS specifically because, unlike plain accesses or fixed-arithmetic atomics, it is strong enough to build all of them. *Tempting wrong answer*: "fetch_add is just as universal as CAS" — fetch-and-add has a finite consensus number; it cannot express arbitrary conditional commits. **Feed-forward**: §22.4.

22.12 Machine-specific by design. A correct write-up reports: (a) the plain `-00` counter falls short of 4,000,000 (on the book's machine ~ 1.36 million, losing ~ 2.6 million) and varies run to run, while `atomic_fetch_add` is exactly 4,000,000 every run; the `-02` plain version may "pass" because the optimizer collapses the loop to one add per thread, shrinking the race window (still undefined behavior). (b) The disassembly shows the atomic form carries a `lock` prefix (`lock xadd`) absent from the plain increment. (c) The CAS-loop atomic max agrees with the true maximum under many threads. Label compiler/CPU; timing figures are illustrative. *Tempting wrong answer*: "the `-02` plain version works, so the bug was the optimizer's fault" — the program was always racy; `-02` only hid it. **Feed-forward**: §22.1, §22.2, §22.3.

Chapter 23 — Solutions

23.1 That the system as a whole always makes progress: out of all contending threads, some thread completes its operation in a bounded number of steps, so the structure can never be globally stuck. *Tempting wrong answer:* "every thread is guaranteed to complete" — that is wait-free; lock-free permits an individual thread to starve while others succeed. **Feed-forward:** §23.1.

23.2 Push: the successful CAS that installs the new node as `top`. Pop: the successful CAS that swings `top` to `old->next`. *Tempting wrong answer:* "the `malloc` (for push) or the return (for pop)" — the item is neither published nor removed until the CAS commits; that is the instant the operation takes effect. **Feed-forward:** §23.2.

23.3 A CAS checks only that a location still holds the same *bits* you read; if the value went $A \rightarrow B \rightarrow A$, the bits match and the CAS succeeds, even though the object those bits denote may have been removed, freed, and replaced. *Tempting wrong answer:* "the same value means nothing changed" — it means the value is *equal now*, not that it never changed. **Feed-forward:** §23.4.

23.4

```
push(v): n = new node(v)
        old = load(top)
        do { n.next = old } while (!CAS(&top, &old, n))
pop():  old = load(top)
        while (old) { if (CAS(&top, &old, old.next)) return old }
        return NULL
```

(a) A failed CAS means another thread changed `top` first; the failure reloads `old` with the current `top`, so the loop re-points `n.next` (push) or re-reads `old.next` (pop) and retries against fresh state. (b) The successful CAS is the single instant the operation takes effect for all threads — the linearization point — so concurrent operations order by whose CAS won. (c) Lock-free, not wait-free: the system always advances because some CAS always succeeds, but a persistently unlucky thread can keep losing the race and never complete. *Method-selection cue:* "update a shared pointer only if it hasn't moved" → CAS retry loop. **Feed-forward:** §23.2.

23.5 (a) It cannot be one CAS because enqueue must both link the new node (`tail.next: NULL -> n`) and move `Tail`, two separate locations; no single-word CAS updates both. (b) After the first CAS but before the second, a second thread can see `Tail` lagging — pointing at a node whose `next` is no longer `NULL`; "helping" means that thread advances `Tail` to `tail.next` itself (via CAS) rather than waiting for the enqueueer, so no one stalls on the laggard. (c) The dummy node guarantees the list is never empty, so `Head` and `Tail` always reference a real node and there are no null-pointer boundary cases at either end. *Tempting wrong answer:* "the second thread must wait for the enqueueer to finish" — helping is precisely what avoids that wait. **Feed-forward:** §23.3.

23.6 Thread A reads `top = A`, reads `A.next = B`, and is preempted before its CAS. Now: another thread pops A (`top = B`), pops B (`top = C`), and frees both; then it allocates a new node that `malloc` places at A's old address, sets some `next`, and pushes it, so `top = A` again — but this A is a different logical node and `A.next` is now C or garbage. A resumes and does `CAS(&top, A, B)`: it sees `top == A` (same bits), succeeds, and sets `top = B` — a node that was already popped and freed. The stack is now corrupt (dangling or lost nodes). A tagged pointer prevents it: the version counter incremented on each of those pops/pushes makes the pair `(A, version)` differ, so A's CAS fails. Hazard pointers prevent it too: A and B could not have been freed while A still referenced them. *Tempting wrong answer*: "the CAS is buggy" — the CAS is correct; the bug is freeing and reusing memory under it. **Feed-forward**: §23.4.

23.7 (a) Under low contention a CAS usually succeeds on the first try, so a push/pop is one atomic op with no blocking, no context switch, and no `futex` system call — cheaper than a mutex's lock/unlock (hence ~176 vs ~359 ns/op here); it spends its time on the single CAS rather than on scheduler round-trips. (b) Under very high contention on one hot pointer, most CASes fail and retry, and each attempt drags the cache line to the current core in exclusive state — a coherence storm — so throughput can fall below even the lock's, as work is wasted on failed attempts. (c) Elimination (pairing a push with a concurrent pop so both skip the stack) or per-thread sub-stacks cut contention; they trade strict LIFO/global ordering and simplicity for scalability. *Tempting wrong answer*: "lock-free is always faster" — under heavy contention a lock can win because it serializes instead of burning cycles on retries. **Feed-forward**: §23.2, and Ch 3.

23.8 Step 4: Hazard-pointer pop: (1) read `t = load(top)`; (2) publish `t` in this thread's hazard slot; (3) re-check `load(top) == t` — if it changed, restart, because a node could have been retired between (1) and (2); (4) now safe to read `next = t.next` and attempt `CAS(&top, t, next)`; (5) on success, `t` is unlinked — retire it to a per-thread list and free it later, only once no thread's hazard slot references it. This makes use-after-free impossible because a node any thread is dereferencing is published as hazardous and therefore never freed; and it makes ABA impossible because a node cannot be freed and its address reused while it is still referenced, so `top` cannot cycle back to a recycled `t` under a stalled reader. *Tempting wrong answer*: "just free in pop after the CAS" — another thread may still hold that node as its own `old`; the re-check-after-publish is what closes the window. **Feed-forward**: §23.4.

23.9 (a) **Lock-free** — an audio callback with a hard deadline must never block on a lock a slower thread holds. (b) **Mutex** — rarely contended and read once at startup; a lock is simpler and correctness is easier to reason about. (c) **Lock-free** (of a signal-safe kind) or a lock-free single-slot handoff — a signal handler must not take a lock the interrupted thread may already hold (Chapter 18's `async-signal-safety`). (d) **Mutex** — large, complex operations with modest contention favor the simplicity of a lock; making them lock-free would be intricate for little gain. *Method-selection cue*: must-never-block or signal context → lock-free; simplicity with tolerable contention → mutex. **Feed-forward**: §23.1, and Ch 18.

23.10 (a) **False** — lock-free code is built *from* atomic instructions (CAS); it means no locks and a system-progress guarantee, not no atomics. (b) **False** — a lock-free structure cannot deadlock; retries mean some other thread is succeeding (making progress), which is the opposite of deadlock, though one thread may starve. (c) **False** — the dummy node is a correctness device that removes null-boundary cases, not an optional optimization. (d) **True** — that is exactly the ABA problem: equal bits do not prove the object was untouched. (e) **True** — wait-free (every thread completes) implies lock-free (some thread completes) but not conversely. *Tempting wrong answer:* calling (b) true because "infinite retries look like being stuck" — the system is still making progress; only the unlucky thread is not. **Feed-forward:** §23.1, §23.3, §23.4.

23.11 (a) Helping solves the two-step enqueue: because linking the node and swinging `Tail` are separate CASes, an enqueueer can be preempted between them, leaving `Tail` lagging; a later thread advances `Tail` itself so it need not wait for the preempted enqueueer. (b) If every thread can finish any pending operation it encounters, then a thread preempted mid-update never blocks progress — whoever runs next completes (or helps complete) the operation, so the system always advances, which is the lock-free guarantee. (c) A lock has no helping: a preempted holder stalls everyone until it is rescheduled. Helping buys immunity to that at the cost of extra CASes (advancing another's tail) and more intricate invariants (every intermediate state must be well-defined and completable by anyone). *Tempting wrong answer:* "helping is just a nicety" — it is what makes the structure non-blocking despite the multi-step update. **Feed-forward:** §23.3.

23.12 Machine-specific by design. A correct write-up reports: (a) the eight-thread Treiber stack pushes and pops the full item count with none lost or duplicated (on the book's machine 800,000 in and 800,000 out, matching checksums). (b) Against a mutex-guarded stack under contention, ns/op for both, with the note that which wins depends on contention (here lock-free ~176 vs mutex ~359 ns/op). (c) The tagged-pointer version compiles with `-mcx16 -latomic`, passes the same correctness test, and the disassembly shows `lock cmpxchg16b`. Label compiler/CPU; throughput figures are illustrative. *Tempting wrong answer:* "the leaking version is fine since the test passes" — it passes only because it never frees; adding `free` without safe reclamation reintroduces ABA. **Feed-forward:** §23.2, §23.4.

Chapter 24 — Solutions

24.1 Sequential consistency: all cores observe memory operations as if interleaved into a single global order that respects each thread's program order. Real machines violate it because of the per-core **store buffer** (a store is buffered locally before becoming globally visible). *Tempting wrong answer:* "caches cause it, and cache coherence would fix it" — coherence keeps a single value per location consistent; the reordering here is the store buffer delaying visibility, which coherence does not prevent. **Feed-forward:** §24.1.

24.2 A data race is two accesses to the same location from different threads, at least one a write, not ordered by synchronization (happens-before), with at least one non-atomic. The standard says such a program has **undefined behavior** — no meaning at all, and the compiler may assume it cannot occur. *Tempting wrong answer:* "a data race just gives a torn or stale value" — it is undefined behavior, strictly worse than an unspecified value. **Feed-forward:** §24.2.

24.3 It guarantees atomicity (no torn or lost updates) and coherence on the same object; it does *not* impose any ordering on other memory operations and creates no synchronizes-with edge, so it cannot make surrounding writes visible to another thread. *Tempting wrong answer:* "relaxed still orders the writes right before it" — it orders nothing but the atomic's own value. **Feed-forward:** §24.3.

24.4 (a) Under SC the two stores take some global order; say `x = 1` is first. Then thread 1's later `r1 = x` must read 1 (the store is already globally visible), so `r1 == 0` is impossible — and symmetrically if `y = 1` is first, `r0 == 0` is impossible; hence both zero cannot happen. (b) On x86 each store enters a private store buffer and is not yet visible to the other core, so each thread's load reads the other variable's old value 0 before the store drains. (c) Make both stores (and ideally the loads) `seq_cst`; on x86 the `seq_cst` store compiles to a locked `xchg` that drains the store buffer, forbidding the store-load reorder. *Tempting wrong answer:* "making the loads acquire fixes it" — acquire/release does not forbid store-load reordering; only `seq_cst` (or an explicit full fence) does. **Feed-forward:** §24.1, §24.2, §24.4.

24.5 *Sequenced-before* is program order within a single thread. *Synchronizes-with* is a cross-thread edge, e.g. from a release-store to the acquire-load that reads its value. *Happens-before* is (roughly) the transitive closure of sequenced-before and synchronizes-with. Two conflicting accesses not ordered by happens-before are a data race because nothing establishes which is visible to the other — their relative order is undefined, so the standard declares it undefined behavior. *Tempting wrong answer:* "happens-before is just program order" — program order (sequenced-before) is only the within-thread part; the cross-thread edges come from synchronizes-with. **Feed-forward:** §24.2, §24.3.

24.6

```

producer:  payload = ...;          consumer:  while (load_acquire(&flag) != R) {}
           store_release(&flag, R);           use payload;

```

(a) The publishing store is `release`; the observing load is `acquire`. (b) When the `acquire`-load reads the value the `release`-store wrote, the store synchronizes-with the load, so everything sequenced-before the `release` (the payload writes) happens-before everything sequenced-after the `acquire` (the payload reads) — the payload is guaranteed visible. (c) With both `relaxed` there is no synchronizes-with and no happens-before, so the consumer may see the flag set but the payload stale (undefined in general); it may still pass on x86 because the hardware does not reorder store-store or load-load, so the writes happen to reach memory in order. *Method-selection cue*: "publish data behind a flag" → `release` on the publish, `acquire` on the observe. **Feed-forward**: §24.3, §24.4.

24.7 (a) x86-TSO never reorders store-store or load-load, so `acquire` (a load) and `release` (a store) need no barrier — they compile to plain `mov`, and only the compiler is restrained from reordering across them; `seq_cst` additionally needs store-load ordering, which the hardware *does* relax, so its store needs a real fence. (b) The `xchg` is a locked read-modify-write that drains the store buffer before proceeding, so the store is globally visible before the following load — forbidding the store-buffering outcome that a plain `mov` allows. (c) ARM/POWER reorder store-store and load-load, so `acquire` and `release` require explicit barrier instructions there; hence the weakest correct ordering matters for portable performance — an unnecessary `seq_cst` costs a fence on every architecture. *Tempting wrong answer*: "`seq_cst` is free like `acquire`/`release` because x86 is strong" — the `seq_cst` store still pays for the `xchg`. **Feed-forward**: §24.4.

24.8 Step 4: Make the pointer-publishing store `memory_order_release` and the reader's pointer load `memory_order_acquire`. Then, when the reader's `acquire`-load observes the published pointer, the `release`-store synchronizes-with it, so all the node's field writes (sequenced-before the `release` in the writer) happen-before the reader's field reads (sequenced-after the `acquire`) — the fields are guaranteed visible, no garbage. The bug hid on x86 because x86-TSO does not reorder store-store, so the field writes reached memory before the pointer even with `relaxed` ordering; ARM reorders store-store, exposing it. *Tempting wrong answer*: "add a sleep or a second read to let the writes catch up" — timing hacks do not create happens-before; only the `release`/`acquire` pairing does. **Feed-forward**: §24.3, §24.4.

24.9 (a) **relaxed** — a counter whose ordering relative to other data is irrelevant; only the final sum matters. (b) **release** on the publish and **acquire** on the consumer's read — it must make the buffer's writes visible. (c) **relaxed** — a stop flag whose late observation is harmless needs only atomicity, not ordering. (d) **seq_cst** — Dekker needs store-load ordering across the two flags, which only `seq_cst` (or a full fence) provides. *Method-selection cue*: no cross-variable ordering → `relaxed`; publish/observe data → `release`/`acquire`; need store-load / global order → `seq_cst`. **Feed-forward**: §24.3, §24.4.

24.10 (a) **False** — a data race is undefined behavior, not merely an unspecified value. (b) **True** — on x86 a release-store and a relaxed store both compile to a plain `mov`. (c) **False** — relaxed orders nothing but the atomic's own value; it does not make surrounding writes visible in order. (d) **True** — that is the DRF-SC guarantee. (e) **False** — `seq_cst`'s store costs a locked `xchg` on x86; acquire/release do not. *Tempting wrong answer*: calling (e) true because "x86 is a strong model" — strong enough to make acquire/release free, but the `seq_cst` store still needs a fence. **Feed-forward**: §24.2, §24.4.

24.11 (a) By promising SC only to race-free programs, the standard lets the compiler and CPU reorder and cache freely wherever no synchronization observes it, while a program that synchronizes correctly still sees the simple sequentially consistent model — the best of both. (b) "Undefined" (not "unspecified") lets the compiler assume races never occur and optimize on that basis (e.g. keeping a shared variable in a register across a supposedly racy read); under a weaker "unspecified value" definition the compiler would have to reload shared memory conservatively and could not make those transformations. (c) A data race is undefined behavior exactly like a buffer overflow or signed overflow (Chapter 10): the optimizer assumes it cannot happen, so a racy program may work in testing and break under a new compiler, optimization level, or CPU — "it worked when I tested it" proves nothing about a program with UB. *Tempting wrong answer*: "undefined is just pedantry; in practice you get a stale value" — the compiler is entitled to transformations that make the failure far worse than a stale value. **Feed-forward**: §24.2, and Ch 10.

24.12 Machine-specific by design. A correct write-up reports: (a) with relaxed atomics the store-buffering outcome `r0 == r1 == 0` occurs a nonzero number of times (on the book's machine ~349,000 of two million, about 1 in 5), and switching to `seq_cst` drops it to zero. (b) The disassembly shows relaxed store, release store, and acquire load as plain `mov`, while the `seq_cst` store is a locked `xchg` — the one that differs. (c) The release/acquire message-passing handoff verifies a large number of payloads (e.g. one million) with zero stale reads. Label compiler/CPU; the reordering counts are hardware-dependent. *Tempting wrong answer*: "zero SC violations means my relaxed publish is fine" — on x86 you may see zero even for buggy relaxed publishing; the counts are hardware-dependent, and ARM would differ. **Feed-forward**: §24.1, §24.4.

Chapter 25 — Solutions

25.1 (1) **Creation/teardown cost**: creating a thread makes the kernel allocate a stack and task structure and schedule it, and joining tears that down — hundreds of microseconds, dwarfing a small task. (2) **Oversubscription**: more runnable threads than cores forces the scheduler to time-slice them, paying context switches and cache disruption without adding parallelism. *Tempting wrong answer*: “threads are bad because of locks” — the cost here is creation and oversubscription, independent of any locking.

Feed-forward: §25.1.

25.2 It is the **producer-consumer** pattern (Chapter 21): producers enqueue tasks, worker-consumers dequeue and run them. The work queue uses a **mutex** (mutual exclusion over the shared queue's state) and a **condition variable** (workers block on it when the queue is empty and are signalled when a task is enqueued, so they wait without spinning). *Tempting wrong answer*: “it just needs a mutex” — without a condition variable an empty-queue worker either busy-waits or has no way to be woken. **Feed-forward**: §25.2.

25.3 A **future** is a handle to a value that may not be computed yet (returned immediately when a task is submitted). `.get()` on an unfinished task *blocks* until it completes and then returns the result; on an already-finished task it returns the result immediately. *Tempting wrong answer*: “a future is the result value” — it is a handle/placeholder for a result that may still be pending, not the value itself. **Feed-forward**: §25.3.

25.4 A pool creates its w workers once at startup and keeps them alive, so the ~ 220 us creation cost is paid w times total, not per task; every one of the N tasks reuses an already-running worker, so its cost is just enqueue + dequeue + run. In the measurement, thread-per-task cost 217 us/task while the pool cost 1.04 us/task — so on the order of 216 of the 217 us per task (well over 99%) was thread creation/teardown overhead, not useful work; the pool eliminates it by never creating a thread per task. *Tempting wrong answer*: “the pool is faster because it uses more threads” — it uses *fewer* (a fixed w); the win is amortized creation, not more threads. **Feed-forward**: §25.1, §25.2.

25.5 (a) The mutex protects only the queue; holding it while running `t.fn()` would let just one worker be inside the pool at a time, serializing all the work and destroying parallelism. (b) A `while` loop re-checks the predicate after waking because condition-variable waits can wake spuriously and because another worker may have taken the item first; an `if` would proceed on an empty queue (Chapter 21). (c) The worker *blocks* in `cond_wait` when the queue is empty, releasing the CPU until a signal arrives, so an idle pool consumes no CPU. *Tempting wrong answer*: “run under the lock so the task is safe” — the task's own data safety is its concern; the lock guards the queue, and holding it over the task serializes the pool. **Feed-forward**: §25.2.

25.6 Each worker owns a double-ended queue (**deque**). The owner pushes and pops from one end (the bottom) as a local, uncontended operation; a **thief**, when its own deque is empty, steals from the *other* end (the top) of a randomly chosen victim. Because owner and thieves touch opposite ends, the common local path needs no contention with thieves. Blumofe–Leiserson: a well-structured (fully-strict) computation with work T_1 and span T_∞ runs in expected $O(T_1/P + T_\infty)$ time on P processors; near-linear speedup needs enough parallelism, $T_1/T_\infty > P$. *Tempting wrong answer*: “a thief steals from the same end the owner uses” — that would contend on every local push/pop; thieves take the opposite (top) end. **Feed-forward**: §25.4.

25.7 (a) Every worker must hold the one global lock for time c to dequeue; the lock is a serial resource, so the pool cannot dequeue faster than one task per c . Once P is large enough that $P \cdot t$ of work demands dequeues faster than the lock can grant them, throughput saturates at roughly $1/c$ tasks per unit time — the shared queue is the bottleneck (an Amdahl serial fraction, next chapter). (b) Per-worker dequeues let each worker enqueue/dequeue locally with no shared lock; stealing is rare, so the serial bottleneck largely disappears. (c) Coarser tasks mean more useful work t per dequeue c , so the lock is taken less often relative to work done, improving scaling; too coarse and you have too few tasks to balance load and fill all workers (poor parallelism, idle cores). *Tempting wrong answer*: “just add more workers” — past saturation, more workers only add lock contention. **Feed-forward**: §25.2, §25.4.

25.8 Step 4: This is **pool deadlock** (thread-pool starvation). It is a genuine deadlock: the 4 worker threads are the only resource that can run tasks, and all 4 are *held* by request tasks that are *waiting* on futures for subtasks; those subtasks need a free worker to run, but none exists — a circular wait between “workers held by waiters” and “subtasks needing a worker.” Fix 1 (structure): run blocking/dependent work in a *separate* pool from the CPU work, or use a work-stealing fork-join runtime whose join executes the pending subtask on the current thread instead of blocking. Fix 2 (expression): don't have a worker block on a future for same-pool work — compose dependencies with continuations/callbacks the runtime can schedule, so no worker sits idle holding a thread. *Tempting wrong answer*: “just make the pool a bit bigger” — any fixed size w deadlocks once w tasks block on un-started subtasks; size only shifts the load at which it hangs. **Feed-forward**: §25.4.

25.9 (a) **Right** — CPU-bound non-blocking tasks want about one worker per hardware thread; more just oversubscribes. (b) **Change it** — I/O-bound tasks spend most time blocked, so a larger pool (or async I/O) keeps cores busy while threads wait. (c) **Change it** — millions of tiny recursive subtasks want a work-stealing fork-join pool whose join doesn't block a worker, not a plain shared-queue pool. (d) **Change it** — separate the pools: a hardware-sized pool for the short CPU tasks and a distinct pool (or async) for the long blocking DB calls, so blocking work can't starve CPU work. *Method-selection cue*: CPU-bound $\rightarrow$ $\sim$ hardware-thread count; blocking $\rightarrow$ larger/separate pool or async; recursive fork-join $\rightarrow$ work-stealing runtime. **Feed-forward**: §25.2, §25.4.

25.10 (a) **False** — parallelism is capped at the number of cores; extra threads oversubscribe. (b) **True** — the work queue is a producer-consumer buffer under a mutex + condition variable. (c) **False** — an idle pool blocks in `cond_wait` and uses no CPU. (d) **True** — work-stealing replaces the one shared queue with per-worker dequeues (plus occasional stealing). (e) **False** — `get()` blocks only until *that task* finishes (or returns immediately if it already has), not until the whole pool is idle. *Tempting wrong answer:* marking (a) true because “each task has its own thread” — having a thread is not having a core to run it on. **Feed-forward:** §25.1–§25.4.

25.11 (a) The owner pushes/pops the bottom, so a thief taking from the top rarely touches the same slot — low contention — and the top holds the *oldest* task, which in a fork-join tree is the highest, coarsest-grained subtree; stealing a large task means one steal hands the thief a lot of independent work, so steals are infrequent and load balances well. (b) In $O(T_1/P + T_\infty)$, T_1/P is the total work spread perfectly over P processors, and T_∞ is the span — the critical path that must run sequentially no matter how many processors you have. If T_∞ is close to T_1 (little parallelism), the T_∞ term dominates and adding processors barely helps — the same wall Amdahl's law describes with a serial fraction. *Tempting wrong answer:* “more processors always cuts the time by that factor” — only the T_1/P term shrinks; the span T_∞ is a floor. **Feed-forward:** §25.4, and Ch 26.

25.12 Machine-specific by design. A correct write-up reports: (a) a per-thread `create+join` cost (on the book's WSL2 machine ~220 us; bare-metal Linux is typically far lower — label your environment). (b) The pool runs the same tiny tasks at a small fraction of the thread-per-task time (here 1.04 us/task vs 217 us/task, ~200×), because creation is amortized. (c) A fork-join sum showing near-linear speedup at small P (here ~1.98× at 2) that bends below linear as P grows (here ~4.38× at 8). Warm up and average; absolute numbers are hardware/OS dependent. *Tempting wrong answer:* “my pool is broken because 8 threads didn't give 8×” — sub-linear scaling is expected and is exactly Chapter 26's subject, not a bug. **Feed-forward:** §25.1–§25.4.

Chapter 26 — Solutions

26.1 $\text{speedup}(P) = 1/(s + (1-s)/P)$. As $P \rightarrow \infty$ the $(1-s)/P$ term vanishes and speedup approaches $1/s$ — the ceiling set by the serial fraction. *Tempting wrong answer:* “speedup grows linearly with P ” — only the parallel part divides by P ; the serial part is a floor.

Feed-forward: §26.1.

26.2 False sharing is when two threads write independent variables that happen to occupy the same cache line, so each write invalidates the other's cached copy and the line ping-pongs between cores — paying the cost of sharing while sharing no data. The granularity is the **cache line**, 64 bytes on a typical x86 machine. *Tempting wrong answer:* “it's caused by sharing a variable” — no variable is shared; the collision is at cache-line granularity, not variable granularity. **Feed-forward:** §26.4.

26.3 A lock's critical section runs one thread at a time, so if every thread must pass through the same lock on the hot path, that section executes serially — it is a serial fraction, and Amdahl caps speedup at one over it. *Tempting wrong answer:* “the lock only adds a small constant overhead” — it serializes an entire fraction of the work, and Amdahl's ceiling applies to that fraction. **Feed-forward:** §26.3.

26.4 $s = 0.20$. (a) Limit = $1/0.20 = 5\times$ with infinite cores. (b) At $P = 4$: $1/(0.2 + 0.8/4) = 1/0.4 = 2.5\times$; at $P = 16$: $1/(0.2 + 0.8/16) = 1/0.25 = 4.0\times$. (c) Quadrupling cores (4→16) improved speedup only from $2.5\times$ to $4.0\times$ — a factor of 1.6, not 4 — so “just add cores” gives sharply diminishing returns once the serial fraction dominates. *Tempting wrong answer:* “16 cores gives $4\times$ the speedup of 4 cores” — it gives $1.6\times$; the serial 20% throttles it. **Feed-forward:** §26.1.

26.5 (a) The entire work is inside the one critical section, so it is effectively 100% serial — Amdahl ceiling $1\times$, meaning even in the best case adding threads cannot speed it up. (b) On top of that, the lock's own **cache line** (and the contended counter's line) must be transferred to whichever core acquires it next, so it bounces between cores; that coherence traffic and the wait/hand-off overhead grow with the number of threads, pushing the time *above* the single-thread time — negative scaling. (c) Fix 1: shard into per-thread counters combined once at the end — serial fraction drops toward 0. Fix 2: a lock-free atomic or shrinking the critical section — reduces the serialized portion and the contended line. *Tempting wrong answer:* “it's slow because the mutex syscall is expensive” — uncontended mutexes are cheap userspace operations; the cost here is serialization plus cache-line bouncing under contention. **Feed-forward:** §26.3.

26.6 (a) Version B (padded) is faster — about **10×** in the §26.4 measurement (2.56 s padded vs 0.25 s padded). In version A both counters share a cache line, so each thread's write invalidates the line in the other core, forcing a re-fetch every increment (the line ping-pongs); in B each counter owns its line, so writes stay core-local. (b) Neither has a data race: each thread writes only its own counter, and no location is accessed by two threads — the collision is at the cache-line level, not the variable level, so there is no shared variable at all. (c) Fix:

```
struct { long v; char pad[64 - sizeof(long)]; } c[2]; /* one counter per line */
/* or */ _Alignas(64) long c0; _Alignas(64) long c1; /* C11 */
/* or GCC: */ long c[2] __attribute__((aligned(64))); /* + stride to a line each */
```

Method-selection cue: hot per-thread data written in a tight loop → give each thread's datum its own cache line. **Feed-forward:** §26.4.

26.7 (a) Efficiency = speedup/**P** = $1/(\mathbf{P} \cdot \mathbf{s} + (1 - \mathbf{s}))$. As **P** grows the **P**·**s** term grows, so efficiency falls monotonically toward 0 for any **s** > 0 — you keep more cores idle relative to the serial floor. (b) When efficiency has dropped so far that the marginal core does less useful work than it costs (money, power, coordination), more cores stop being cost-effective — well before the 1/**s** ceiling. (c) Under Gustafson the problem grows with **P**, so the parallel work per core stays roughly constant and the serial fraction of the *larger* job shrinks; efficiency can stay high because you are not dividing a fixed parallel part into ever-thinner slices. *Tempting wrong answer:* “efficiency is constant if the program is well written” — for a fixed problem with any serial part, efficiency always falls with **P**; only weak scaling avoids it. **Feed-forward:** §26.1, §26.2.

26.8 Step 4: This is **false sharing**. It gets worse with more threads because each additional core writing its slot adds another participant fighting over the same 64-byte line — more invalidations and cross-core transfers per unit time, so coherence traffic (not useful work) grows with the thread count. The fix is to give each thread's counter its own cache line: pad the element to 64 bytes (`struct { long v; char pad[56]; } counts[N];`) or align each counter with `_Alignas(64) / __attribute__((aligned(64)))` so no two threads' hot counters share a line. The §26.4 measurement showed this is worth roughly a **10×** speedup (2.56 s → 0.25 s) for the same arithmetic. *Tempting wrong answer:* “add a lock around each slot to make it thread-safe” — it is already correct; a lock would add real contention on top of the false sharing and make it worse. **Feed-forward:** §26.4.

26.9 (a) **Serial fraction (Amdahl)** — a plateau at $\sim 10\times$ with no lock implies $\sim 10\%$ inherently serial work; confirm by measuring the sequential portion's share of runtime ($\sim 1/10$). (b) **Lock contention** — throughput dropping with a profiler pointing at `pthread_mutex_lock` is the signature; confirm with the lock's contention/wait counters. (c) **False sharing** — per-thread array, no locks, poor scaling; confirm by padding each element to its own cache line and seeing the slowdown vanish (or via a perf counter for cache-coherence/HITM events). (d) **Serial fraction (Amdahl)** — a fixed 2 s single-threaded merge is a literal serial section; confirm it stays constant as the parallel phase's core count changes. *Method-selection cue*: plateau & no lock $\rightarrow$ Amdahl; time in lock $\rightarrow$ contention; per-thread array & no lock $\rightarrow$ false sharing. **Feed-forward**: §26.1, §26.3, §26.4.

26.10 (a) **False** — a fixed problem is capped at $1/s$ by its serial fraction, no matter the core count. (b) **False** — the laws answer different questions (fixed problem vs problem grown with the machine); they are consistent. (c) **False** — false sharing is a performance bug, not a data race; each thread writes only its own variable, so behavior is well-defined. (d) **True** — a contended hot lock serializes work and bounces a cache line, so adding threads can slow the program (negative scaling). (e) **False** — padding changes only memory layout and timing, never the computed results. *Tempting wrong answer*: calling (c) true because “threads touch nearby memory” — touching nearby memory is not a race; a race needs conflicting access to the *same* location. **Feed-forward**: §26.1–§26.4.

26.11 (a) Fix the total work $W = \text{serial} + \text{parallel}$ with serial share s ; time on P is $s \cdot W + (1-s)W/P$, so speedup $= 1/(s + (1-s)/P) \rightarrow 1/s$ — Amdahl. If instead the parallel work grows with P (fixed wall-clock), the work done is $s + (1-s)P$ relative to one processor's serial-plus-one-share, giving scaled speedup $P - s(P-1)$ — Gustafson, near-linear. (b) “Linear to 1000 cores” is plausible under Gustafson if they grew the simulation with the machine (weak scaling); it does not contradict Amdahl, which governs a *fixed*-size run. Ask: “Did the problem size stay fixed as you added cores, or grow with them?” (c) The span T_∞ is the critical path that must run sequentially no matter how many processors you have — exactly the role of Amdahl's serial fraction; in $O(T_1/P + T_\infty)$ the span is the irreducible floor. *Tempting wrong answer*: “the $1000\times$ claim disproves Amdahl” — it measures weak scaling, a different regime, not a violation. **Feed-forward**: §26.1, §26.2, and Ch 25.

26.12 Machine-specific by design. A correct write-up reports: (a) measured speedups that track Amdahl's $1/(s + (1-s)/P)$ for the measured s and bend toward the $1/s$ ceiling (on the book's machine $s \approx 0.089$, ceiling $\sim 11\times$, 16 threads $\approx 6.4\times$). (b) The single-mutex counter showing flat or negative scaling (here 1.80 s at 1 thread rising to 7.06 s at 8). (c) A packed-vs-padded false-sharing slowdown of several times (here $\sim 10\times$). Label CPU and cache-line size; numbers are hardware-dependent. *Tempting wrong answer*: “my parallel loop is buggy because 8 threads gave $4\times$ ” — sub-linear scaling is expected from the serial fraction and memory effects, not necessarily a bug. **Feed-forward**: §26.1, §26.3, §26.4.

Chapter 27 — Solutions

27.1 (1) Every value has an owner. (2) There is exactly one owner at a time. (3) When the owner goes out of scope, the value is dropped (its resources freed). *Tempting wrong answer*: “a value can have several owners as long as they cooperate” — rule 2 forbids it; shared access comes later via borrowing, not multiple ownership. **Feed-forward**: §27.2.

27.2 “Dropped” means Rust automatically runs the value's cleanup at the point its owner's scope ends; for a heap-owning value it corresponds to calling `free` on the buffer in C — but inserted by the compiler, deterministically, not by hand. *Tempting wrong answer*: “drop is like a garbage collector reclaiming it later” — drop happens at a fixed, known point (end of scope), not whenever a collector runs. **Feed-forward**: §27.2.

27.3 No — `s1` has been **moved from** and is invalid; the compiler rejects any use of it with a compile error (`error[E0382]`, borrow/use of moved value). *Tempting wrong answer*: “`s1` is still usable because assignment copies” — for an owning type assignment moves, not copies. **Feed-forward**: §27.3.

27.4 (a) After `char *t = s;` and `free(s);` the `printf("%s", t)` is a **use-after-free** (reading a freed buffer through the alias), and `free(t)` is a **double-free** (the same block returned twice). (b) It compiles because C's type system does not track ownership or validity — a `char*` carries no information about whether its target is still allocated or who may free it. (c) In Rust the analogous `let t = s;` moves ownership from `s` to `t`, so `s` immediately becomes invalid; there is never a state with two live owners, so no second owner exists to free the buffer again — the double-free is unrepresentable. *Tempting wrong answer*: “Rust catches it by inserting a runtime check before the second free” — it is a compile-time rejection; there is no second free to check. **Feed-forward**: §27.1, §27.3.

27.5 Rule 2 (one owner) means only one binding is ever responsible for the buffer, so it is freed exactly once — preventing the double-free. Rule 3 (drop at end of scope) means that single owner's cleanup runs automatically when its scope ends, so the free is never forgotten — preventing the leak. Rule 1 ties every value to an owner so rules 2 and 3 always apply. No GC is needed because the drop point is known at compile time and inserted there. *Tempting wrong answer*: “you still need to remember to drop it” — drop is automatic at scope exit; forgetting is not possible for an owned value. **Feed-forward**: §27.2.

27.6 (a) **Move**: ownership transfers, source becomes invalid, no heap buffer duplicated (just the small stack record is bit-copied and the source invalidated). **clone**: allocates and deep-copies the heap buffer, source stays valid, two independent owners of two buffers. **Copy**: bit-copy of a stack-only value, source stays valid, no heap involved. (b) `String` owns a heap buffer, so a copy would create two owners of one buffer (unsafe) — hence it moves; `i32` is entirely on the stack with no resource to double-free, so it implements `Copy` and is duplicated harmlessly. (c) Reaching for `.clone()` to silence a move error is wrong when you only need to *read* the value — you should borrow a reference instead of paying for a deep copy. *Method-selection cue*: owns a resource → moves; plain stack data → `Copy`; need a real second copy → `clone`; only need to look → borrow. **Feed-forward**: §27.3.

27.7 (a) The move/drop machinery is resolved at **compile time** — the compiler decides where drops go and rejects use-after-move statically; the ownership *checking* adds no run-time cost. (b) drop-at-end-of-scope is **deterministic** — memory is freed at a known program point — whereas a tracing GC reclaims at unpredictable times; determinism means no GC pauses and tighter, more predictable peak memory. (c) A move only bit-copies a small stack record and invalidates the source (essentially free); `.clone()` allocates a new buffer and copies the contents (real time and memory proportional to the data). *Tempting wrong answer*: “ownership adds runtime overhead like reference counting” — plain ownership is entirely compile-time; ref-counting (`Rc/Arc`) is a separate, opt-in mechanism. **Feed-forward**: §27.2, §27.3.

27.8 Step 4: Fix A (no clone): have `process` take a *reference* — `process(&cfg)` — so it borrows rather than moves; `cfg` stays owned by the caller and `log(&cfg)` (or `log(cfg)` last) works. This is right whenever the functions only need to read/inspect the value. Fix B (clone): if `process` genuinely needs to *take ownership* (e.g. to store or consume `cfg`) and you still need it afterward, pass `process(cfg.clone())` and keep the original for `log`. Decide by asking whether `process` must own the value: if not, borrow (Fix A); if it must and you also need the original, clone (Fix B). Blindly cloning is the wrong default because it silently deep-copies data you may only be reading, discarding performance for no reason. *Tempting wrong answer*: “always clone — it's simplest” — it builds needless copies into hot paths; borrowing is the usual right answer. **Feed-forward**: §27.3, and Ch 28 (borrowing).

27.9 (a) **Move** — `String` owns a heap buffer and is not `Copy`; `a` becomes invalid. (b) **Copy** — `i32` implements `Copy`; `a` stays usable. (c) **Move** — `Vec<u8>` owns a heap buffer; passing by value moves it and the caller can no longer use it. (d) **Copy** — a tuple of `Copy` types (`i32, bool`) is itself `Copy`, so it copies and `a` stays usable. *Method-selection cue*: owns heap → move; all-stack `Copy` types → `copy`. **Feed-forward**: §27.3.

27.10 (a) **False** — Rust has no garbage collector; safety comes from compile-time ownership. (b) **False** — `s1` is moved from and invalid; only `s2` is valid. (c) **True** — move errors are compile-time and add no runtime cost. (d) **False** — `clone` deep-copies into a new buffer (both valid); a move transfers ownership and invalidates the source (no buffer copied). (e) **False** — ownership controls aliasing, which is also the basis for preventing data races. *Tempting wrong answer*: marking (a) true because “safe languages use GC” — Rust is the counterexample; it is safe without one. **Feed-forward**: §27.1–§27.4.

27.11 (a) Moving a value into a thread transfers ownership to that thread; by rule 2 the original thread is no longer an owner and cannot touch the value, so two threads cannot both own-and-mutate it — the accidental sharing that causes races is structurally excluded. (b) A use-after-free is one alias freeing/mutating memory another alias still uses; a data race is two threads accessing the same location with at least one write, unsynchronized — both are “two paths to the same data, at least one mutating.”

Ownership attacks that shared root by ensuring a single responsible owner, rather than patching use-after-free and races as separate problems. (c) Ownership-by-move alone won't let you *use* a value without giving it away, so you can't, e.g., let two functions read the same buffer without moving/cloning; you need **borrowing** (references), which must guarantee that references never outlive their value and that you never have a mutable reference aliased with any other reference — keeping aliasing safe without transferring ownership. *Tempting wrong answer:* “ownership by itself is enough to write real programs” — without borrowing you'd move or clone constantly; borrowing is what makes it practical. **Feed-forward:** §27.4, and Ch 24, Ch 28.

27.12 Toolchain-specific. A correct write-up reports: (a) the C snippet compiles with no warning and, at run time, aborts on the double-free (glibc “double free detected”, SIGABRT) or, under `-fsanitize=address`, is flagged as a heap-use-after-free — demonstrating the C type system misses it. (b) With `rustc`, the `let s2 = s1; println!("{}", s1);` program fails to compile with error[E0382] on the `println!` line (use of moved value); without `rustc`, name that exact line and explain the move at `let s2 = s1;` invalidated `s1`. (c) For three variables, correctly label each assignment move or copy from its type (owning type → move; Copy stack type → copy). *Tempting wrong answer:* “the C program is fine because it compiled” — compiling proves nothing about UB; it aborts (or corrupts) at run time. **Feed-forward:** §27.1, §27.3.

Chapter 28 — Solutions

28.1 A **shared reference** `&T` is a read-only borrow; any number may exist at once, and you may read but not mutate the value through one. A **mutable reference** `&mut T` is a read-write borrow; exactly one may exist, and while it lives no other reference (of either kind) to the value may. *Tempting wrong answer*: “`&mut` is just a pointer, like `&` but you can write” — it is that plus an *exclusivity* guarantee the compiler enforces. **Feed-forward**: §28.2, §28.3.

28.2 `&String` borrows: passing `&owner` lends a reference without moving the value, so the caller keeps ownership and `owner` stays valid. `fn length(s: String)` takes by value, which *moves* the `String` into the function and invalidates `owner`. *Tempting wrong answer*: “both work because the function only reads” — reading is irrelevant; taking by value moves regardless of what the body does. **Feed-forward**: §28.1.

28.3 The rule: at any time a value may be borrowed by *any number of shared references* (`&T`) or by *exactly one mutable reference* (`&mut T`), never both. It forbids two mutable borrows at once (E0499) and a mutable borrow while a shared borrow is live (E0502). *Tempting wrong answer*: “you can mix one `&mut` with a few `&` if they don't overlap in code” — if they are simultaneously live, that is exactly E0502. **Feed-forward**: §28.3.

28.4 (a) The bug is the `printf("via p", *p)` line: after `realloc` may have moved the block, `p` is a dangling pointer, so the dereference is a use-after-free. (b) Plain `gcc -O2 -Wall` gives no warning because C's type system does not track whether `p` still points at live memory after `realloc`; only AddressSanitizer, at run time, reports **heap-use-after-free**. (c) In Rust `p = &v[0]` is a shared borrow of `v`, and `v.push(4)` needs `&mut v`; the mutable borrow cannot coexist with the live shared borrow, so it is E0502 at compile time. The checker does not need to know `push` reallocates — the rule “no mutable borrow while a shared borrow is live” is sufficient. *Tempting wrong answer*: “Rust catches it because it checks the reallocation at run time” — it is a static rejection with no runtime check. **Feed-forward**: §28.4.

28.5 (a) Two references to one value with at least one mutable would let one path write while another reads or writes — a torn or unexpectedly-changed read, or a pointer invalidated under the other, all in one thread. (b) That is the definition of a data race (Chapter 24) minus the threads: two accesses to one location, at least one a write, unsynchronized. (c) Exclusivity lets the compiler assume nothing else can see or change the value through the `&mut`, which makes mutation safe and enables optimizations (no aliasing to worry about). *Tempting wrong answer*: “exclusivity is just a style rule” — it is a soundness requirement; without it Rust could not rule out aliasing-plus-mutation. **Feed-forward**: §28.3, §28.4.

28.6 (a) Compiles — many shared references are allowed. (b) E0499 — two mutable borrows at once. (c) E0502 — mutable borrow while a shared borrow is live. (d) Compiles — because the shared borrow has ended (last use) before the mutable borrow starts, they do not overlap (non-lexical lifetimes). *Method-selection cue*: count simultaneously-live borrows; many & ok, one &mut alone ok, any overlap of a writer with anything else is rejected. **Feed-forward**: §28.2, §28.3.

28.7 (a) A &T is, at the machine level, just a pointer — the same pointer C would use — carrying no extra runtime data (for a sized T). (b) The borrow checking happens entirely at compile time; it adds zero runtime cost. (c) RefCell enforces the same one-writer-XOR-many-readers rule at run time, so it carries a borrow-flag it checks and updates on each borrow (a small time cost and a possible panic on violation); paying it is justified only when the borrow structure cannot be proven statically but you still need shared mutability. *Tempting wrong answer*: “references are cheap because they're reference-counted” — plain references are not counted at all; that is Rc/Arc, a separate mechanism. **Feed-forward**: §28.2, §28.3.

28.8 Step 4: Fix A (separate the passes): during a read-only pass, collect what you want to add into a new vector (`let additions: Vec<i32> = v.iter().map(|x| 2 * x).collect();`), then, after that shared borrow ends, `v.extend(additions)` under a single mutable borrow. Fix B (in-place, if that were the task): use `for x in v.iter_mut() { *x *= 2; }`, which is one exclusive borrow and is allowed because there is no aliasing. Cloning the entire vector to iterate a copy while mutating the original is the wrong default because it pays a full O(n) copy to hide a real hazard; separating read from write costs nothing extra and is what the rule is pointing at. *Tempting wrong answer*: “just `v.clone()` and iterate that” — needless copy, and it silently changes semantics (you iterate the old contents). **Feed-forward**: §28.4.

28.9 (a) Allowed — many shared references may coexist (no writer). (b) Forbidden — a mutable and a shared borrow simultaneously is E0502. (c) Forbidden — two mutable borrows is E0499. (d) Allowed — the shared borrows have ended, so the mutable borrow is the only live one. *Method-selection cue*: the only legal states are “N readers, 0 writers” or “0 readers, 1 writer.” **Feed-forward**: §28.3.

28.10 (a) **False** — at most one &mut may exist at a time, regardless of timing (E0499). (b) **False** — a shared reference is read-only; mutation needs &mut. (c) **False** — borrow checking is compile-time; a dereference is just a pointer load. (d) **True** — that is the iterator-invalidation hazard; it is E0502. (e) **False** — a reference does not own the value and drops nothing; only the owner drops it. *Tempting wrong answer*: marking (c) true because “safety must cost something at run time” — the borrow rule's cost is paid entirely at compile time. **Feed-forward**: §28.2–§28.4.

28.11 (a) A use-after-free through an aliased pointer is one alias mutating/freeing memory another alias still reads; a data race is two threads accessing one location, one writing, unsynchronized — both are “two paths to one value, at least one writing.” The single rule “aliasing XOR mutation” forbids that combination, attacking the shared root rather than the two symptoms. (b) Unlimited shared references are safe because with no writer no reader can observe a changing value — reads never conflict with reads. Access becomes a problem only when a write is introduced, which the rule permits solely under exclusivity. (c) To extend “no data races” across threads you need the Send/Sync traits (and, for shared mutation, synchronization such as `Arc<Mutex<T>>`): the borrow rule governs one value in one thread, but sending references between threads requires the type system to also track which types are safe to share or move across threads. *Tempting wrong answer*: “the borrow rule alone already prevents all data races” — it prevents intra-thread aliasing; cross-thread safety needs Send/Sync on top. **Feed-forward**: §28.4, and Ch 24; the concurrency chapter ahead.

28.12 Toolchain-specific. A correct write-up reports: (a) the C snippet builds with no warning under `-O2 -Wall` and, rebuilt with `-fsanitize=address`, reports **heap-use-after-free** when the old pointer is dereferenced after a moving `realloc` — the C type system misses the dangling alias. (b) With `rustc`, the `let p = &v[0]; v.push(4); println!("{}", p)` program fails with `error[E0502]` because `push`'s mutable borrow conflicts with the live shared borrow `p`; without `rustc`, name those two conflicting borrows. (c) The compiling version ends the shared borrow (use `p` before the `push`, or index `v[0]` directly and copy the value out) so no borrow overlaps the mutable one. *Tempting wrong answer*: “the C program is fine because it compiled without warning” — no warning proves nothing about UB; the sanitizer shows the use-after-free. **Feed-forward**: §28.4.

Chapter 29 — Solutions

29.1 A **lifetime** is a compile-time name for the region of the program over which a reference is valid (bounded by its referent's scope). The second borrow rule it enforces: a reference may never be used after the value it refers to has been dropped. *Tempting wrong answer:* “a lifetime is how long the value lives, which you can set” — it only *describes* existing scope facts; it cannot make anything live longer. **Feed-forward:** §29.1, §29.3.

29.2 `x` is a stack local dropped when the function returns, so the returned `&x` would point at reclaimed storage — a dangling reference / use-after-scope. *Tempting wrong answer:* “it's fine because 5 is a small constant” — the value's size is irrelevant; its storage is gone once the frame is popped. **Feed-forward:** §29.1.

29.3 The three `'a`s assert: there is some region `'a` such that both inputs `x` and `y` are valid for at least `'a`, and the returned reference is valid for `'a` — so the result may not outlive the shorter-lived of the two inputs. *Tempting wrong answer:* “`'a` makes both inputs live equally long” — it constrains the output's validity relative to the inputs; it does not change how long anything lives. **Feed-forward:** §29.2.

29.4 (a) Both return (directly or through a struct field) the address of a stack local whose storage is reclaimed when the function returns, so dereferencing it in the caller is a use-after-scope. (b) gcc's `-Wreturn-local-addr` follows simple direct data flow and catches `return &local;`, but gives up when the address is stored in a struct and returned indirectly, so that version builds clean. (c) Rust's rule is not a heuristic — the borrow checker rejects *every* reference used past its referent's drop; the analogous Rust function is `error[E0597]`, “does not live long enough.” *Tempting wrong answer:* “the struct version is safe because it compiled without warning” — ASan shows stack-use-after-scope; no warning proves nothing. **Feed-forward:** §29.1, §29.3.

29.5 (a) The containment condition: for every use of a reference, the region over which the reference is used must be contained within the region its referent is alive. (b) You can't fix it by annotating a longer lifetime because a lifetime only describes existing scope facts — no name can extend a value past its owner's drop. (c) The two structural fixes: return an *owned* value (move ownership out — right when the caller should own the result), or borrow from something that genuinely outlives the use (an input reference or an outer-scope value — right when the data lives elsewhere). *Tempting wrong answer:* “annotate it `'static`” — that only compiles if the data really is static, which a fresh local isn't. **Feed-forward:** §29.3, §29.4.

29.6 The rules: (1) each elided input reference gets its own lifetime; (2) one input lifetime is assigned to all elided outputs; (3) if one input is `&self`/`&mut self`, its lifetime is assigned to all elided outputs. (a) `first` has one input reference, so rule 2 assigns its lifetime to the output — no annotation needed. (b) `get` is a method with `&self`, so rule 3 gives the output `self`'s lifetime. (c) `longest` has two input references and neither is `self`, so rules 2 and 3 don't apply and the compiler can't tell which input the output borrows — you must write `'a`. *Method-selection cue*: one input ref → elided; method with `self` → elided; multiple non-self input refs feeding an output ref → annotate. **Feed-forward**: §29.2, §29.4.

29.7 (a) Lifetimes are checked during borrow-checking (compile time) and are *erased* before code generation. (b) A `&'a T` has the same runtime representation as the C pointer it replaces — a bare address, no extra data (for sized `T`). (c) A GC answers “still valid?” by tracing/reachability at run time (time cost, pauses); reference counting answers it by per-reference counter updates at run time (time and space); Rust answers it once at compile time, paying nothing at run time. *Tempting wrong answer*: “lifetimes are a runtime tag on each reference” — they are purely static and erased. **Feed-forward**: §29.3.

29.8 Step 4: Correct fix — change the return type to an owned `String` and return `s` by value: `fn greeting() -> String { String::from("hi") }`; ownership moves out to the caller, so there is no dangling reference. Alternative — if the caller already owns suitable data, take it as an input reference and return a slice of that: `fn greeting(src: &str) -> &str`, borrowing from something that outlives the call. Reaching for `'static` is wrong because a freshly built `String` is not static data; `'static` would either fail to compile or force an unintended constraint. *Tempting wrong answer*: “annotate the return `&'static str`” — the local isn't static, so this doesn't fix the dangle. **Feed-forward**: §29.3, and Ch 27 (moving ownership out).

29.9 (a) **Elided** — one input reference; rule 2 assigns its lifetime to the output. (b) **Required** — a struct with a reference field must name its lifetime (`&'a u8`), else E0106. (c) **Dangling-reference error** — returning a reference to a local is E0597. (d) **Required** — two non-self input references feeding an output reference; elision can't choose, so annotate `'a`. *Method-selection cue*: reference field → annotate; return-a-local → error; multiple input refs to one output ref → annotate; single input ref → elided. **Feed-forward**: §29.2, §29.4.

29.10 (a) **False** — `'static` asserts the referent already lives for the whole program; it does not make an arbitrary value live forever. (b) **False** — lifetimes are erased before code generation; a reference is a bare pointer with no extra data. (c) **False** — a struct with a `&str` field must name a lifetime, else E0106. (d) **True** — E0597 fires precisely when a reference's use would extend past its referent's drop. (e) **False** — elision handles most signatures; explicit lifetimes are the exception. *Tempting wrong answer*: marking (b) true because “the compiler must store the lifetime somewhere” — it stores it only during checking, then erases it. **Feed-forward**: §29.3, §29.4.

29.11 (a) A struct storing a `String` owns its data — it costs an allocation but carries no lifetime parameter and may be used freely. A struct storing `&'a str` avoids the allocation but gains a lifetime parameter and may not outlive the `str` it borrows — every use is constrained by that borrow. (b) A self-referential struct is hard because the borrowing field's lifetime would have to name the struct's own storage, which moves (Chapter 27) — after a move the internal pointer would dangle, so plain lifetimes cannot express it safely (it requires pinning or indirection). (c) The own-vs-borrow choice is exactly C's “who owns this buffer, and who may free it” question (Chapters 9, 27); lifetimes make that question explicit in the type and let the compiler check the answer, instead of leaving it to convention. *Tempting wrong answer:* “always borrow to avoid allocations” — borrowing threads lifetimes through everything and forbids outliving the source; owning is often the simpler, correct choice. **Feed-forward:** §29.4, and Ch 27.

29.12 Toolchain-specific. A correct write-up reports: (a) under `-O2 -Wall` the direct `return &local;` warns with `-Wreturn-local-addr` while the struct-laundered version builds without warning, and `-fsanitize=address` flags the latter as **stack-use-after-scope** at run time. (b) With `rustc`, `fn dangle() -> &str { let s = String::from("hi"); &s }` fails with `error[E0597], “s does not live long enough”`; without `rustc`, name the `&s` return line and explain `s` is dropped at the closing brace. (c) The longest function needs `<'a>` because of two input references; a struct with a `&str` field needs `<'a>` too, and omitting it is `E0106`, “missing lifetime specifier.” *Tempting wrong answer:* “the struct C version is fine — no warning” — ASan proves the use-after-scope. **Feed-forward:** §29.1, §29.4.

Chapter 30 — Solutions

30.1 A **generic function** is written once over a type parameter `T` and reused at many concrete types. A **trait** names a set of behaviors a type can implement. A **trait bound** (`T: Trait`) requires a generic's type parameter to implement a trait, and in return lets the body use that trait's operations. *Tempting wrong answer*: “a trait is a base class you inherit from” — a trait is a set of required behaviors a type opts into, not an inheritance hierarchy. **Feed-forward**: §30.2, §30.3.

30.2 Monomorphization is the compiler emitting a separate specialized copy of a generic for each concrete type it is used with. After it, there is no type parameter at run time — each call goes to an ordinary type-specific function with direct, inlinable operations — so there is no runtime dispatch. *Tempting wrong answer*: “generics carry a type tag checked at run time” — the type is resolved and erased at compile time. **Feed-forward**: §30.2.

30.3 `error[E0277]` — “the trait bound `T: Trait` is not satisfied” — meaning you used a type that doesn't implement a trait in a place that required it. *Tempting wrong answer*: “it's a type-mismatch error, so cast the type” — the type isn't mismatched; it lacks the required behavior, so you implement the trait or relax the bound. **Feed-forward**: §30.3.

30.4 (a) `qsort` is generic in that it sorts arrays of any element type via `void*`; its two costs are lost type safety (element type erased to `void*`) and an indirect call through the comparator function pointer per comparison. (b) Passing, say, a `double` comparator while sorting `ints` compiles because both are `void*` to `qsort`; at run time the comparator reinterprets the bytes wrongly and produces a garbage ordering, and the compiler cannot catch it because the types were erased. (c) The `qsort` function-pointer mechanism corresponds to a Rust **trait object** (dynamic dispatch); **generics with monomorphization** give the duplication-free, fully-checked, fully-inlined alternative. *Tempting wrong answer*: “`qsort` is type-safe because you cast inside the comparator” — the cast is unchecked; a wrong-typed comparator still compiles. **Feed-forward**: §30.1, §30.4.

30.5 (a) The caller must supply a type `T` that implements `Summary`. (b) The body may call `summarize` (and any other `Summary` method) on `item` — impossible without the bound, since otherwise the compiler couldn't know `T` has it. (c) Rust checks the bound at the generic's *definition*, so the generic is verified once against its stated requirements; C++ checks at instantiation, so an unmet requirement surfaces as errors deep inside the template body, per call site. *Tempting wrong answer*: “the bound is just documentation” — it is enforced; violating it is `E0277`. **Feed-forward**: §30.3.

30.6 (a) Static dispatch (generics) resolves at compile time — a direct, inlinable call, zero dispatch cost; dynamic dispatch (trait objects) resolves at run time through a vtable — an indirect call per method, not inlinable. (b) A trait object is a pair of pointers: one to the value, one to the vtable of the trait's method implementations for that concrete type. (c) Dynamic dispatch is required for a heterogeneous collection (e.g. `Vec<Box<dyn Draw>>`) or run-time-chosen implementations; static dispatch is preferable for a hot loop over a single known type where inlining matters. *Method-selection cue*: known type(s) at compile time and speed matters → static; heterogeneity or run-time choice → dynamic.

Feed-forward: §30.4.

30.7 (a) Monomorphization costs **binary size** (and compile time): a separate code copy per concrete type, which dynamic dispatch avoids by keeping one copy. (b) Dynamic dispatch costs an **indirect call per method invocation** (an indirect branch, not inlinable) that a monomorphized generic does not. (c) For a small method in a tight loop over a homogeneous collection, choose static dispatch — inlining and the absence of an indirect call dominate; for a large plugin interface with many implementations, choose dynamic dispatch — one code copy and run-time flexibility outweigh per-call cost, and monomorphizing over many types would bloat the binary. The deciding factor is per-call hot-path cost vs. code-size/heterogeneity. *Tempting wrong answer*: “always use dynamic dispatch to keep binaries small” — that pays indirect calls everywhere, including hot paths where static dispatch is free. **Feed-forward:** §30.2, §30.4.

30.8 Step 4: Corrected signature: `fn largest<T: PartialOrd + Copy>(list: &[T]) -> T` — the minimal bounds the body uses (`PartialOrd` for `>`, `Copy` for copying elements out). Bounding at the definition, not the call site, is what lets the compiler verify the generic *once*: the body is checked against the declared requirements, and every call is then checked to supply a satisfying type, so there is no per-call re-checking of the body (the C++ failure mode). Casting or wrapping to silence `E0277` is wrong because the error isn't a type mismatch — it's that the body needs a capability the type must actually provide; the fix is to state the bound (or implement the trait), not to cast. *Tempting wrong answer*: “add `T: Any` and `downcast`” — that throws away static checking to paper over a missing bound.

Feed-forward: §30.3.

30.9 (a) **Static** — a hot numeric loop wants inlining and no indirect call; the type is known. (b) **Dynamic** — one `Vec` of mixed shape types needs a trait object (`Box<dyn Draw>`); a single monomorphized type can't hold different concrete types. (c) **Static** — one concrete type; monomorphization gives a direct call and no code bloat. (d) **Dynamic** — plugins chosen at run time can't be known at compile time, so a trait object is required. *Method-selection cue*: heterogeneity or run-time choice forces dynamic; otherwise static wins on cost. **Feed-forward:** §30.4.

30.10 (a) **False** — a generic is monomorphized; there is no type parameter or dispatch at run time. (b) **True** — a trait object dispatches through a vtable at run time. (c) **False** — E0277 means the type doesn't implement a required trait; fix by implementing it or relaxing the bound. (d) **False** — monomorphization can make binaries *larger* (a copy per type); it trades size for speed. (e) **False** — a trait object is still fully type-checked (it can only hold a type that implements the trait), unlike C's `void*`. *Tempting wrong answer:* marking (e) true by analogy to `void*` — the analogy holds for the dispatch mechanism, not for type safety. **Feed-forward:** §30.2–§30.4.

30.11 (a) `Copy` is a trait; whether `let y = x;` copies or moves is decided by whether `T` implements it — exactly the same mechanism as `T: PartialOrd` deciding whether `>` is allowed: a capability named as a trait, required in the type, checked by the compiler. (b) `Option<T>` and `Result<T, E>` are generic enums; being generic, they carry the success value's type and force the caller to pattern-match/handle the absent or error case, making “did you check the return?” a compile-time question. (c) `T: Send` would mean “a value of `T` is safe to transfer to another thread”; requiring it at a thread-spawn boundary means the compiler refuses to move a non-thread-safe value across threads, so the aliasing-plus-mutation that causes a data race can't be set up — a compile error instead of undefined behavior. *Tempting wrong answer:* “traits are only for shared method behavior like `summarize`” — they also encode capabilities and guarantees (`Copy`, `Send`, `Sync`). **Feed-forward:** §30.5; the error-handling and concurrency chapters ahead, and Ch 24.

30.12 Toolchain-specific. A correct write-up reports: (a) the `qsort` demo prints the sorted array (e.g. `1 2 5 7 9`), and passing a comparator for a different element type still compiles because `qsort`'s parameters are `void*` — demonstrating type erasure. (b) With `rustc`, `largest<T: PartialOrd + Copy>` compiles at two types; removing the bounds yields `error[E0277]` on the comparison/copy lines (the `PartialOrd/Copy` bound is missing); without `rustc`, name those lines and bounds. (c) A `Vec<Box<dyn Draw>>` can hold two different implementors because each is a trait object dispatched via vtable; a single monomorphized generic `Vec<T>` is one concrete type and cannot hold two different types at once. *Tempting wrong answer:* “the wrong-typed `qsort` comparator is fine because it compiled” — it compiles but reinterprets bytes wrongly at run time. **Feed-forward:** §30.1, §30.4.

Chapter 31 — Solutions

31.1 `Option<T>` represents a value that may be absent — variants `Some(T)` and `None` — replacing the null pointer. `Result<T, E>` represents an operation that may fail with a reason — variants `Ok(T)` and `Err(E)`. *Tempting wrong answer:* “`None` is Rust's null pointer” — `None` is a variant the type system tracks, not a pointer value that inhabits every reference type; you cannot dereference it. **Feed-forward:** §31.2, §31.3.

31.2 On an expression `x` of type `Result<T, E>`, `x?` evaluates to the inner `T` if `x` is `Ok(T)`, and otherwise performs an early *return* of the `Err(E)` from the enclosing function (converting the error type through `From` if needed). *Tempting wrong answer:* “`?` panics on `Err`” — that is `unwrap()`; `?` returns the error to the caller instead. **Feed-forward:** §31.4.

31.3 `Result` is for *recoverable* failures the caller should handle; `panic!` is for *unrecoverable* bugs or violated invariants, and it unwinds the stack and aborts the thread. `unwrap()` is a shortcut for “if this is `Err`, treat it as a bug and `panic!`”. *Tempting wrong answer:* “`panic!` is just Rust's exception you catch upstream” — it is meant for unrecoverable conditions, not routine control flow; recoverable errors use `Result`. **Feed-forward:** §31.4.

31.4 (a) The program did not check `fopen`'s return value; C signals that failure by returning a `NULL` pointer (and setting `errno`). (b) `-Wall` did not warn because checking a returned sentinel is a convention the compiler does not enforce; the warning it emitted was the unrelated `-Wunused-result` for ignoring `fgets`'s own `warn_unused_result` return. (c) The `NULL` `fp` is dereferenced inside `fgets`, causing a segmentation fault (exit 139); dereferencing an invalid pointer is undefined behavior (Chapter 10), so the compiler may assume it never happens and no diagnostic is required. *Tempting wrong answer:* “the warning means the compiler caught the `NULL` bug” — the warning was about `fgets`'s return, not the unchecked `fp`. **Feed-forward:** §31.1, §31.3.

31.5 (a) The success value lives inside the `Ok(T)` variant, so to reach a usable value you must write code that also names the `Err` case (a `match`, a combinator, or `?`). (b) `#[must_use]` makes the compiler *warn* if you discard the whole `Result` without using it — the enforcement C's convention lacks. (c) `unwrap()` is the closest analog to skipping C's check because it reaches the success value while turning any `Err` into a `panic!`; it is at least visible and greppable, unlike C's silent fall-through. *Tempting wrong answer:* “`unwrap()` is safe because it can't segfault” — it can't corrupt memory, but it still crashes the thread on a recoverable error. **Feed-forward:** §31.3.

31.6 (a) `?` propagates an `Err` by returning it from the enclosing function, but `load` returns `()`, which has nowhere to put a propagated error — so the operation is ill-typed and reported as E0277. (b) Corrected: `fn load(path: &str) -> Result<(), io::Error> (with Ok() at the end)`. (c) Adding `.unwrap()` compiles but changes the semantics: a missing file — a recoverable error — now `panic!`s and aborts the thread instead of being handed to the caller who could recover. *Tempting wrong answer*: “E0277 means `File::open` returned the wrong type” — the operand type is fine; the enclosing function's return type is what disallows `?`. **Feed-forward**: §31.4.

31.7 (a) In C the worst case is undefined behavior at run time — the NULL dereference and segfault of §31.1, or worse, silent corruption. (b) In Rust, dropping the `Result` is caught at compile time as a `#[must_use] warning` (the value is unused, no crash), whereas reflexively `unwrap()`ing it downgrades a recoverable error to a run-time `panic!` that aborts the thread — bad, but a controlled unwind, not UB. (c) Ordered by severity: compile-time diagnostic (safest, before the program runs) < warning < controlled `panic!` < runtime UB (worst). This ordering matters because a system you cannot easily debug in production fails most cheaply when the compiler catches the omission, and most expensively when it is silent undefined behavior. *Tempting wrong answer*: “a Rust panic is as bad as a C segfault” — a panic is a defined, controlled unwind with a message; UB is undefined and may corrupt silently. **Feed-forward**: §31.1, §31.4.

31.8 Step 4: Corrected: `fn first(path: &str) -> Result<String, io::Error> { let s = std::fs::read_to_string(path)?; Ok(s) }` — the function now returns a `Result`, so the `io::Error` that `?` may propagate has somewhere to go, and the success value is wrapped in `Ok`. This is preferable to `.unwrap()` because a missing or unreadable file is a routine, recoverable condition the caller may want to handle (fall back to defaults, report to the user); `unwrap()` would `panic!` and abort instead. *Tempting wrong answer*: “return `String` and `unwrap()` so the signature stays simple” — that hides a recoverable failure behind a crash. **Feed-forward**: §31.4.

31.9 (a) **Option** — a key's presence is a plain absent/present question with no error reason (or `Result` only if the lookup itself can fail). (b) **Result** — malformed input is an expected, recoverable failure carrying a reason. (c) **panic!** is acceptable — an out-of-range index you just proved in-range would be a bug, so a panic (or plain indexing, which panics) documents the invariant; alternatively return `Result` if the proof is not airtight. (d) **Result** — a user-specified file may legitimately be missing or unreadable; the caller should recover. *Method-selection cue*: absence with no reason → `Option`; expected failure with a reason → `Result`; a violated invariant/bug → `panic!`. **Feed-forward**: §31.2–§31.4.

31.10 (a) **False** — `None` is an `Option` variant the type system tracks, not a null pointer that inhabits every reference type. (b) **False** — `?` is only valid in a function returning `Result` or `Option` (else E0277). (c) **False** — dropping a `Result` is a `#[must_use] warning`, not an error. (d) **False** — `unwrap()` on `Err` `panic!`s; `unwrap_or` is the one that returns a default. (e) **True** — `?` converts the error type through the `From` trait on the way out. *Tempting wrong answer*: marking (c) true by analogy to (b) — dropping a value is a warning, using `?` in the wrong context is the hard error. **Feed-forward**: §31.2–§31.4.

31.11 (a) Using `?` states a requirement on the enclosing function: “this function can absorb a propagated error” — i.e. its return type implements the residual/conversion machinery (built on `From`). That is a trait requirement, so its violation is the same `E0277` (“trait bound not satisfied”) as an unmet bound in Chapter 30. (b) Because `Option<T>` and `Result<T, E>` are *generic*, one `?` operator is defined over all their instantiations at once — it works uniformly on `Result<File, io::Error>`, `Result<String, ParseError>`, and every other fallible return without per-type code, exactly the reuse monomorphization gives. (c) A worker's result would travel as a `Result<Output, E>` (“the work, or why it failed”); a thread panic! is observed by the parent as an `Err` from `JoinHandle::join`, carrying the panic payload. *Tempting wrong answer*: “`?`'s error is a special language feature unrelated to traits” — it is built on the `From`/residual traits, which is why the diagnostic is a trait-bound error. **Feed-forward**: §31.5; the concurrency chapter ahead, and Chapter 30.

31.12 Toolchain-specific. A correct write-up reports: (a) `cerr.c` compiles (with only the `-Wunused-result` warning) and, on a missing file, terminates with SIGSEGV (exit status 139); adding `if (fp == NULL) { perror(...); return 1; }` converts the crash into a clean error message and non-crashing exit. (b) With `rustc`, `read_config` returning `Result<String, io::Error>` returns the error on a missing path (no crash); changing its return type to `()` produces `error[E0277]` on the `?` line. Without `rustc`, name the `?` line and the return type as the cause. (c) Replacing a `?` with `.unwrap()` changes a propagated, recoverable error into a thread-aborting panic!. *Tempting wrong answer*: “since `cerr.c` compiled, the `NULL` path is fine” — it compiles but is undefined behavior at run time. **Feed-forward**: §31.1, §31.4.

Chapter 32 — Solutions

32.1 `Send` means a value of the type is safe to *move* (transfer) to another thread; `Sync` means it is safe to *share* a reference to it across threads. The relationship: `T: Sync` exactly when `&T: Send`. *Tempting wrong answer*: “`Send` is for moving and `Sync` is for locking” — `Sync` is about shared references being sendable, not about acquiring locks. **Feed-forward**: §32.2.

32.2 A **data race** is two threads accessing the same memory concurrently, at least one writing, with no synchronization ordering the accesses. In C (and C++) it is *undefined behavior* (Chapter 24). In safe Rust it cannot occur: the type system (the `Send/Sync` bounds plus the borrow rules) rejects the unsynchronized sharing at compile time. *Tempting wrong answer*: “a data race just gives a wrong answer” — in C it is UB, so the program has no defined meaning at all, not merely a wrong number. **Feed-forward**: §32.1, §32.3.

32.3 `move` forces the closure to take *ownership* of the variables it uses rather than borrowing them, so the data lives as long as the thread. Without it, a closure that borrows a local and is handed to a thread that may outlive the borrow draws error[E0373] (“closure may outlive the current function”). *Tempting wrong answer*: “`move` copies the captured data” — it transfers ownership (a move), which for non-Copy types invalidates the original. **Feed-forward**: §32.3.

32.4 (a) It compiled because C has no rule against two threads touching shared mutable state; the data race is invisible to the compiler. (b) Each `counter++` is load-add-store; when two threads read the same value before either stores, one increment overwrites the other and is lost, so the total is below 4000000, and because the interleaving is nondeterministic the exact loss differs each run. (c) The C fix is a mutex around every access to `counter`; it is fragile because it relies on you remembering the lock at *every* access site — one omission silently reintroduces the race, whereas Rust makes the unsynchronized access unrepresentable. *Tempting wrong answer*: “the total is low because the loop ran too few times” — the loop ran fully; updates were lost to interleaving. **Feed-forward**: §32.1, §32.4.

32.5 (a) `Send` rules out moving into the new thread any value that is not safe to transfer across threads — including any setup that would give two threads unsynchronized mutable access. (b) `'static` rules out the closure holding a borrow of data owned by the parent frame, which could dangle once the parent returns — the lifetime rule of Chapter 29. (c) An `Rc` is not `Send` (non-atomic refcount), so capturing it in a `spawn` closure leaves the closure non-`Send` and the bound fails. *Tempting wrong answer*: “`'static` means the thread runs forever” — `'static` bounds the *lifetime of borrows* the closure may hold, not the thread's duration. **Feed-forward**: §32.2.

32.6 (a) E0373 reports that a closure would outlive a value it borrows — an instance of the Chapter 28–29 borrowing/lifetime rules applied to a closure escaping to a thread; E0277 reports an unsatisfied Send (or Sync) bound — an instance of the Chapter 30 trait-bound check applied to Send. (b) Fix E0373 with `move` (take ownership); fix E0277 by using the thread-safe type (`Arc` for `Rc`, an atomic for a plain counter). (c) `unsafe impl Send` merely *asserts* thread-safety the type does not have, disabling the check and reintroducing the race — an unearned promise, not a fix. *Tempting wrong answer*: “E0373 and E0277 are the same borrow error” — one is a lifetime/borrow problem, the other a trait-bound problem, with different fixes. **Feed-forward**: §32.3.

32.7 (a) The `Mutex` costs a lock acquire/release per update; as thread count rises, **lock contention** (Chapter 26) grows and can serialize the threads, throttling throughput. (b) The channel costs a send/receive and the move of the value; moving ownership down the channel means the sender can no longer touch the value, so there is nothing to race on and no lock to contend. (c) For a high-contention counter, an **atomic** (Chapter 22, e.g. a lock-free `fetch_add`) avoids the lock entirely; you would still prefer the `Mutex` when each update touches *multiple* fields that must change together atomically, which a single atomic cannot express. *Tempting wrong answer*: “channels are always faster than a mutex” — for a single hot counter an atomic or even a mutex may beat channel overhead; it depends on the access pattern. **Feed-forward**: §32.4.

32.8 Step 4: Replace `Rc` with `Arc`: `let c = Arc::new(Mutex::new(0));` and, inside the loop, `let cc = Arc::clone(&c);` before the `move` closure. This is a real fix because `Arc`'s reference count is updated with atomic operations, so it is genuinely safe to send/share across threads and therefore `Send + Sync` — the bound is now satisfied by fact, not by fiat. `unsafe impl Send` for ... would instead assert `Send` for the non-atomic `Rc`, switching off the check while the underlying refcount race remains — silencing the diagnostic without removing the bug. *Tempting wrong answer*: “wrap the `Rc` in a `Mutex` too” — the problem is the non-atomic refcount of `Rc` itself, not the interior value; you need `Arc`. **Feed-forward**: §32.3.

32.9 (a) **Prevented** — sharing an `i32` mutably across threads requires either a non-`Sync` share or aliasing a `&mut`, both rejected; you must use a `Mutex`/atomic. (b) **Not prevented** — deadlock from opposite lock-ordering is a logical bug, not a data race; the type system does not analyze lock order. (c) **Prevented** — `Rc` is not `Send`, so sending it across threads is E0277. (d) **Not prevented** — holding a guard across a slow call is safe (no race) but a contention/granularity problem the compiler does not flag. *Method-selection cue*: data races (unsynchronized shared mutation) are prevented; higher-level bugs (deadlock, contention, granularity) are not. **Feed-forward**: §32.3, §32.4.

32.10 (a) **False** — `Send/Sync` are marker (auto) traits with no methods, derived structurally. (b) **False** — Rust's `Mutex` owns the data; the only way to the value is through `lock()`, so lock and data are not separable. (c) **True** — `Arc` is `Rc` with an atomic reference count, making it `Send + Sync`. (d) **False** — it prevents data races, not deadlock. (e) **False** — sending *moves* ownership; the sender can no longer use the value. *Tempting wrong answer*: marking (b) true by analogy to C's separate `pthread_mutex_t` — Rust deliberately bundles the data inside the `Mutex`. **Feed-forward**: §32.2–§32.4.

32.11 (a) `Send/Sync` extend aliasing-XOR-mutation across threads: sharing a `&mut` between two threads would create two mutable paths to one value at once (aliased mutation), which is exactly what the borrow rule forbids and what would be a data race — so it does not type-check. (b) All three are *trait-bound* failures: an unmet generic bound (Chapter 30), a `?` in a function that cannot absorb the error (Chapter 31, an unmet residual/`From` bound), and a non-`Send` type crossing threads (here) are all “a required trait is not implemented,” which is `E0277`. (c) The C volumes made you prove by hand that memory is freed once and used while valid, references never dangle, and shared state is never raced on; Rust turns each such proof into a type-system rule the compiler checks (ownership, borrowing, lifetimes, `Send/Sync`) — same machine, same generated code, proof automated. *Tempting wrong answer*: “the reused error code is a coincidence” — it reflects that thread-safety, error propagation, and generic reuse are all expressed as trait requirements. **Feed-forward**: §32.5; Chapters 27–31.

32.12 Toolchain-specific. A correct write-up reports: (a) at `-00` the totals are below `4000000` and differ run to run (e.g. ~ 2.1 – 2.8 million), since each `++` is a real load/add/store with a wide race window; at `-02` the optimizer may collapse the loop into a single register update, shrinking the window so the race often does not manifest — a reminder that a data race is UB and its observable effect depends on optimization and scheduling. (b) A `pthread_mutex_t` around `counter++` makes the total exactly `4000000`, deterministically. (c) With `rustc`, the `Arc<Mutex<i32>>` version is deterministic; the `Rc` version fails with `E0277` on the `spawn` capturing the non-`Send Rc`. Without `rustc`, name that line. *Tempting wrong answer*: “`-02` printing `4000000` proves the race is gone” — it proves only that the window closed this time; the program is still UB. **Feed-forward**: §32.1, §32.4.

Chapter 33 — Solutions

33.1 “Measure before you optimize”: locate the real cost by measurement before spending effort, because bottlenecks are rarely where intuition says. A **profiler** tells you *which* code spends the runtime (find the target); a **benchmark** compares two versions of a known piece of code honestly (evaluate a fix). *Tempting wrong answer*: “profiling and benchmarking are the same thing” — profiling finds the hot spot across a whole program; benchmarking times a specific candidate. **Feed-forward**: §33.1, §33.4.

33.2 Because the loop's result is never used, so at `-O2` the optimizer proves the loop has no observable effect and deletes it entirely (Chapter 10) — nothing runs, so the timer reads `0.0000 s`. *Tempting wrong answer*: “the CPU is just that fast” — a billion iterations are not free; they did not happen. **Feed-forward**: §33.2.

33.3 A single run is disturbed by the scheduler, cache/TLB state, frequency scaling, and other tenants, so its time is noisy; running many trials and reporting the minimum (or median) gives the run least disturbed by outside interference — closest to the code's intrinsic cost — and you report the spread too. *Tempting wrong answer*: “average the runs” — the mean is dragged up by one-directional outliers (interference only slows a run), so the minimum is a more principled floor. **Feed-forward**: §33.3.

33.4 (a) At most 6.4% — perfectly optimizing a 6.4%-of-runtime part caps its whole-program payoff at 6.4% (Amdahl). (b) $0.30 \times 93.6\% = 28.1\%$ of total runtime. (c) It violates §33.1 because effort was spent on the part that mathematically cannot matter ($\leq 6.4\%$) while the 93.6% dominant cost went untouched — and the target was chosen by how the code “looks” rather than by a profile. *Tempting wrong answer*: “optimizing parse is fine because every bit helps” — Amdahl bounds “every bit” at 6.4%, a poor return versus `compute`. **Feed-forward**: §33.1, §33.4.

33.5 (a) When the loop's result is never observed, the compiler may delete it — the Chapter 10 principle that computations with no observable effect can be removed. (b) Make the result observable: print it as a checksum, or write it to a `volatile` (or a compiler-opaque sink). (c) Because the debug (`-O0`) build's costs are not the release build's — you must measure the code as it will actually be compiled and shipped, or the number describes a program you are not deploying. *Tempting wrong answer*: “add a `printf` inside the loop” — that dominates the timing with I/O; consume the result once, after the loop. **Feed-forward**: §33.2.

33.6 (a) An instrumenting profiler inserts counting/timing code at function entry and exit (exact call counts); a sampling profiler interrupts the program many times per second and records the program counter, building a statistical map. (b) Instrumentation perturbs timing and can badly distort tiny, hot functions whose own cost is dwarfed by the added instrumentation; sampling has negligible overhead but is statistical (needs enough samples for accuracy). (c) For a first look at a large program without recompiling, reach for the **sampling** profiler — low overhead, no rebuild, and it points at the dominant cost. *Method-selection cue*: exact counts on specific calls → instrumenting; low-overhead whole-program map → sampling. **Feed-forward**: §33.4.

33.7 (a) Three sources: the OS scheduler preempting the process (as in trial 4); cache/TLB and branch-predictor state differing between runs; and CPU frequency scaling (turbo/thermal) plus contention from other tenants. (b) Interference only ever makes a run *slower* — there is no mechanism that makes identical work finish faster than its intrinsic cost — so the minimum is the closest estimate of that intrinsic floor, whereas the mean is pulled upward by outliers. (c) The median is better when you care about *typical* behavior under realistic interference (e.g. tail-aware service work) rather than the best case; alongside either, always report the spread (range or a percentile) so the reader sees the variance. *Tempting wrong answer*: “discard the outlier and average the rest” — arbitrary trimming hides real variance; report the minimum/median and the spread instead. **Feed-forward**: §33.3.

33.8 Step 4: Accumulate the hash results and consume the accumulator once after the loop, and vary the input so the call is not loop-invariant, e.g. `unsigned long acc = 0; for (i=0; i<N; i++) acc ^= hash(key[i]); // vary input then printf("%lu\n", acc);` (or write `acc` to a volatile). This makes the work observable so the optimizer cannot delete it, and the varying `key[i]` defeats constant-folding/hoisting. Before trusting the number: time with `clock_gettime(CLOCK_MONOTONIC, ...)` at the shipping optimization level, warm up once, run several trials, and report the minimum and spread. *Tempting wrong answer*: “mark hash volatile” — consume the *result*; a volatile qualifier on the function is not the fix. **Feed-forward**: §33.2, §33.3.

33.9 (a) **Profile** — you do not yet know which part is slow; a profile finds it. (b) **Benchmark** — the hot function is confirmed; benchmark the two candidates head to head (honestly). (c) **Neither yet** — first profile to see if that 2% module even matters; per Amdahl a rewrite of 2% buys $\leq 2\%$, so likely do not optimize it at all. (d) **Benchmark** (or re-profile) — measure total runtime before and after to confirm the change actually moved the whole program, not just the microbenchmark. *Method-selection cue*: unknown target → profile; compare known candidates or confirm a whole-program win → benchmark; small share → do not optimize yet. **Feed-forward**: §33.1, §33.4.

33.10 (a) **False** — `0.0000 s` almost always means the loop was deleted (dead-code elimination), so nothing was measured. (b) **False** — `-00` costs differ from the shipped `-02` build; benchmark what you ship. (c) **False** — a sampling profiler needs no recompilation; it samples the running program. (d) **True** — Amdahl caps a 4%-of-runtime part's whole-program payoff at 4%. (e) **False** — even on an idle machine, cache state, frequency scaling, and cold-start effects vary a single run; run several. *Tempting wrong answer:* marking (e) true because “idle means no interference” — idle removes some noise, not cache/frequency/warm-up variation. **Feed-forward:** §33.2–§33.4.

33.11 (a) A loop can be slow at a low instruction count because of **cache misses** (Chapter 3) stalling on memory, **branch mispredictions** flushing the pipeline, and **memory-bandwidth** limits or poor data layout (Chapters 4–5) — the instruction is cheap; the data it touches is not. (b) To test a cache-friendlier layout, hold everything else constant (same inputs, same optimization level, same machine, same measurement harness), change only the layout, make the result observable so the loop runs, and run many trials reporting the minimum and spread for both layouts — then compare. (c) It is the scientific method: the profile is an observation, “the layout causes cache misses” is a hypothesis, changing one thing is a controlled experiment, and re-measuring tests the prediction — changing several things at once, or not re-measuring, forfeits the ability to attribute the result. *Tempting wrong answer:* “the profile already tells you why it's slow” — it tells you *where*, not *why*; the cause is a separate hypothesis to test. **Feed-forward:** §33.3, §33.4; the cache and data-layout chapters ahead.

33.12 Platform-specific. A correct write-up reports: (a) the unused-result loop times near `0.0000 s` (deleted), and after printing the variable it times to its real cost (a fraction of a second to seconds depending on *N* and hardware), with the difference attributed to dead-code elimination. (b) Five `CLOCK_MONOTONIC` trials of the honest version, with a minimum below the others and a spread (max – min) of some percent — e.g. a ~20% span is unremarkable. (c) With `perf/gprof`, the profiler agrees that the heavier function dominates; without them, hand timing shows the split (e.g. ~94% vs ~6%) and the write-up names the sampling profiler it would use and that it adds per-line/whole-program attribution without recompilation. *Tempting wrong answer:* “report the first run's time” — the first run pays cold-cache/warm-up costs; report the minimum of several. **Feed-forward:** §33.2–§33.4.

Chapter 34 — Solutions

34.1 Spatial locality: if you access one memory address, you are likely to access nearby addresses soon (so contiguous, small-stride access is cheap). **Temporal locality:** if you access an address, you are likely to access the *same* address again soon (so recently-used data should still be cached). Good spatial locality: `for (i) s += a[i];` (stride 1). Poor spatial locality: `for (i) s += a[i*1024];` (each access a new line). *Tempting wrong answer:* swapping the two definitions — “spatial” is about neighbors in space (address), “temporal” about revisiting over time. **Feed-forward:** §34.2, §34.4.

34.2 One fetch, in the best case: the first `double` pulls in its whole 64-byte line, and the next 7 `doubles` ($7 \times 8 = 56$ bytes) live in that same line, so they hit cache with no new fetch — 8 `doubles` fill exactly one 64-byte line. *Tempting wrong answer:* “8 fetches, one per `double`” — the hardware moves a whole line, not a single element. **Feed-forward:** §34.2.

34.3 Because if a DRAM access costs $\sim 100\times$ an L1 hit, a loop that misses cache is stalled waiting on memory for most of its time, so its speed is set by *how it touches memory*, not by how many instructions it runs — two loops with identical instruction counts can differ by an order of magnitude in time. *Tempting wrong answer:* “instruction count still decides it” — the CPU idles through the miss regardless of how few instructions remain. **Feed-forward:** §34.1, §34.3.

34.4 (a) Only the *order* of memory access differs: row-major strides by one element (each 64-byte line serves 8 `double` accesses), column-major strides by a full row (each access lands in a fresh line and, since the matrix is 512 MiB, that line is evicted long before its 8 elements are revisited). (b) 1 of 8 — the column walk uses one `double` per fetched line and discards the other 7, and does not return to them before eviction. (c) Swap the loop nesting so the inner loop strides by one, or transpose the matrix so the column access becomes contiguous. *Tempting wrong answer:* “the column version does more arithmetic” — it does exactly the same 67M additions; only locality differs. **Feed-forward:** §34.3.

34.5 (a) At stride 1 each 64-byte line serves 16 `int` accesses, amortizing the fetch 16 ways; as the stride grows the accesses served per line fall (8, 4, 2), so each access carries more of the line-fetch cost and the per-access time climbs. (b) At a 64-byte stride each access lands in a distinct line and shares none with its neighbors — spatial locality is fully defeated, the first stride at which no two accesses reuse a line. (c) Larger strides also defeat the hardware prefetcher (it can no longer predict the next line) and stress the TLB (Chapter 4), each access touching a new page more often. *Tempting wrong answer:* “the array got bigger” — the array and the access count are fixed; only the stride changed. **Feed-forward:** §34.2.

34.6 (a) Separately `malloc'd` nodes are scattered across the heap, so `next` jumps to an unpredictable, distant address — each node likely a fresh cache line (and page), with no spatial locality between consecutive nodes. (b) Store the nodes in a contiguous array (or a pool/arena allocator) so traversal strides through adjacent memory, restoring spatial locality and letting the prefetcher work. (c) Both turn a logically simple traversal into one that lands in a new cache line on nearly every step — the pointer chase and the column walk are the same “stride across lines” pathology. *Tempting wrong answer:* “small nodes mean good locality” — small but scattered is still one line per node. **Feed-forward:** §34.2, §34.3; the layout chapter ahead.

34.7 (a) While $S < C$ the array fits in cache, so after the first pass later passes mostly hit and the miss rate is low; as S grows past C , each pass evicts what the next pass needs, so the miss rate jumps toward one-miss-per-line. (b) Blocking processes a chunk of size $\leq C$ completely — doing all the work that touches it — before moving on, so each loaded line is reused many times while still cached, keeping the working set within a cache level. (c) The arithmetic operation count is unchanged, but blocking removes the memory stalls that dominated the time, so the loop stops waiting on DRAM and becomes limited by the arithmetic it was always doing — memory-bound to compute-bound with no fewer operations. *Tempting wrong answer:* “blocking does fewer operations” — it does the same operations in a cache-friendlier order. **Feed-forward:** §34.4.

34.8 Step 4: Ship `for (i) for (j) sum += A[i][j]` — the inner loop strides by one element along the contiguous (row-major) dimension, so each 64-byte line serves 16 `float` accesses and spatial locality is perfect, while the other nesting strides a full 16 KiB row per access and misses almost every time. To measure: compile both at the shipping optimization level (`-O2`), time each as the minimum of several `clock_gettime(CLOCK_MONOTONIC)` trials, hold everything else constant (same matrix, same machine, same harness), make the sum observable so neither loop is deleted (Chapter 33), warm up once, and report the minimum and spread for each; the row-order loop should be several times faster. *Tempting wrong answer:* “time each once and compare” — a single run is noisy; run several and report the minimum. **Feed-forward:** §34.3; Chapter 33.

34.9 (a) **Spatial** — column order strides across lines, wasting most of each fetched line. (b) **Temporal** — the data is reused every scan but a 100 MiB set overflows the 24 MiB L3, so it is evicted before reuse. (c) **Neither (both good)** — a 1 KiB array fits in L1 and is reused, so spatial and temporal locality are both excellent. (d) **Spatial** — scattered nodes give one line per node with no neighbor reuse. *Tempting wrong answer:* calling (b) “spatial” — the scan is contiguous (good spatial); the problem is the set does not stay cached (temporal). **Feed-forward:** §34.2, §34.4.

34.10 (a) **False** — reading one byte fetches its whole 64-byte cache line. (b) **False** — identical instruction counts can differ by an order of magnitude if locality differs (the §34.3 19.5×). (c) **False** — the compiler often may *not* reorder loops (it can change results or be illegal), so the traversal order is your responsibility. (d) **True** — a working set that fits in L2 is hit by repeated passes. (e) **False** — blocking reorders access to improve reuse; it does not reduce the arithmetic operation count. *Tempting wrong answer:* marking (c) true because “the optimizer handles everything” — loop interchange is limited and not guaranteed. **Feed-forward:** §34.2–§34.4.

34.11 (a) Hypothesis: the loop is memory-bound because of poor locality (striding across lines or a working set larger than cache); test it by improving locality (transpose/reorder, or shrink the working set) and re-measuring, or with a cache-miss counter where available. (b) Wall-clock experiment: run the loop over inputs of growing size and watch for a sharp per-element slowdown as the working set crosses a cache size (a memory-bound signature), or compare the loop against a locality-fixed version; if the fixed version is much faster, it was memory-bound, if not, compute-bound. (c) Transposing/reordering is a one-line, low-risk change that attacks the actual cause (locality) and is easy to re-measure, whereas rewriting in assembly is high-effort, high-risk, and does nothing about a memory stall the arithmetic was never causing. *Tempting wrong answer:* “assembly is always faster” — assembly cannot beat the memory ceiling that poor locality created; fix locality first. **Feed-forward:** §34.4; Chapter 33.

34.12 Platform-specific. A correct write-up reports: (a) the machine's real L1/L2/L3 sizes and line size from `/sys` (e.g. 64-byte lines). (b) The row-major sum timed several times and the column-major sum timed several times, both as minima, with a ratio of a few× to >10× depending on hardware, explained by the column walk using one element per fetched line. (c) A stride sweep whose per-access cost has a knee at the line size (contiguous cheap, one-line-stride expensive). *Tempting wrong answer:* “report the first, cold run” — the first run pays cold-cache costs; report the minimum of several (Chapter 33). **Feed-forward:** §34.2–§34.3.

Chapter 35 — Solutions

35.1 AoS (array-of-structs): one array whose elements each hold all a record's fields together — `struct P{float x,y,vx,vy;} p[N];`. **SoA (struct-of-arrays):** a separate contiguous array per field — `float x[N],y[N],vx[N],vy[N];`. *Tempting wrong answer:* “SoA is just an array of structs each holding arrays” — no; each *field* gets its own flat array, so one field's values are contiguous. **Feed-forward:** §35.2.

35.2 Useful bytes per fetched line = of every 64-byte line the hardware pulls in, how many bytes the loop actually reads. Summing one field of a 64-byte AoS struct reads 8 bytes (one double) of each fetched line and ignores the other 56, so 7 of every 8 bytes of memory traffic are wasted. *Tempting wrong answer:* “it uses the whole line because it fetched it” — fetching is not using; the other fields are dragged in and discarded. **Feed-forward:** §35.2, §35.3.

35.3 False sharing: two threads write two *different* variables that happen to sit in the *same* cache line, so the coherence protocol bounces that line between cores on every write even though no data is truly shared. Fix: **padding** (or alignment) so each thread's hot data sits on its own cache line. *Tempting wrong answer:* “add a lock” — a lock addresses real contention; false sharing has none, and a lock would make it worse. **Feed-forward:** §35.4.

35.4 (a) AoS uses 8 of 64 bytes per line; SoA uses 64 of 64 — an 8× better useful-bytes ratio, and since the loop is limited by memory traffic, ~8× less wasted traffic predicts a large speedup. (b) The measured 4.08× is below 8× because the hardware prefetcher partly hides AoS's waste (it streams ahead), and other fixed costs (the add, loop overhead) do not shrink — so the memory saving does not translate one-for-one. (c) A loop that reads *all* fields of one element before moving on: then AoS puts the whole record in one fetched line and uses all of it, while SoA would touch several arrays (several lines) per element. *Tempting wrong answer:* “SoA must give exactly 8×” — byte ratio is an upper bound, not a guarantee. **Feed-forward:** §35.2.

35.5 (a) Split into hot `struct{float x,y;} hot[N];` and cold `struct{char name[..]; Inv inventory; Ach achievements;} cold[N];`, indexed in parallel. (b) The per-frame loop now streams the dense `hot[]` array — every fetched line is full of `x,y` — instead of striding a fat struct and wasting the line on cold fields, even though no field was deleted. (c) Cost: touching a player's full state now means two arrays (worse locality for the rare all-fields access), and the code is more complex and easy to desynchronize; not worth it when the hot loop reads most fields anyway, or when the struct is already small. *Tempting wrong answer:* “always hot/cold split” — it hurts loops that need all fields of one element. **Feed-forward:** §35.3.

35.6 (a) Alignment: `double b` must sit on an 8-byte boundary, so the compiler pads after `char a` (7 bytes) and after `char c` (7 bytes) to keep the struct's size a multiple of its alignment — about 24 bytes, not 10. (b) Reorder largest-to-smallest: `{ double b; char a; char c; }` — `b` takes 8, `a` and `c` pack into the next 2 bytes, then 6 bytes of tail padding to a multiple of 8: 16 bytes. (c) At 16 vs 24 bytes, 50% more elements fit per 64-byte line, so a sweep over a million of them fetches fewer lines and runs faster. *Tempting wrong answer*: “the struct is 10 bytes, the size of its fields” — alignment padding makes it larger. **Feed-forward**: §35.3.

35.7 (a) From the cache-coherence protocol: the four counters share one line, so each thread's write invalidates the line in the other three cores' caches, forcing the line to bounce (ping-pong) between cores — pure coherence traffic, not data sharing. (b) Padding puts each counter on its own line, so a thread's writes never invalidate another's line; it trades ~56 wasted bytes per counter (harmless here) for eliminating the bounce, an 8× win. (c) It would shrink to nothing: with one thread there is no other core to bounce the line to, so packed and padded run the same — false sharing is a purely concurrent cost. *Tempting wrong answer*: “the writes themselves are slow” — the increment is trivial; the cost is the line moving between cores. **Feed-forward**: §35.4.

35.8 Step 4: Choose the **hot/cold split** (a packed `{cx,cy,cz,r}` array plus a parallel cold array): the intersection loop reads exactly those four fields together per sphere, so a 32-byte hot record gives it two spheres per 64-byte line with everything it needs, while full SoA would make it touch four separate arrays (four lines) per sphere — worse for a loop that uses all four hot fields at once. (Full SoA would win only if the loop swept a single field.) Measure it the Chapter 33 way: compile original and split at the shipping level (-O2), same scene, same machine, make the result observable, warm up, and report the minimum and spread of several `clock_gettime(CLOCK_MONOTONIC)` trials for each; the split should be several times faster. *Tempting wrong answer*: “full SoA is always best” — here the loop reads four fields together, so a packed hot record beats four separate arrays. **Feed-forward**: §35.2, §35.3; Chapter 33.

35.9 (a) **SoA** — one field swept across many records is pure spatial locality when that field is its own array. (b) **AoS** — processing one record's full state at a time wants the whole record in one line. (c) **Hot/cold split** — 3 hot fields streamed every frame go in a dense array, the 30 cold ones elsewhere. (d) **Pad to a cache line** (a layout change, not AoS/SoA) — per-thread counters must not share a line, or false sharing bites. *Tempting wrong answer*: calling (d) “SoA” — the issue is cross-thread line sharing, fixed by padding. **Feed-forward**: §35.2–§35.4.

35.10 (a) **False** — SoA wins for one-field sweeps but AoS wins when a loop uses all fields of one element. (b) **True** — reordering fields to reduce alignment padding shrinks the struct without removing a field. (c) **False** — false sharing is two threads writing *different* variables that share a cache line, not the same variable. (d) **False** — padding wastes bytes single-threaded but prevents false sharing across threads, an 8× win in §35.4. (e) **False** — same operation count, different layout, can change runtime by several× (the AoS/SoA 4×). *Tempting wrong answer*: marking (e) true from RAM-model intuition — locality, not operation count, sets the time. **Feed-forward**: §35.2–§35.4.

35.11 (a) Chapter 34 (reorder accesses over a fixed layout) fits when you cannot change the data structure — e.g. iterating a given matrix in the right order; this chapter (change the layout) fits when you own the types — e.g. AoS→SoA for a hot field sweep; use both when you both pick the layout and choose the loop order over it (SoA arrays, swept contiguously). (b) A vectorizer loads 8 contiguous elements into one register; SoA makes a field's values contiguous, so one wide load fills a vector, while AoS scatters them across struct-sized gaps and cannot. (c) Converting everything before profiling optimizes code that may be nowhere near the dominant cost (Chapter 33/Amdahl), risks pessimizing loops that wanted AoS, and spends effort with no measured payoff — profile first, convert the hot path, re-measure. *Tempting wrong answer:* “SoA everywhere is strictly better” — it hurts all-fields-per-element loops and ignores Amdahl. **Feed-forward:** §35.2; Chapters 33, 36.

35.12 Platform-specific. A correct write-up reports: (a) an AoS field sum and an SoA field sum timed as minima of several trials, with a ratio of a few× explained by the useful-bytes-per-line difference (struct size vs field size). (b) the ratio growing as the struct grows (16→128 bytes), because a bigger struct wastes more of each line in AoS while SoA is unchanged. (c) if threads are available, a packed-counters vs padded-counters run with a several× slowdown from false sharing, vanishing at one thread. *Tempting wrong answer:* “time each layout once” — single runs are noisy; report minima of several (Chapter 33). **Feed-forward:** §35.2, §35.4.

Chapter 36 — Solutions

36.1 SIMD (Single Instruction, Multiple Data): one instruction operates on a whole vector of elements at once, on a single core. A 256-bit AVX register holds 8 floats (32 bits each) or 4 doubles (64 bits each), and one AVX arithmetic instruction processes all of them — 8 or 4 — in one step. *Tempting wrong answer:* “SIMD means multiple cores” — it is width on one core, not more cores. **Feed-forward:** §36.1.

36.2 Contiguity (the elements adjacent in memory so one load fills a vector register) and **independence** (no lane depends on another's result within the vector, so no loop-carried dependency). Contiguity is supplied by Chapter 35's struct-of-arrays layout. *Tempting wrong answer:* “the loop just has to be short” — length is irrelevant; layout and independence are what matter. **Feed-forward:** §36.1, §36.3.

36.3 (a) `-O3` (gcc does not vectorize at `-O2` in this version); more generally `-ftree-vectorize`. (b) `-march=native`, letting it use the widest vectors the build CPU supports (AVX2 here). (c) `-fopt-info-vec-optimized` (or `-fopt-info-vec-missed` to see why a loop did *not* vectorize). *Tempting wrong answer:* “`-O2` is enough” — it is not, in this gcc; you need `-O3`. **Feed-forward:** §36.2.

36.4 (a) Because vectorization depends on the build: `-O2` emitted scalar code, `-O3` used 128-bit SSE2 (4-wide), `-O3 -march=native` used 256-bit AVX2 (8-wide) — same source, three code generations. (b) Vectorizing also let the compiler use fused multiply-add (one instruction for `a*b+c`) and unroll the loop, cutting per-iteration overhead, so the gain is lane count times these extra wins. (c) Confirming with `-fopt-info-vec` gives evidence the loop vectorized (and how wide), whereas assuming it did risks shipping a “vectorized” loop the compiler actually left scalar. *Tempting wrong answer:* “ $13\times$ proves a bug in the model” — no; FMA and unrolling legitimately push it past raw lane count. **Feed-forward:** §36.2.

36.5 (a) Vectorizing a sum groups the additions differently (e.g. 8 partial sums combined at the end), and floating-point addition is not associative, so the rounded result can differ from the strict left-to-right order the source specifies; the standard forbids the compiler from changing your result, so it declines. (b) `-ffast-math` (or `-funsafe-math-optimizations`) permits reassociation; it also allows other unsafe assumptions (no NaNs/infinities, reordering that changes rounding), which can break code that relies on strict IEEE behavior. (c) The dot product is **memory-bandwidth-bound** — it streams arrays far larger than cache, so it waits on DRAM, not the adder; that is a bandwidth ceiling, unrelated to the reduction/associativity issue. *Tempting wrong answer:* “the small gain is the reduction still being scalar” — once `-ffast-math` vectorizes it, bandwidth still caps it. **Feed-forward:** §36.3, §36.4.

36.6 (a) **Yes** — elementwise, contiguous, independent lanes. (b) **No** — loop-carried dependency: `a[i]` needs `a[i-1]`, so lanes cannot be computed in parallel. (c) **No (or poorly)** — a heavy data-dependent branch per element breaks the uniform lane operation (though light branches can sometimes be masked/predicated). (d) **No** at plain `-O3` — a floating-point reduction needs `-ffast-math` to reassociate. *Tempting wrong answer:* marking (b) yes because “it’s a simple loop” — simplicity is irrelevant; the dependency is fatal. **Feed-forward:** §36.3.

36.7 (a) Arithmetic intensity = floating-point operations performed per byte moved to/from memory. (b) Dot product: ~ 2 flops (a multiply and an add) per two 4-byte loads = 8 bytes, so ~ 0.25 flop/byte — very low, so once the arrays exceed cache it is limited by how fast DRAM delivers bytes: bandwidth-bound. (c) The §36.2 kernel: ~ 11 flops per 4-byte element already resident in L1, so ~ 2.75 flops/byte and almost no DRAM traffic — the core’s arithmetic throughput binds: compute-bound. (d) SIMD helps loops that are compute-bound (high intensity); for bandwidth-bound (low intensity) loops the fix is locality/layout, not more lanes. *Tempting wrong answer:* “more flops per element always means SIMD helps” — it also depends on whether the bytes are in cache (intensity counts bytes moved). **Feed-forward:** §36.4.

36.8 Step 4: $1.2\times$ is expected, not a bug: the loop is bandwidth-bound (tiny arithmetic intensity, 100 MiB arrays streamed from DRAM), so the memory bus caps it and widening the add from 1 lane to 8 cannot help — the adder was never the bottleneck. Confirm the diagnosis by timing the plain `-O3 -march=native` auto-vectorized version (checking `-fopt-info-vec` that it vectorized): it should match $\sim 1.2\times$, showing the intrinsics are fine and memory is the ceiling. The real speedup must come from moving fewer bytes — better locality and layout (Chapters 34–35), e.g. fusing this with adjacent passes so the data is touched once, or shrinking the element type — not from more lanes. *Tempting wrong answer:* “keep debugging the intrinsics” — they are correct; the model, not the code, was wrong. **Feed-forward:** §36.4.

36.9 (a) **Large** — compute-bound, cache-resident, independent lanes: SIMD’s ideal case. (b) **None (to speak of)** — a 1 GiB copy is pure bandwidth, so wide loads/stores are already bandwidth-limited; lanes do not add throughput. (c) **Large** — cache-resident and arithmetic-heavy, so the vector unit binds, not memory (a reduction needs `-ffast-math` to vectorize). (d) **None** — a loop-carried dependency is inherently serial; it cannot vectorize as written. *Tempting wrong answer:* calling (b) large “because wider loads move more” — the bus, not the load width, is already the cap. **Feed-forward:** §36.3, §36.4.

36.10 (a) **False** — 8 lanes bound arithmetic throughput at up to $8\times$, but bandwidth-bound loops gain far less and compute-bound ones can exceed $8\times$ via FMA/unrolling. (b) **False** — gcc will not vectorize a floating-point reduction at `-O3` without `-ffast-math` (associativity). (c) **True** — `-march=native` targets the build machine’s widest supported vectors. (d) **False** — SIMD helps compute-bound loops, not bandwidth-bound ones. (e) **False** — an AoS field is strided, so it cannot be loaded into a vector in one go; packed SoA can. *Tempting wrong answer:* marking (a) true from the lane-count intuition — the ceiling, not the lane count, sets the gain. **Feed-forward:** §36.1–§36.4.

36.11 (a) SIMD needs contiguous data (one wide load fills a register) and data that reaches the core fast (or the vector unit starves), which are exactly the layout of Chapter 35 and the locality of Chapter 34 — without them the vectorizer either cannot fire or gains nothing. (b) Ask in order: what is this kernel bound by — arithmetic, cache misses, or bandwidth (measure scalar vs vectorized, vary input size to find cache cliffs)? If cache/bandwidth-bound, fix locality (Chapter 34) and layout (Chapter 35) first; if compute-bound, vectorize (Chapter 36) and confirm with `-fopt-info-vec`; re-measure after each change (Chapter 33). (c) Because the true limit is usually a specific resource (cache, bandwidth, vector width), and raising a ceiling you are not under does nothing — “fewer instructions” ignores that a memory-bound loop is idle waiting on bytes regardless of instruction count. *Tempting wrong answer*: “always vectorize the hot loop” — pointless if it is bandwidth-bound; find the binding ceiling first. **Feed-forward**: Chapters 33–36.

36.12 Platform-specific. A correct write-up reports: (a) an L1-resident compute kernel at scalar / `-O3` / `-O3 -march=native`, with throughputs rising and a speedup at or above the lane count, and the `-fopt-info-vec-optimized` line naming the vector width. (b) a cache-overflowing dot product that barely vectorizes without `-ffast-math` and gains little even with it, explained by low arithmetic intensity (bandwidth-bound). (c) the toolchain version, CPU, and widest vector width, noting these are wall-clock throughputs (no hardware counters) if `perf` is unavailable. *Tempting wrong answer*: “report the first run” — cold-start noise; report the minimum of several trials (Chapter 33). **Feed-forward**: §36.2–§36.4.